

Computer Architecture I Mid-Term I

Chinese Name: _____

Pinyin Name: _____

Student ID: _____

E-Mail ... @shanghaitech.edu.cn: _____

Question:	1	2	3	4	5	6	7	8	9	10	11	12	Total
Points:	1	3	11	11	10	12	7	16	5	10	7	7	100
Score:													

- This test contains 21 numbered pages, including the cover page, printed on both sides of the sheet.
- We will use Gradescope for grading, so only answers filled in at the obvious places will be used.
- Use the provided blank paper for calculations and then copy your answer here.
- Please turn **off** all cell phones, smartwatches, and other mobile devices. Remove all hats and headphones. Put everything in your backpack. Place your backpacks, laptops and jackets out of reach.
- Unless told otherwise always assume a 32bit machine.
- The total estimated time is 120 minutes.
- You have 120 minutes to complete this exam. The exam is closed book; no computers, phones, or calculators are allowed. You may use two A4 pages (front and back) of handwritten notes in addition to the provided green sheet.
- There may be partial credit for incomplete answers; write as much of the solution as you can. We will deduct points if your solution is far more complicated than necessary. When we provide a blank, please fit your answer within the space provided.
- Do **NOT** start reading the questions/ open the exam until we tell you so!

1 1. **First Task (worth one point): Fill in you name**

Fill in your name and email on the front page and your ShanghaiTech email on top of every page (without @shanghaitech.edu.cn) (so write your email in total 21 times).

2. Various Questions**3** (a) Name the 6 Great Ideas in Computer Architecture as taught in the lectures.

Solution:

1. Abstraction (Layers of Representation/Interpretation)
2. Moore's Law (Designing through trends)
3. Principle of Locality (Memory Hierarchy)
4. Parallelism
5. Performance Measurement and Improvement
6. Dependability via Redundancy

3. Number Representation**3** (a) Given the number `0x811F00FA`. It can be interpreted as:

a binary number: _____

four unsigned bytes: _____

four two's complement bytes: _____

Solution:

1000 0001 0001 1111 0000 0000 1111 1010;

129, 31, 0, 250; -127, 31, 0, -6;

4

- (b) A **quarter** is a single byte split into the following fields (1 sign, 3 exponent, 4 mantissa): **SEEE MMMM**. It has all the properties of **IEEE 754** (including denormal numbers, NaNs and $\pm\infty$) just with different ranges, precision and representations. For a **quarter**, the bias of the exponent is 3, and the implicit exponent for denormal numbers are -2.

What is the largest number smaller than ∞ ?

In binary _____

In decimal _____

Which negative denormal number is closest to 0?

In binary _____

In decimal _____

Solution: 01101111; 15.5; 10000001; $-\frac{1}{64}$.

4

- (c) What is the value of **q1**, **q2**, **c**, **d**?

Hint Rounding mode: round toward even/0.

```
1 quarter q1, q2, q3, c, d;
2 q1 = -0.25;
3 q2 = -4.0;
4 q3 = 0.125;
5 c = q1 + (q2 + q3);
6 d = (q1 + q2) + q3;
```

q1 in binary _____

q2 in binary _____

c in decimal _____

d in decimal _____

Solution: 10010000; 11010000; -4.25; -4.50.

4. C Basics

5

(a) Memory of C

```
1 #include <stdlib.h>
2
3 int main() {
4     static int p = 5;
5
6     char *str = _____;
7     /* some other codes, and you can skip it. */
8     return 0;
9 }
```

1. You need to allocate a string `str` containing `p` characters. Write the code above (please use `malloc`).

Solution: `char *str = malloc(sizeof(char) * (p + 1));`

2. Fill in the correct memory section based on what the given C expressions evaluate to.

`&p` _____

`&str` _____

`str` _____

Solution: static, stack, heap.

3

(b) Catch bugs!

1. When you want to debug with GDB, what flag you will put in your compilation?

Solution: -g

2. Write down some essential commands in GDB. Example: Start your program:
run/r

Set break point: _____

Show next line (stepping into function calls): _____

Solution: break/b, step/s

3

(c) C programming: Reverse singly linked list. For example, convert $1 \rightarrow 2 \rightarrow 3 \rightarrow \text{NULL}$ to $3 \rightarrow 2 \rightarrow 1 \rightarrow \text{NULL}$. (You may not need all of the lines)

```

1 #include <stdio.h>
2 #include <stdlib.h>
3
4 /* Definition for singly-linked list. */
5 struct ListNode {
6     int val;
7     struct ListNode *next;
8 };
9
10 /* Given the head of a singly linked list, reverse
11    the list, and return the head of reversed list.*/
12 struct ListNode *reverse_list(struct ListNode *head) {
13     struct ListNode *prev = NULL;
14     struct ListNode *curr = head;
15     struct ListNode *next = head;
16     while (curr) {
17         next = curr->next;
18
19         _____
20
21         _____
22
23         _____
24
25         _____
26
27         _____
28
29         _____
30     }
31     return prev;
32 }

```

Solution:

```
curr->next = prev;  
prev = curr;  
curr = next;
```

5. Byte-Swap Operation

Assuming we are in a **32bit, little endian** system. *Little Dragon* receives a 4-byte integer num , he wants to swap the value of num 's i^{th} byte and j^{th} byte ($i, j \in \{0, 1, 2, 3\}$, $i \neq j$) to get a new number!

3

- (a) **Idea I:** *Little Dragon* wants to directly retrieve the i^{th} and j^{th} byte of num , then swap them.

First of all, define a MACRO to get the i^{th} byte of num . Read the following C code, then help *Little Dragon* to fill in the blank lines (Line 4 and 10) so the output should be **0x34**. When defining the MACRO, use **&, |, ^, ~, >>, <<** operators **only**. Remember to write a meaningful MACRO such that *Little Dragon* can reuse it again (directly return **0x34** is not allowed)!

```
1 #include <stdio.h>
2 #include <stdint.h>
3
4 #define GET_BYTE(num, ind) _____
5
6 int main(){
7     int number, index;
8     int8_t byte;
9     number = 0x12345678;
10
11     index = _____; /* index is one of {0, 1, 2, 3} */
12     byte = GET_BYTE(number, index);
13     printf("%#x\n", byte); /* should print 0x34 */
14     return 0;
15 }
```

Write your answer above.

Solution: `((num) >> ((ind) << 3)) & 0xFF); 2.`

Notice that there should be brackets around **num** and **ind**!

4

- (b) **Idea II:** An alternative way to fetch the i^{th} byte is *Union*. *Little Dragon* wrote the following code, but he is a little confused about the concept of *little endian* and *big endian*. Help him answer the questions below!

```

1 #include <stdio.h>
2 #include <stdint.h>
3
4 /* Tip on union: data type that stores its members
5    in the same memory location */
6 typedef union {
7     struct {
8         uint8_t byte0;
9         uint8_t byte1;
10        uint8_t byte2;
11        uint8_t byte3;
12    } bytes;
13    int all_bits;
14 } MyInt;
15
16 int main() {
17     MyInt intA;
18     intA.all_bits = 0x12345678;
19     printf("%#x, %#x\n", intA.bytes.byte1, intA.bytes.byte3);
20     return 0;
21 }

```

What is the expected output (in hexadecimal format) of Line 19:

- if the system is **little endian**?

- if the system is **big endian**?

Solution:

0x56, 0x12;

0x34, 0x78.

3

- (c) **Idea III:** *Little Dragon* is fascinated in playing with bitwise operations. He wrote the following function in C.

```

1 void byte_xor(int num, int a, int b) {
2     char *ret_val = (char *) &num;
3
4     ret_val[b] ?? ret_val[a];
5     ret_val[a] ?? ret_val[b];
6     ret_val[b] ?? ret_val[a];
7
8     printf("%#x\n", num);
9 }

```


What operators are expected to substitute the ?? in Line 4, 5, and 6, such that the result of `byte_xor(0x1133CCFF, 1, 3)` will be `0xCC3311FF`?

A. `&=`, `&=`, `&=`

B. `&=`, `^=`, `^=`

C. `|=`, `^=`, `~=`

D. `^=`, `^=`, `^=`

Solution: D

6. RISC-V programming

In this question, you are asked to implement a simple recursive function in RISC-V. The function takes a decimal number as input, then outputs its octal representation encoded as decimal digits. For example, if the input to this function is 100, then the output would be 144.

The recursive function implemented in C is given below:

```
1 int find_octal(unsigned int decimal) {
2     if (decimal == 0) {
3         return 0;
4     } else {
5         return decimal % 8 + 10 * find_octal(decimal / 8);
6     }
7 }
```

A skeleton of RISC-V code is given below.

DO NOT fill in them immediately. Do some warm-ups first!

```
1 find_octal:
2     addi    sp, sp, -8
3     sw      ra, 4(sp)
4     sw      s0, 0(sp)
5
6     beq     a0, x0, _____
7
8     _____ # set s0 to something
9
10    _____ # set a0 to something
11
12    jal     ra, _____ # recursive call
13
14    _____
15    mul     a0, t0, a0
16
17    _____ # a0 = ???
18 postamble:
19
20    _____ # Restore ra
21
22    _____ # restore ...
23
24    _____ # restore ...
25 end:
26    jr      ra
```

- 2 (a) Translate the following RISC-V instructions into machine code.

sw ra, 4(sp) _____

andi s0, a0, 7 _____

Solution:

0x00112223

0x00757413

- 2 (b) What is one pseudo instruction in the RISC-V code above? How can you change it into one base instruction?

Pseudo instruction: _____

After your change: _____

Solution:

jr ra

jalr x0 ra

- 8 (c) Fill in the missing code above.

Solution:

```

beq    a0, x0, postamble
andi   s0, a0, 7
srli   a0, a0, 3
jal     ra, find_octal
li      t0, 10 / addi    t0, x0, 10
add     a0, a0, s0
lw      ra, 4(sp)
lw      s0, 0(sp)
addi    sp, sp, 8

```

7. RISC-V Basic

- 5 (a) Write a function in RISC-V code to return 0 if the input 32-bit float is an infinite value, else a non-zero value. The input and output will be stored in **a0**, as usual. Do not use pseudo instructions!

is_not_infinity:

ret # <= Return instruction

Solution:

```

1 is_not_infinity:
2     lui a1 0x80000
3     addi a1 a1 -1
4     and a0 a0 a1
5     lui a1 0x7F800
6     xor a0 a0 a1
7     ret

```

2 (b) True or False.

1. Let `a0` point to the start of an array `x`. `lw s0, 4(a0)` will always load `x[1]` into `s0`.
2. After calling a function and having that function return, the `t` registers may have been changed during the execution of the function, while `a` registers cannot.

1	2

Solution:

1. False. This only holds for data types that are four bytes wide, like `int` or `float`. For data-types like `char` that are only one byte wide, `4(a0)` is too large of an offset to return the element at index 1, and will instead return a `char` further down the array (or some other data beyond the array, depending on the array length).
2. False. `a0` and `a1` registers are often used to store the return value from a function, so the function can set their values to the its return values before returning. So all `a` registers can be used temporary registers, as specified in the green card.

8. CALL

Answer the following questions with regard to the following C program.

```

1 #include <stdio.h>
2
3 int main(int argc, char *argv[]) {
4     if (argc == (1 + 1)) {
5         printf("Hello, %s.\n", argv[1]);
6     } else {
7         printf("Goodbye.\n");
8     }
9
10    return 0;
11 }

```

- 8 (a) Select which stage of CALL is responsible for the following actions. Please fill your answer (A, B, C or D) in the table below.

A. Compiler B. Assembler C. Linker D. Loader

1. Removes all pseudo instructions.
2. Provide the address to the string "Goodbye.\n".
3. Remove most duplicate instructions in the program in order to optimize the program.
4. Put arguments in the address of `argv` so that the program could read from it.
5. Incorporating dynamic libraries so that the program could call `printf` in the C standard library.
6. Creates the symbol table so that we can know the address to the function `main` in future stages.
7. The parser is used to determine the operator precedence in `argc == (1 + 1)`.
8. Determine the jump address the `if` statement is jumping to.

1	2	3	4	5	6	7	8

Solution: B; C; A; D; D; B; A; B.

- 8 (b) True or False. Please fill your answer (T or F) in the table below.

1. Pseudo instructions are not allowed in the output of compiler.
2. Statically-linked libraries are incorporated into the program during the load stage.
3. Dynamically-linked libraries are incorporated into the program during the link stage.
4. The interpreted program (like Python) runs way faster than a compiled one (like C) in most cases.

5. The assembler takes two passes over the code to resolve PC-Relative target addresses.
6. Copying arguments passed to the program onto the stack is done during the linking stage.
7. Assembler can always provide the correct immediate value when translating all `la` instructions.
8. Compiling stage is the one most often responsible for code optimization.


1	2	3	4	5	6	7	8

Solution: F; F; F; F; T; F; F; T.

9. Logic

5


- (a) The circuit shown below can be simplified. Please **write** down the boolean expression that exactly corresponds to the circuit shown (no simplification). Then simplify this prepossession step by step, applying one rule at a time. Then **draw** the circuit according to the simplified boolean expression using the minimum number of **one-** or **two-input** logic gates.



Solution:



$$\begin{aligned}
 \overline{A}\overline{B}C + (A + \overline{B})(\overline{C}\overline{D}) &= \overline{A} + B + \overline{C} + (A + \overline{B})(\overline{C} + \overline{D}) \\
 &= \overline{A} + B + \overline{C} + A\overline{C} + A\overline{D} + \overline{B}\overline{C} + \overline{B}\overline{D} \\
 &= \overline{A} + B + \overline{C} + A\overline{D} + \overline{B}\overline{D} \\
 &= \overline{A} + B + \overline{C} + \overline{D}
 \end{aligned}$$


One possible circuit:



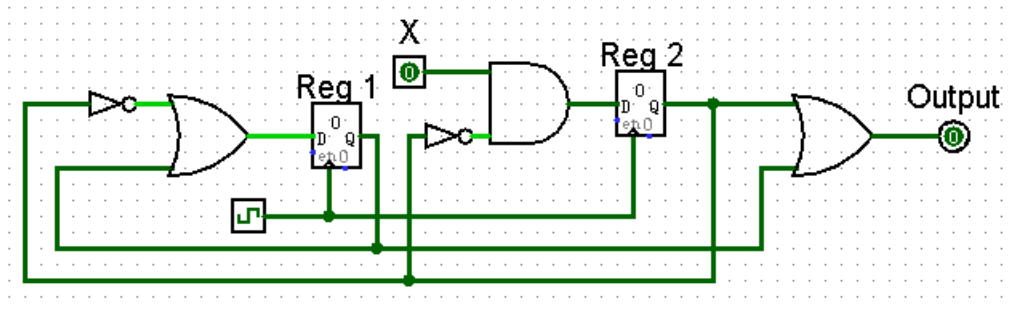
6

-



Solution:

- 4 (b) Consider the following circuit. Assume the clock has a frequency of 50 MHz, all gates have a propagation delay of 6 ns, X changes 10 ns after the rising edge of clk, Reg1 and Reg2 have a clk-to-q delay of 1 ns.



What is the **longest possible setup time** such that there are no setup time violations?

What is the **longest possible hold time** such that there are no hold time violations?

Solution:

The clock period is $\frac{1}{50 \times 10^6} s = 20 \text{ ns}$.

Reg1 longest possible setup time: the path is *the output of Reg2* \rightarrow NOT \rightarrow OR, with a delay of $1 \text{ ns} + 6 \text{ ns} + 6 \text{ ns} = 13 \text{ ns}$. So $20 - 13 = 7 \text{ ns}$.

Reg2 longest possible setup time: the path is *X changes* \rightarrow AND, with a delay of $10 \text{ ns} + 6 \text{ ns} = 16 \text{ ns}$. So $20 - 16 = 4 \text{ ns}$.


So longest setup time: $\min(7 \text{ ns}, 4 \text{ ns}) = 4 \text{ ns}$.

Reg1 longest possible hold time: the path is *the output of Reg1* \rightarrow OR, with a delay of $1 \text{ ns} + 6 \text{ ns} = 7 \text{ ns}$.

Reg2 longest possible hold time: the path is *the output of Reg2* \rightarrow NOT \rightarrow AND, with a delay of $1 \text{ ns} + 6 \text{ ns} + 6 \text{ ns} = 13 \text{ ns}$.

So longest hold time: $\min(7 \text{ ns}, 13 \text{ ns}) = 7 \text{ ns}$.

11. FSM



- 3 (a) Fill in the truth table for the FSM above.

state bit1	state bit0	input	next state bit1	next state bit0	output
0	0	0			
0	0	1			
0	1	0			
0	1	1			
1	0	0			
1	0	1			

Solution:

state bit1	state bit0	Input	next state bit1	next state bit0	Output
0	0	0	0	0	0
0	0	1	0	1	0
0	1	0	1	0	0
0	1	1	0	1	0
1	0	0	0	0	0
1	0	1	0	1	1


- 1 (b) What does the given FSM output with the input bit string '0100101010'?

Solution: It outputs '0000001010'.

- 1 (c) What does the given FSM implement (Describe when the FSM will output 1)?

Solution: When the FSM receives '101', it outputs 1.

- 2 (d) Draw a FSM that outputs 1 when it receives two or more successive '0'.



Solution:

