



信息科学与技术学院

School of Information Science and Technology

CS 110

Computer Architecture

Intro to C I

Instructors:

Siting Liu & Chundong Wang

Course website: [https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/
Spring-2023/index.html](https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html)

School of Information Science and Technology (SIST)

ShanghaiTech University

2023/2/6

Course Info

- HW1 due Feb. 16th!
- Team (Lab & project) partners are required to be within the same lab session! Decide before Feb. 11th!
- Labs & Projects must be done in a group of two students. Please let the TA know immediately if you cannot find a partner. It is not allowed to change your lab-mate after this week.
- Lab 1 is available and in this week's Lab session
- Lab & Discussion starts this week.
- <https://piazza.com/shanghaitech.edu.cn/spring2023/cs110> (access code: **uutib6ruvql**)

Review

- Moore's Law; Amdahl's Law; Dennard scaling
- Binary system: Integer (Unsigned & Signed integer)
 - We use "0x" as the prefix of hexadecimal number
- We use 2's complement to represent signed integer in modern computer: easy arithmetic (addition & subtraction)
- C: **portable** (GPU/CUDA, DSP, MCU, 寒武纪MLU/Bang C, etc.) and **efficient** (used in building UNIX/MATLAB/python, etc.); utilized to understand how computer works in this course.
- Compiler first step: C pre-processing (text editing for further compiling steps)

Function-Like Macro

- #define MAG(x, y) (sqrt((x)*(x) + (y)*(y)))

```
#include <stdio.h>
```

```
%clang/gcc -E introC_1_0.c
```

```
#include <math.h>
```

```
#define MAG0(x, y) sqrt(x*x + y*y)
```

```
#define MAG(x, y) (sqrt((x)*(x) + (y)*(y)))
```

```
#define MAG2(x,y) ({double a=x; double b=y; sqrt(a*a + b*b);})
```

```
#define MSG "Hello \
```

```
World!\n"
```

```
int main() {
```

```
#ifdef MSG
```

```
    printf(MSG /* "hi!\n" */);
```

```
#endif
```

```
    printf("%f\n",MAG(3.0,4.0));
```

```
    double i=2, j=3, k0, k1, k2, k3;
```

```
    double c=2, d=3;
```

```
    k0=MAG0(i+1,j+1);
```

```
    k1=MAG(i+1,j+1);
```

```
    k2=MAG(++i,++j);
```

```
    k3=MAG2(++c,++d);
```

```
    printf("%f\n",k0);
```

```
    printf("%f\n",k1);
```

```
    printf("%f\n",k2);
```

```
    printf("%f\n",k3);
```

```
    return 0;
```

```
}
```

```
k0=sqrt(i+1*i+1 + j+1*j+1);
```

```
k1=(sqrt(((i+1)*(i+1) + (j+1)*(j+1))));
```

```
k2=(sqrt(((++i)*(++i) + (++j)*(++j))));
```

```
k3=({double a=++c; double b=++d; sqrt(a*a + b*b);});
```

=> Convention: put parenthesis EVERYWHERE!

CPP Macro II

- Avoid using macros whenever possible
- NO or very tiny speedup.
- Instead use C functions – e.g. inline function:

```
double mag(double x, double y);
```

```
double inline mag(double x, double y)
```

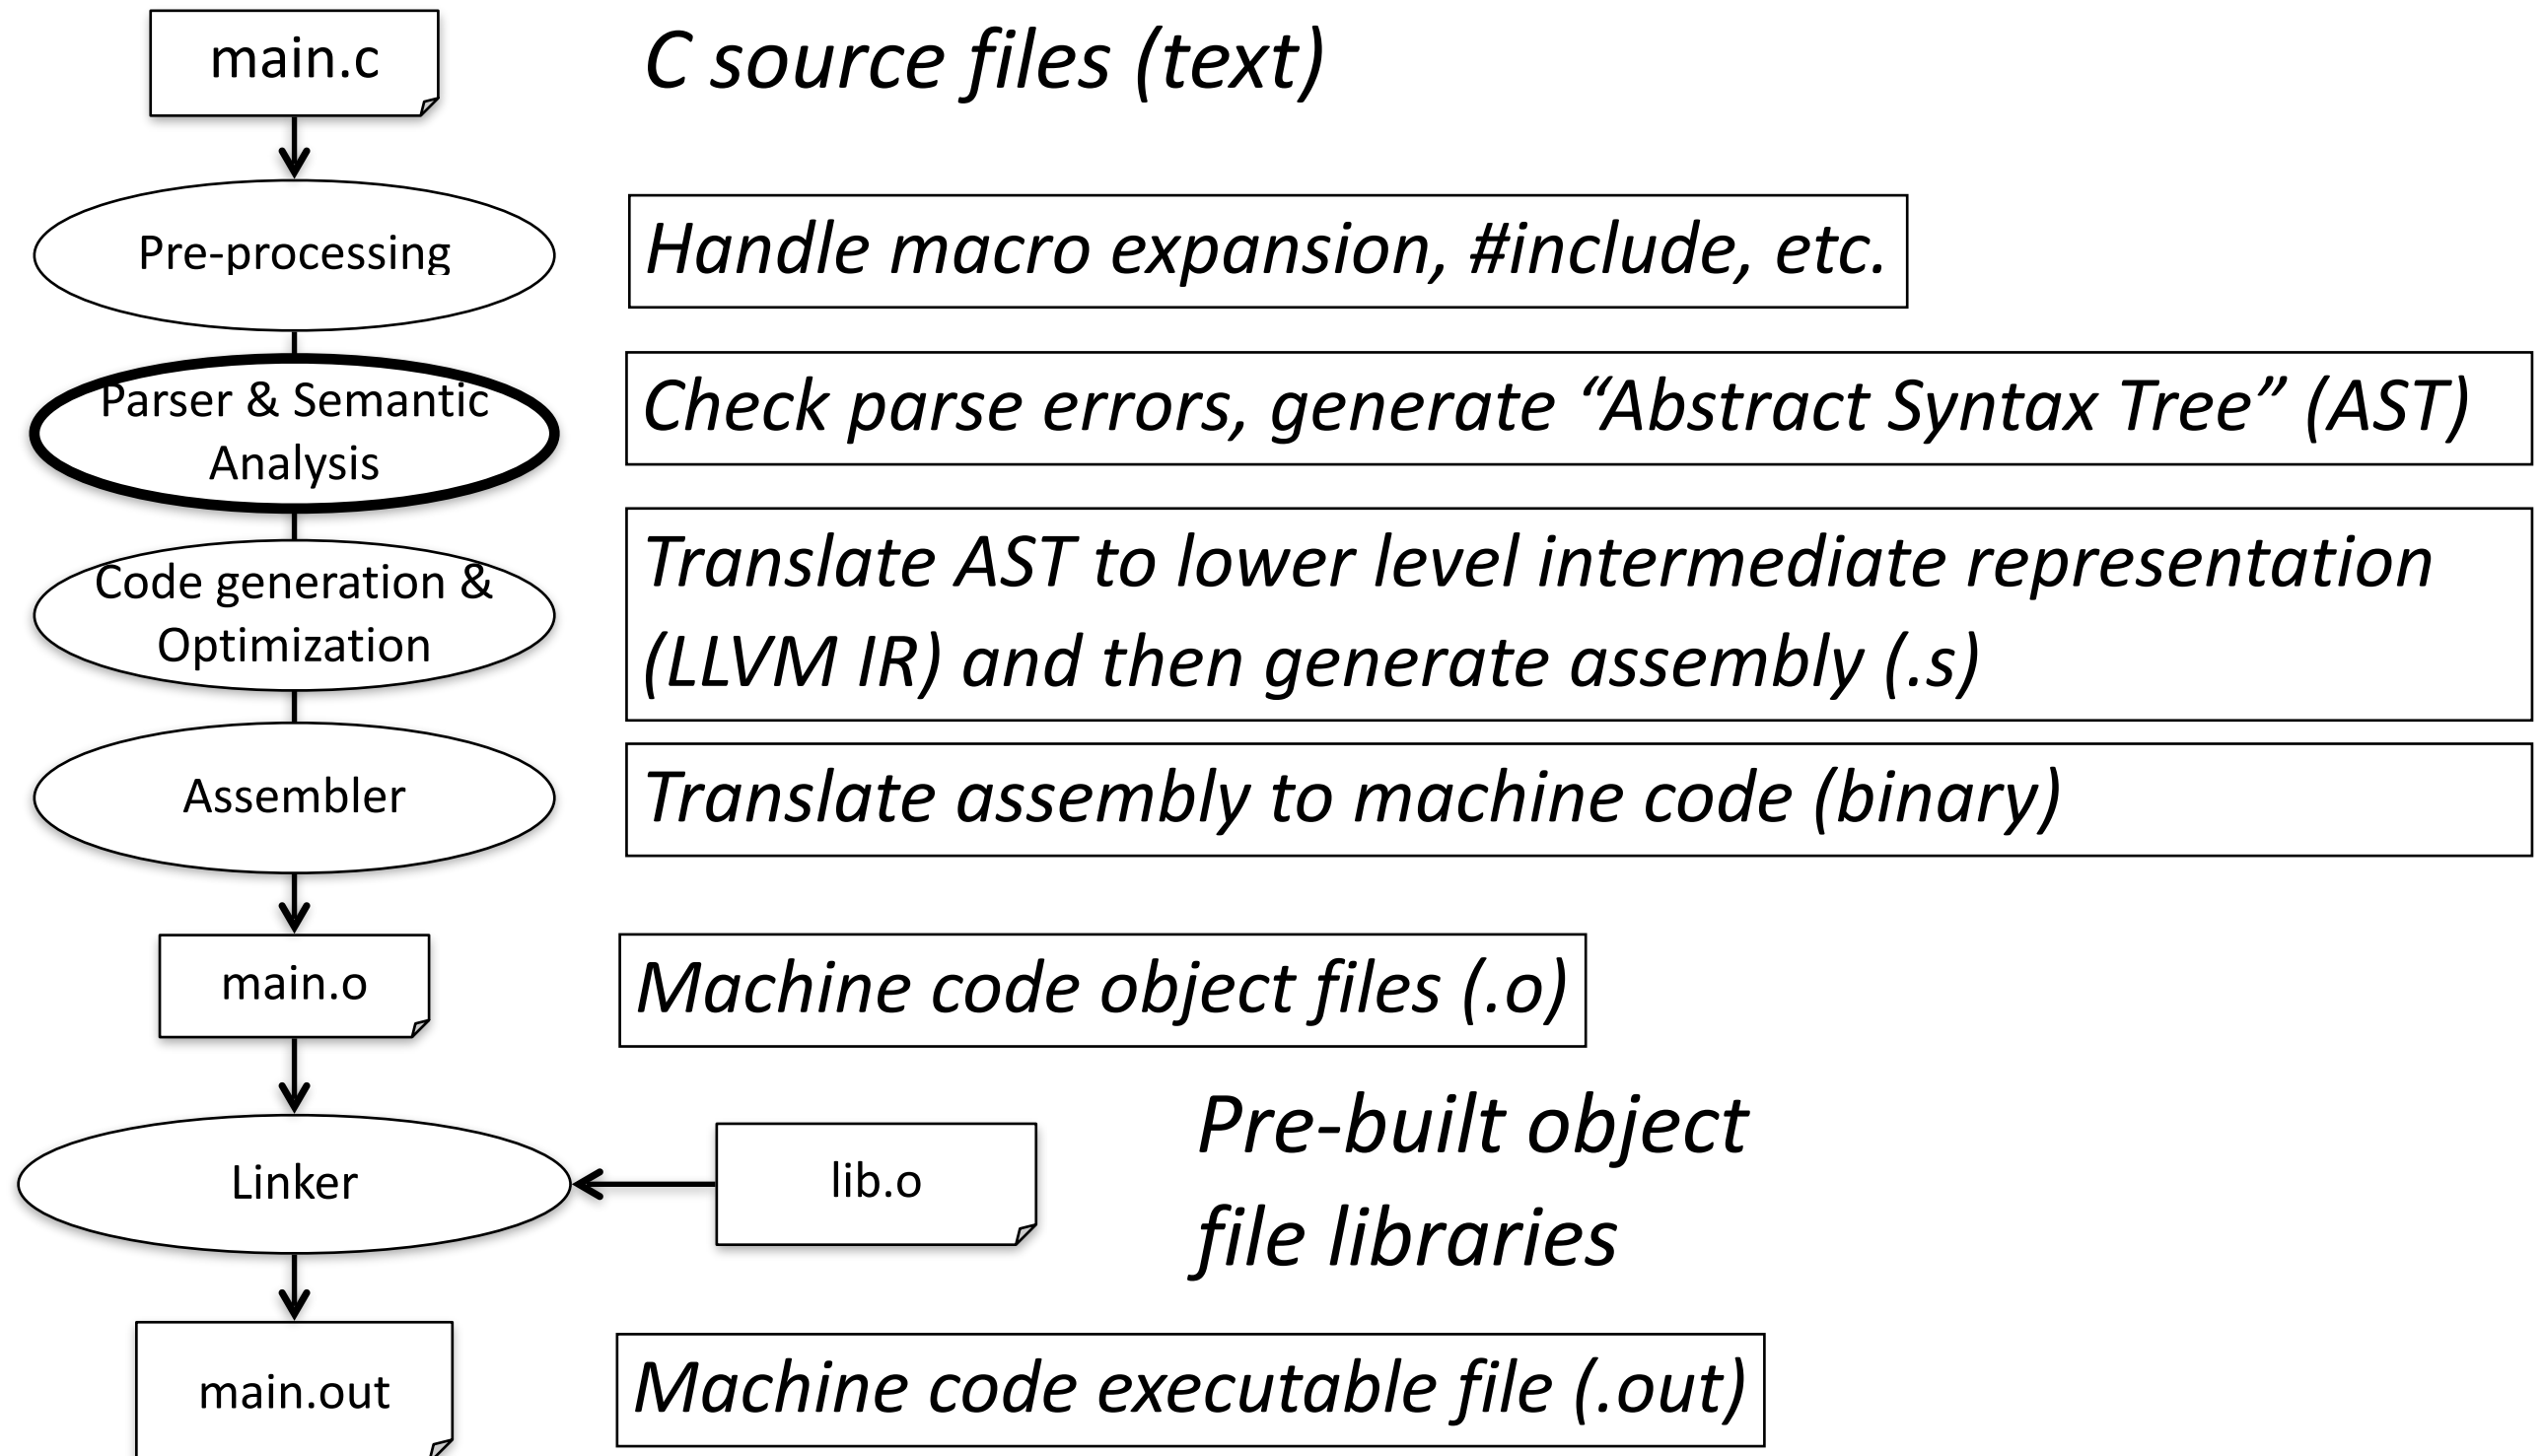
```
{ return sqrt( x*x + y*y ); }
```

Read more...

RTFM: <https://gcc.gnu.org/onlinedocs/cpp/Macros.html>

https://chunminchang.gitbooks.io/cplusplus-learning-note/content/Appendix/preprocessor_macros_vs_inline_functions.html

Outline



Parser & Semantic Analysis

- Recognize each code word as a “token” (identifiers/symbols, C keywords, constant, comma, semicolon, etc.)
- Record the location of each token

```
%clang -fsyntax-only -Xclang -dump-tokens introC_1_1.c
```

```
#include <stdio.h>
int main() { //compute 1234 + 4321
    int x = 1234, y = 4321;
    int z = x+y;
    printf("z=%d/n", z);
    return 0;
}
```

```
int 'int' [StartOfLine] [LeadingSpace] Loc=<introC_1_1.c:3:3>
identifier 'x' [LeadingSpace] Loc=<introC_1_1.c:3:7>
equal '=' [LeadingSpace] Loc=<introC_1_1.c:3:9>
numeric_constant '1234' [LeadingSpace] Loc=<introC_1_1.c:3:11>
comma ',' Loc=<introC_1_1.c:3:15>
identifier 'y' [LeadingSpace] Loc=<introC_1_1.c:3:17>
equal '=' [LeadingSpace] Loc=<introC_1_1.c:3:19>
numeric_constant '4321' [LeadingSpace] Loc=<introC_1_1.c:3:21>
semi ';' Loc=<introC_1_1.c:3:25>
int 'int' [StartOfLine] [LeadingSpace] Loc=<introC_1_1.c:4:3>
identifier 'z' [LeadingSpace] Loc=<introC_1_1.c:4:7>
equal '=' [LeadingSpace] Loc=<introC_1_1.c:4:9>
identifier 'x' [LeadingSpace] Loc=<introC_1_1.c:4:11>
plus '+' Loc=<introC_1_1.c:4:12>
identifier 'y' Loc=<introC_1_1.c:4:13>
semi ';' Loc=<introC_1_1.c:4:14>
identifier 'printf' [StartOfLine] [LeadingSpace] Loc=<introC_1_1.c:5:3>
l_paren '(' Loc=<introC_1_1.c:5:9>
string_literal '"z=%d/n"' Loc=<introC_1_1.c:5:10>
comma ',' Loc=<introC_1_1.c:5:18>
identifier 'z' Loc=<introC_1_1.c:5:19>
r_paren ')' Loc=<introC_1_1.c:5:20>
semi ';' Loc=<introC_1_1.c:5:21>
return 'return' [StartOfLine] [LeadingSpace] Loc=<introC_1_1.c:6:3>
numeric_constant '0' [LeadingSpace] Loc=<introC_1_1.c:6:10>
semi ';' Loc=<introC_1_1.c:6:11>
r_brace '}' [StartOfLine] Loc=<introC_1_1.c:7:1>
```

Lexer

Parser & Semantic Analysis

- Organize tokens as “AST” tree
- Report errors

```
% clang -fsyntax-only -Xclang -ast-dump introC_1_1.c
```

```
#include <stdio.h>
```

```
int main() { //compute 1234 + 4321
```

```
    int x = 1234, y = 4321;
```

```
    int z = x+y;
```

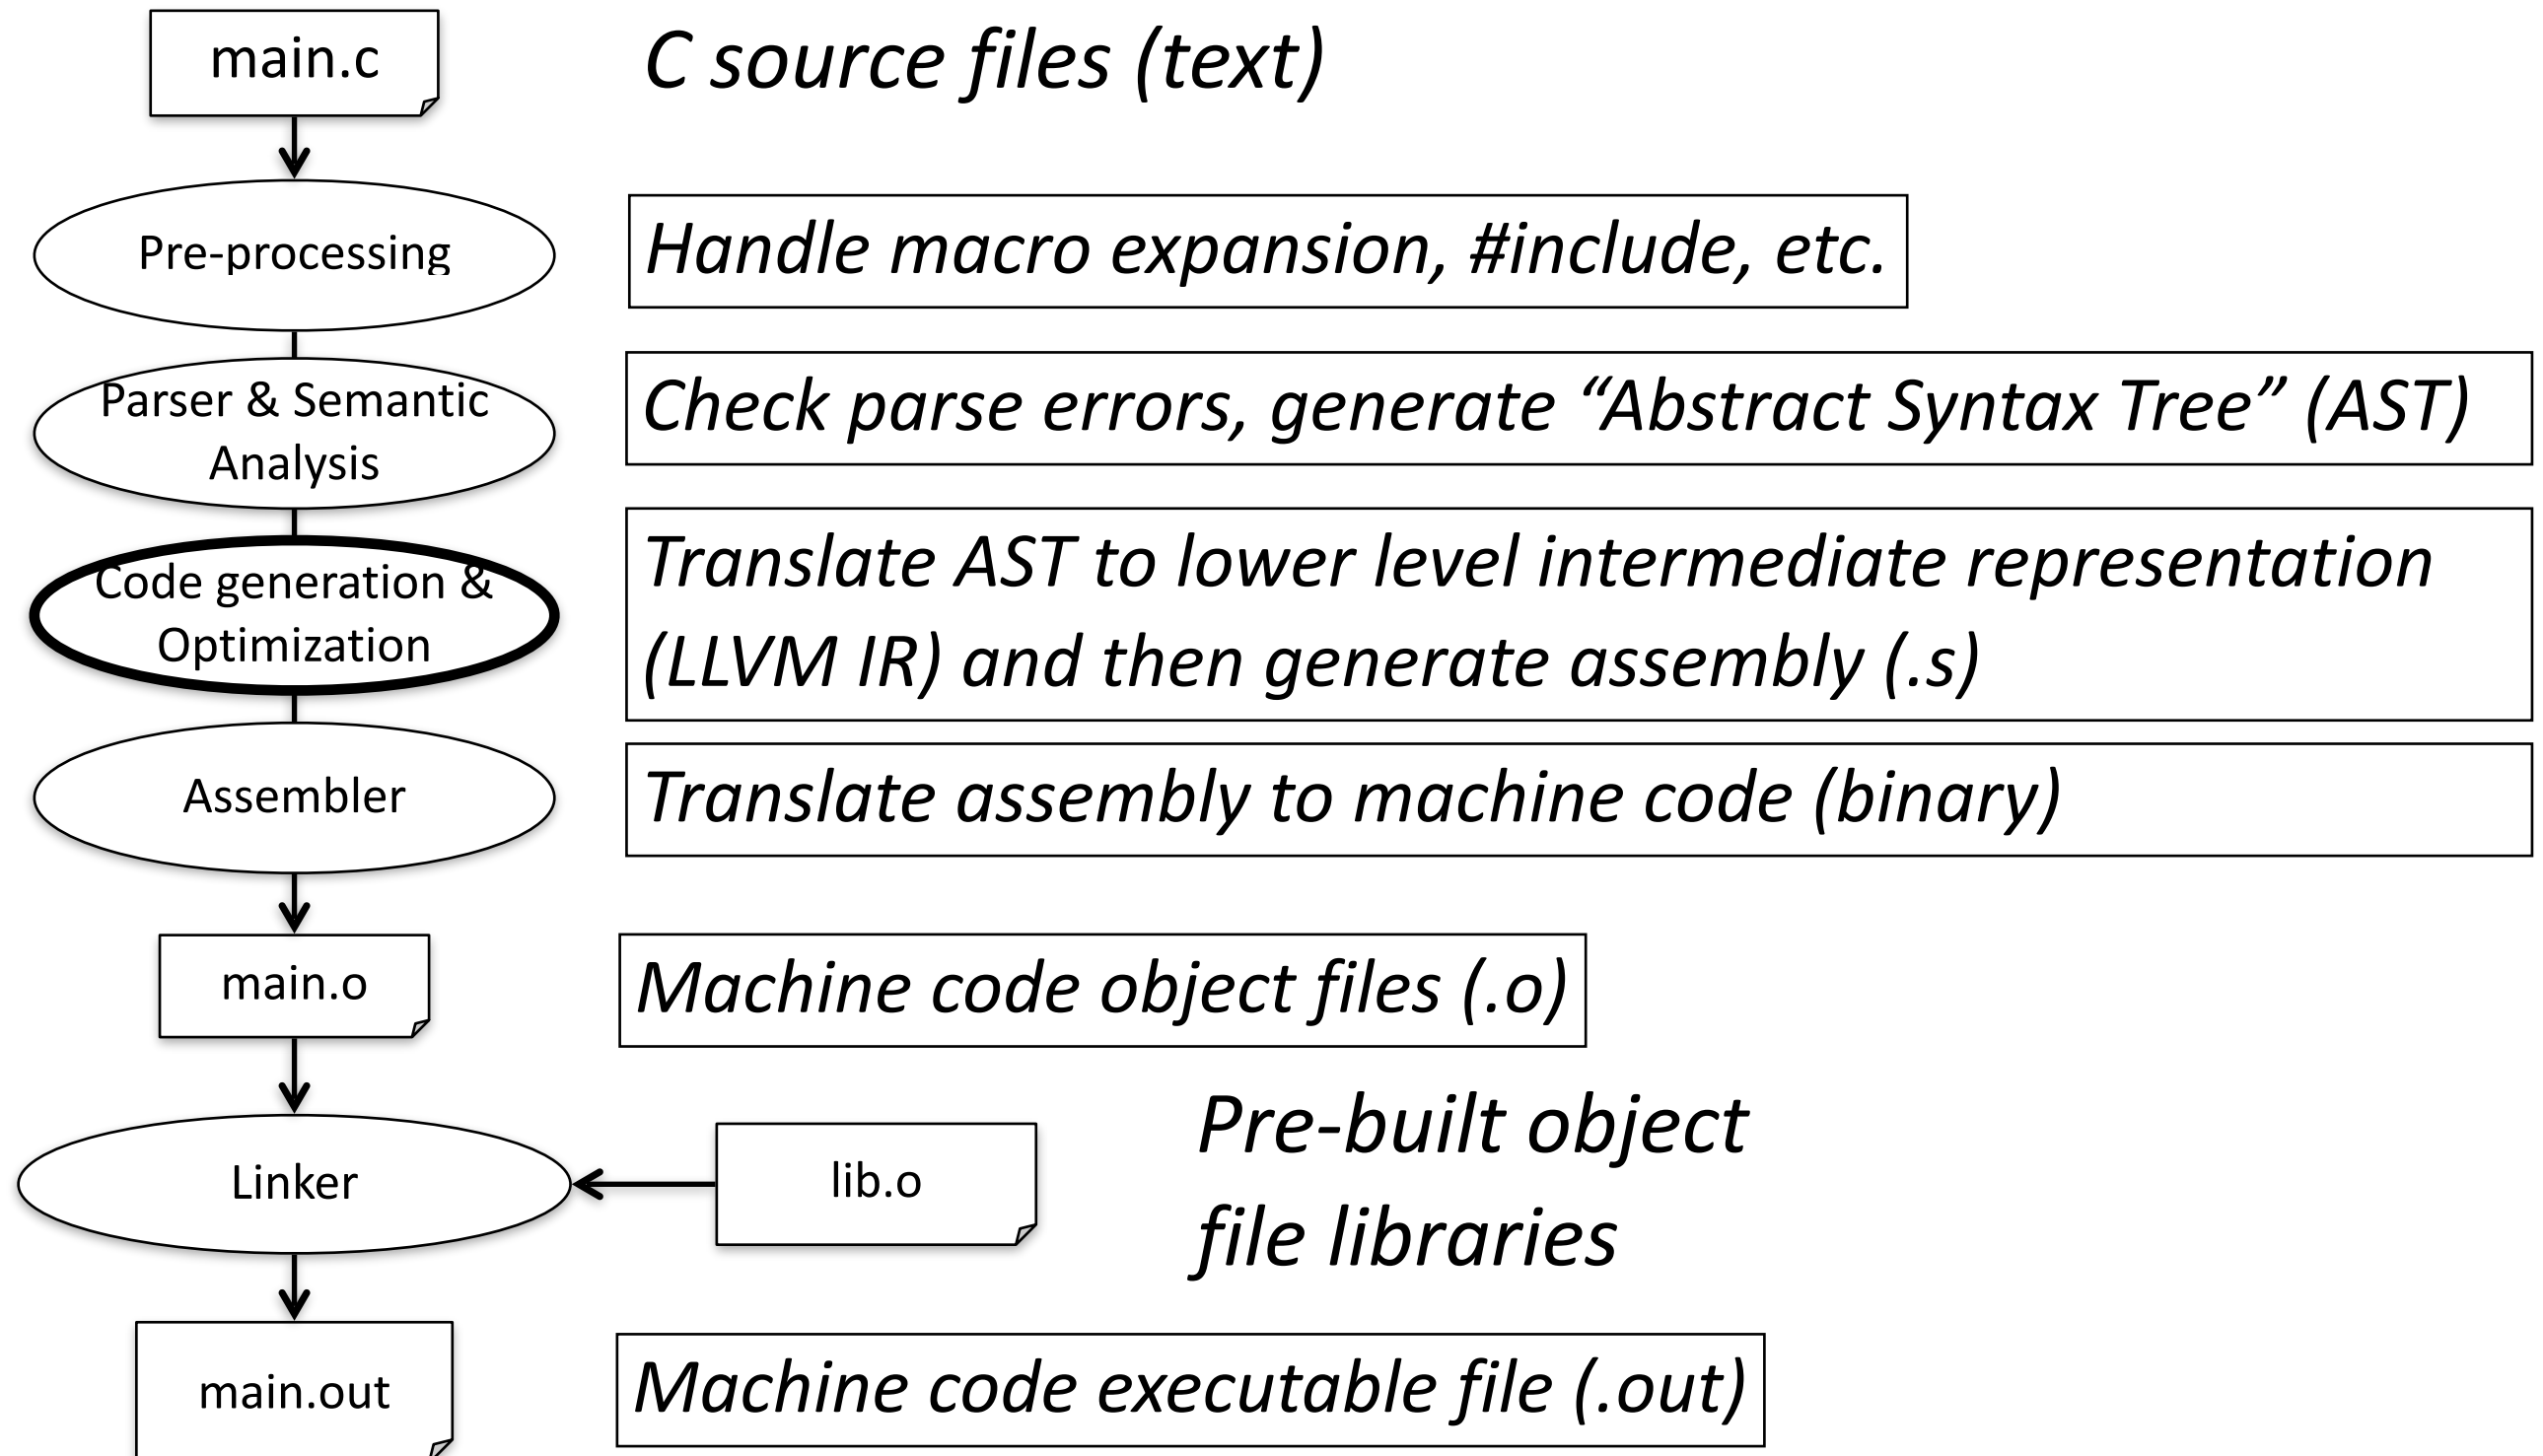
```
    printf("z=%d/n", z);
```

```
    return 0;
```

```
}
```

```
-FunctionDecl 0x1590f8600 <introC_1_1.c:2:1, line:7:1> line:2:5 main 'int ()'
-CompoundStmt 0x1590f8aa0 <col:12, line:7:1>
-DeclStmt 0x1590f87f8 <line:3:3, col:25>
  -VarDecl 0x1590f86b8 <col:3, col:11> col:7 used x 'int' cinit
    -IntegerLiteral 0x1590f8720 <col:11> 'int' 1234
  -VarDecl 0x1590f8758 <col:3, col:21> col:17 used y 'int' cinit
    -IntegerLiteral 0x1590f87c0 <col:21> 'int' 4321
-DeclStmt 0x1590f8920 <line:4:3, col:14>
  -VarDecl 0x1590f8828 <col:3, col:13> col:7 used z 'int' cinit
    -BinaryOperator 0x1590f8900 <col:11, col:13> 'int' '+'
      -ImplicitCastExpr 0x1590f88d0 <col:11> 'int' <LValueToRValue>
        -DeclRefExpr 0x1590f8890 <col:11> 'int' lvalue Var 0x1590f86b8 'x' 'int'
      -ImplicitCastExpr 0x1590f88e8 <col:13> 'int' <LValueToRValue>
        -DeclRefExpr 0x1590f88b0 <col:13> 'int' lvalue Var 0x1590f8758 'y' 'int'
-CallExpr 0x1590f89f8 <line:5:3, col:20> 'int'
  -ImplicitCastExpr 0x1590f89e0 <col:3> 'int (*)(const char *, ...)' <FunctionToPointerDeca
y>
  -DeclRefExpr 0x1590f8938 <col:3> 'int (const char *, ...)' Function 0x1590dd388 'printf
' 'int (const char *, ...)'
```


C Compilation Simplified Overview (more later in course)



Code Generation & Optimization

- Generate intermediate representation (IR)
 - LLVM IR for clang/LLVM
 - GIMPLE for gcc

```
%clang -S -emit-llvm introC_1_1.c -o introC_1_1.ll
```

```
#include <stdio.h>
int main() { //compute 1234
    int x = 1234, y = 4321;
    int z = x+y;
    printf("z=%d/n", z);
    return 0;
}
```


```
; ModuleID = 'introC_1_1.c'
source_filename = "introC_1_1.c"
target datalayout = "e-m:o-i64:64-i128:128-n32:64-S128"
target triple = "arm64-apple-macosx12.0.0"

@.str = private unnamed_addr constant [7 x i8] c"z=%d/n\00", align 1

; Function Attrs: noinline nounwind optnone ssp uwtable
define i32 @main() #0 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %3 = alloca i32, align 4
    %4 = alloca i32, align 4
    store i32 0, i32* %1, align 4
    store i32 1234, i32* %2, align 4
    store i32 4321, i32* %3, align 4
    %5 = load i32, i32* %2, align 4
    %6 = load i32, i32* %3, align 4
    %7 = add nsw i32 %5, %6
    store i32 %7, i32* %4, align 4
    %8 = load i32, i32* %4, align 4
    %9 = call i32 @printf(i8* getelementptr inbounds ([7 x i8]
0), i32 %8)
    ret i32 0
}
```

Components of Computers

- Von Neumann Architecture
 - *First Draft of a Report on the EDVAC*



By John von Neumann - <https://archive.org/stream/firstdraftofrepooovonn#page/n1/mode/2up>, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=26685284>

Central arithmetic (CA)

Central control (CC)

Central Processing Unit
(CPU)

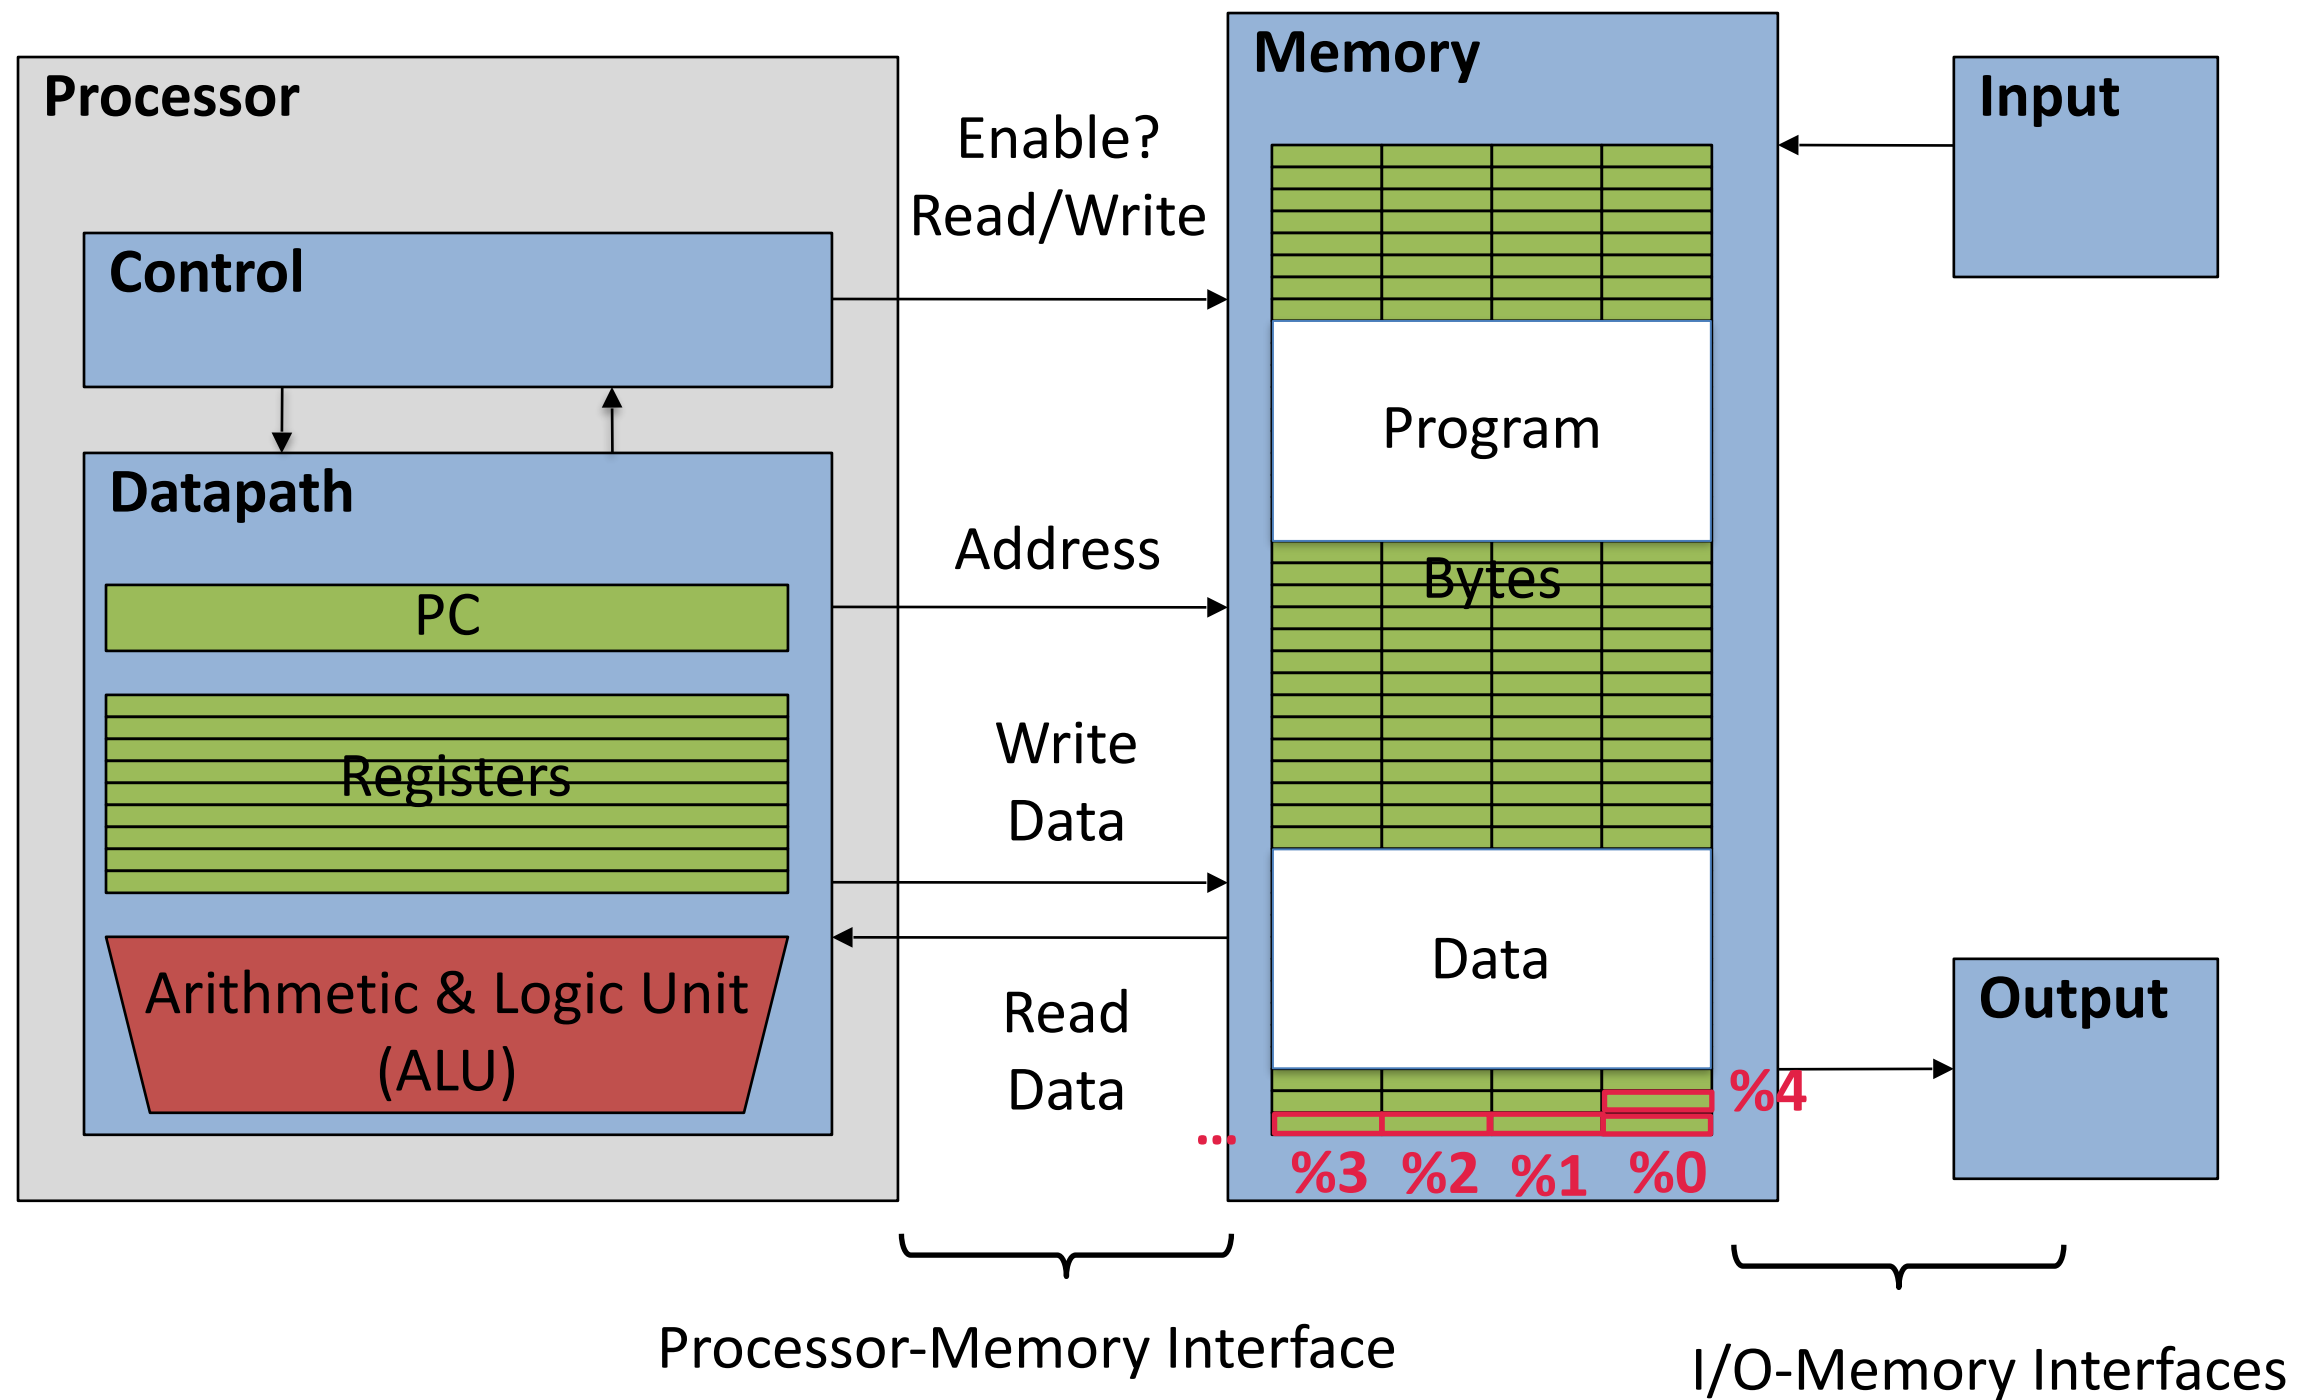
Memory (M)
(Data & Program/Instructions)

Input (I)

Output (O)

External memory (R)

Components of Computers



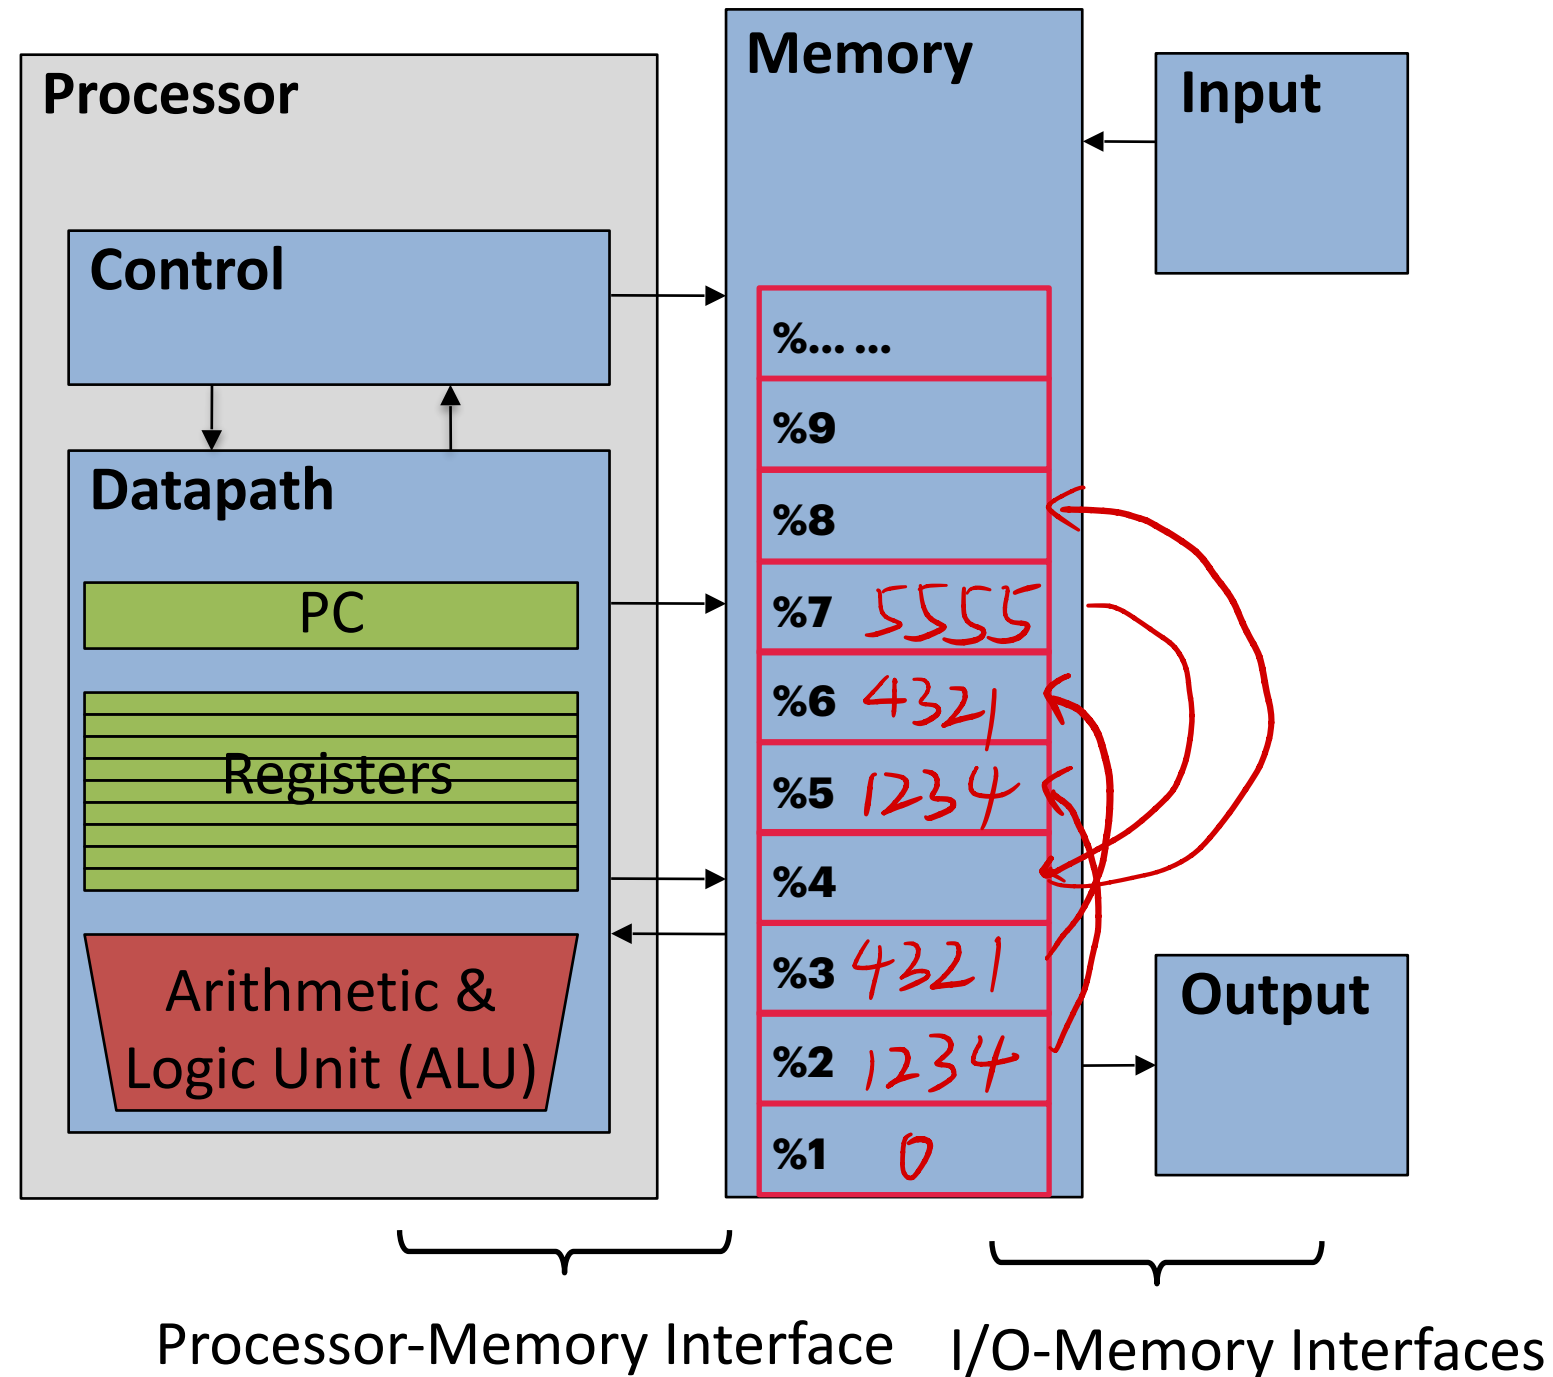
IR Implication on Hardware

```
#include <stdio.h>
int main() { //compute 1234 + 4321
    int x = 1234, y = 4321;
    int z = x+y;
    printf("z=%d/n", z);
    return 0;
}
```


Original
code

LLVM IR

```
define i32 @main() #0 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %3 = alloca i32, align 4
    %4 = alloca i32, align 4
    store i32 0, i32* %1, align 4
    store i32 1234, i32* %2, align 4
    store i32 4321, i32* %3, align 4
    %5 = load i32, i32* %2, align 4
    %6 = load i32, i32* %3, align 4
    %7 = add nsw i32 %5, %6
    store i32 %7, i32* %4, align 4
    %8 = load i32, i32* %4, align 4
    %9 = call i32 (i8*, ...) @printf(i8*
    getelementptr inbounds ([7 x i8], [7
    x i8]* @.str, i64 0, i64 0), i32 %8)
    ret i32 0
}
```



Optimization



IR to Assembly

```
% clang -S introC_1_1.c -o introC_1_1.s
```

```
#include <stdio.h>
int main() { //compute 1234 + 4321
    int x = 1234, y = 4321;
    int z = x+y;
    printf("z=%d/n",z);
    return 0;
}
```

Original
code

```
define i32 @main() #0 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %3 = alloca i32, align 4
    %4 = alloca i32, align 4
    store i32 0, i32* %1, align 4
    store i32 1234, i32* %2, align 4
    store i32 4321, i32* %3, align 4
    %5 = load i32, i32* %2, align 4
    %6 = load i32, i32* %3, align 4
    %7 = add nsw i32 %5, %6
    store i32 %7, i32* %4, align 4
    %8 = load i32, i32* %4, align 4
    %9 = call i32 @printf(i8*
getelementptr inbounds ([7 x i8], [7 x i8]*
@.str, i64 0, i64 0), i32 %8)
    ret i32 0
}
```

LLVM IR

```
.section __TEXT,__text,regular,pure_instructions
.build_version macos, 12, 0      sdk_version 13, 1
.globl _main                      ; -- Begin function

main
.p2align 2
_main:                             ; @main
    .cfi_startproc
; %bb.0:
    sub sp, sp, #48
    stp x29, x30, [sp, #32]        ; 16-byte Folded Spill
    add x29, sp, #32
    .cfi_def_cfa w29, 16
    .cfi_offset w30, -8
    .cfi_offset w29, -16
    mov w8, #0
    str w8, [sp, #12]              ; 4-byte Folded Spill
    sturwzr, [x29, #-4]
    mov w8, #1234
    sturw8, [x29, #-8]
    mov w8, #4321
    sturw8, [x29, #-12]
    ldurw8, [x29, #-8]
    ldurw9, [x29, #-12]
    add w8, w8, w9
    str w8, [sp, #16]
    ldr w9, [sp, #16]
    ...
    add sp, sp, #48
    ret
```

ARM Assembly
(Hardware abstraction)

Translated to machine code
defined by ISA

IR to Assembly to Machine Code

```
% clang -c introC_1_1.c -o introC_1_1.o
```

```
% objdump -d introC_1_1.o
```

Disassembly of section __TEXT,__text:

0000000000000000 <ltmp0>:

```
0: ff c3 00 d1 : sub sp, sp, #48
4: fd 7b 02 a9 : stp x29, x30, [sp, #32]
8: fd 83 00 91 : add x29, sp, #32
c: 08 00 80 52 : mov w8, #0
10: e8 0f 00 b9 : str w8, [sp, #12]
14: bf c3 1f b8 : stur wzr, [x29, #-4]
18: 48 9a 80 52 : mov w8, #1234
1c: a8 83 1f b8 : stur w8, [x29, #-8]
20: 28 1c 82 52 : mov w8, #4321
24: a8 43 1f b8 : stur w8, [x29, #-12]
28: a8 83 5f b8 : ldur w8, [x29, #-8]
2c: a9 43 5f b8 : ldur w9, [x29, #-12]
30: 08 01 09 0b : add w8, w8, w9
34: e8 13 00 b9 : str w8, [sp, #16]
38: e9 13 40 b9 : ldr w9, [sp, #16]
3c: e8 03 09 aa : mov x8, x9
40: e9 03 00 91 : mov x9, sp
44: 28 01 00 f9 : str x8, [x9]
48: 00 00 00 90 : adrp x0, 0x0 <ltmp0+0x48>
4c: 00 00 00 91 : add x0, x0, #0
50: 00 00 00 94 : bl 0x50 <ltmp0+0x50>
54: e0 0f 40 b9 : ldr w0, [sp, #12]
58: fd 7b 42 a9 : ldp x29, x30, [sp, #32]
5c: ff c3 00 91 : add sp, sp, #48
60: c0 03 5f d6 : ret
```

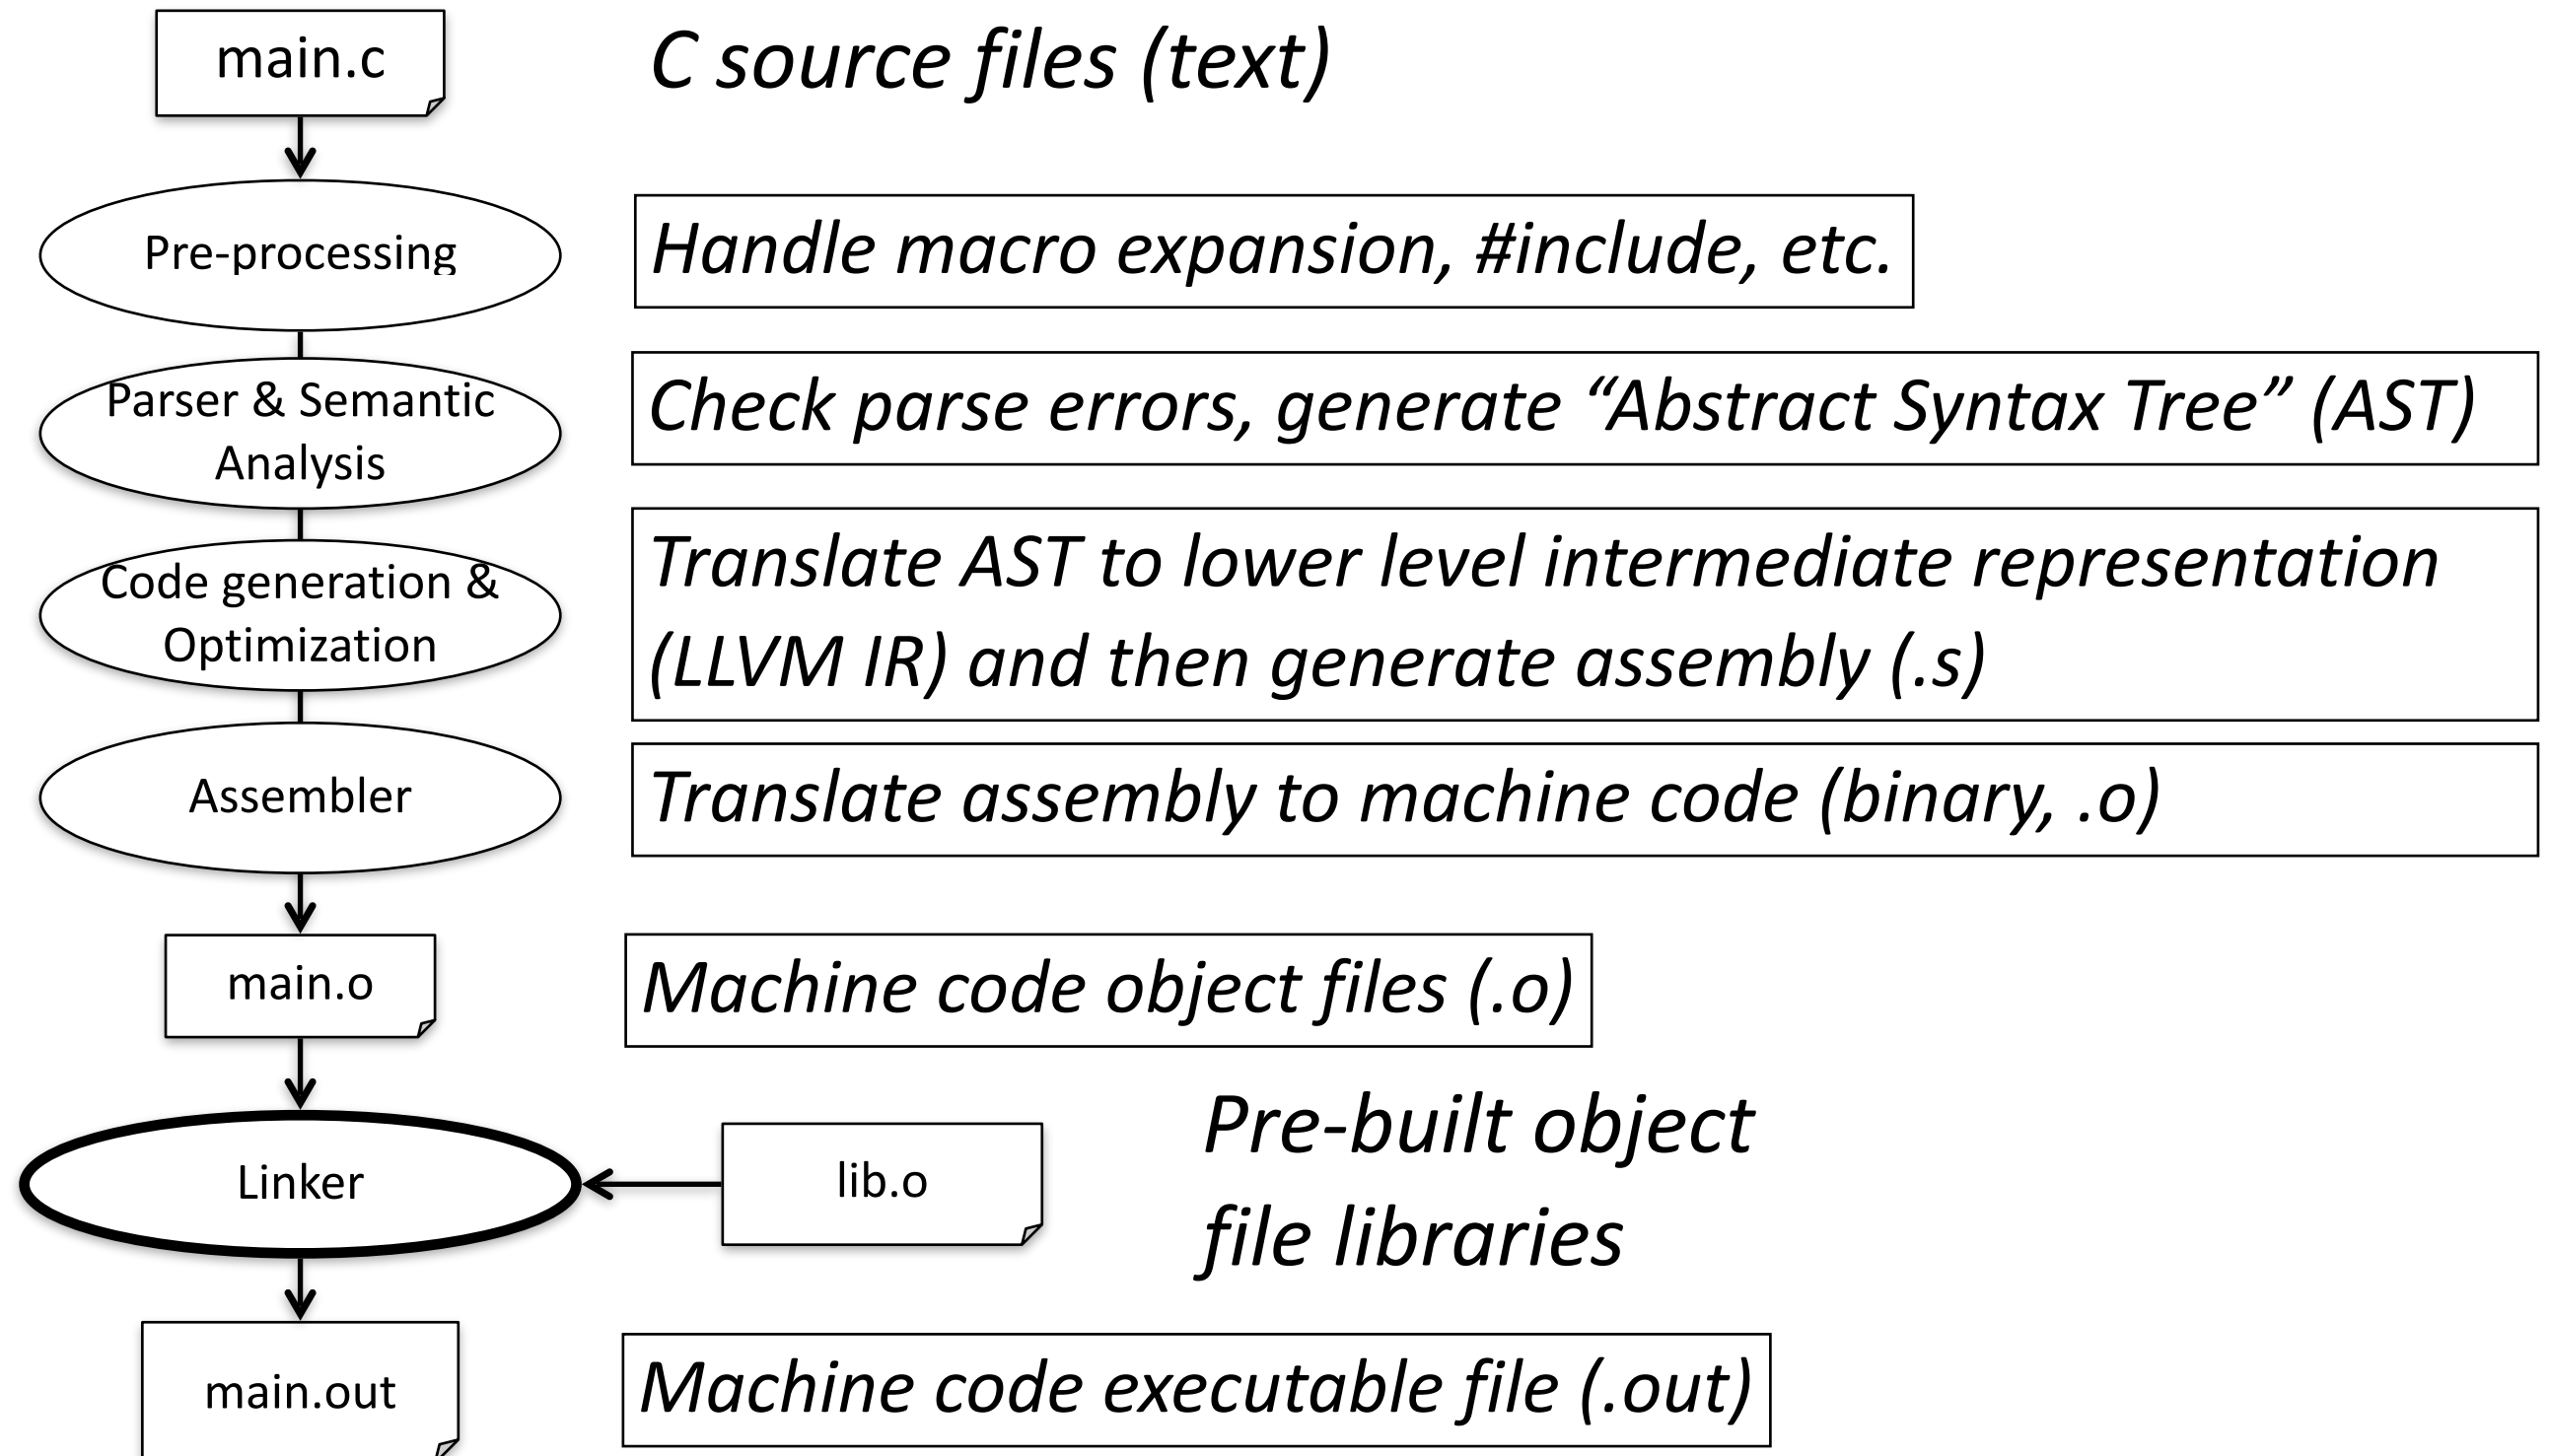
Machine Code

(Stored program/
instructions)

[→12


ARM Assembly
(ARM ISA)

C Compilation Simplified Overview



Organization of Computers

- Von Neumann Architecture
 - a.k.a. Princeton architecture
 - Uniform memory for data & program/instruction
- Harvard Architecture
 - Separated memory for data & program
 - E.g. MCU, DSP, L1 Cache

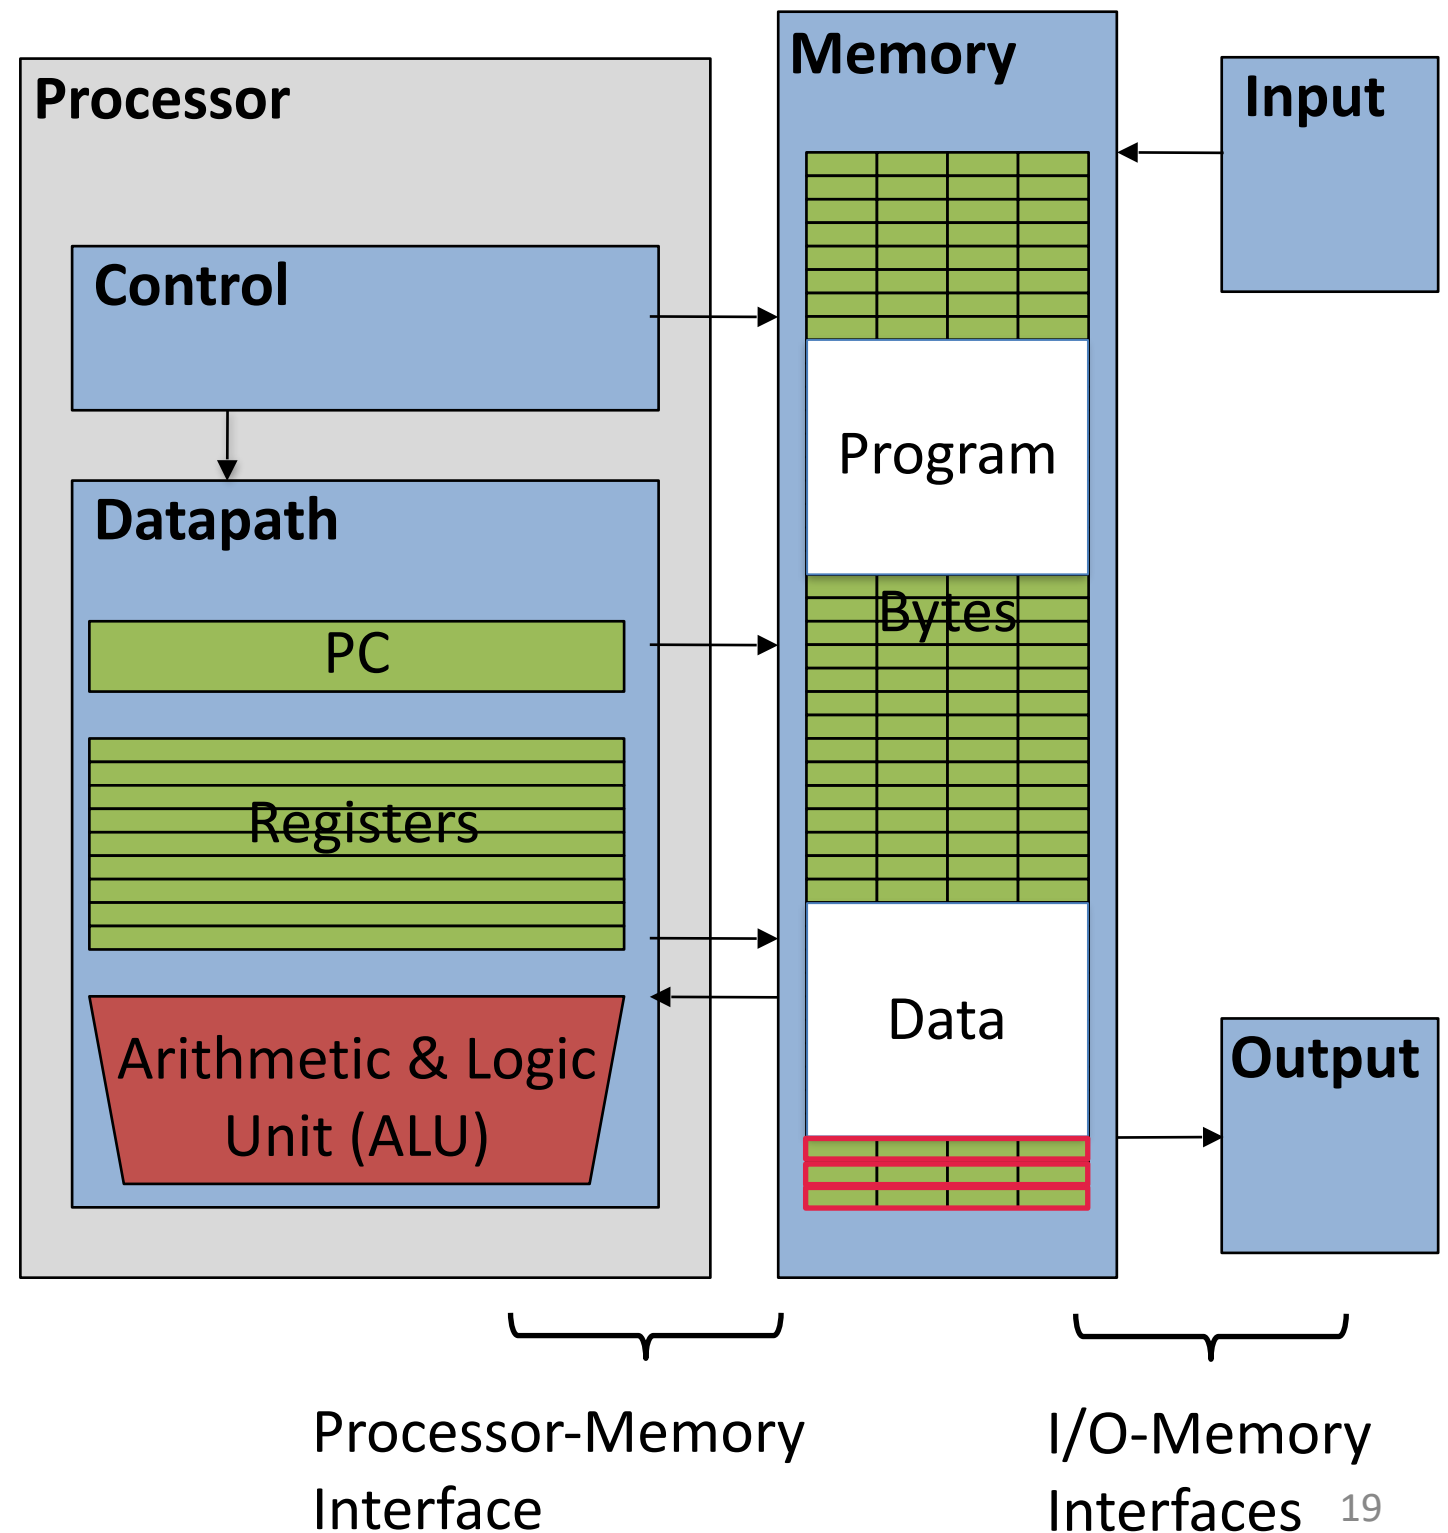


Wrap-it-up

- From C to machine code (clang *.c → *.out & ./*.out)
 - Pre-processing (macro, **function-like macro**, text editing, #include)
 - Use “()” whenever necessary, or use “function” directly
 - Parser & Semantic Analysis (tokenization & generate AST, basic operations)
 - Translate to IR & optimize (computer components)
 - Translate to assembly and then machine code, executed by hardware (**Covered in future lectures**)
 - Clang manual: <https://releases.llvm.org/14.0.0/tools/clang/docs/UsersManual.html>
 - GCC: <https://gcc.gnu.org/>

Wrap-it-up

- Von Neumann Architecture
- Harvard Architecture
- Stored-program computer



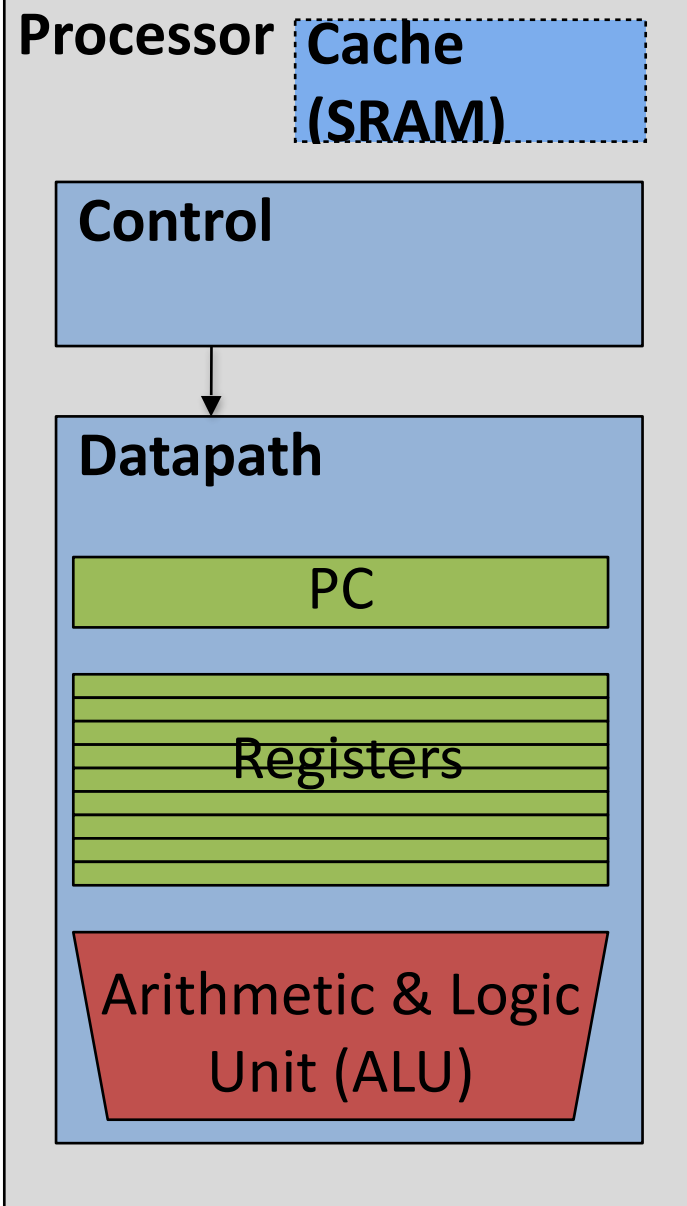
Real Stuff

Intel i7 12700 4.90 GHz




<https://maj191.com/product/intel-core-i7-12700-3-6ghz-cpu-25mb-cache-lga1700-tray/>

https://www.ocinside.de/review/intel_core_i7_12700k/3/




DIMM DDR4 5066 MHz



Kingston
8/16/32G
SDRAM

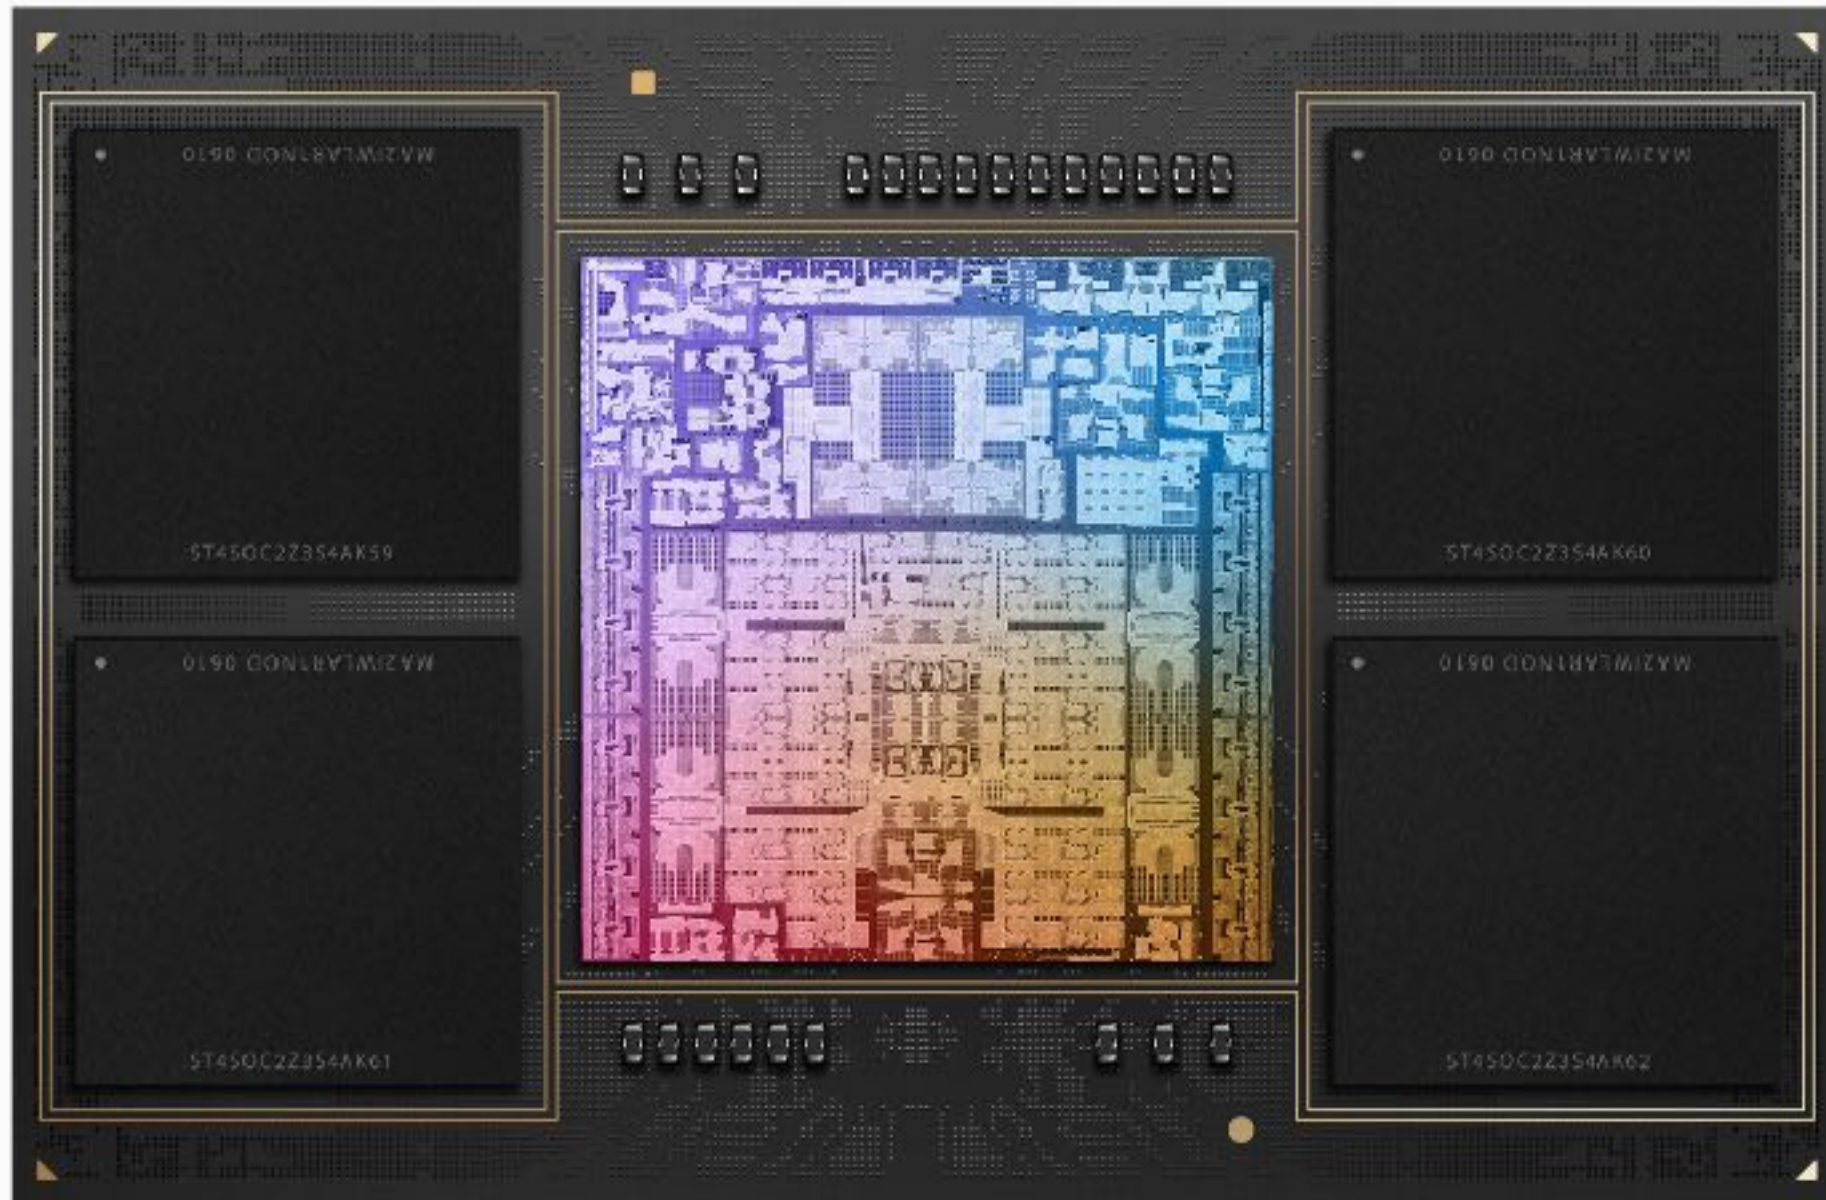


Memory



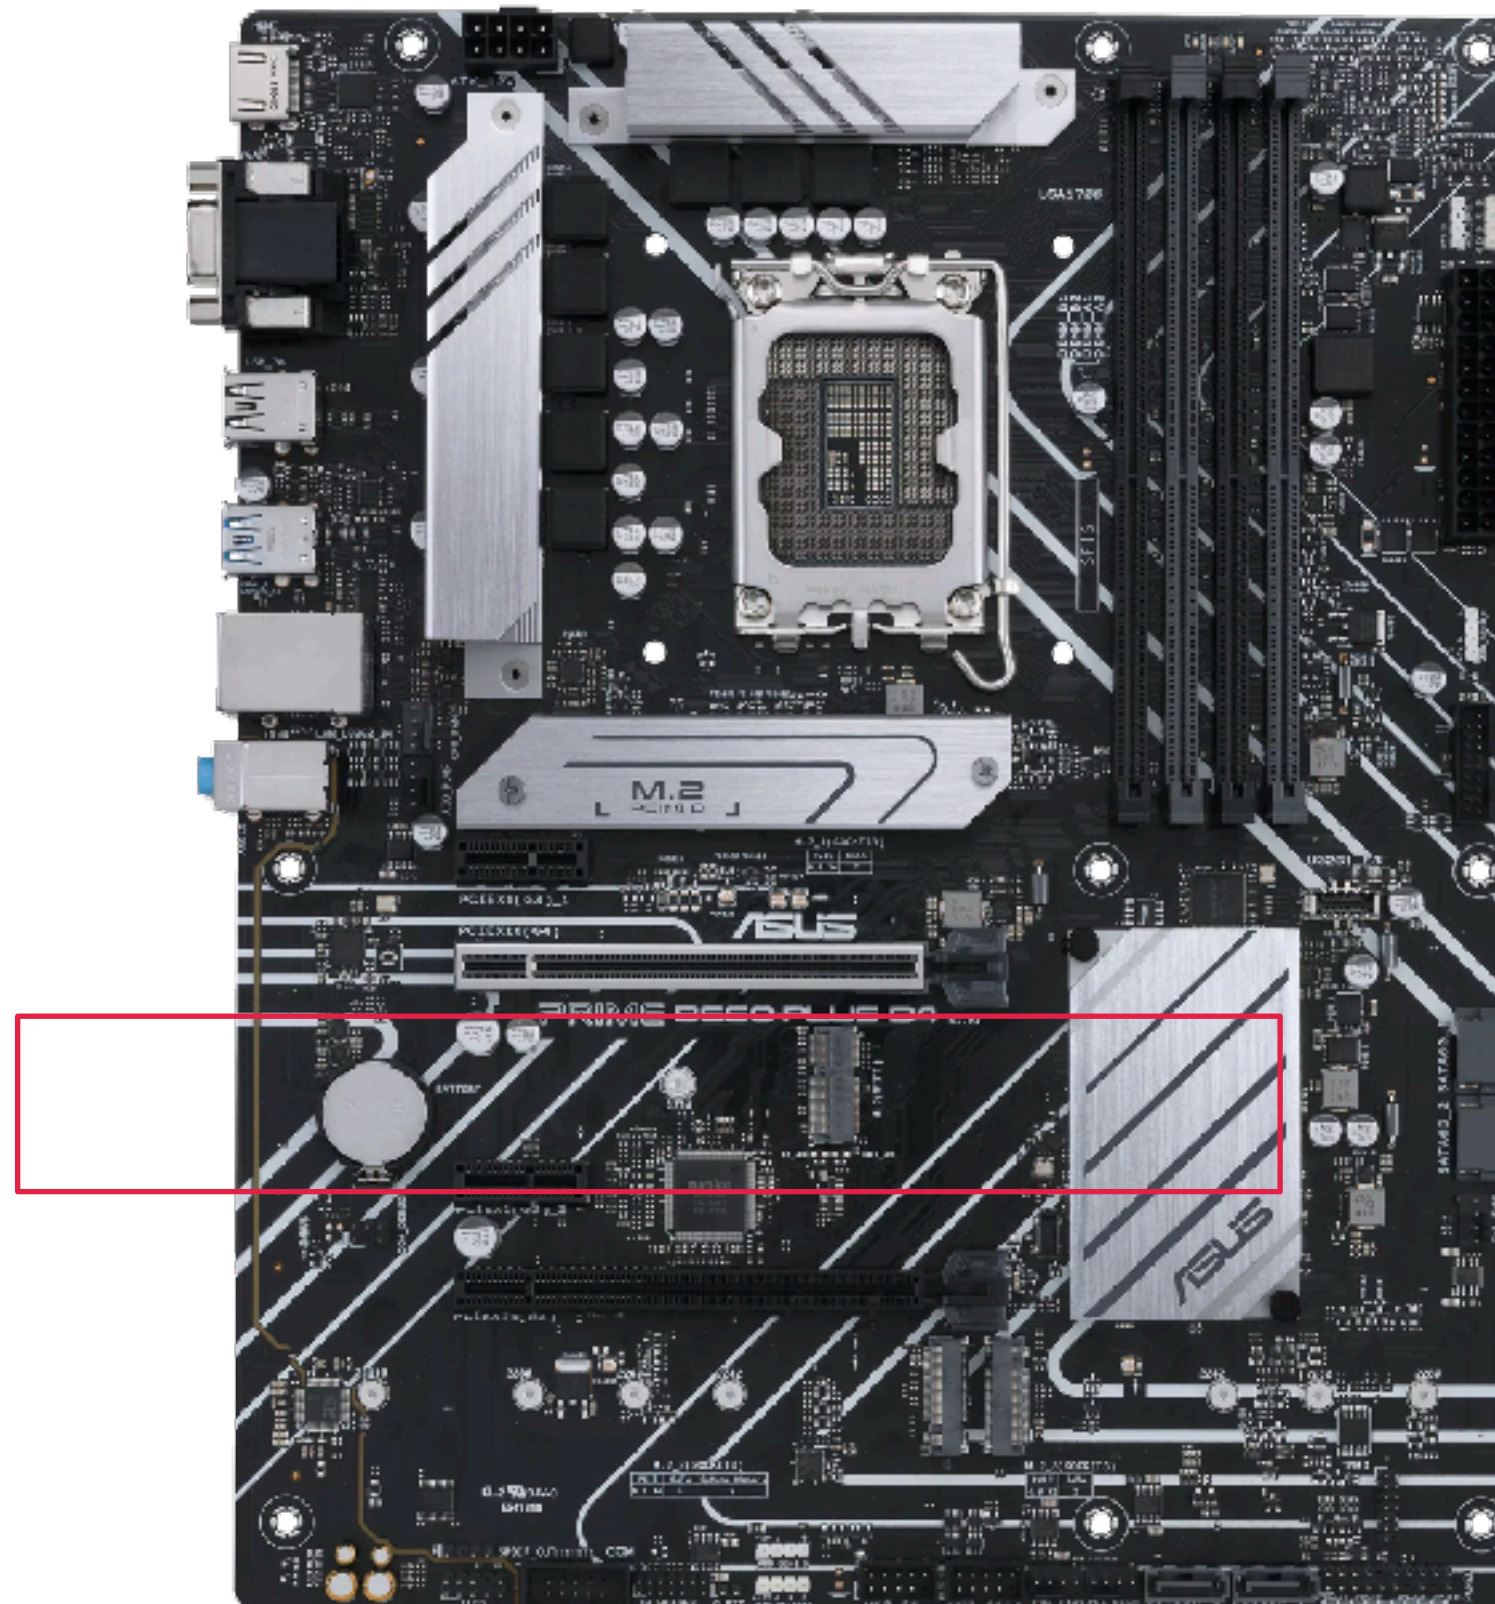
<https://www.asus.com/motherboards-components/motherboards/prime/prime-b660-plus-d4/>

Real Stuff



Credit to Apple. Apple M2 Max, with 96GB unified RAM


Real Stuff



Real Stuff



SSD vs. HDD



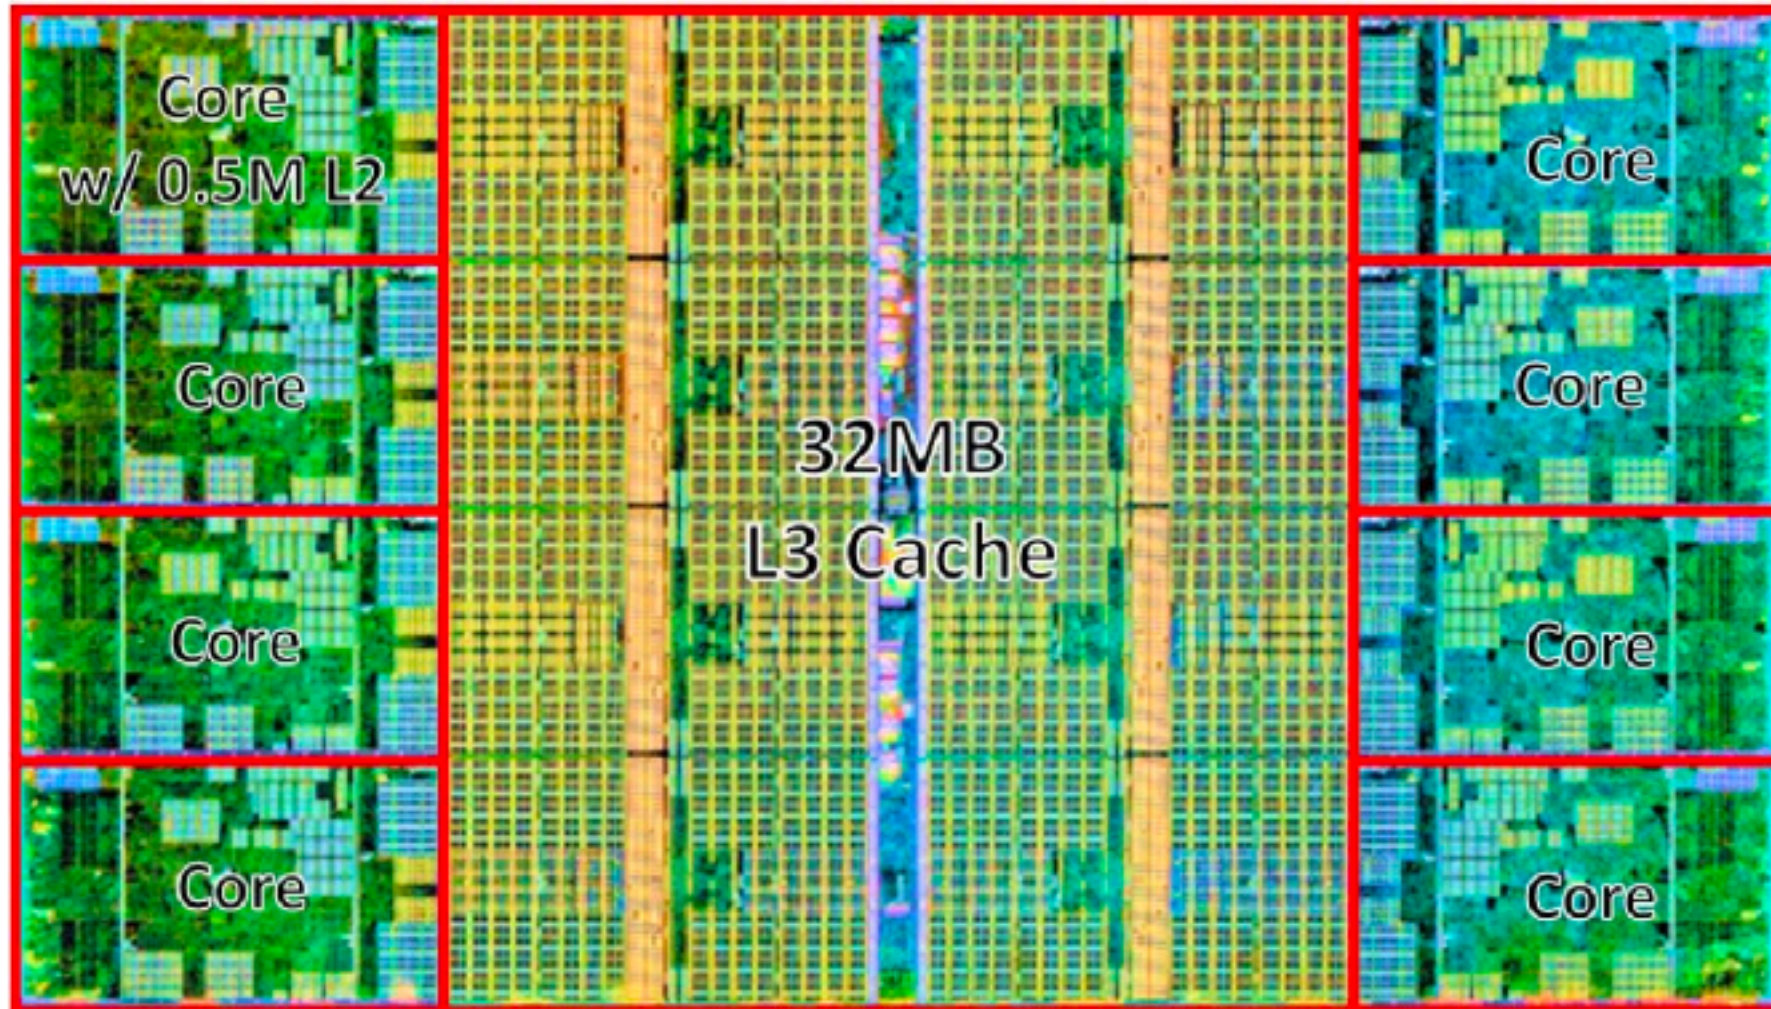
From Seagate



<https://www.samsung.com/us/computing/memory-storage/solid-state-drives/ssd-970-pro-nvme-m2-512gb-mz-v7p512bw/>
<https://www.seagate.com/in/en/products/hard-drives/barracuda-hard-drive/>

Real Stuff— Inside a CPU

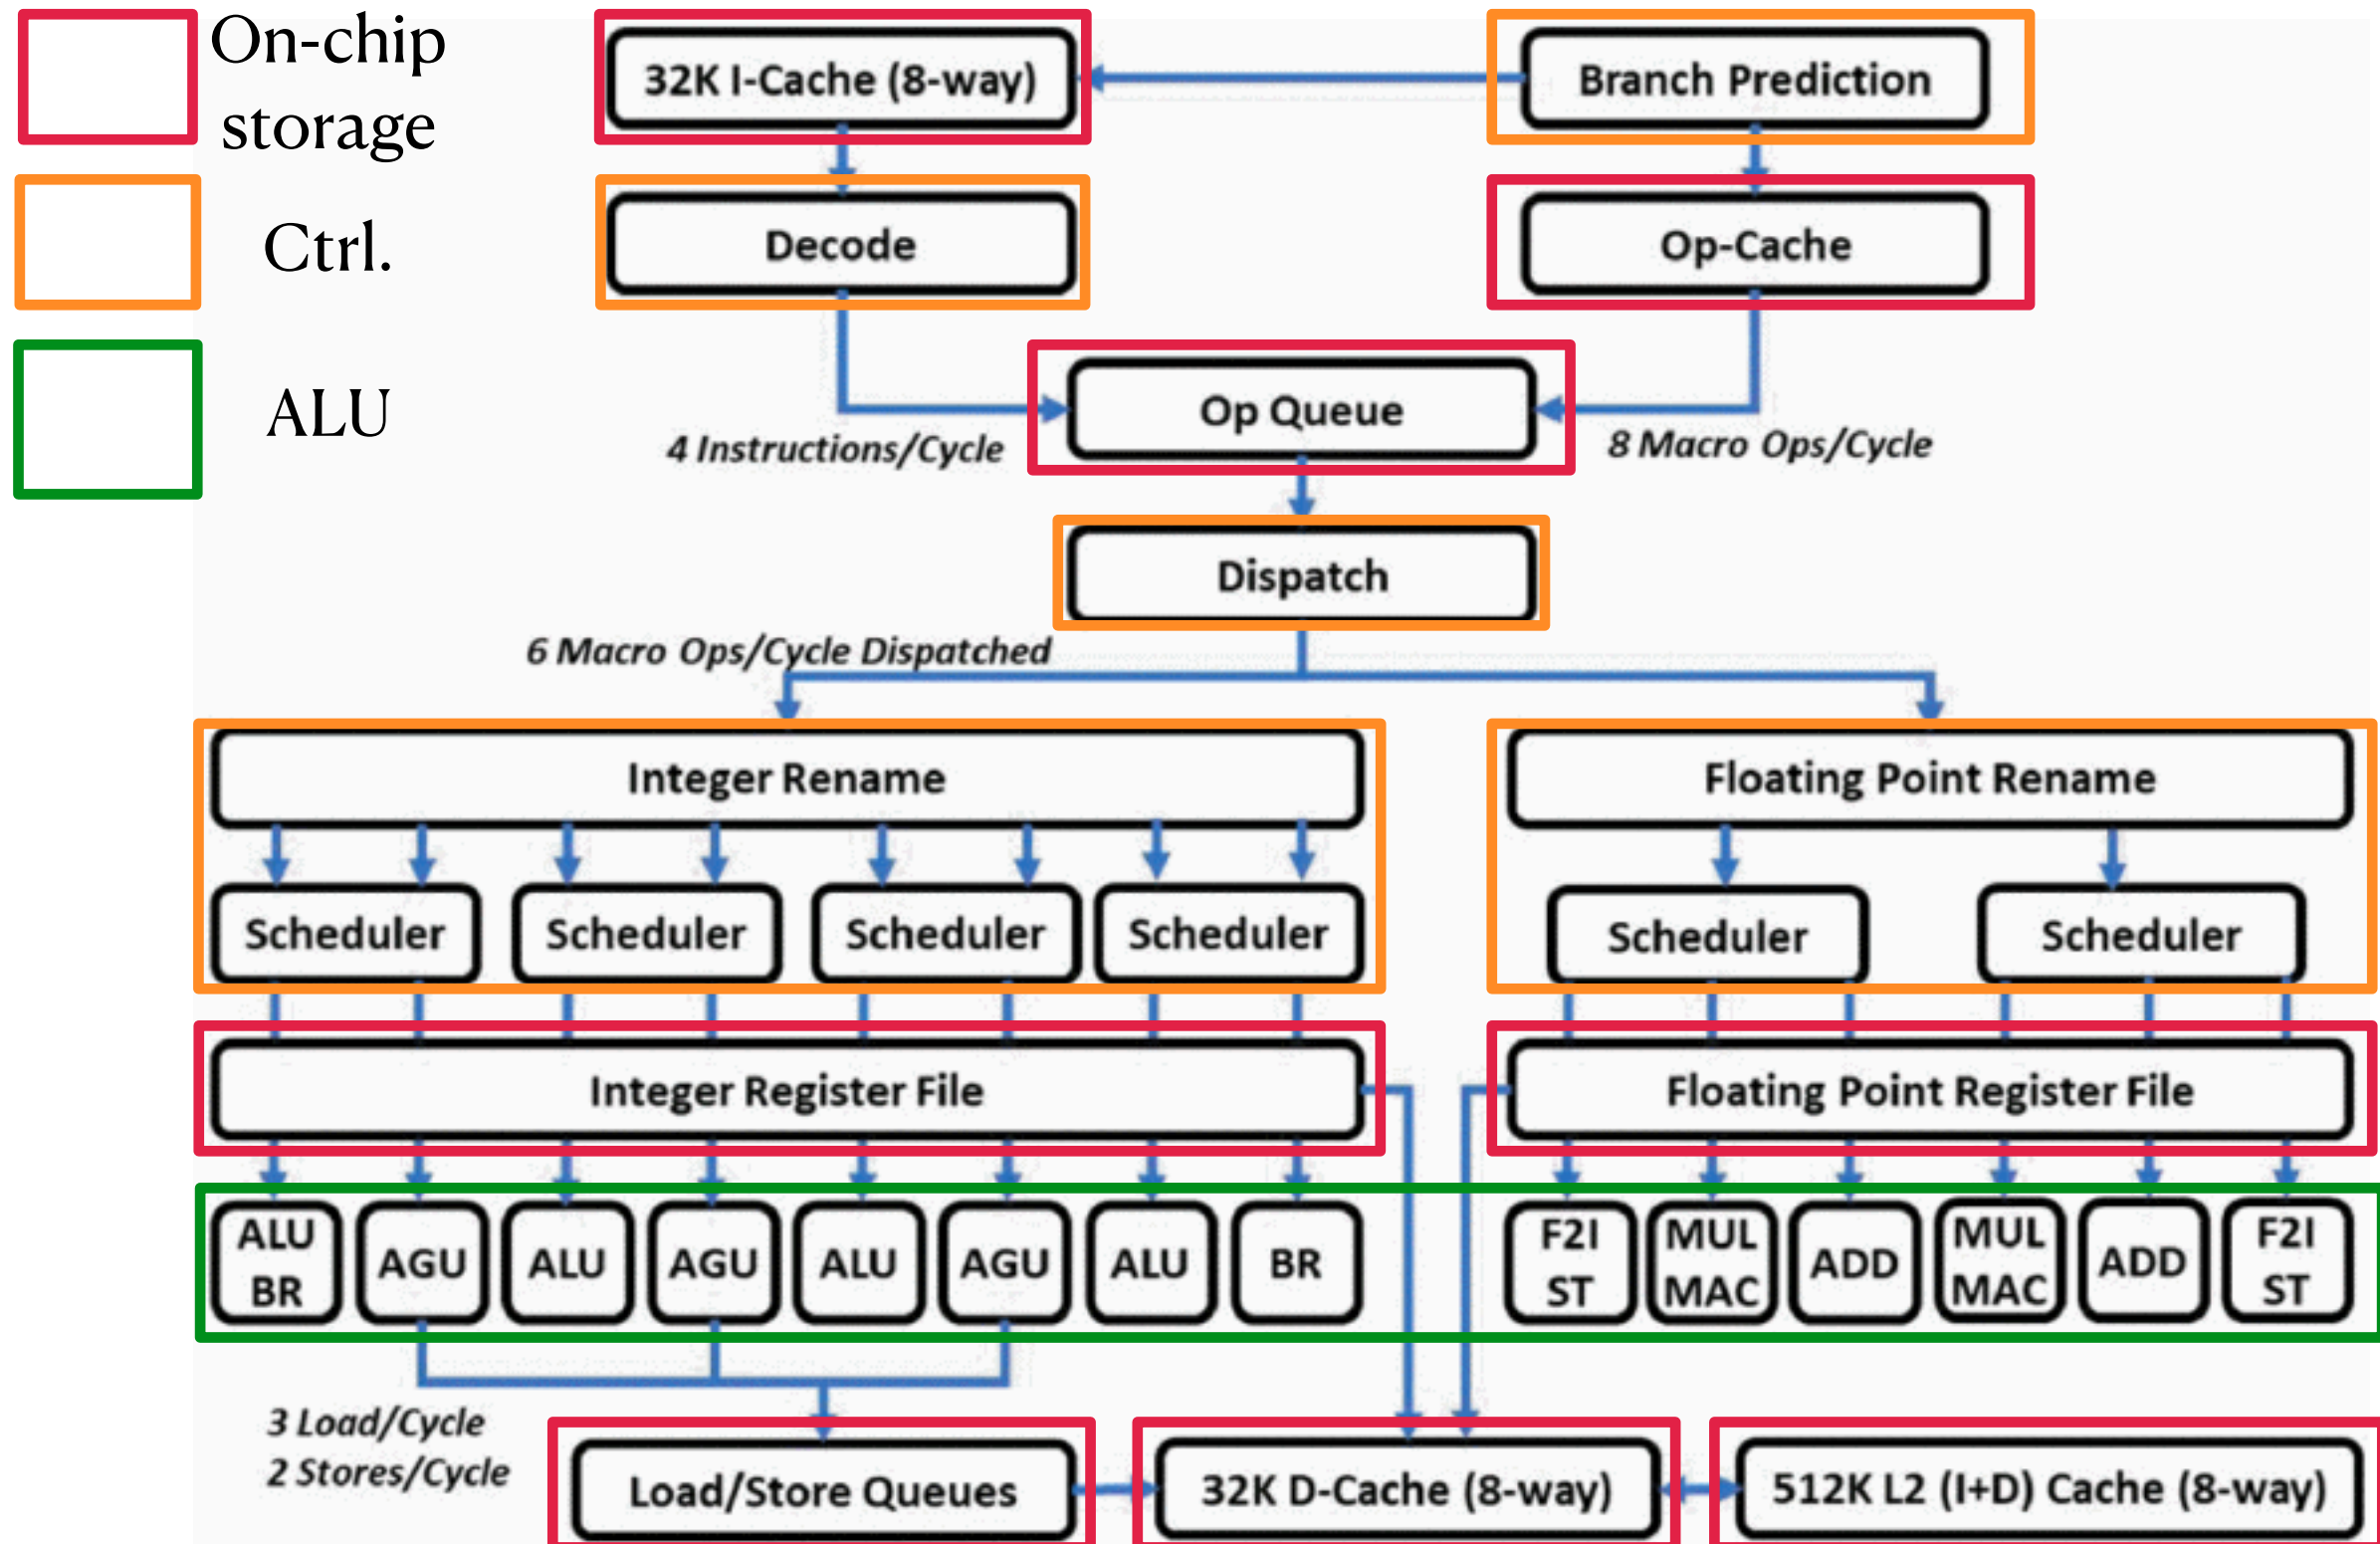
AMD Zen 3 8-core CPU, 7 nm process, 4.08B transistors in 68 mm²



T. Burd *et al.*, "Zen3: The AMD 2nd-Generation 7nm x86-64 Microprocessor Core," *2022 IEEE International Solid- State Circuits Conference (ISSCC)*, San Francisco, CA, USA, 2022, pp. 1-3.

Real Stuff— Inside a CPU

AMD Zen 3, 7 nm process, a single core



Back to C

- Typical C program

The diagram shows a typical C program with several annotations and arrows pointing to specific parts of the code:

- Comments**: Points to the first line of the program: `// Created by Siting Liu on 2023/2/5.`
- Preprocessing elements (header/macro)**: Points to the `#include <stdio.h>` line.
- Variables**: Points to the `const char * argv[]` part of the `main` function signature.
- Functions**: Points to the `main` function signature.
- Statements**: Points to the `return 0;` statement.

```
// Created by Siting Liu on 2023/2/5.
//
#include <stdio.h>
int main(int argc, const char * argv[]) {
    // insert code here...
    printf("Hello, World!\n");
    return 0;
}
```

- Must C program start with `main()`? (RTFM)

Variables

- Typed Variables in C

```
int    variable1    = 2;  
float  variable2    = 1.618;  
char   variable3    = 'A';
```

Must declare the type of data a variable will hold;

Initialize, otherwise it holds garbage

| Type | Description | Examples |
|-----------------|--------------------------------------|------------------|
| int | integer numbers, including negatives | 0, 78, -1400 |
| unsigned int | integer numbers (no negatives) | 0, 46, 900 |
| long | larger signed integer | -6,000,000,000 |
| (un)signed char | single text character or symbol | 'a', 'D', '?' |
| float | floating point decimal numbers | 0.0, 1.618, -1.4 |
| double | greater precision/big FP number | 10E100 |

C89 standard defines a lot of “Undefined Behavior”s. It means the code may produce unpredictable behavior. It may

- Produce different results on different computers/OS;
- Produce different results among multiple runs;
- Very difficult to re-produce and debug



Integers

- Typed Variables in C

| Language | sizeof(int) |
|----------|---|
| Python | ≥ 32 bits (plain ints), infinite (long ints) |
| Java | 32 bits |
| C | Depends on computer; 16 or 32 or 64 bits |

- C: int should be integer type that target processor works with most efficiently
- Generally: $\text{sizeof(long long)} \geq \text{sizeof(long)} \geq \text{sizeof(int)} \geq \text{sizeof(short)}$
 - Also, short ≥ 16 bits, long ≥ 32 bits
 - All could be 64 bits

Integer Constants

$\text{long int} \quad \left\{ \begin{array}{l} 0 \dots 0110 \\ -0 \dots 1 \dots 0000 \end{array} \right. \Rightarrow$
 $\uparrow \uparrow$ bit 31

```

#include <stdio.h>
int main() {
    printf((6-2147483648)>(6)?"T\n":"F\n");
    printf((6-0x80000000)>(6)?"T\n":"F\n");
    return 0;
}
```

 $8 \times 16^7 = 2^{31}$

Semantics: The value of a decimal constant is computed base 10; that of an octal constant base 8; that of a hexadecimal constant base 16. The lexically first digit is the most significant.

The type of an integer constant is the first of the corresponding list in which its value can be represented. Unsuffix decimal **int, long int, unsigned long int**; unsuffix octal or hexadecimal: **int, unsigned int, long int, unsigned long int**; suffixed by the letter u or U: **unsigned int, unsigned long int**; suffixed by the letter l or L: **long int, unsigned long int**; suffixed by both the letters u or U and l or L: **unsigned long int**.

Range of each type defined in <limits.h> (INT_MAX, INT_MIN)

Consts. and Enums. in C

- Constant is assigned a typed value once in the declaration; value can't change during entire execution of program

```
const float golden_ratio = 1.618;  
const int days_in_week = 7;
```

- You can have a constant version of any of the standard C variable types
- Enums: a group of related integer constants. Ex:

```
enum cardsuit {CLUBS, DIAMONDS, HEARTS, SPADES};  
enum color {RED, GREEN, BLUE};
```

C Syntax: Variable Declarations

- All variable declarations must appear before they are used (e.g., at the beginning of the block)
- A variable may be initialized in its declaration; if not, it holds garbage!
- Examples of declarations:
 - **Correct:**

```
{  
    int a = 0, b = 10;  
    ...  
}
```
 - **Incorrect:**

```
for (int i = 0; i < 10; i++)  
    }
```

Newer C standards are more flexible about this...

C Syntax: True or False

- What evaluates to FALSE in C?
 - 0 (integer)
 - NULL (a special kind of pointer: more on this later)
- No explicit Boolean type
- What evaluates to TRUE in C?
 - Anything that isn't false is true
 - Same idea as in Python: only 0 or empty sequences are false, anything else is true!

C operators

- arithmetic: +, -, *, /, %
- assignment: =
- augmented assignment: +=, -=, *=, /=, %=, &=, |=, ^=, <<=, >>=
- bitwise logic: ~, &, |, ^
- bitwise shifts: <<, >>
- boolean logic: !, &&, ||
- equality testing: ==, !=
- subexpression grouping: ()
- order relations: <, <=, >, >=
- increment and decrement: ++ and --
- member selection: ., ->
- conditional evaluation: ? :

Make sure you understand each operator!

Typed C Functions

- You need to declare the return type of a function when you declare it (plus the types of any arguments)
- You also need to declare functions before they are used

- Usually in a separate header file, e.g.

```
int number_of_people();  
float dollars_and_cents();  
int sum(int x, int y);
```

```
int number_of_people()  
{ return 3; }
```

- `void` type means “returns nothing”

```
float dollars_and_cents ()  
{ return 10.33; }
```

```
int sum (int x, int y)  
{ return x + y; }
```

Summary

- C preprocessing
- How does C work
- Basic C elements
- Variables and functions