

Course Info

- Lab 4 will be released after class (10 a.m.), get yourself prepared before going to lab sessions!
- Project 1.1 available, and will be marked in lab sessions. Deadline March 13th.
- HW3 this week, keep an eye on piazza.
- Discussion next week on CALL



信息科学与技术学院

School of Information Science and Technology

CS 110

Computer Architecture

CALL

Instructors:

Siting Liu & Chundong Wang

Course website: [https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/
Spring-2023/index.html](https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html)

School of Information Science and Technology (SIST)

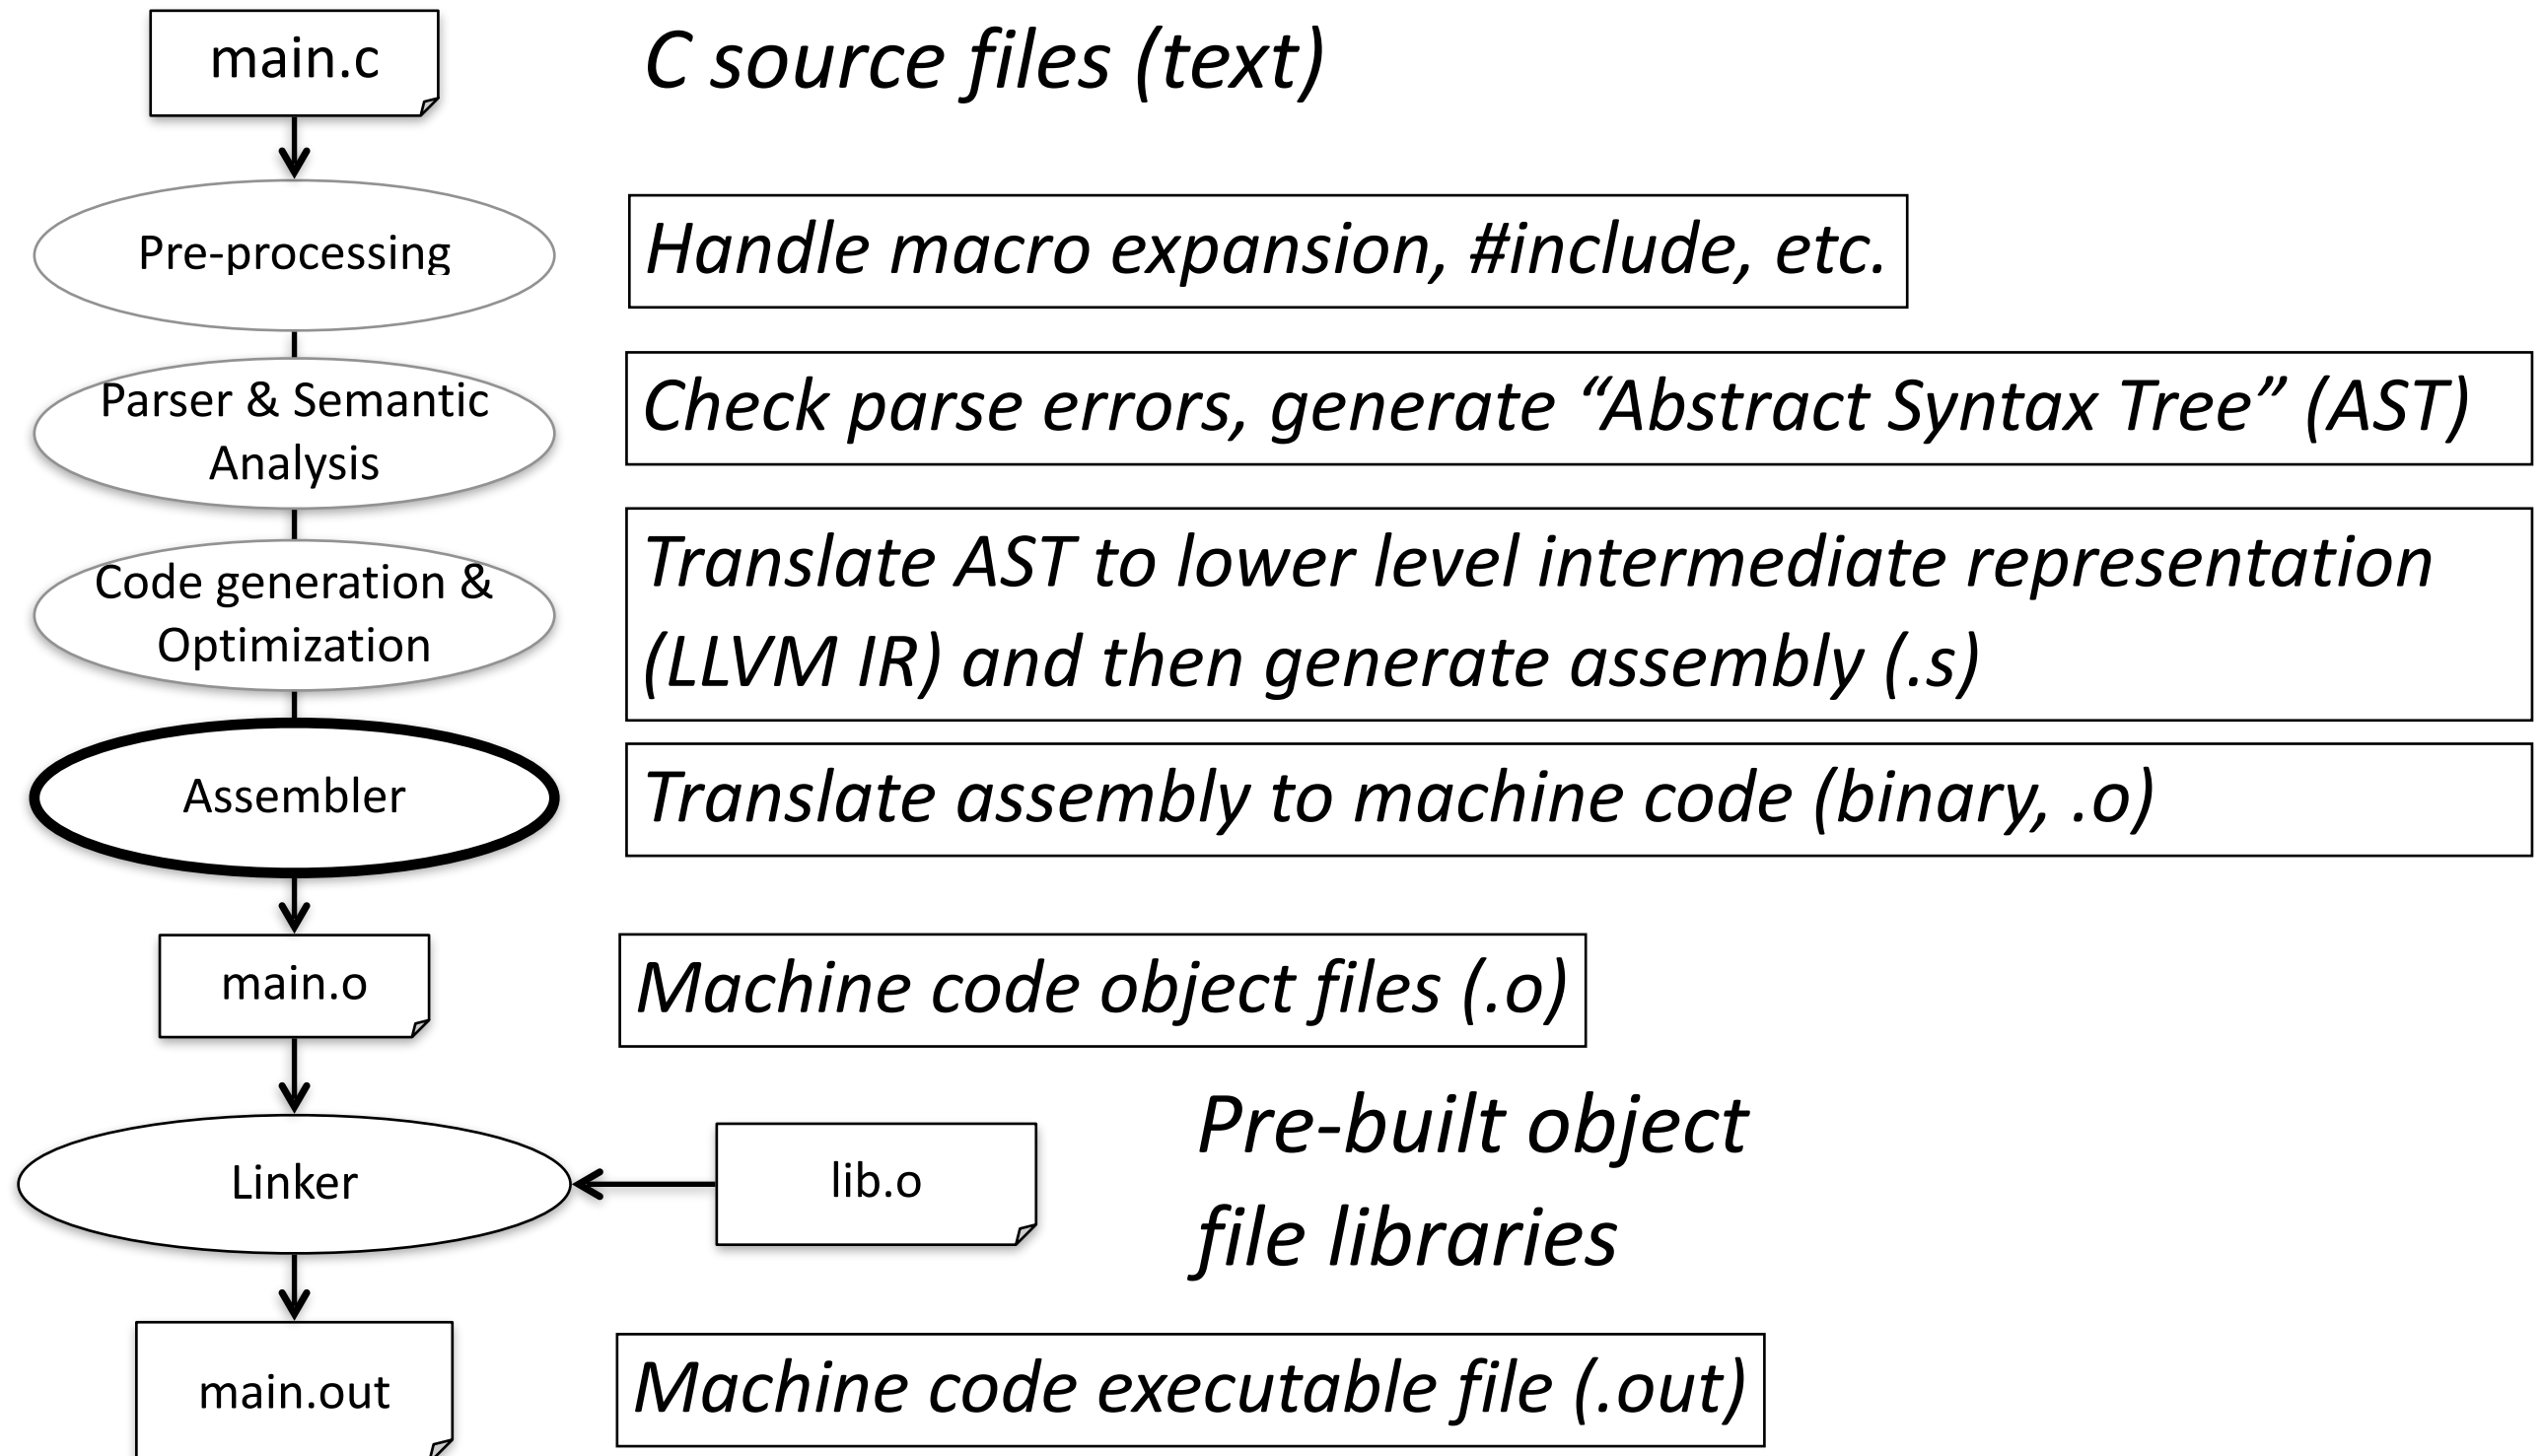
ShanghaiTech University

2023/2/6

CALL

Compiler
Assembler
Linker
Loader

C Compilation Simplified Overview



CALL

Compiler

Assembler

Linker

Loader

Assembler 汇编器

- Input: assembly language code (generated by compiler, usually contains pseudo-instructions)
- Output: object code, information tables
- Reads and uses directives 指令
- Replace pseudo-instructions
- Produce machine language
- Creates object file

Directives 指令

- Give directions to assembler, but do not produce machine instructions
 - `.text`: Subsequent items put in user text segment (instructions)
 - `.data`: Subsequent items put in user data segment (binary rep of data in source file)
 - `.globl sym`: declares sym global and can be referenced from other files
 - `.ascii str`: Store the string str in memory and null-terminate it
 - `.word w1...wn`: Store the n 32-bit quantities in successive memory words

Pseudo-instruction Examples

Assembler	
Pseudo	Real
nop	addi x0, x0, 0
not rd, rs	xori rd, rs, -1
beqz rs, offset	beq rs, x0, offset
bgt rs1, rs2, offset	blt rs2, rs1, offset
j offset	jal x0, offset
ret	jalr x0, x1, offset
call offset (too big to jal)	auipc x6, offset[31:12] jalr x1, x6, offset[11:0]
tail offset (too far to j)	auipc x6, offset[31:12] jalr x0, x6, offset[11:0]
li/la rd imm/label	lui rd <hi20bits> (too large) addi rd, x0, <low12bits>
mv rs1, rs2	addi rs1, rs2, 0

Tail

- Nested procedures

a₁
↑
x11

a₀
↑
x10

caller

```
int fact (int n, int prod) {  
    if (n>1) return fact(n-1, prod*n);  
    else return (prod);  
}
```

```
fact: addi    t0,x0,1  
      ble     x11,t0,Exit  
      mul     x10,x10,x11  
      addi    x11,x11,-1  
      jalr    x0,fact  
Exit: addi    x10,x0,x10  
      jalr    x0,0(x1)
```

Producing Machine Language (1/3)

- Simple Case
 - Arithmetic, Logical, Shifts, and so on
 - All necessary info is within the instruction already
- What about Branches?
 - PC-Relative (e.g., `beq/bne` and `jal`), position-independent code (PIC), within one file
 - So once pseudo-instructions are replaced by real ones, we know by how many instructions to branch

Producing Machine Language (2/3)

- “Forward Reference” problem
 - Branch instructions can refer to labels that are “forward” in the program:

```

3 instructions forward →
L1: addi t2, zero, 9      # t2 = 9
    slt  t1, zero, t2    # 0 < t2? Set t1
    beq  t1, zero, L2     # NO! t2 <= 0; Go to L2
    addi t2, t2, -1      # YES! t2 > 0; t2--
3 instructions back →
                        j  L1      # Go to L1
L2:
```

- Solved by taking two passes over the program
 - First pass remembers position of labels (symbol table)
 - Second pass uses label positions to generate code

Producing Machine Language (3/3)

- What about jumps (j, jal)?
 - Jumps within a file are PC relative (and we can easily compute):
 - Just count the number of instructions between target and jump to determine the offset: position-independent code (PIC)
 - Jumps to other files we can't
- What about references to static data/external functions/multiple files?
 - la gets broken up into lui and addi
 - These require the full 32-bit address of the data
- These can't be determined yet, so we create two tables ...

Symbol Table

- List of “items” in this file that may be used by other files
- What are they?
 - Labels: function calling; `.global` directive
 - Data: anything in the `.data` section; variables which may be accessed across files

Relocation Table

- List of “items” whose (absolute) address this file needs later.
What are they?
 - Any external label jumped to: `jal`, `jalr`
 - External (including lib files)
 - Any piece of data in static section
 - Such as the `la` instruction
E.g., for `lw/sw` base register

Summary: Object File Format

- object file header: size and position of the other pieces of the object file
- text segment: the machine code
- data segment: binary representation of the static data in the source file
- relocation information: identifies lines of code that need to be fixed up later (by linker)
- symbol table: list of this file's labels and static data that can be referenced
- debugging information
- A standard format is ELF (except MS)
http://www.skyfree.org/linux/references/ELF_Format.pdf

CALL

Compiler

Assembler

Linker


Loader

Linker (1/3)


- Input: Object code files, information tables (e.g., `<your C code>.o`, `libc.o` for RISC-V)
- Output: Executable code (e.g., `a.out` for RISC-V)
- Combines several object (.o) files into a single executable (“linking”)
- Enable separate compilation of files
 - Changes to one file do not require recompilation of the whole program
 - Linux source > 20 M lines of code!
 - Old name “Link Editor” from editing the “links” in jump and link instructions

Linker (2/3)


.o file 1



.o file 2



a.out



Linker (3/3)

- Step 1: Take text segment from each .o file and put them together; Take data segment from each .o file, put them together, and concatenate this onto end of text segments
- Step 2: Determine the addresses of data and instruction labels
- Step 3: Resolve references
 - Go through Relocation Table; handle each entry
 - That is, fill in all absolute addresses

Three Types of Addresses

- PC-Relative Addressing (`beq`, `bne`, `jal`)
 - Never need to relocate (PIC: position independent code)
- External Function Reference (usually `jal`)
 - Always relocate
- Static Data Reference (often `auipc`/`addi`)
 - Always relocate
 - RISC-V often uses `auipc` rather than `lui` so that a big block of stuff can be further relocated as long as it is fixed relative to the `pc`

Absolute Addresses in RISC-V

- Which instructions need relocation editing?
 - J-format: jump and link: ONLY for external jumps

xxxxxx	rd	jal
---------------	-----------	------------

- I-,S- Format: Loads and stores to variables in static area, relative to global pointer

xxx	gp		rd	lw	
xx	rs1	gp		x	sw

- What about conditional branches?

xx	rs1	rs2		x	beq bne
-----------	------------	------------	--	----------	--------------------------

- PC-relative addressing preserved even if code moves

Resolving References (1/2)

- Linker knows:
 - Length of each text and data segment
 - Ordering of text and data segments
- Linker calculates:
 - Absolute address of each label to be jumped to and each piece of data being referenced

Resolving References (2/2)

- To resolve references:
 - search for reference (data or label) in all “user” symbol tables
 - if not found, search library files
(for example, `printf`, `malloc`)
 - once absolute address is determined, fill in the machine code appropriately
- Output of linker: executable file containing text and data (plus header)

Static vs. Dynamic Linking

- What we've described is the traditional way: statically-linked approach
 - The library is now part of the executable, so if the library updates, we don't get the fix (have to recompile if we have source)
 - It includes the entire library even if not all of it will be used
 - Executable is self-contained
- An alternative is dynamically linked libraries (DLL), common on Windows (.dll) & UNIX (.so) & MacOS (.dylib) platforms

Dynamically linked libraries

- Space/time issues
 - + Storing a program requires less disk space
 - + Sending a program requires less time
 - + Executing two programs requires less memory (if they share a library)
 - – At runtime, there's time overhead to do link
- Upgrades
 - + Replacing one file (libXYZ.so) upgrades every program that uses library "XYZ"
 - – Having the executable isn't enough anymore
 - Thus "containers": We hate dependencies, so we are just going to ship around all the libraries and everything else as part of the 'application'

Overall, dynamic linking adds quite a bit of complexity to the compiler, linker, and operating system. However, it provides many benefits that often outweigh these

CALL

Compiler
Assembler
Linker
Loader

Loader Basics

- Input: Executable Code
(e.g., a.out for RISC-V)
- Output: (program run)
- Executable files are stored on disk
- When one is run, loader's job is to load it into memory and start it running
- In reality, loader is the operating system (OS)
 - loading is one of the OS tasks

Loader ... what does it do?

- Reads executable file's header to determine size of text and data segments
- Creates new address space for program large enough to hold text and data segments, along with a stack segment
- Copies instructions and data from executable file into the new address space
- Copies arguments passed to the program onto the stack
- Initializes machine registers
 - Most registers cleared, but stack pointer assigned address of 1st free stack location
- Jumps to start-up routine that copies program's arguments from stack to registers & sets the PC
 - If main routine returns, start-up routine terminates program with the exit system call

Question

At what point in process are all the machine code bits generated for the following assembly instructions:

1) `add x6, x7, x8`

2) `jal x1, fprintf` → linker 时才找到

A: 1) & 2) After compilation

B: 1) After compilation, 2) After assembly

☒ C: 1) After assembly, 2) After linking

D: 1) After assembly, 2) After loading

E: 1) After compilation, 2) After linking

Answer

At what point in process are all the machine code bits determined for the following assembly instructions:

1) `add x6, x7, x8`

2) `jal x1, fprintf`

C: (1) After assembly, (2) After linking

Example

```
#include <stdio.h>
int main()
{
    printf("Hello, %s\n", "world");
    return 0;
}
```

C -> assembly -> obj. -> exe.

Example

```
main:
    addi    sp, sp, -16
    Store double sd    ra, 8(sp)
    sd      s0, 0(sp)
    addi    s0, sp, 16
    lui     a5, %hi(.LC0)
    addi    a1, a5, %lo(.LC0)
    lui     a5, %hi(.LC1)
    addi    a0, a5, %lo(.LC1)
    call    printf
    li      a5, 0
    mv      a0, a5
    ld      ra, 8(sp)
    ld      s0, 0(sp)
    addi    sp, sp, 16
    jr      ra
    .size   main, .-main
    .ident  "GCC: (g2ee5e430018-dirty)"

.attribute stack_align, 16 word
.text
.section .rodata
.align 3
.LC0:
.string    "world"
.align 3
.LC1:
.string    "Hello, %s\n"
.text
.align 1
.globl main
.type main, @function
```

C -> assembly -> obj. -> exe.

Example

000000000000000000 <main>:

```
0: 1141          add sp,sp,-16
2: e406          sd  ra,8(sp)
4: e022          sd  s0,0(sp)
6: 0800          add s0,sp,16
8: 000007b7      lui  a5,0x0
c: 00078593      mv   a1,a5
10: 000007b7     lui  a5,0x0
14: 00078513     mv   a0,a5
18: 00000097     auipc ra,0x0
1c: 000080e7     jalr  ra,#18 <main+0x18>
20: 4781          li   a5,0
22: 853e          mv   a0,a5
24: 60a2          ld   ra,8(sp)
26: 6402          ld   s0,0(sp)
28: 0141          add  sp,sp,16
2a: 8082          ret
```

RVC included

Address
placeholder

gcc -c

C -> assembly -> obj. -> exe.

Example

objdump -r hello.o

RELOCATION RECORDS FOR [.text]:

OFFSET	TYPE		VALUE
000000000000000008	R_RISC	V_HI20	.LC0
000000000000000008	R_RISC	V_RELAX	*ABS*
00000000000000000c	R_RISC	V_LO12_I	.LC0
00000000000000000c	R_RISC	V_RELAX	*ABS*
000000000000000010	R_RISC	V_HI20	.LC1
000000000000000010	R_RISC	V_RELAX	*ABS*
000000000000000014	R_RISC	V_LO12_I	.LC1
000000000000000014	R_RISC	V_RELAX	*ABS*
000000000000000018	R_RISC	V_CALL_PLT	printf
000000000000000018	R_RISC	V_RELAX	*ABS*

C -> assembly -> obj. -> exe.

Example

00000000000101ac <main>:


101ac: 1141	add sp,sp,-16
101ae: e406	sd ra,8(sp)
101b0: e022	sd s0,0(sp)
101b2: 0800	add s0,sp,16
101b4: 67f5	lui a5,0x1d
101b6: a4078593	add a1,a5,-1472 # 1ca40 <__clzdi2+0x46>
101ba: 67f5	lui a5,0x1d
101bc: a4878513	add a0,a5,-1464 # 1ca48 <__clzdi2+0x4e>
101c0: 146000ef	jal 10306 <printf>
101c4: 4781	li a5,0
101c6: 853e	mv a0,a5
101c8: 60a2	ld ra,8(sp)
101ca: 6402	ld s0,0(sp)
101cc: 0141	add sp,sp,16
101ce: 8082	ret

and the other libs

C -> assembly -> obj. -> exe.

In Conclusion...

- Compiler converts a single HLL file into a single assembly language file.
- Assembler removes pseudo-instructions, converts what it can to machine language, and creates a checklist for the linker (relocation table). A .s file becomes a .o file.
 - Does 2 passes to resolve addresses, handling internal forward references
- Linker combines several .o files and resolves absolute addresses.
 - Enables separate compilation, libraries that need not be compiled, and resolves remaining addresses
- Loader loads executable into memory and begins execution.



In Conclusion...

- Compiler converts a single HLL file into a single assembly language file.
- Assembler removes pseudo-instructions, converts what it can to machine language, and creates a checklist for the linker (relocation table). A .s file becomes a .o file.
 - Does 2 passes to resolve addresses, handling internal forward references
- Linker combines several .o files and resolves absolute addresses.
 - Enables separate compilation, libraries that need not be compiled, and resolves remaining addresses
- Loader loads executable into memory and begins execution.

