

Course Info

- Lab 6 next week, prepare before lab sessions! Keep an eye on piazza.
- Project 1.2 ddl March 31st. Project 2.1 released next week!
- This week discussion on ALU & FSM. Next week discussion on datapath.
- Mid-term I solution & score released. If you have questions about the solution, feel free to ask on Piazza; If you have questions regarding your marks, email the instructors **BEFORE April 2nd**. We will get back to you ASAP.
- Any regrade request after April 2nd **WOULD NOT** be considered

Course Info

- HW4 released. Submit your paper homework to the box below (at SIST 3-322). DDL April 7th.
- Remember to add your name. You have only one chance to submit and cannot be withdrawn.





信息科学与技术学院

School of Information Science and Technology

CS 110

Computer Architecture

Pipeline

Instructors:

Siting Liu & Chundong Wang

Course website: [https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/
Spring-2023/index.html](https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html)

School of Information Science and Technology (SIST)

ShanghaiTech University

2023/3/23

ISA

**temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;**

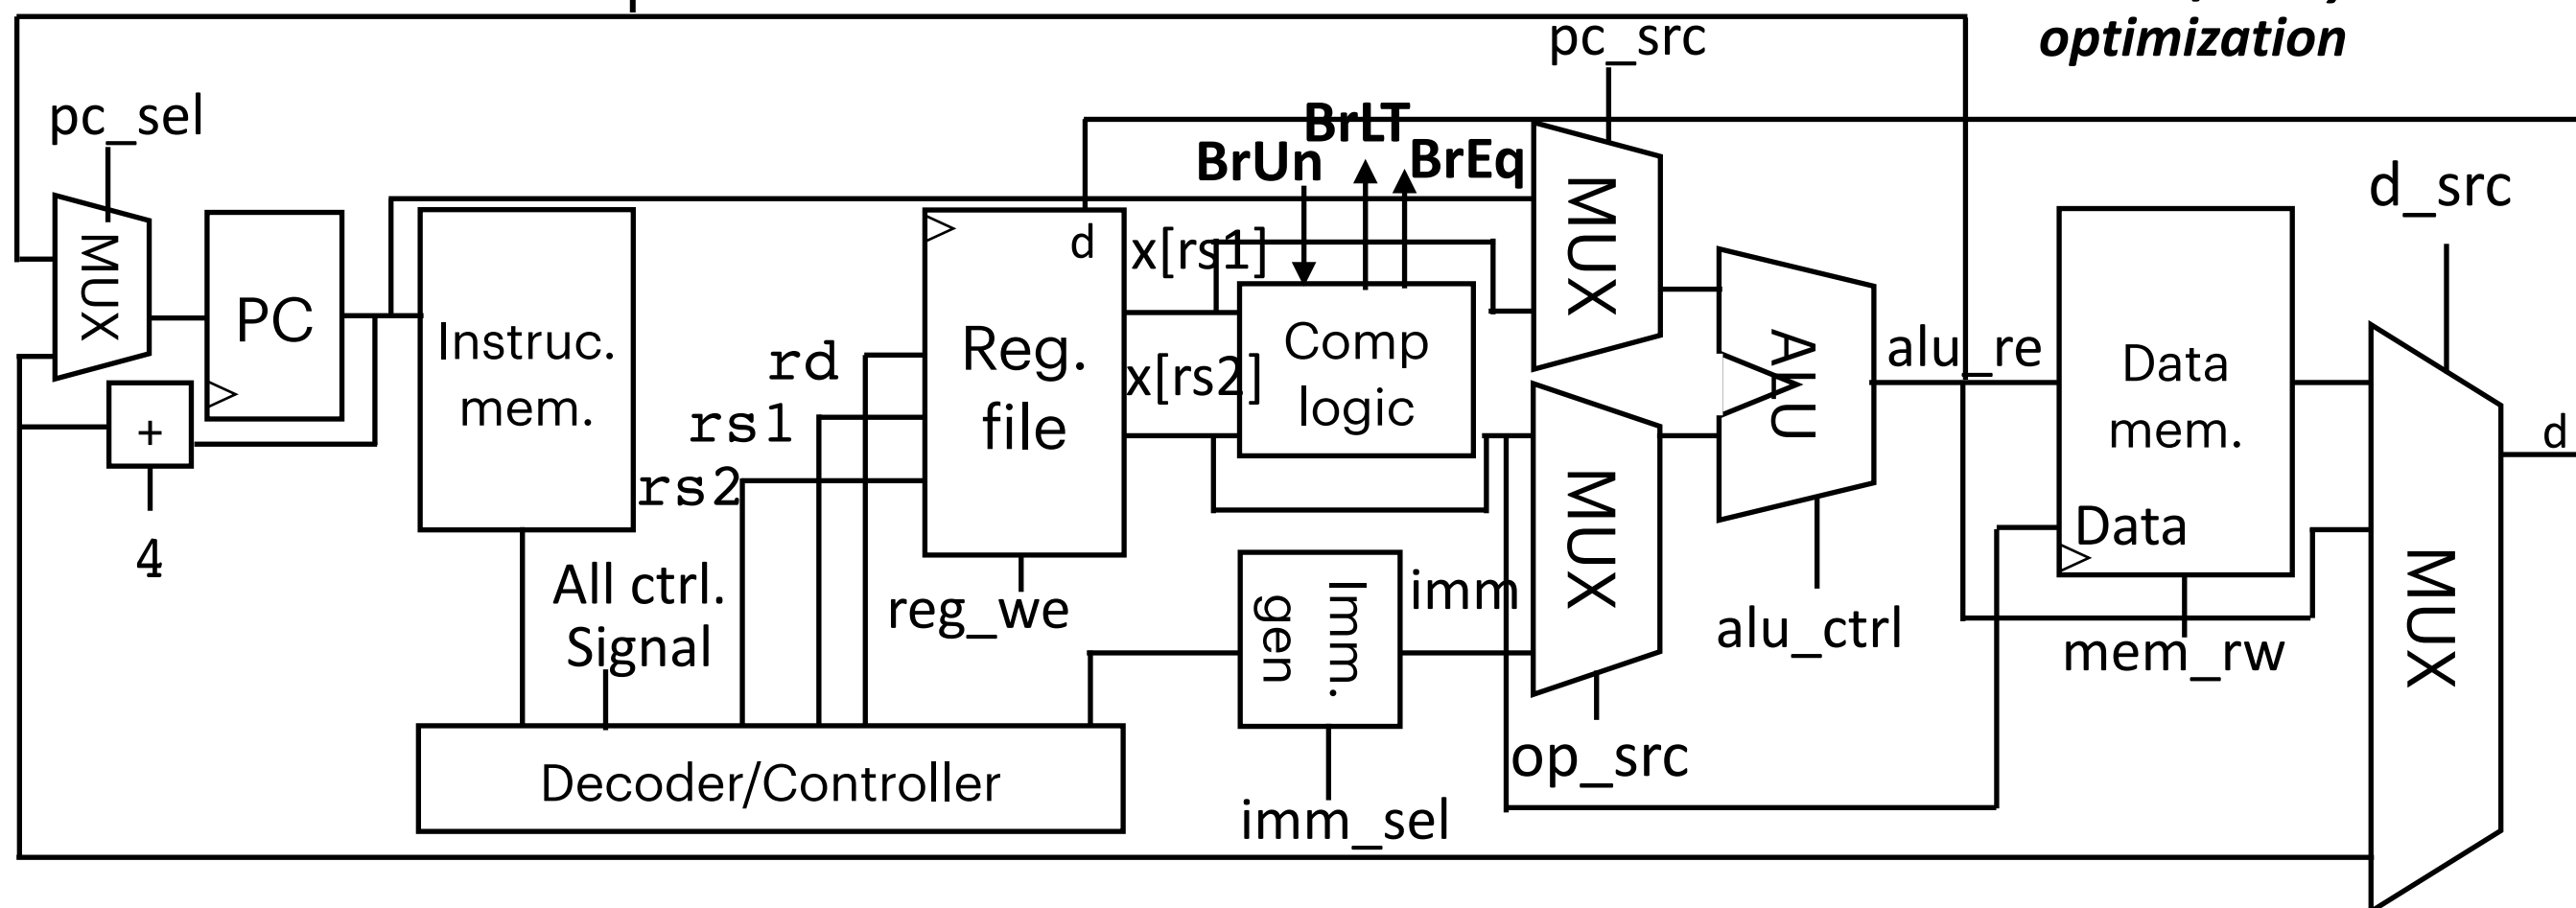
Anything can be represented
as a *number*,
i.e., data or instructions

0000 1001 1100 0110 1010 1111 0101 1000
1010 1111 0101 1000 0000 1001 1100 0110
1100 0110 1010 1111 0101 1000 0000 1001
0101 1000 0000 1001 1100 0110 1010 1111



**Architecture
Implementation**

**Single-core simple
CPU w/o any
optimization**

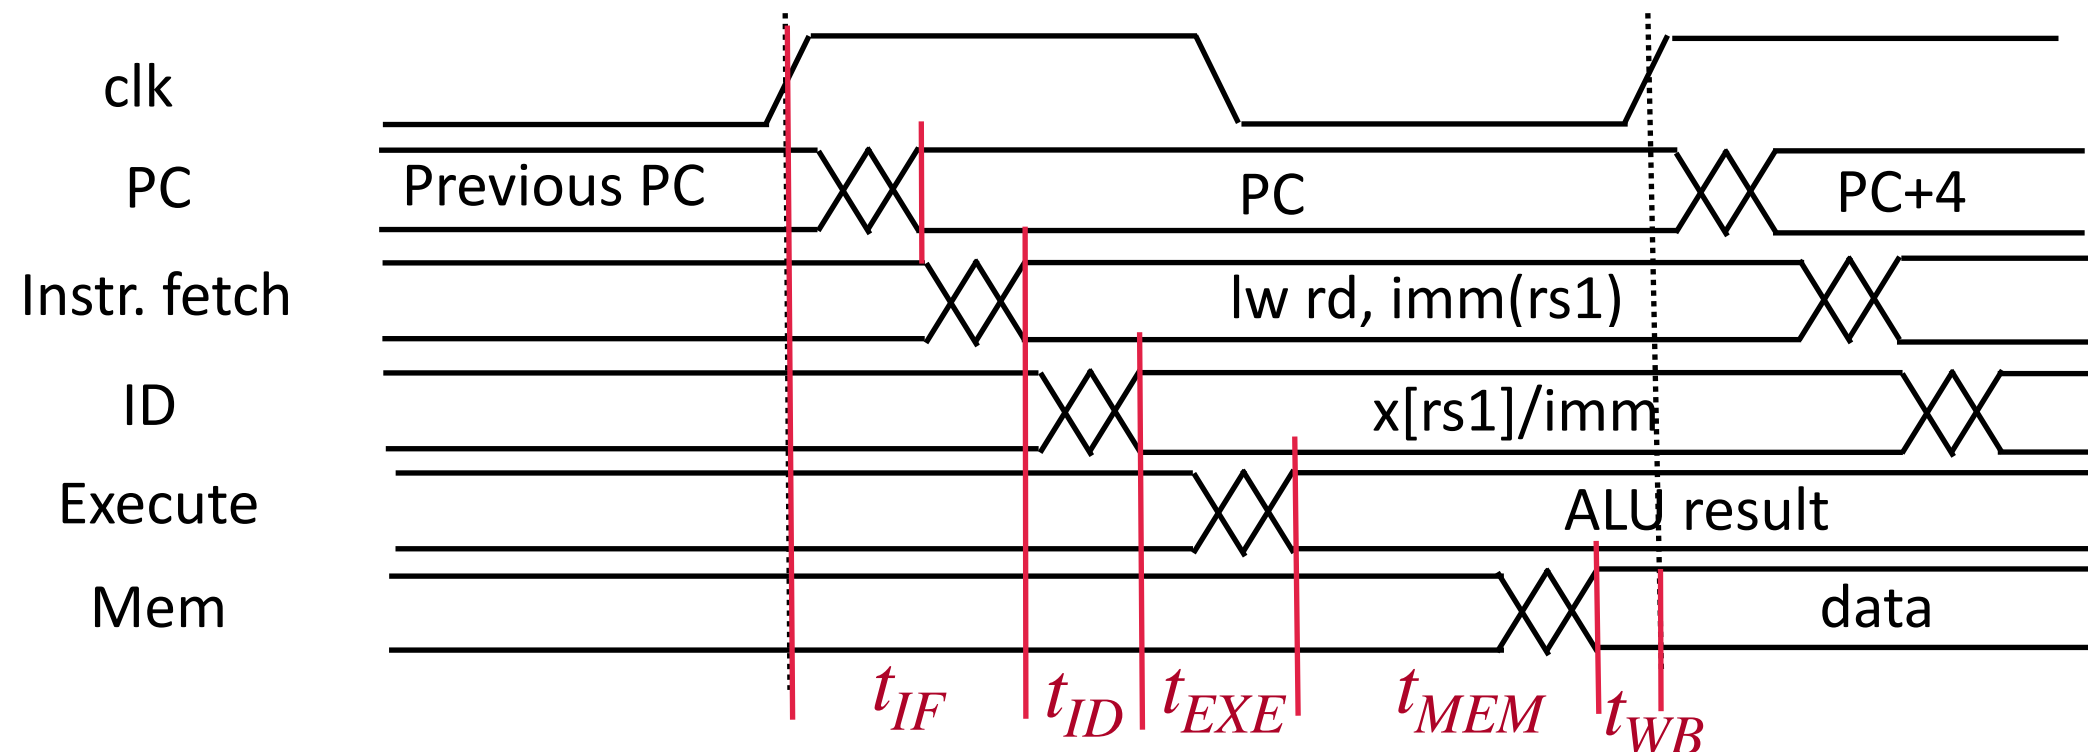


Summary

- We have built a processor!
 - Capable of executing all RISC-V instructions in one cycle each
 - Not all units (hardware) used by all instructions
 - Critical path changes for different instructions
 - $T_{clk} = \text{Max delay, } t_{IF} + t_{ID} + t_{EX} + t_{MEM} + t_{WB}$
 - Max Frequency = 1/Max delay
- 5 Phases of execution
 - IF, ID, EX, MEM, WB
 - Not all instructions are active in all phases
- Controller specifies how to execute instructions
 - Implemented as ROM or logic

Bottleneck

Instru.	IF = 300 ps	ID = 100 ps	EX = 200 ps	MEM = 300 ps	WB = 100 ps	Total
add	X	X	X		X	700 ps
beq	X	X	X			600 ps
jal	X	X	X		X	700 ps
lw	X	X	X	X	X	1000 ps
sw	X	X	X	X		900 ps



Great Ideas in Computer Architecture

- Abstraction (Layers of Representation / Interpretation)
- Moore's Law
- Principle of Locality/Memory Hierarchy
- Parallelism
- Performance Measurement and Improvement
- Dependability via Redundancy

“Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\overset{\text{指令数}}{\text{Instructions}}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}}$$

$$\frac{\text{Instructions}}{\text{Program}}$$

- ISA-relevant
- Compiler
- What task, complexity
- Etc.

load & store



RISC vs. CISC

load

alu

store

alu → 内存

“Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}}$$

$$\frac{\text{Cycles}}{\text{Instruction}} \quad \text{CPI} \quad \text{cycles per instruction}$$

- Microarchitecture implementation or circuit design
- ISA

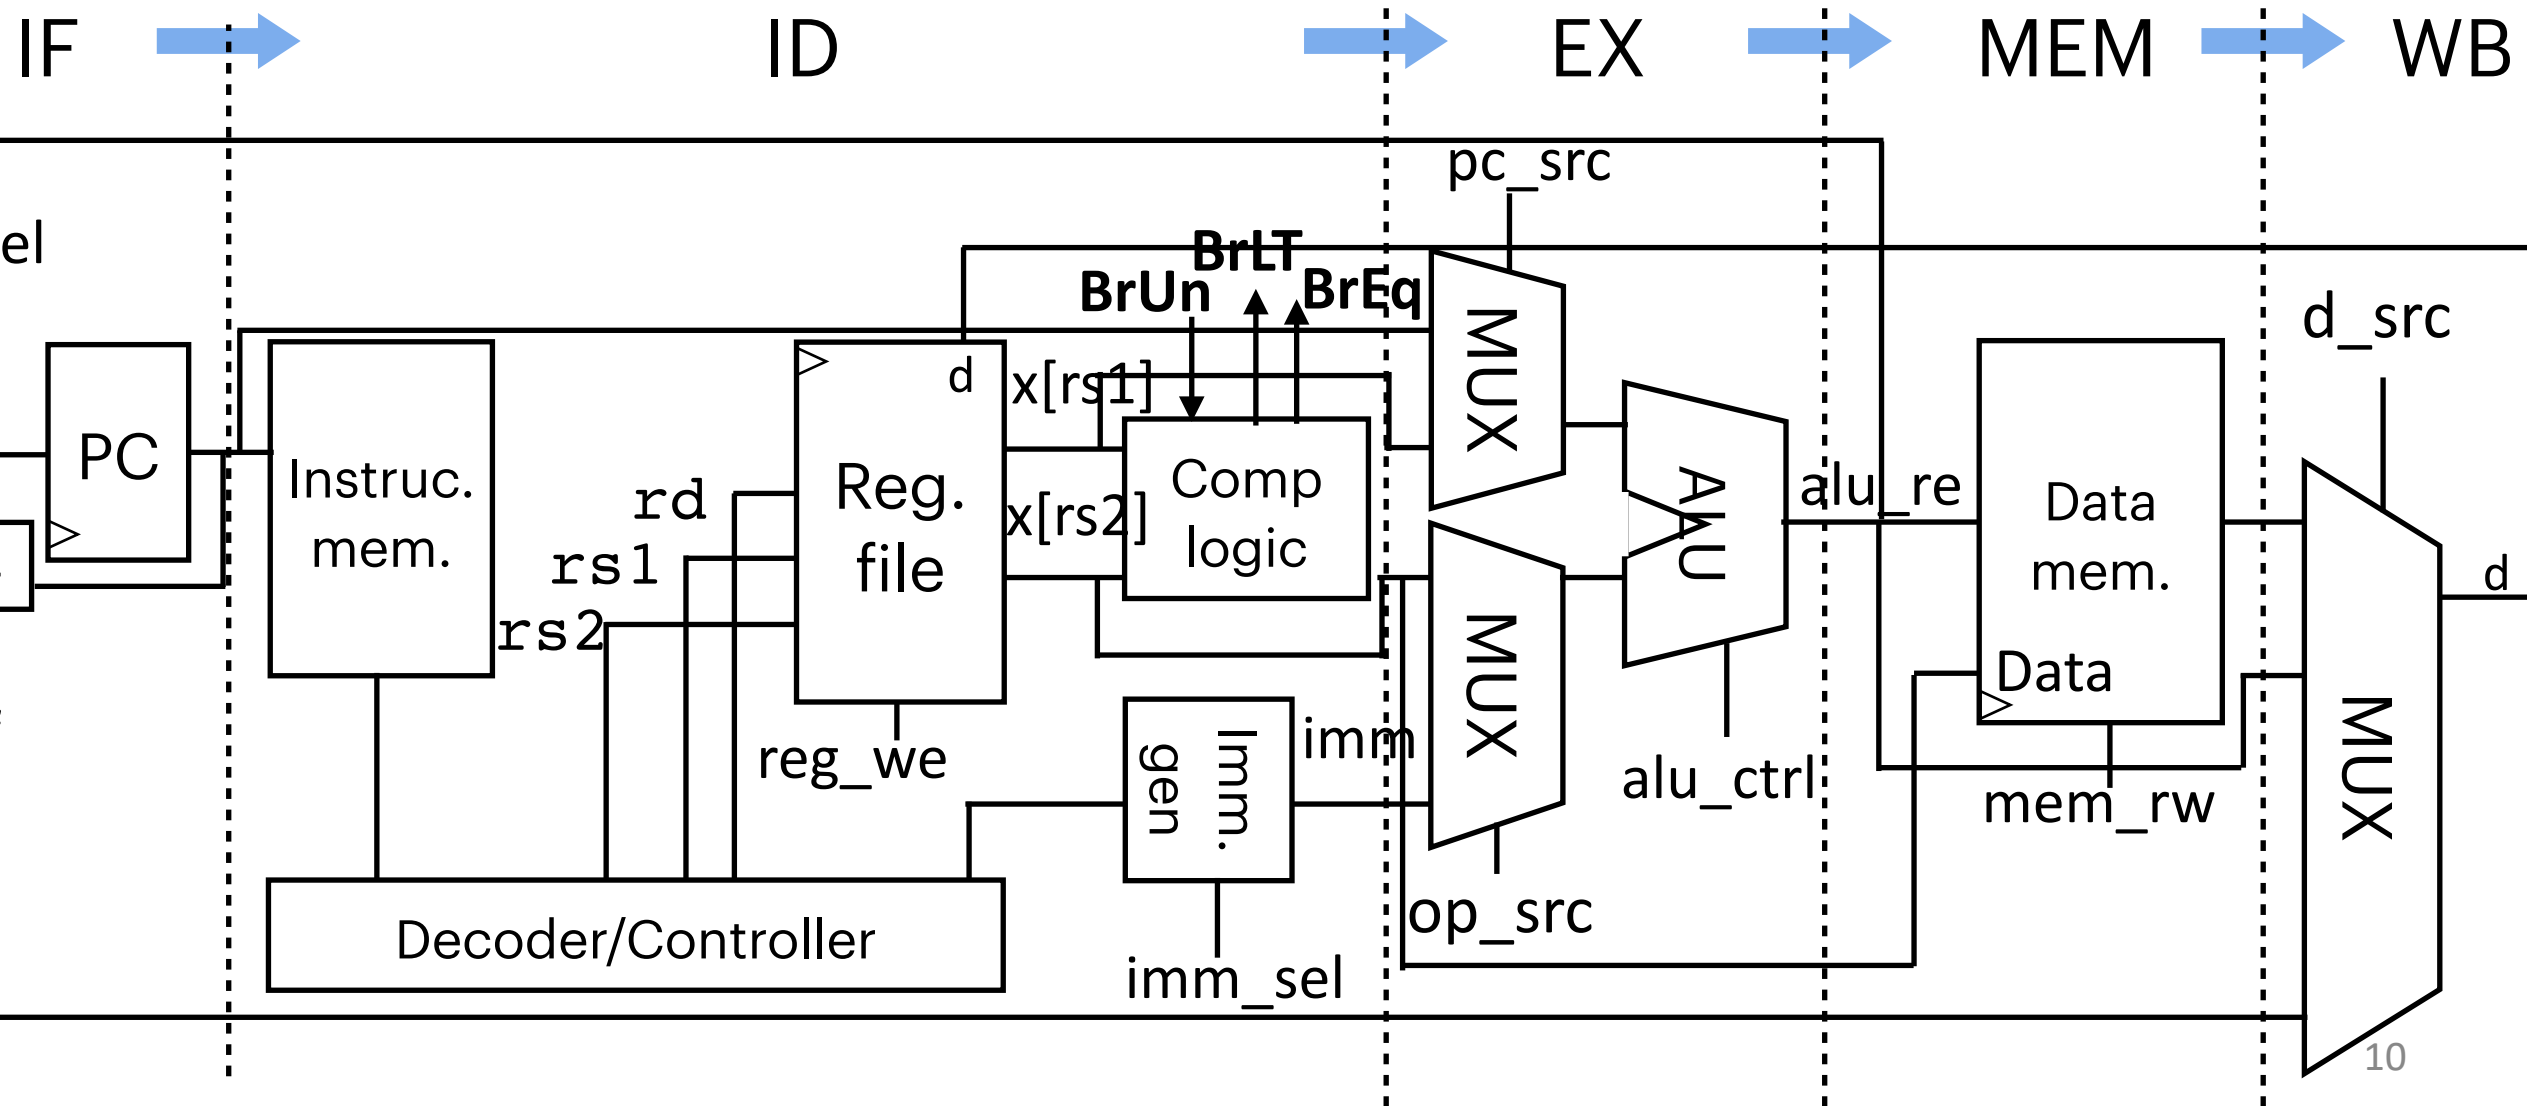
Timing Diagram (Consider delays)

Instru.	IF = 300 ps	ID = 100 ps	EX = 200 ps	MEM = 300 ps	WB = 100 ps	Total
lw	X	X	X	X	X	1000 ps

$\frac{\text{Cycles}}{\text{Instruction}}$ (CPI)

$\frac{\text{Time}}{\text{Cycle}}$ (Critical path; technology; TPD etc.)

Power
 $P \propto cvf^2$



Timing Diagram (Consider delays)

假设代码共1000行

一个周期 (cycle) 执行一条指令 $\Rightarrow CPI = 1$

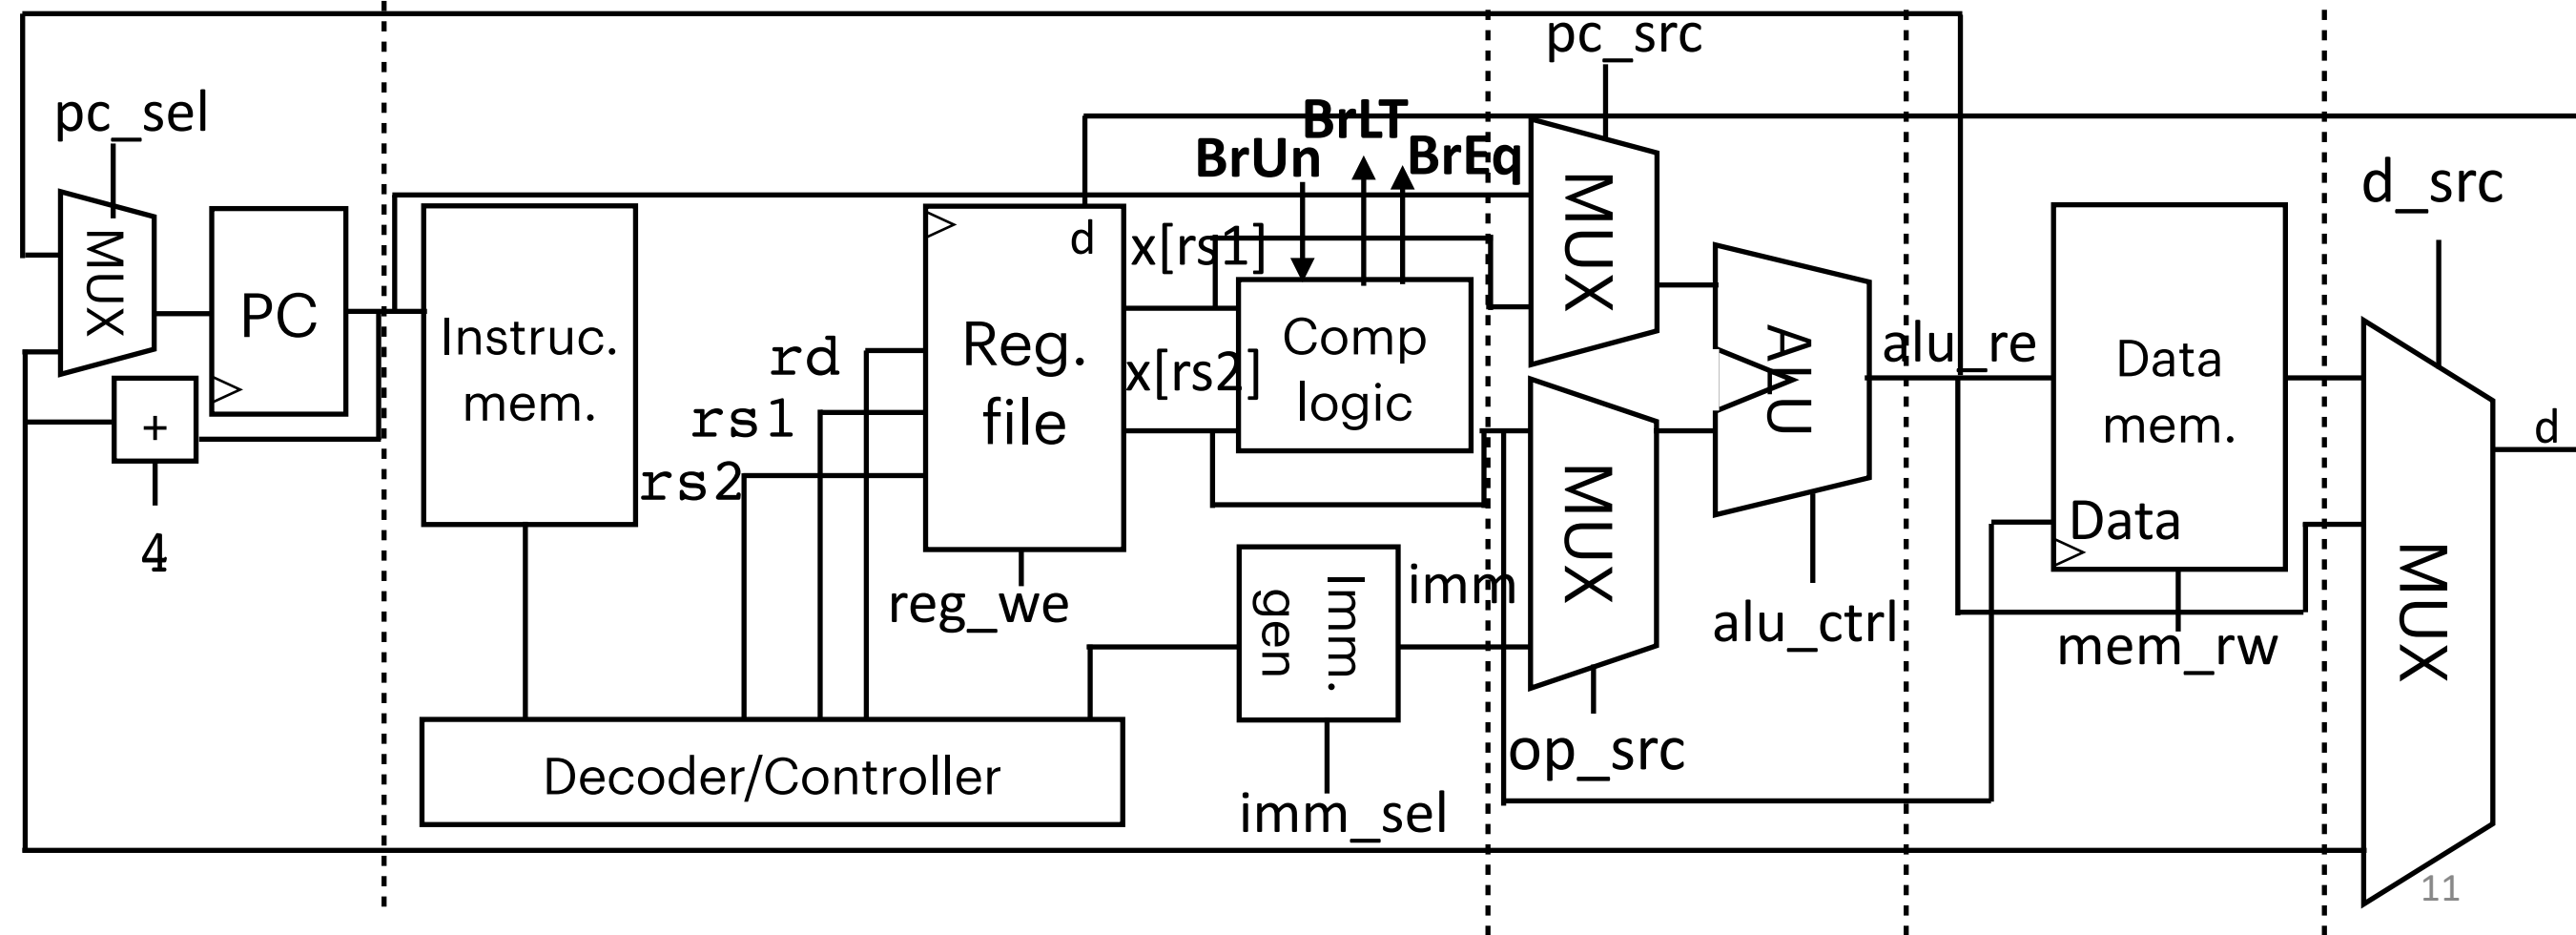
Instru.	IF = 300 ps	ID = 100 ps	EX = 200 ps	MEM = 300 ps	WB = 100 ps	Total
lw	X	X	X	X	X	1000 ps

性能

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}}$$

$$\frac{\text{Time}}{\text{program}} = \frac{1000}{1} \cdot 1.100 \text{ ps} = 1.1 \mu\text{s}$$

IF → ID → EX → MEM → WB



Pipeline: Improve Time/Program

- Instruction: PCR test: 1. scan QR code; 2. get the tube; 3. sample
- 5 seconds for each step:
- Single-cycle (once for all)



1000人 \Rightarrow 1000条指令

1个人重复1次 \Rightarrow CPI = 1

1000 \times

1

\times

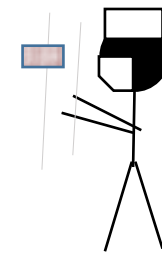
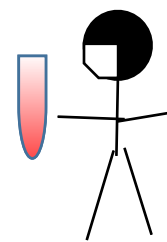
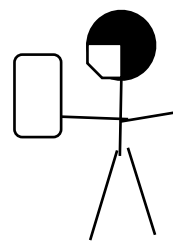
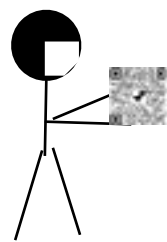
15S

= 15000S

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Clock cycle}}.$$

Multi-cycle

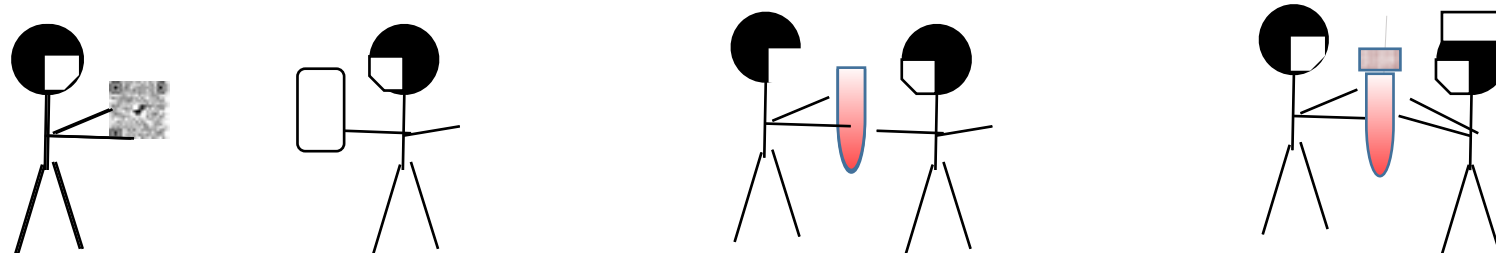
- Instruction: PCR test: 1. scan QR code; 2. get the tube; 3. sample
- 5 seconds for each step
- Multi-cycle or multi-step



$$\frac{\text{Time}}{\text{Program}} = \frac{\overset{1000}{\text{Instructions}}}{\text{Program}} \times \frac{\overset{3}{\text{Clock cycles}}}{\text{Instruction}} \times \frac{\overset{5}{\text{Time}}}{\text{Clock cycle}}.$$

Pipeline

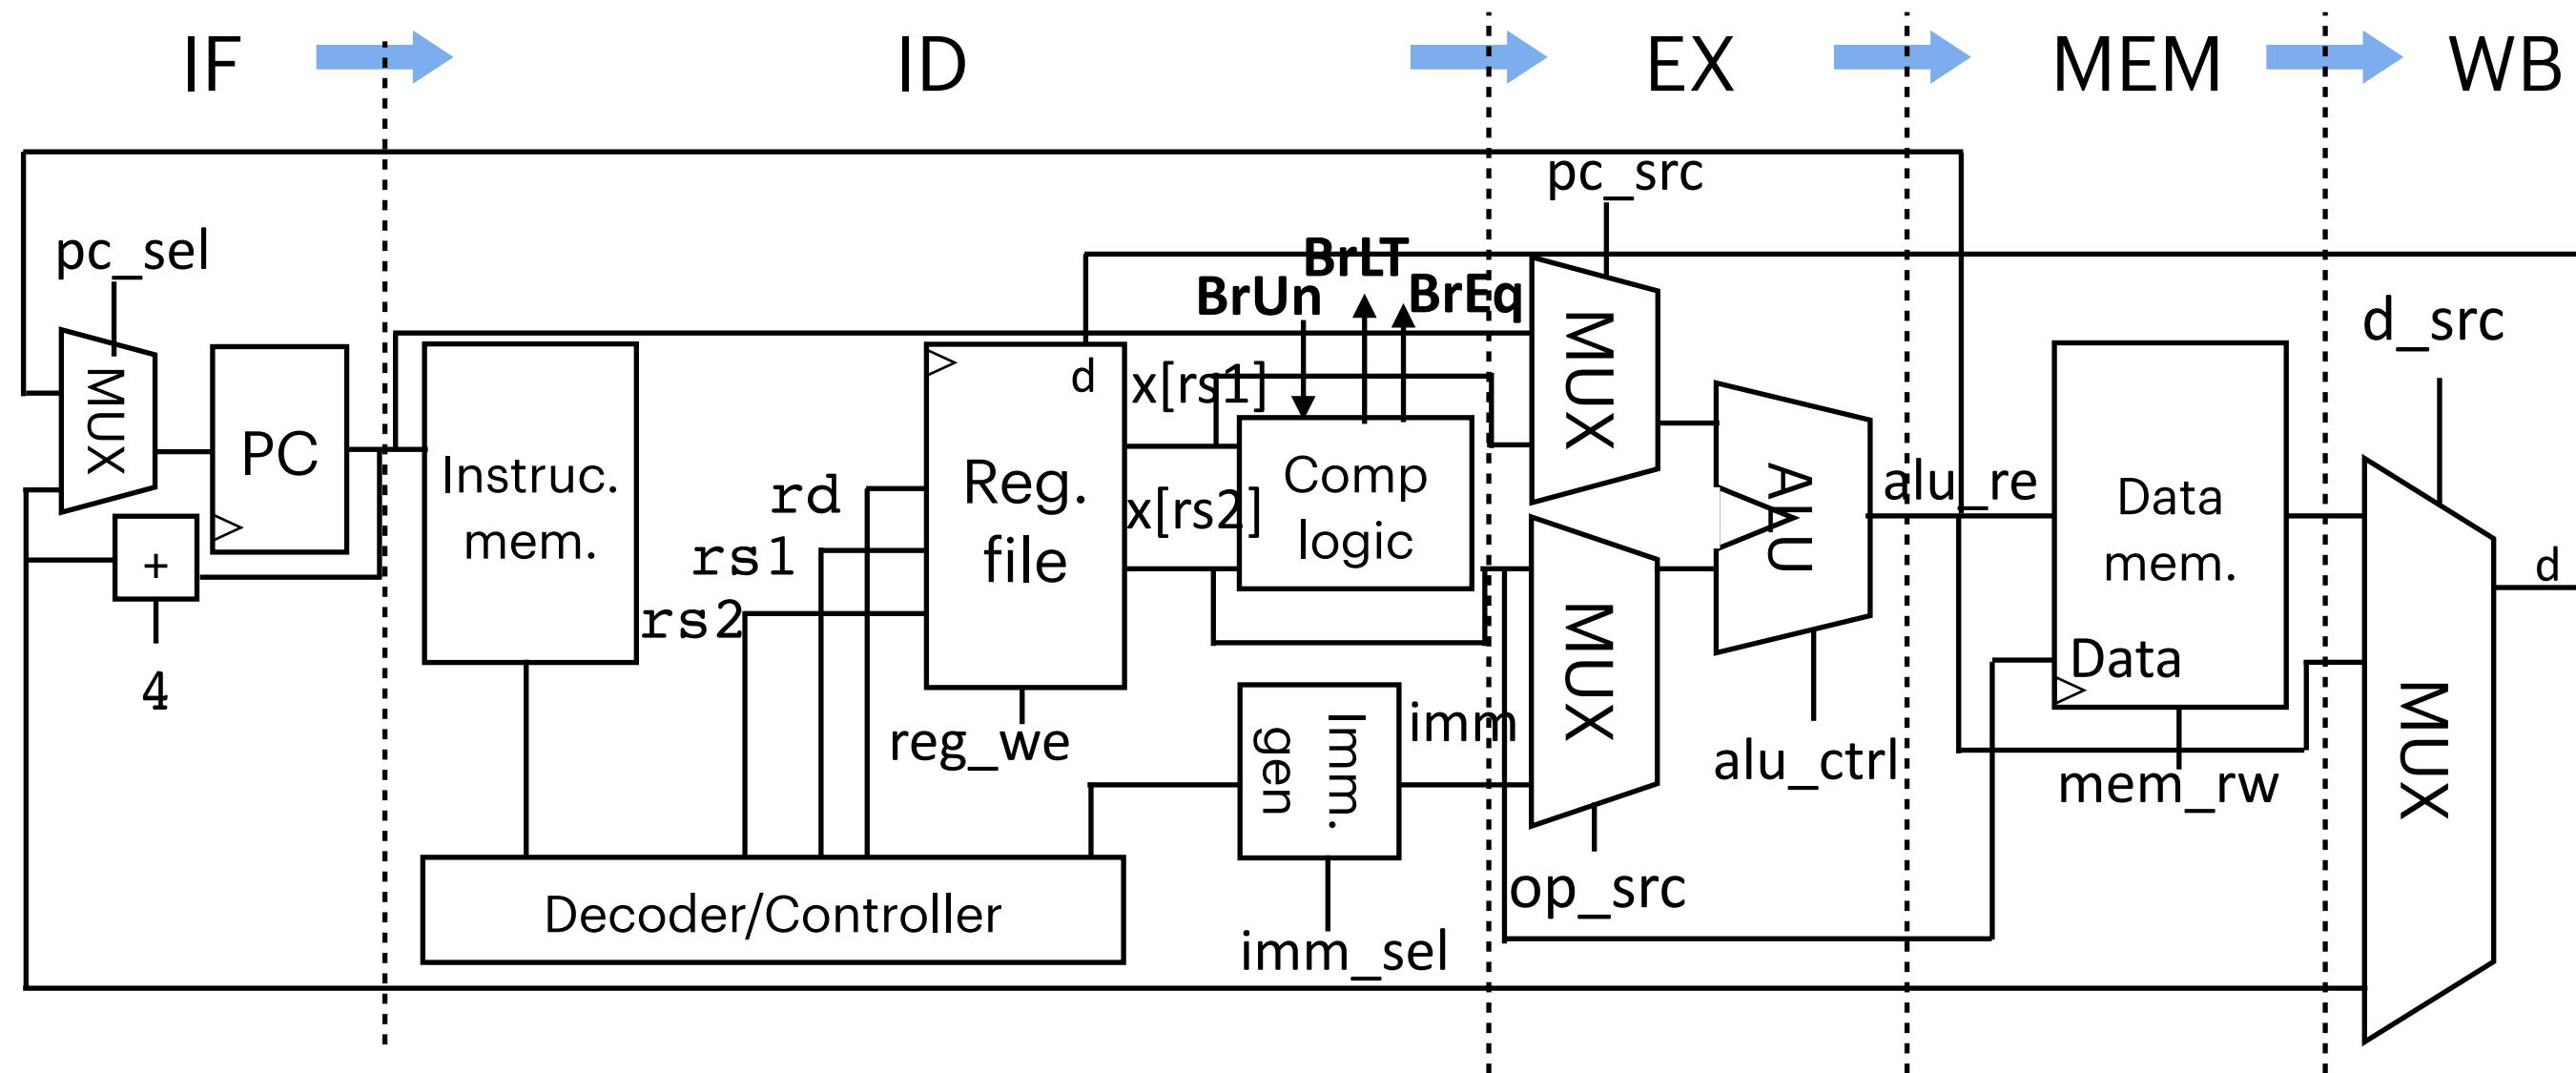
- PCR test in pipeline



$$\frac{\text{Time}}{\text{Program}} = \frac{\overset{1000}{\text{Instructions}}}{\text{Program}} \times \frac{\overset{1}{\text{Clock cycles}}}{\text{Instruction}} \times \frac{\overset{5}{\text{Time}}}{\text{Clock cycle}}.$$

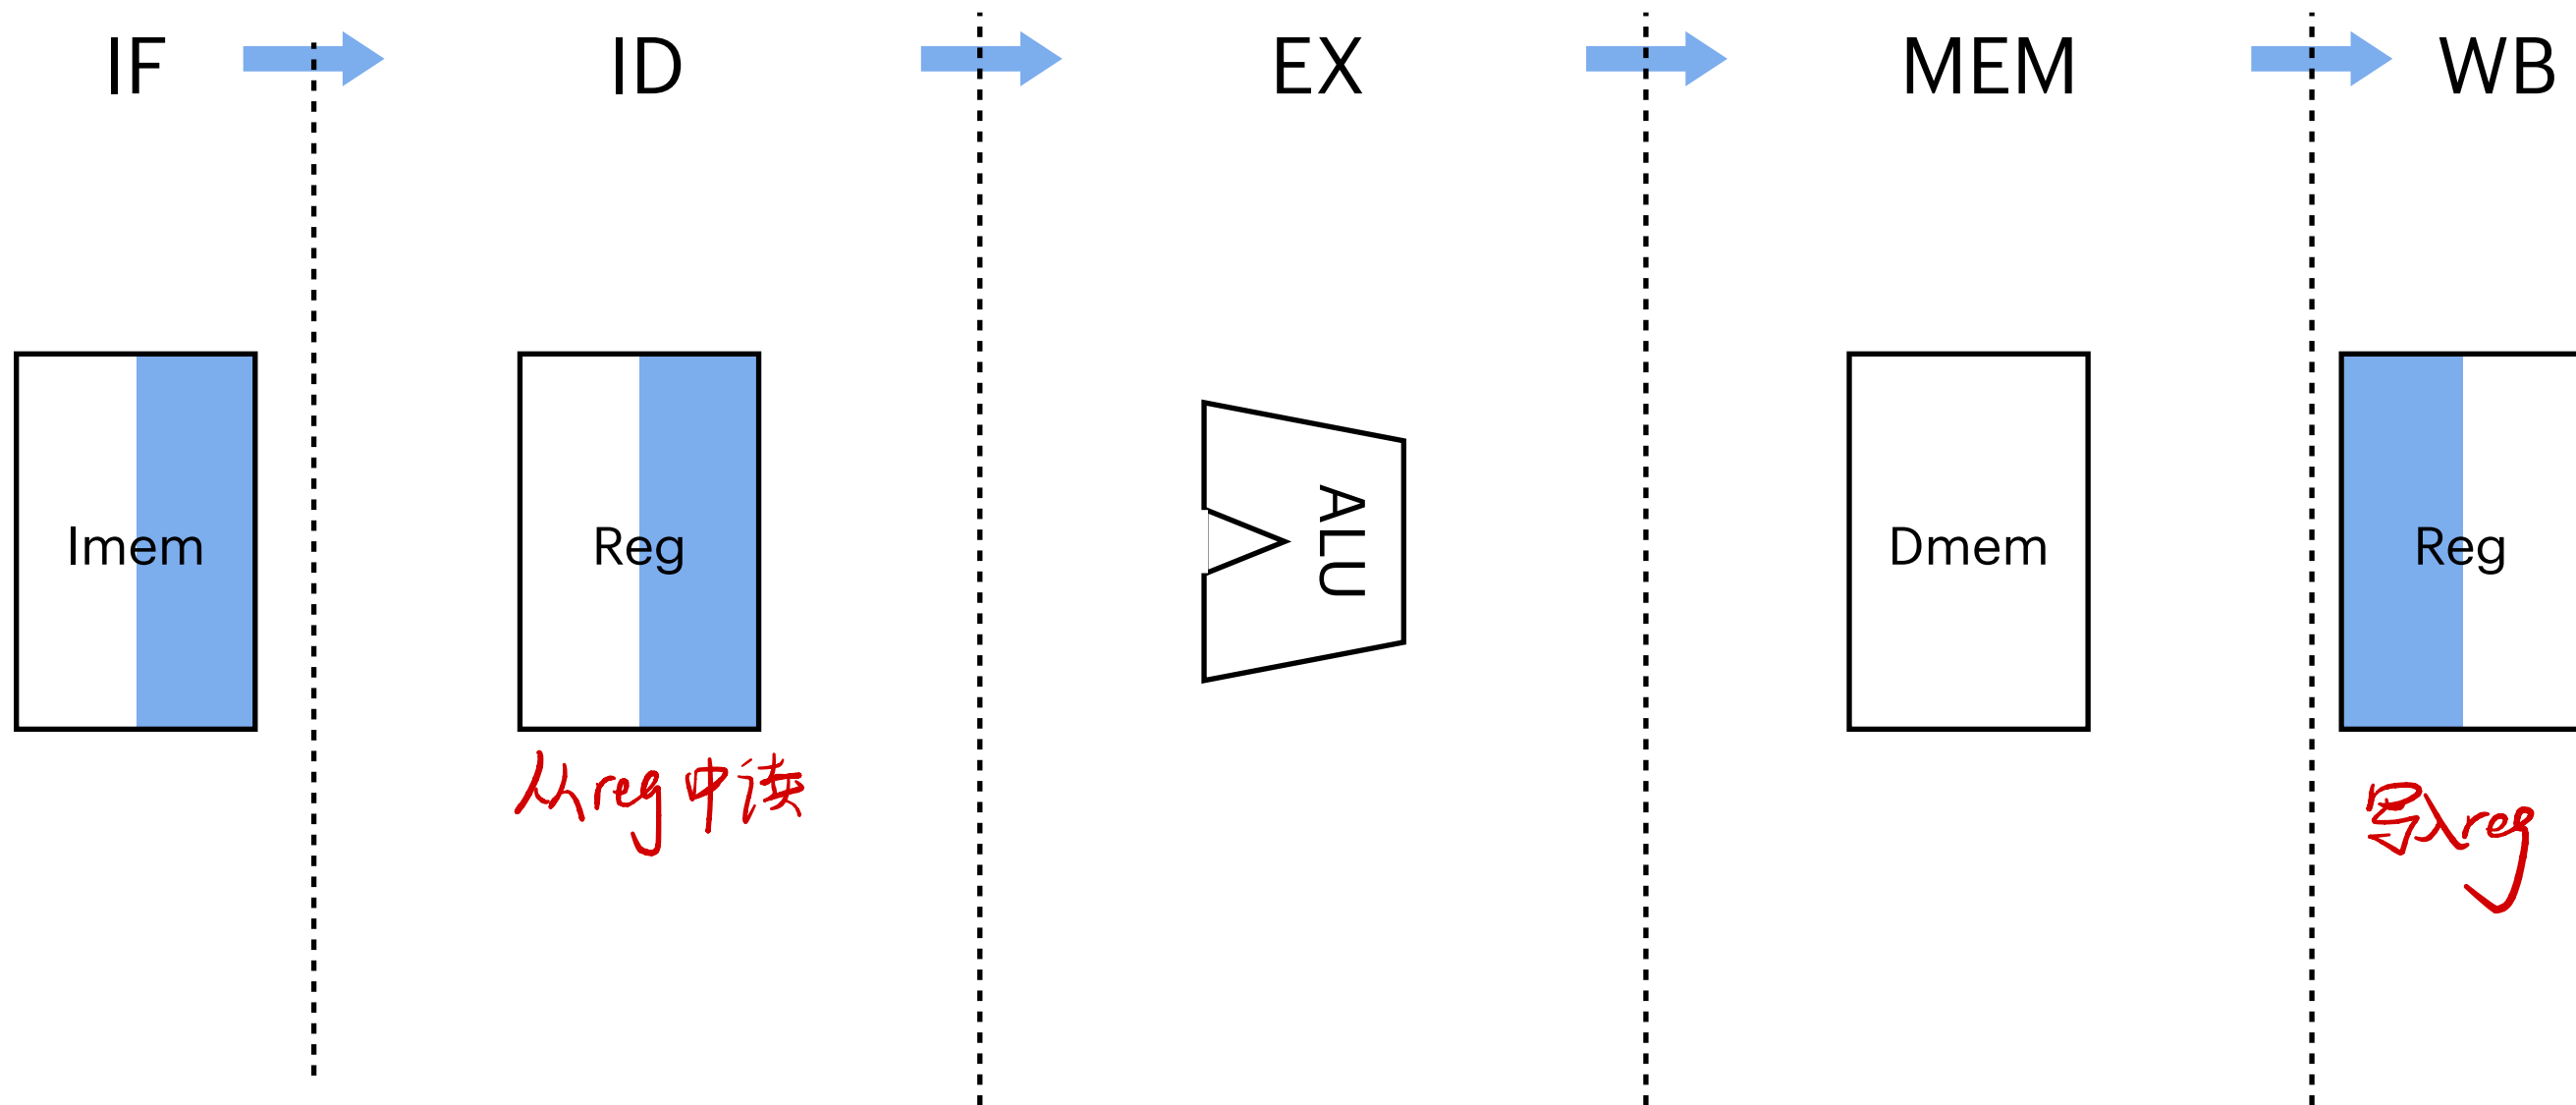
↓
系统整体来看
 $\frac{3}{3} = 1$

Analogy in our CPU Design

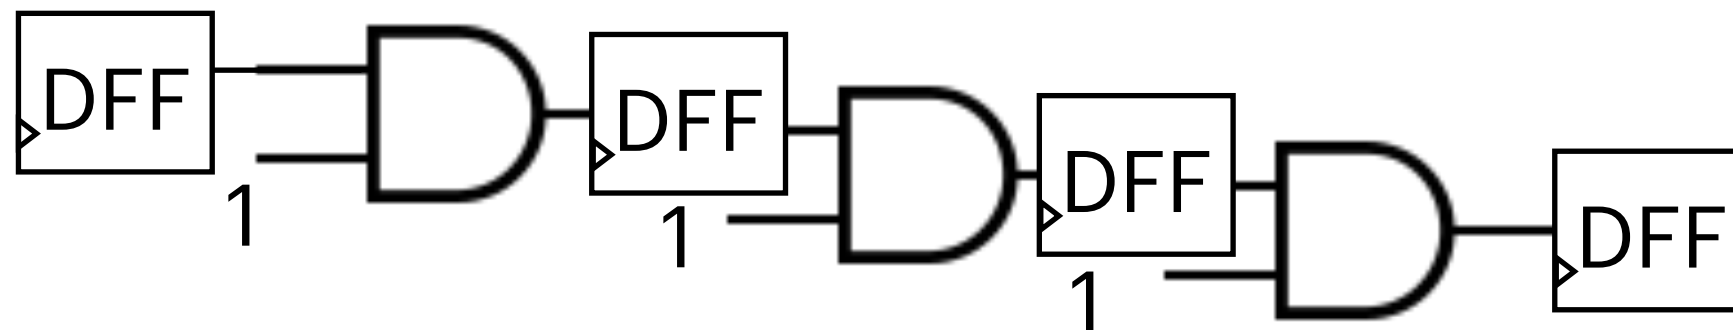
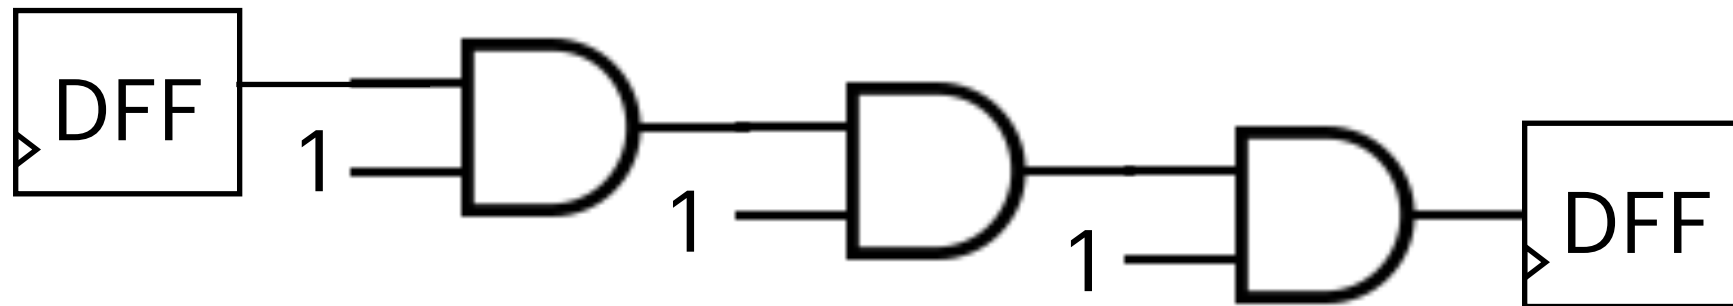


$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Clock cycle}}.$$

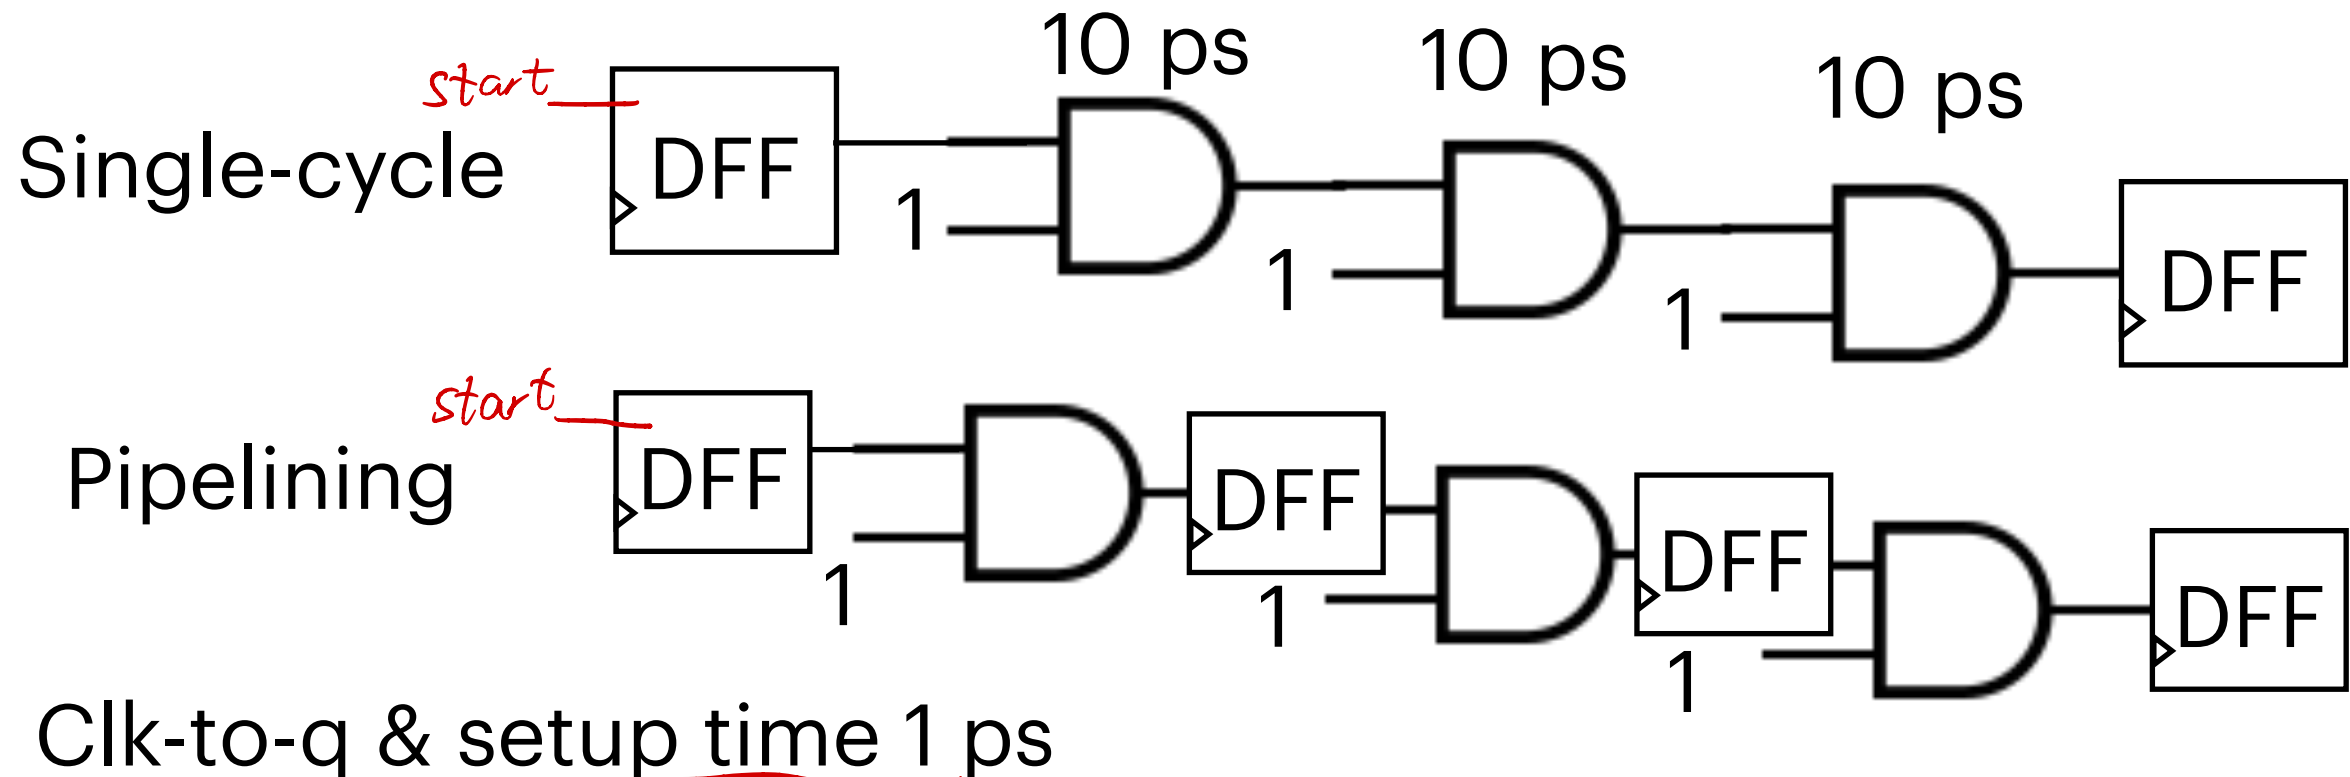
Simplify the Model using Symbols



Recall DFFs

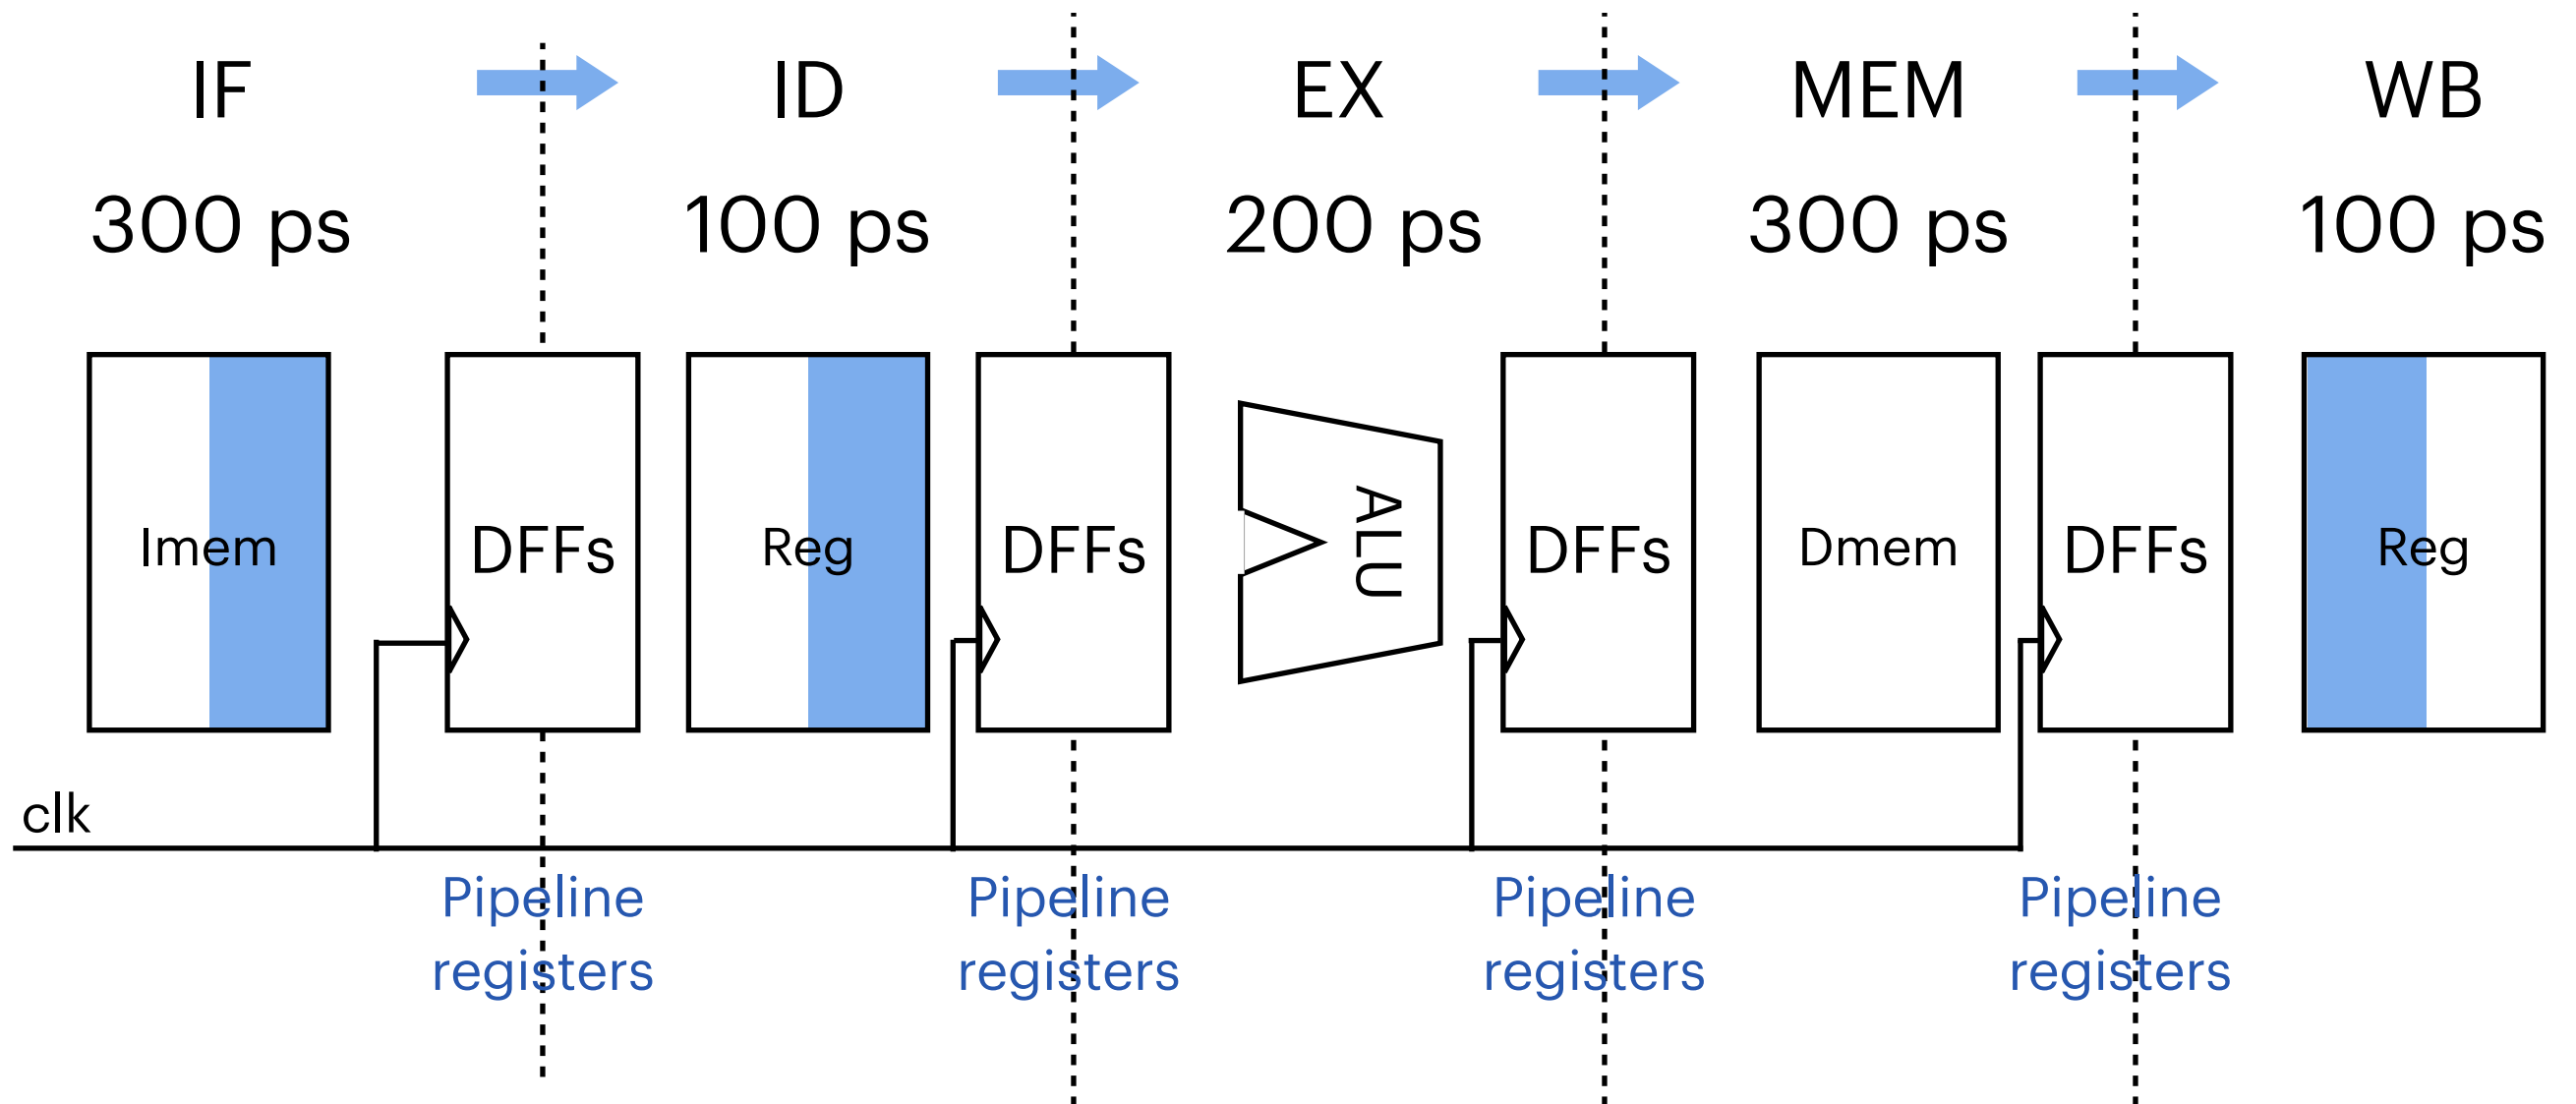


Recall DFFs



	Single-cycle	Pipelining
T_{clk}	32 ps	12 ps
Clock rate/frequency	$\frac{1}{32 ps}$	$\frac{1}{12 ps}$
Computation time	32 ps	36 ps
Clock cycle <u>per output</u>	1	1
Relative speed	1 (设为1)	$\frac{32}{12}$
Hardware	2 DFF + 3 AND	4 DFF + 3 AND

Insert Pipeline Registers

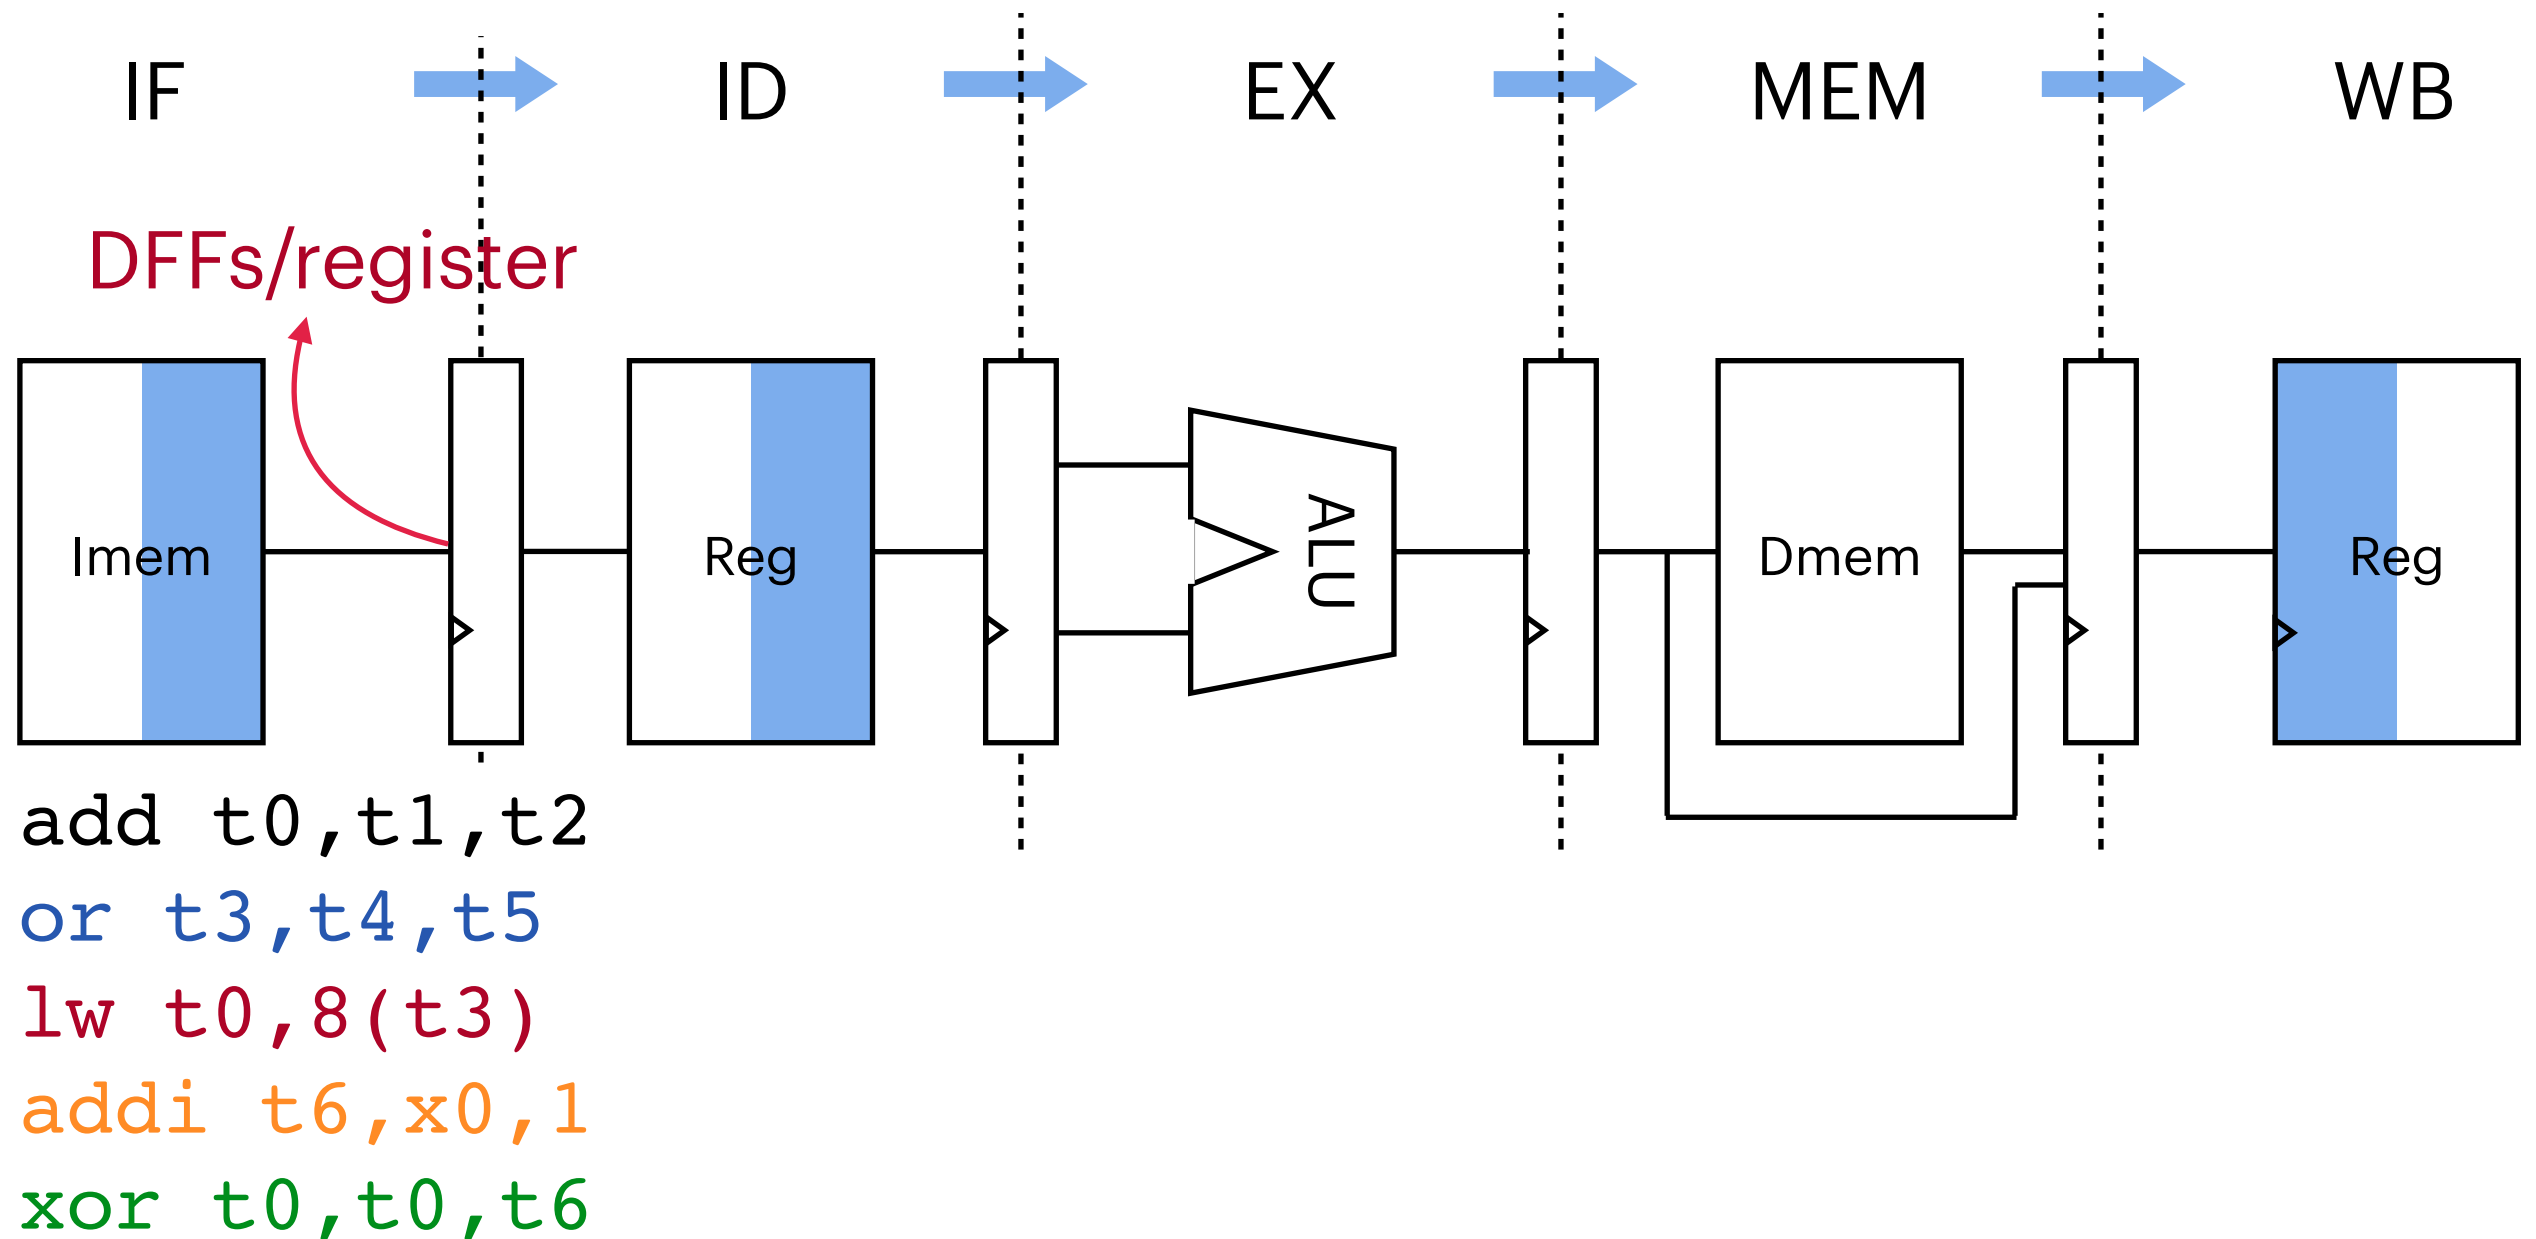


$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}}$$

300ps

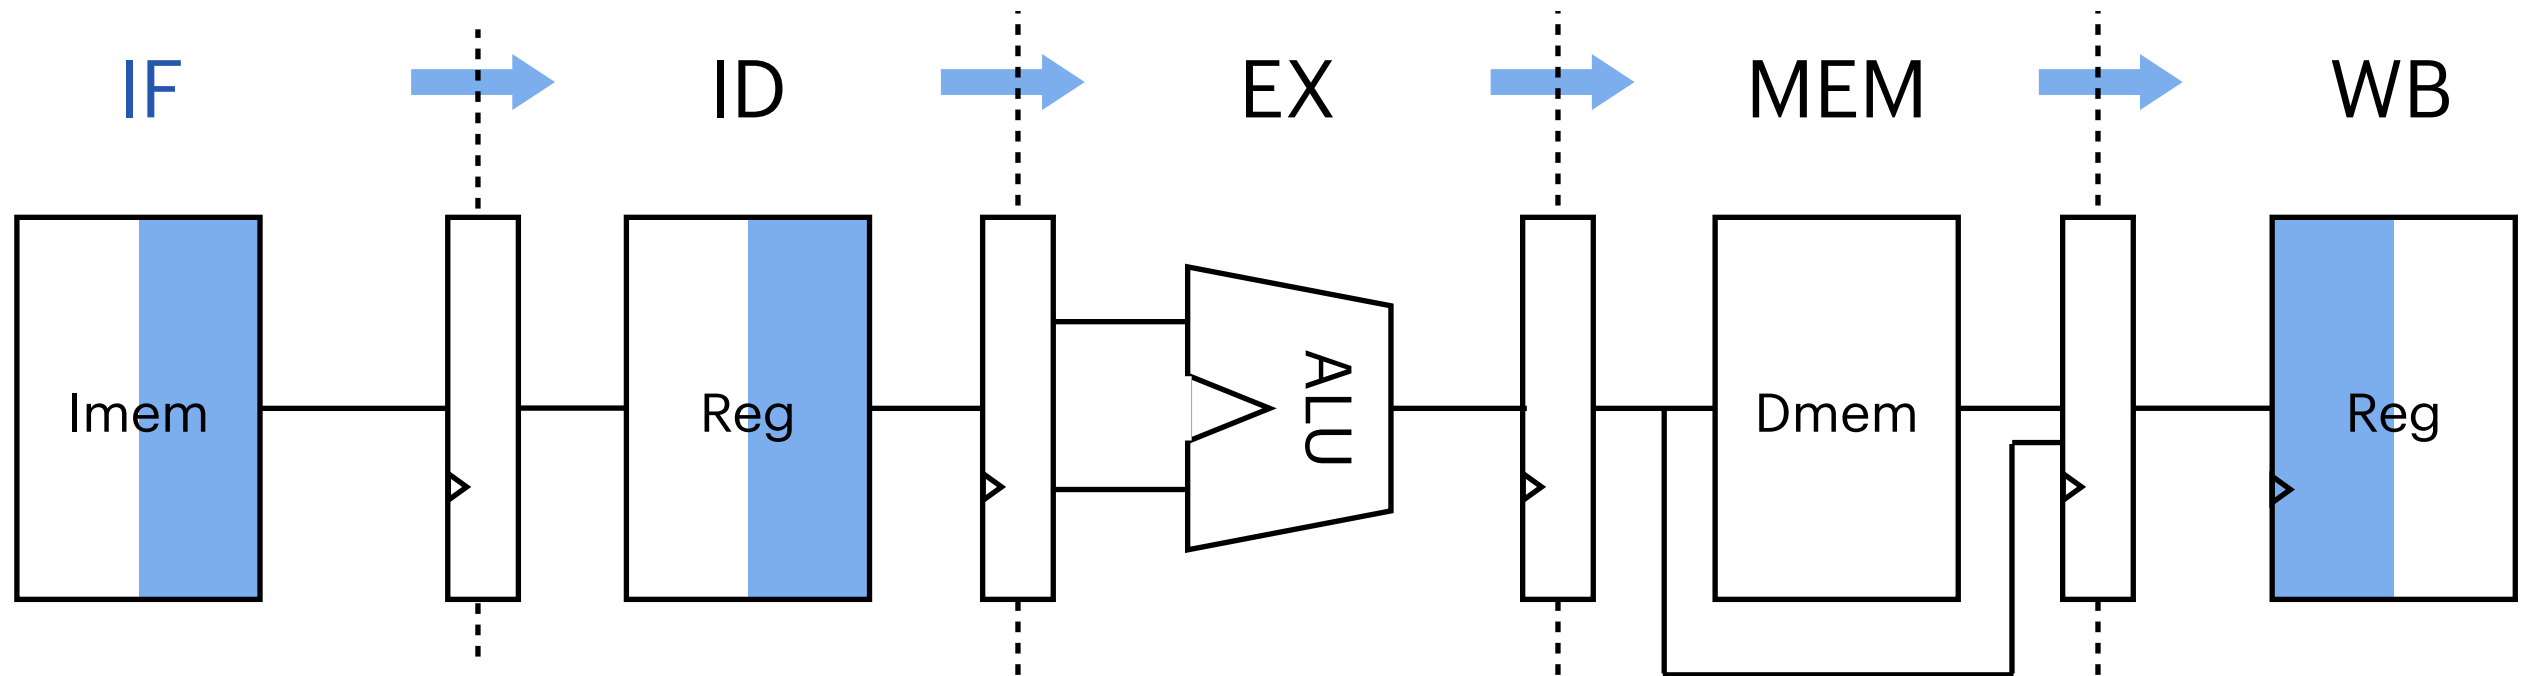
$$T_{clk} = \text{Max Delay} = \max(t_{IF}, t_{ID}, t_{EX}, t_{MEM}, t_{WB})$$

Insert Pipeline Registers



$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}}$$

Insert Pipeline Registers



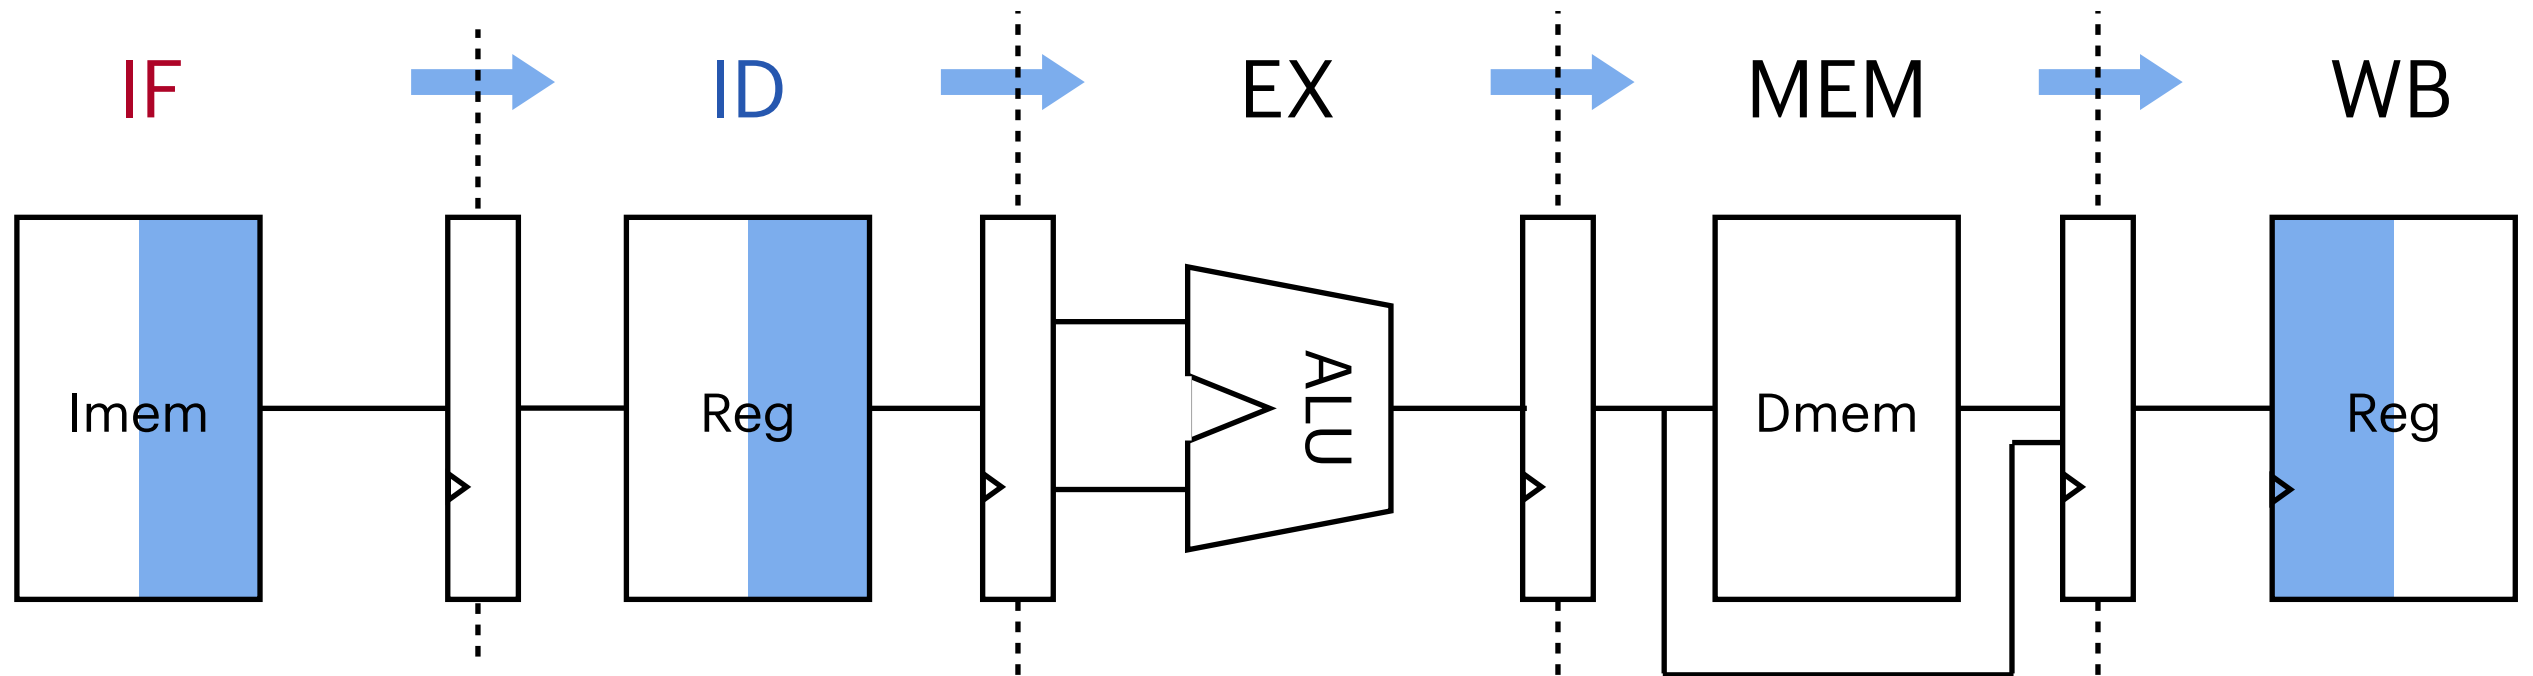
Clock cycle 1

IF add

Clock cycle 2 (PC + 4)

IF or **ID** add

Insert Pipeline Registers



Clock cycle 1

IF add

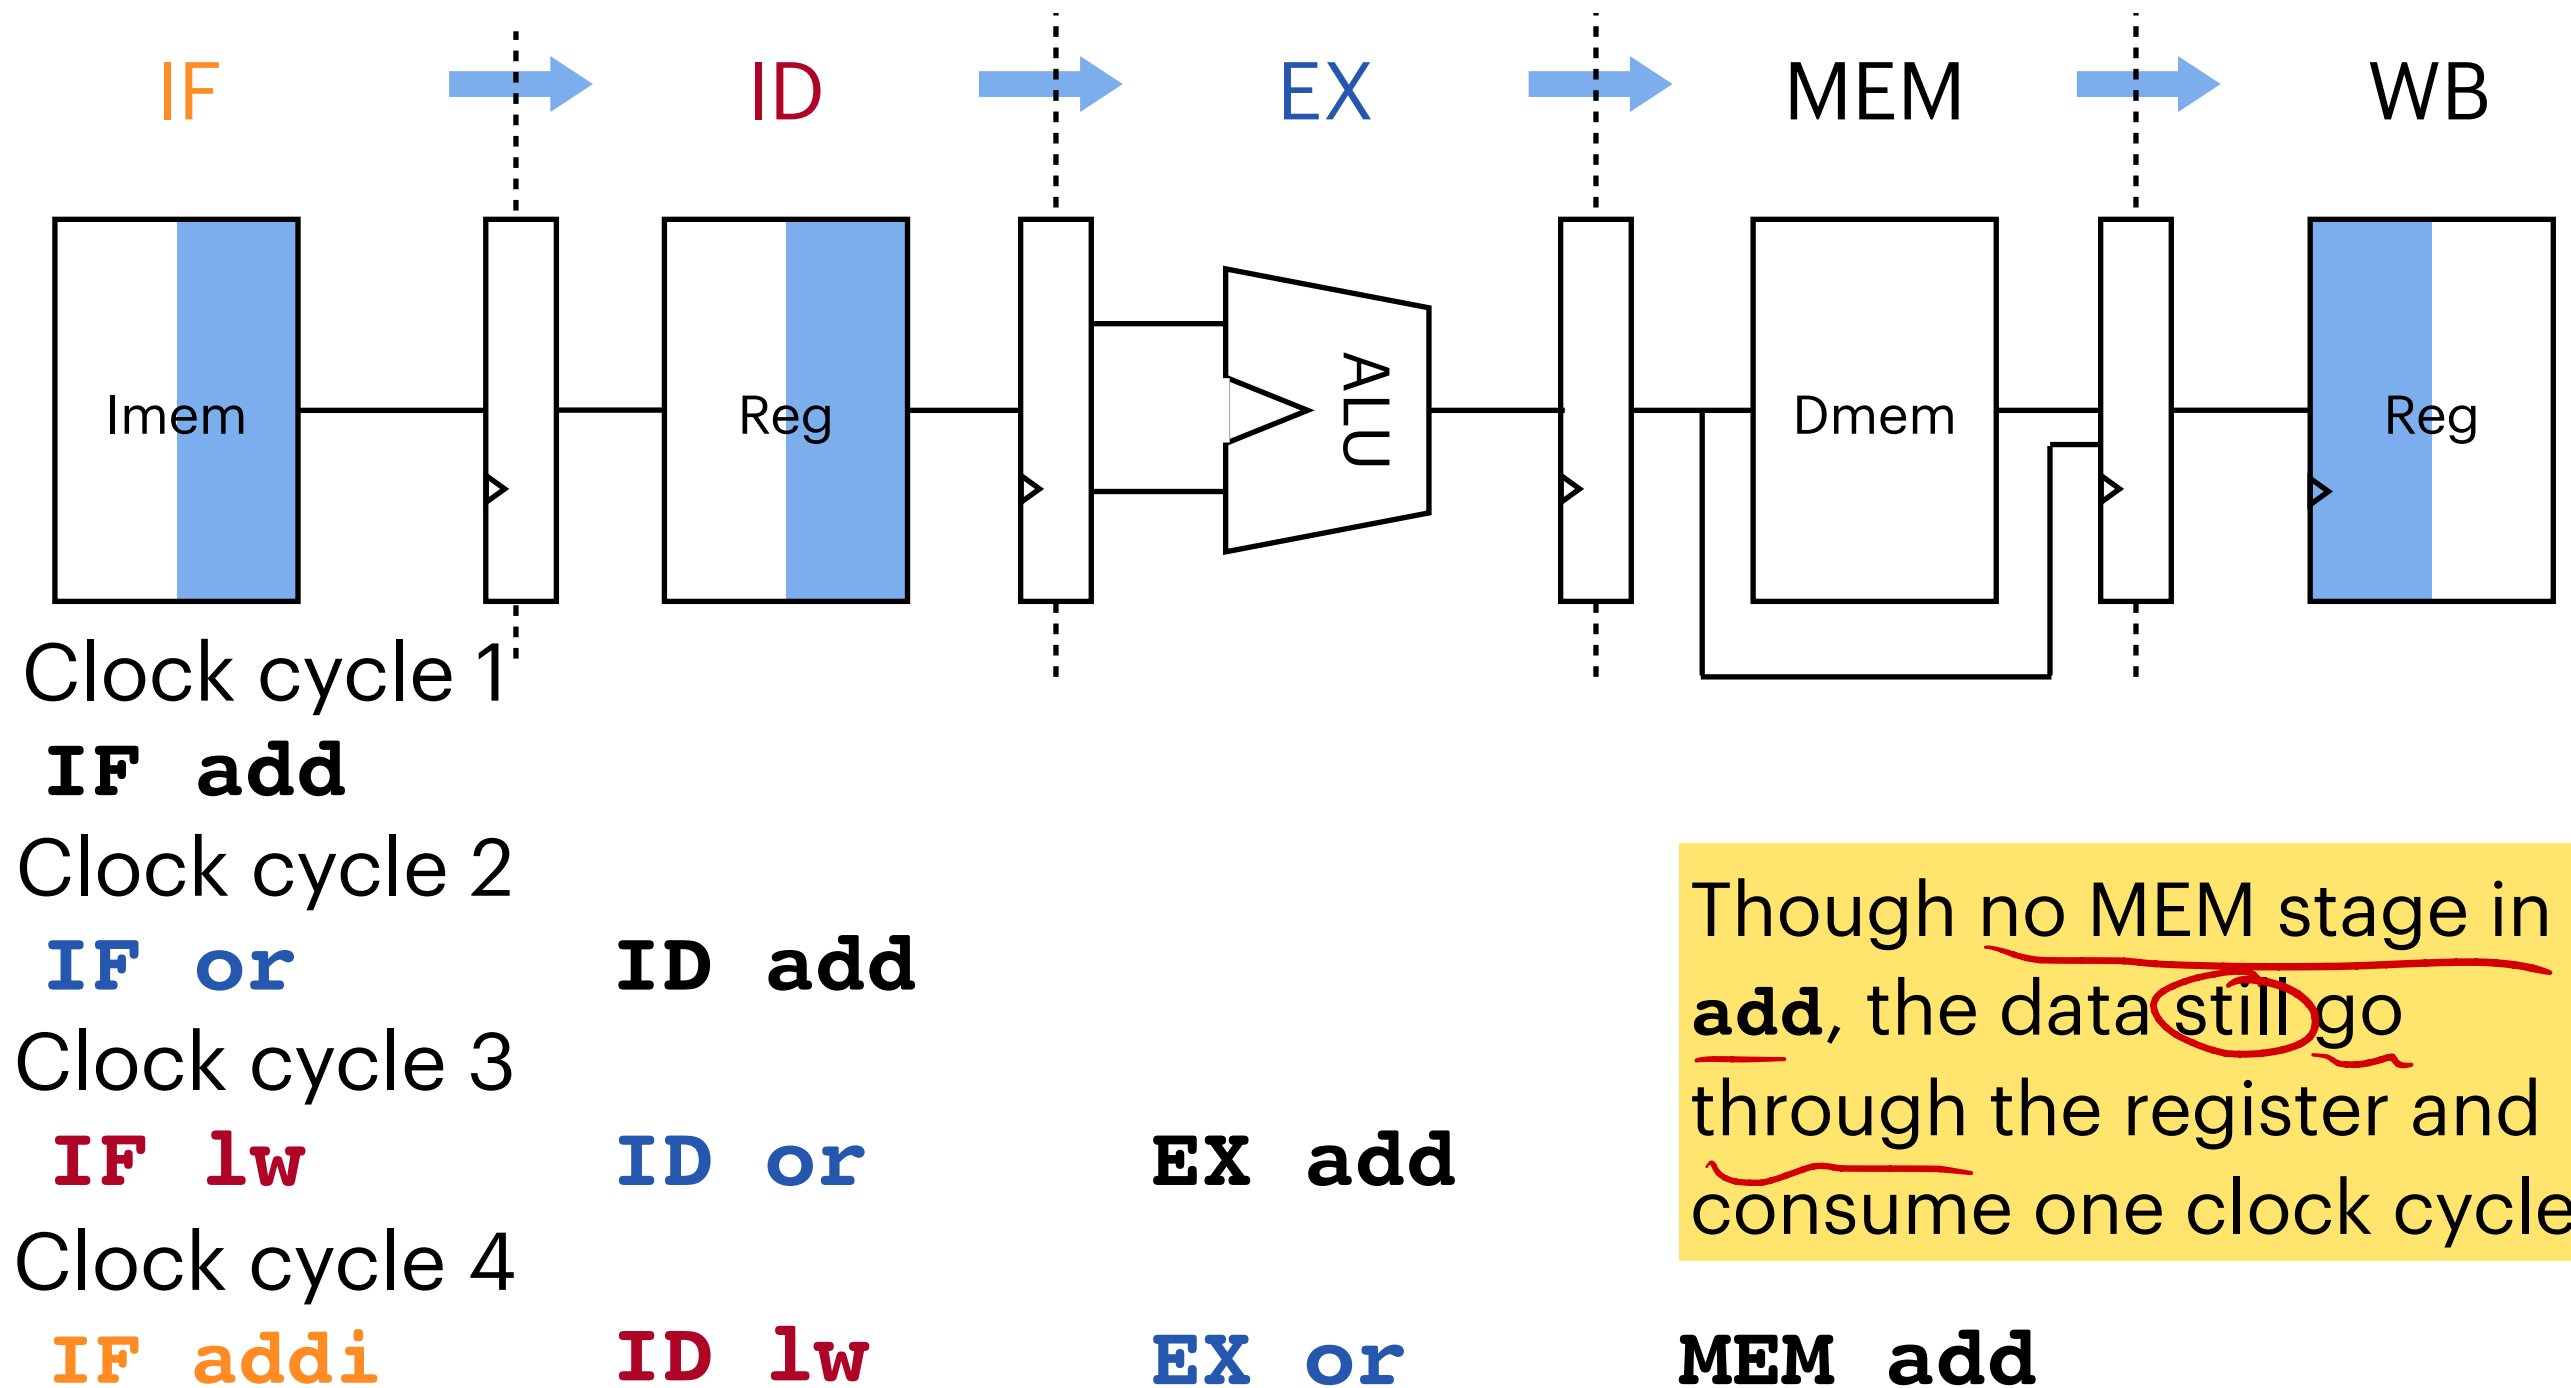
Clock cycle 2 (PC + 4)

IF or **ID** add

Clock cycle 3 (PC + 8)

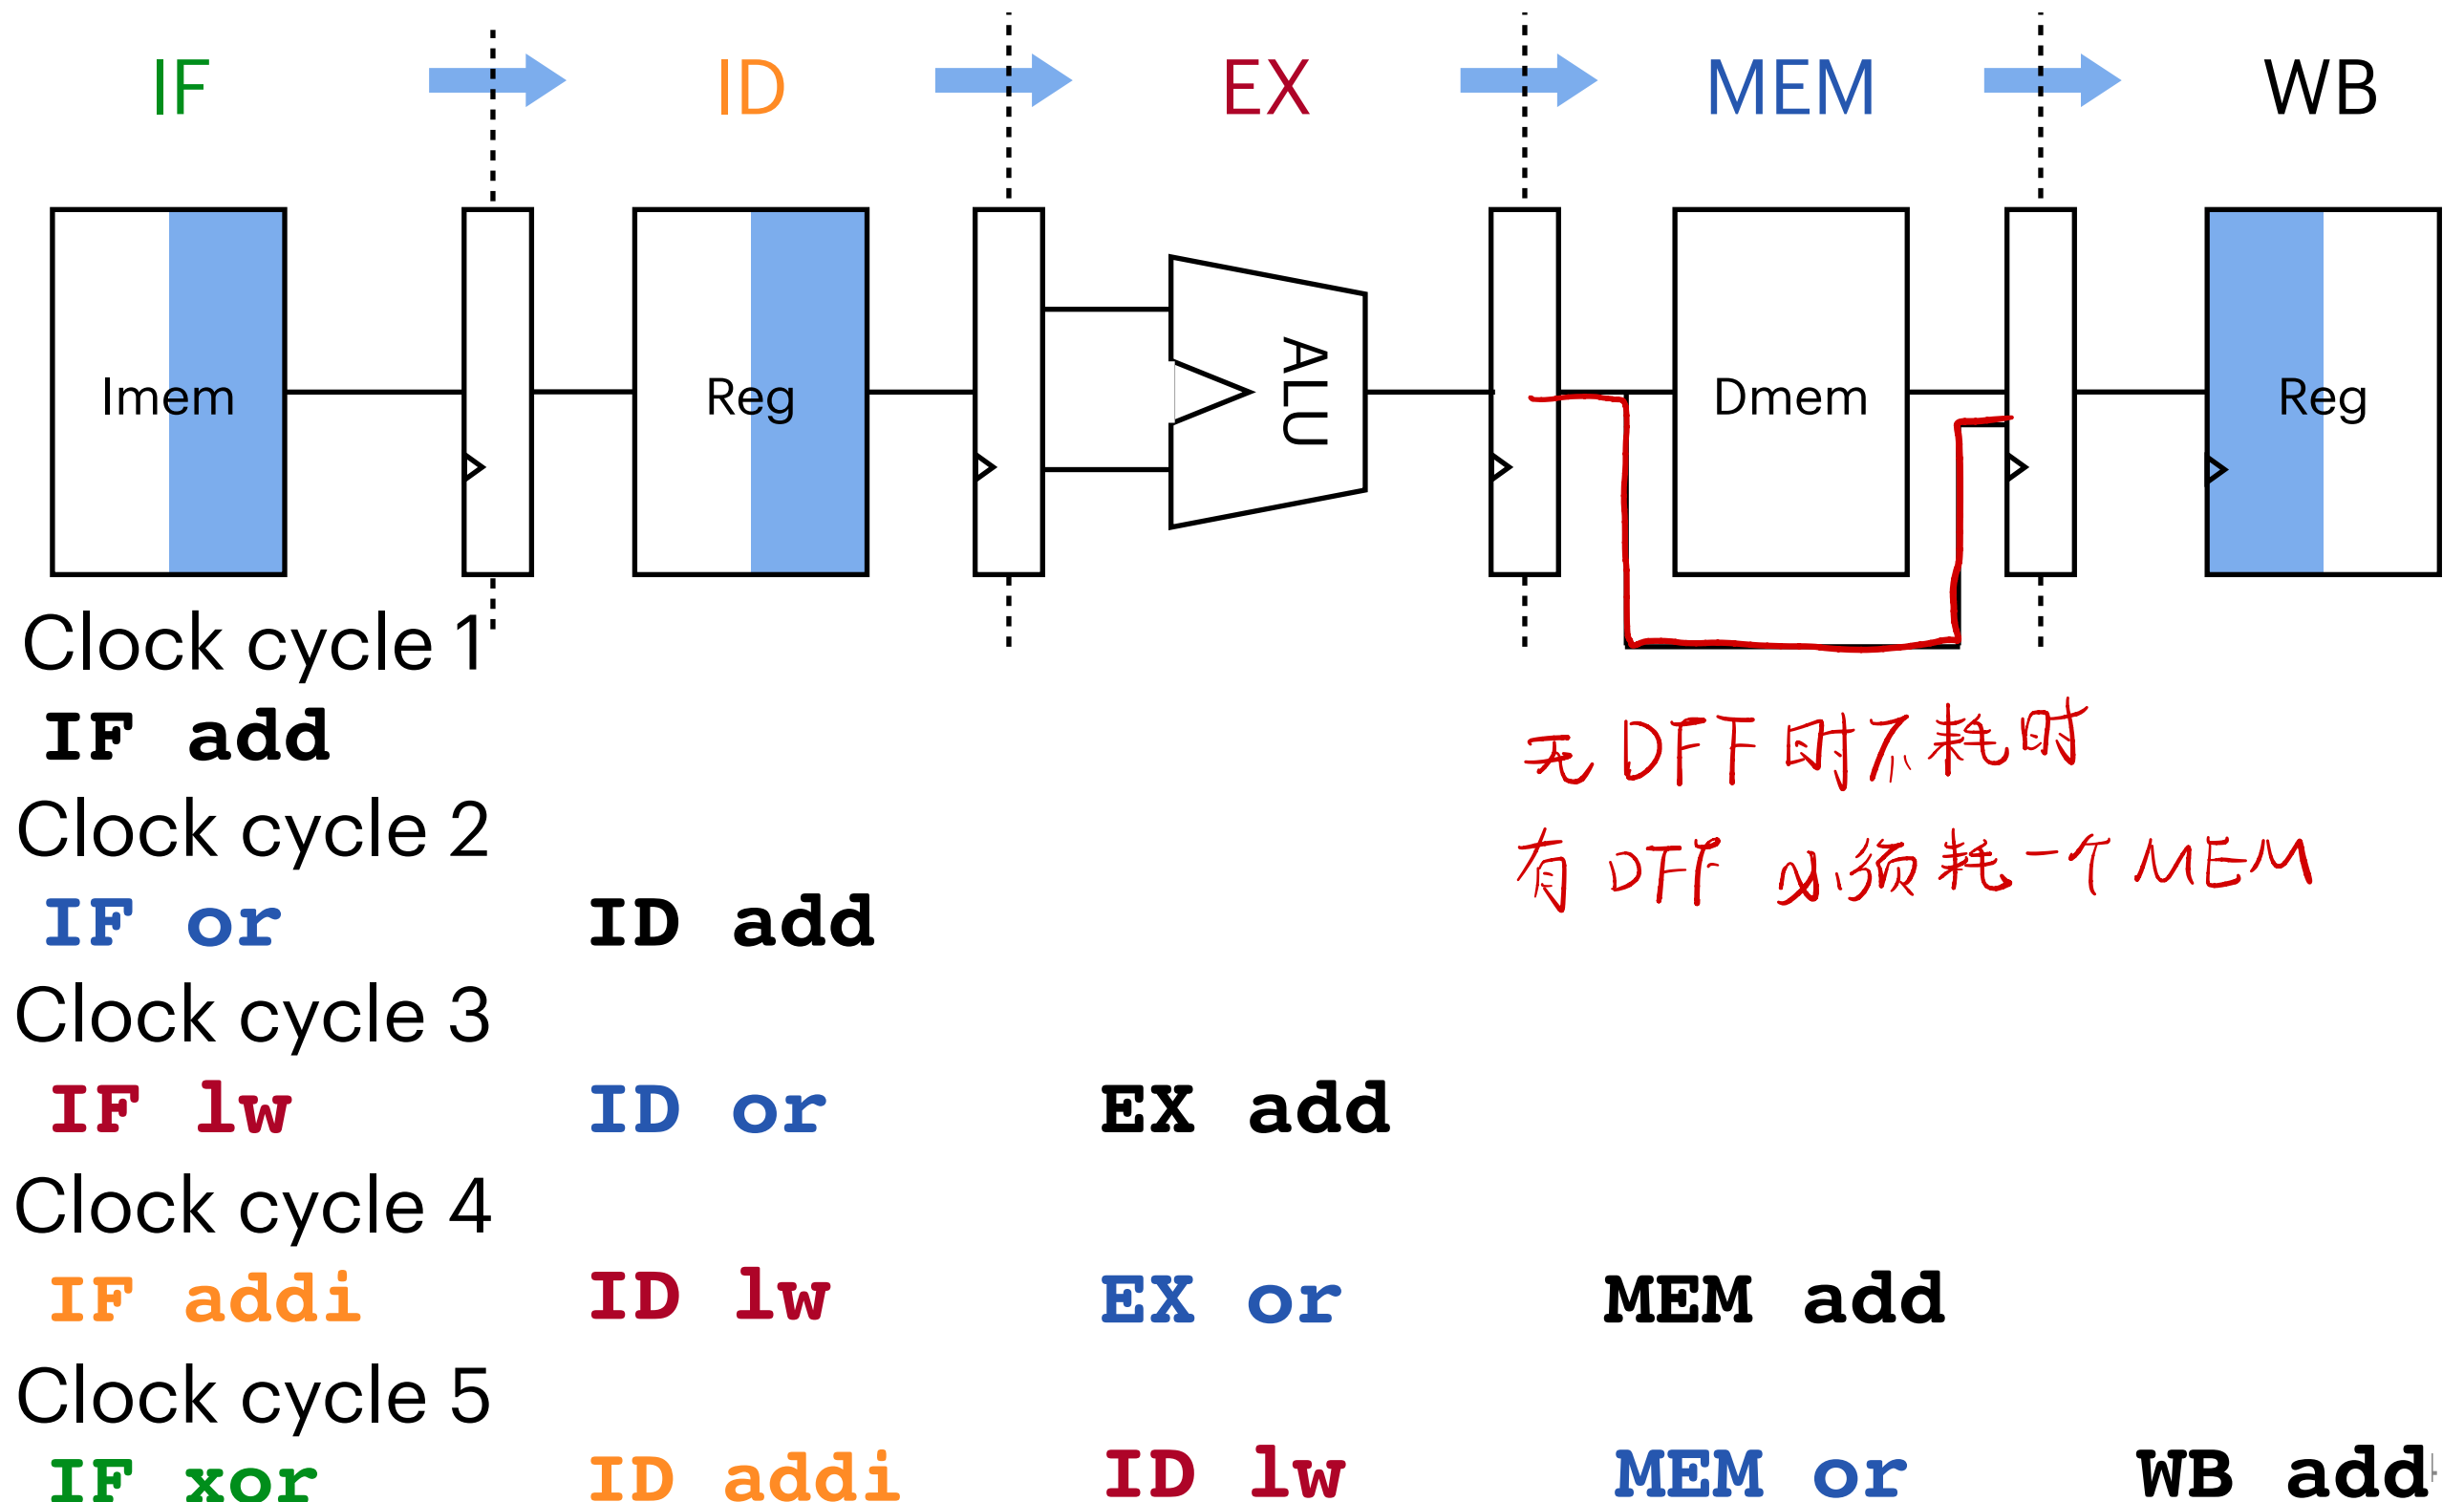
IF lw **ID** or **EX** add

Insert Pipeline Registers

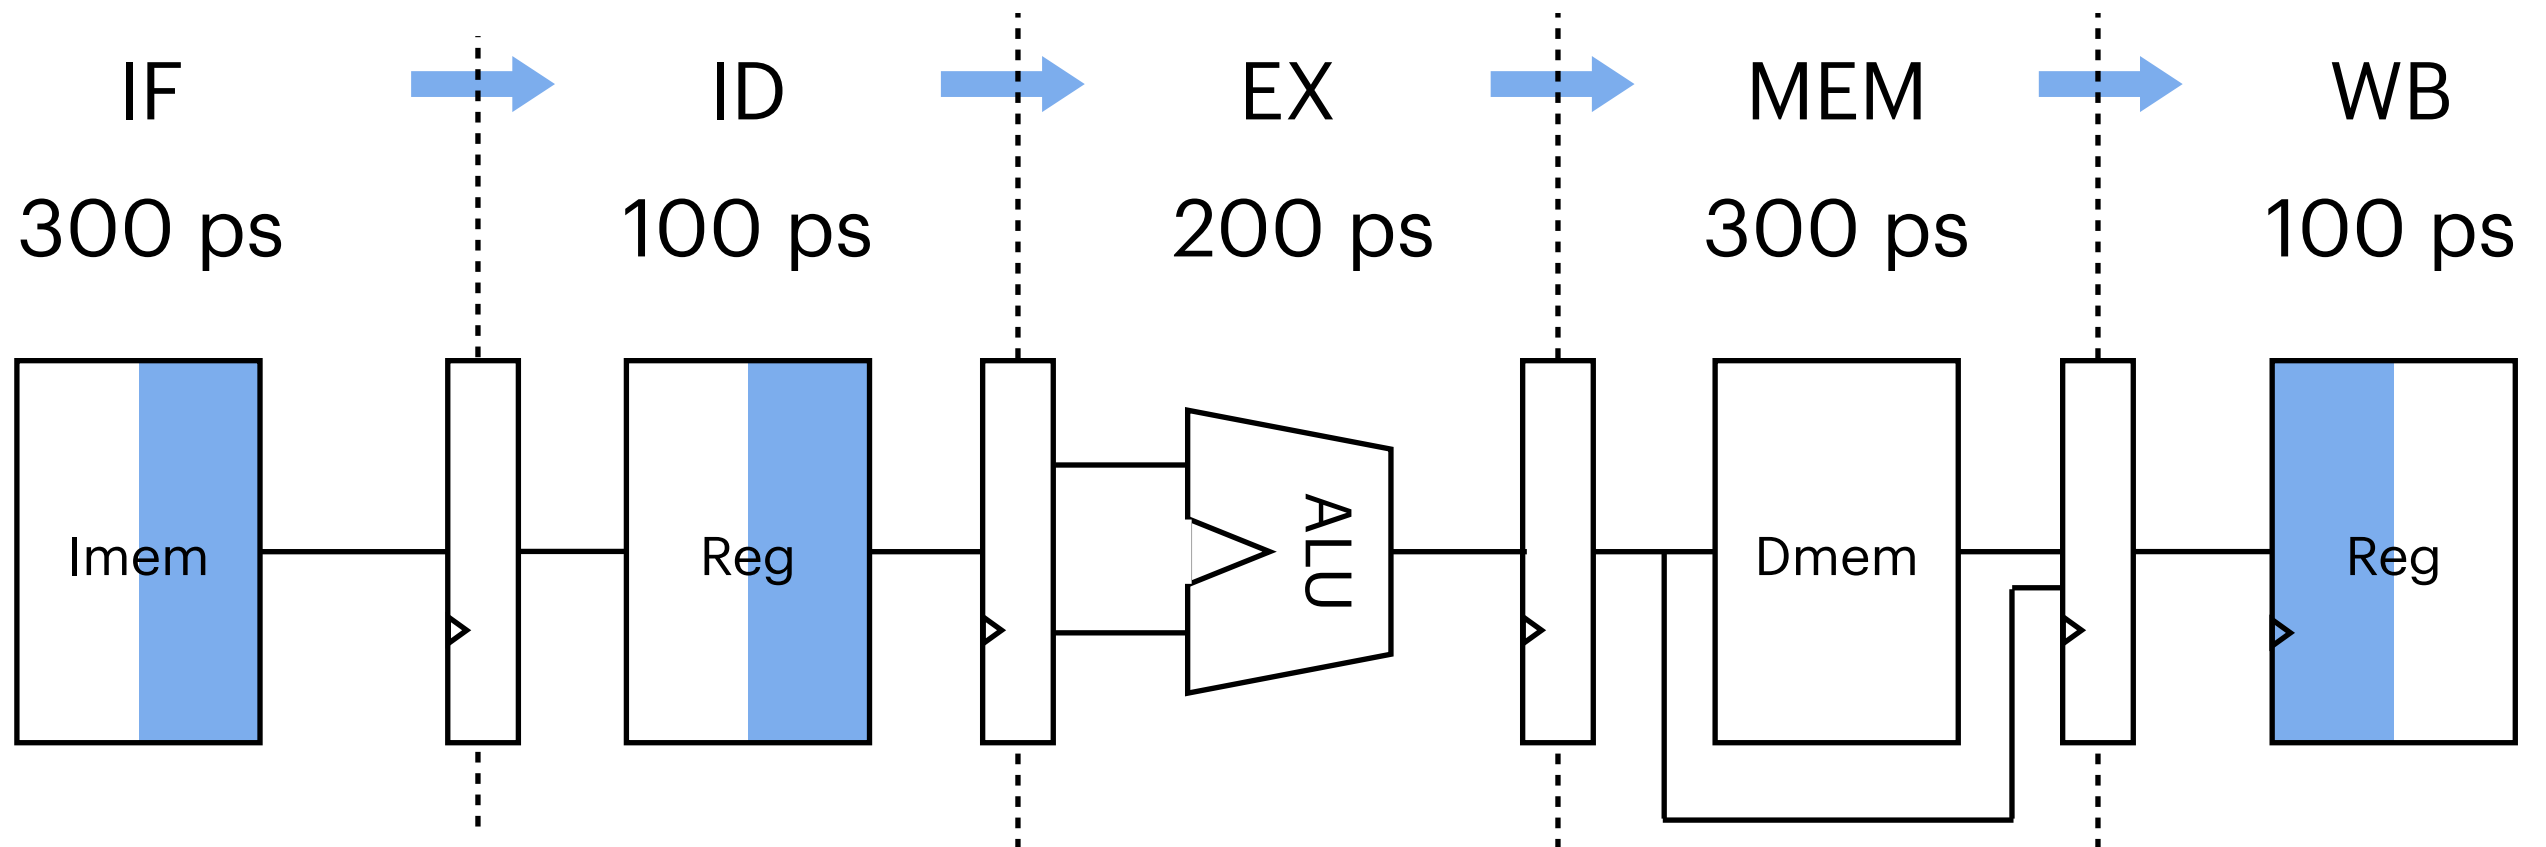


Though no MEM stage in add, the data still go through the register and consume one clock cycle

Insert Pipeline Registers



Insert Pipeline Registers

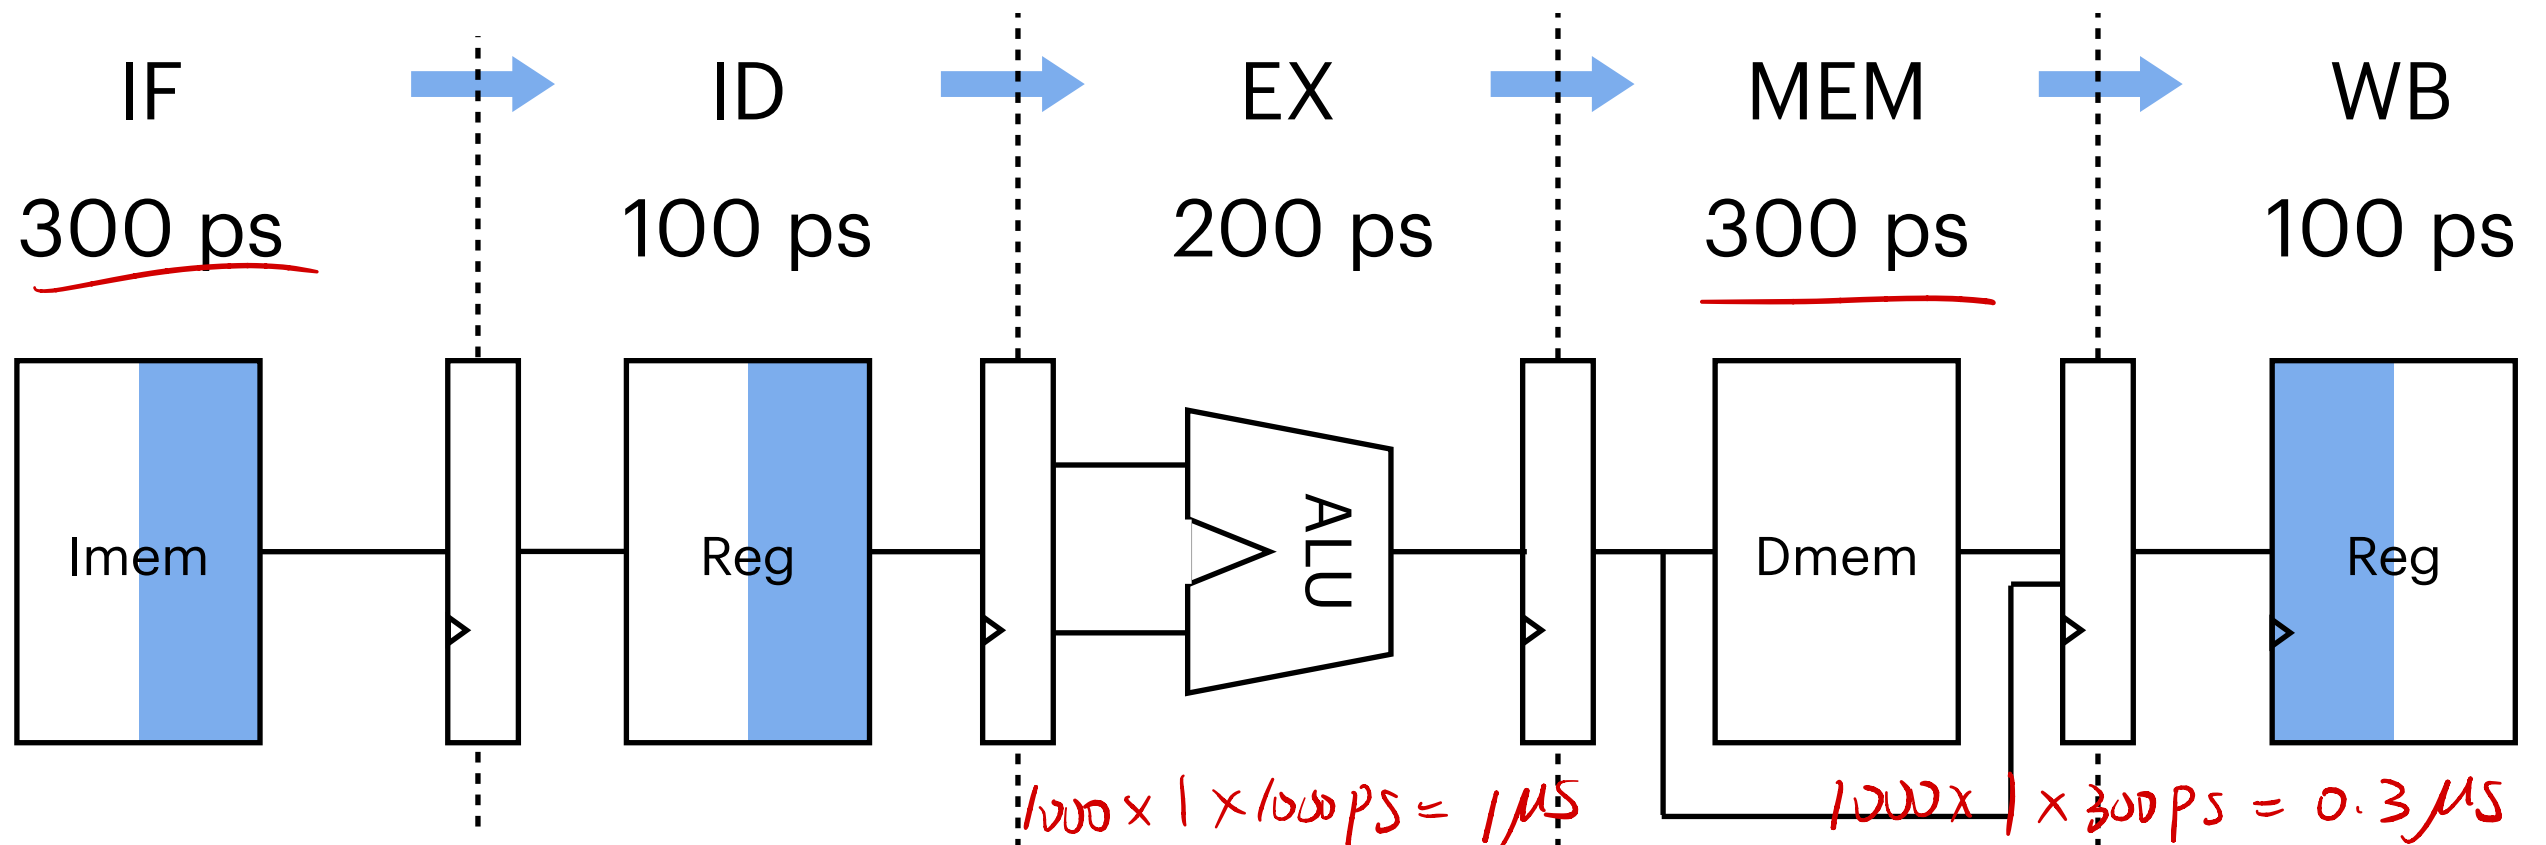


```
add t0,t1,t2
or t3,t4,t5
lw t0,8(t3)
addi t6,x0,1
xor t0,t0,t6
```

... ..

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}}$$

Insert Pipeline Registers



	Single-cycle	Pipelining
T_{clk}	1000 ps	300 ps
Clock rate/frequency	1 GHz	3.33 GHz
Instruction time	1000 ps	$5 \times 300 = 1500 \text{ ps}$
Clock per instruction	1	≈ 1
Relative speed	1x	3.33x
Hardware		

Question

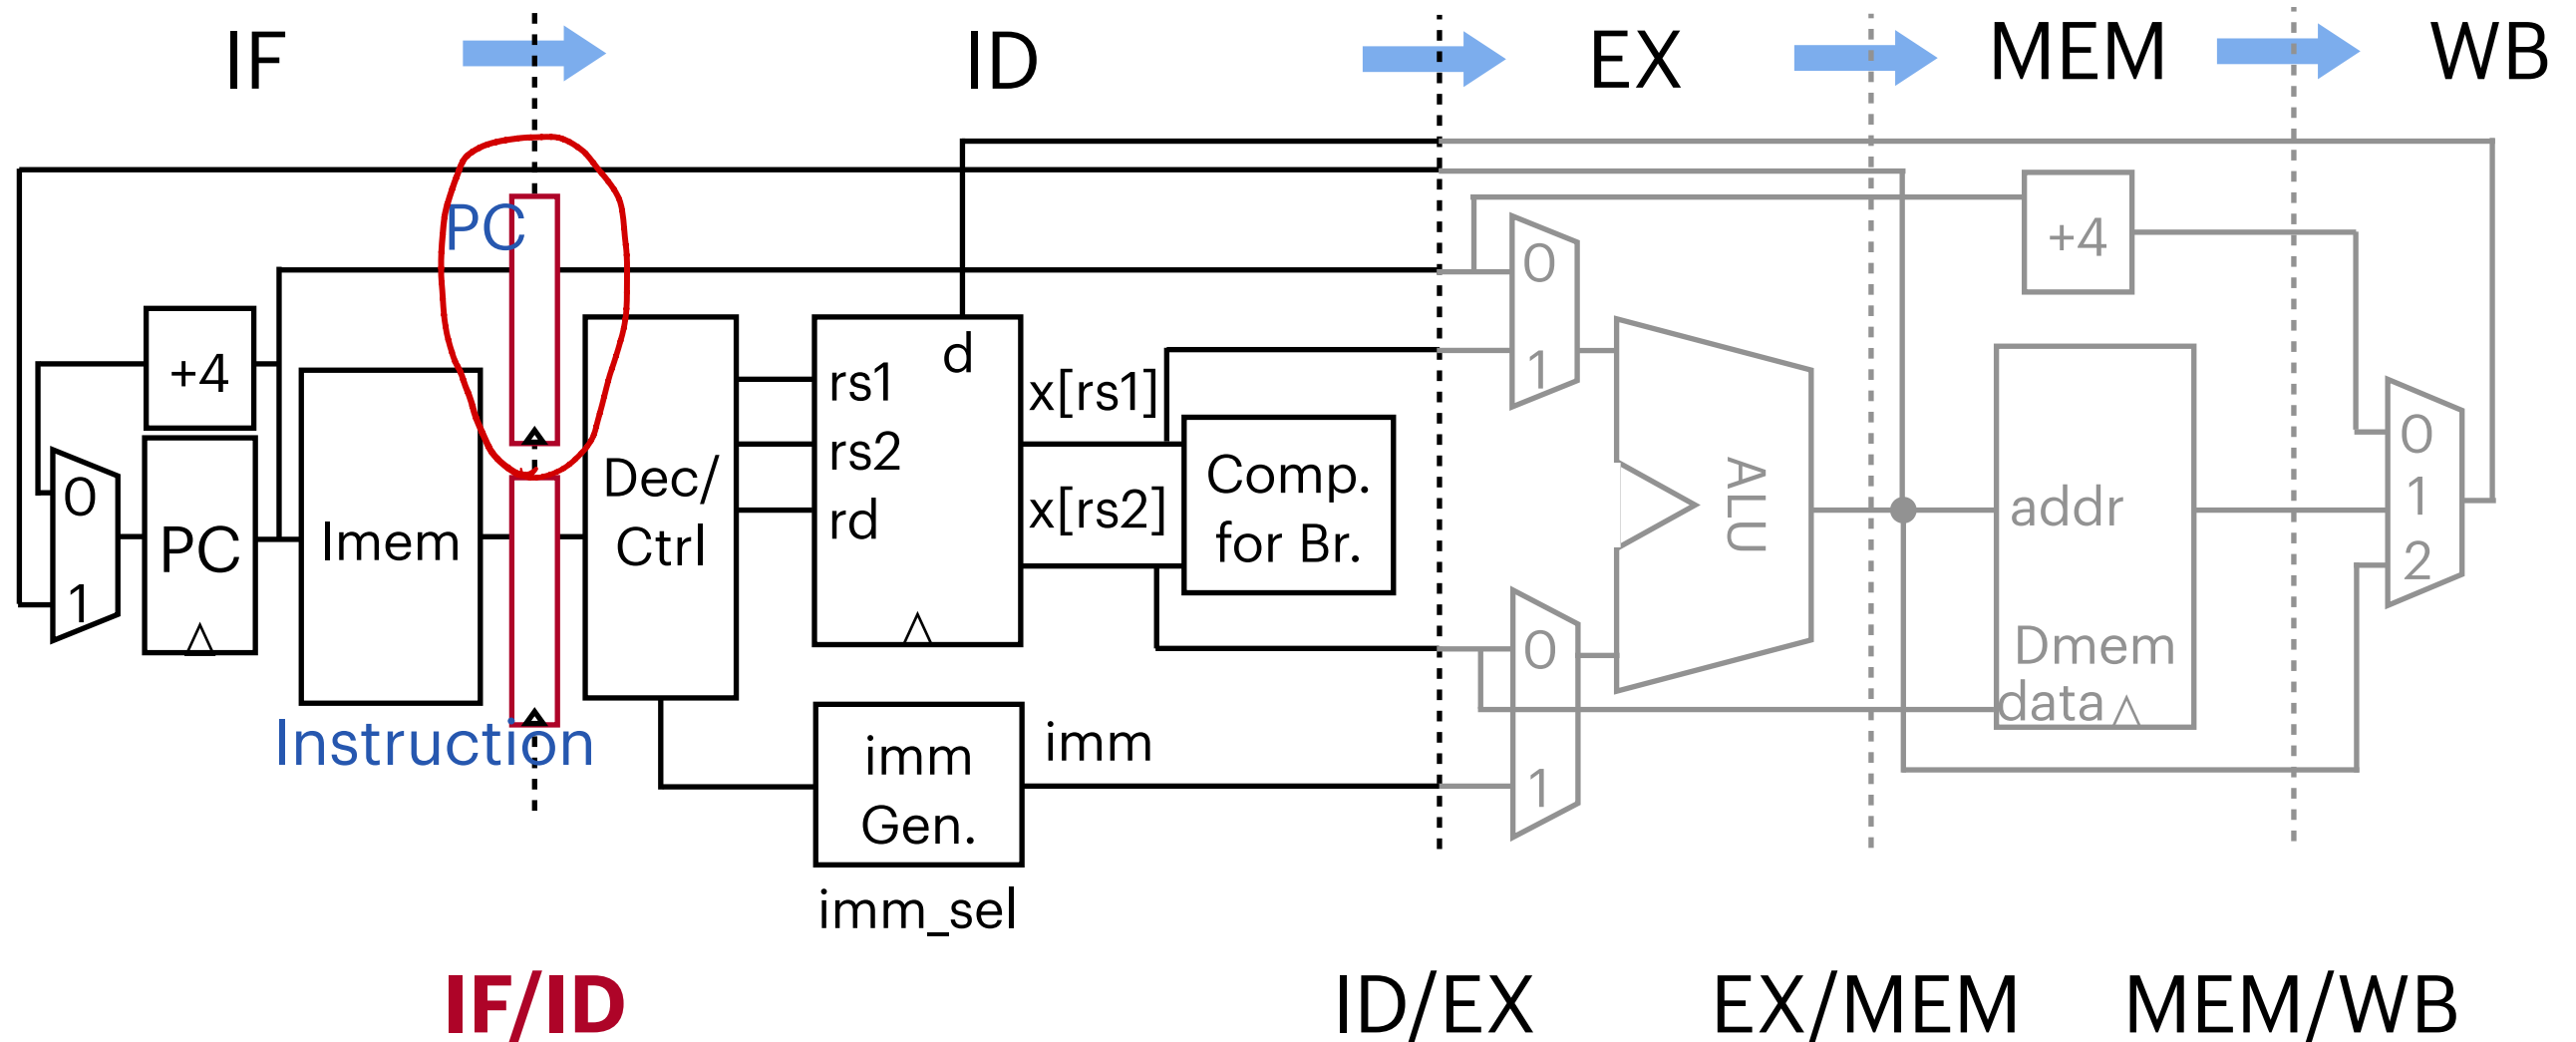
Combinational logic in some stages takes 200 ps and in some 100 ps. Clk-Q delay is 30 ps, and setup-time is 20 ps. What is the maximum clock frequency at which a pipelined design with 10 stages can operate?

X

$$T_{CLK} \leq 200 + 20 + 30 \\ = 250$$

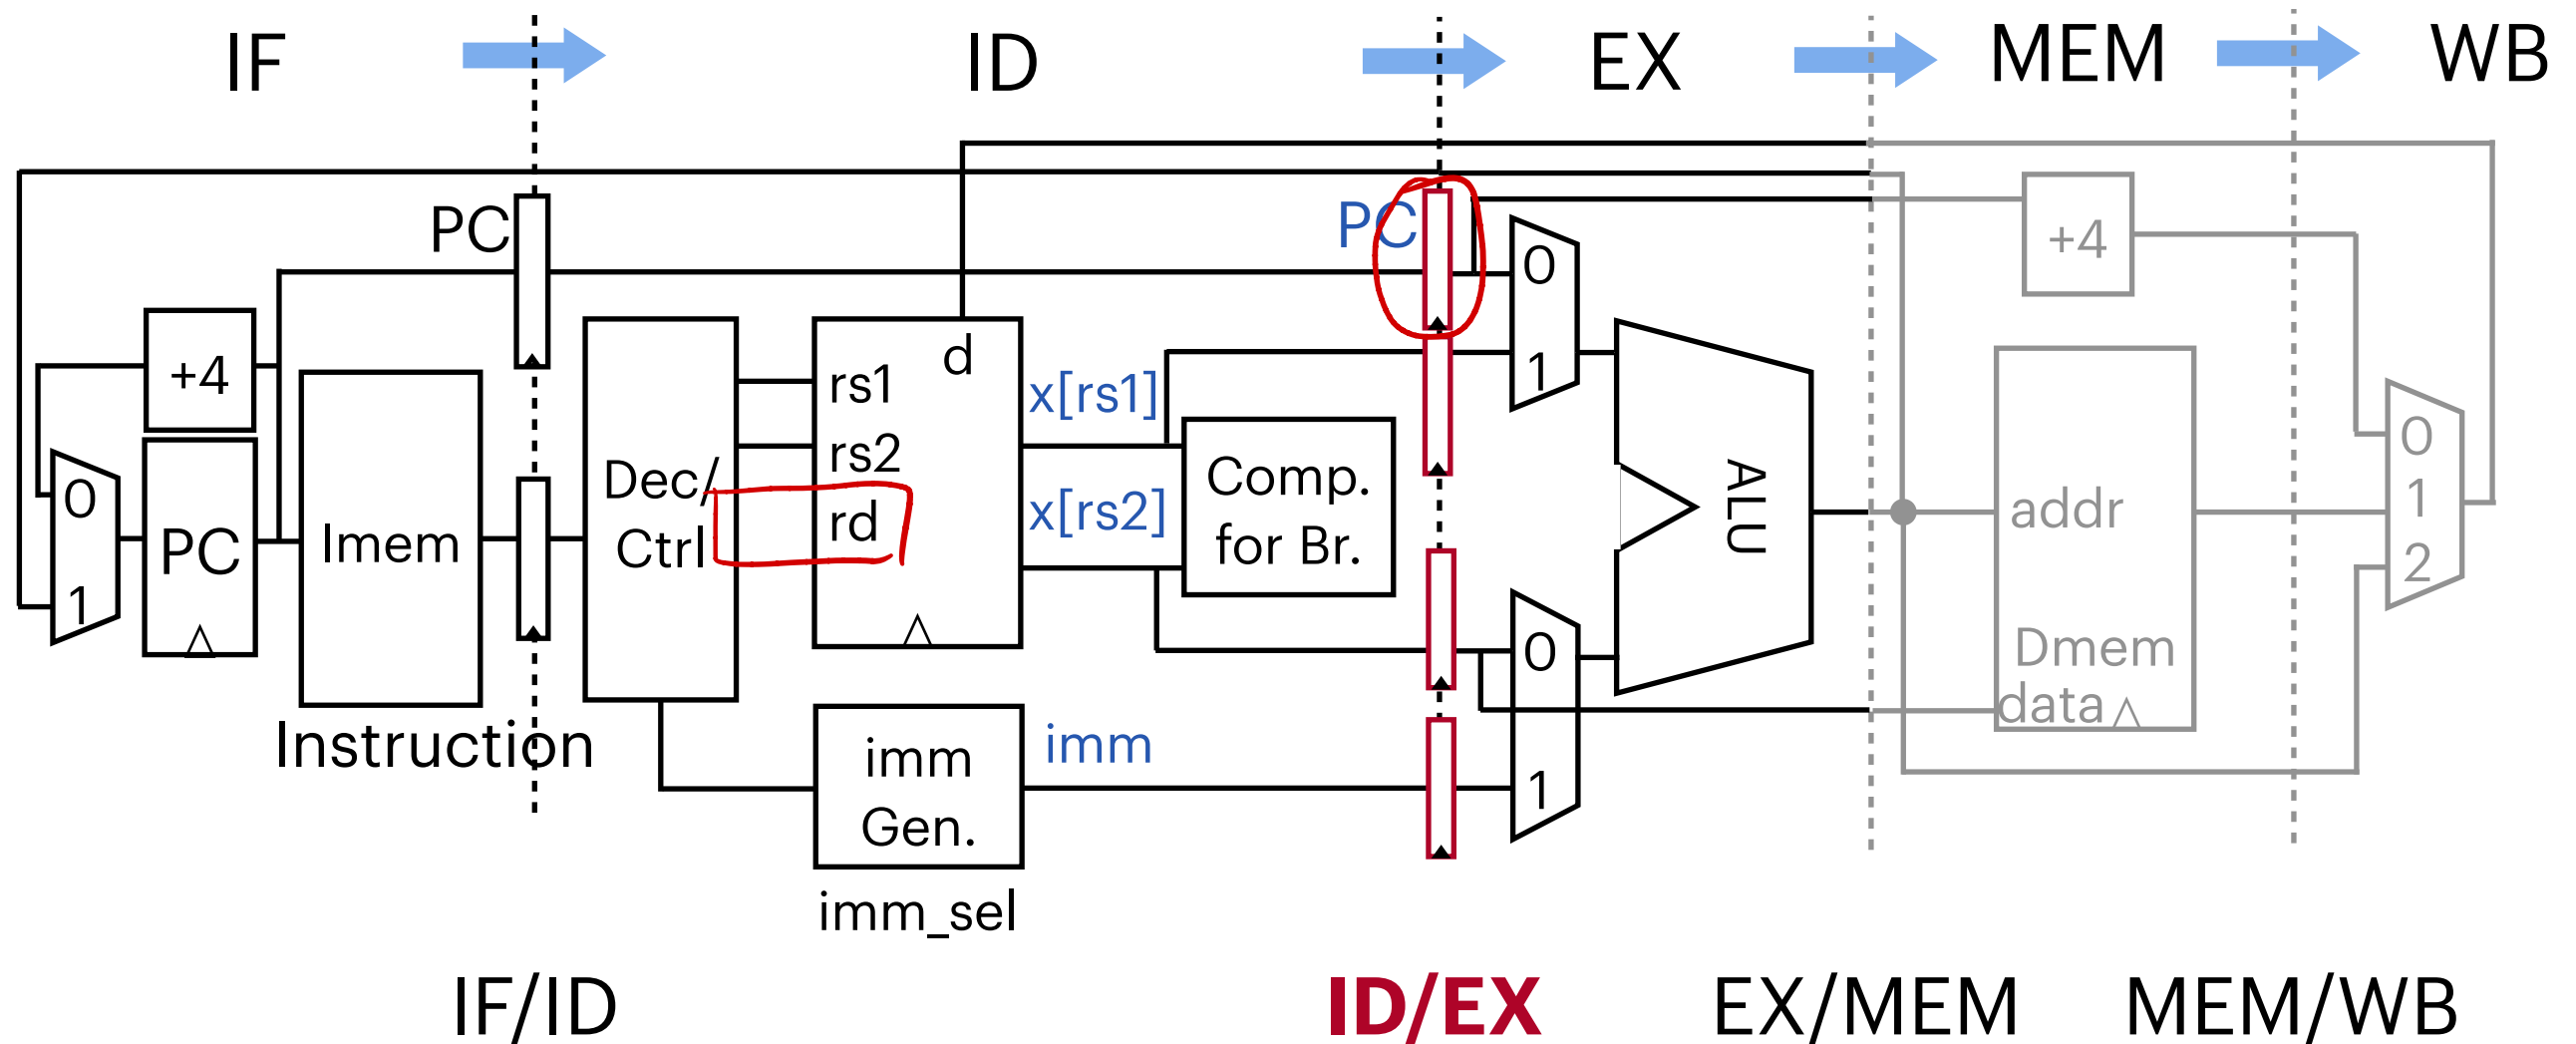
- A: 10GHz
- B: 5GHz
- C: 6.7GHz
- D: 4.35GHz
- ☒ E: 4GHz

Detailed Considerations—Datapath



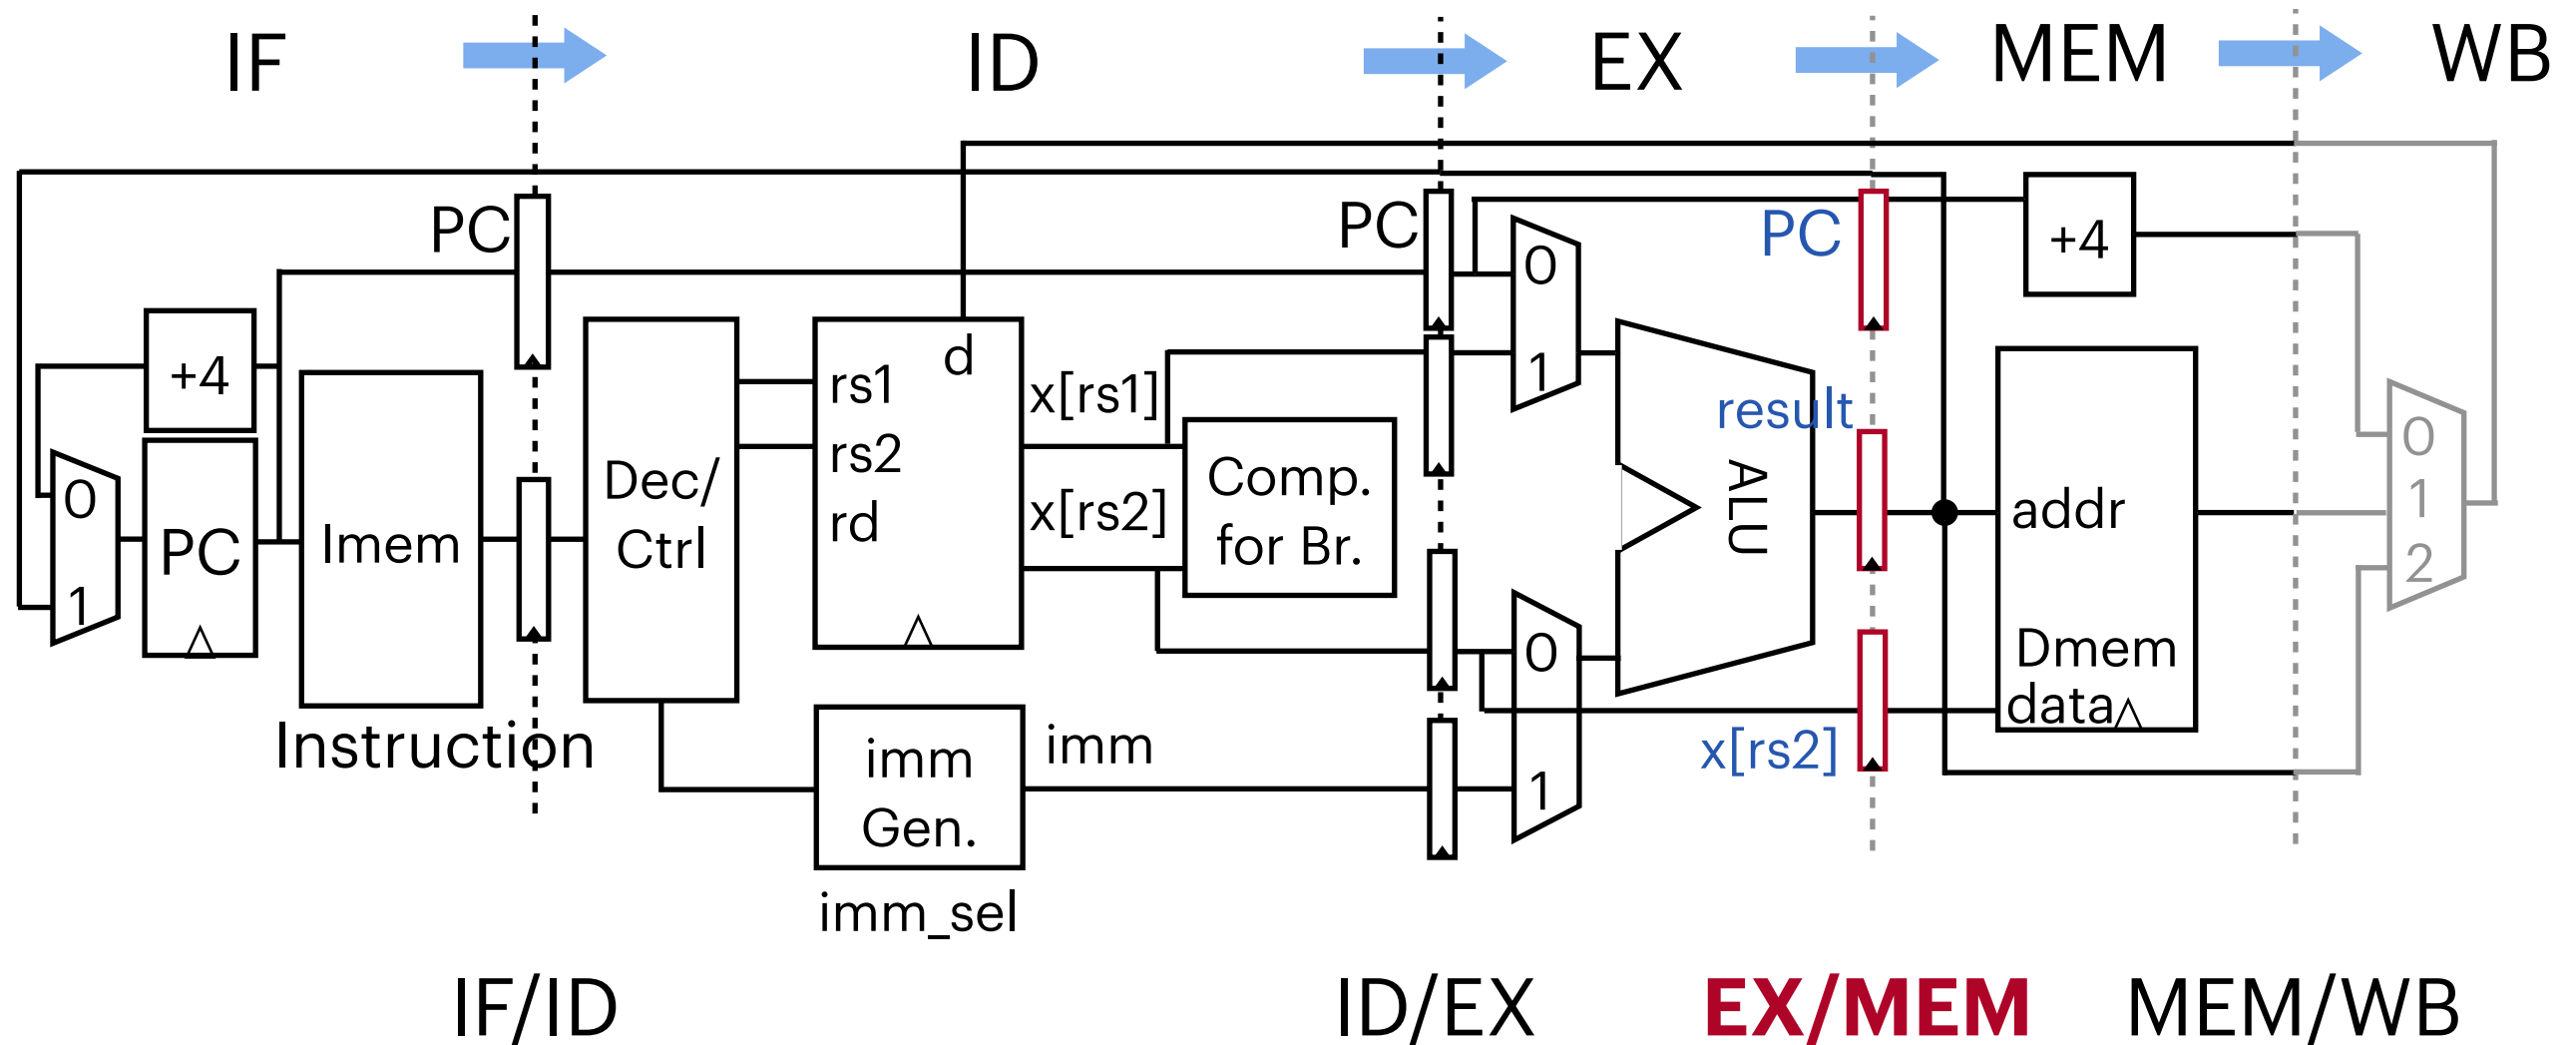
Use pipeline registers to carry instructions/data between stages

ID/EX Pipeline Registers



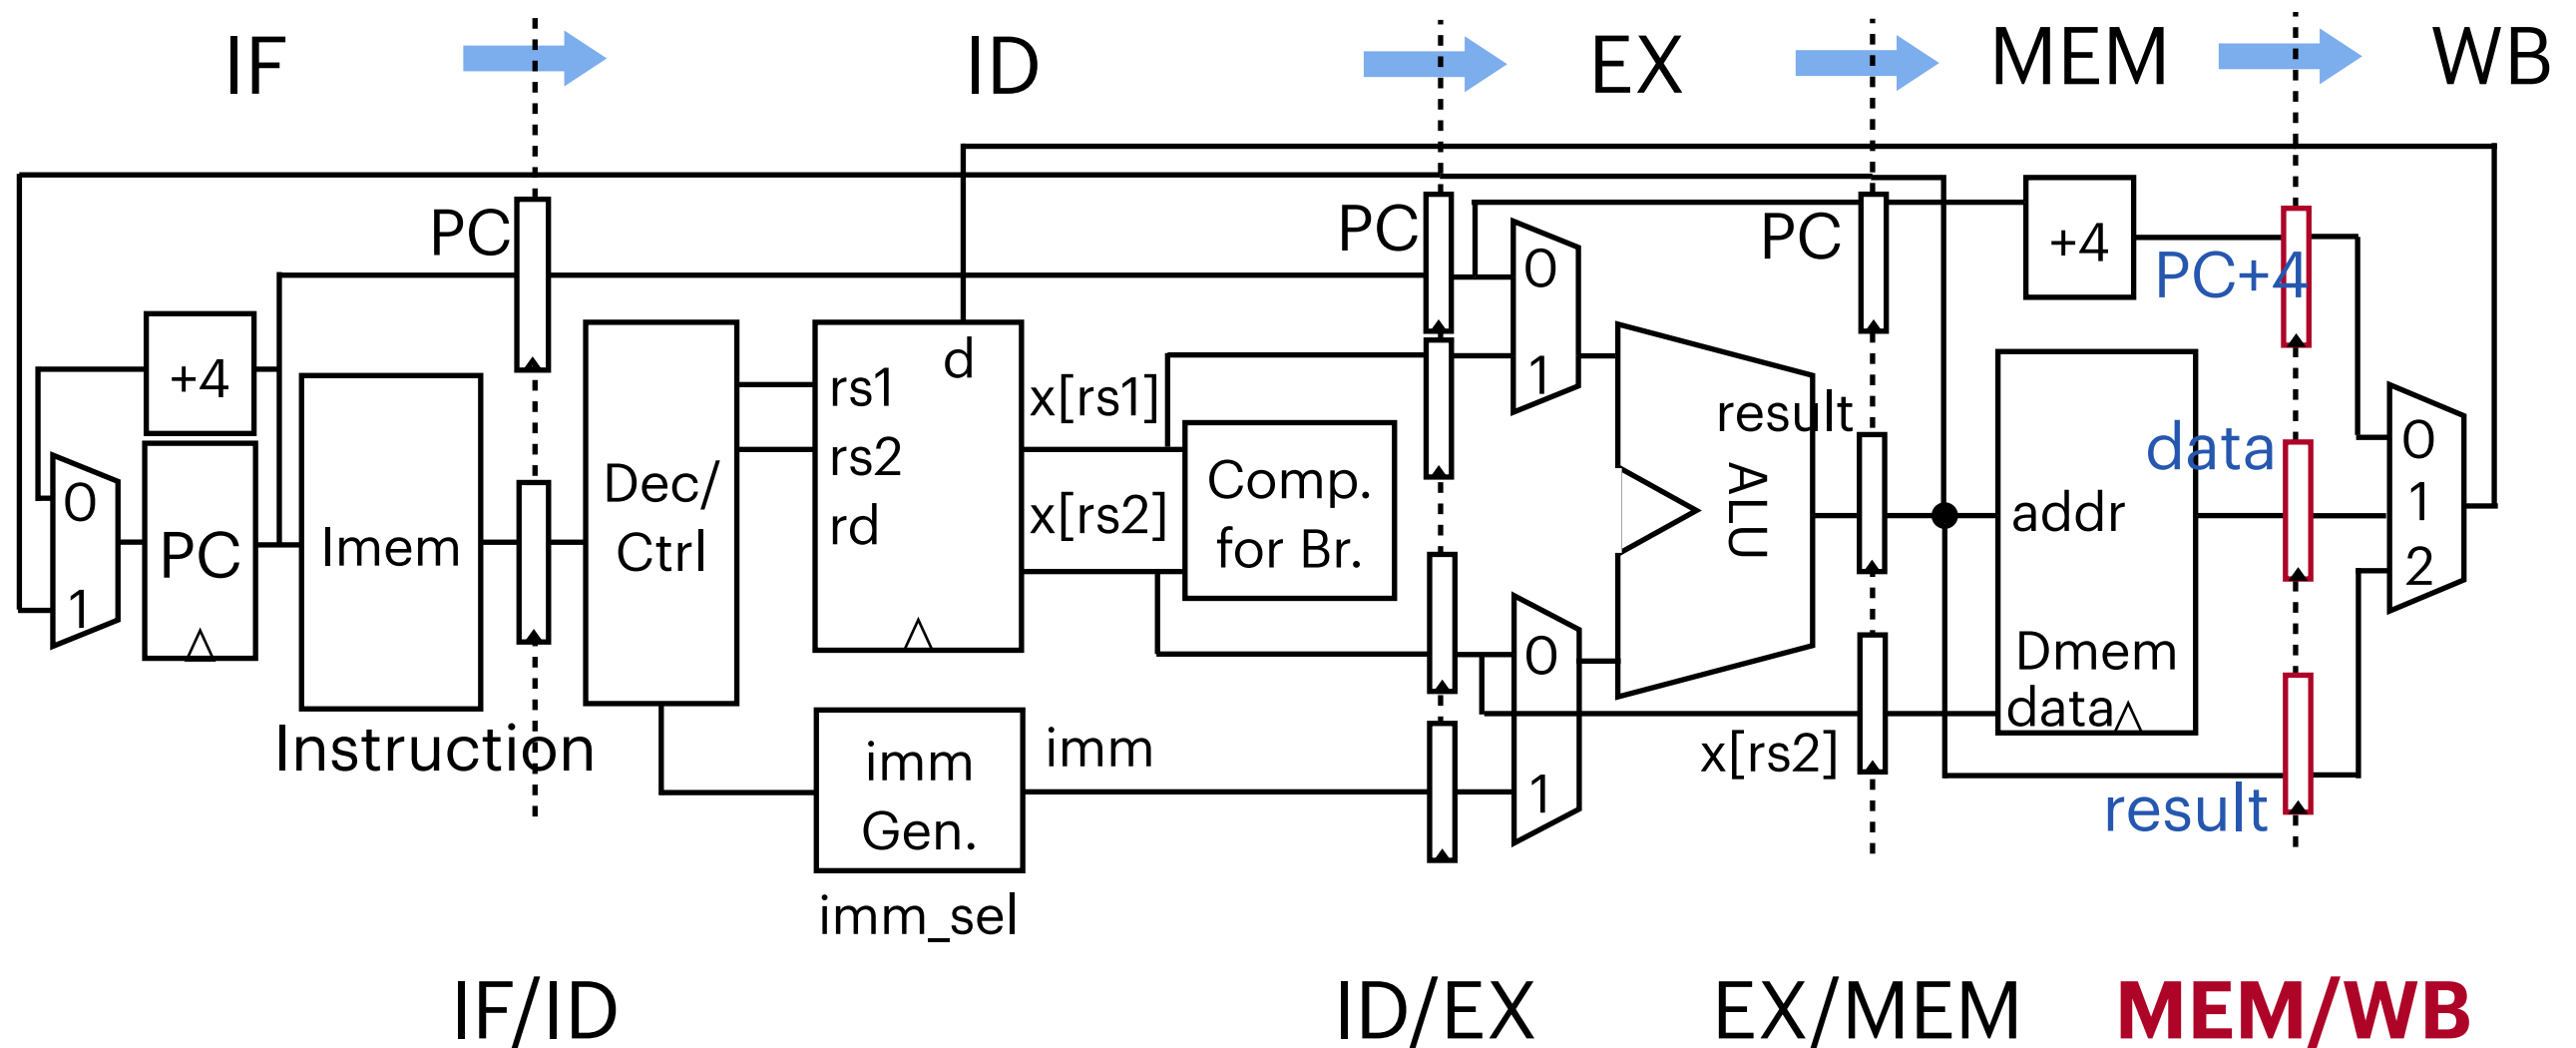
Use pipeline registers to carry instructions/data between stages

EX/MEM Pipeline Registers



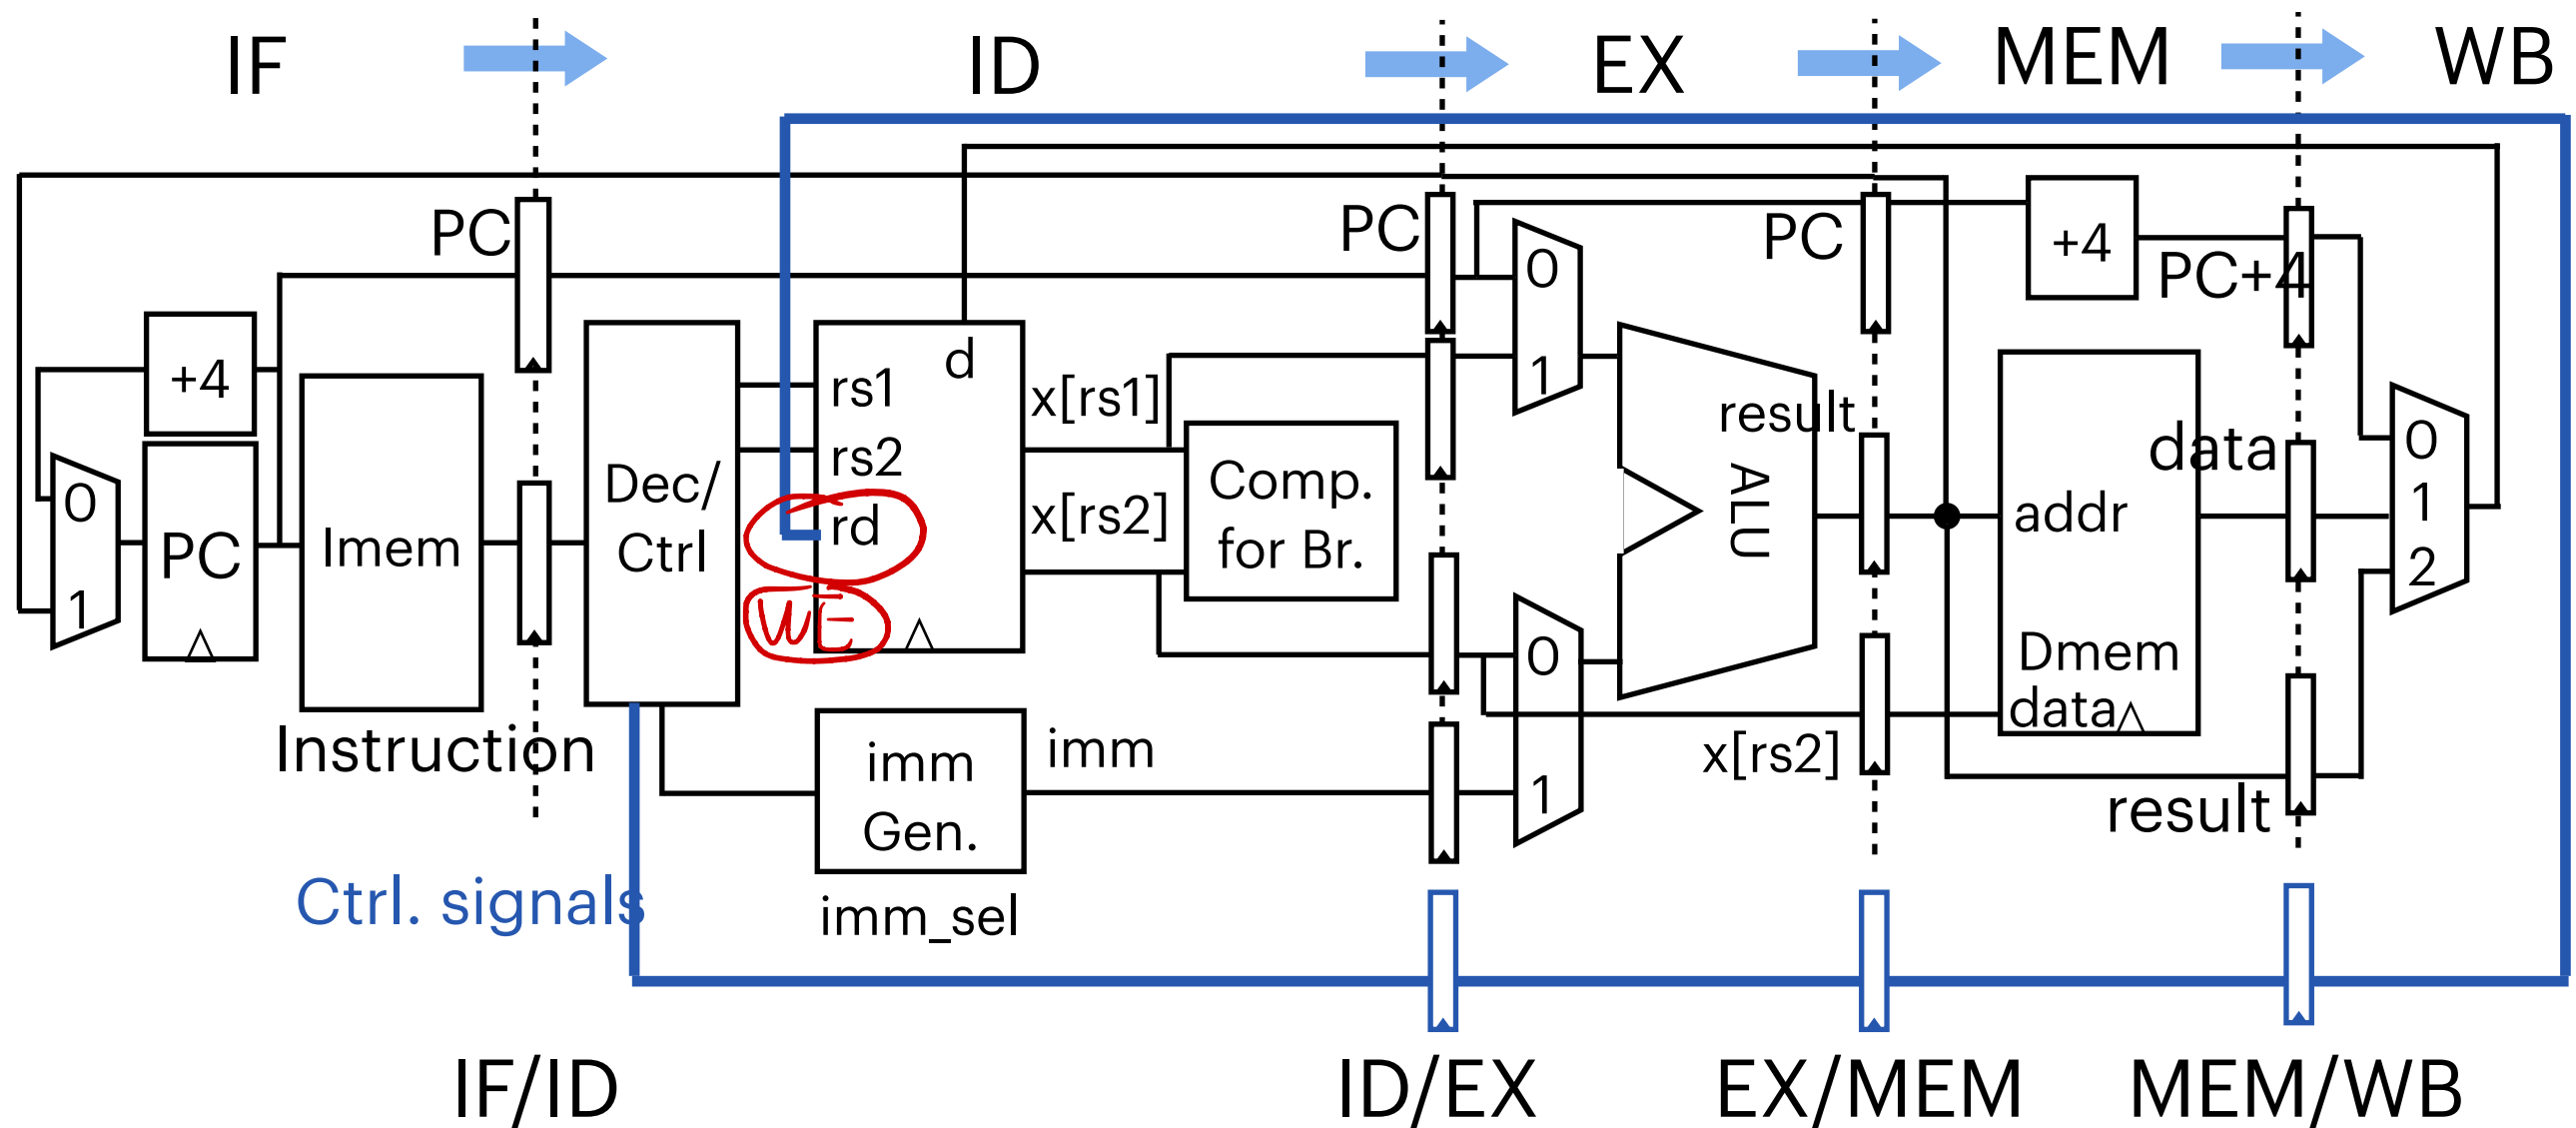
Use pipeline registers to carry instructions/data between stages

MEM/WB Pipeline Registers



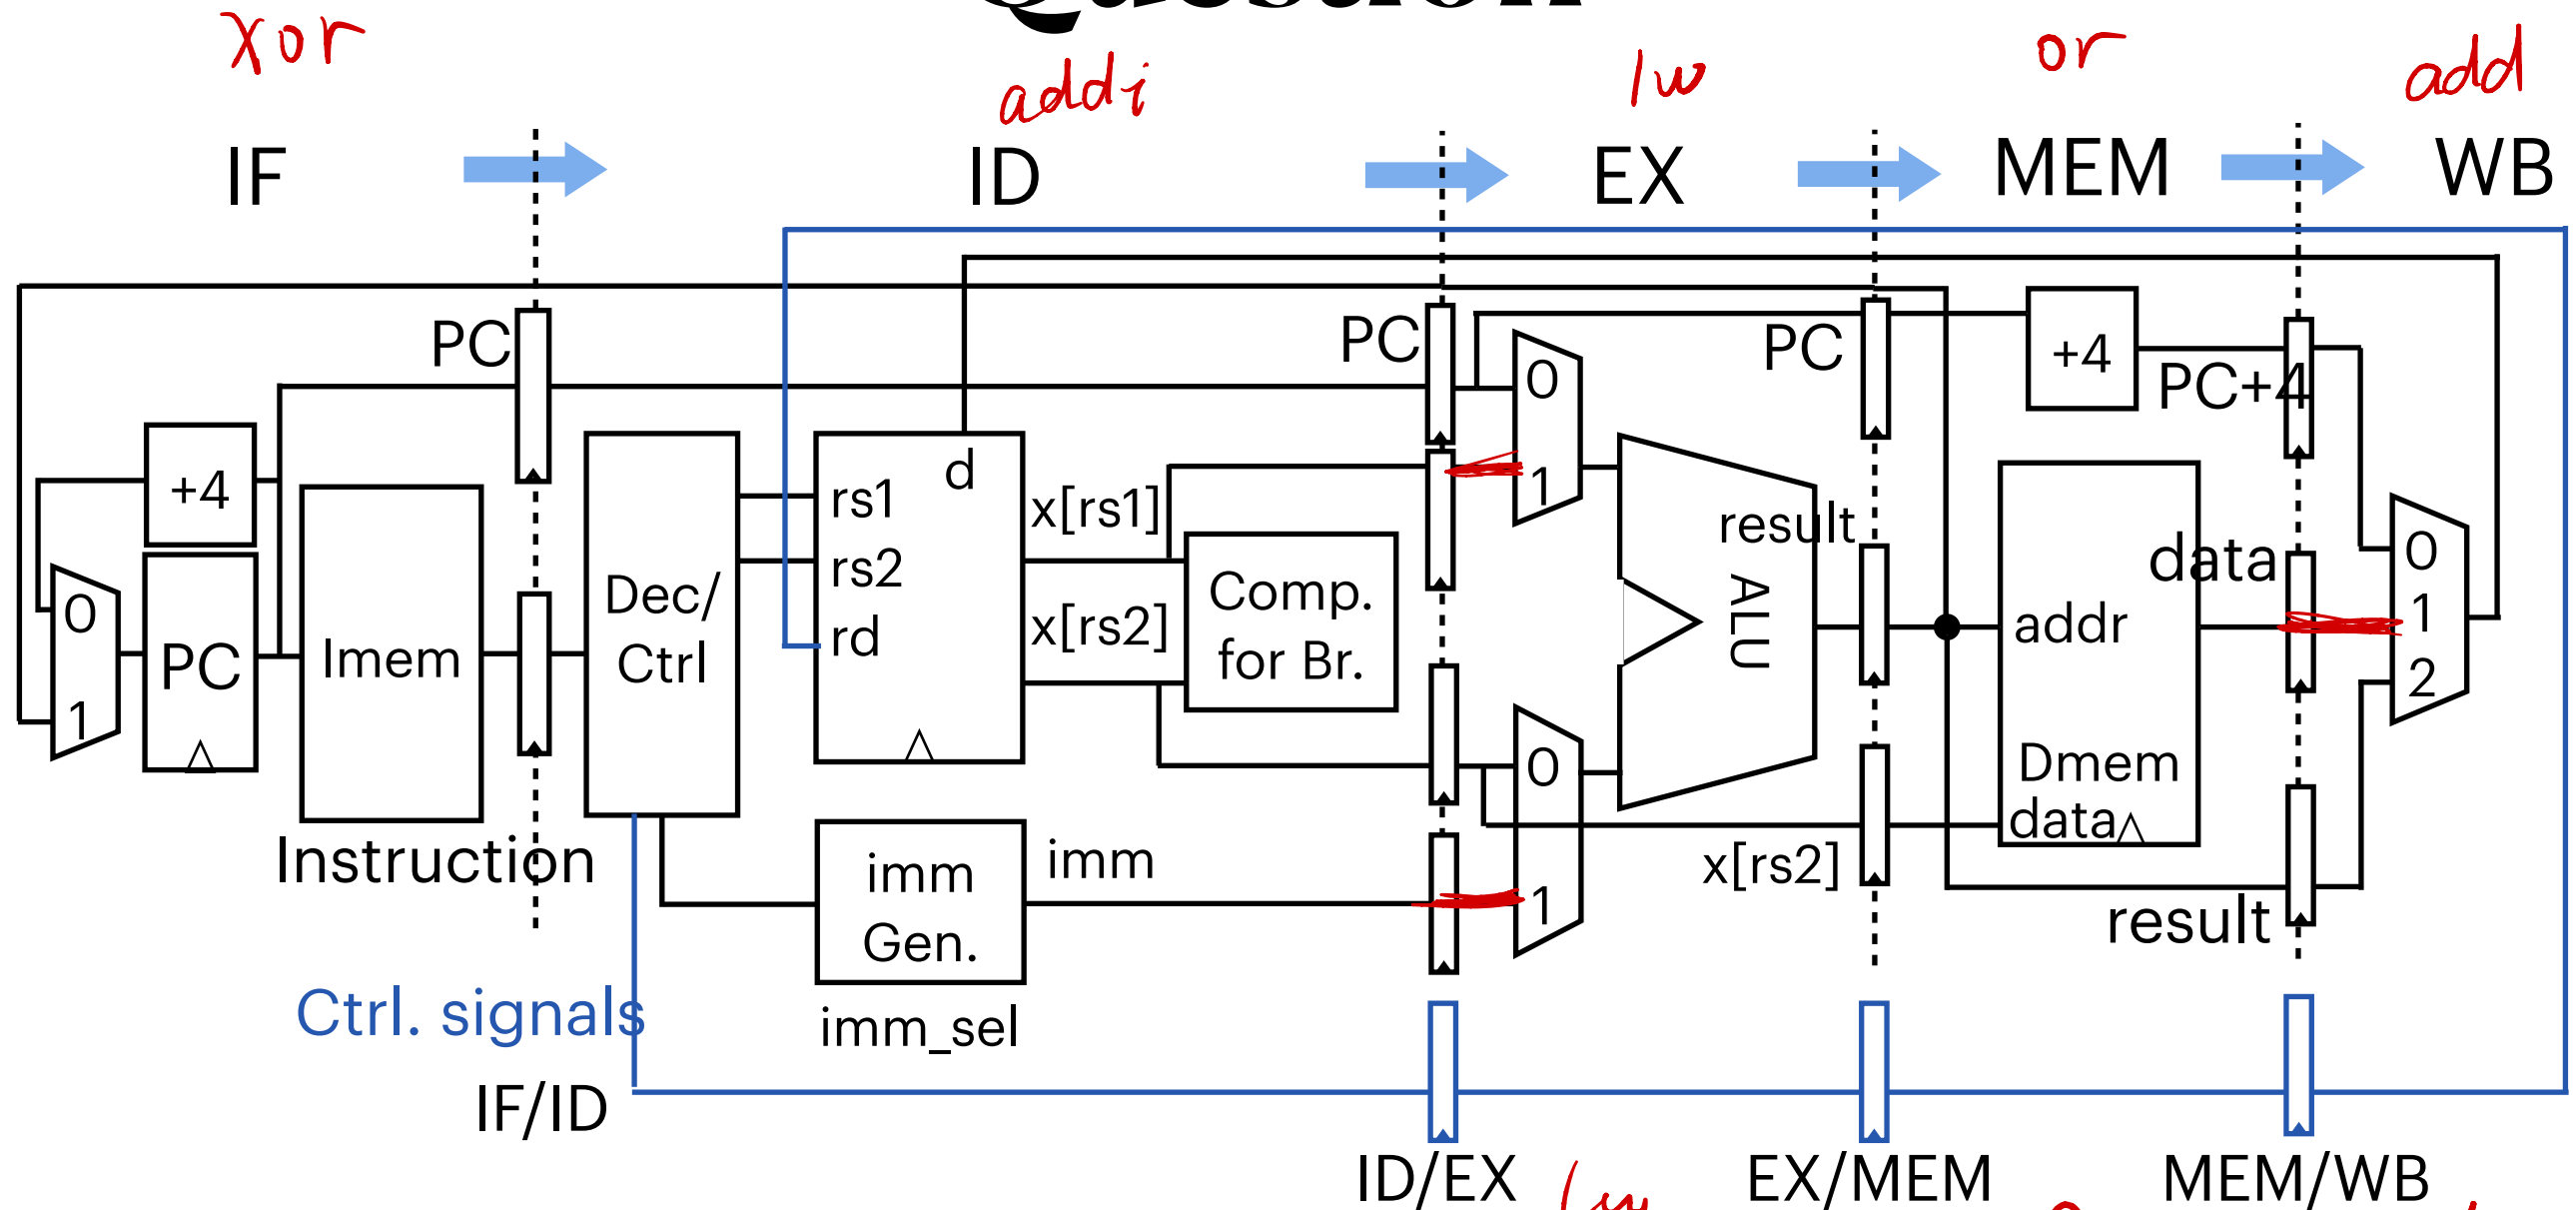
Use pipeline registers to carry instructions/data between stages

Control Signal Pipeline Registers



Use pipeline registers to carry instructions/data between stages

Question



`add t0, t1, t2`
`or t3, t4, t5`
`lw t0, 8(t3)`
`addi t6, x0, 1`
`xor t0, t0, t6`

Q: When add instruction is "WB-ing", which operands are selected for ALU simultaneously?

A: PC; immediate.

B: PC; t4

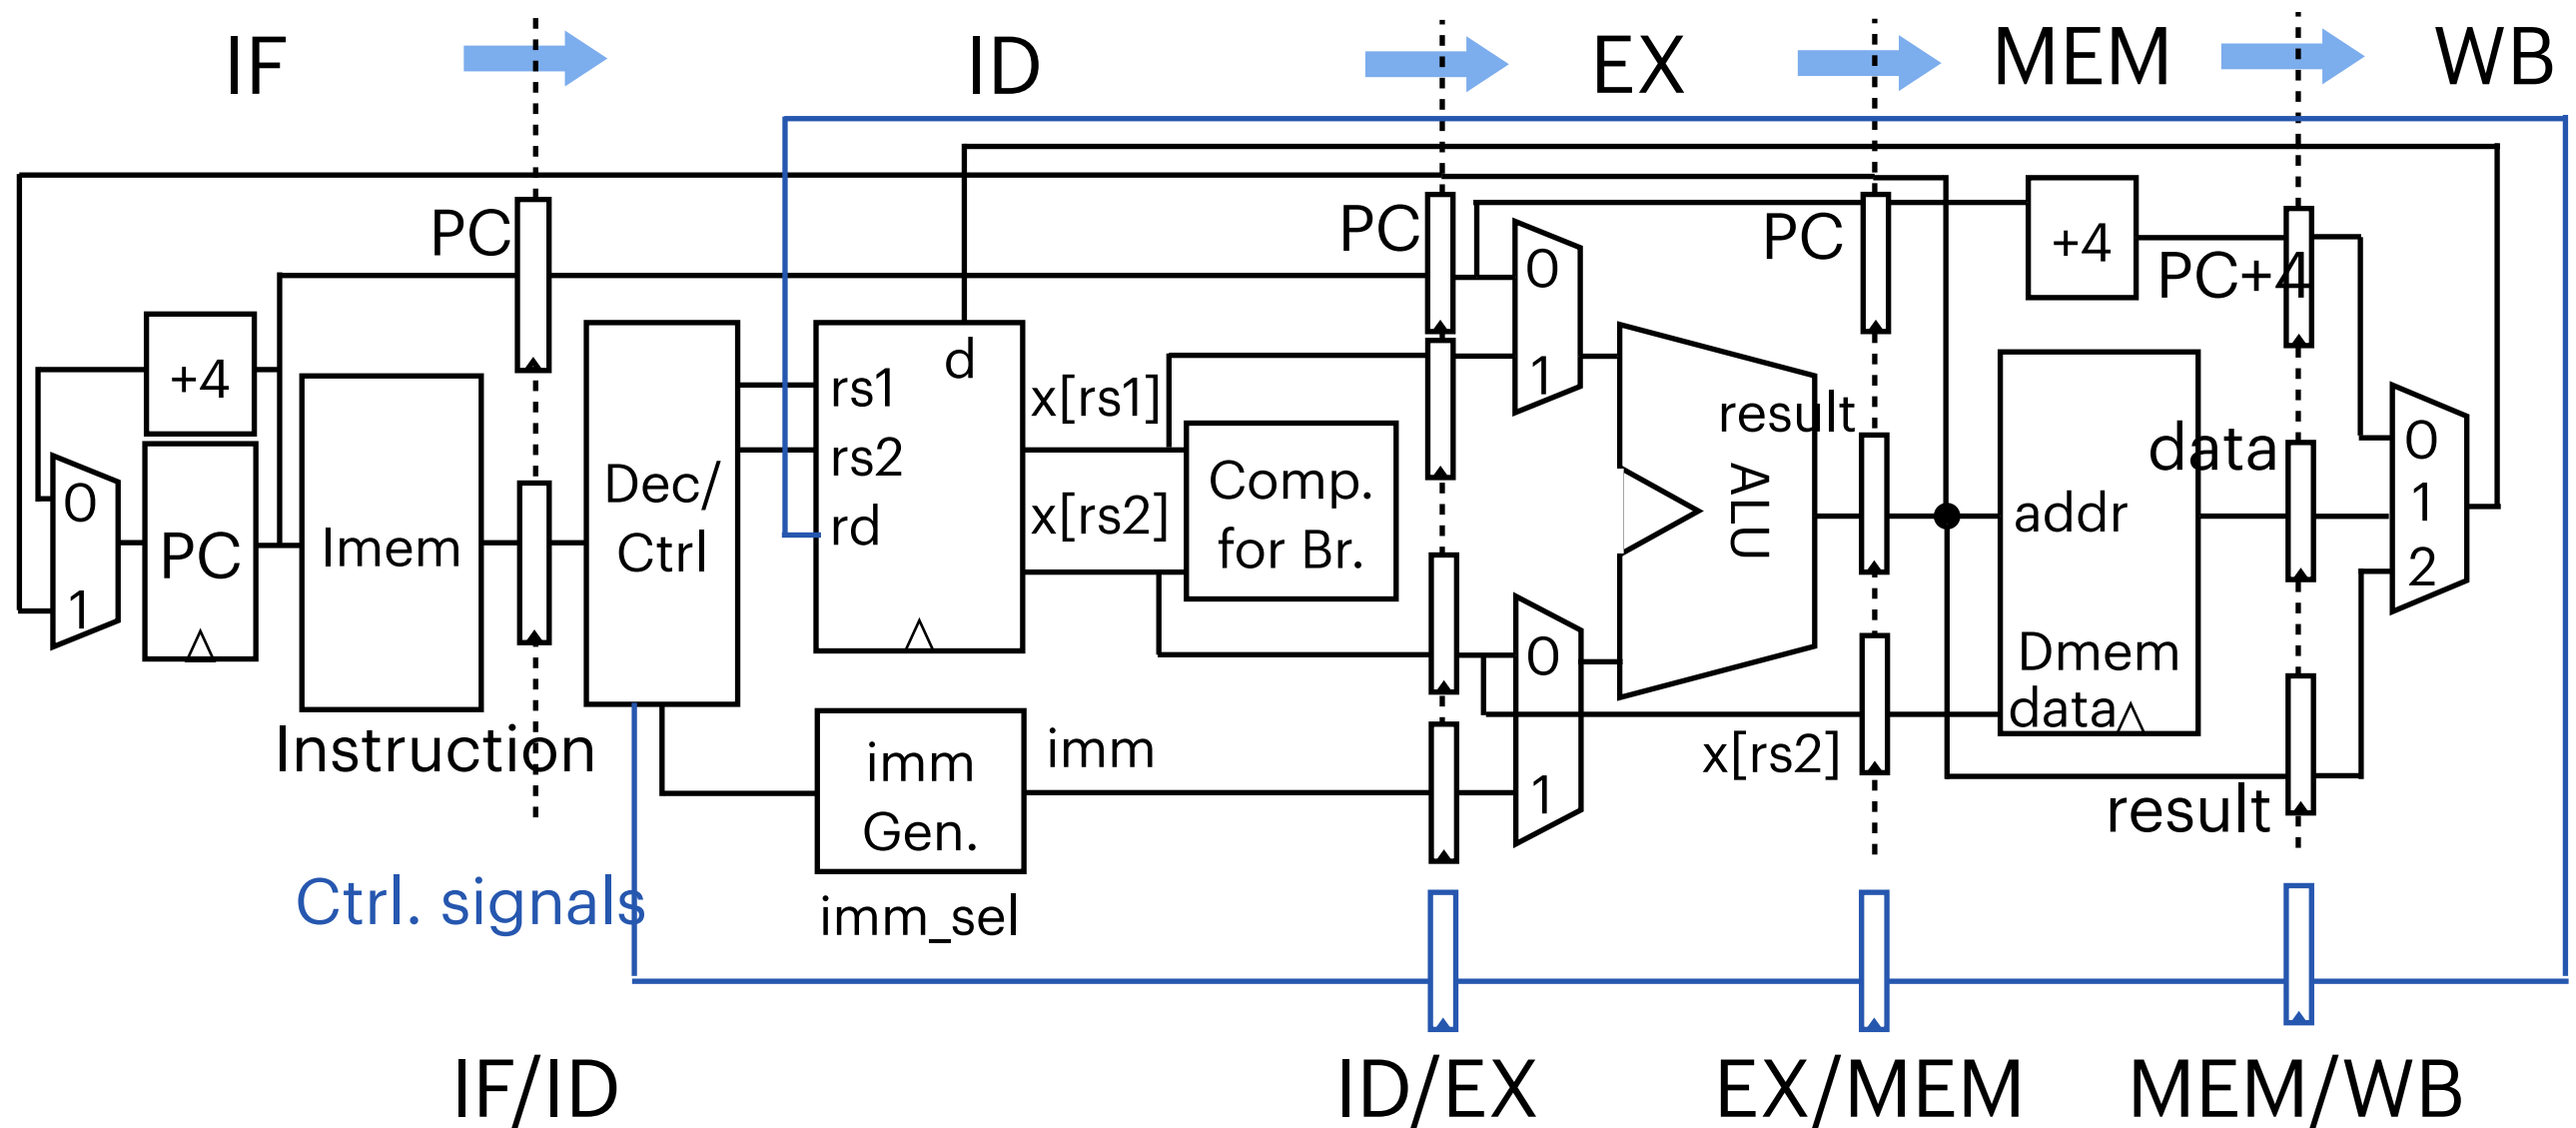
☒ C: t3; immediate.

D: t4; t5

E: Something else.

Without considering correct execution

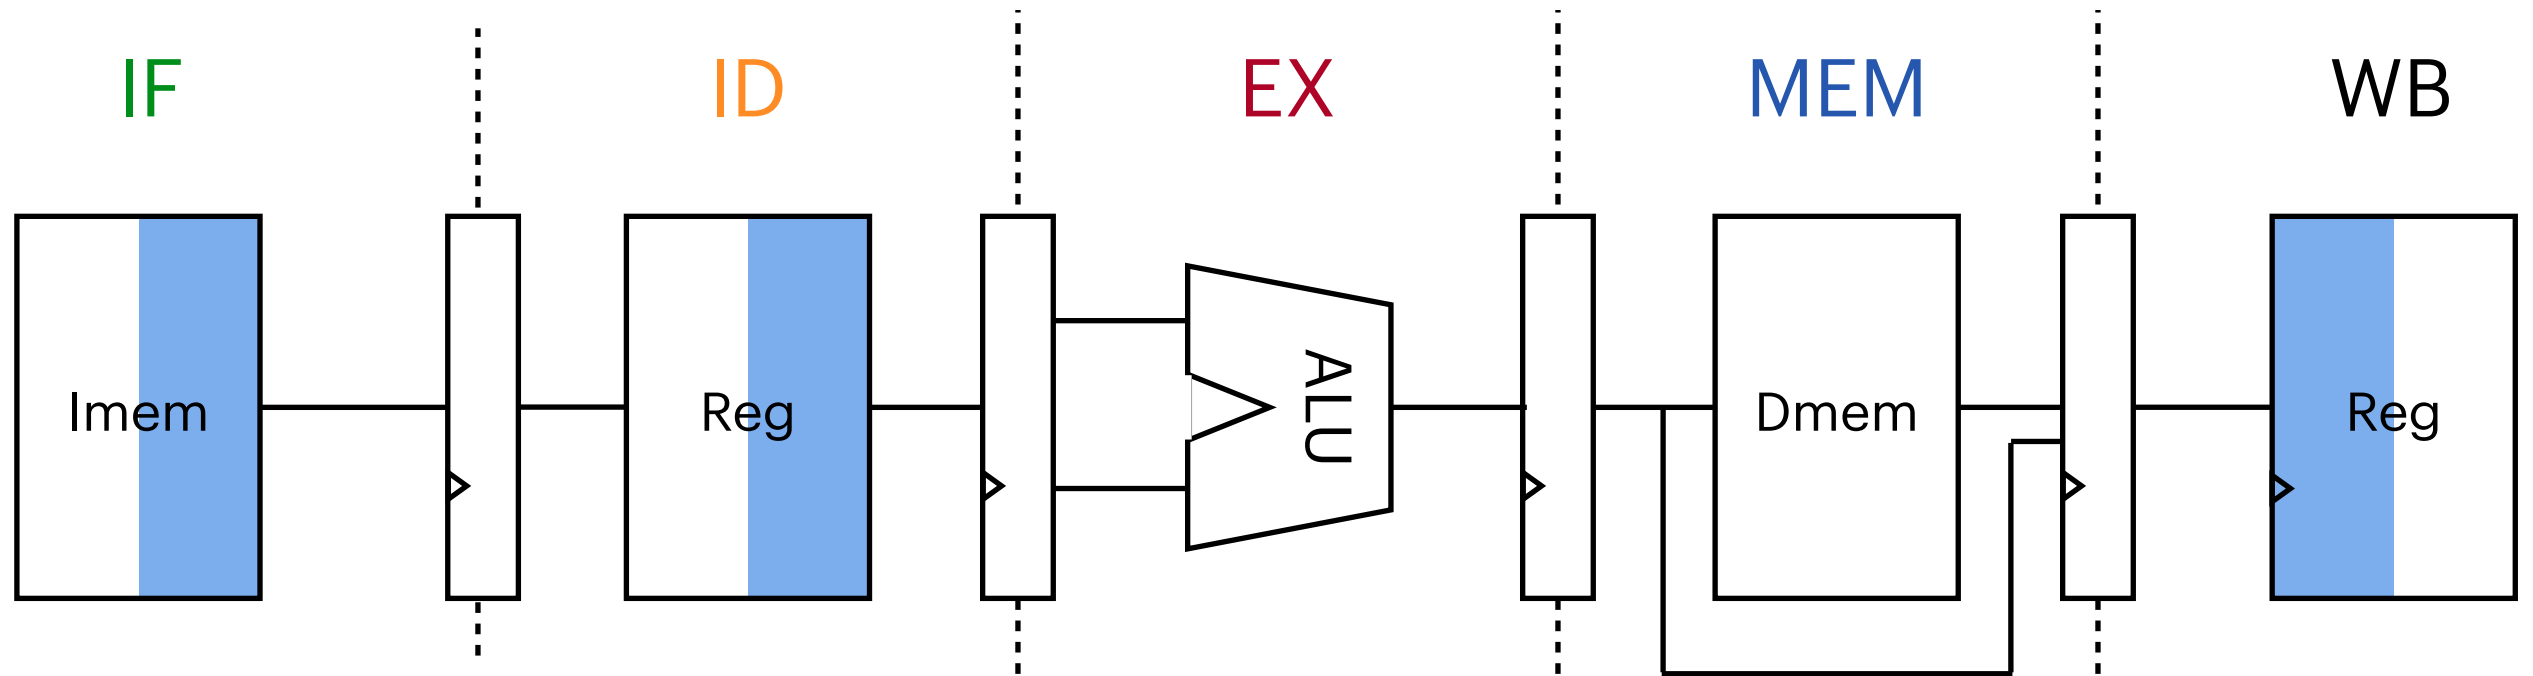
Up to Now



Pipeline Hazards Ahead!!!

冲突

冲突 Pipeline Hazards Ahead!!!



Clock cycle 1

IF add

Clock cycle 2

IF lw

ID add

Clock cycle 3

IF or

ID lw

EX add

Clock cycle 4

IF sw

ID Simultaneous

MEM add

Clock cycle 5

IF sll

ID sw

EX or

MEM lw

WB add⁵

add t0,t1,t2

lw t0,8(t3)

or t3,t4,t5

sw t0,4(t3)

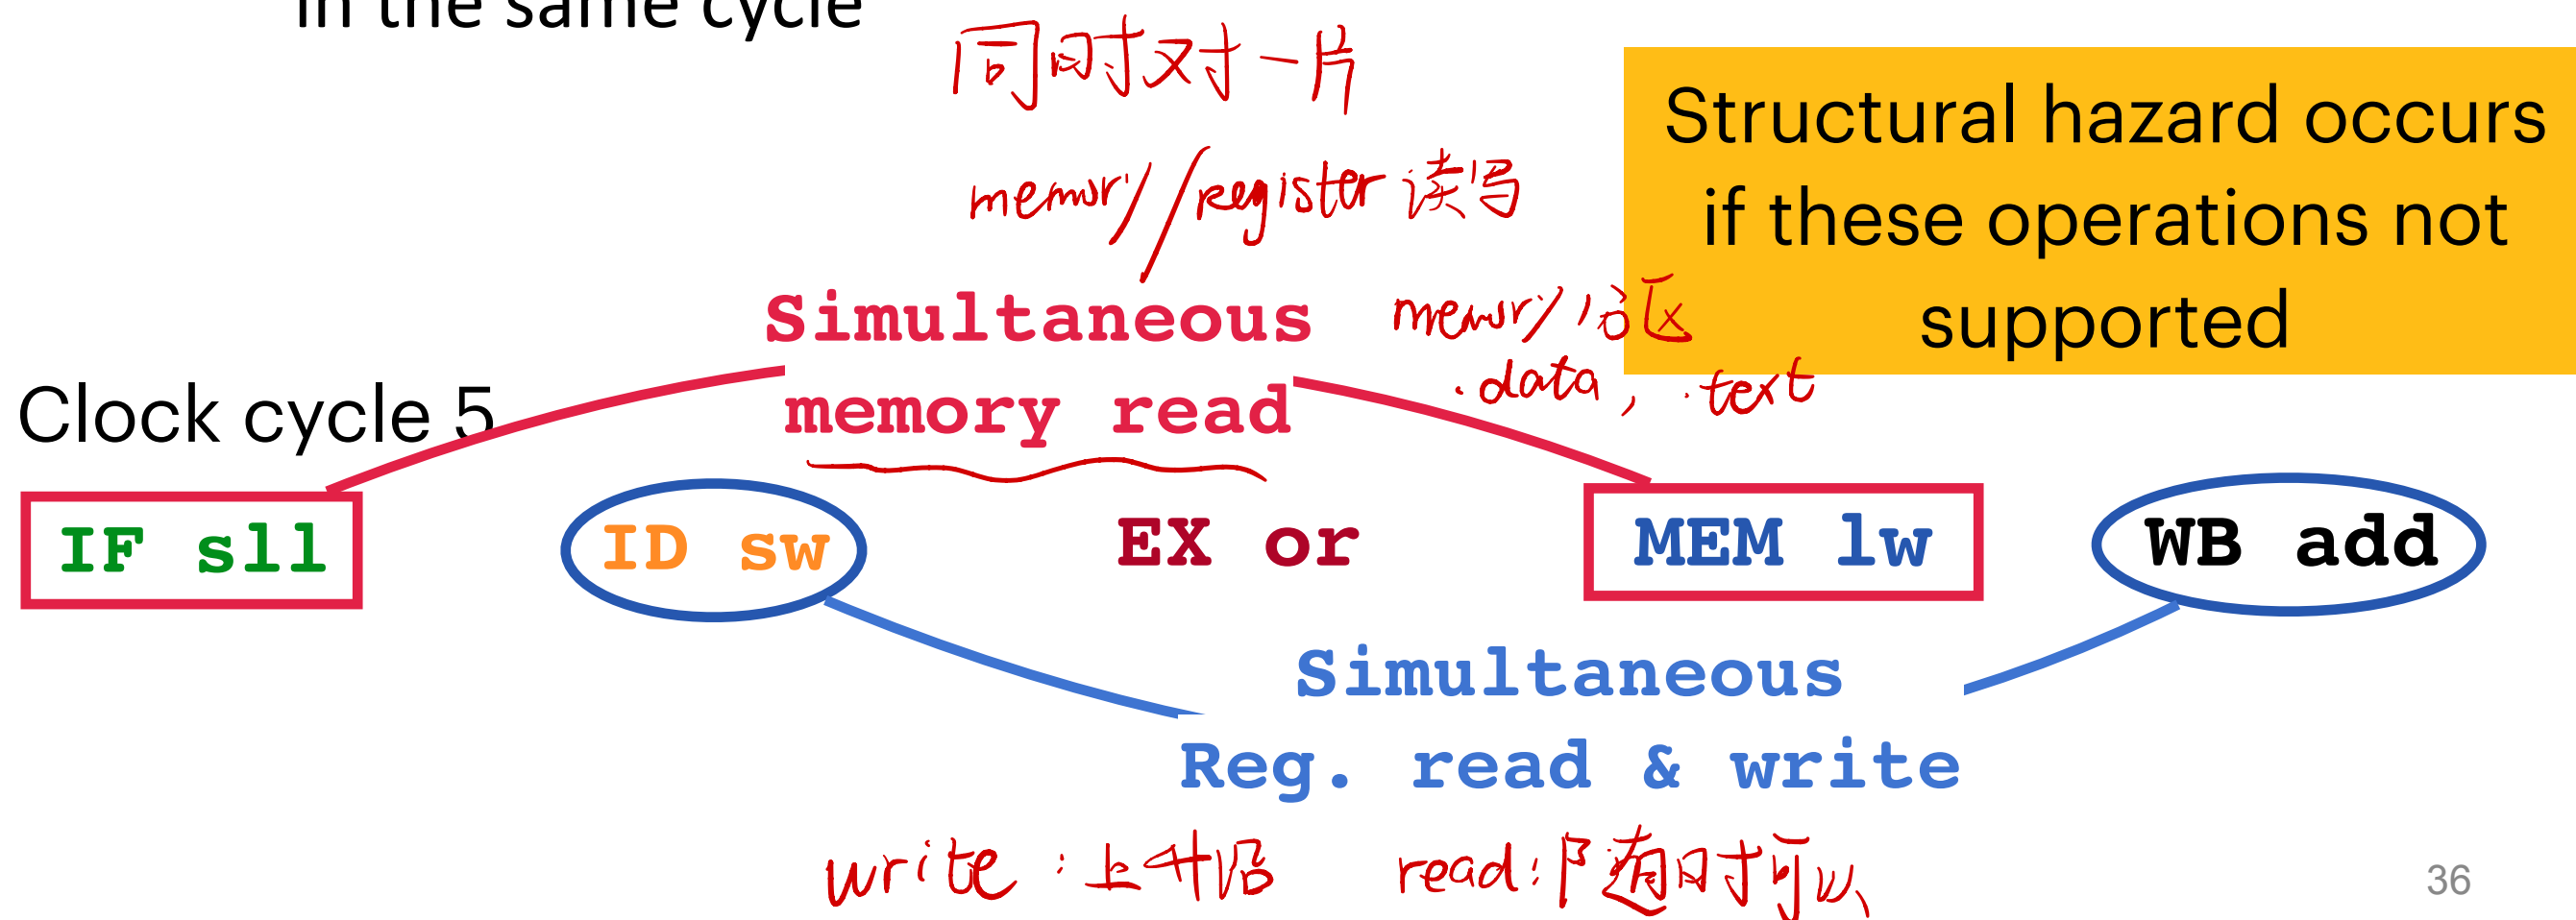
sll t6,t0,t3

Three Types of Pipeline Hazards

A **hazard** is a situation in which a planned instruction cannot execute in the “proper” clock cycle.

1. Structural hazard:

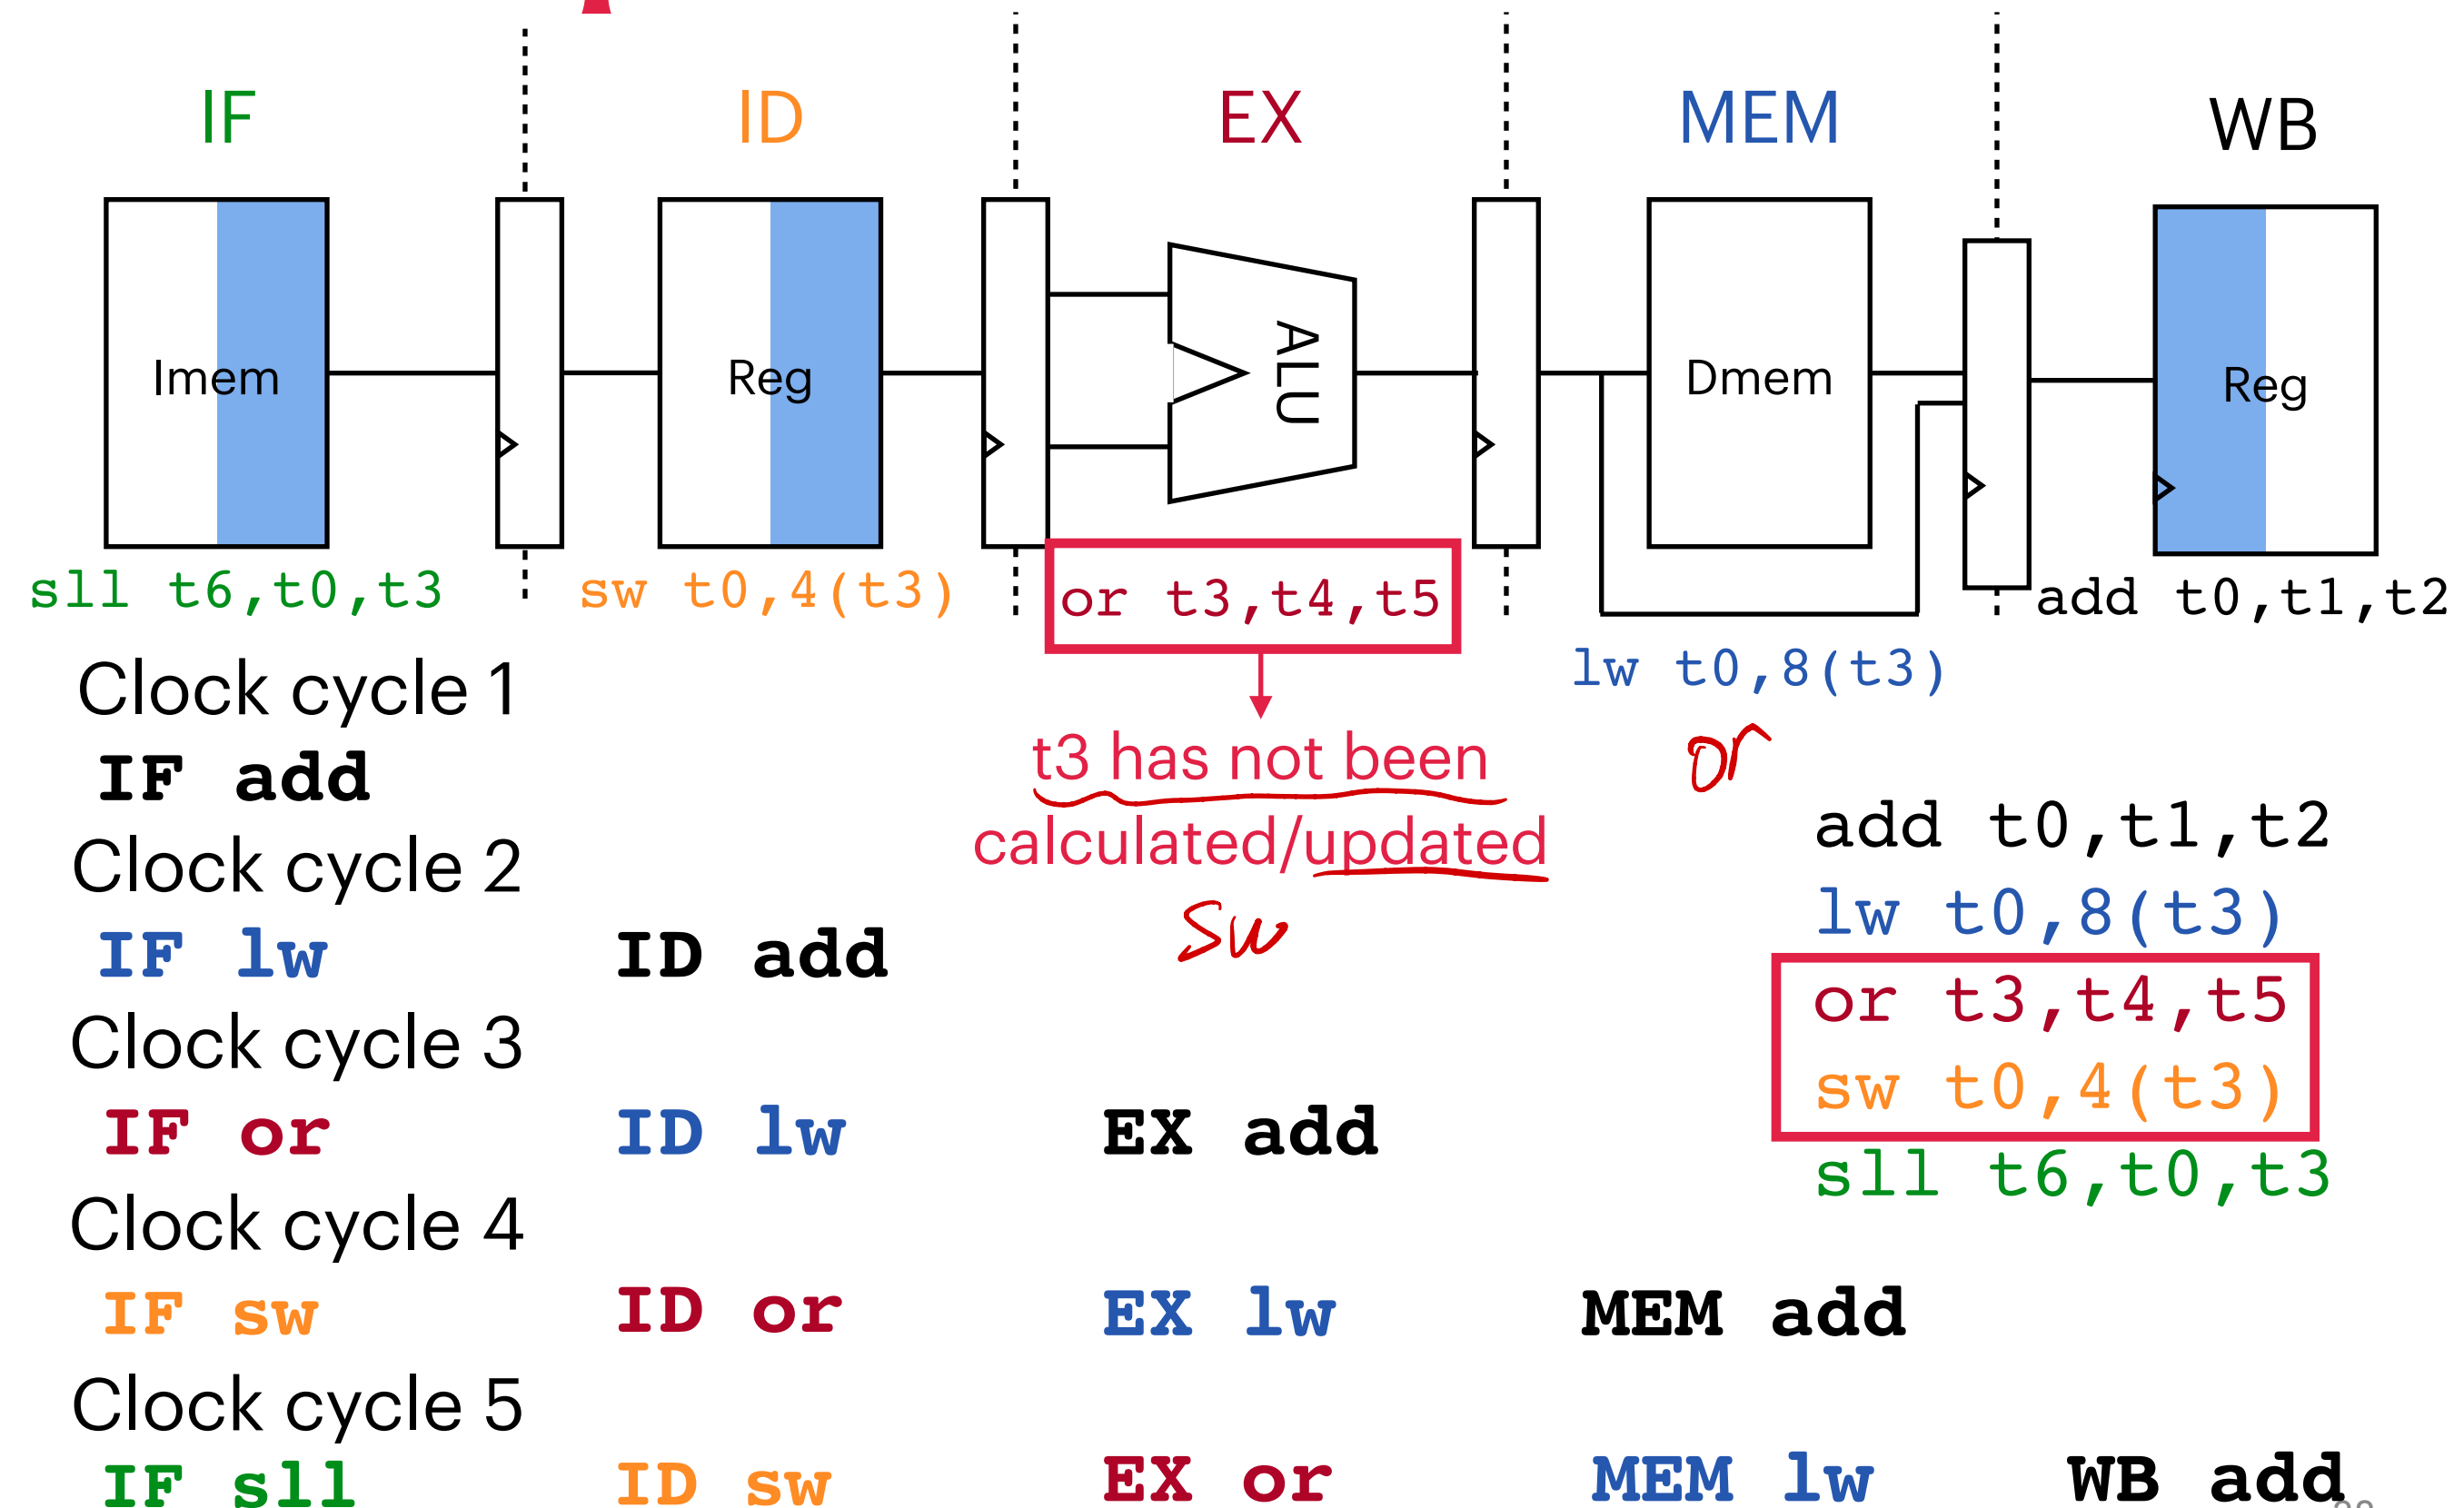
- Hardware does not support access across multiple instructions in the same cycle



Structural Hazards

- Structural hazard:
 - Hardware does not support access across multiple instructions in the same cycle
- Occurs when multiple instructions compete for access to a single physical resource
- Solution 1:
 - Instructions take turns using the resource
 - **Stall** while the resource is busy
- Solution 2:
 - Can always avoid structural hazards by adding more HW 硬件
 - In our current design, *structural hazards are not an issue*
 - Separated instruction and data memory (cache)
 - Synchronized write & unsynchronized read, support simultaneous read/write; Reg. has two read ports

Pipeline Hazards Ahead!!!



Three Types of Pipeline Hazards

A hazard is a situation in which a planned instruction cannot execute in the “proper” clock cycle.

1. Structural hazard:

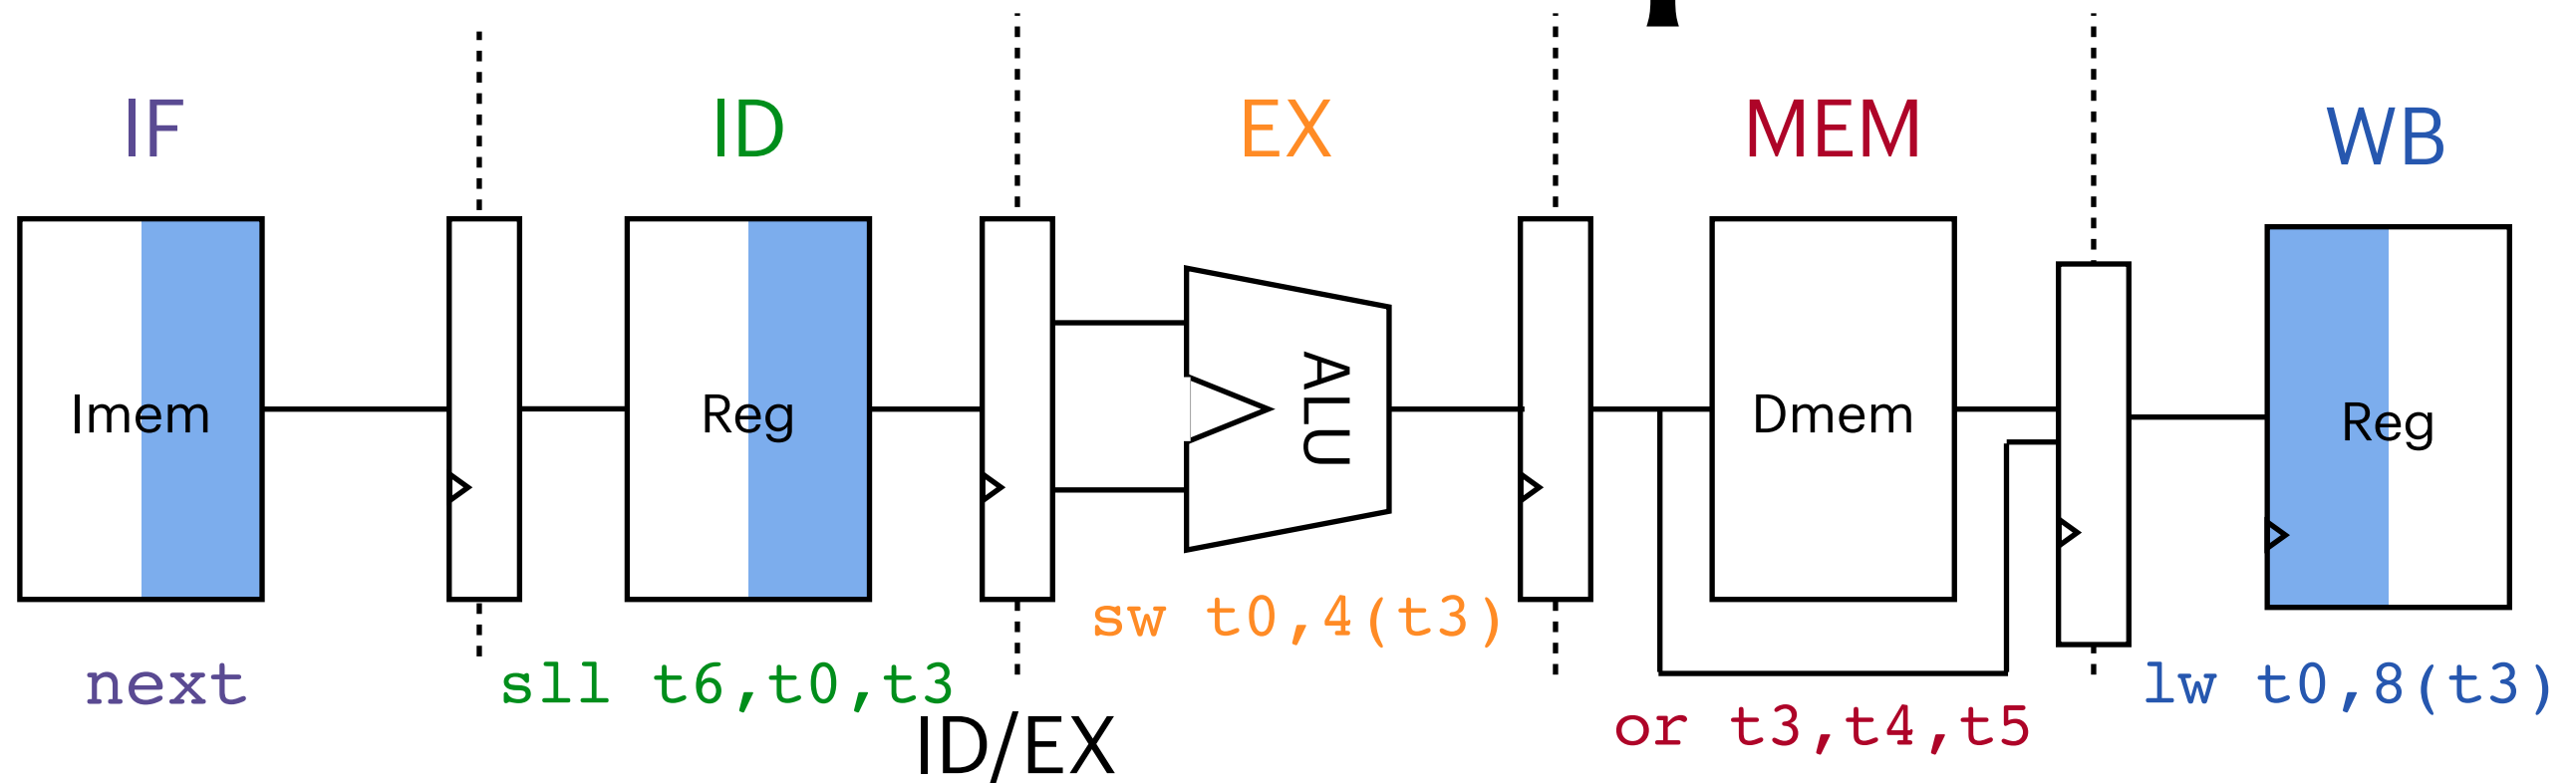
- Hardware does not support access across multiple instructions in the same cycle

2. **Data hazard:** 数据冲突

- Instructions have data dependency
- Occurs when an instruction reads a register before a previous instruction has finished writing to that register

```
add  t0,t1,t2
lw   t0,8(t3)
or   t3,t4,t5
sw   t0,4(t3)
sll  t6,t0,t3
```

t0 as an Example



At the posedge of clock cycle 7, `t0` is updated
 At the same time, ID/EX register captures the
 previous `t0`, which is erroneous

Data hazard

```

add t0, t1, t2
lw t0, 8(t3)
or t3, t4, t5
sw t0, 4(t3)
sll t6, t0, t3
    
```

Clock cycle 5

IF `sll`

ID `sw`

EX `or`

MEM `lw`

WB `add`

Clock cycle 6

IF `next`

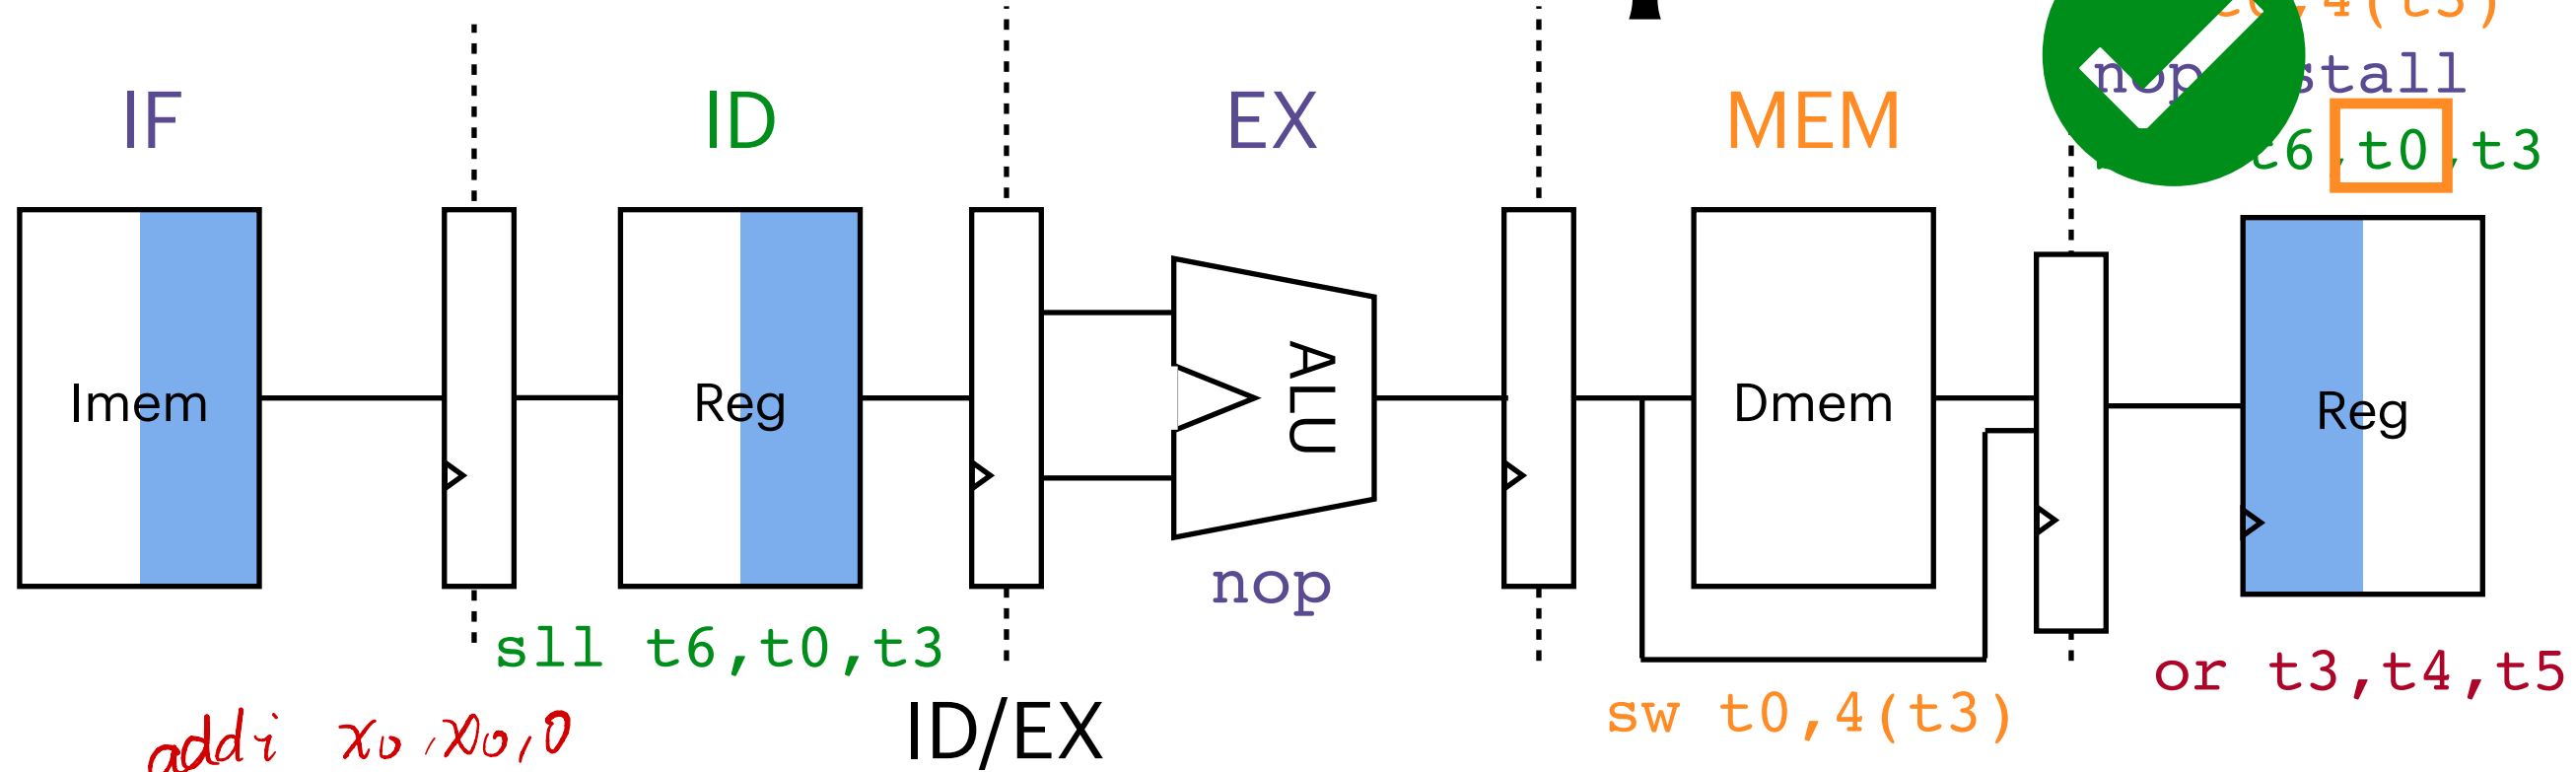
ID `sll`

EX `sw`

MEM `or`

WB `lw`

t0 as an Example



Clock cycle 5

IF **nop** ID **sw** EX **or** MEM **lw** WB **add**

Clock cycle 6

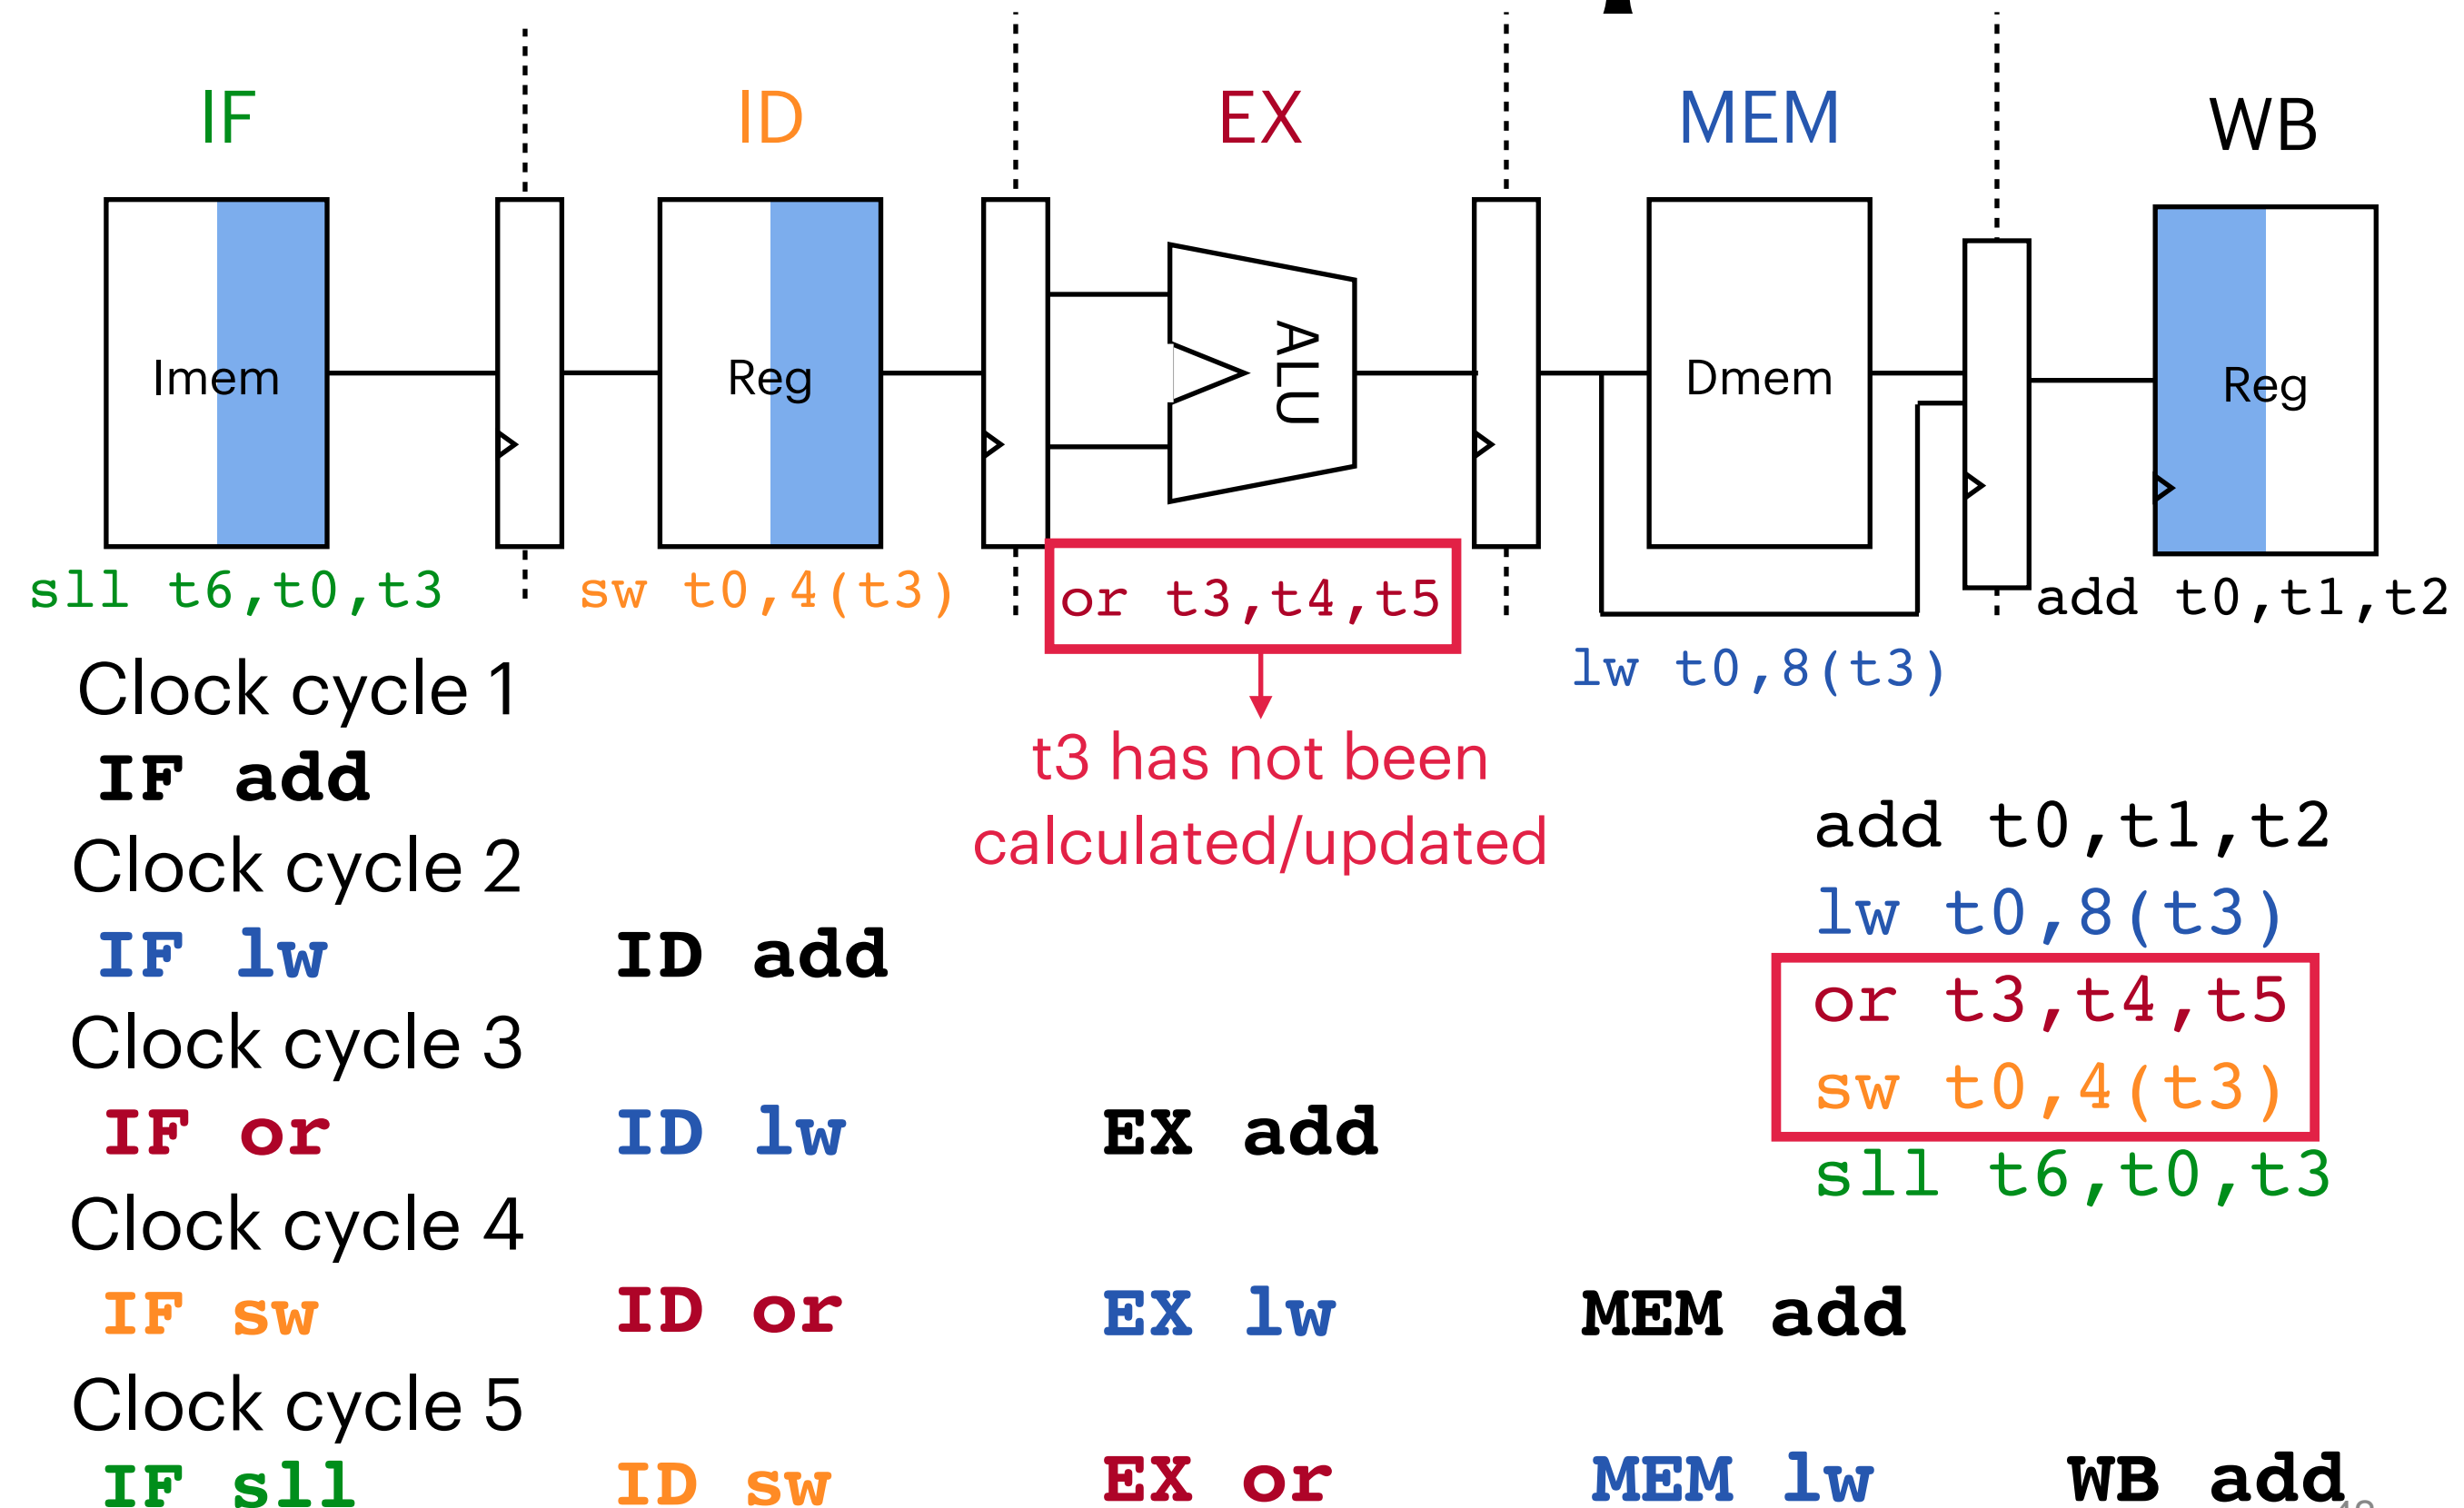
IF **sll** ID **nop** EX **sw** MEM **or** WB **lw**

Clock cycle 7: t0 updated at the beginning of this clock cycle (posedge)

IF **next** ID **sll** EX **nop** MEM **sw** WB **or**

Covered by RegFile write-then-read in the same clock cycle

Another Example



Another Example

add t0,t1,t2

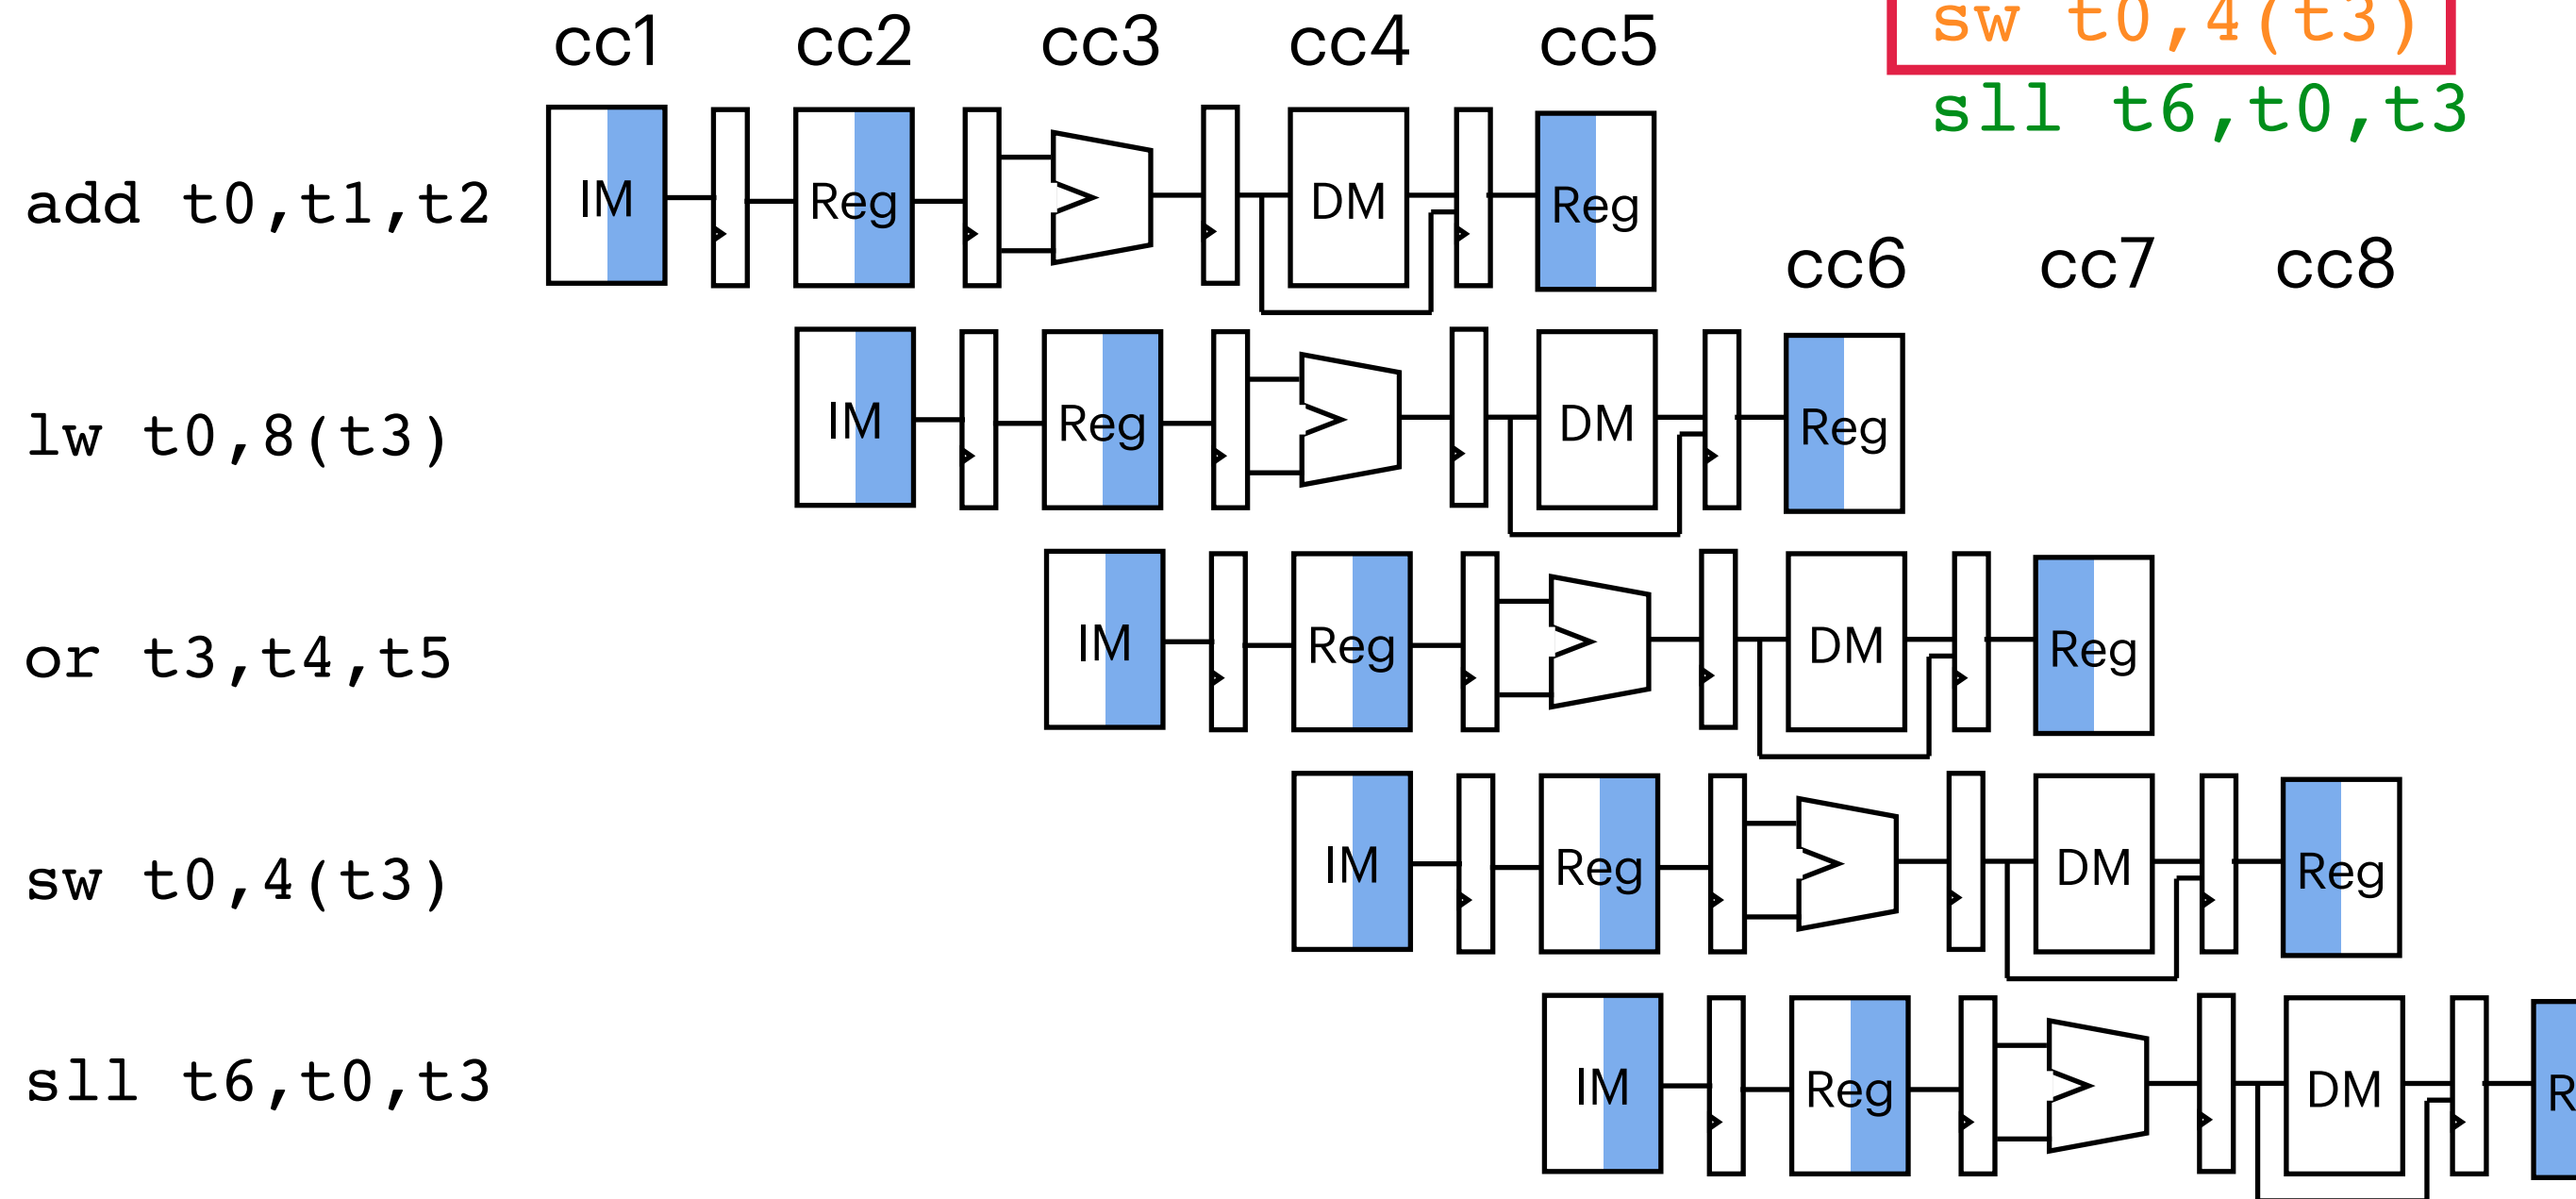
lw t0,8(t3)

or t3,t4,t5

sw t0,4(t3)

sll t6,t0,t3

Expand the pipeline both temporally and spatially



Solution 1

Starting from cc3

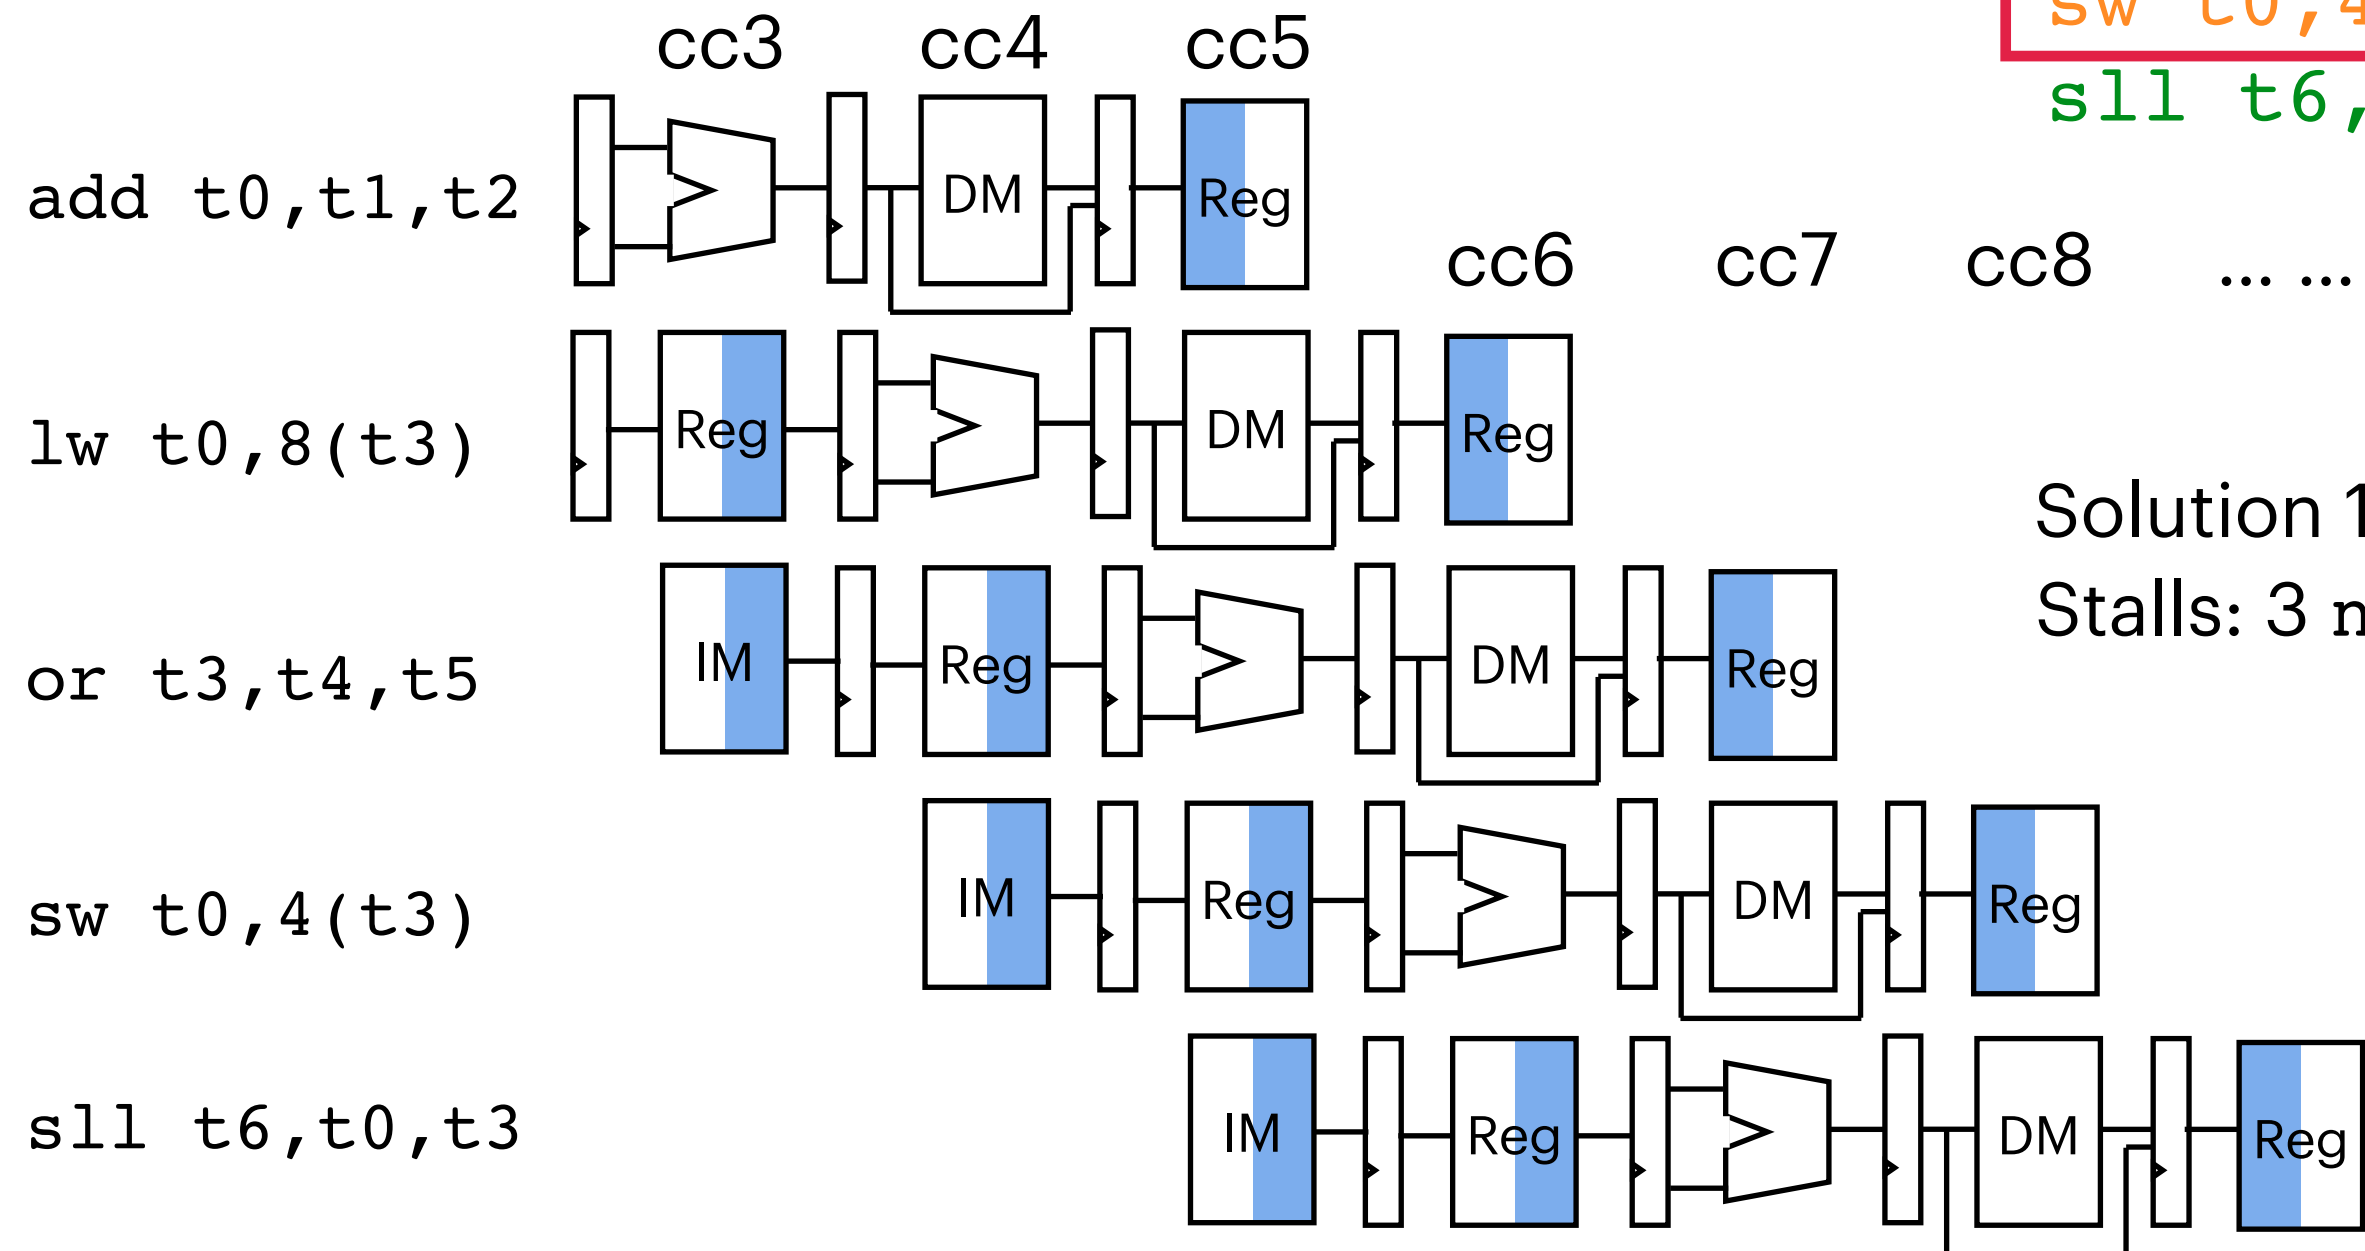
add t0,t1,t2

lw t0,8(t3)

or t3,t4,t5

sw t0,4(t3)

sll t6,t0,t3



Solution 1:

Stalls: 3 nops

t3 value

Prev.

Prev.

Prev.

Prev.

Prev.

New

New

New