

# Course Info

- Lab 7 next week, prepare before lab sessions!
- Project 1.2 ddl soon (March 31<sup>st</sup>).
- Project 2.1 coming soon! April 14<sup>th</sup> ddl.
- Next week discussion on pipeline & superscalar.
- Mid-term I solution & score released. If you have questions about the solution, feel free to ask on Piazza; If you have questions regarding your marks, email the instructors **BEFORE April 2nd**. We will get back to you ASAP.
- Any regrade request after April 2nd **WOULD NOT** be considered.

# Course Info

- HW4 released. Submit your paper homework to the box below (at SIST 3-322). DDL April 7th.
- Remember to add your name. You have only one chance to submit and cannot be withdrawn.





信息科学与技术学院

School of Information Science and Technology

# CS 110

## Computer Architecture

# Hazards & Advanced Techniques

**Instructors:**

**Siting Liu & Chundong Wang**

Course website: <https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html>

**School of Information Science and Technology (SIST)**

**ShanghaiTech University**

2023/3/23

# Review

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}}$$

- We have built a pipelined processor!
  - Each instruction might consumes a longer time, but the overall throughput is improved
  - $T_{clk} = \text{Max delay}(t_{IF}, t_{ID}, t_{EX}, t_{MEM}, t_{WB})$
  - Max Frequency = 1/Max delay
- A hazard is a situation in which a planned instruction cannot execute in the “proper” clock cycle.
  - Structural hazard: does not exist in our current design
  - Data hazard:
    - Solution 1: insert nop/bubble/stall, CPI increased

# Three Types of Pipeline Hazards

A hazard is a situation in which a planned instruction cannot execute in the “proper” clock cycle.

1. Structural hazard:

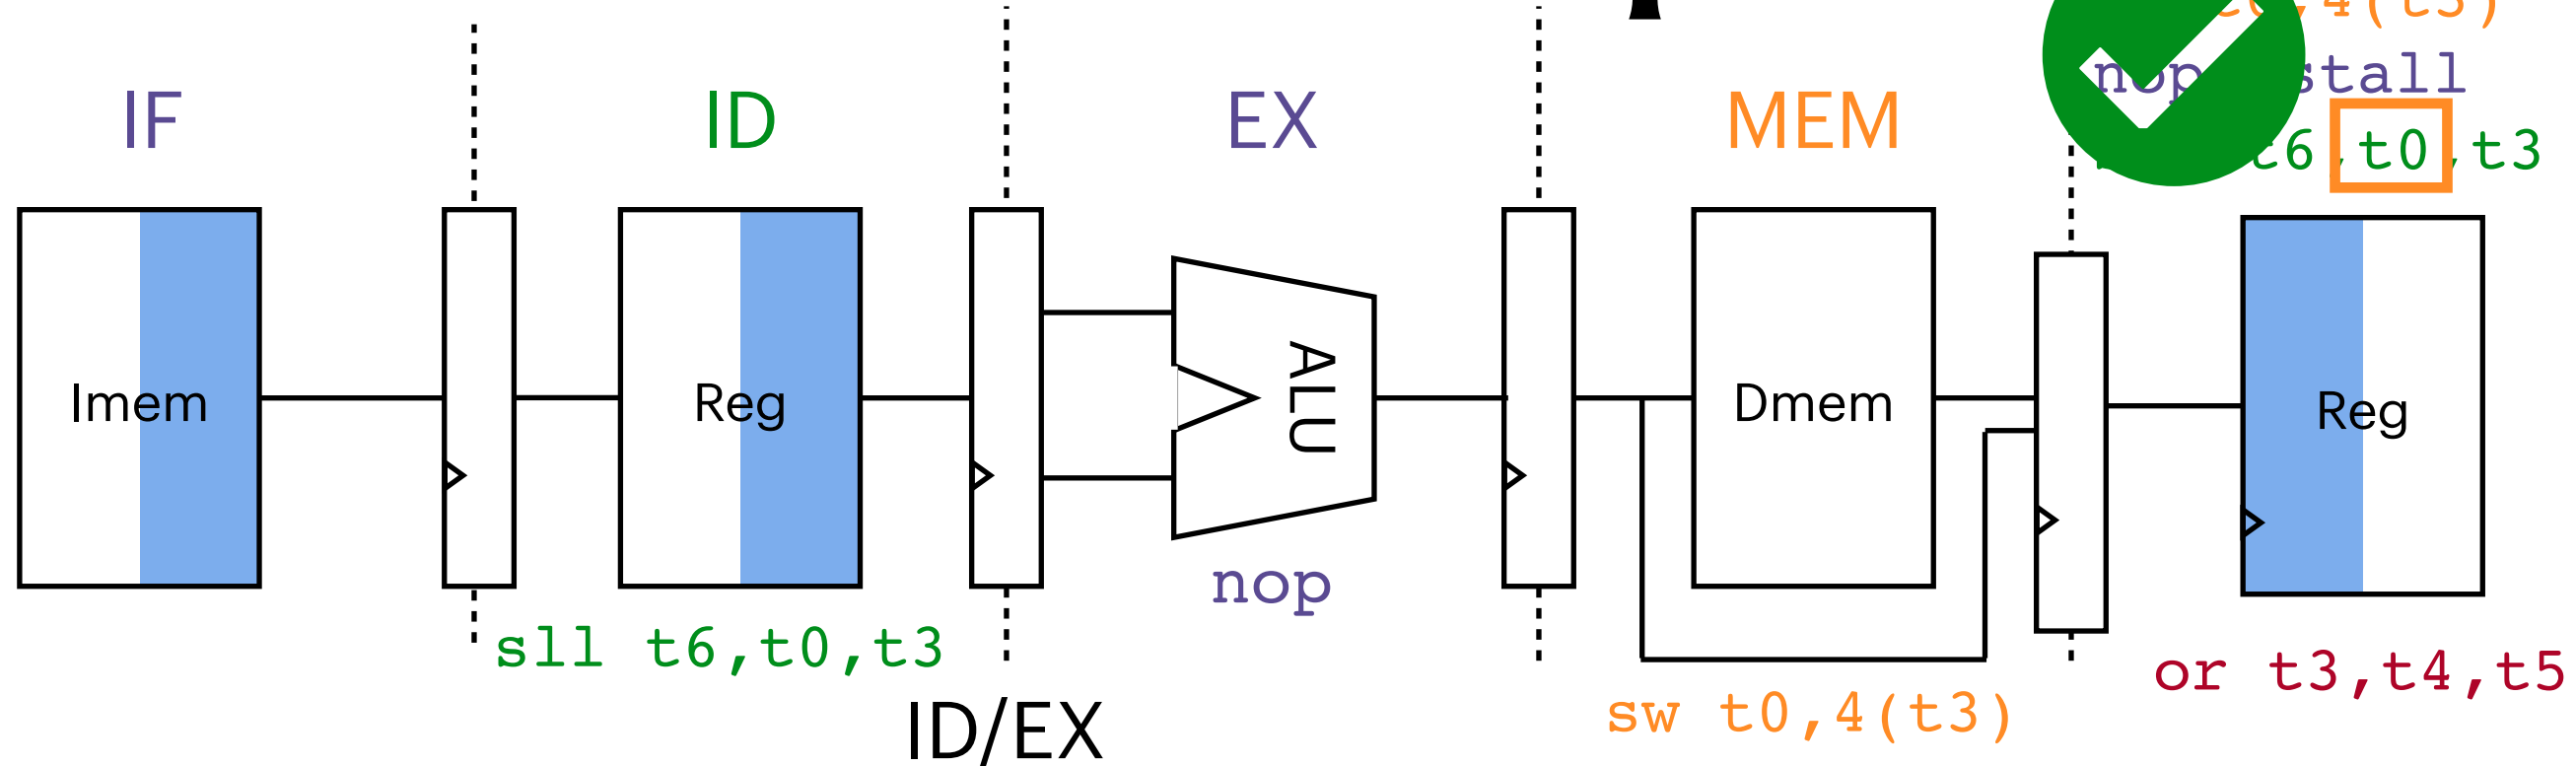
- Hardware does not support access across multiple instructions in the same cycle

2. **Data hazard:**

- Instructions have data dependency
- Occurs when an instruction reads a register before a previous instruction has finished writing to that register

```
add  t0,t1,t2
lw   t0,8(t3)
or   t3,t4,t5
sw   t0,4(t3)
sll  t6,t0,t3
```

# t0 as an Example



Clock cycle 5

IF nop

ID sw

EX or

MEM lw

WB add

Clock cycle 6

IF sll

ID nop

EX sw

MEM or

WB lw

Clock cycle 7: t0 updated at the beginning of this clock cycle (posedge)

IF next

ID sll

EX nop

MEM sw

WB or

Covered by RegFile write-then-read in the same clock cycle

# Solution 1

Starting from cc3

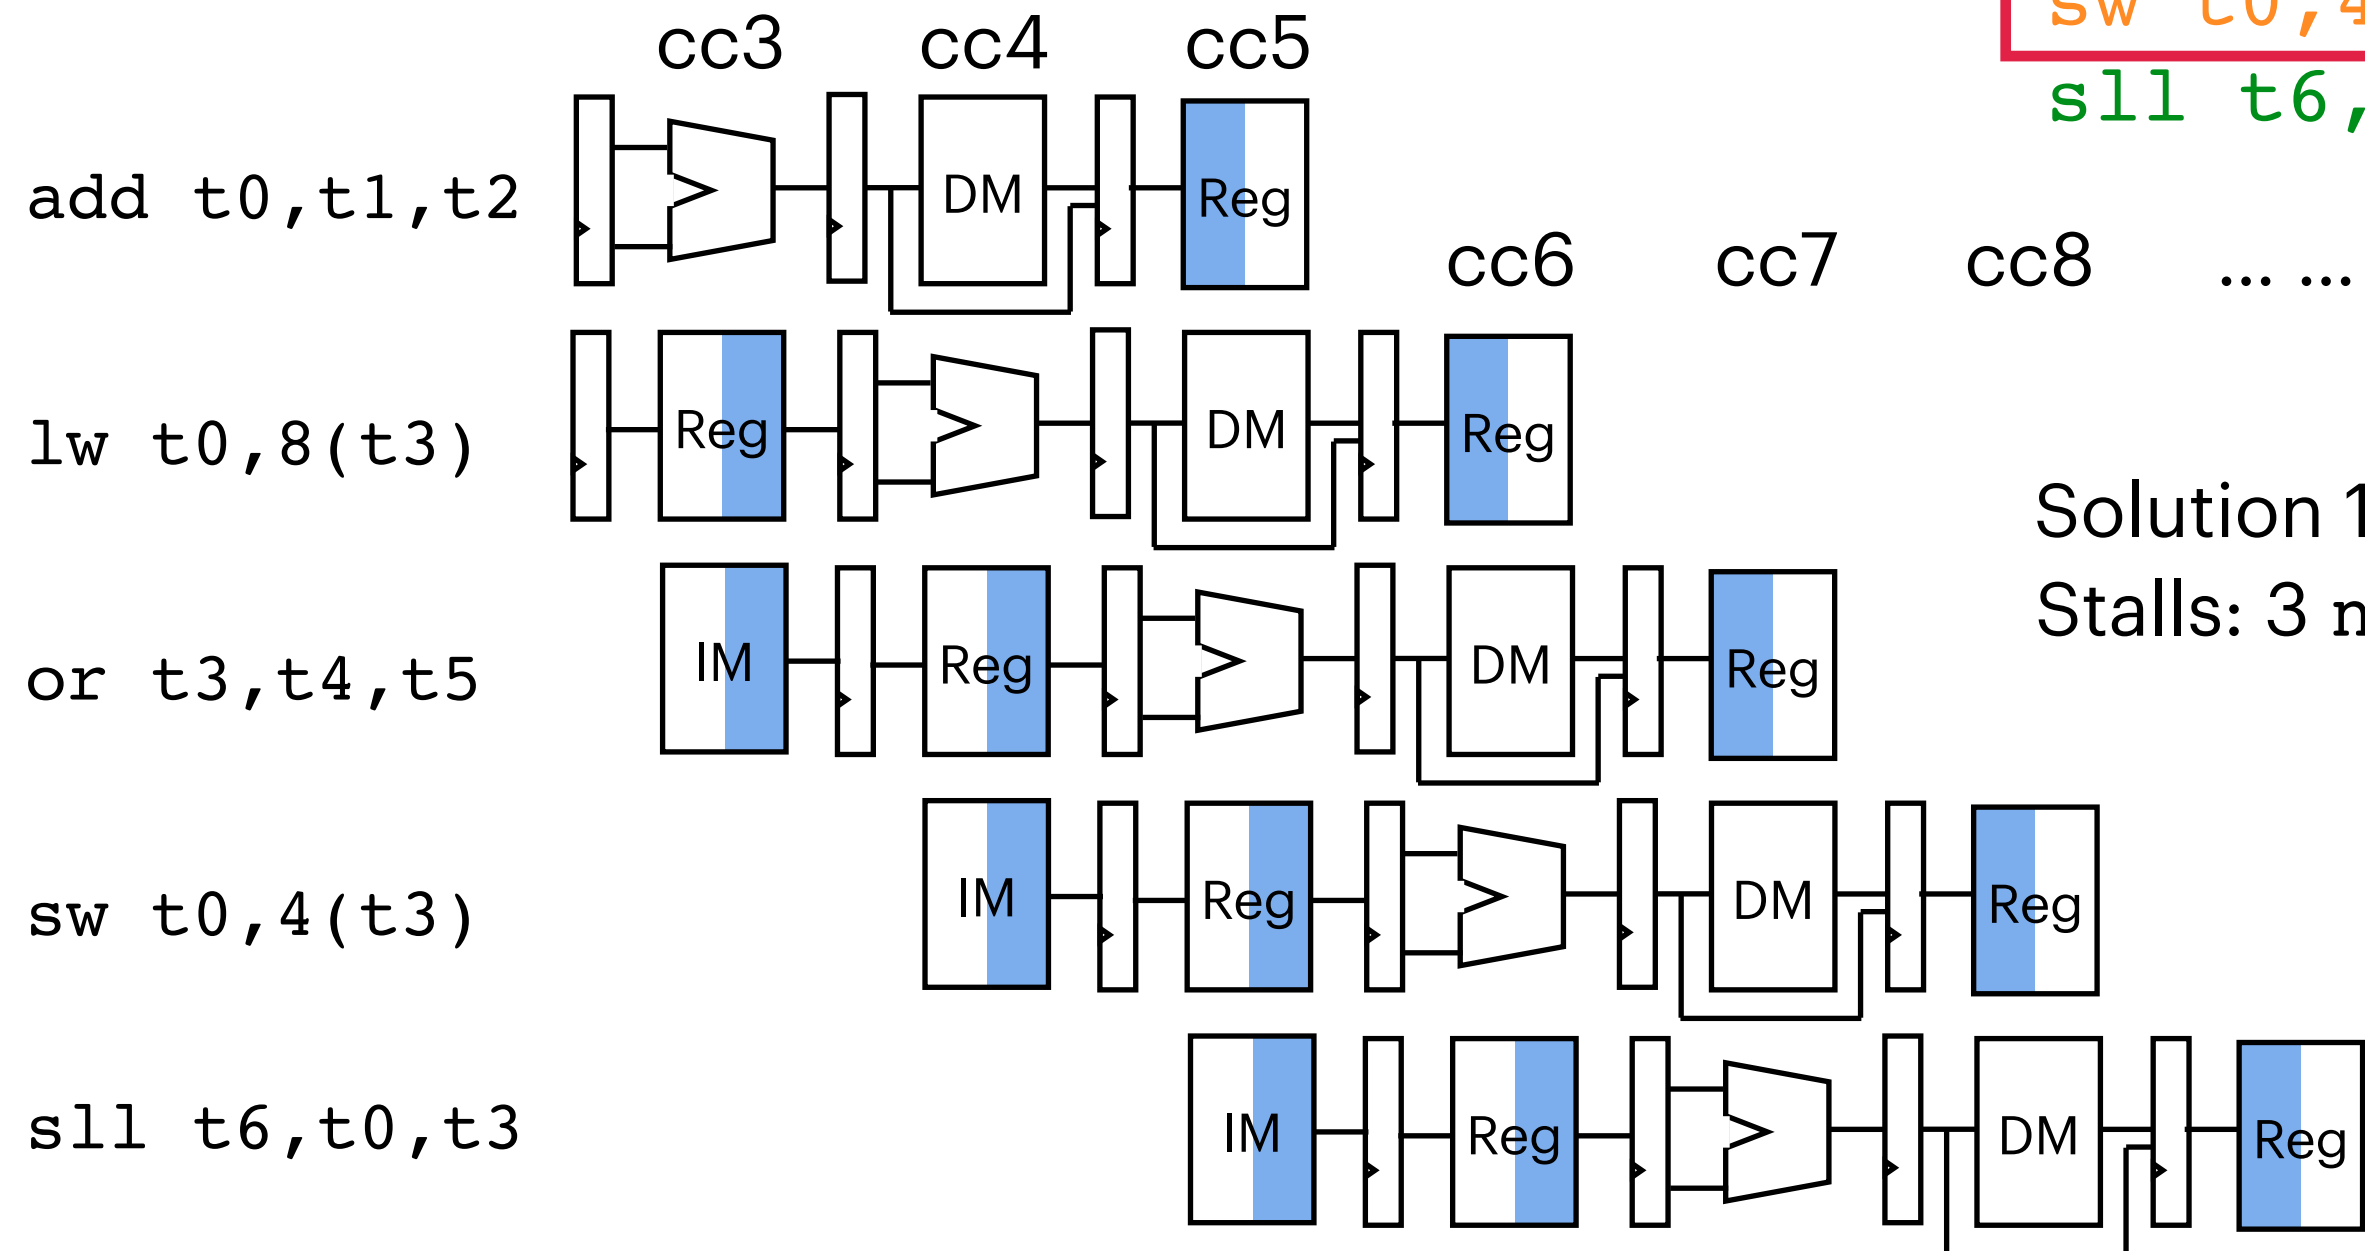
add t0,t1,t2

lw t0,8(t3)

or t3,t4,t5

sw t0,4(t3)

sll t6,t0,t3



Solution 1:

Stalls: 3 nops

t3 value

Prev.

Prev.

Prev.

Prev.

Prev.

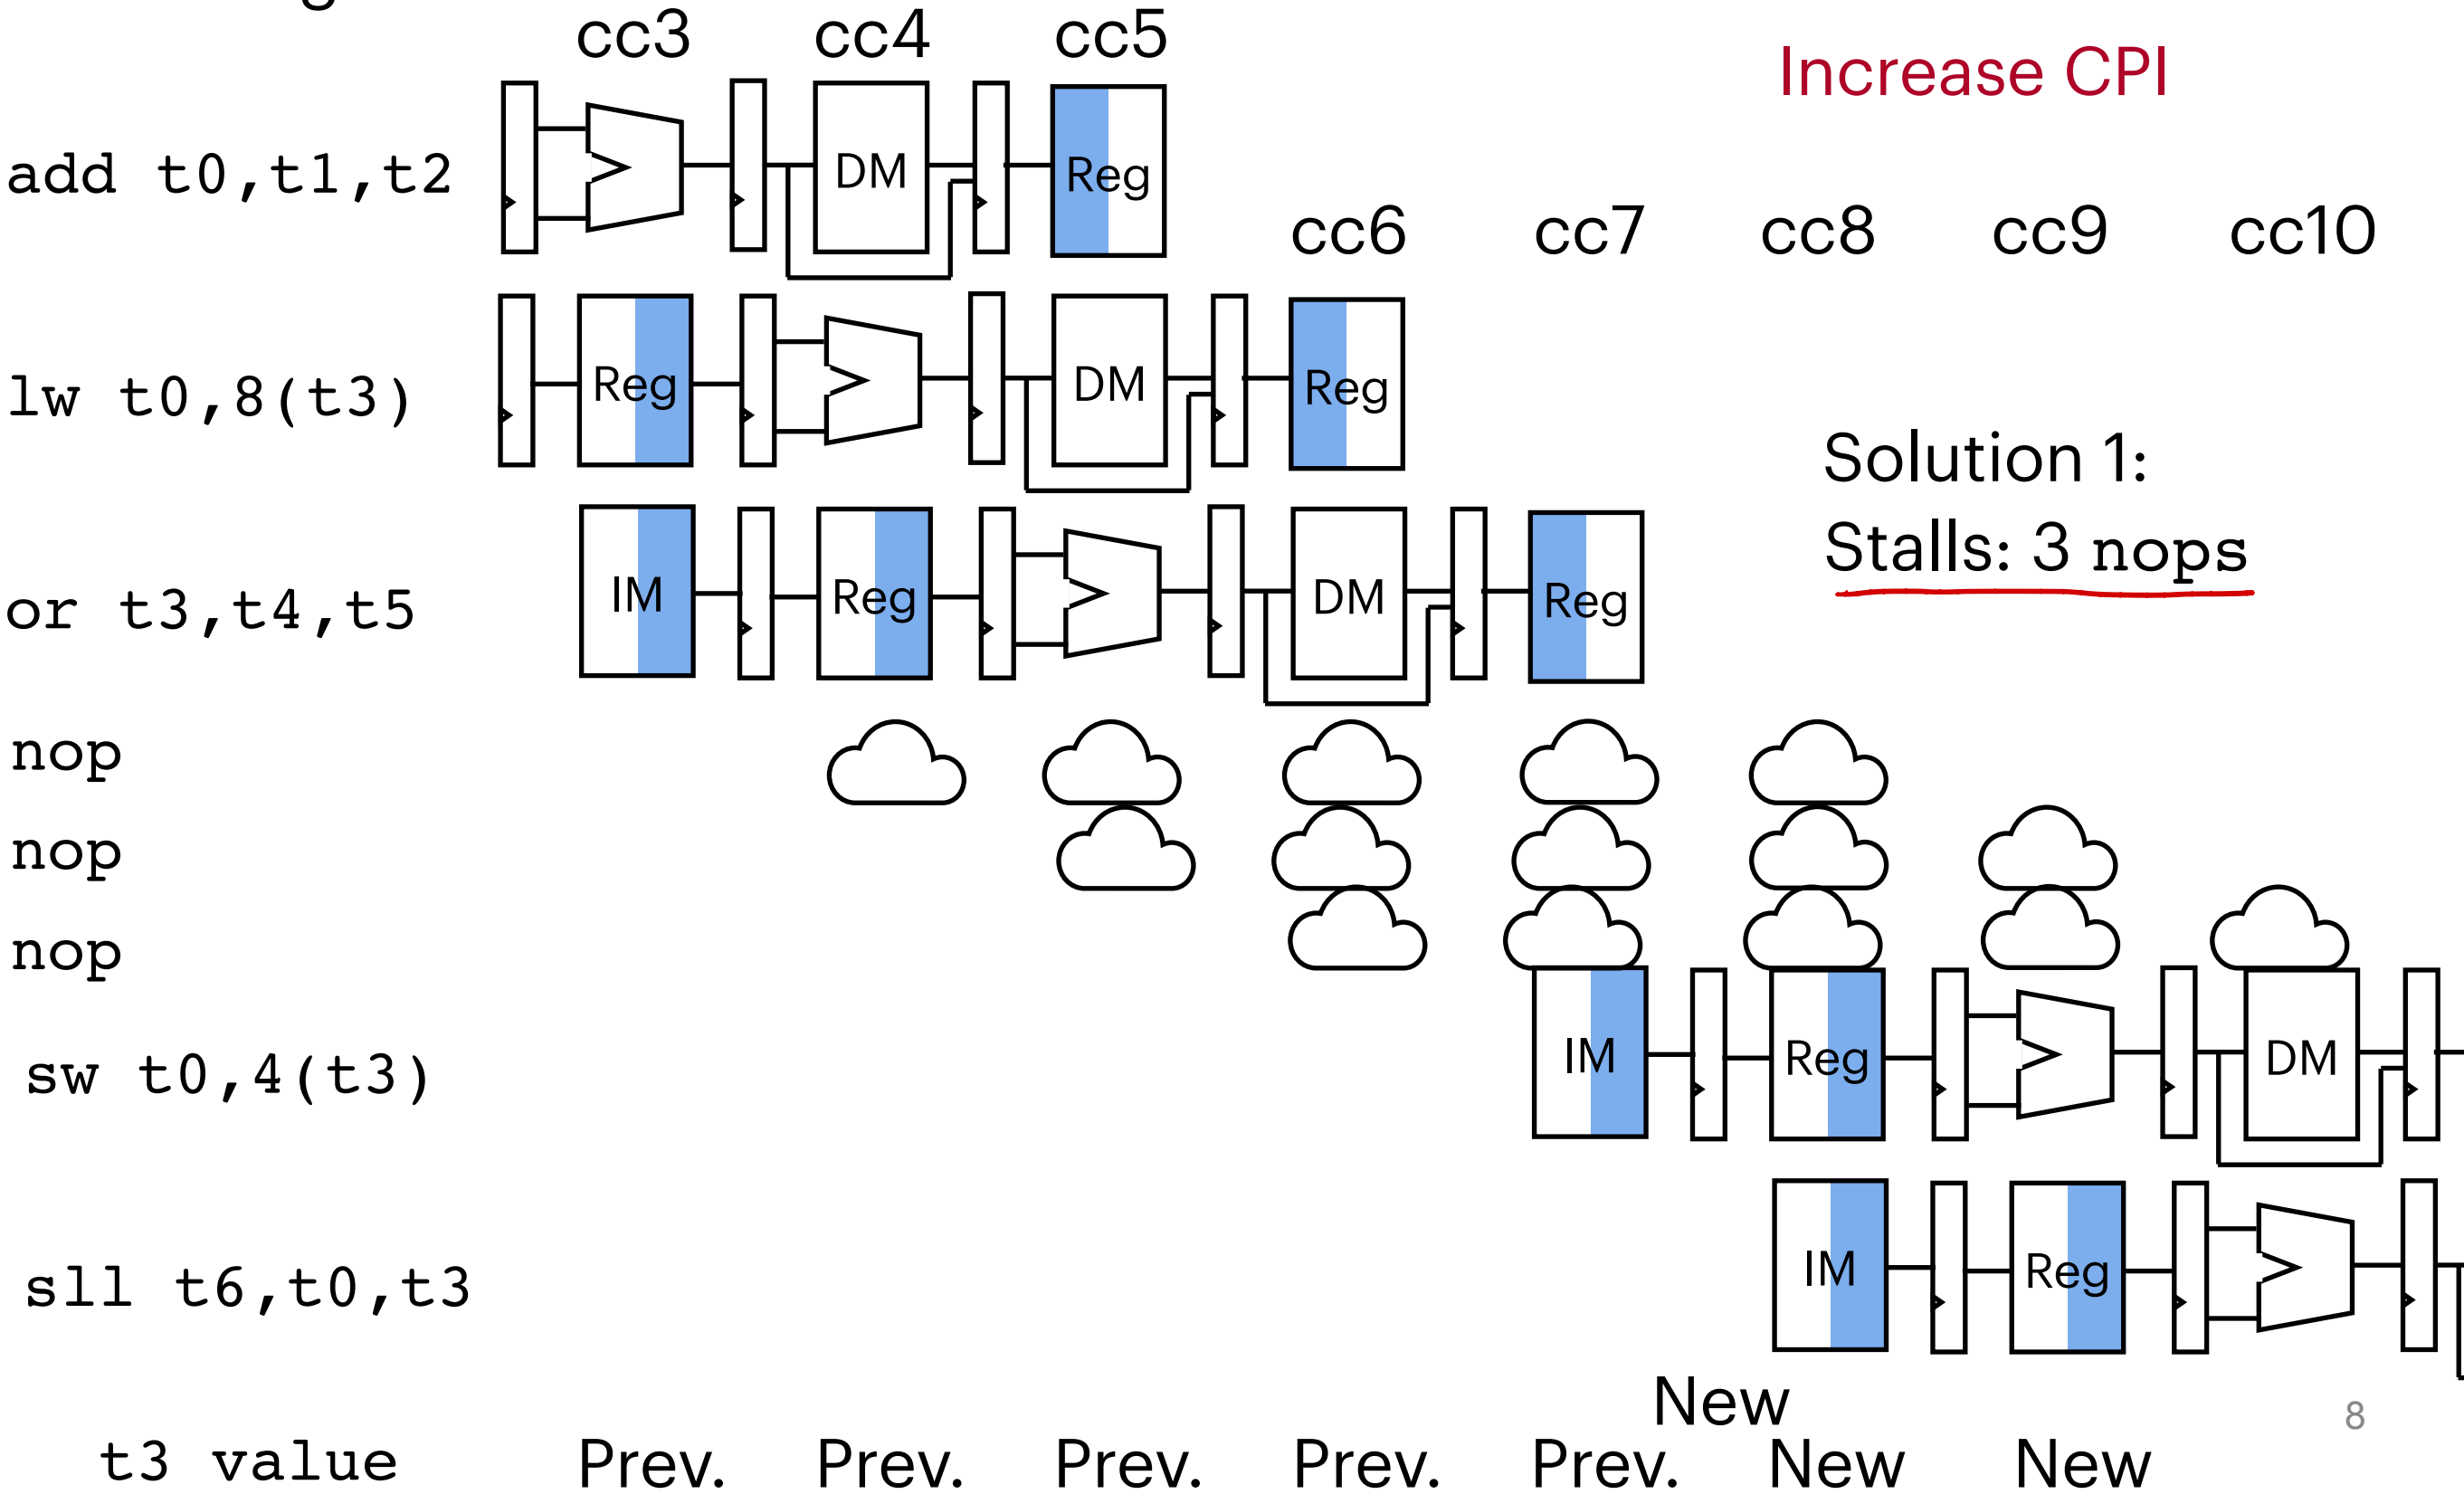
New

New

New

# Stalls and Performance

Starting from cc3





# Solution 2: Forwarding/Bypassing

Starting from cc3

add t0,t1,t2

lw t0,8(t3)

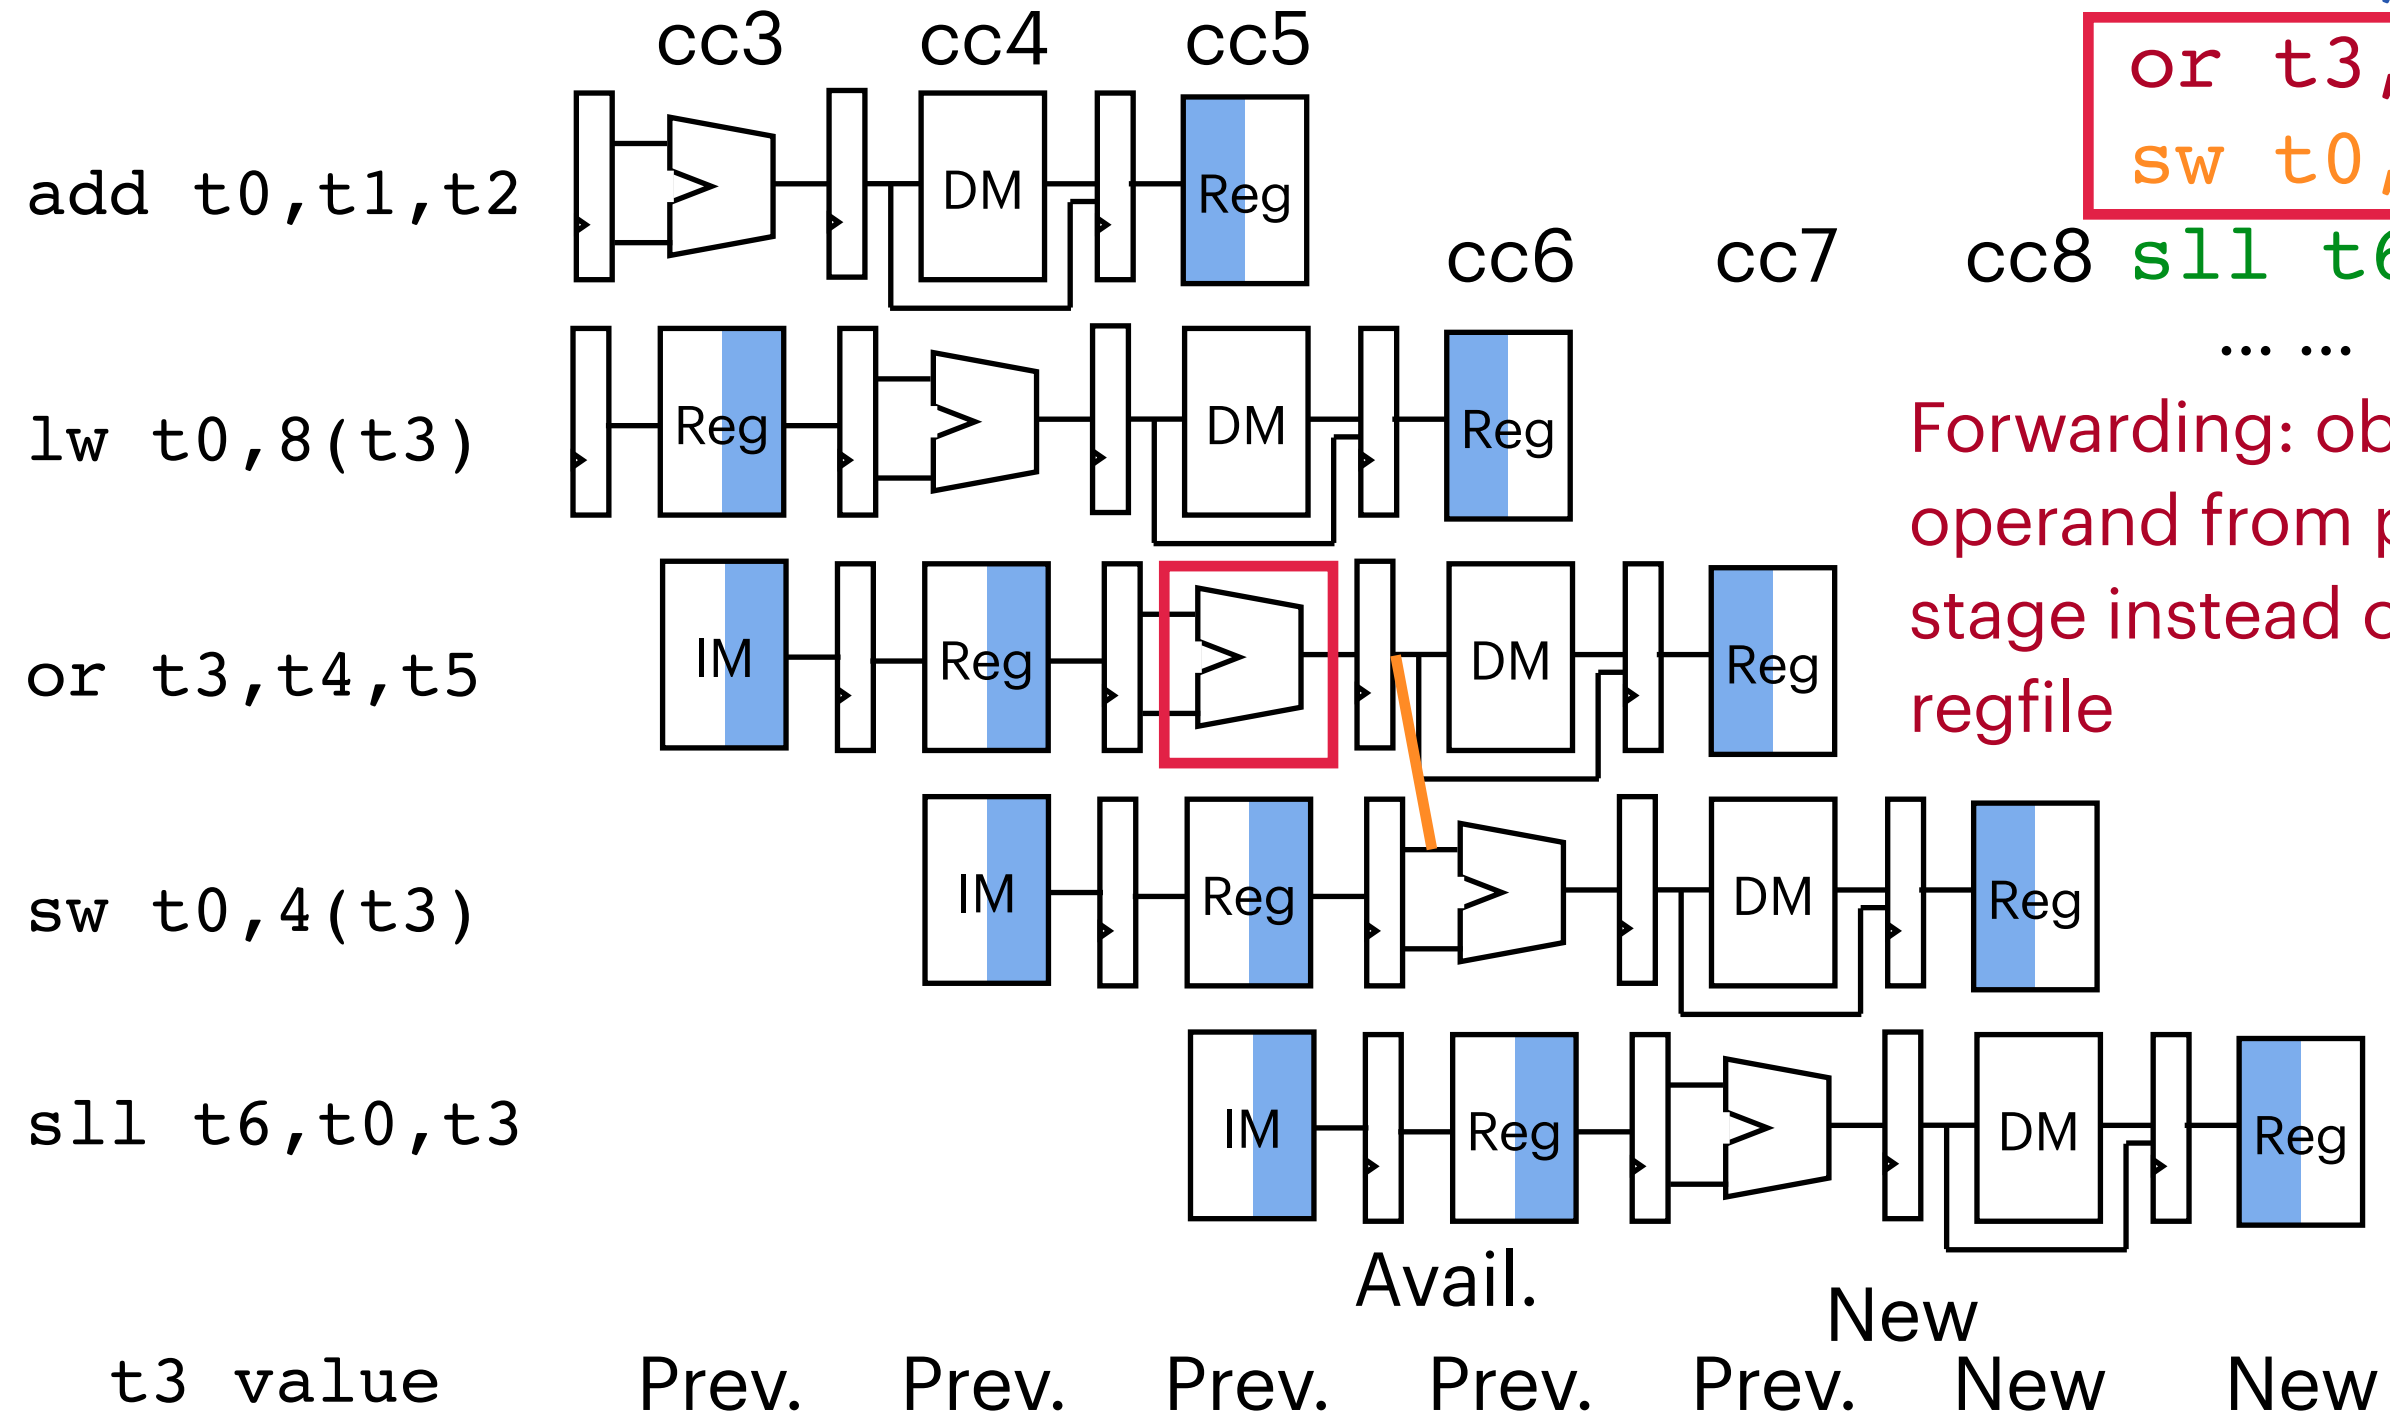
or t3,t4,t5

sw t0,4(t3)

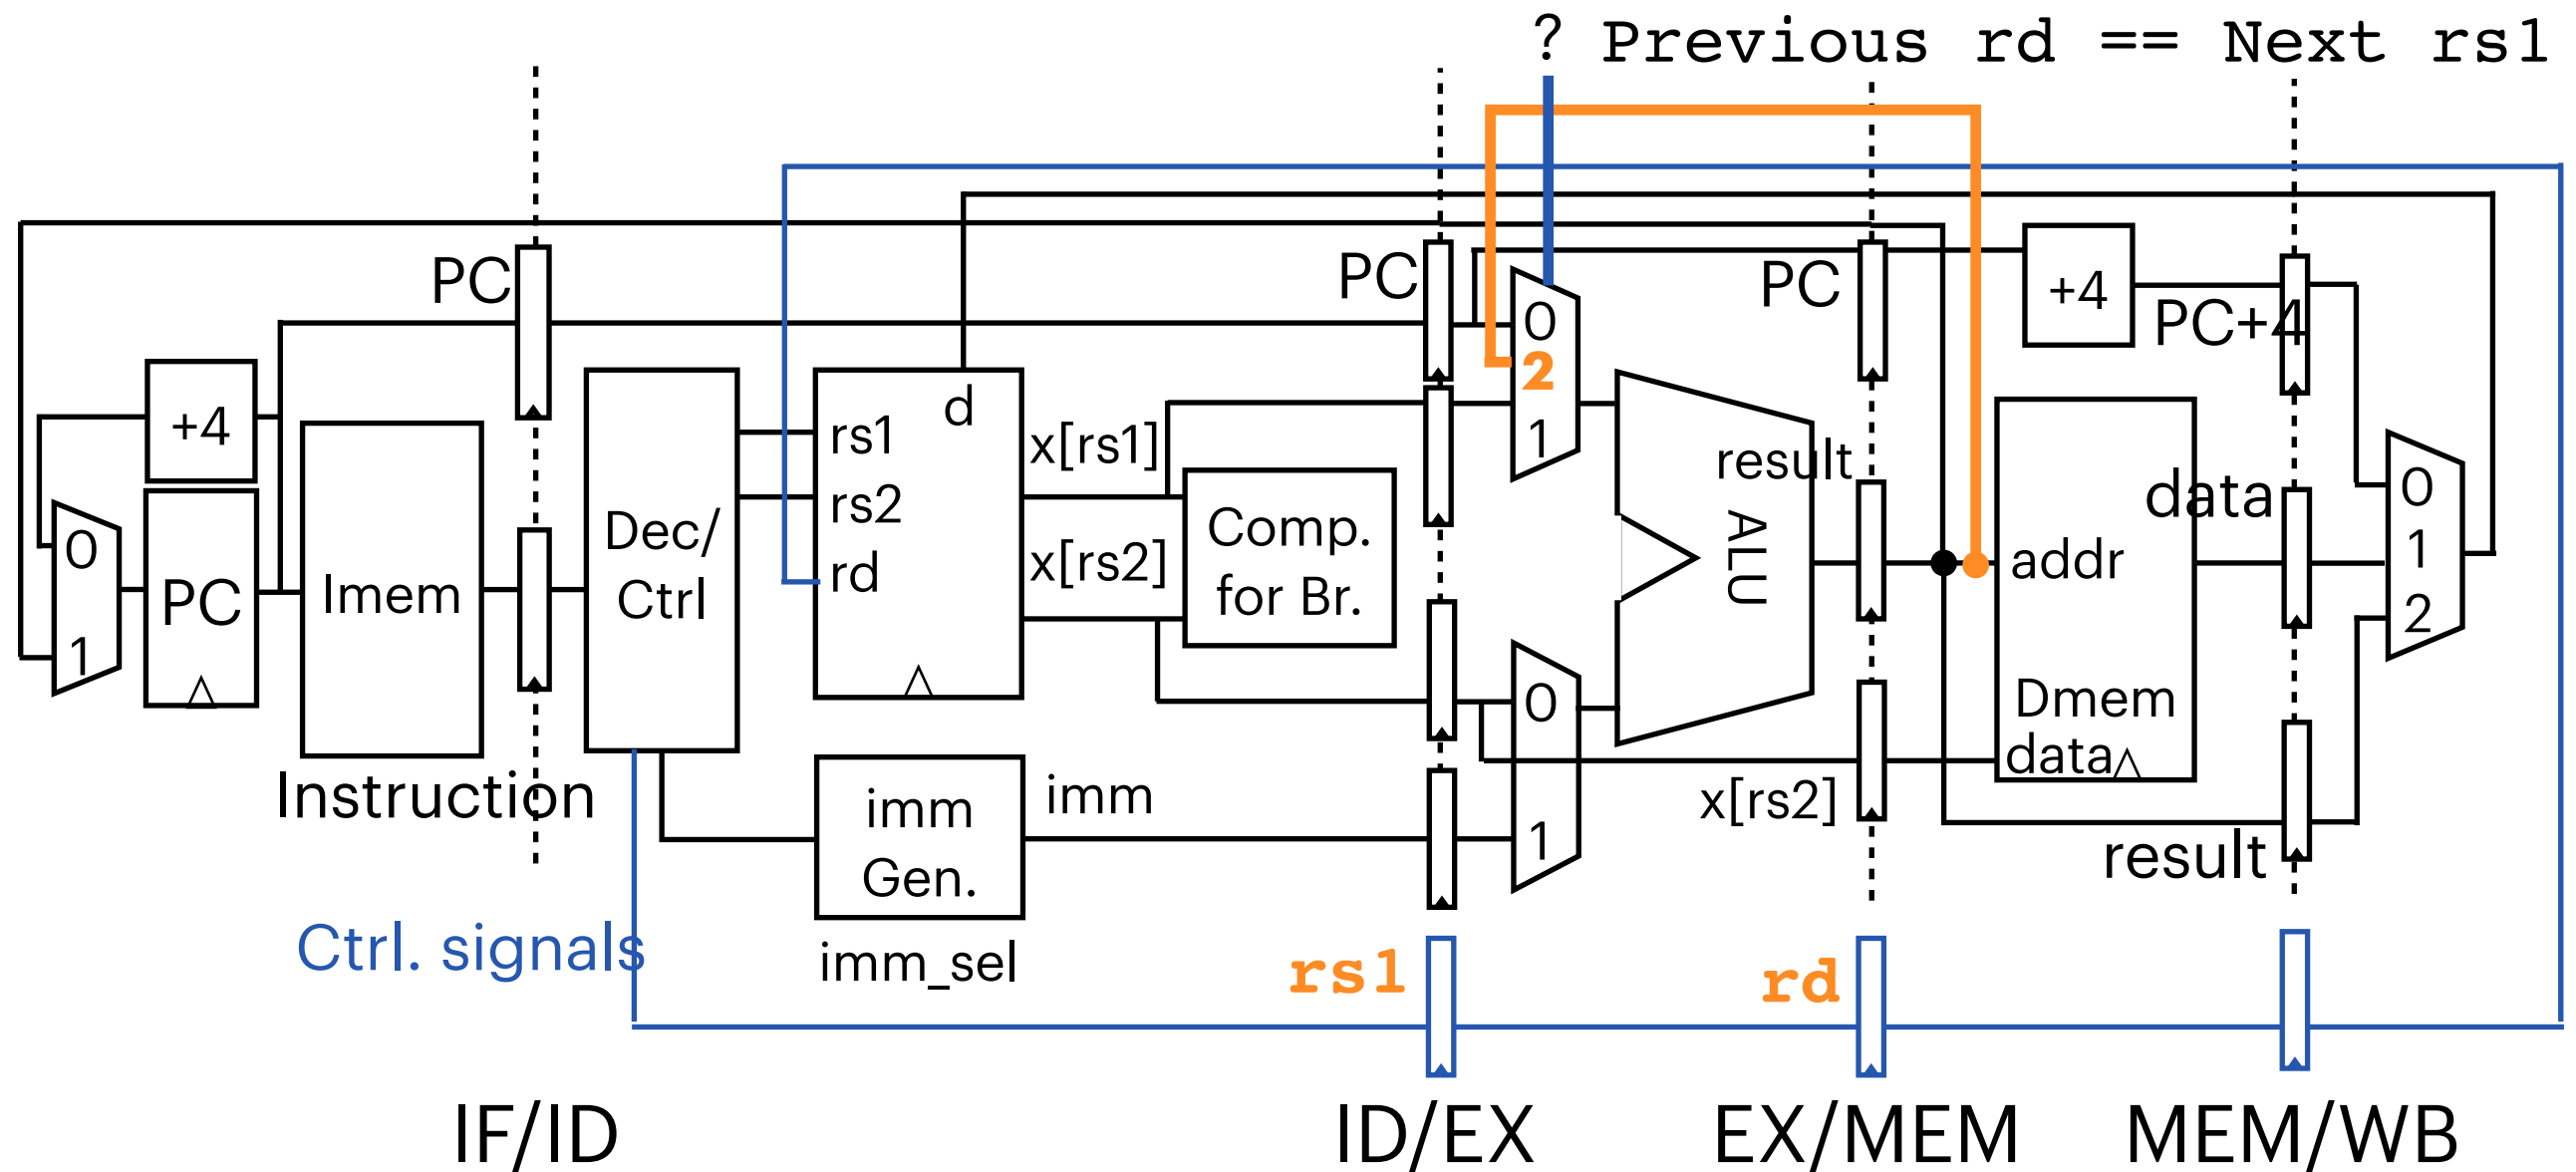
sll t6,t0,t3

...

Forwarding: obtain  
operand from pipeline  
stage instead of the  
regfile

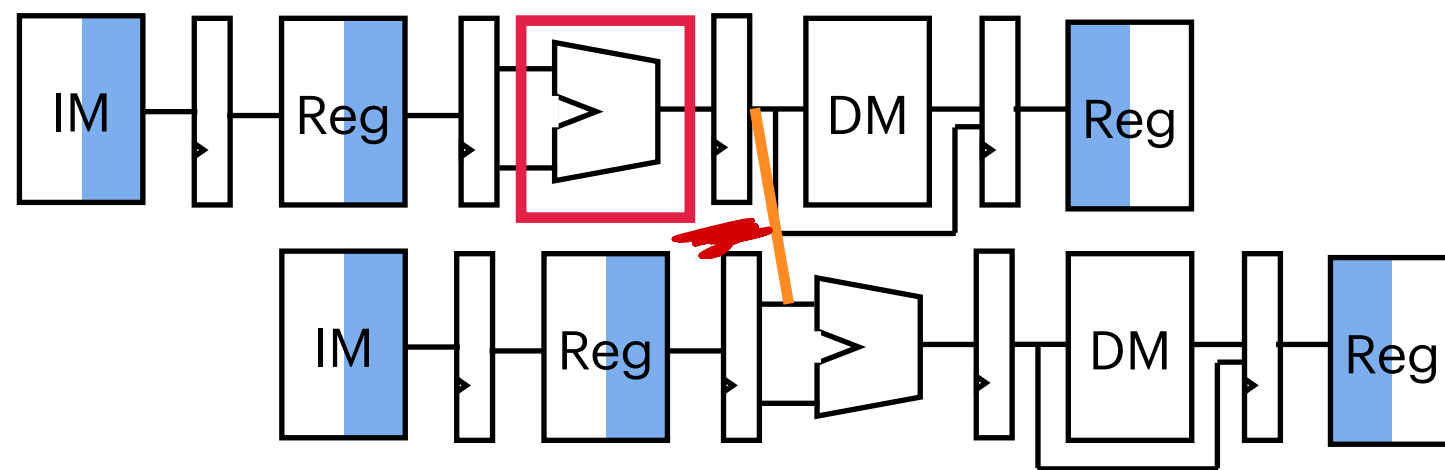


# Adjust the Datapath: Add Extra Connections

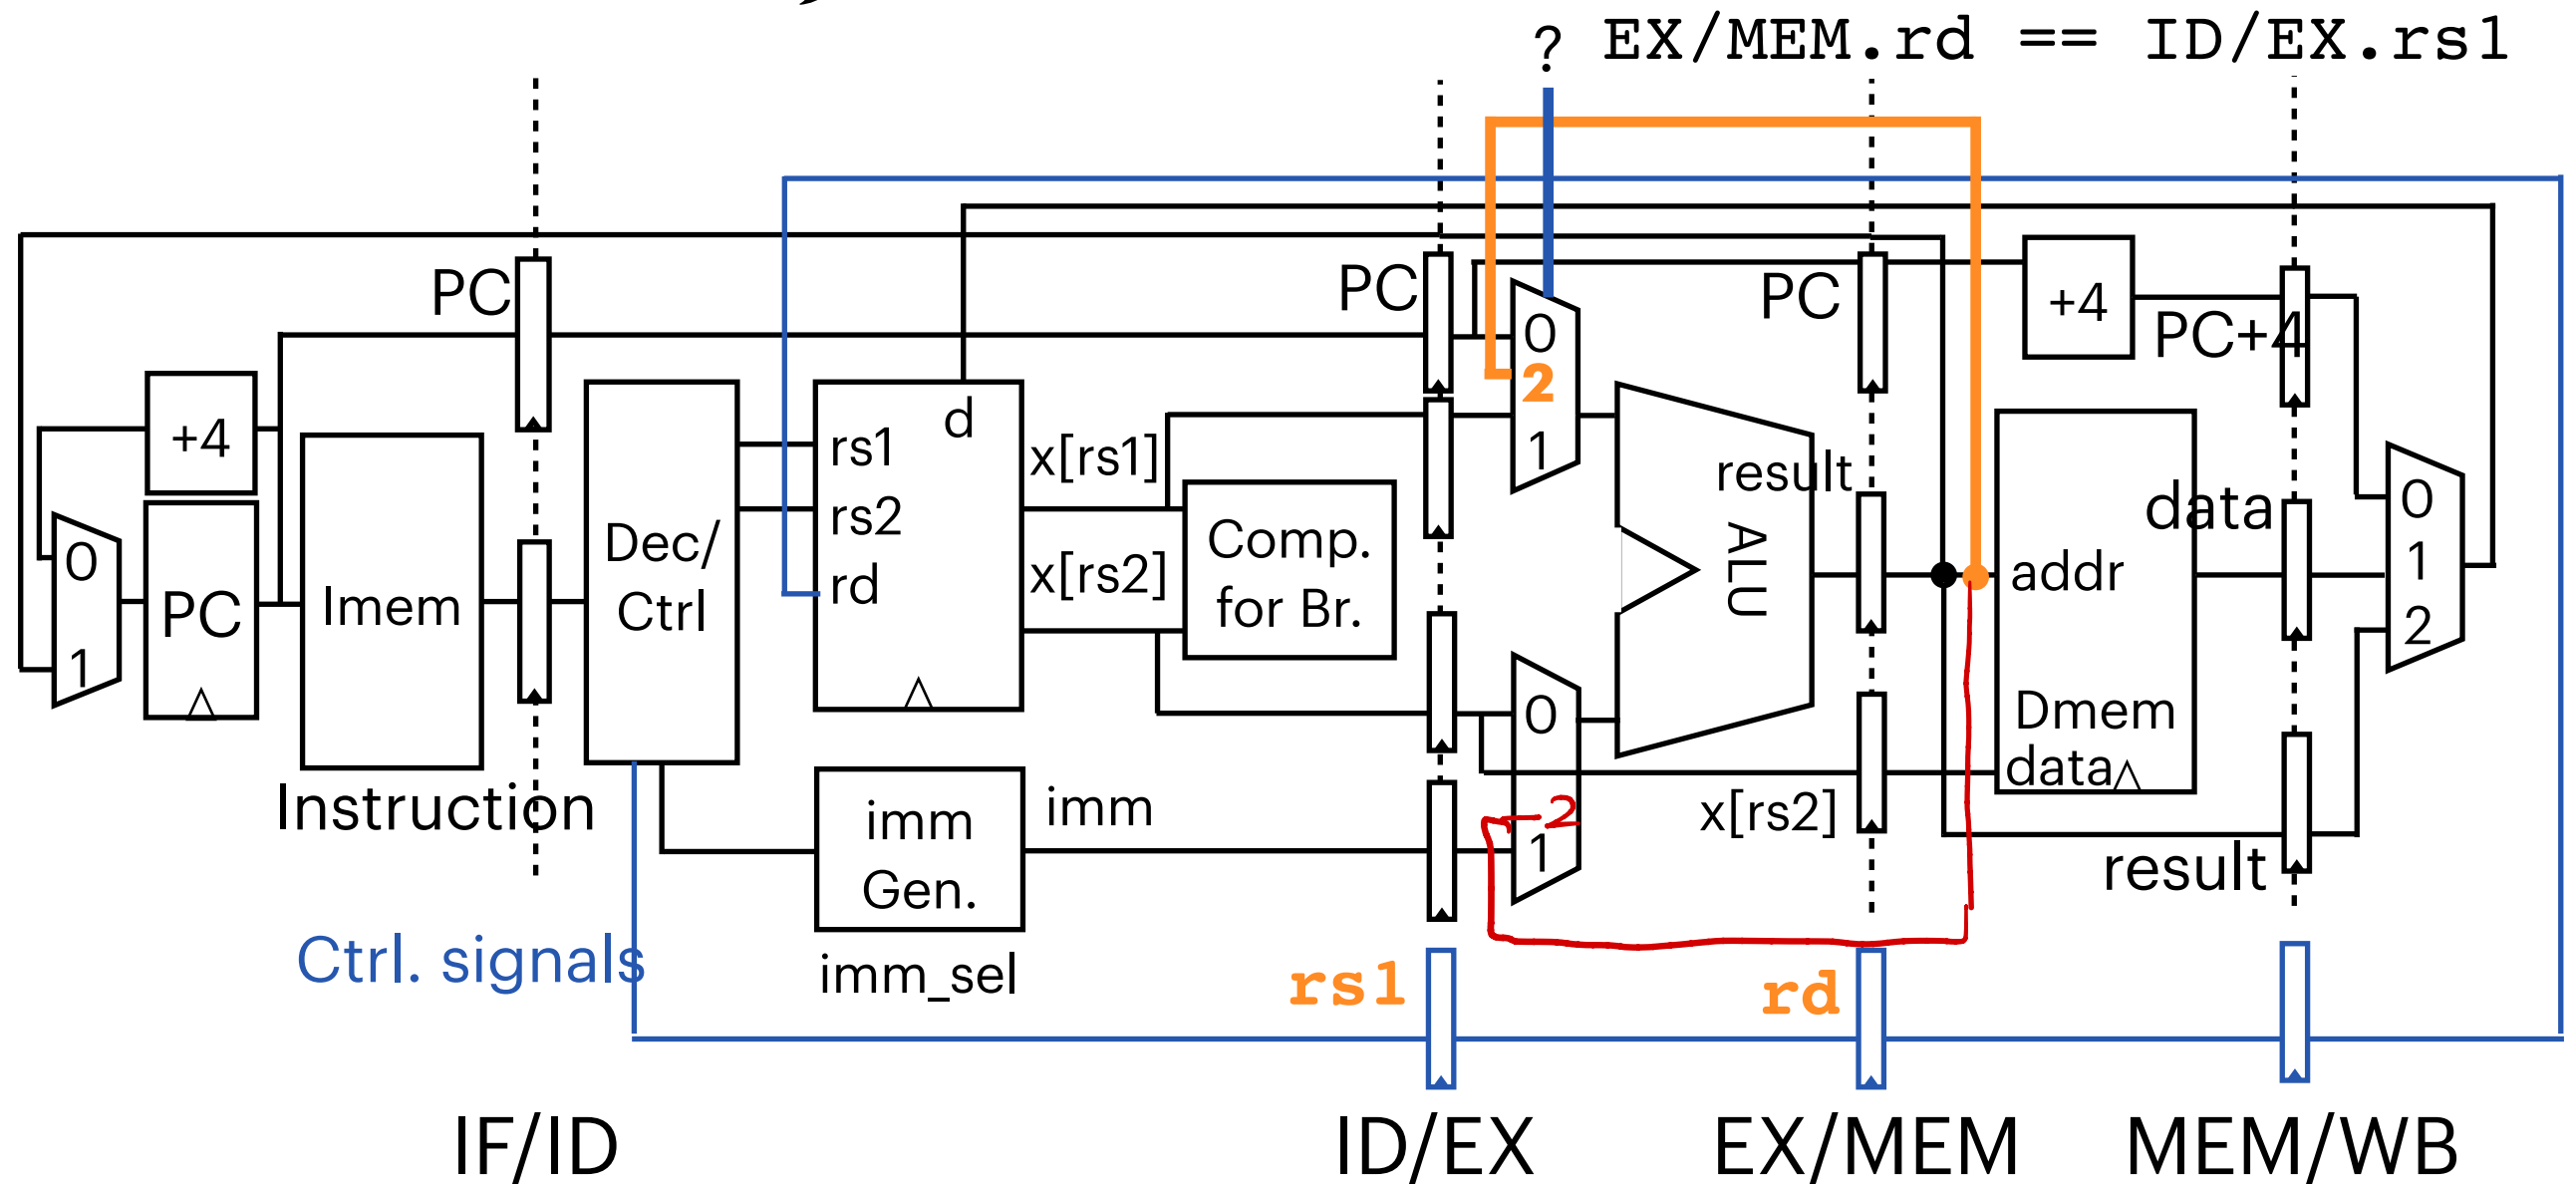


or  $t3, t4, t5$

sw  $t0, 4(t3)$

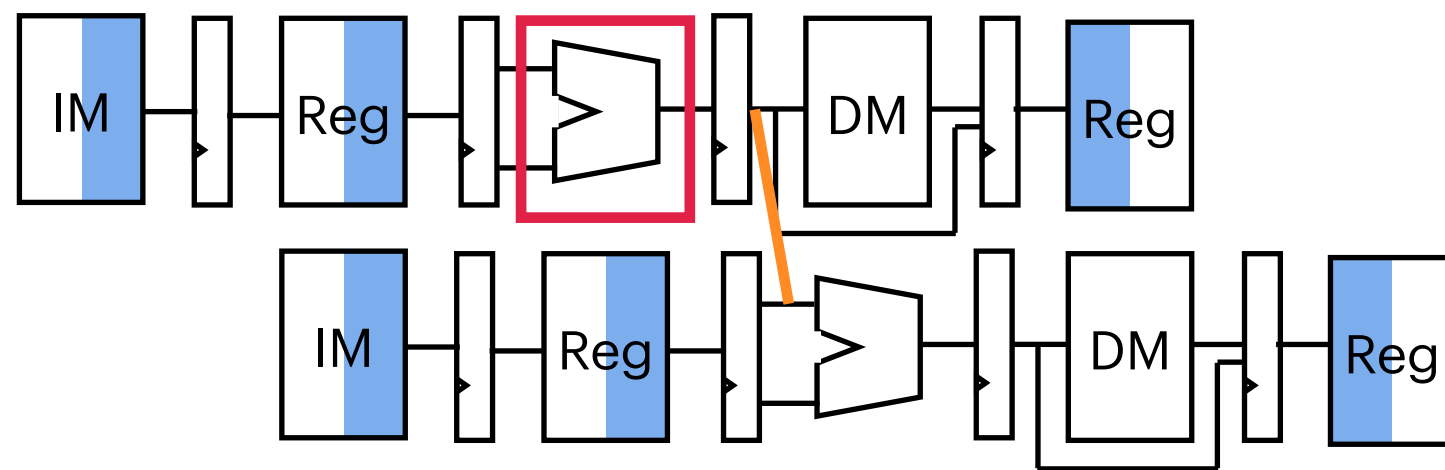


# Adjust the Datapath

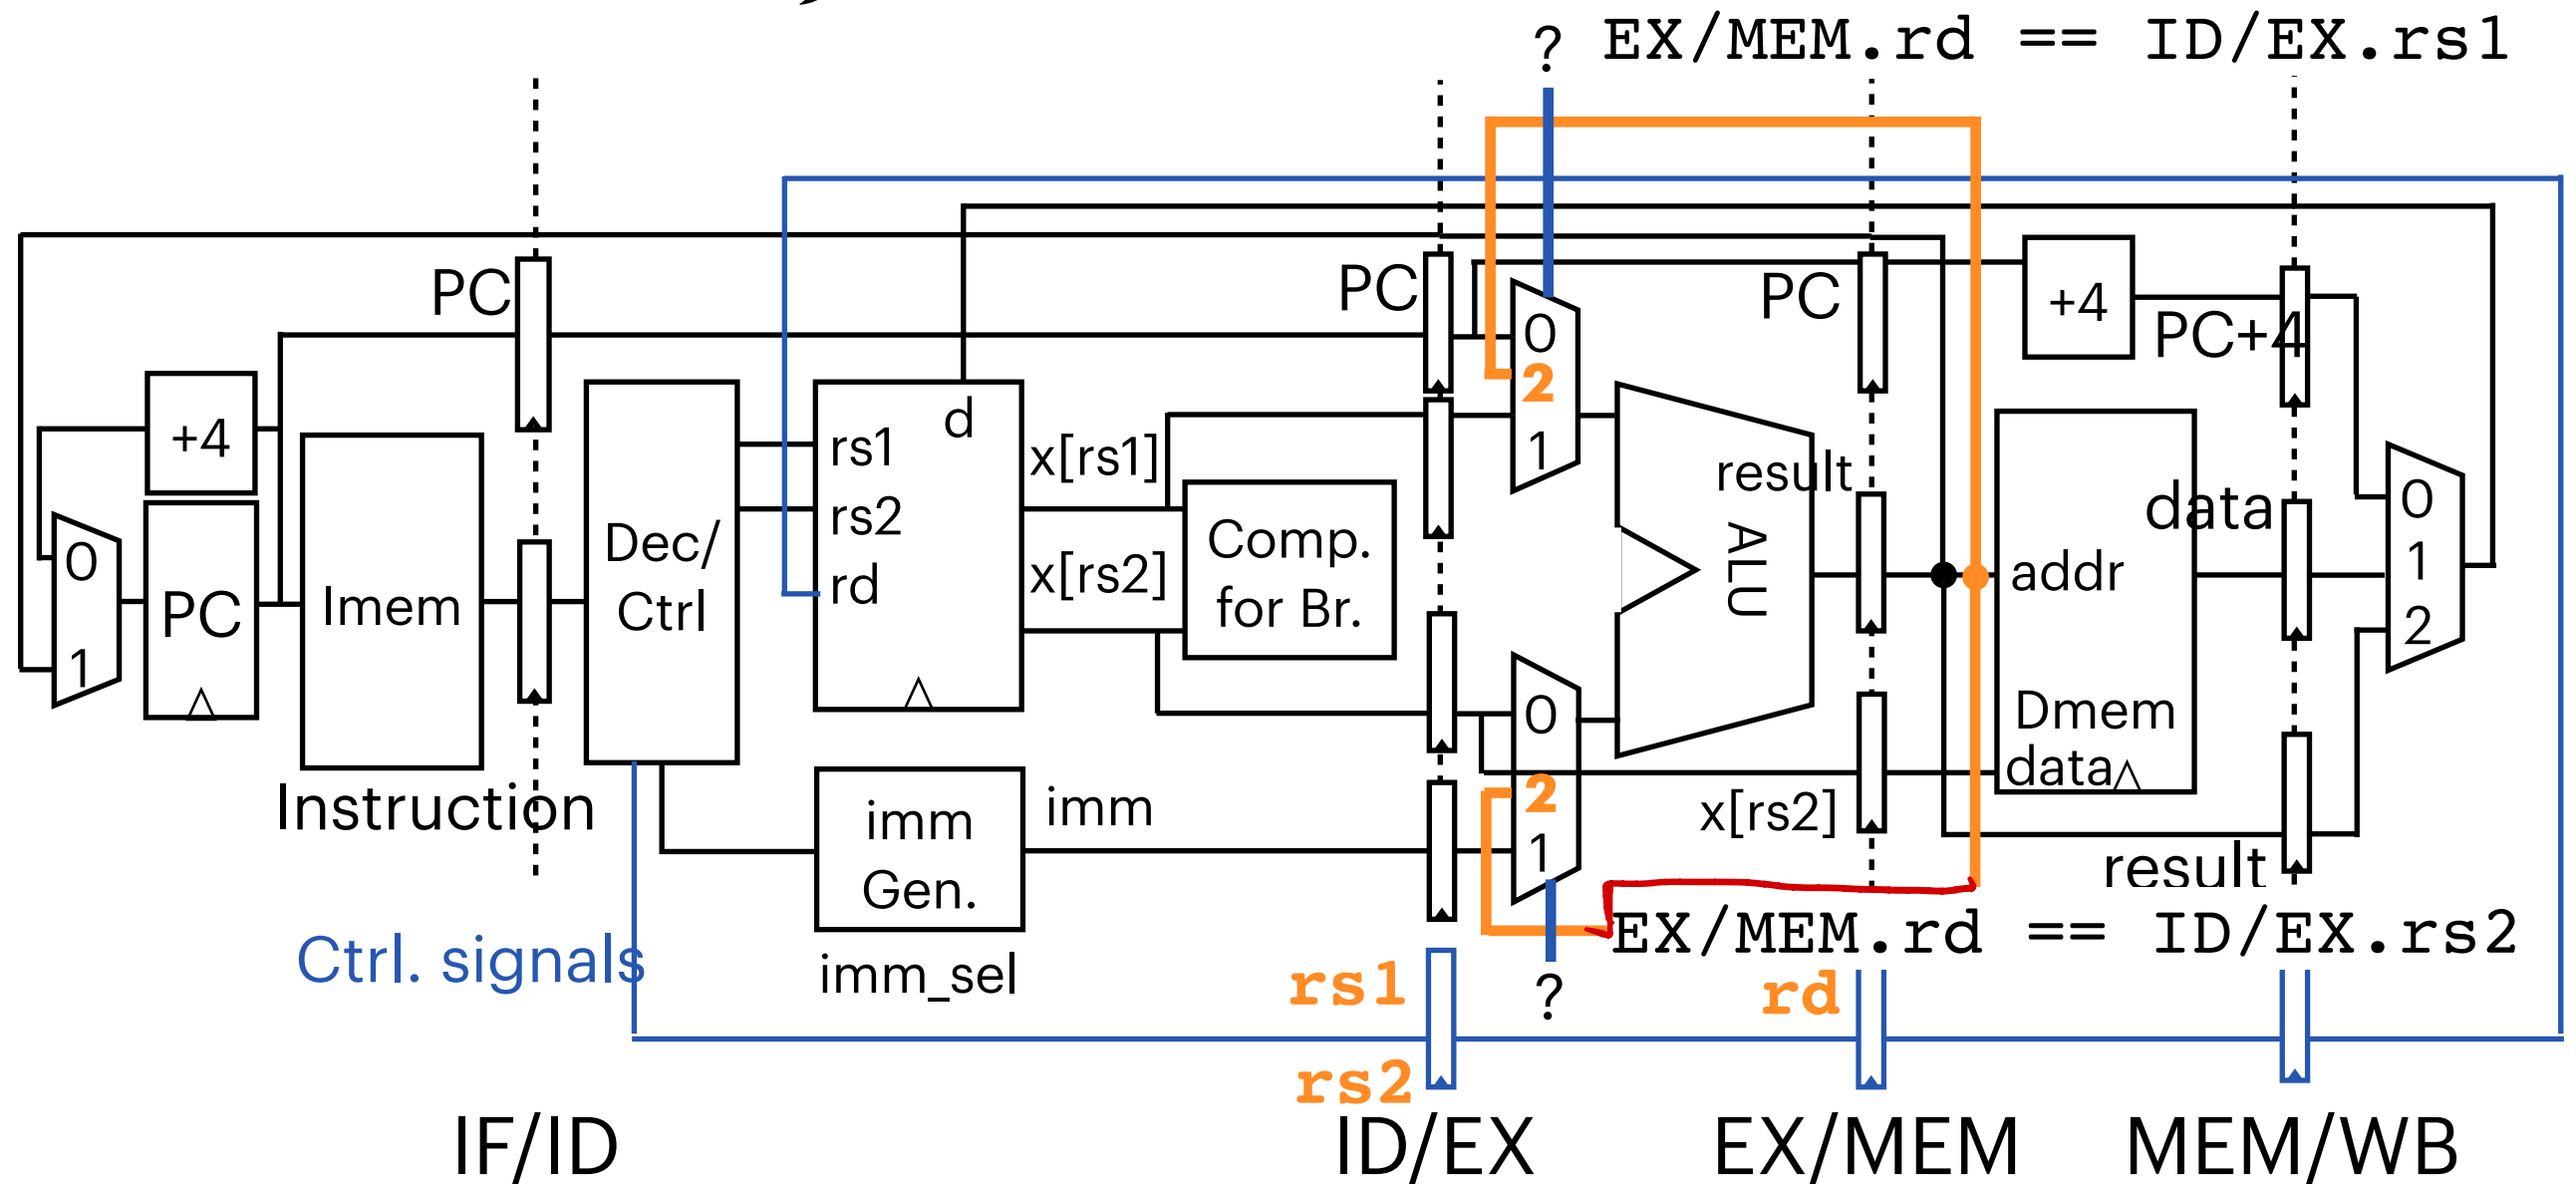


or `t3, t4, t5`

`sw t0, 4(t3)`

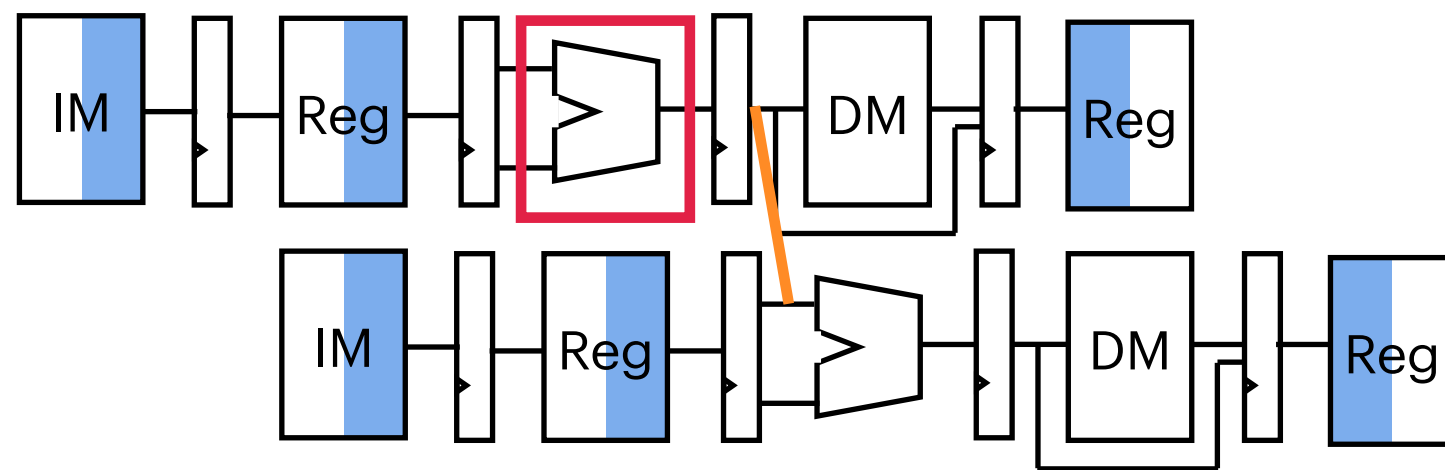


# Adjust the Datapath



or `t3, t4, t5`

`add t0, t2, t3`

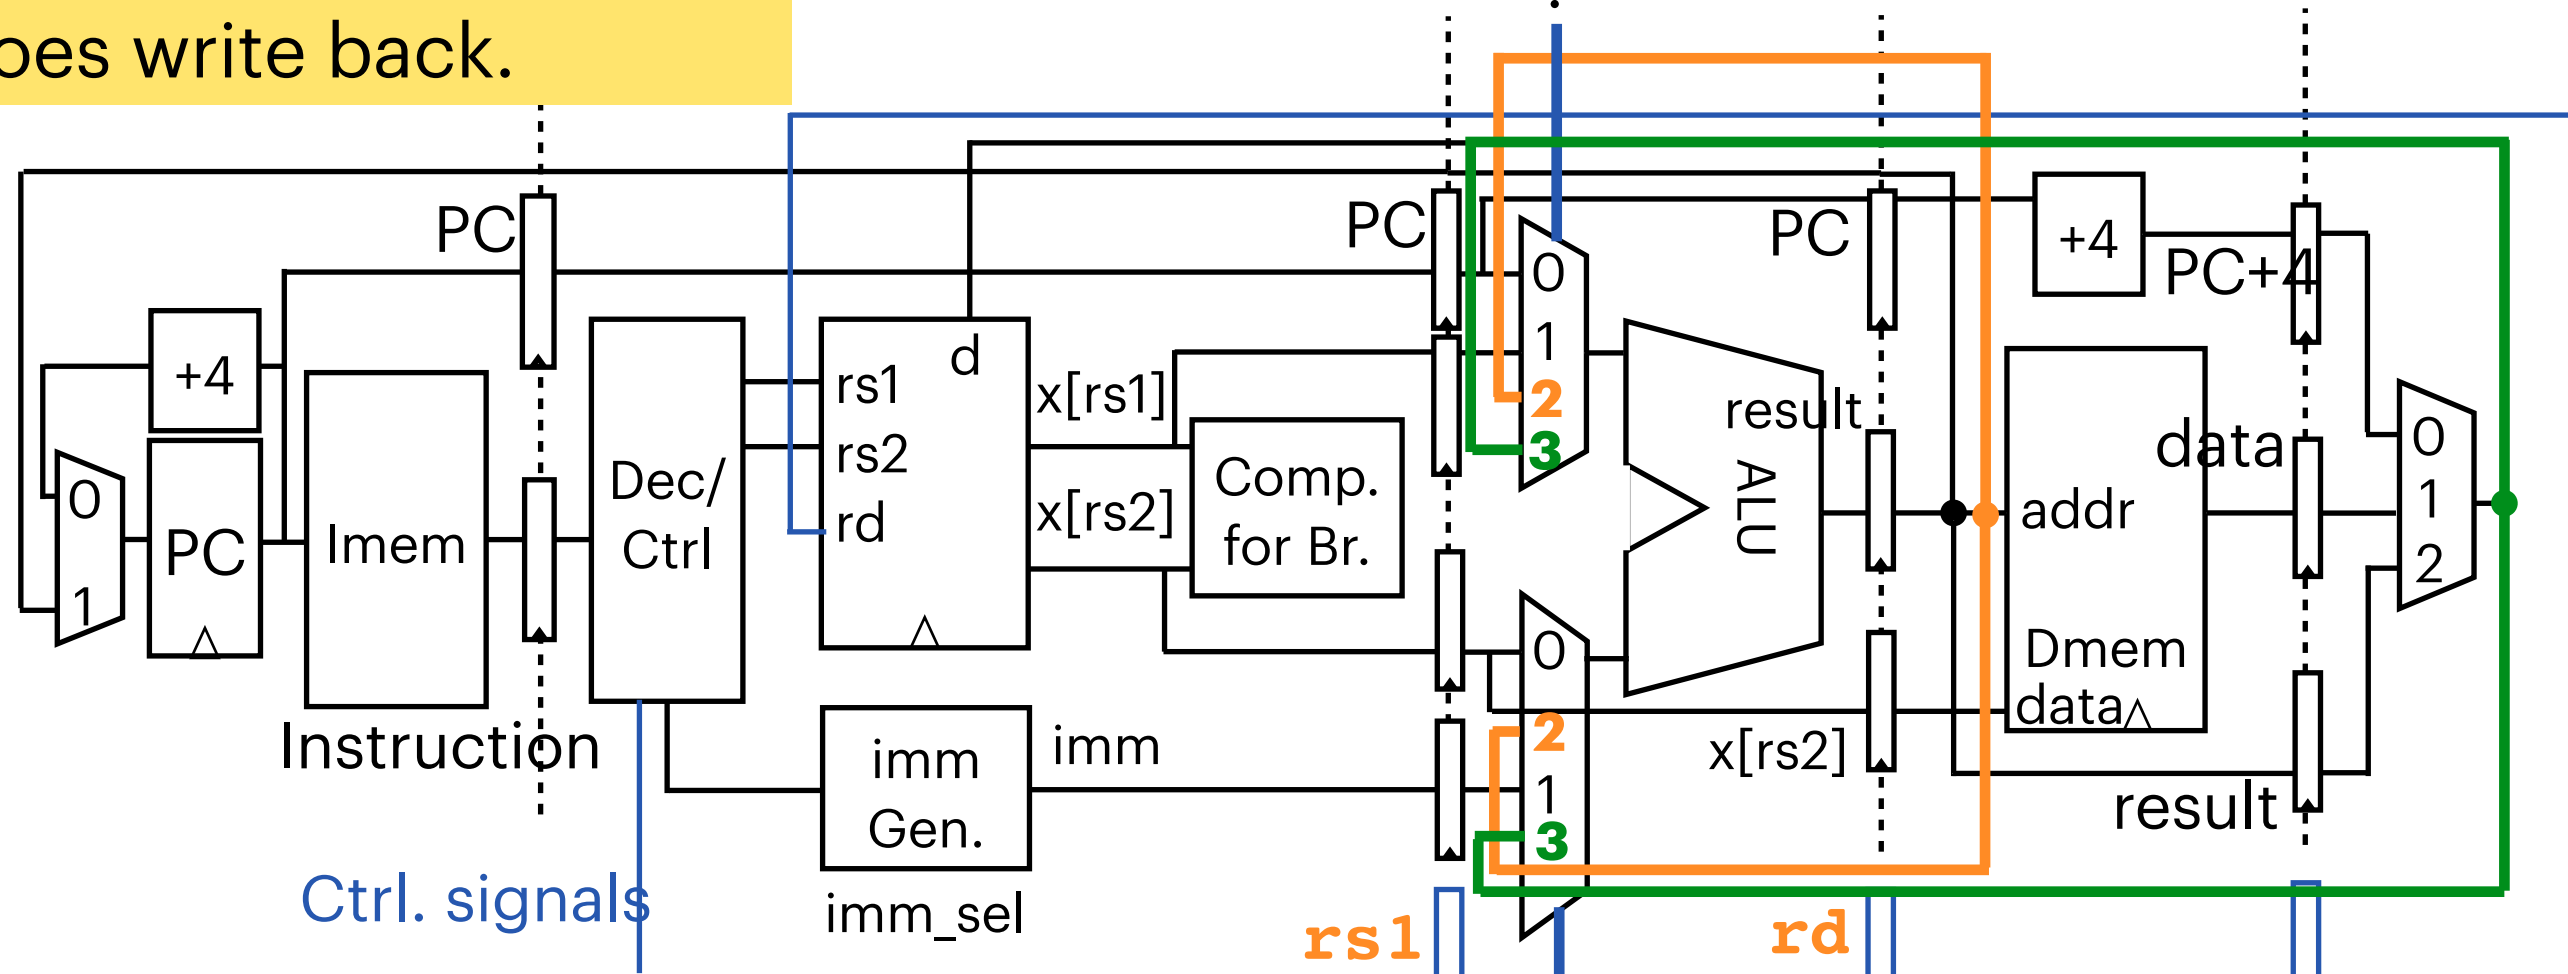


Also have to make sure the previous instruction does write back.

# What if?

MEM/WB.rd == ID/EX.rs1

EX/MEM.rd == ID/EX.rs1



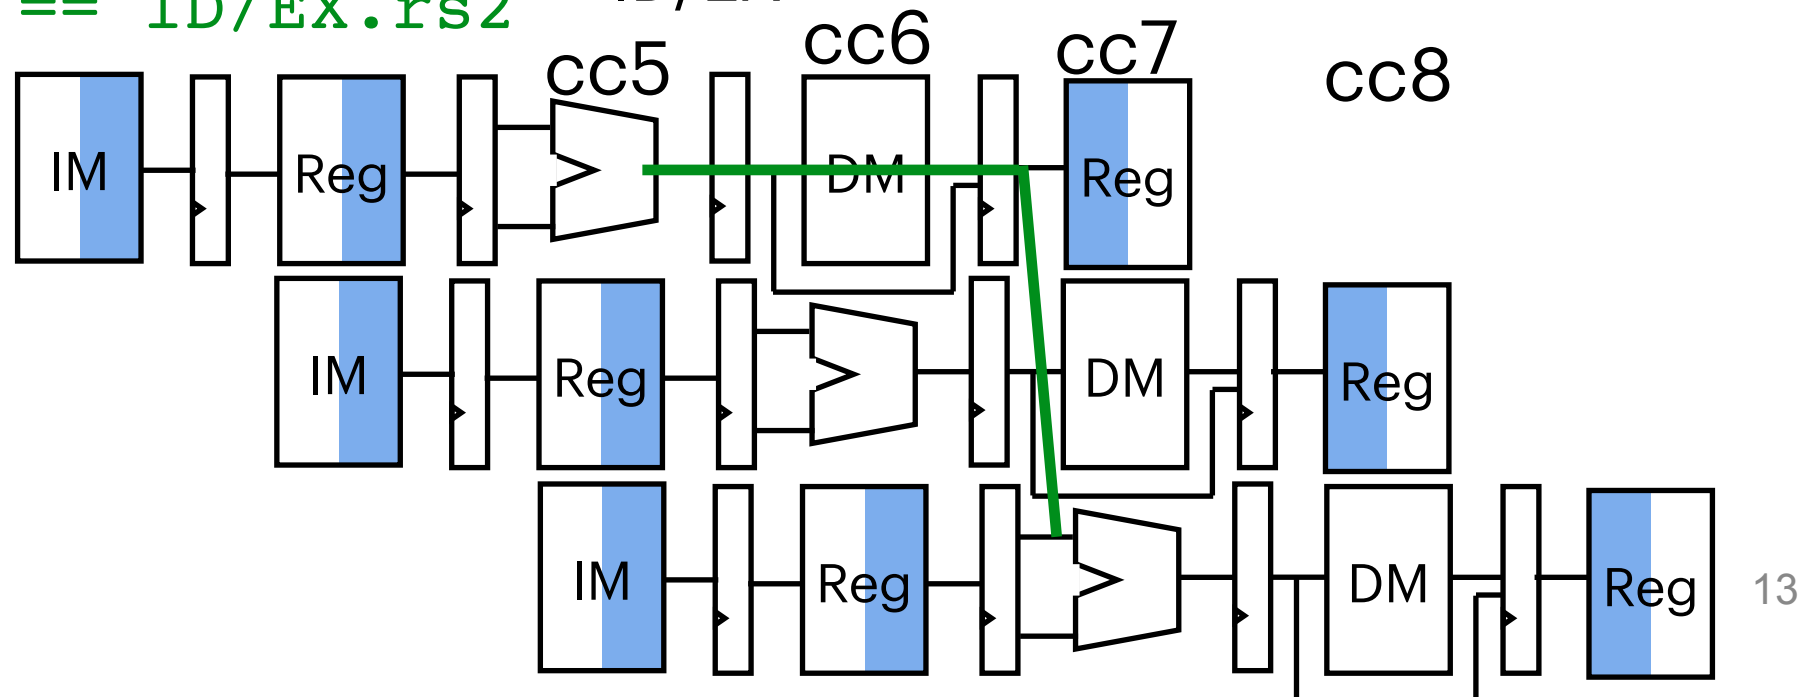
EX/MEM.rd == ID/EX.rs2

MEM/WB.rd == ID/EX.rs2

or t3, t4, t5

One irrelevant instruction  
e.g. sll t1, t2, t6

add t0, t2, t3



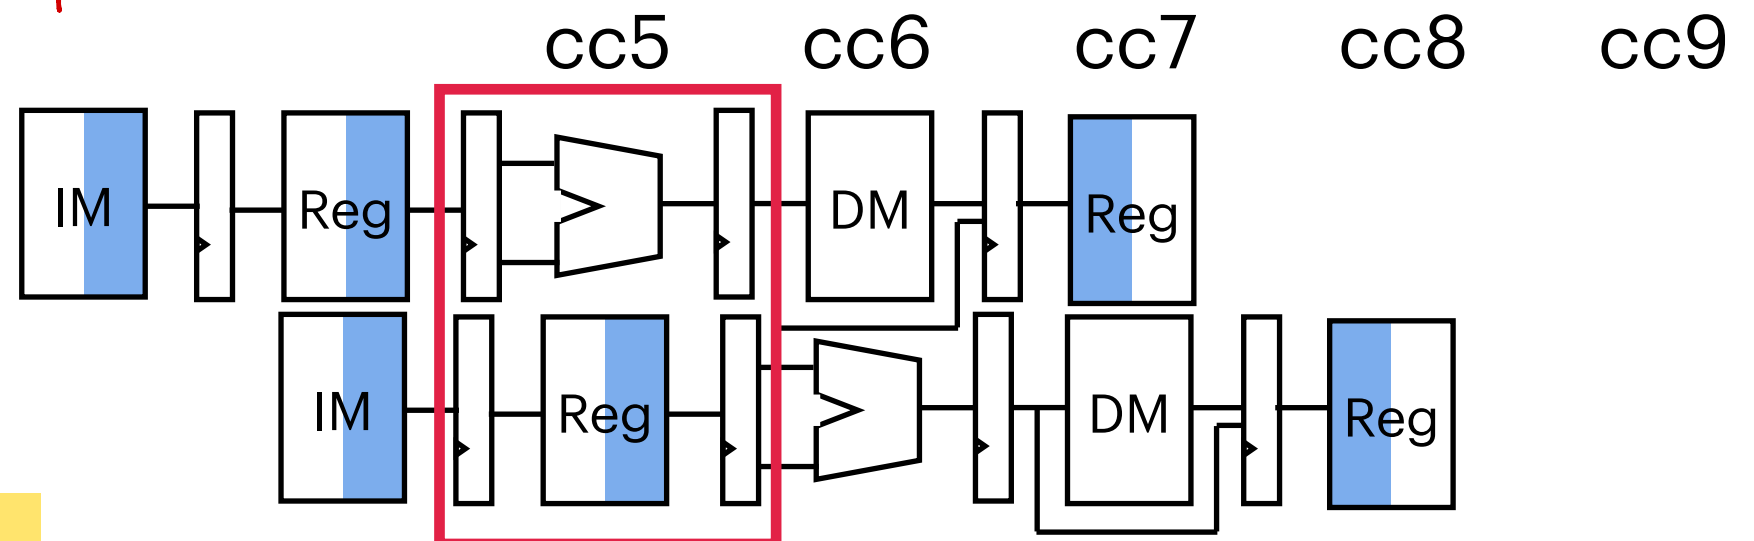
# Extra Considerations

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0			
funct7				rs2			rs1		funct3		rd			opcode		R-type	
imm[11:0]						rs1		funct3		rd			opcode		I-type		
imm[11:5]				rs2			rs1		funct3		imm[4:0]			opcode		S-type	
imm[12]		imm[10:5]			rs2			rs1		funct3		imm[4:1]		imm[11]		opcode	B-type
imm[31:12]										rd			opcode		U-type		
imm[20]		imm[10:1]				imm[11]		imm[19:12]			rd			opcode		J-type	

*imm5 会被当作 rd*

sw t0,28(t1)

add t0,t2,t3



EX/MEM.rd=28

ID/EX.rs2=t3=x28=28

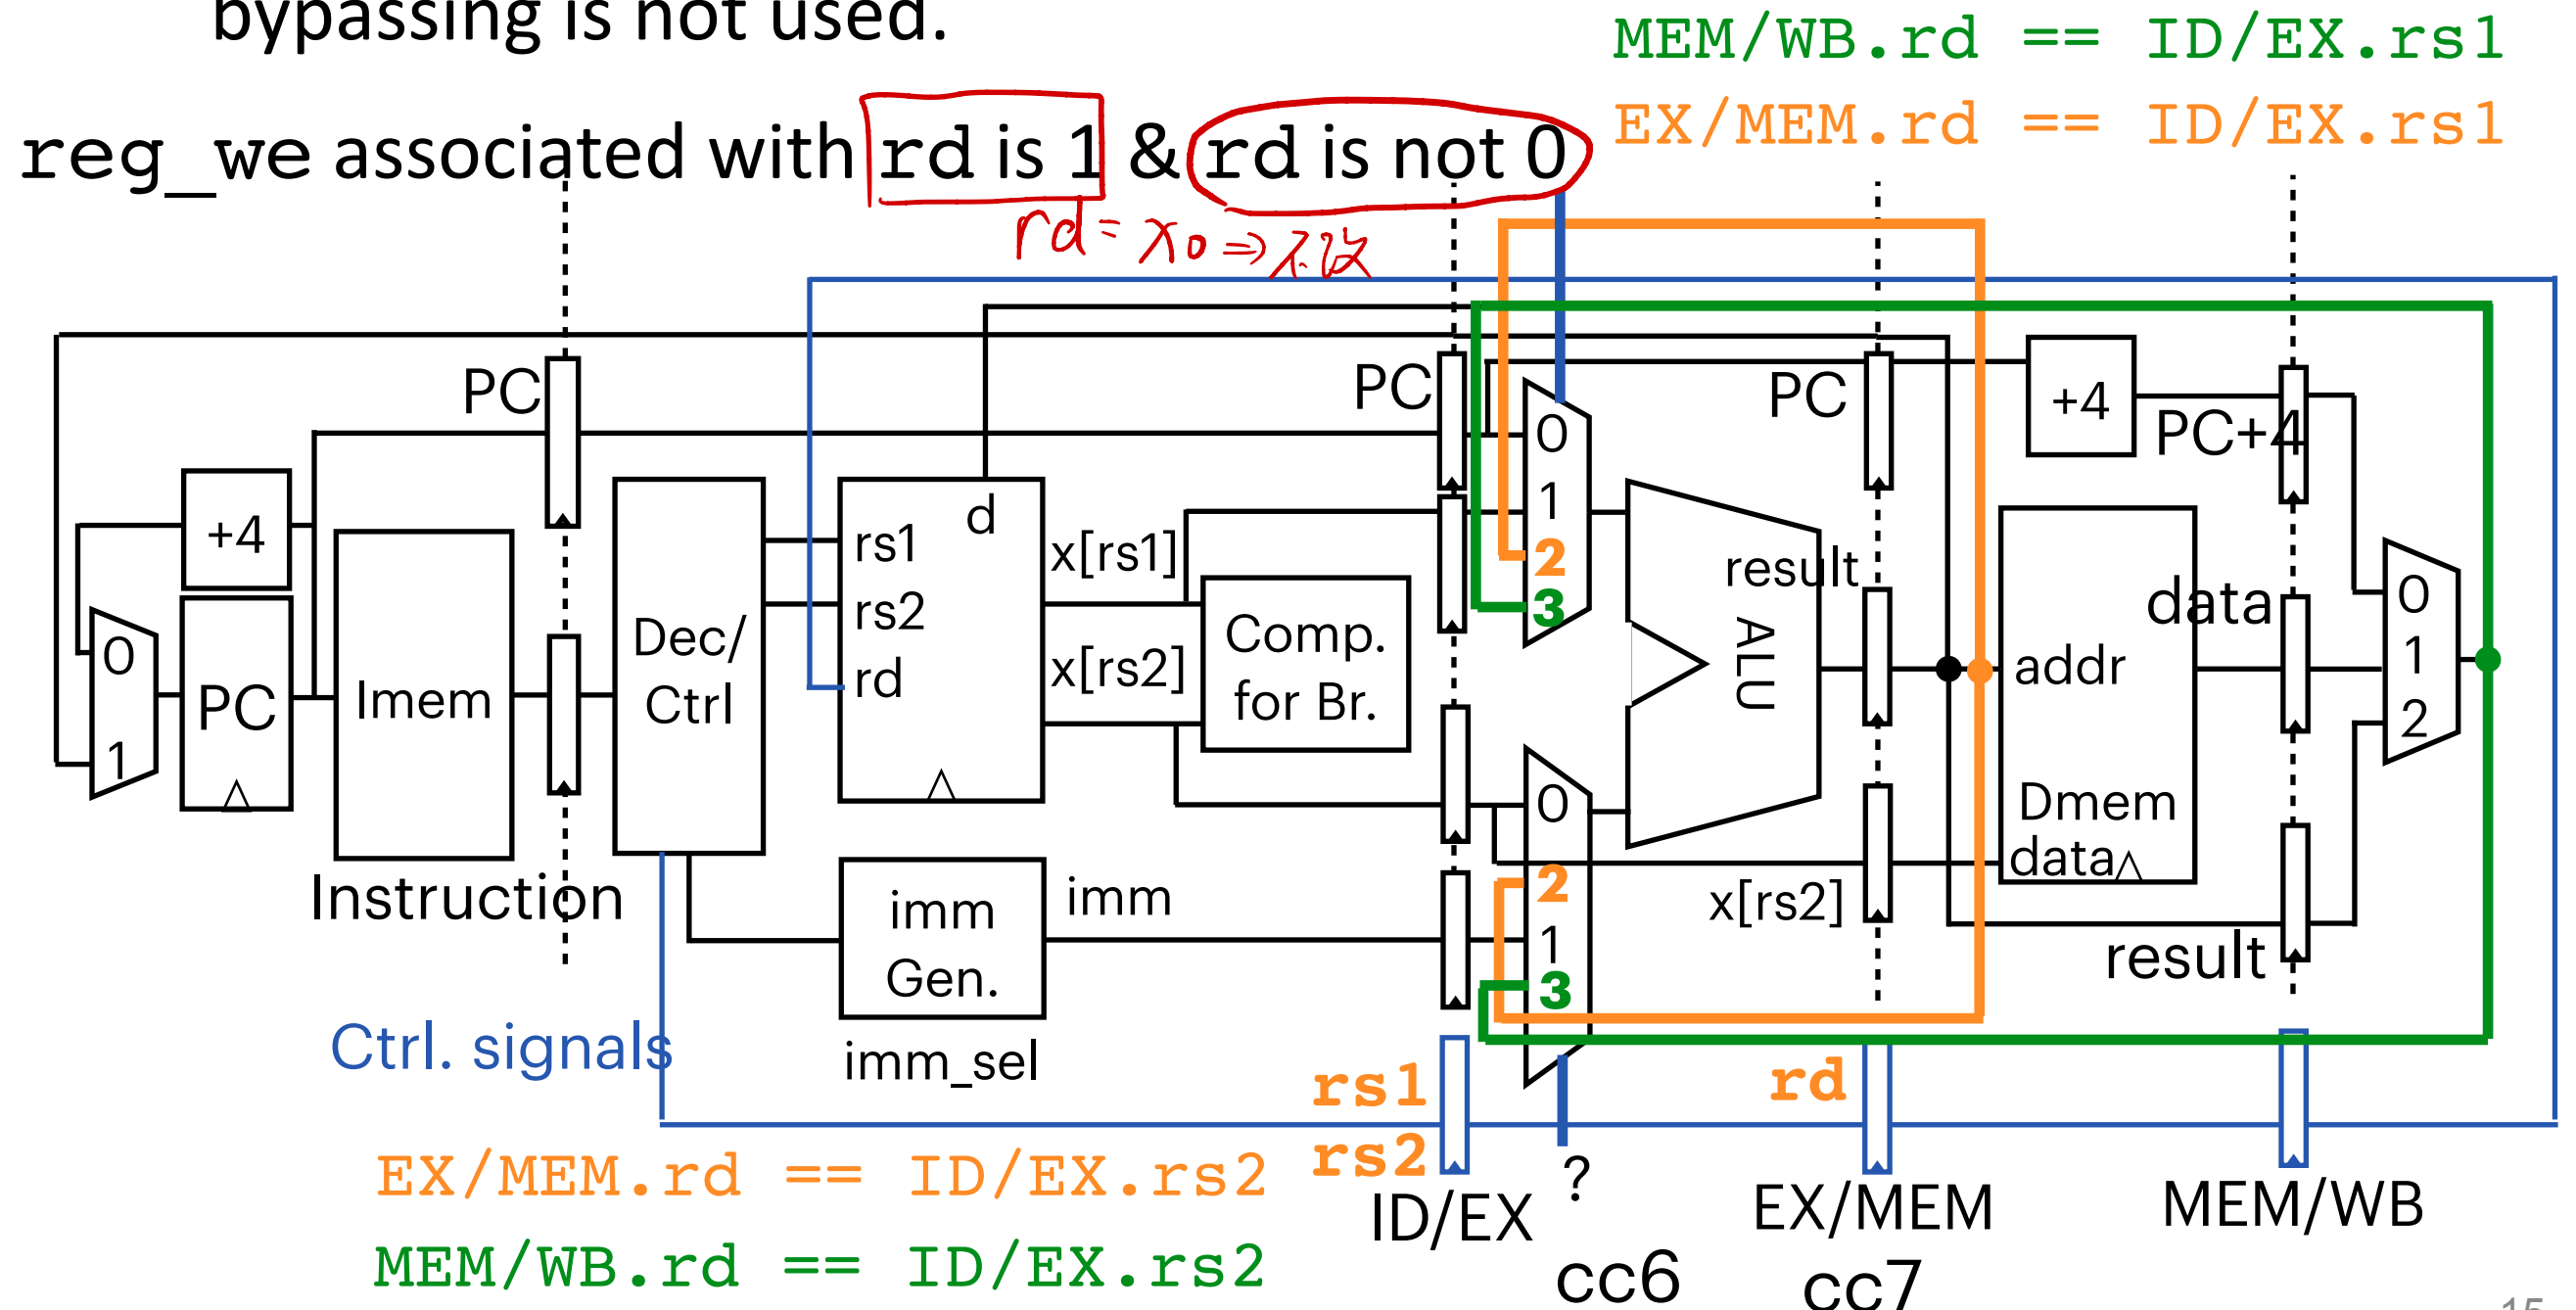


False positive

Also have to make sure the previous instruction does write back.

# What about x0?

x0 register stores zero, and its value won't be modified, i.e., if **rd** is **x0**, there will never be data hazard on **x0**. Thus, forwarding/bypassing is not used.

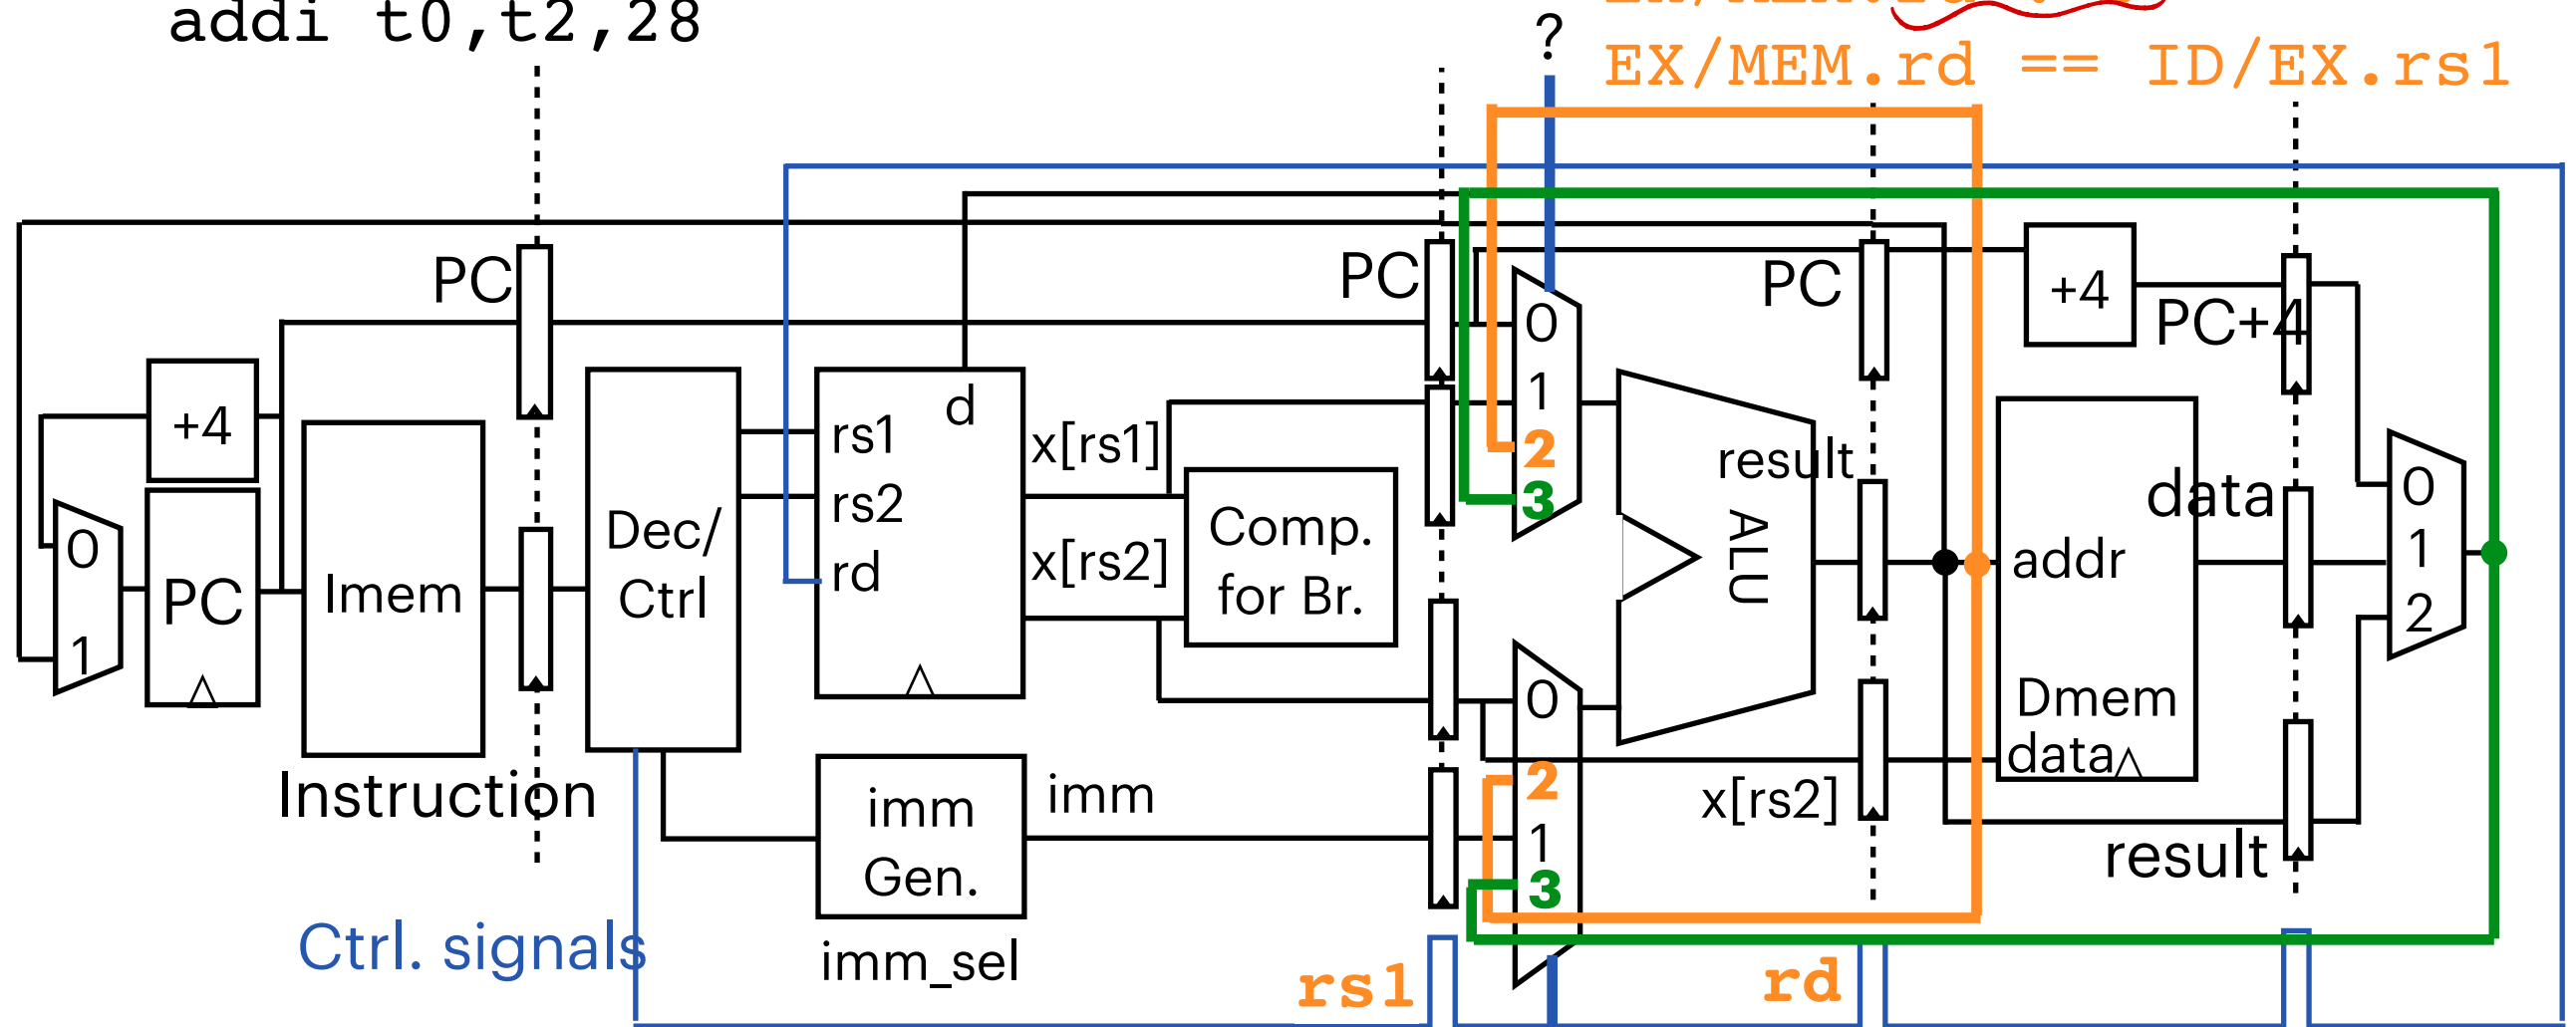


# Extra Considerations

What about the input operands?

Do we have to check either?

```
add t3,t4,t5
addi t0,t2,28
```



EX/MEM.reg\_we == 1

EX/MEM.rd != 0

EX/MEM.rd == ID/EX.rs2

MEM/WB.reg\_we == 1

MEM/WB.rd != 0

MEM/WB.rd == ID/EX.rs1

EX/MEM.reg\_we == 1

EX/MEM.rd != 0

EX/MEM.rd == ID/EX.rs1

ID/FX

EX/MEM

MEM/WB

MEM/WB.reg\_we == 1

MEM/WB.rd != 0

MEM/WB.rd == ID/EX.rs2



# Extra Considerations

What about the input operands?

Do we have to check either?

add t3,t4,t5

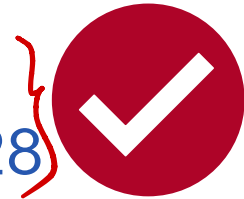
addi t0,t2,28

*imm  
rs2* } 数值相同

×

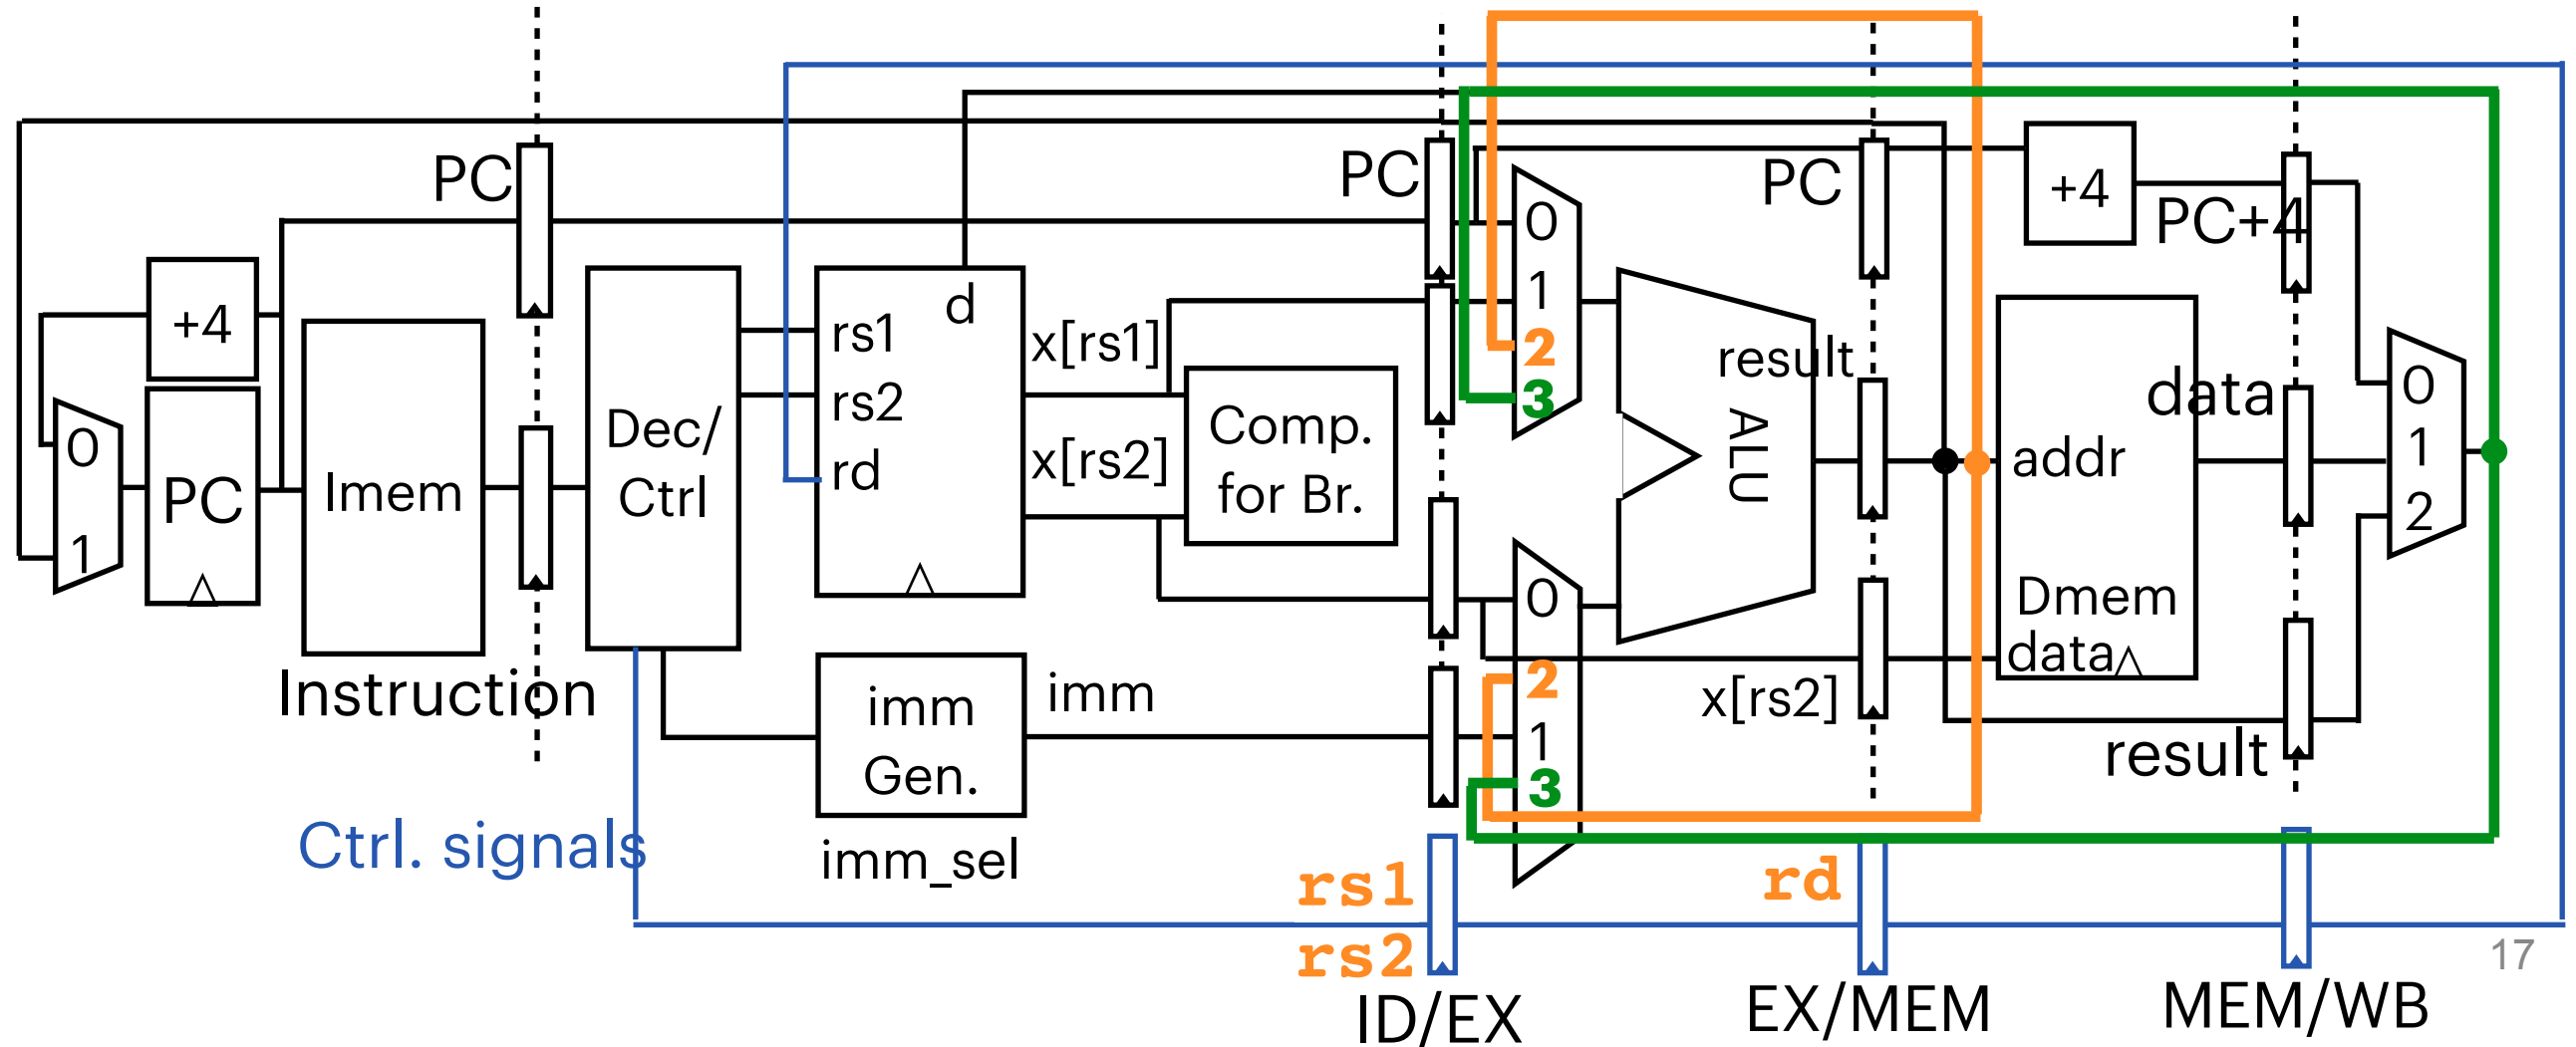
EX/MEM.rd=28

ID/EX.rs2=t3=x28=28



False positive

31	30	25	24	21	20	19	15	14
funct7		rs2				rs1	funct3	
imm[31:11]		rs2				rs1	funct3	
imm[11:5]		rs2				rs1	funct3	
imm[12]	imm[10:5]		rs2				rs1	funct3
imm[31:12]								
imm[20]		imm[10:1]			imm[11]	imm[19:12]		



# Extra Considerations

What about the input operands?

Do we have to check either?

add t3,t4,t5

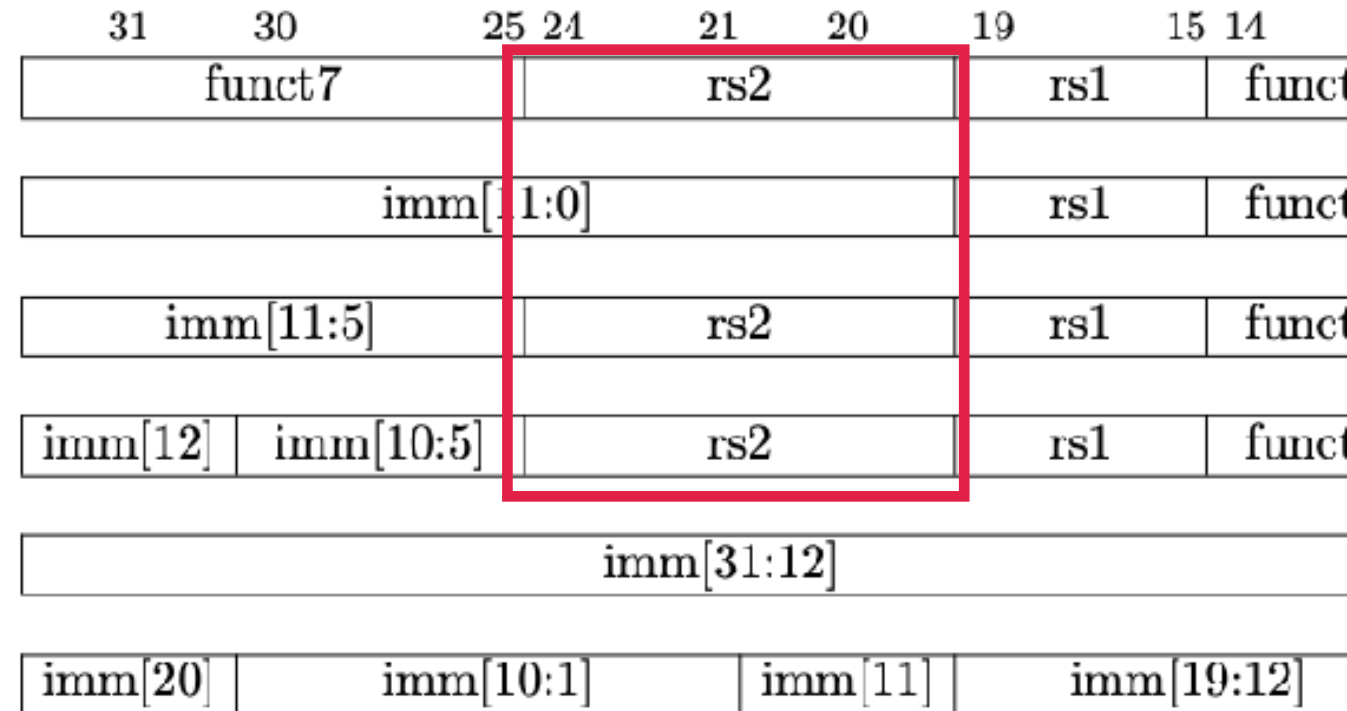
addi t0,t2,28

EX/MEM.rd=28

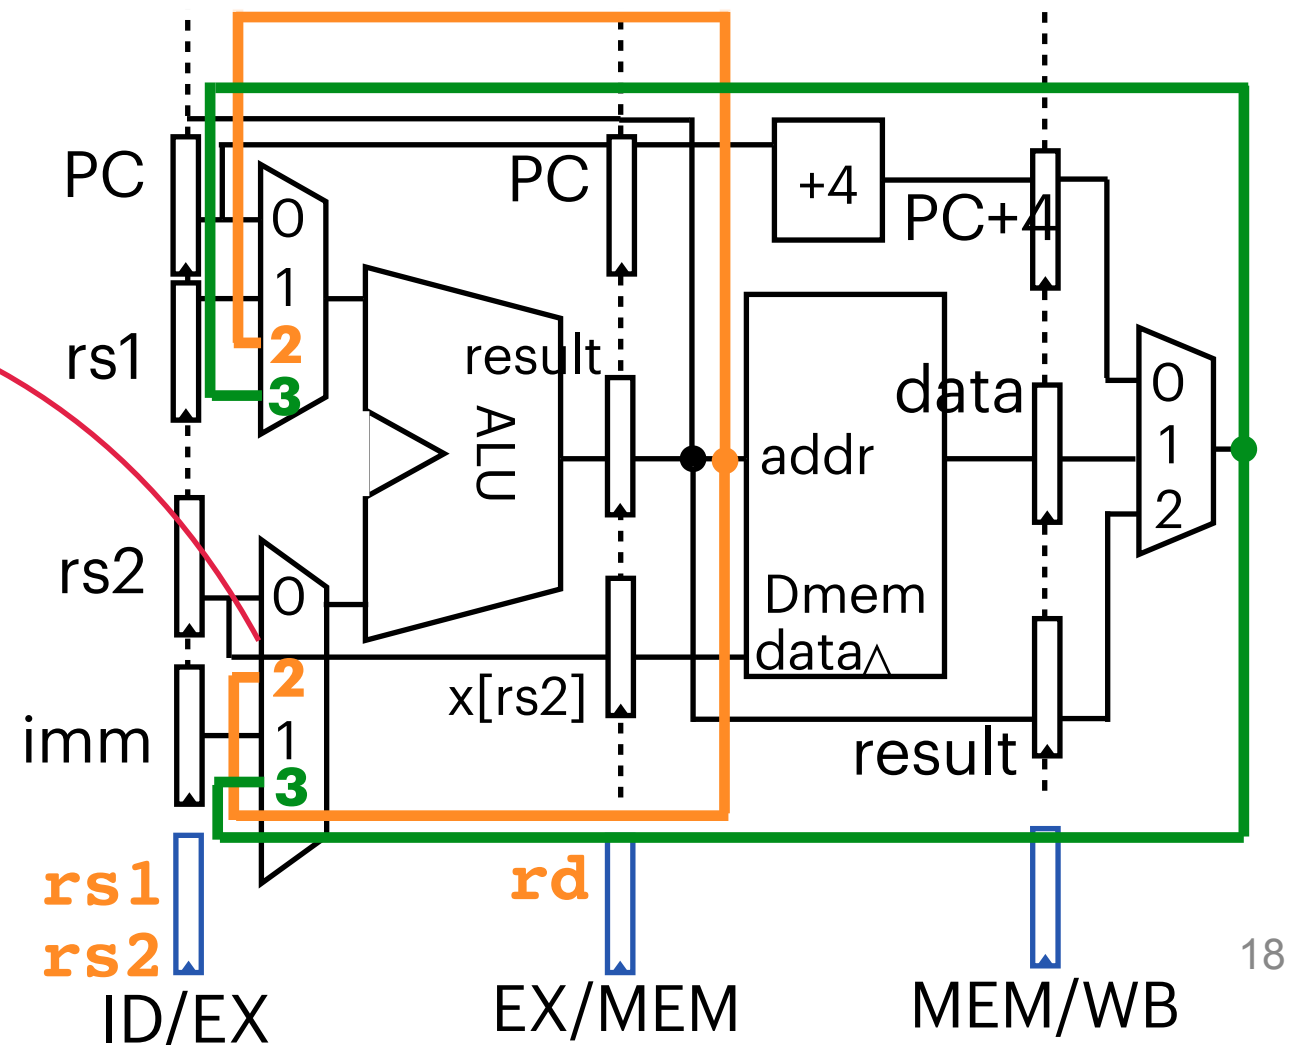
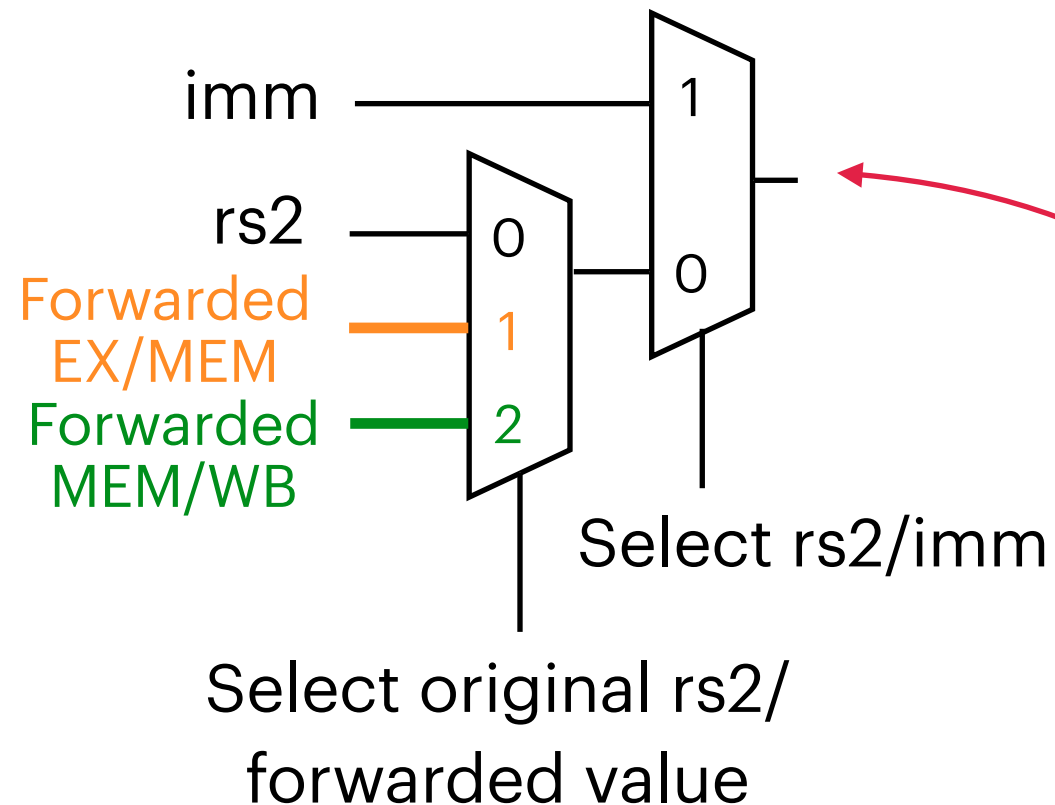
ID/EX.rs2=t3=x28=28



False positive



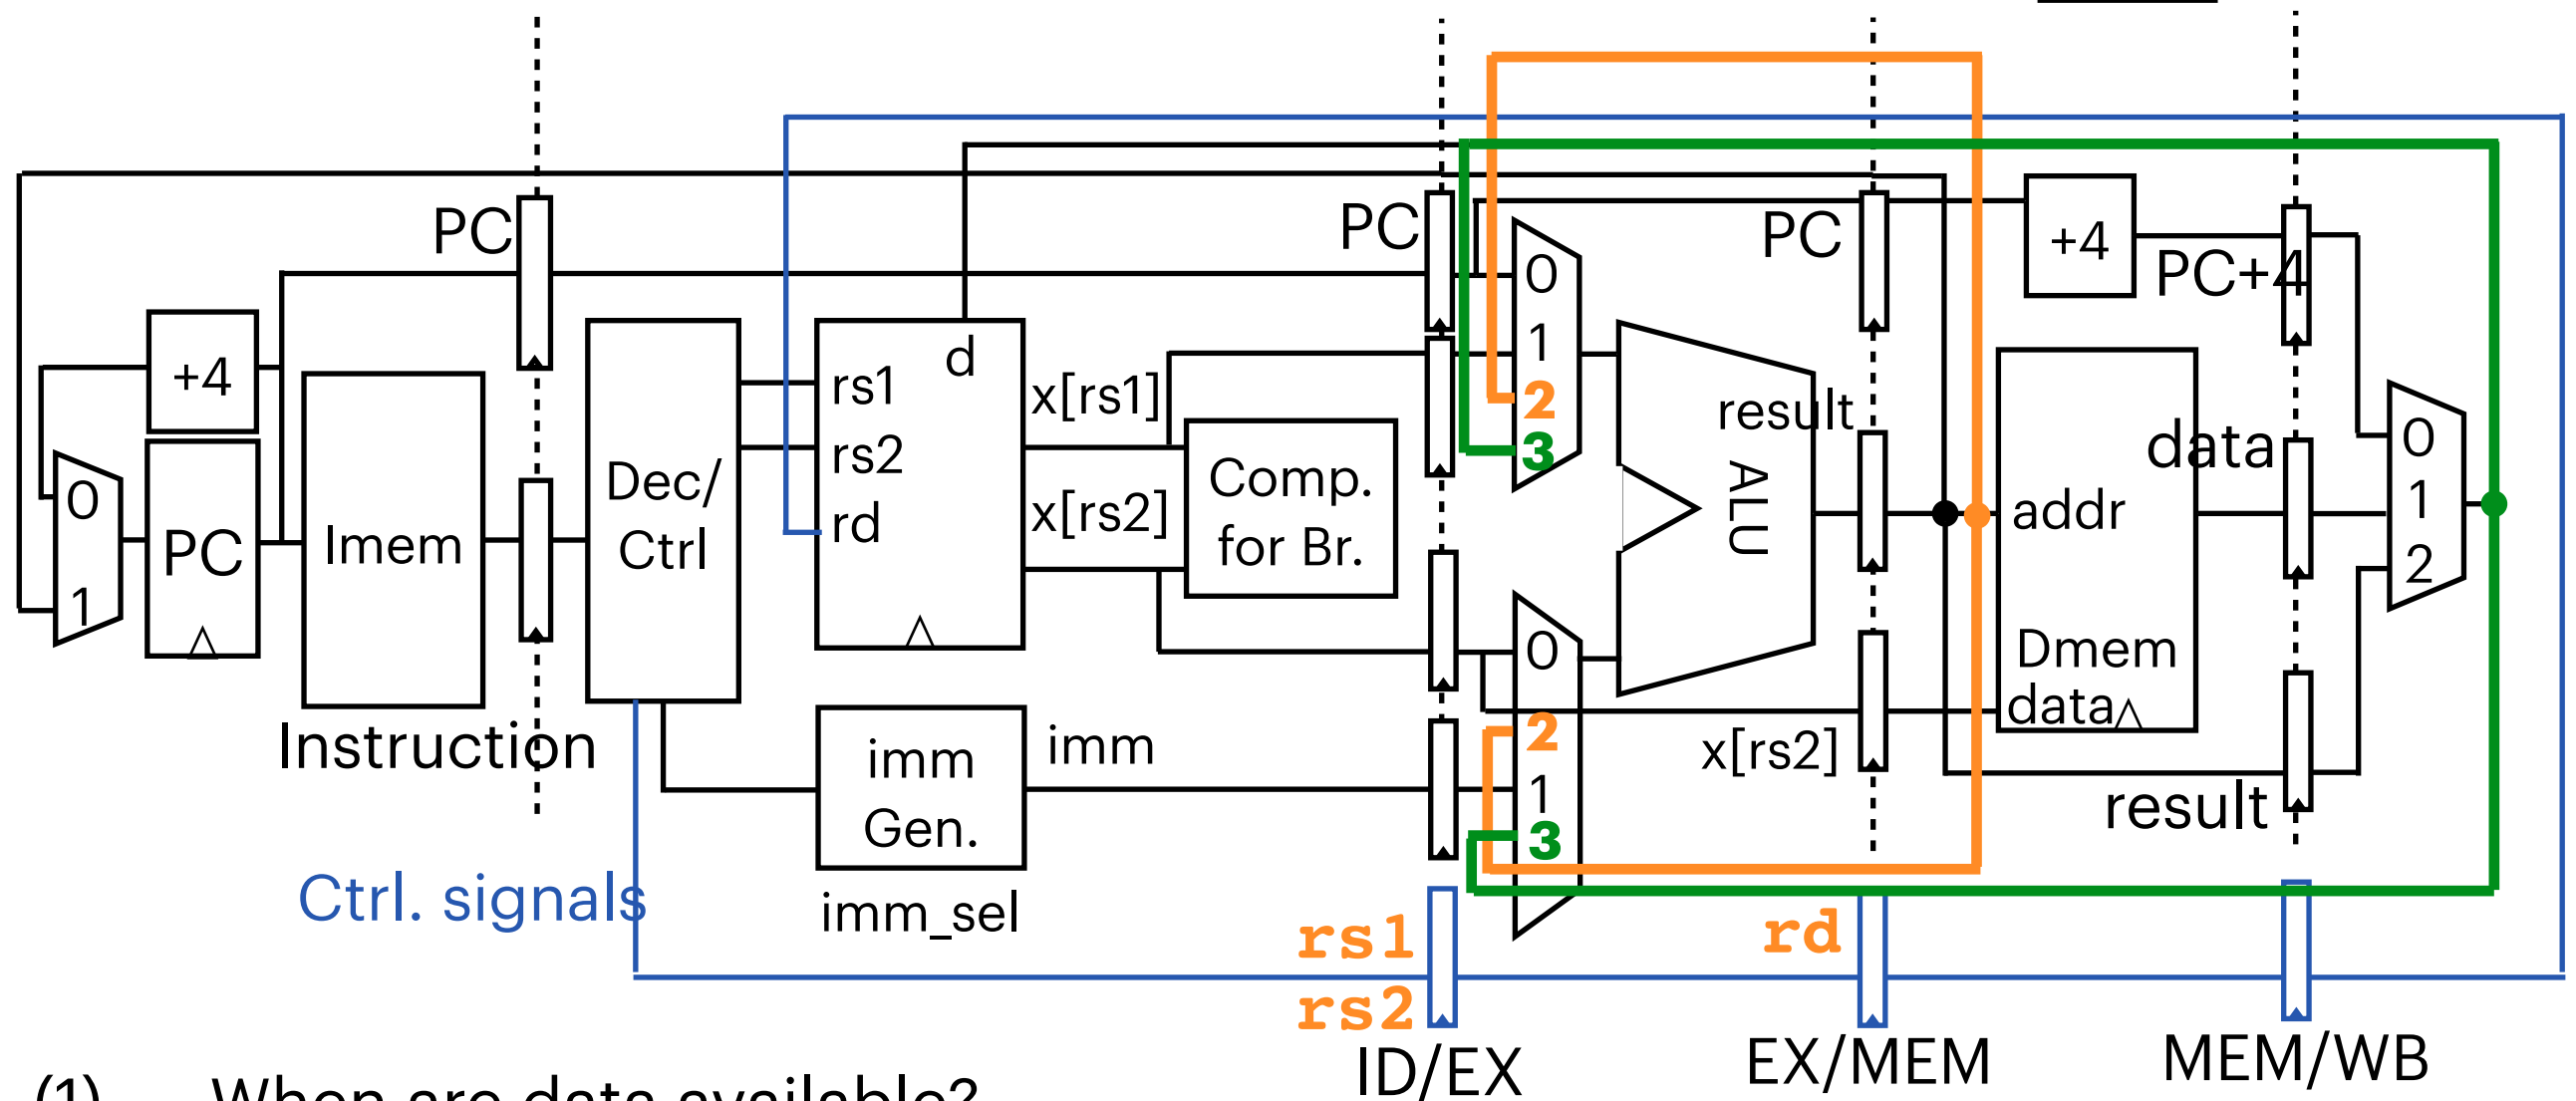
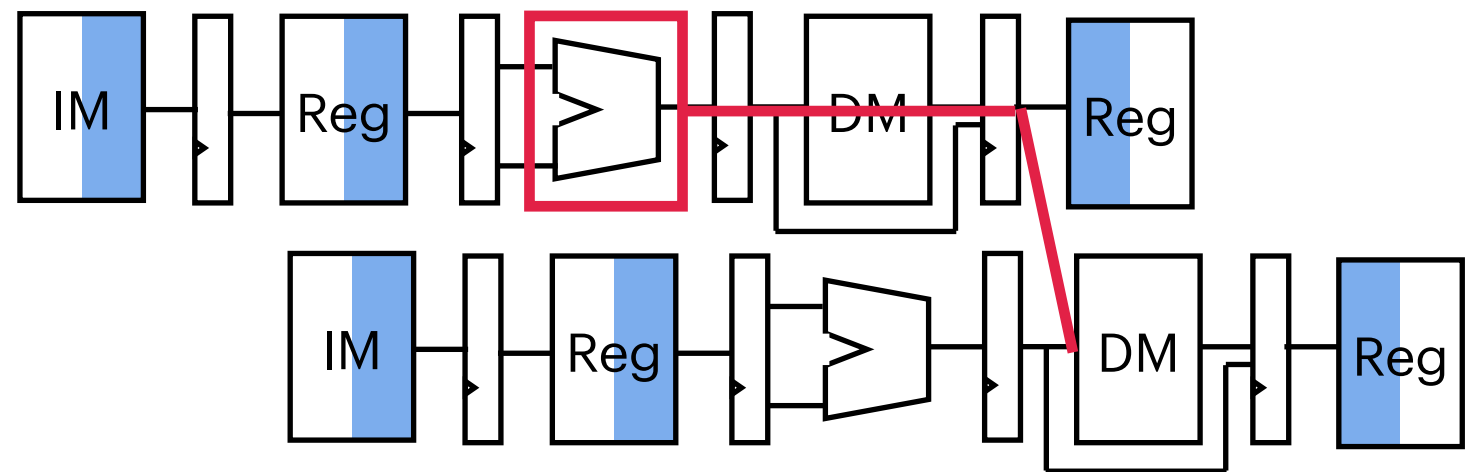
Give priority to rs2/imm judgement



# Forwarding also Resolves sw Hazards

add t3,t1,t2

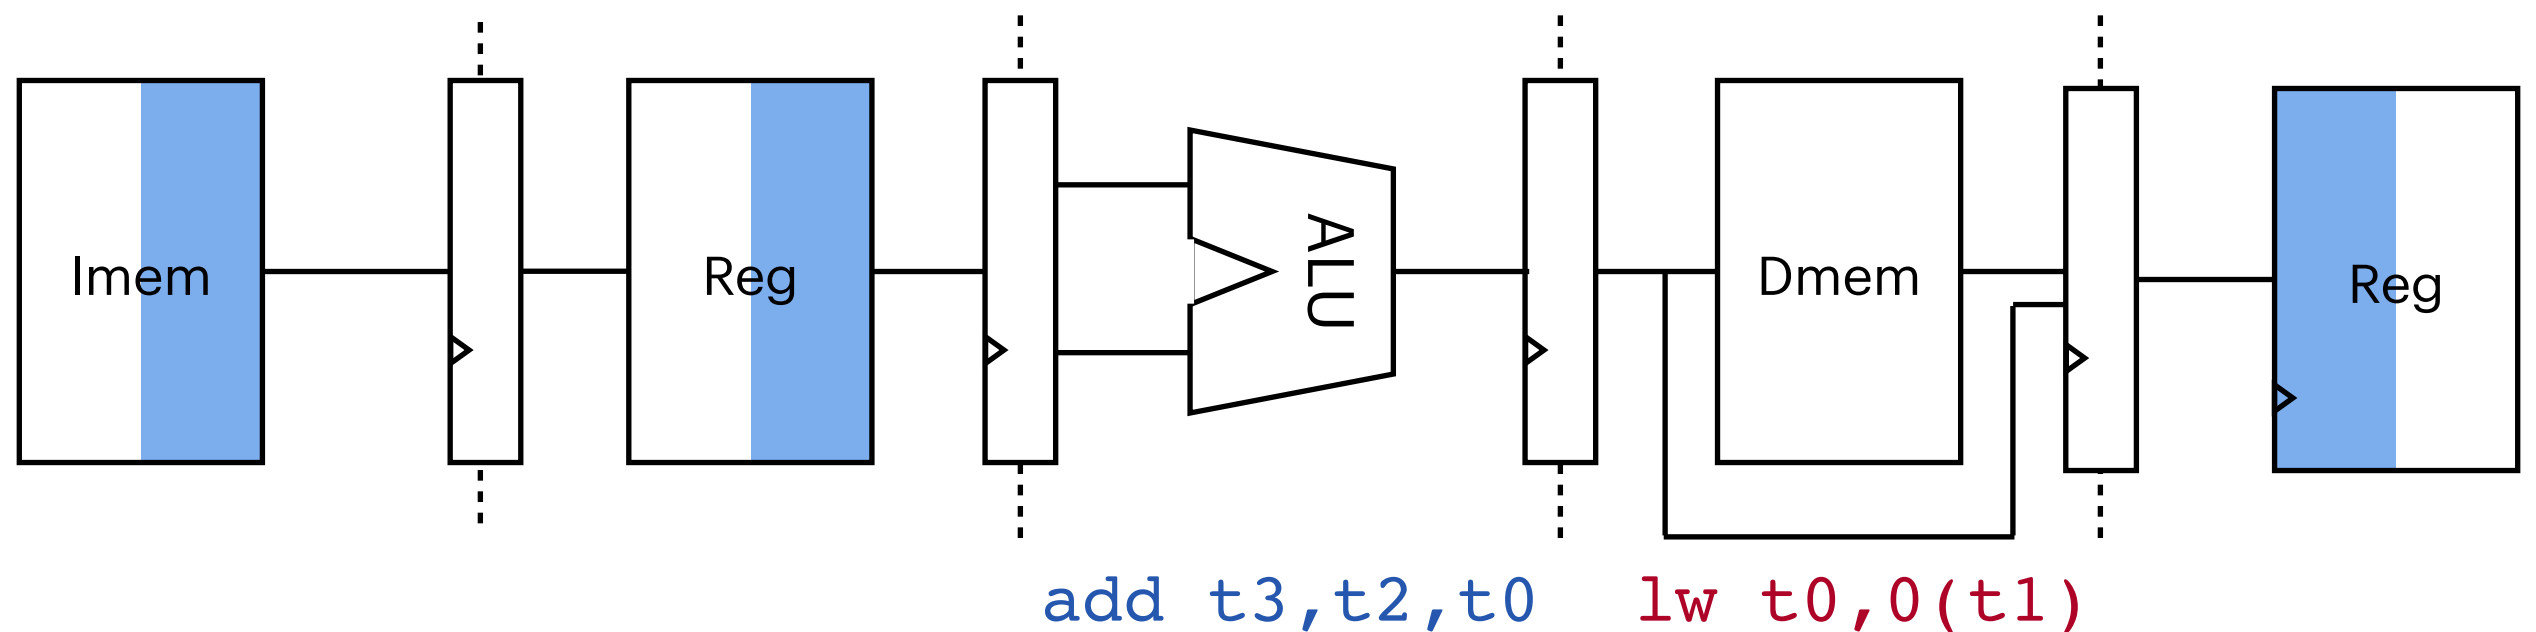
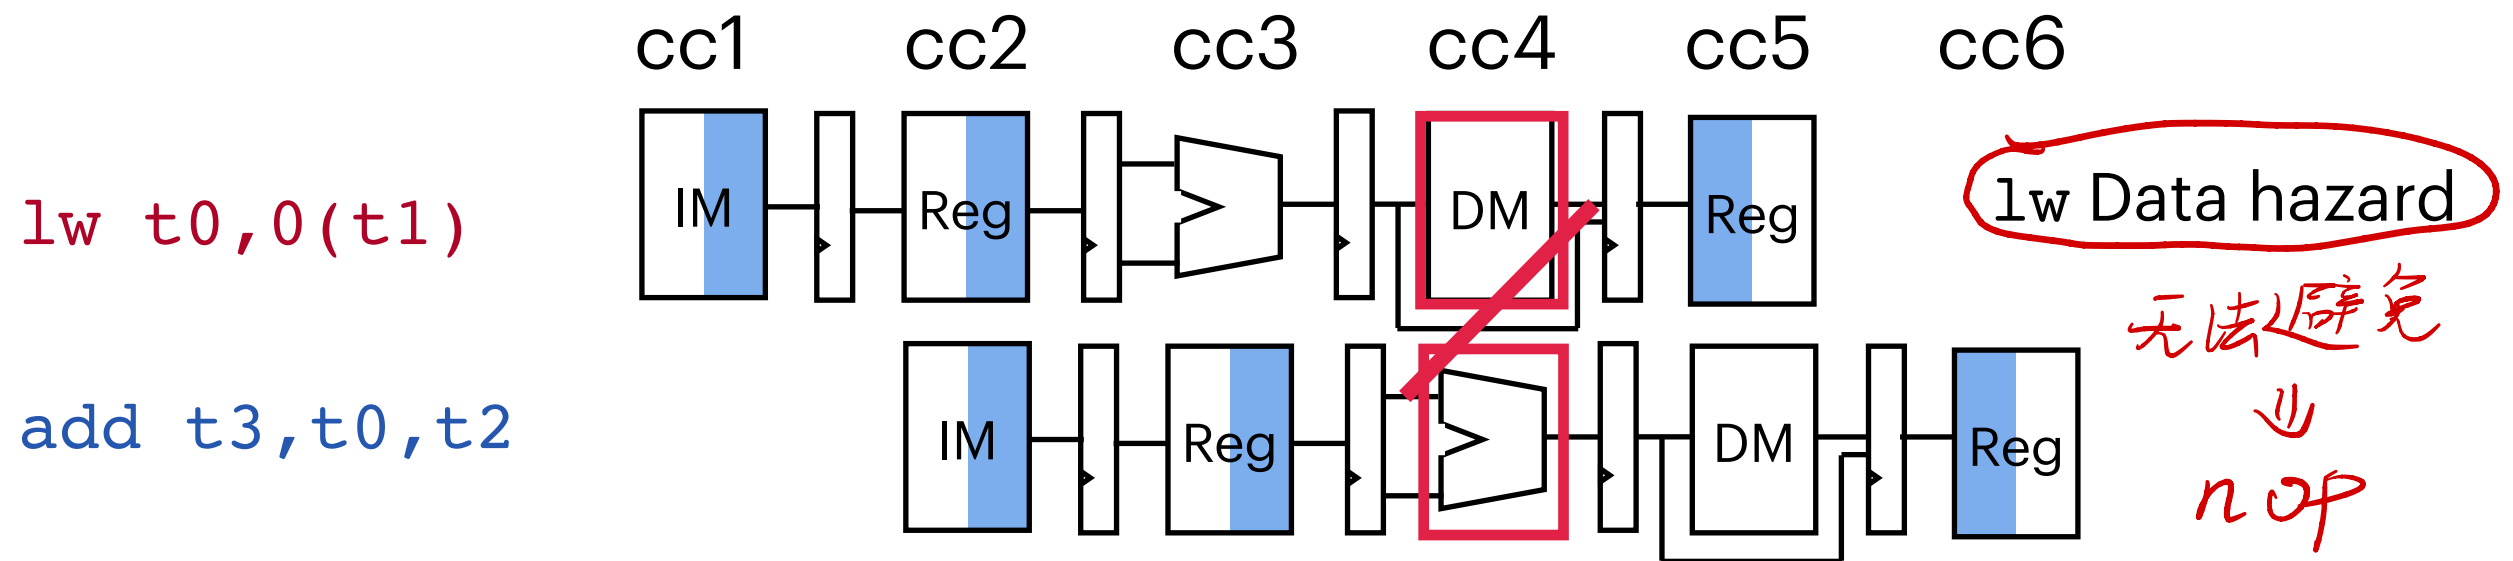
sw t3,t4,~~t5~~ octx)



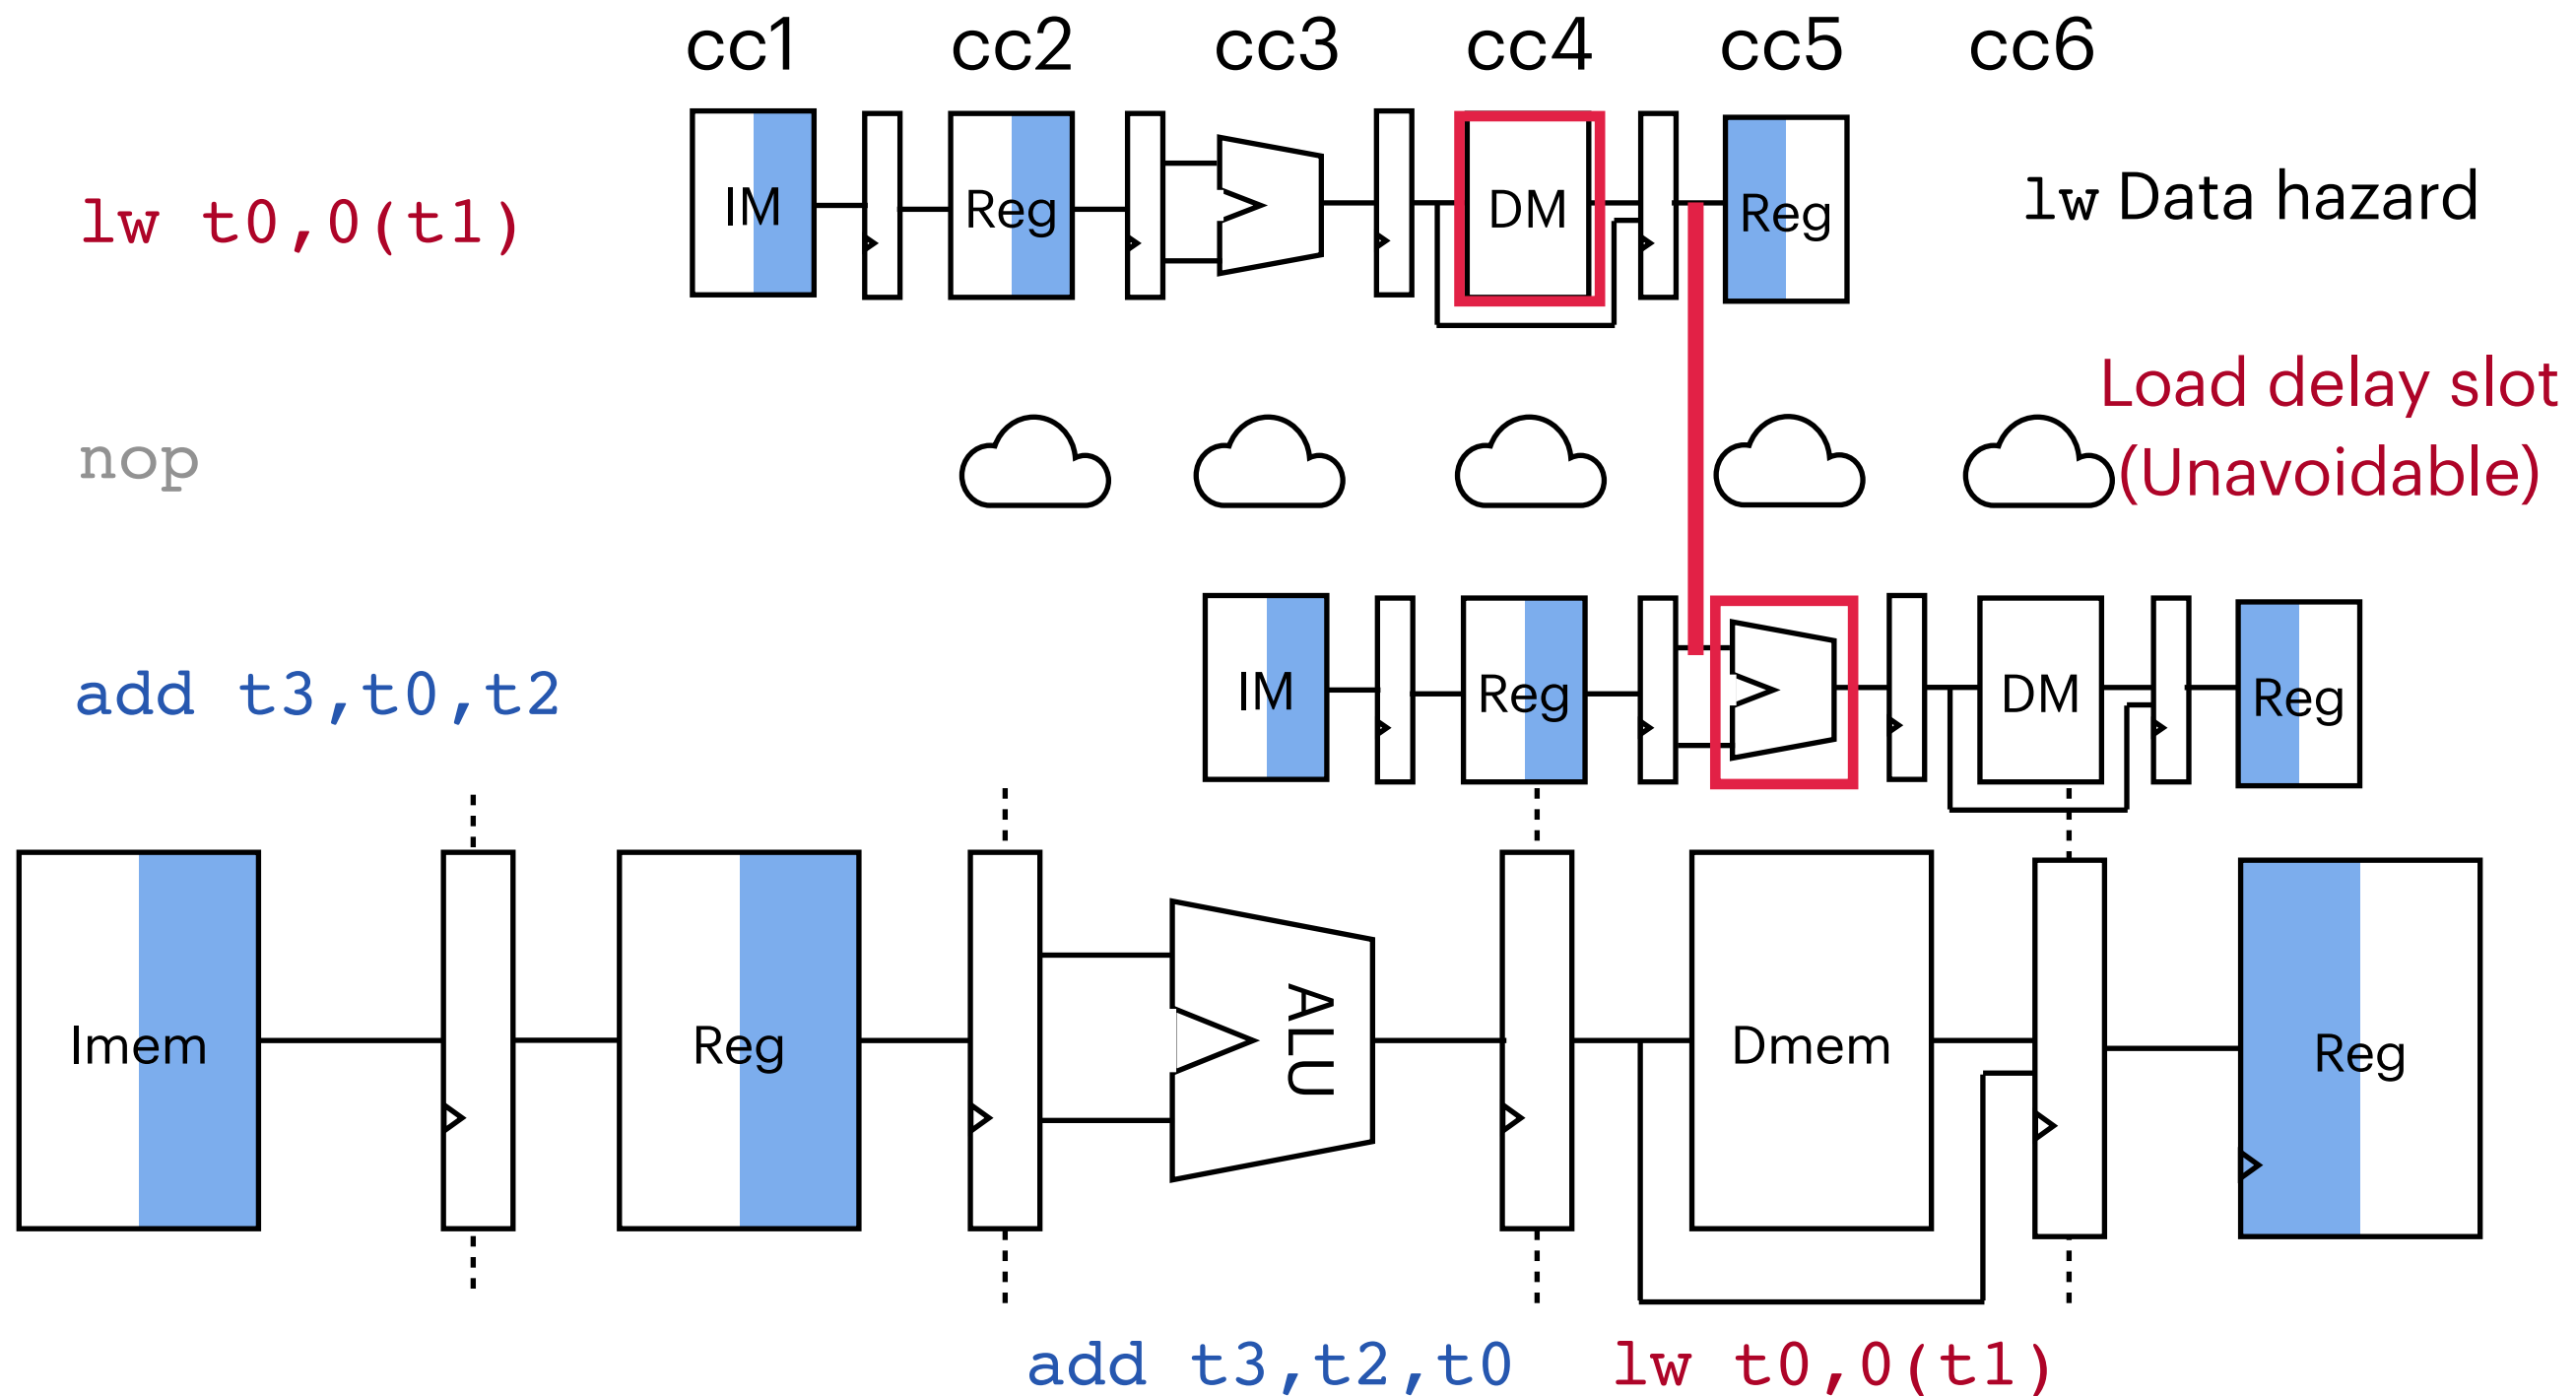
- (1) When are data available?
- (2) When are data consumed?

**Try yourself!**

# Extra Consideration on lw



# Extra Consideration on lw



# Code Scheduling to Avoid Stalls

Assume C code:  $D=A+B$ ;  $E=A+C$

A, B, C addresses are  $0(t_0)$ ,  $4(t_0)$ ,  $8(t_0)$

D, E addresses are  $12(t_0)$ ,  $16(t_0)$

Naïve RISC-V compiler

```
1 lw t1, 0(t0)
2 lw t2, 4(t0) → 1 lw delay
1 add t3, t1, t2
1 sw t3, 12(t0)
2 lw t4, 8(t0) → 1 lw delay
1 add t5, t1, t4
1 sw t5, 16(t0)
```

A smart compiler reorders the instructions:

```
1 lw t1, 0(t0)
1 lw t2, 4(t0)
1 lw t4, 8(t0)
1 add t3, t1, t2
1 sw t3, 12(t0)
1 add t5, t1, t4
1 sw t5, 16(t0)
```

Q: How many clk cycles?

9 CPI > 1 =  $\frac{9}{7}$

7

# Code Scheduling to Avoid Stalls

Code scheduling: With knowledge of the underlying CPU pipeline, the compiler reorders code to improve performance.

# Up-to-Now: Data Hazards

A hazard is a situation in which a planned instruction cannot execute in the “proper” clock cycle.

1. Structural hazard:

- Hardware does not support access across multiple instructions in the same cycle

2. **Data hazard:**

- Instructions have data dependency
- Occurs when an instruction reads a register before a previous instruction has finished writing to that register
- Solution 1: stall
- Solution 2: forwarding/bypassing
- Solution 3: Code scheduling



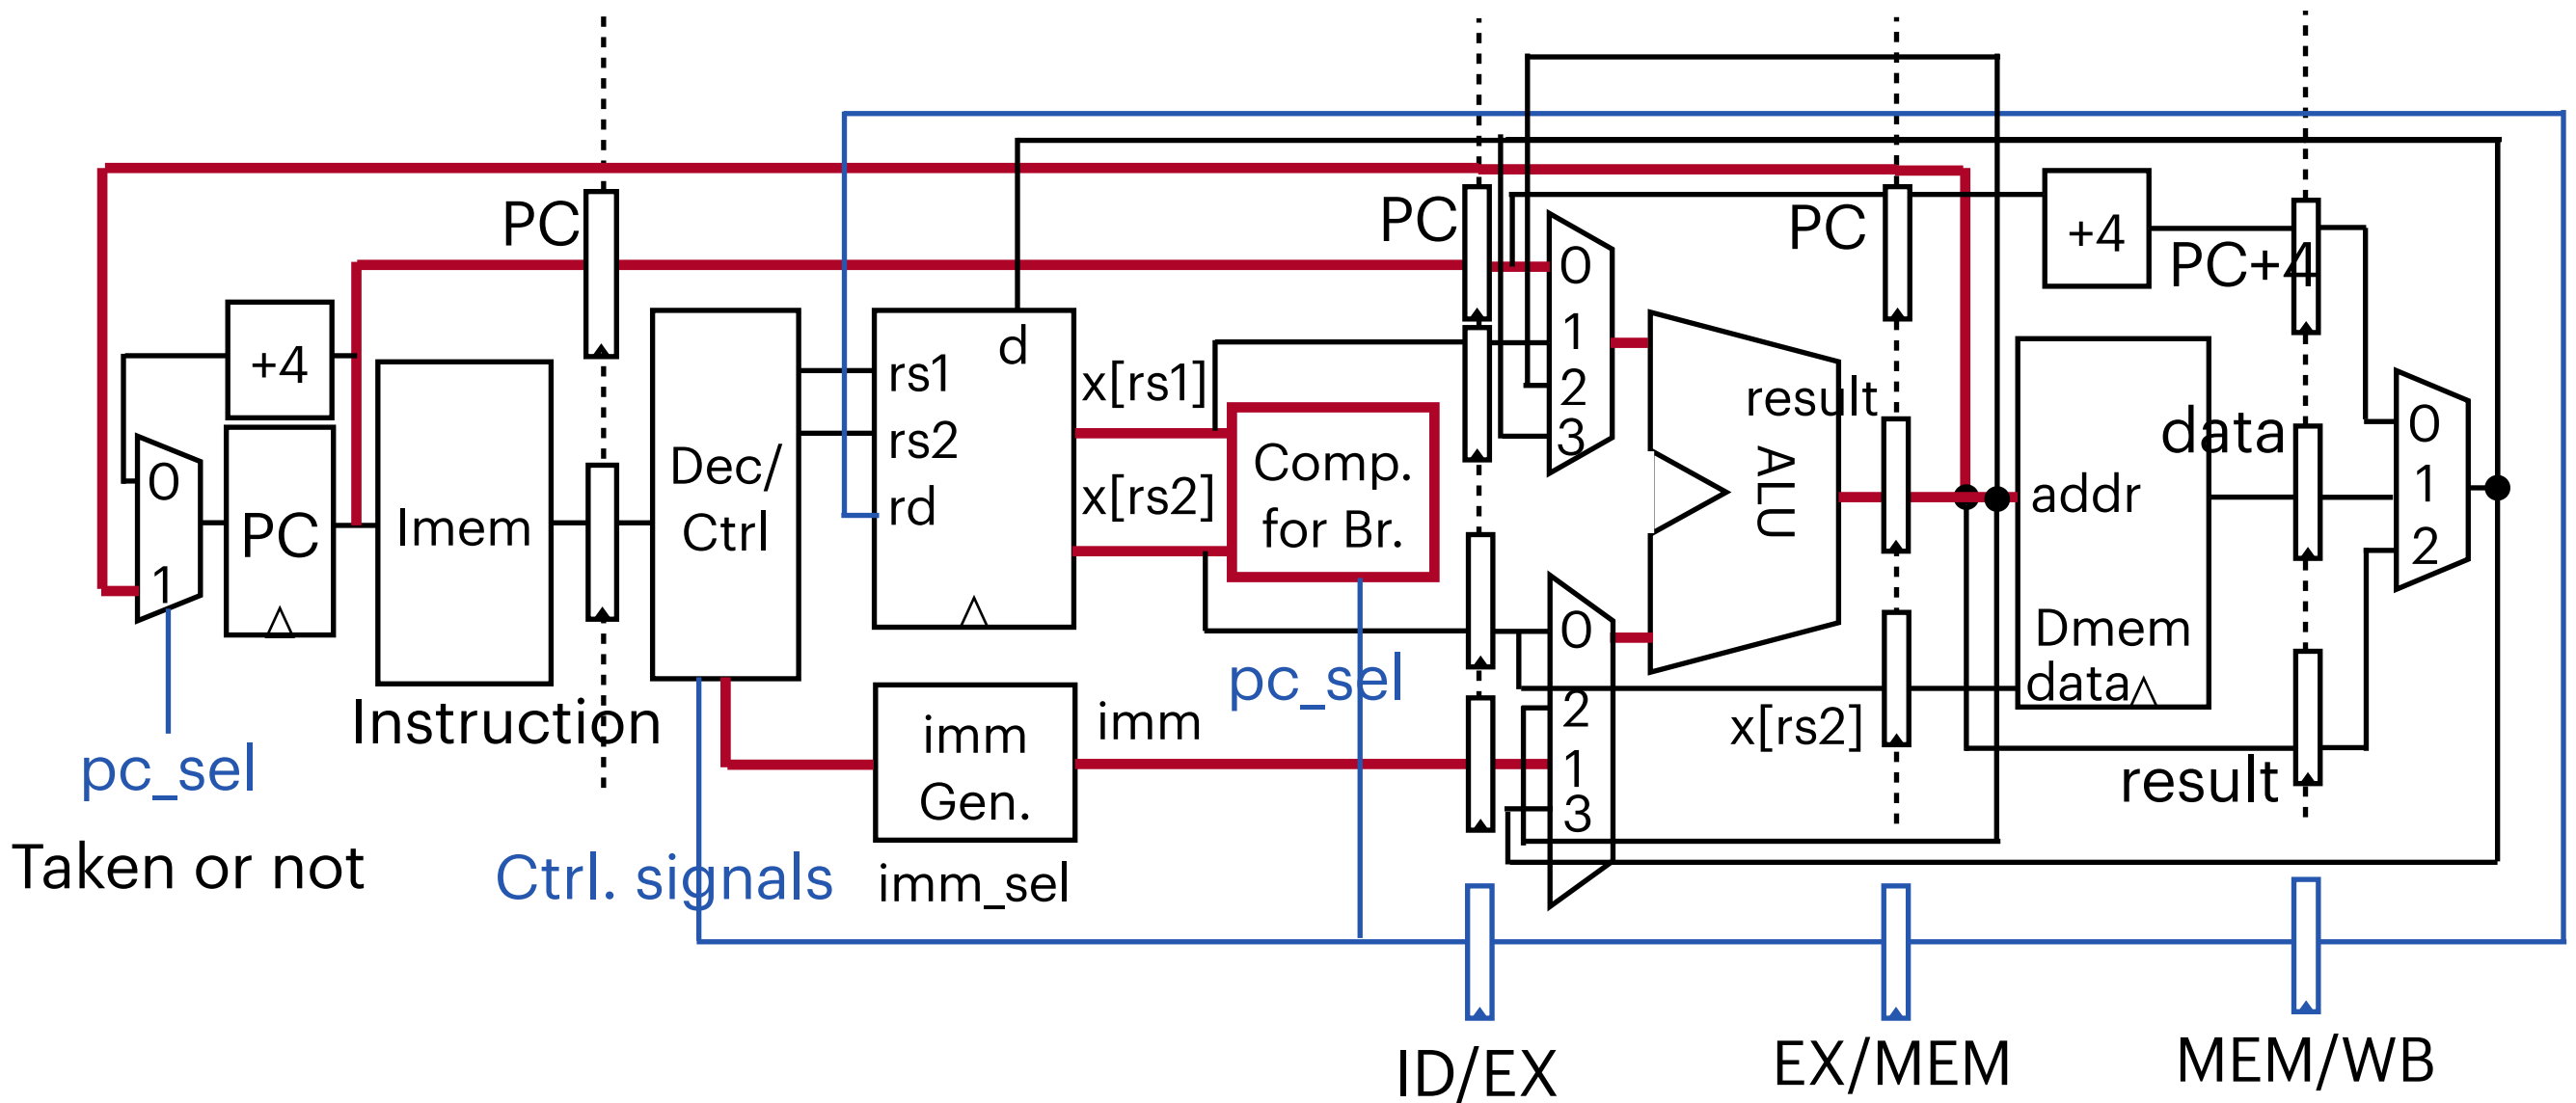
# Control Hazards

A hazard is a situation in which a planned instruction cannot execute in the “proper” clock cycle.

1. Structural hazard:
  - Hardware does not support access across multiple instructions in the same cycle
2. Data hazard:
  - Instructions have data dependency
  - Occurs when an instruction reads a register before a previous instruction has finished writing to that register
3. Control hazard:
  - Flow of execution depends on previous instruction

# Control Hazard Example: Branch

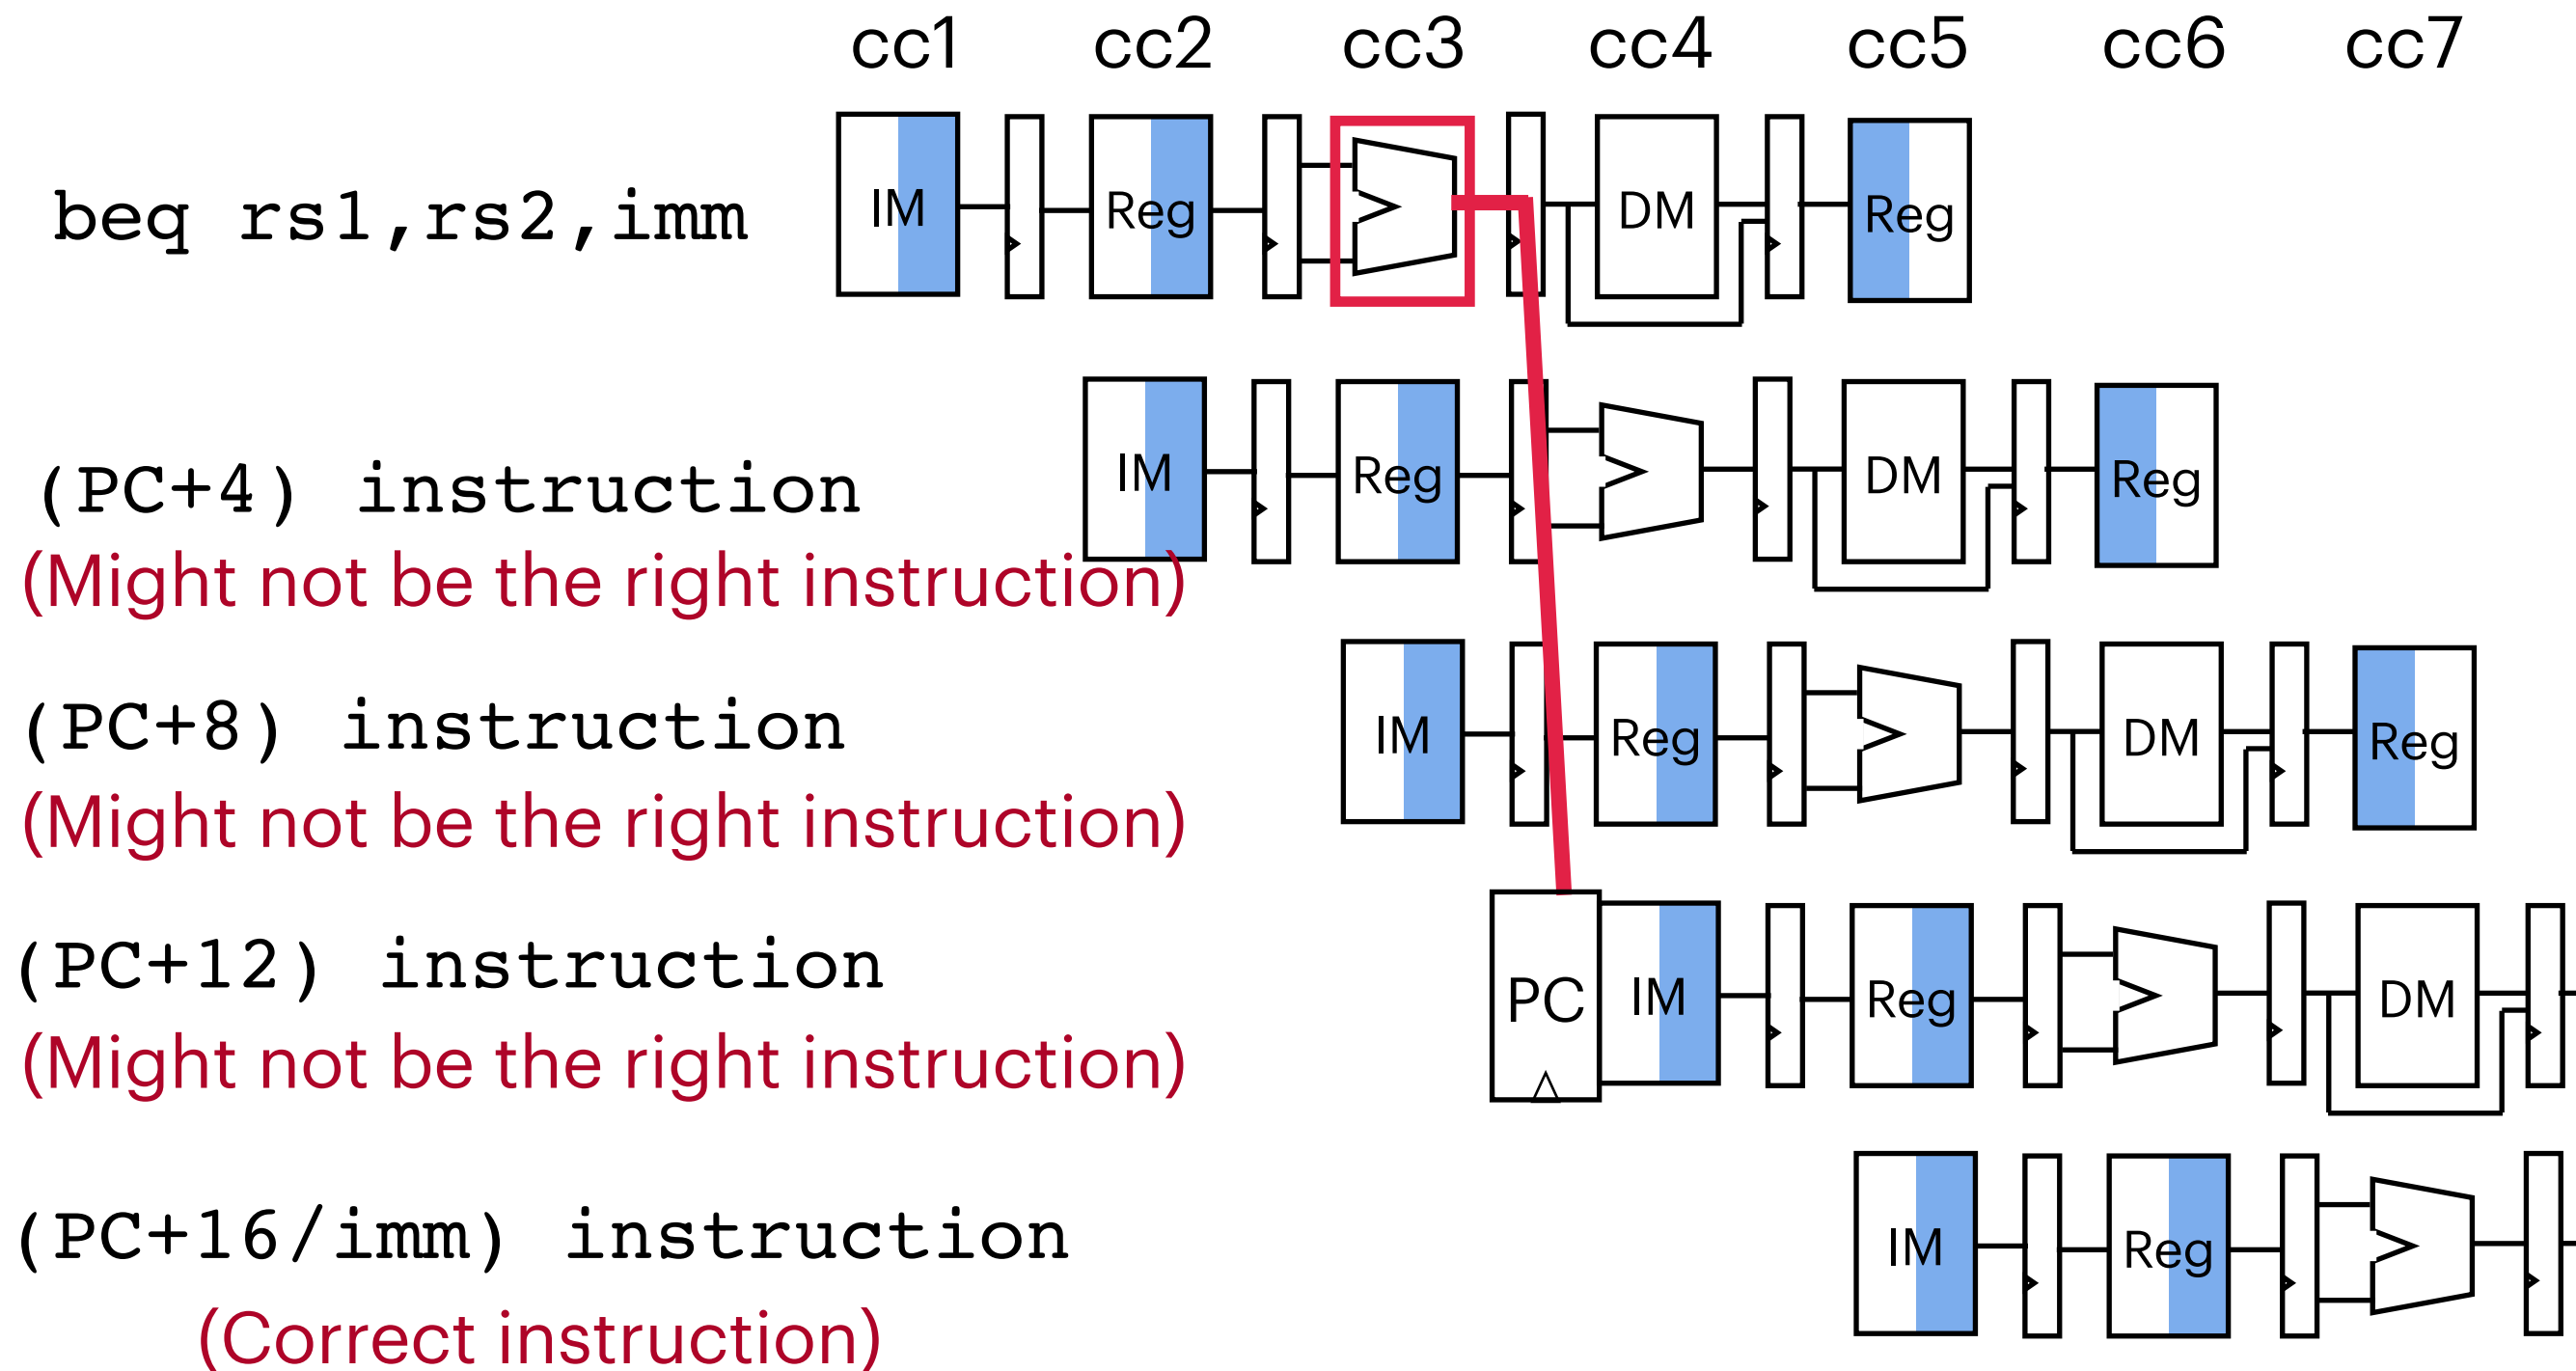
beq rs1,rs2,imm/label



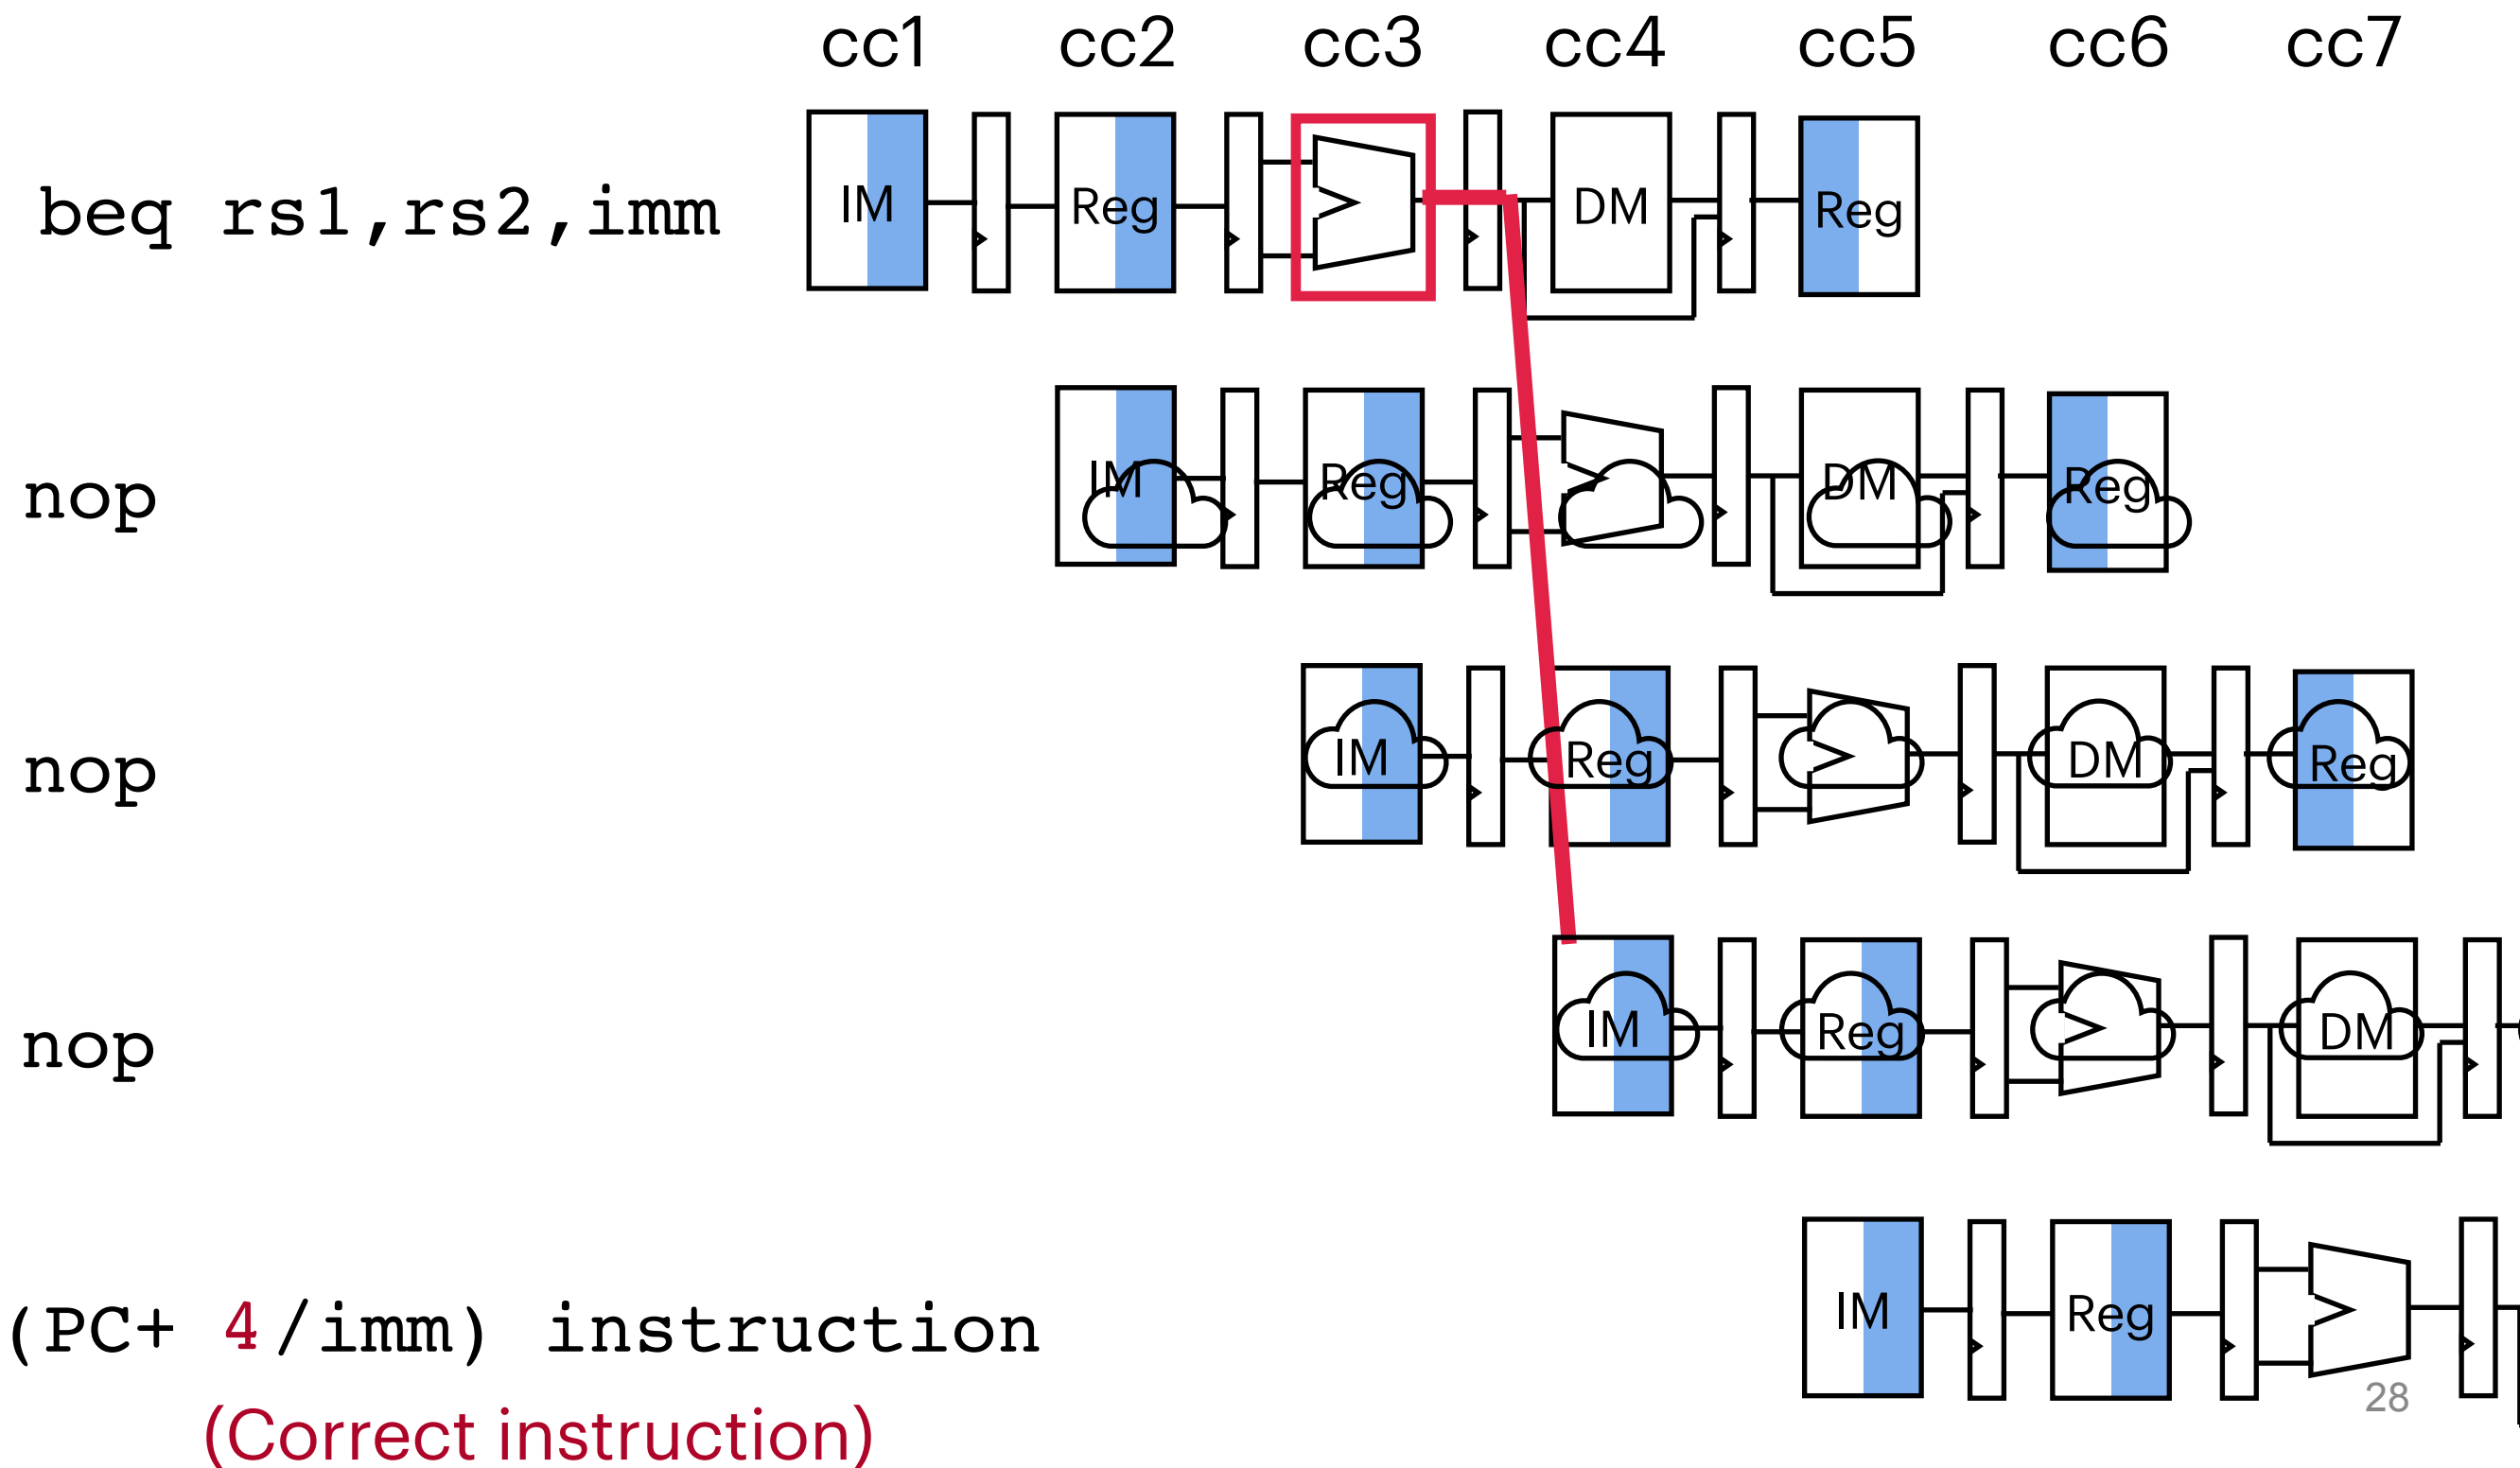
1. rs1/rs2 comparison

2.  $PC = PC + imm$

# Control Hazard Example: Branch

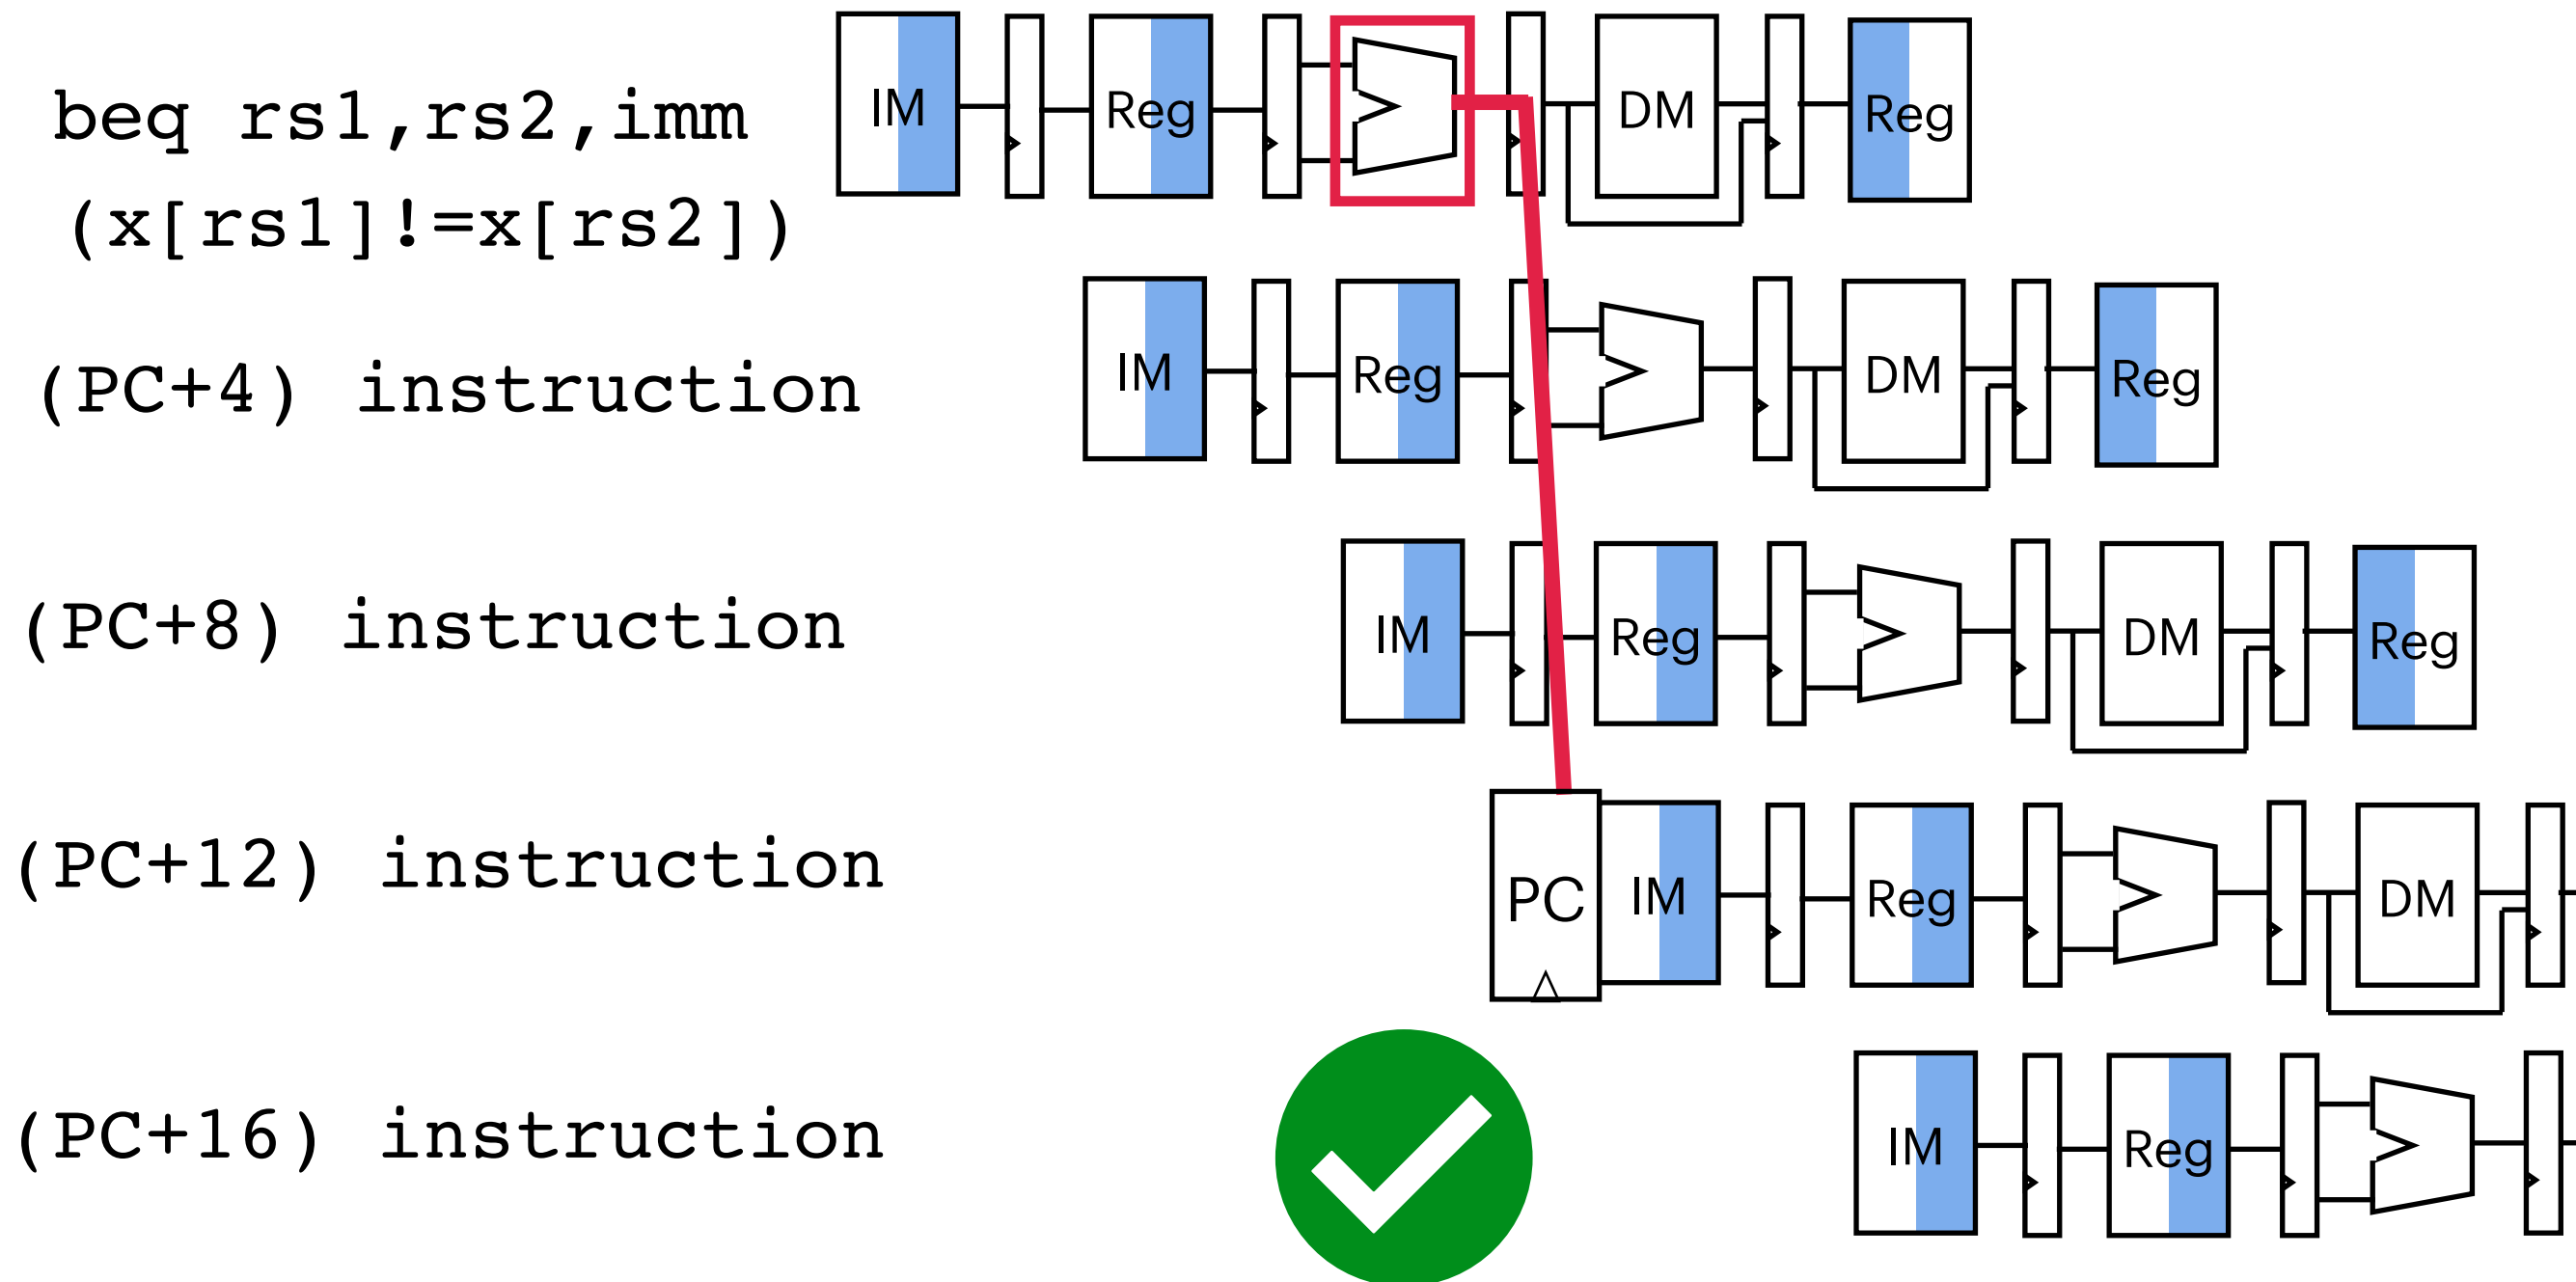


# Control Hazard Example: Branch



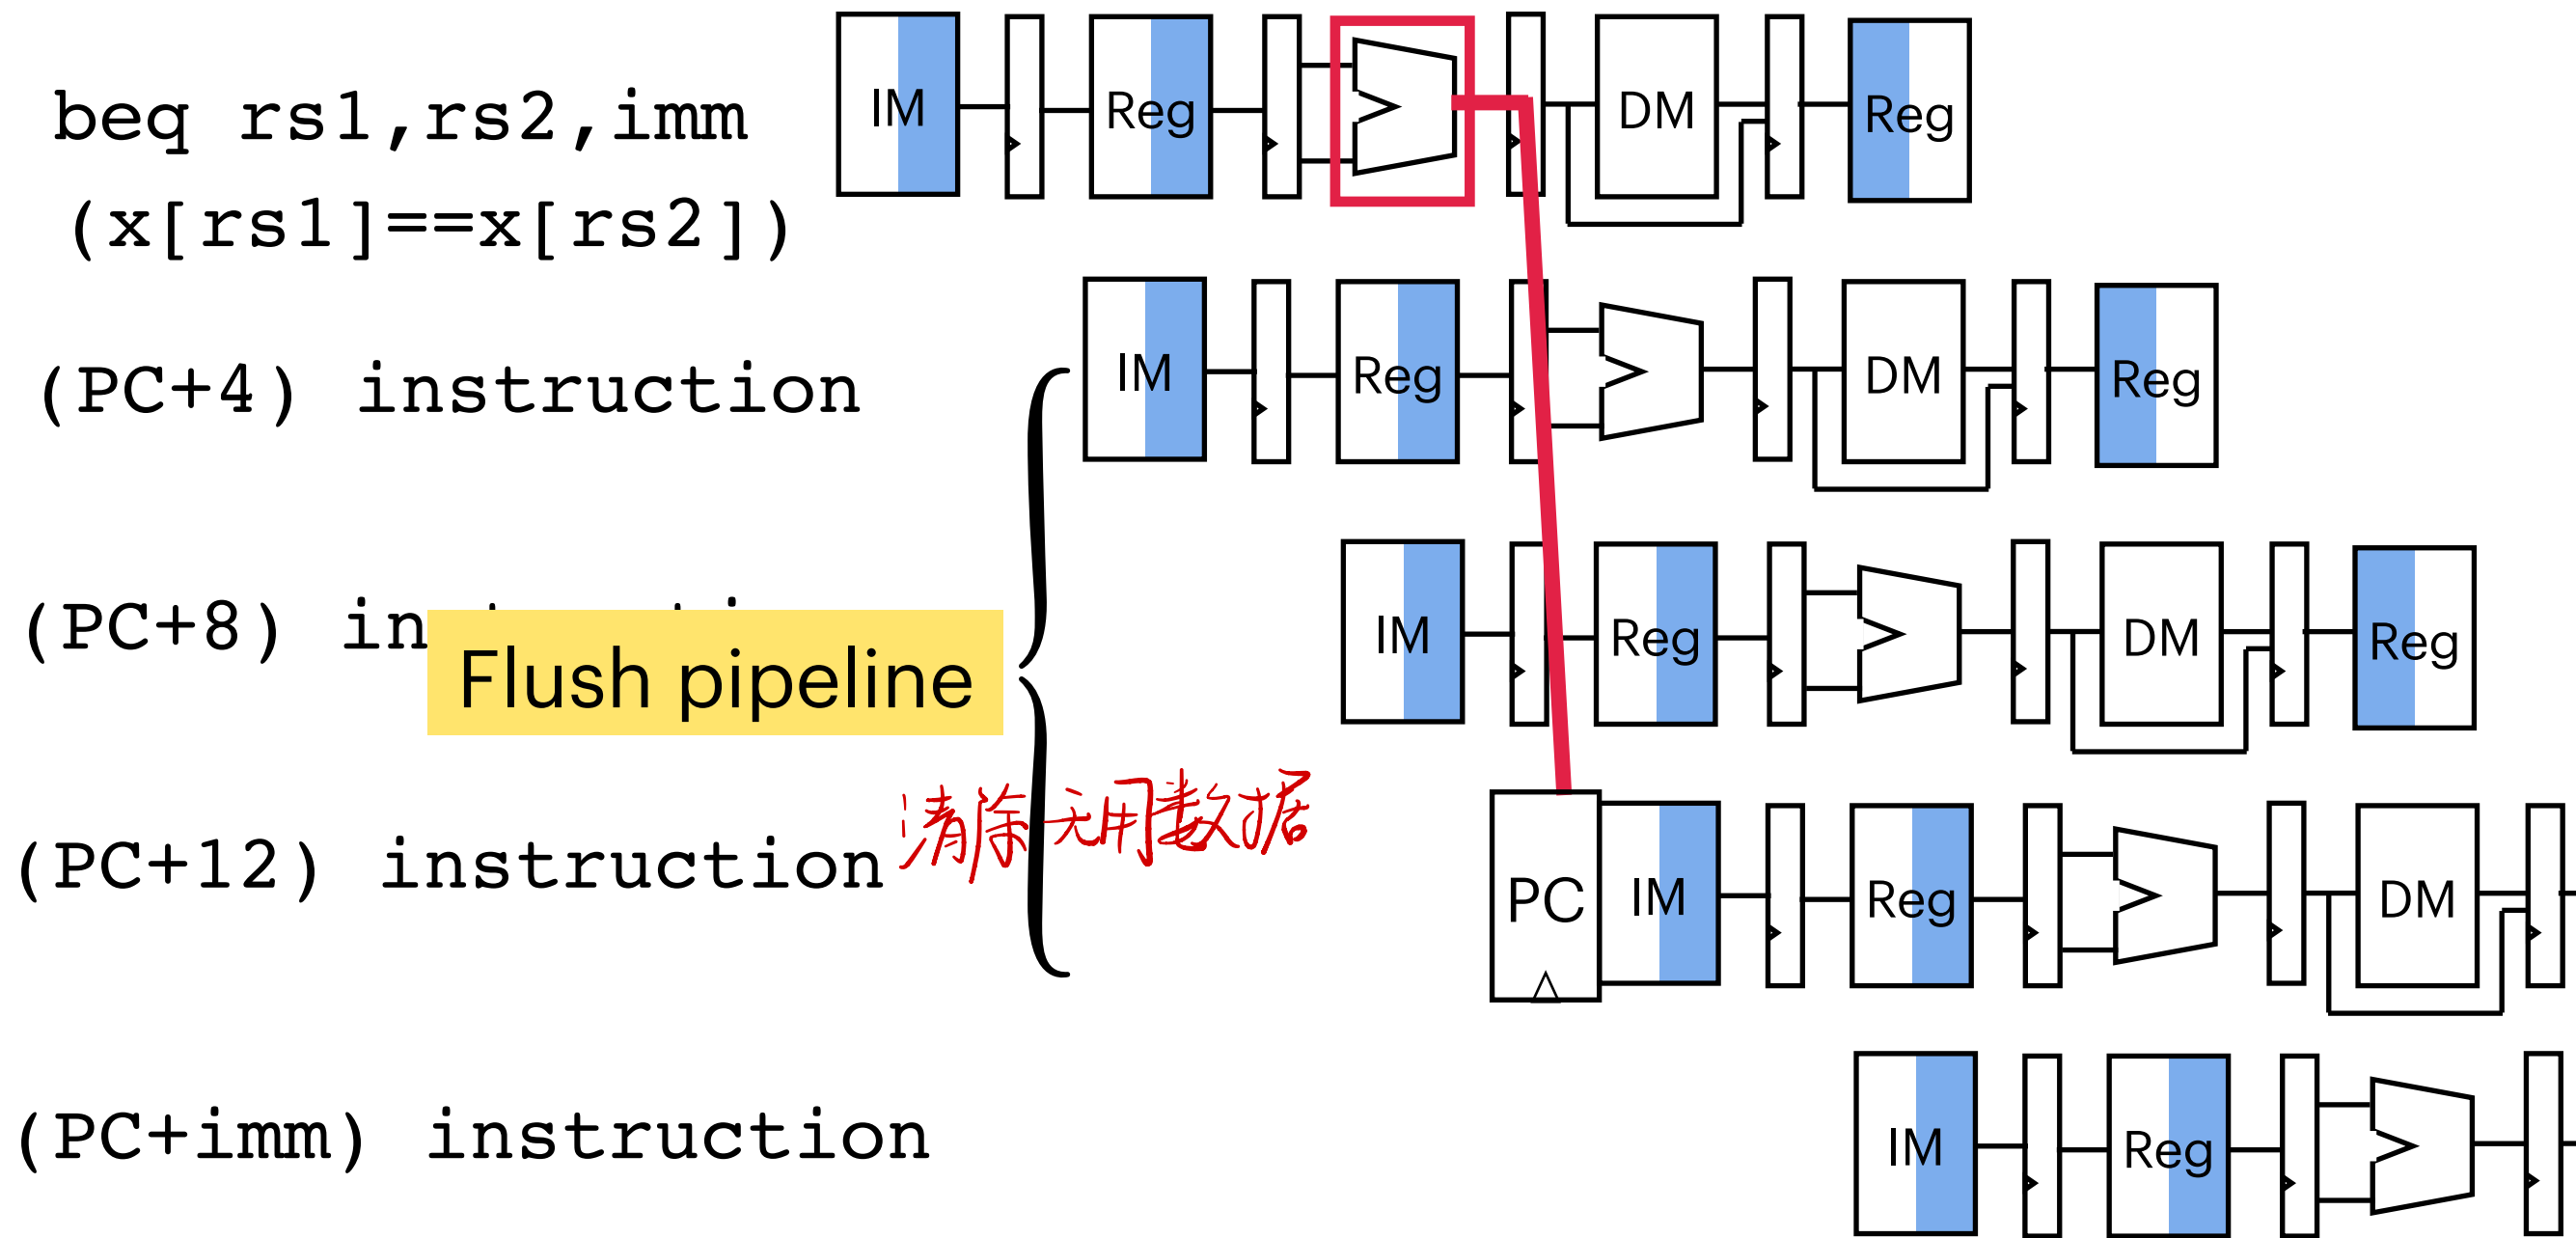
# Solution 1: Branch Prediction

- For example, a naïve predictor: always guess branch not taken.



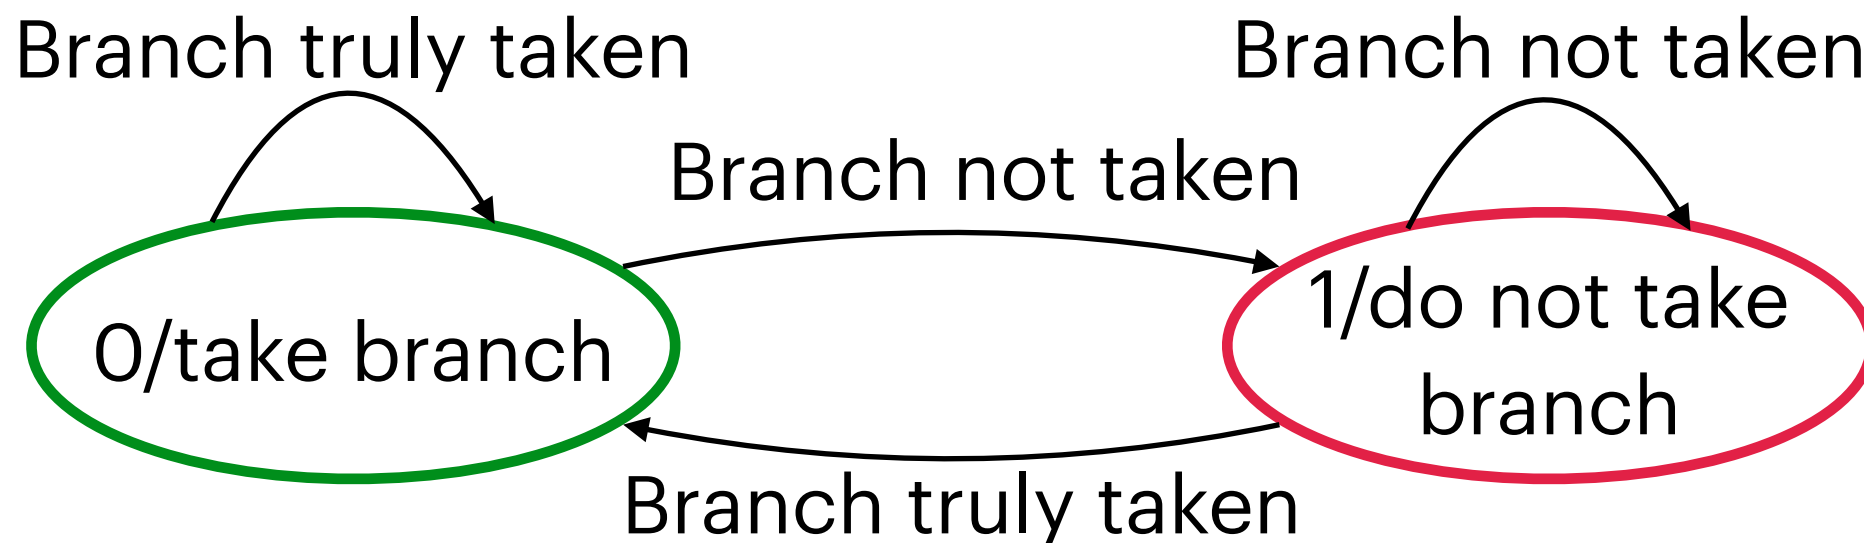
# Solution 1: Branch Prediction

- For example, a naïve predictor: always guess branch not taken.

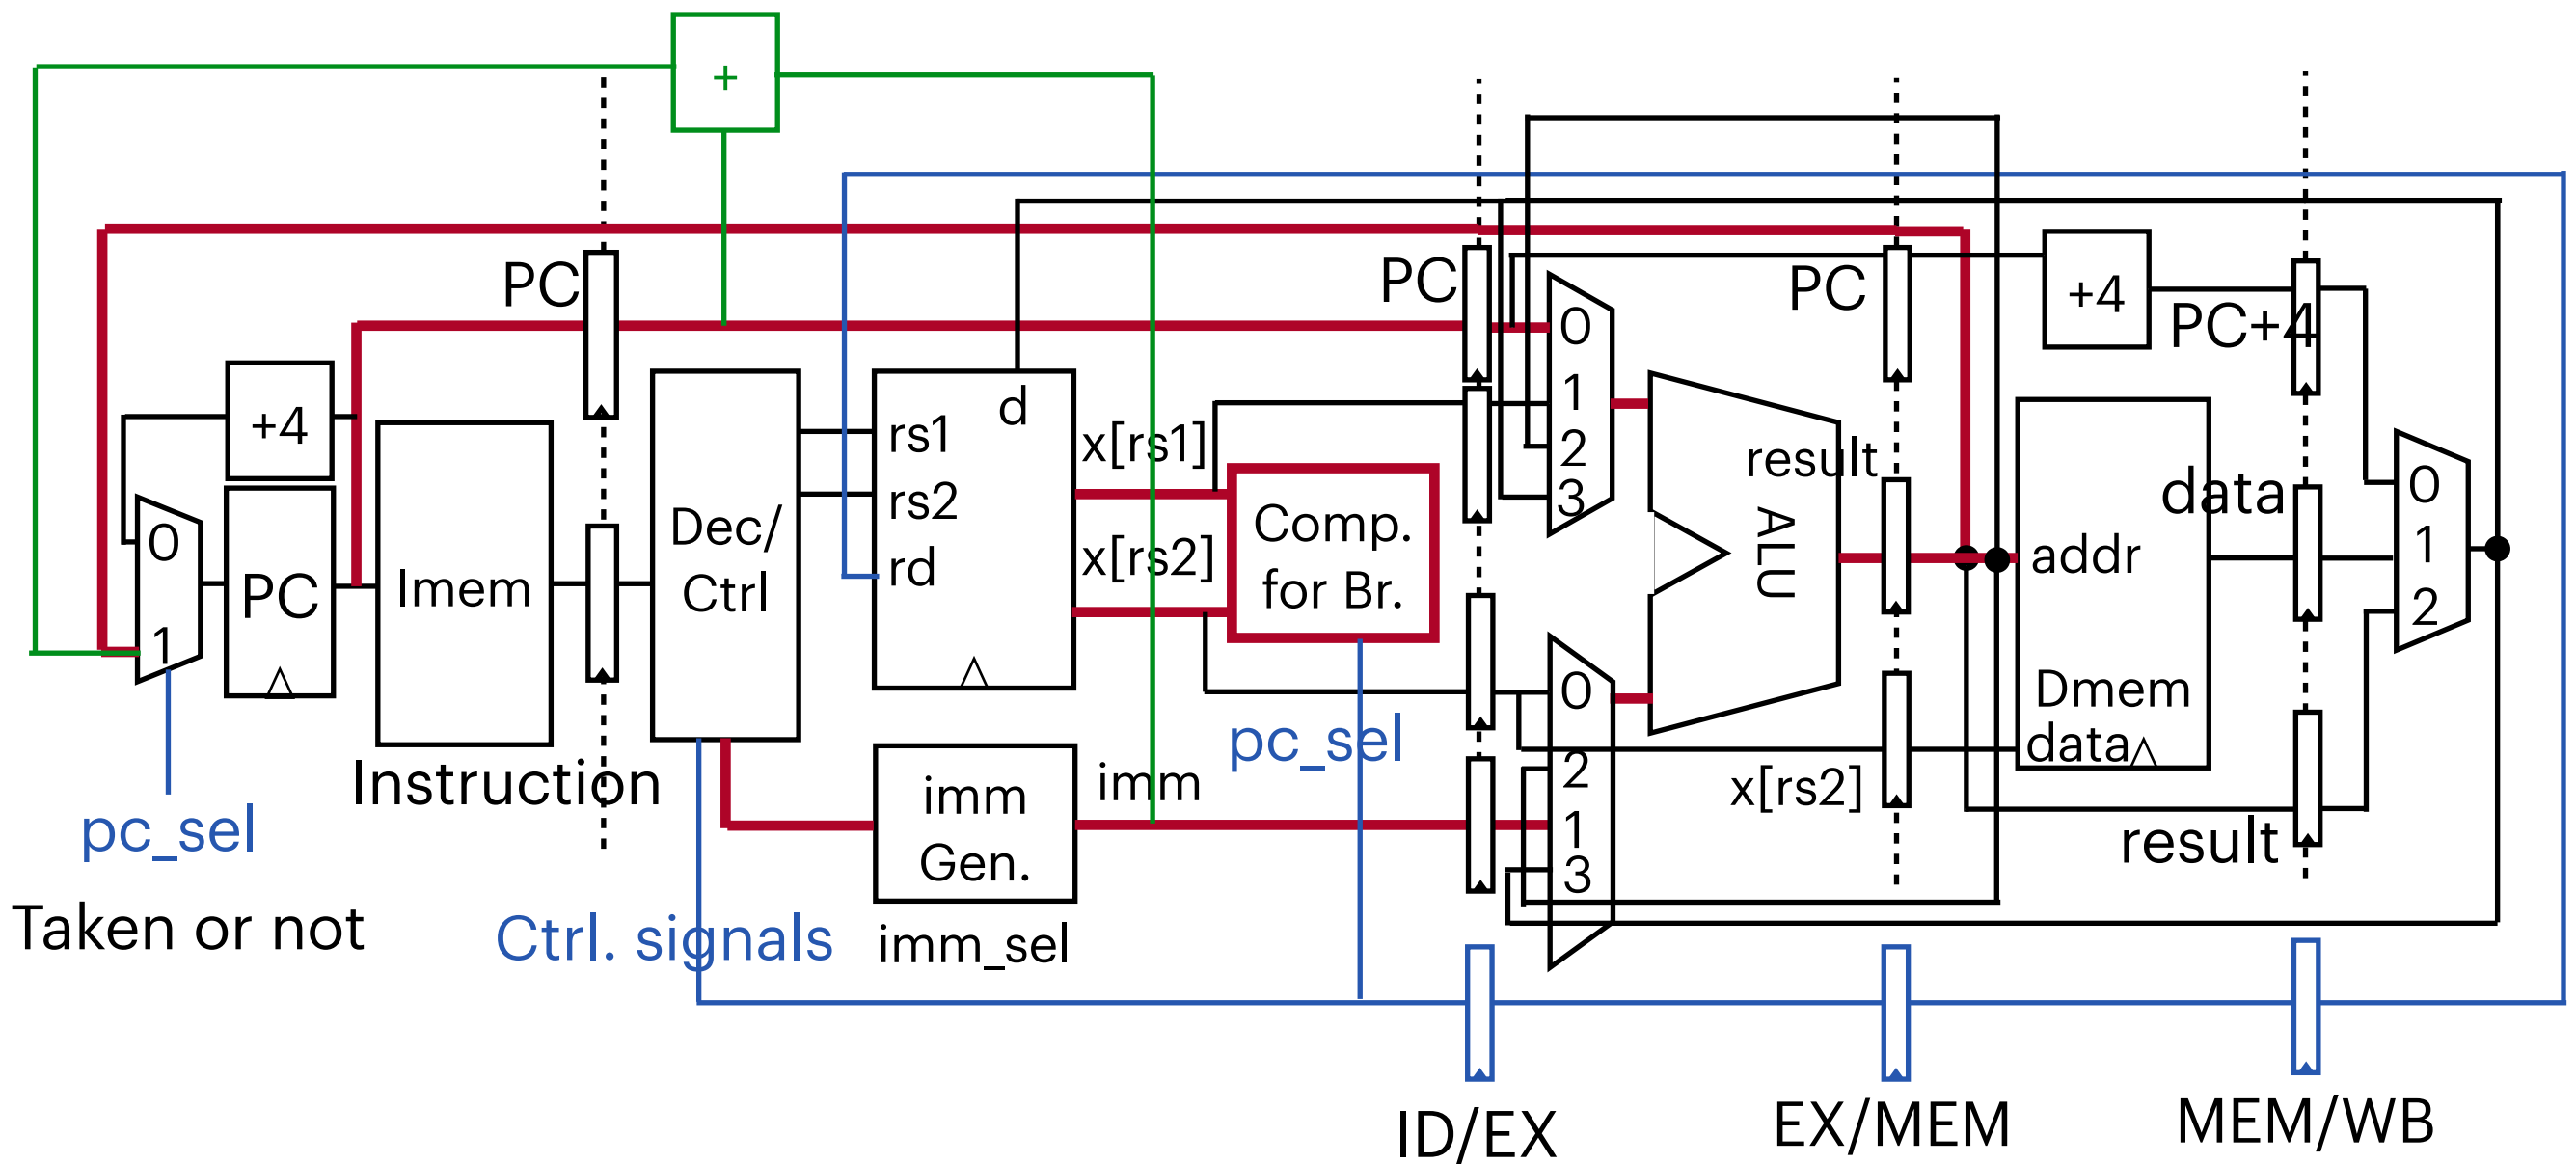


# Solution 1: Branch Prediction

- Other scheme: dynamic branch prediction using runtime info
  - predict based on history
  - Build an FSM to perform the prediction



# Solution 2: Modify the Hardware





# Pipeline Conclusion

- Pipelining increases throughput by overlapping execution of multiple instructions
- Hazards potentially limit performance
  - Maximizing performance requires programmer/compiler assistance/hardware modifications

# Increasing Processor Performance

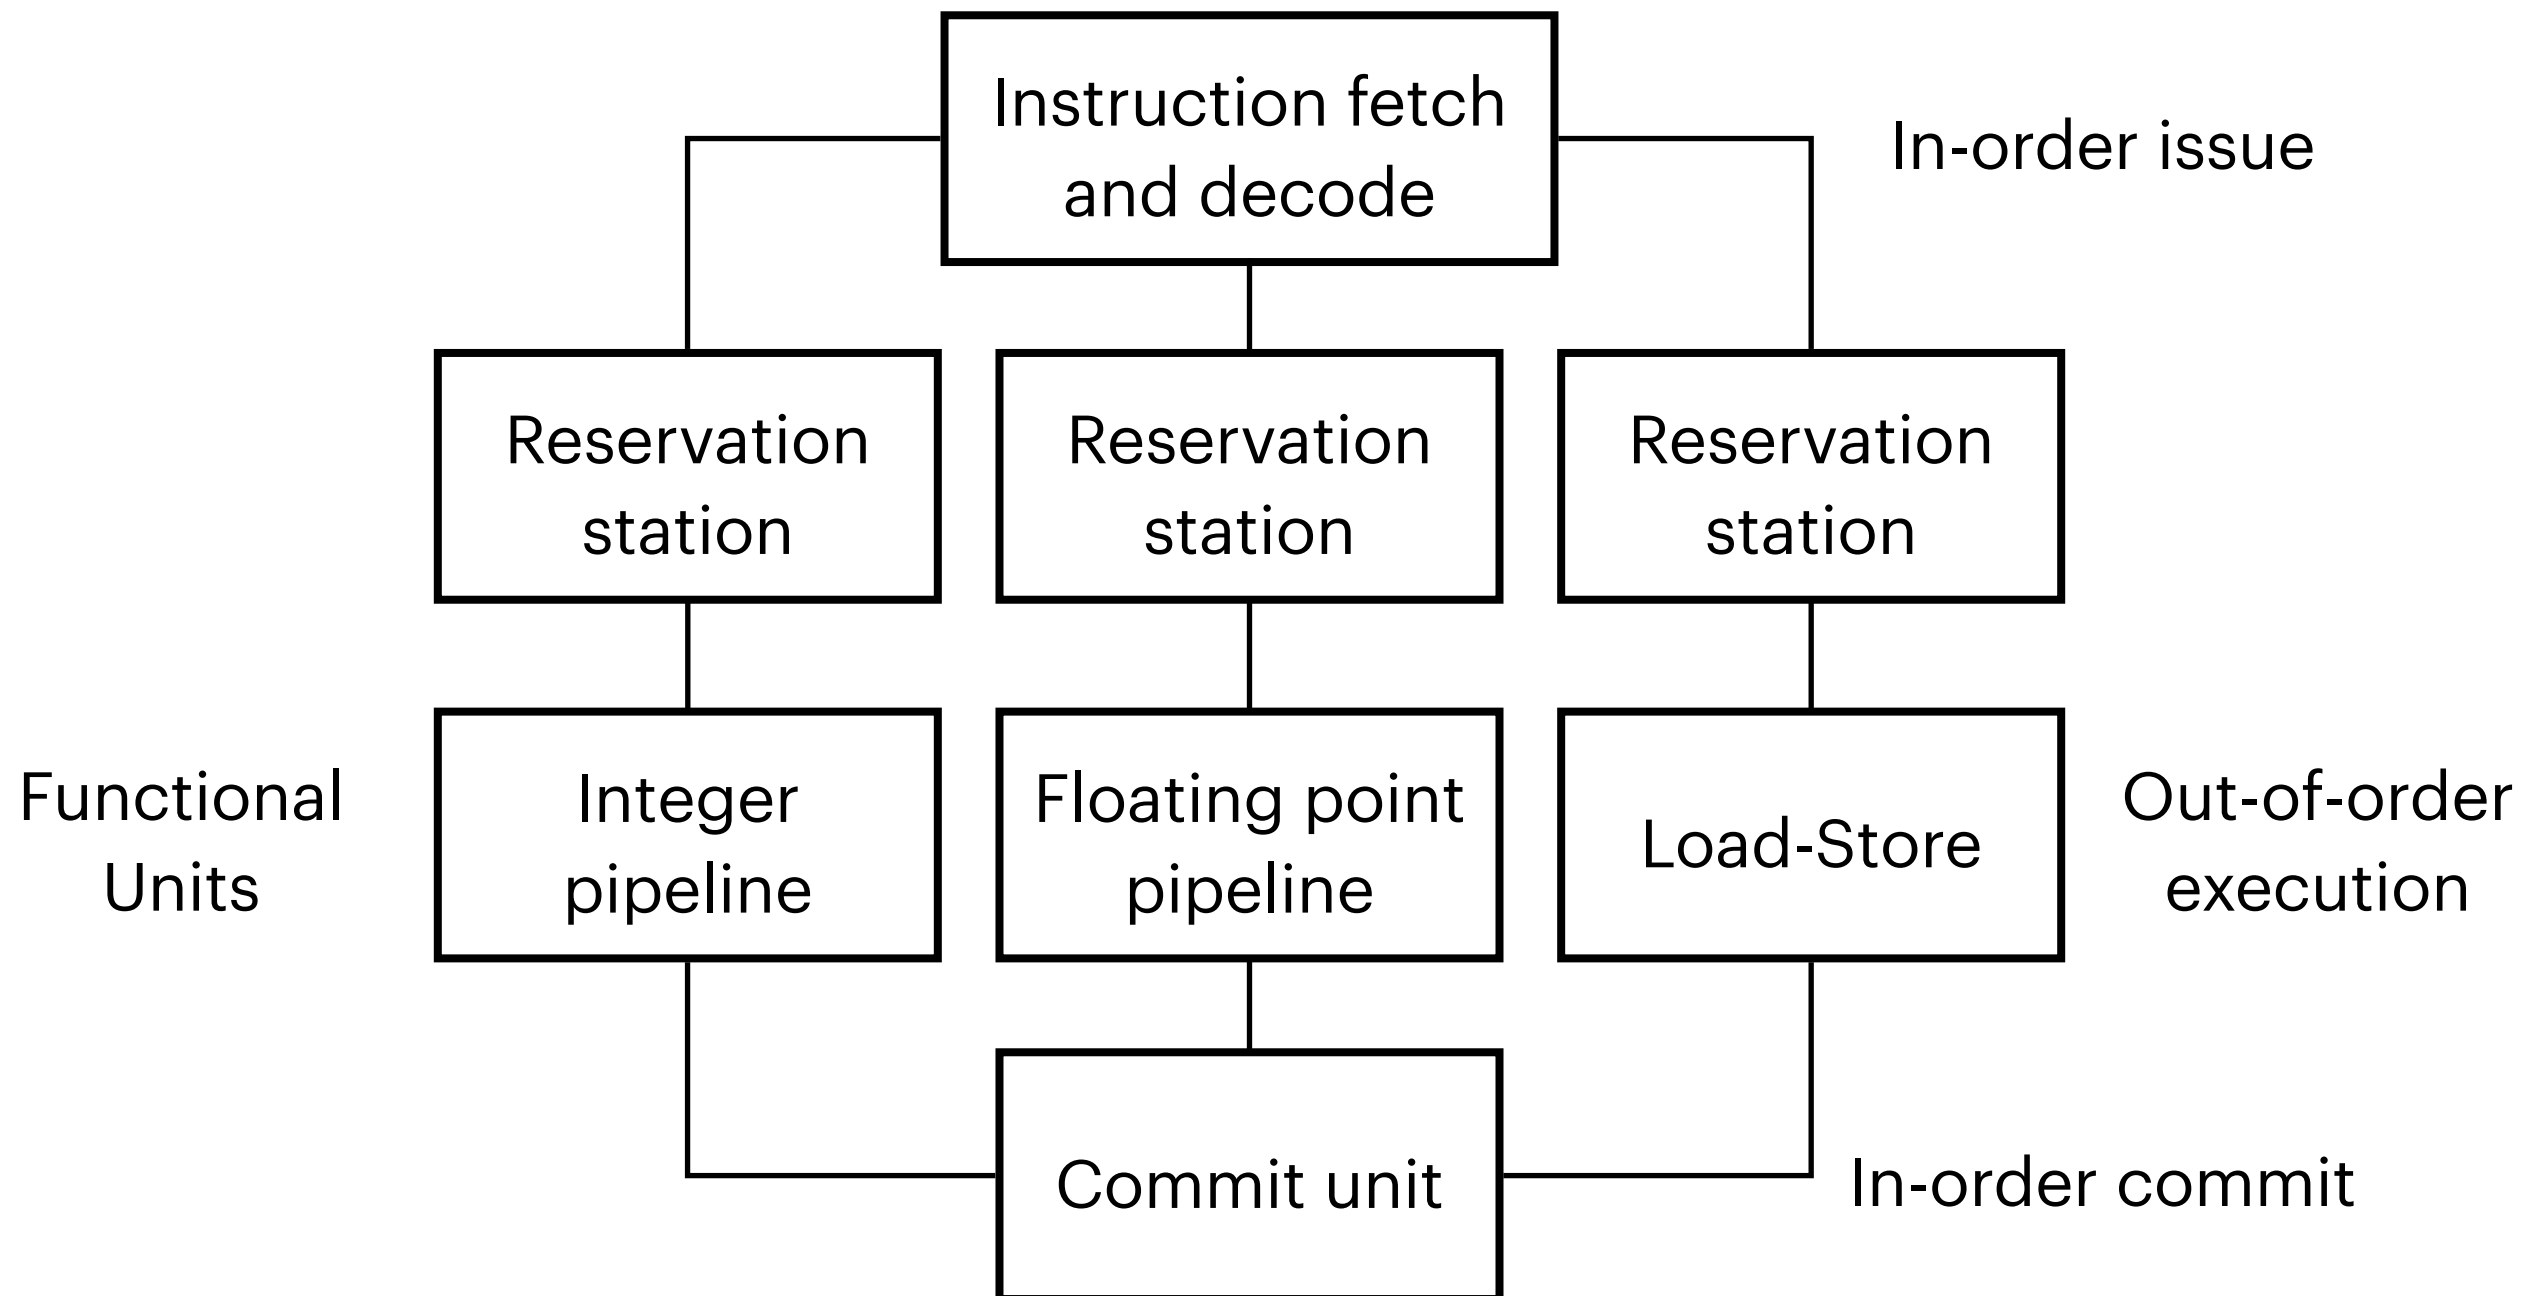
- Pipelining
  - “Overlap” instruction execution
  - Deeper pipeline: 5 => 10 => 15 stages
    - Less work per stage → shorter clock cycle
    - But more potential for hazards
    - Multi-issue processor
- CPI measurement: benchmark to obtain time/program

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Cycles}}{\text{Instruction}} \cdot \frac{\text{Time}}{\text{Cycle}}$$

# Greater Instruction-Level Parallelism (ILP)

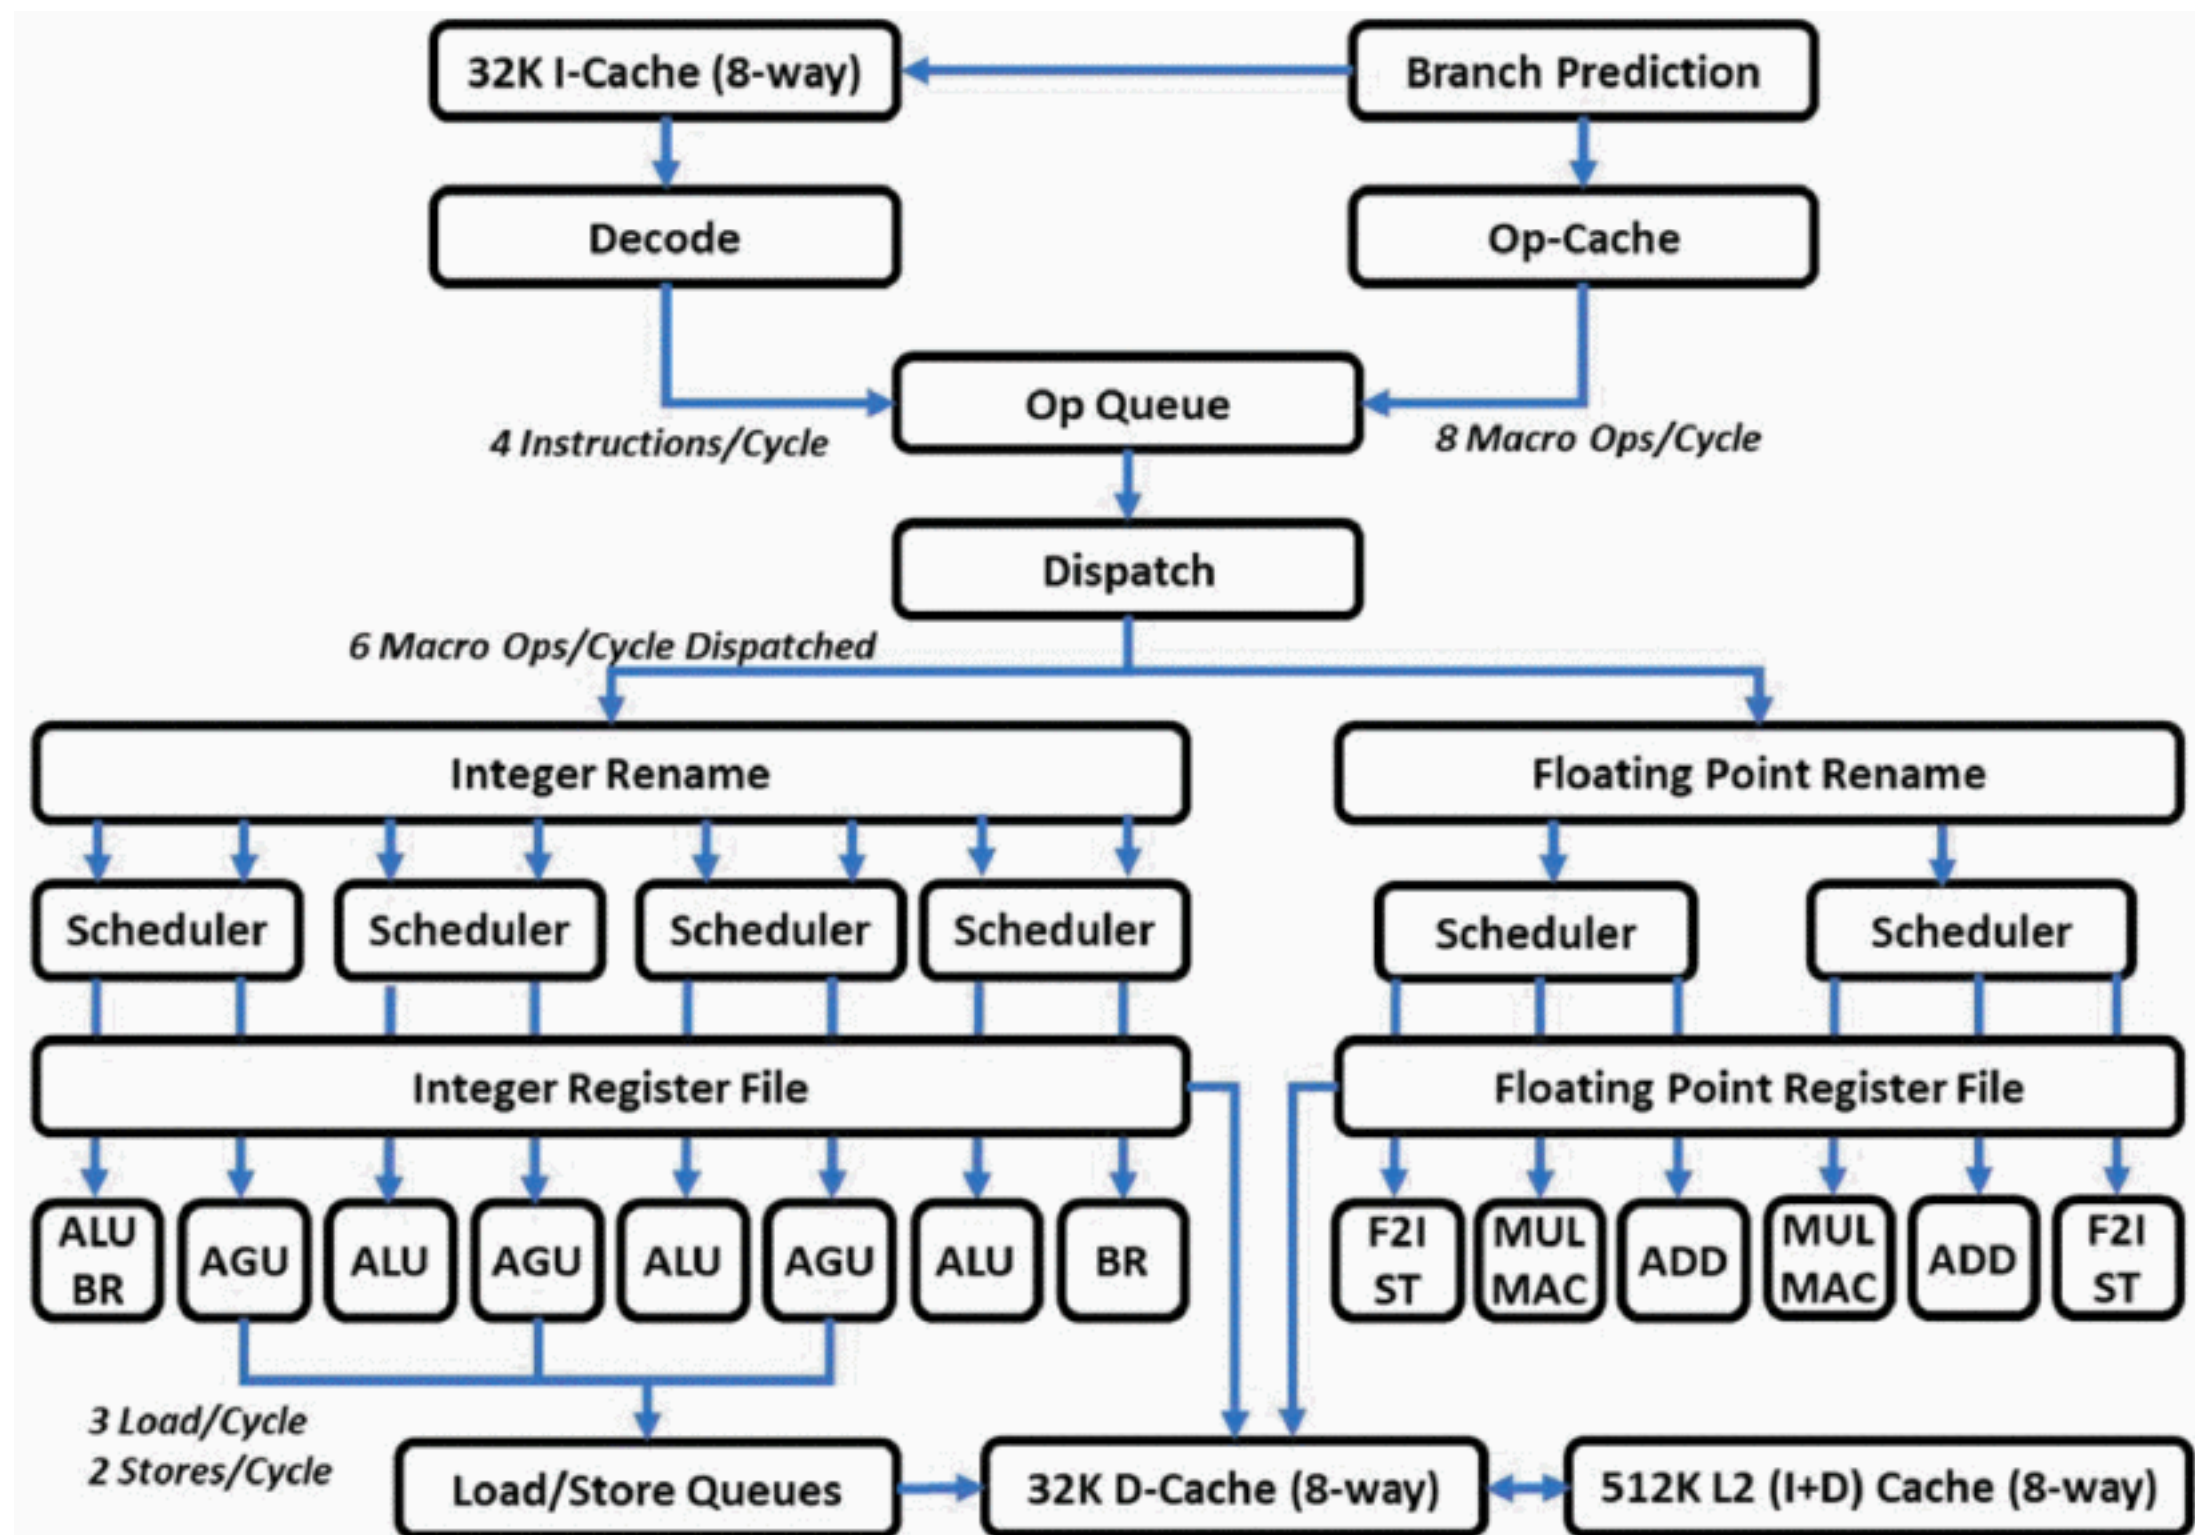
- Multiple issue
  - Replicate pipeline stages => multiple pipelines
  - Start multiple instructions per clock cycle
  - $CPI < 1$ , also use Instructions Per Cycle (IPC)
  - E.g., 4GHz 4-way multiple-issue
    - 16 BIPS, peak  $CPI = 0.25$ , peak  $IPC = 4$
  - But dependencies reduce this in practice
- Superscalar 超标量
- “Out-of-Order” execution
  - Reorder instructions dynamically in hardware to reduce impact of hazards
- More in CAlI & EE219 (about AI chips and HW/SW co-design)

# Examples



Multi-issue/superscalar is not multicore

# Examples: Superscalar (also Multi-issue)



AMD Zen 3, 7 nm process, a single core

# Examples: GPU



NVIDIA A100

Super“vector”  
or

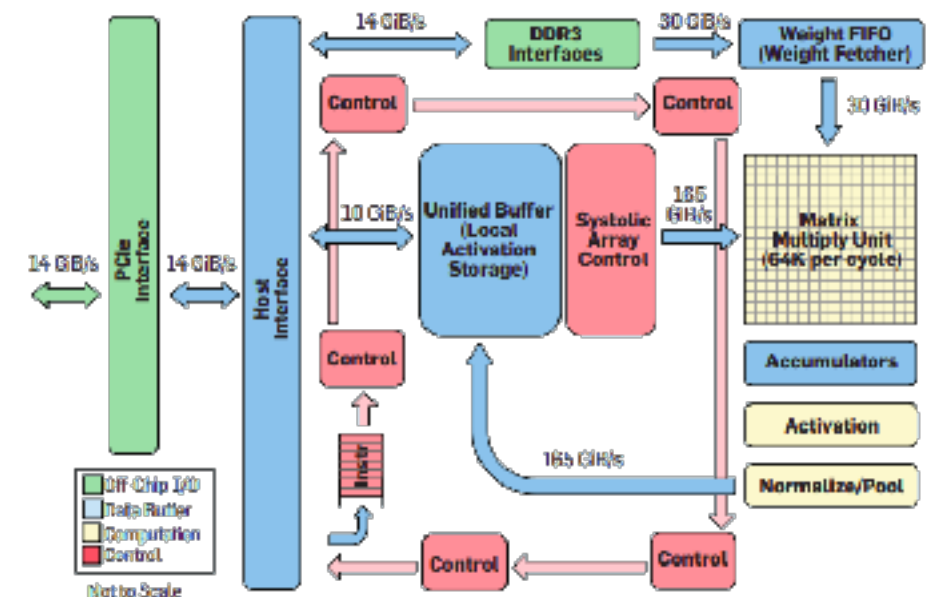
SIMD/MIMD

More on future lectures  
(cover by Prof. Wang)



# Static Multiple Issue

- aka.: Very Long Instruction Word (VLIW)
- Compiler bundles instructions together
- Compiler takes care of hazards
- Used in Google TPU (an AI chip)



ALU or branch	IF	ID	EX	MEM	WB			
Load or store	IF	ID	EX	MEM	WB			
ALU or branch		IF	ID	EX	MEM	WB		
Load or store		IF	ID	EX	MEM	WB		
ALU or branch			IF	ID	EX	MEM	WB	
Load or store			IF	ID	EX	MEM	WB	
ALU or branch				IF	ID	EX	MEM	WB
Load or store				IF	ID	EX	MEM	WB

# CPU Design & Manufacturing

- Specifications defined.
  - Target performance, TPD (power consumption), cost, etc.
- Microarchitecture design
  - Decide ISA, multi-issue/VLIW/superscalar/out-of-order execution, etc.
- Hardware design using hardware description language (HDL)

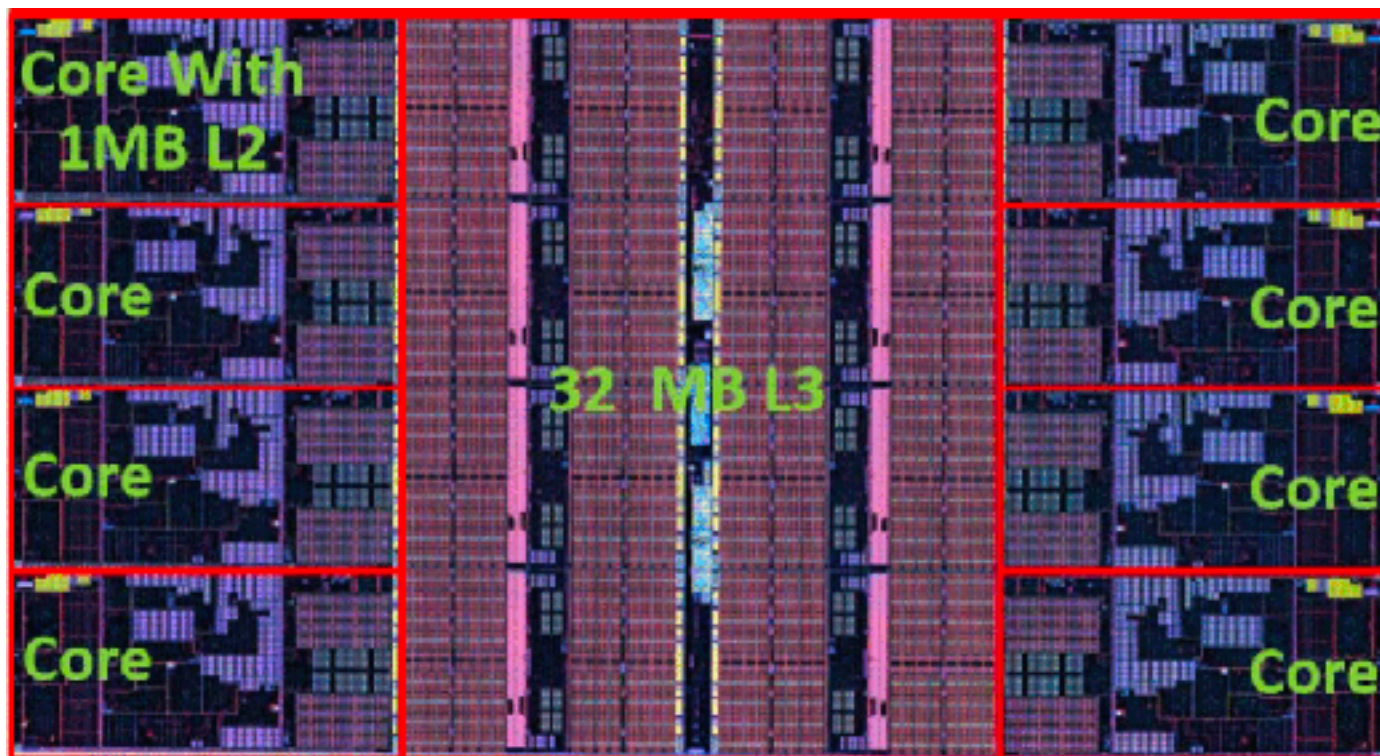
```
module alu(opA, opB, aluop, result, zero);  
  parameter width=32;  
  input [1:0] aluop; input [width-1:0] opA, opB;  
  output reg [width-1:0] result; output reg zero;  
  always @(*) zero = (result == 0);  
  always @(opA, opB, aluop) begin  
    case (aluop) 0: result = opA + opB; 1: result = opA - opB;  
    2: result = opA & opB; 3: result = opA | opB;  
    default: result = 0;  
  endcase  
end  
endmodule
```

An RTL-level description (verilog) of an ALU.

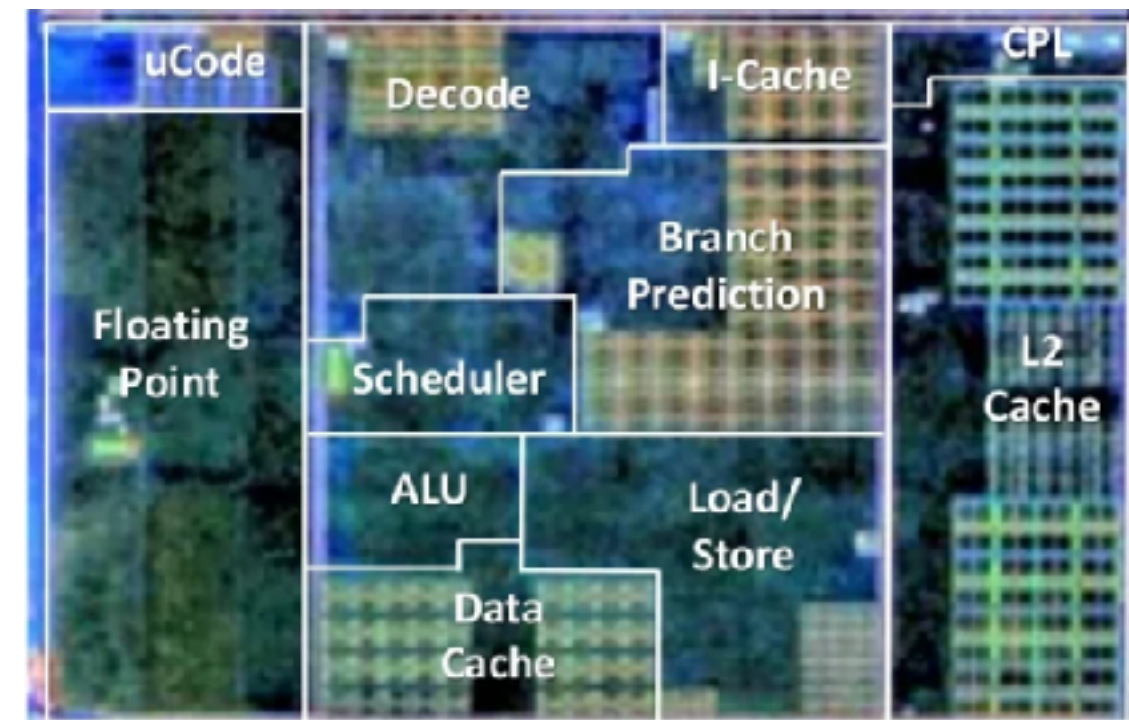


# CPU Design & Manufacturing

- Using EDA tools
  - To generate gate netlist automatically
  - To perform timing analysis
  - To simulate, verify, clock tree generation, etc.
  - To generate layout (GDSII)



AMD Zen 4 <https://ieeexplore.ieee.org/document/10067540>



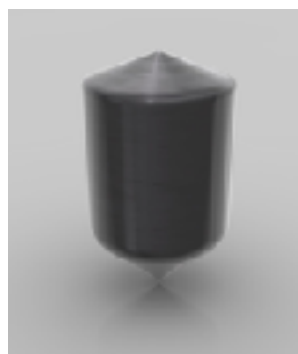
Zen 2: The AMD 7nm Energy-Efficient High-Performance x86-64 Microprocessor Core, 2020 ISSCC.

# CPU Design & Manufacturing

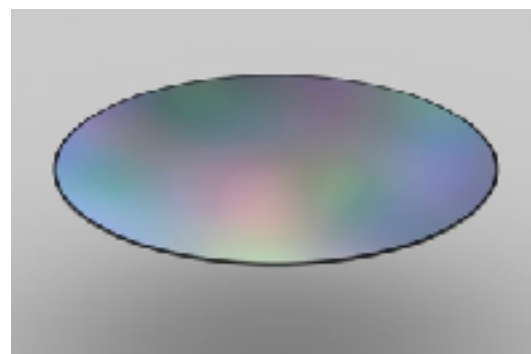
- Manufacturing
  - Different vendors have different strategies (fabless/IDM/foundry)



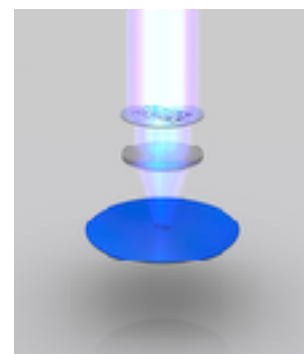
Sand



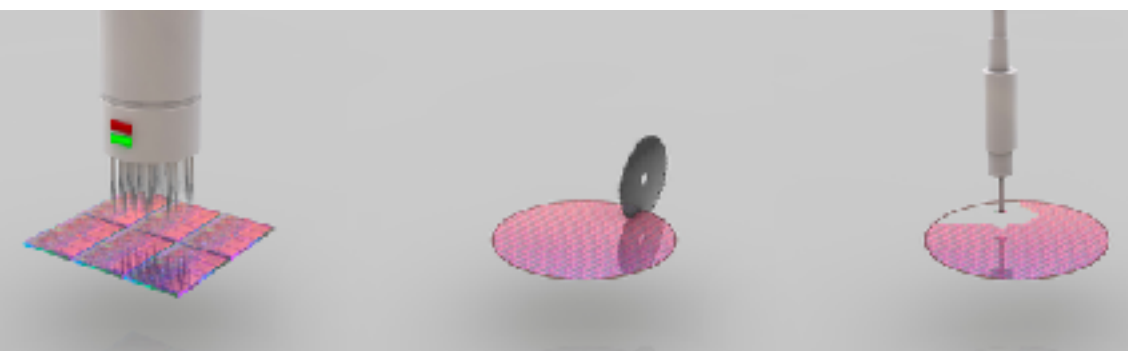
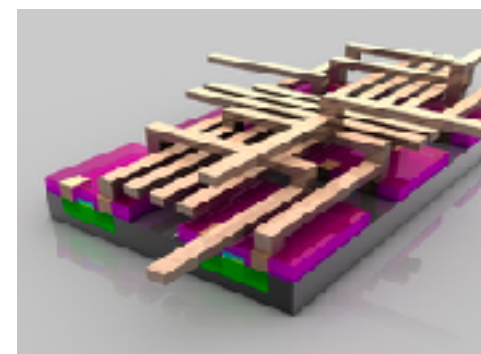
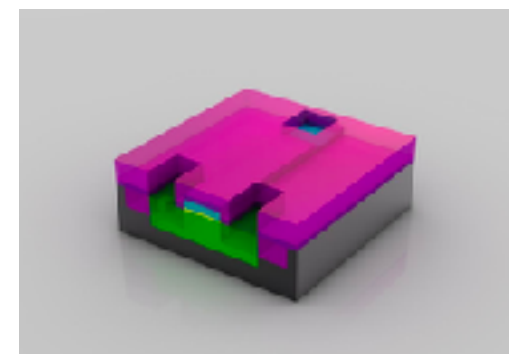
Ingot



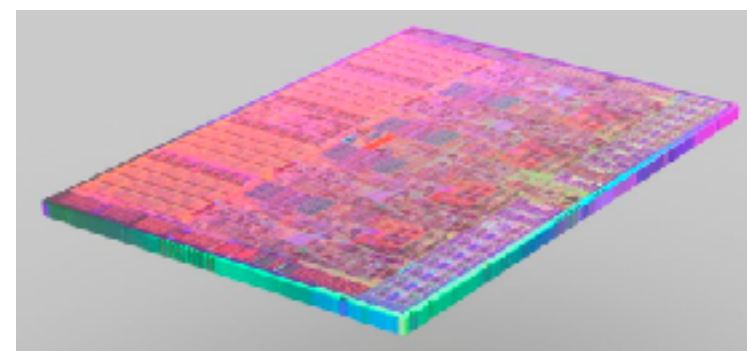
Wafer



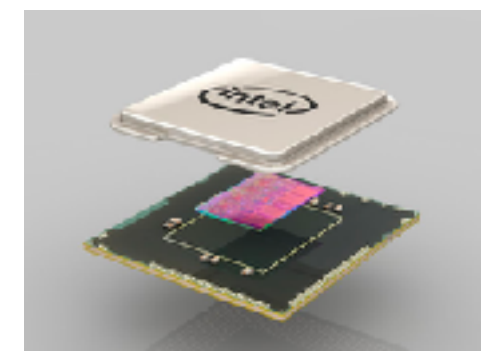
UV light/etching/doping/layering, etc.



Testing/wafer slicing

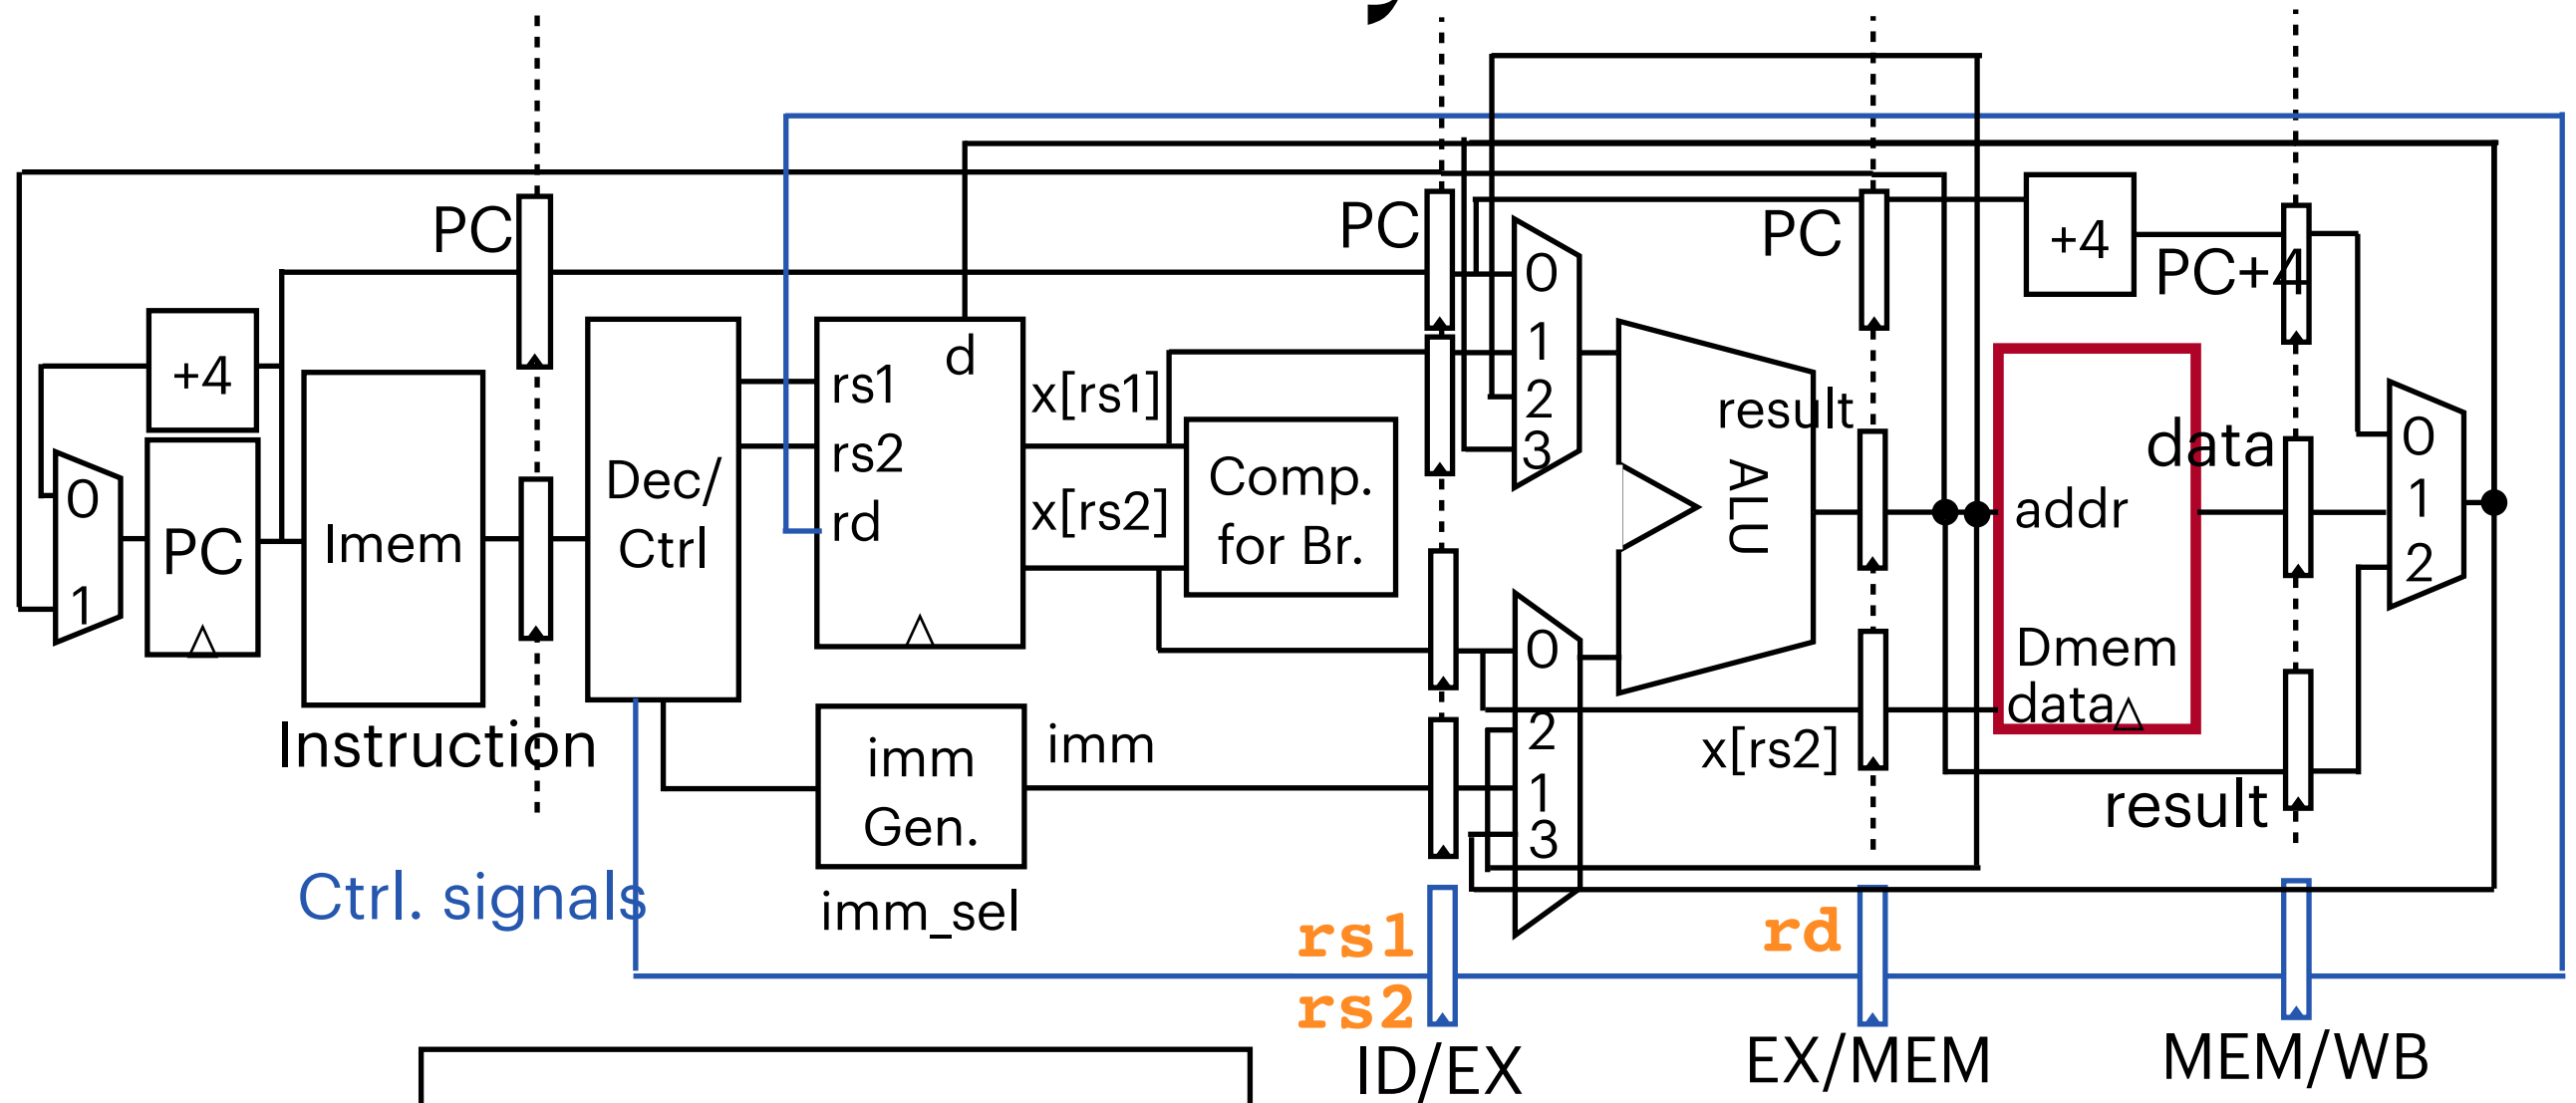


A single die/CPU

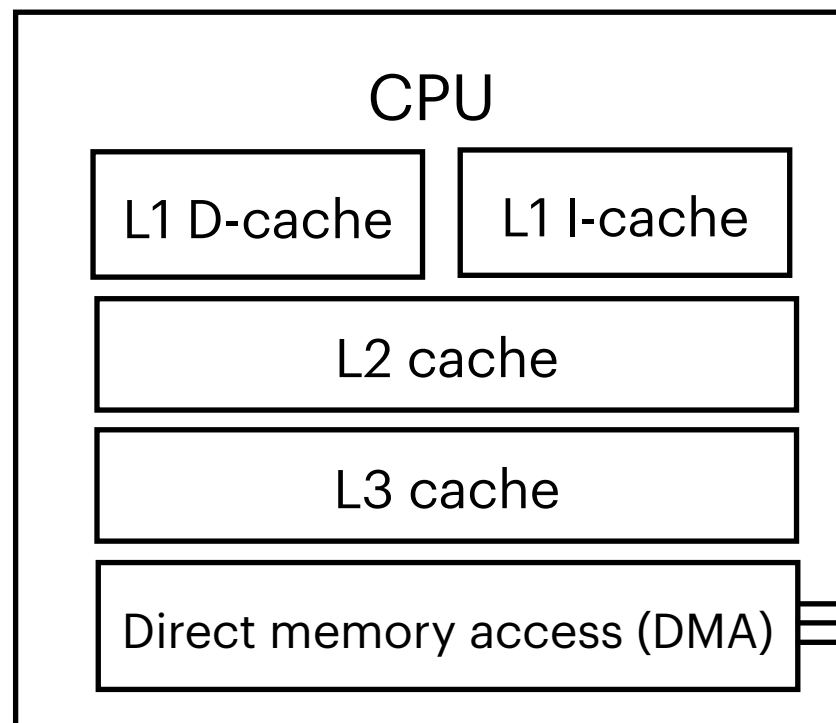


Packaging

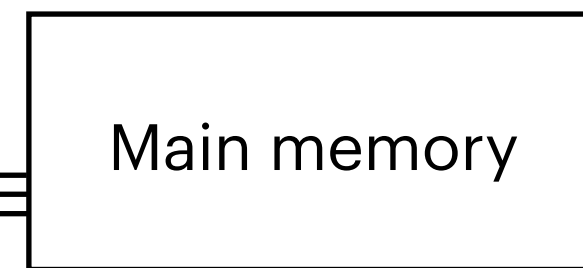
# CPU-Memory Interface



Reality



bus



More on future lectures  
(cover by Prof. Wang)



# Summer Interns Needed

- Possible topics
  - AI chip design (for convolutional neural network accelerator)
  - RISC-V CPU design (refer to <https://ysyx.oscc.cc/>)
- You will learn and practice
  - What you have learnt in the first half CAI course and so on
  - Hardware design using hardware description language (HDL) such as verilog HDL or chisel HDL

