CS 110
Computer Architecture

Caches Part III

Instructors:

Chundong Wang & Siting Liu
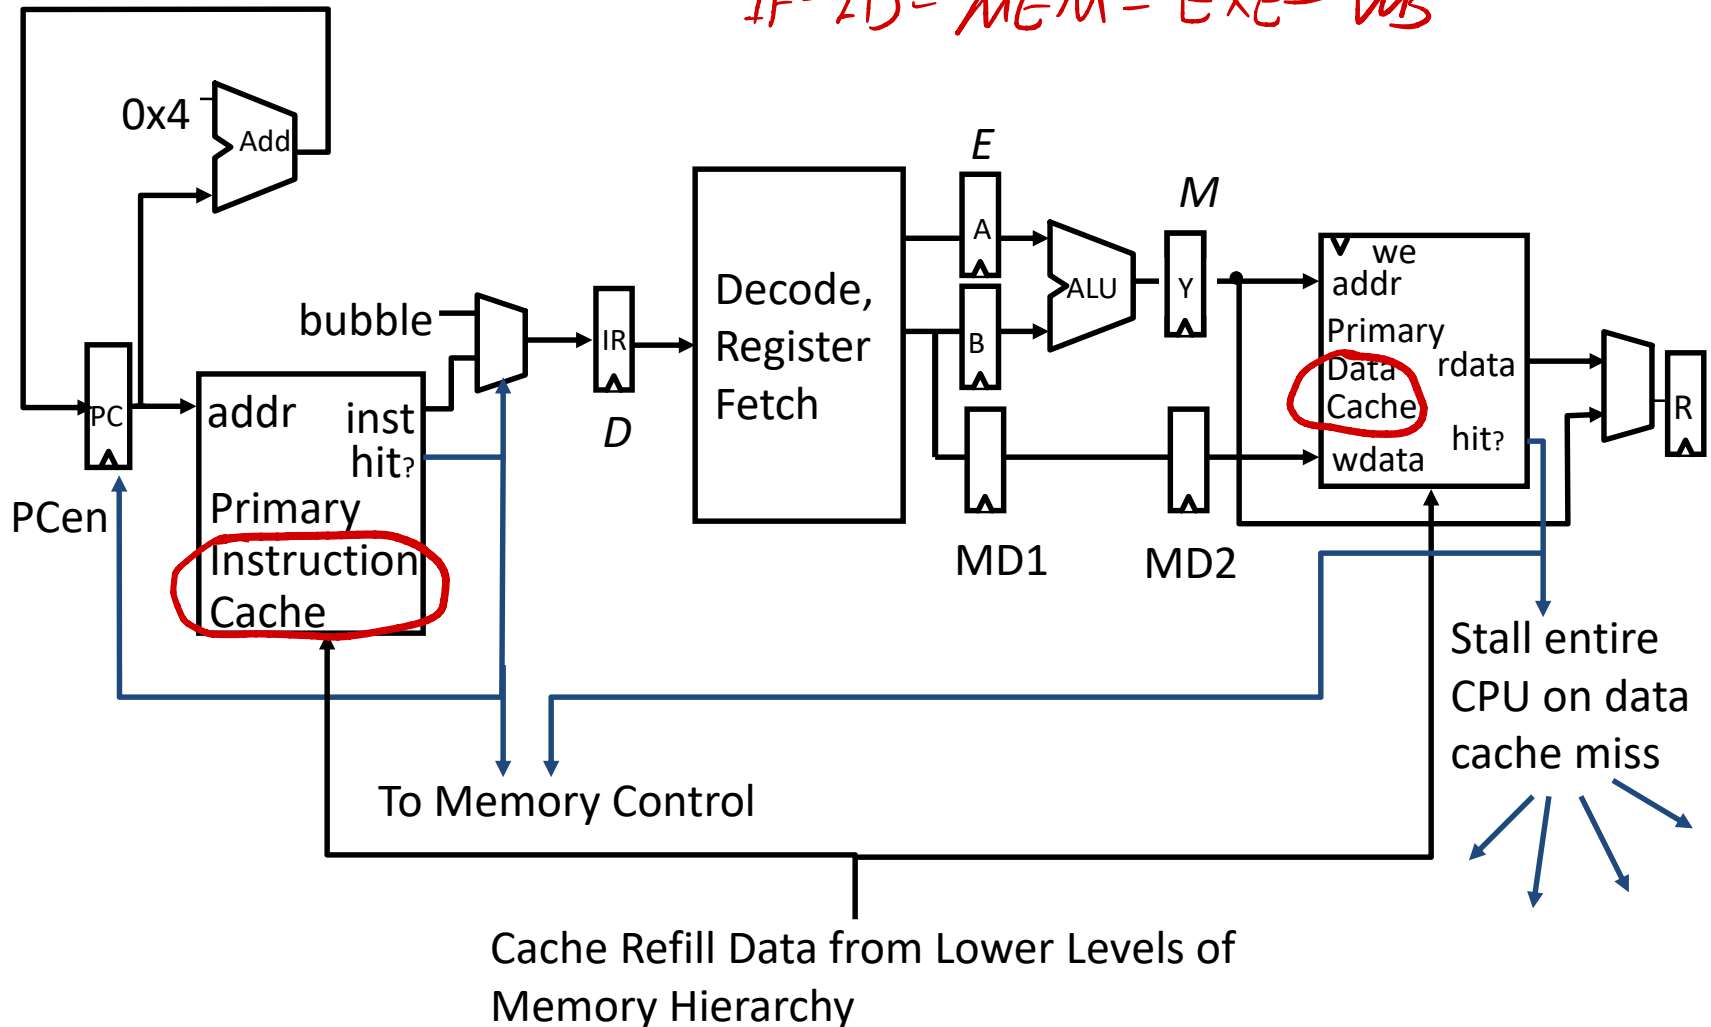
Slides based on UC Berkeley's CS61C

Cache Terms I

- ***Cache***:
  - A small and fast memory used to increase the performance of accessing a big and slow memory
  - Uses ***temporal locality:*** The tendency to reuse data in the same space over time
  - Uses ***spatial locality:*** The tendency to use data at nearby addresses
- Cache ***hit***: The address being fetched is in the cache ☺
- Cache ***miss:*** The address being fetched is not in the cache ☹
- ***Valid bit***: Is a particular entry valid
- Cache ***line flush***: Invalidate and flush one entry
  - e.g., clflush of x86, but newer clwb may not invalidate the entry
- Cache ***flush***: Invalidate all entries
  - e.g., wbinvd of x86

CPU-Cache Interaction

(5-stage pipeline)

IF-ID- MEM- EXE- WB

To Memory Control

Cache Refill Data from Lower Levels of Memory Hierarchy

Stall entire CPU on data cache miss

4

Cache Terms II

- Cache *level*:
  - The order in the memory hierarchy: L1$ is closest to the processor
  - L1 caches may only hold data (Data-cache, D$) or instructions (Instruction Cache, I$)
    - Most L2+ caches are 'unified', can hold both instructions and data
- Cache *capacity*:
  - The total # of bytes in the cache
- Cache *line* or cache *block*:
  - A single entry in the cache
- Cache *block size*:
  - The number of bytes in each cache line

[Handwritten annotations in red:]
离CPU近, 速度快, 容量小 ↑
数据, 指令分开
L2往下↓ 数据, 指令混合
⇒ cache line的 (tag设单上)
Cache line: 交换的最小单元

Cache Terms III
Associativity

- **Number of cache lines**:
  - Cache capacity / block size:
- Cache **associativity**:
  - The number of possible cache lines a given address may exist in.
  - Also the number of comparison operations needed to check for an element in the cache
  - **Direct mapped**: A data element can only be in one possible location (N=1)
  - 组相连 — **N-way set associative**: A data element can be in one of N possible positions
  - 全相连 — **Fully associative**: A data element can be at any location in the cache.
    - Associativity == # of lines
- Total # of cache lines == capacity of cache/line size

- Total # of lines in a set == # ways == N == associativity
- Total # of sets == # of cache lines / associativity

Victim Cache *2nd chance*

- Conflict misses are a pain, but...
  - Perhaps a little associativity can help without having to be a fully associative cache
- In addition to the main cache...   *很小 ⇒ 全相连*
  - Optionally have a very small (16-64 entry) **fully associative** "victim" cache
- Whenever we evict a cache entry
  - Don't just get rid of it, put it in the victim cache
- Now on cache misses...
  - Check the victim cache first, if it is in the victim cache you can just reload it from there

Cache Terms IV
Parts of the Address

- Address is divided into |***TAG*|*INDEX*|*OFFSET*|**
- ***Offset***:   *byte offset*
  - The lowest bits of the memory address which say where data exists within the cache line.
  - It is log2(line/block size)
  - So for a cache with 64B blocks it is 6 bits   *64 Byte = 2^6 Byte*
- ***Index***:
  - The portion of the address which says where in the cache an address may be stored
  - Takes log2(# of cache lines / associativity) bits $\frac{512}{4} = 128$ *line/way*
  - So for a 4-way associative cache with 512 lines it is 7 bits 2^7
- ***Tag***: The portion of the address which must be stored in the cache to check if a location matches
  - # of bits of address - (# of bits for index + # of bits for offset)
  - So with 64-bit addresses it is 51-bit...

Cache Terms V
Writing

- **Eviction**:
  - The process of removing an entry from the cache
- **Write Back**: 写回
  - A cache which only writes data up the hierarchy when a cache line is evicted
  - Instead set a **dirty bit** on cache entries
  - All i7 caches are **write back**

  本层收草, replacement
  发生时才往下写

- **Write Through**: 写穿
  - A cache which always writes to memory
- **Write Allocate**:
  - If writing to memory **not in the cache** fetch it first
  - i7 L2 is Write Allocate
- **No Write Allocate**:
  - Just write to memory without a fetch   写到L2/memory
  - i7 L1 is no write allocate

Replacement Policy

In an associative cache, which line from a set should be evicted when the set becomes full?

1 • Random

5. Least frequently used.

2 • Least-Recently Used (LRU)
   - • LRU cache state must be updated on every access
   - • True implementation only feasible for small sets (2-way)
   - • Pseudo-LRU binary tree often used for 4-8 way

3 • First-In, First-Out (FIFO) a.k.a. Round-Robin
   - • Used in highly associative caches

4 • Not-Most-Recently Used (NMRU)
   - • FIFO with exception for most-recently used line or lines

This is a second-order effect. Why?

(次生灾害)

Replacement only happens on misses

Cache Terms VI
Cache Performance

- **Hit Time**:
  - Amount of time to return data in a given cache: depends on the cache
  - i7 L1 hit time: 4 clock cycles
- **Miss Penalty**:
  - Amount of **additional** time to return an element if its not in the cache: depends on the cache
- **Miss Rate**:
  - Fraction of a **particular program's** memory requests which miss in the cache
- Average Memory Access Time (**AMAT**):
  - Hit time + Miss Rate * Miss Penalty

Understanding Cache Misses: The 3Cs

- Compulsory (cold start or process migration, 1st reference):
  - First access to block impossible to avoid; small effect for long running programs
  - Solution: increase block size (increases miss penalty; very large blocks could increase miss rate)

  导致

  默认只有一个
  变重改变

  eg. Size → ∞, 剩余全在 Cache line 中

- Capacity:

  Cache 不变, way 不变, block Size↑, block 数↓

  - Cache cannot contain all blocks accessed by the program
  - Solution: increase cache size (may increase access time)

- *Conflict (collision):*
  - *Multiple memory locations mapped to the same cache location*
  - *Solution 1: increase cache size*
 - *Solution 2: increase associativity (may increase access time)*

Prefetching...

- Programmer/Compiler: I know that, later on, I will need this data...

  compulsory miss ↓

- Tell the computer to **_prefetch_** the data
  - Can be as an explicit prefetch instruction
  - Or an implicit instruction: **lw x0 0(t0)**
    - Won't stall the pipeline on a cache miss: The processor control logic recognizes this situation

- Allows you to hide the cost of compulsory misses
  - You still need to fetch the data however

How to Calculate 3C's using Cache Simulator

1. *Compulsory*: set cache size to infinity and fully associative, and count number of misses

2. *Capacity*: Change cache size from infinity, usually in powers of 2, and count misses for each reduction in size

    – 16 MB, 8 MB, 4 MB, ... 128 KB, 64 KB, 16 KB

3. *Conflict*: Change from fully associative to n-way set associative while counting misses

    – Fully associative, 16-way, 8-way, 4-way, 2-way, 1-way

14

3Cs Analysis

- Three sources of misses (SPEC2000 integer and floating-point benchmarks)
  - Compulsory misses 0.006%; not visible
  - Capacity misses, function of cache size
  - Conflict portion depends on associativity and cache size

15

Improving Cache Performance

AMAT = Time for a hit + Miss rate x Miss penalty

- Note: miss penalty is **additional** time for cache miss
- Reduce the time to hit in the cache
  - E.g., Smaller cache

  cache ↓   hit time ↓   miss rate ↑

- Reduce the miss rate
  - E.g., Bigger cache
    Longer cache lines (somewhat: improves ability to exploit spatial locality at the cost of reducing the ability to exploit temporal locality)
  - E.g., Better programs!
- Reduce the miss penalty
  - E.g., Use multiple cache levels
- Hit and Miss, 3C

Impact of Larger Cache on AMAT?

- 1) Reduces misses (what kind(s)?)
- 2) Longer Access time (Hit time): smaller is faster
  - Increase in hit time will likely add another stage to the pipeline
- At some point, increase in hit time for a larger cache may overcome the improvement in hit rate, yielding a decrease in performance
- Computer architects expend considerable effort optimizing organization of cache hierarchy – big impact on performance and power!

Increasing Associativity?

- Hit time as associativity increases?
  - Increases, with large step from direct-mapped to >=2 ways, as now need to mux correct way to processor
  - Smaller increases in hit time for further increases in associativity

    ways↑ mux↑ hit time↑

- Miss rate as associativity increases?
  - Goes down due to reduced conflict misses, but most gain is from 1->2->4-way with limited benefit from higher associativities

- Miss penalty as associativity increases?
  - Unchanged, replacement policy runs in parallel with fetching missing line from memory

Increasing #Entries? → Cache 容量

[handwritten: hit time ↑]
[handwritten: hit miss ↓]

- Hit time as #entries increases?
  – Increases, since reading tags and data from larger memory structures

- Miss rate as #entries increases?
  – Goes down due to reduced capacity and conflict misses
  – *Architects rule of thumb: miss rate drops ~2x for every ~4x increase in capacity (only a gross approximation)*

- Miss penalty as #entries increases?
  – Unchanged

At some point, increase in hit time for a larger cache may overcome the improvement in hit rate, yielding a decrease in performance
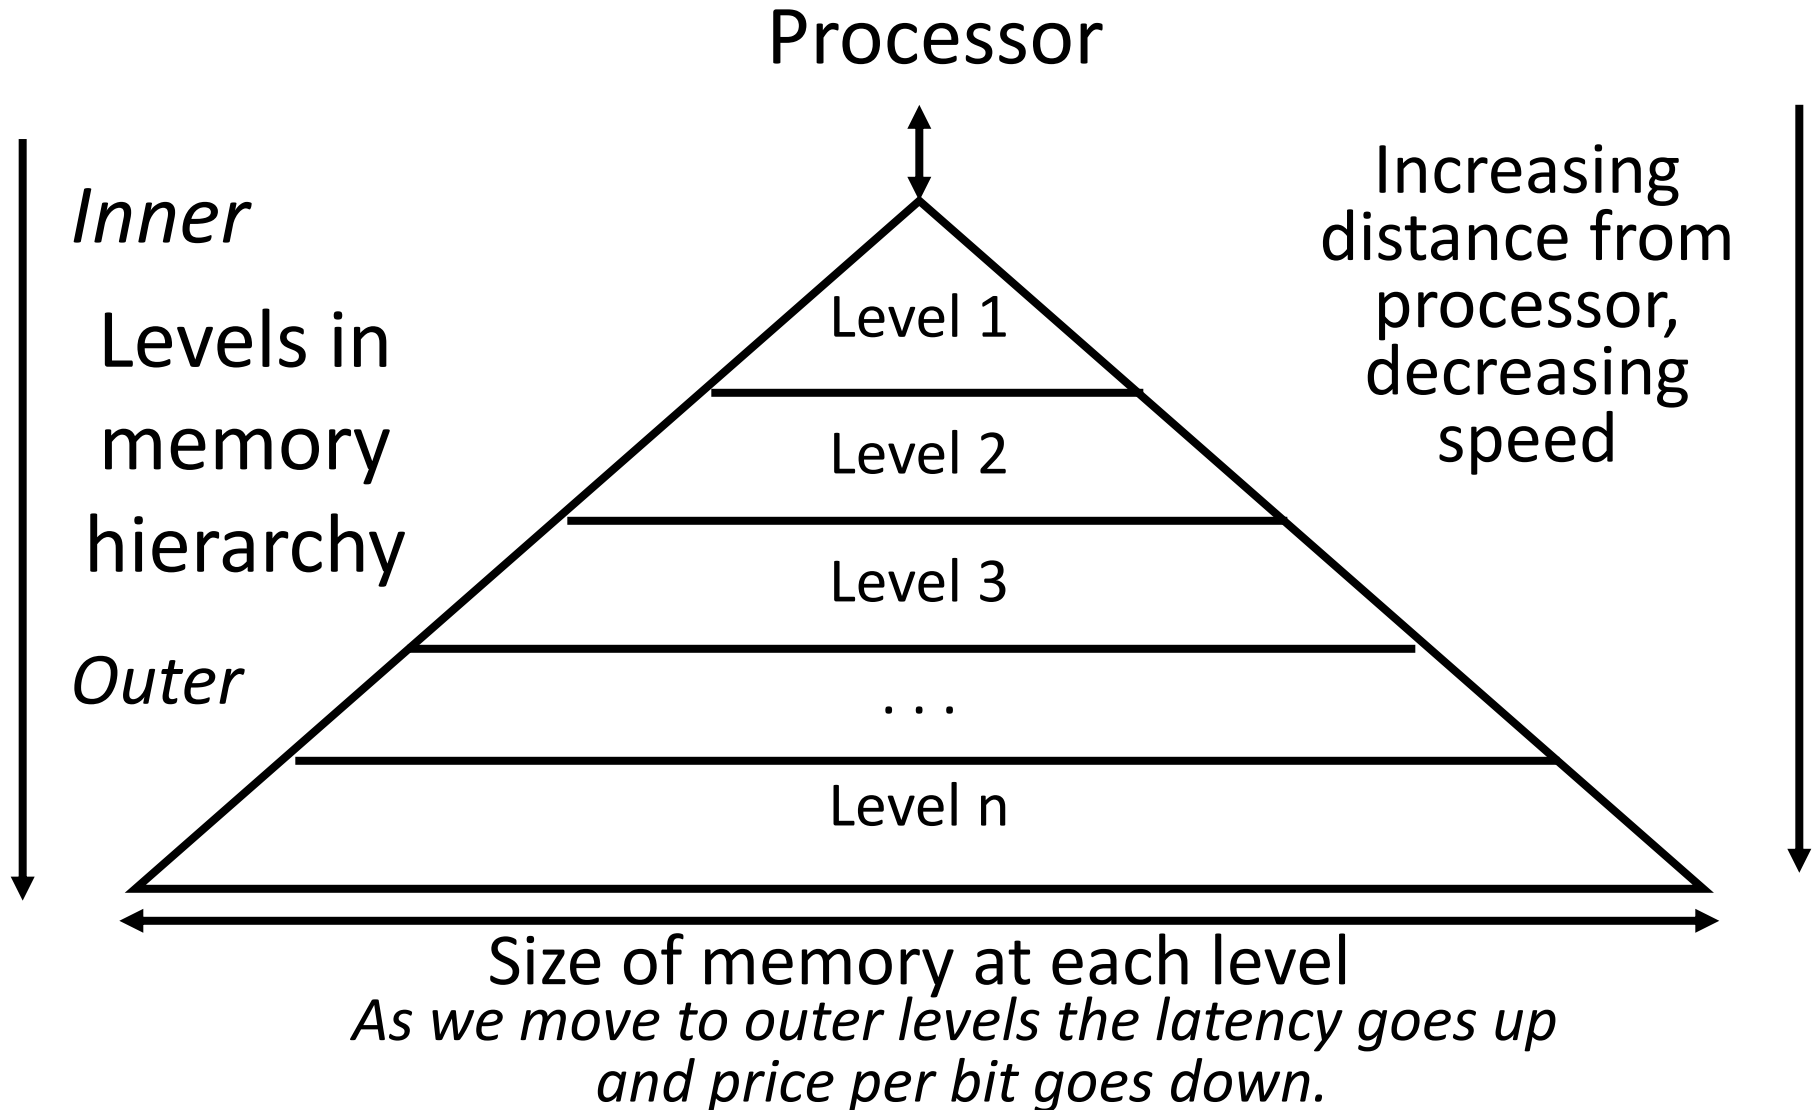
Increasing Block Size?

- Hit time as block size increases?
  - Hit time unchanged, but might be slight hit-time reduction as number of tags is reduced, so faster to access memory holding tags
- Miss rate as block size increases?
  - Goes down at first due to spatial locality, then increases due to increased conflict misses due to fewer blocks in cache
- Miss penalty as block size increases?
  - Rises with longer block size, but with fixed constant initial latency that is amortized over whole block
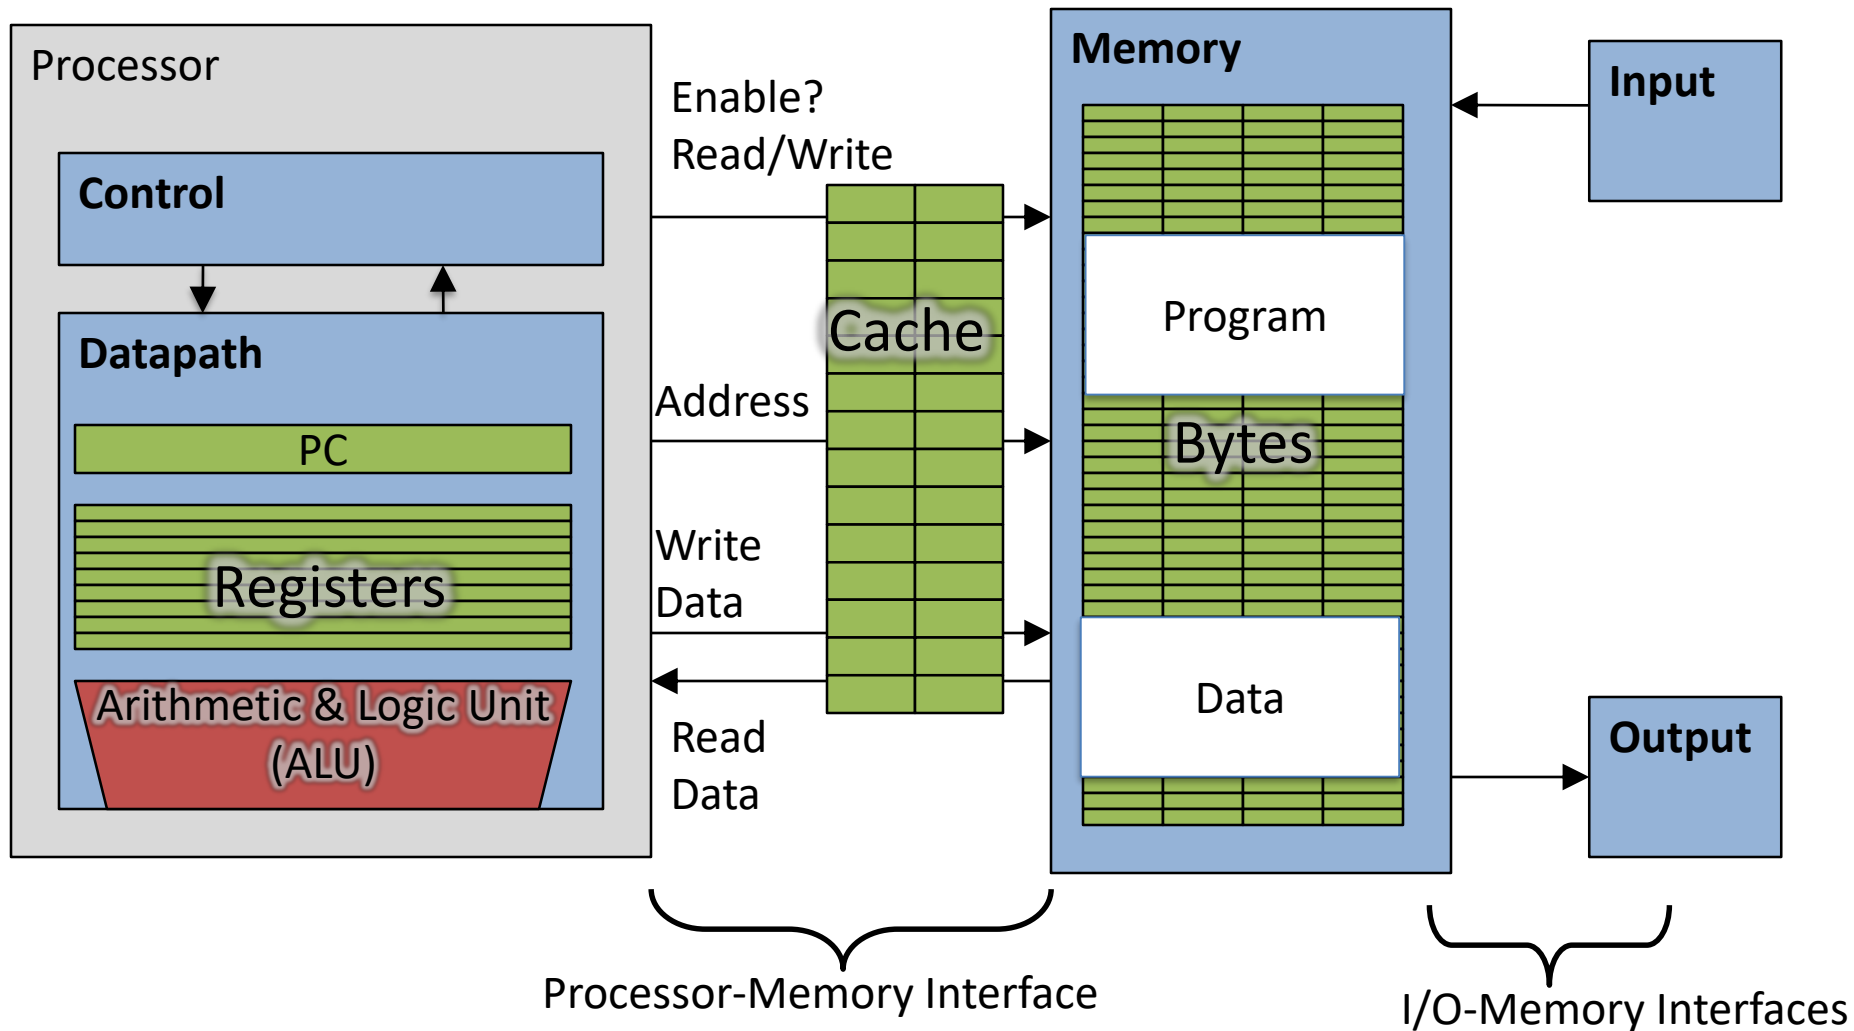
How to Reduce Miss Penalty?

- Could there be locality on misses from a cache?

- Use multiple cache levels!

- With Moore's Law, more room on die for bigger L1 caches and for second-level (L2) cache

- And in some cases even an L3 cache!

- IBM mainframes have ~1GB L4 cache off-chip.

Review: Memory Hierarchy

Processor

Inner

Levels in memory hierarchy

Outer

Level 1

Level 2

Level 3

. . .

Level n

Increasing distance from processor, decreasing speed

Size of memory at each level
As we move to outer levels the latency goes up and price per bit goes down.

Adding Cache to Computer

Processor

Control

Datapath

PC

Registers

Arithmetic & Logic Unit (ALU)

Enable?
Read/Write

Address

Write Data

Read Data

Cache

Memory

Program

Bytes

Data

Input

Output

Processor-Memory Interface

I/O-Memory Interfaces

L1 and L2 Caches

27
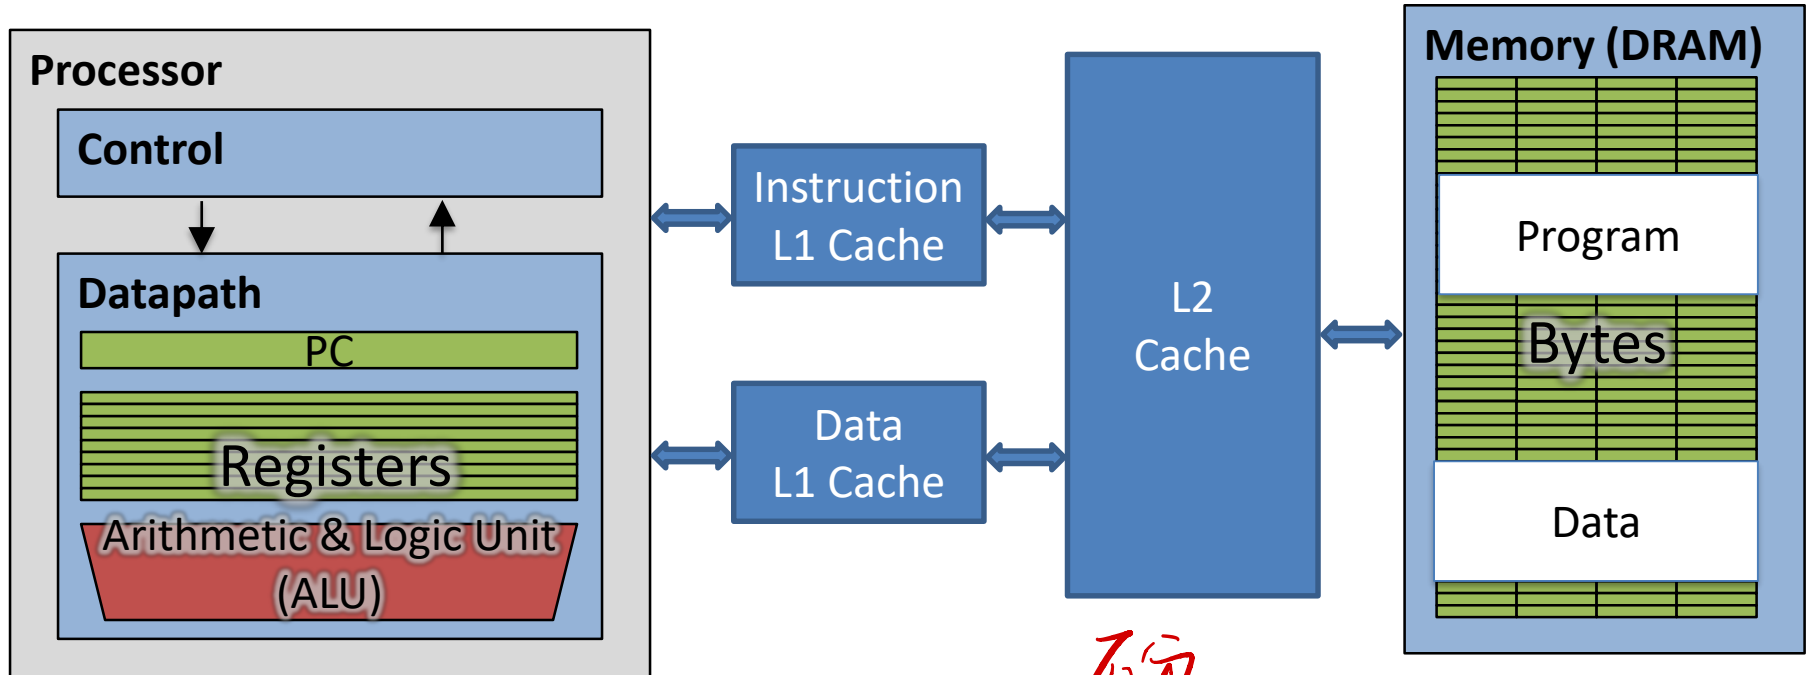
Typical Memory Hierarchy

	On-Chip Components				
Speed (cycles):	½'s	1's	10's	100's	1,000,000's
Size (bytes):	100's	10K's	M's	G's	T's
Cost/bit:	highest				lowest
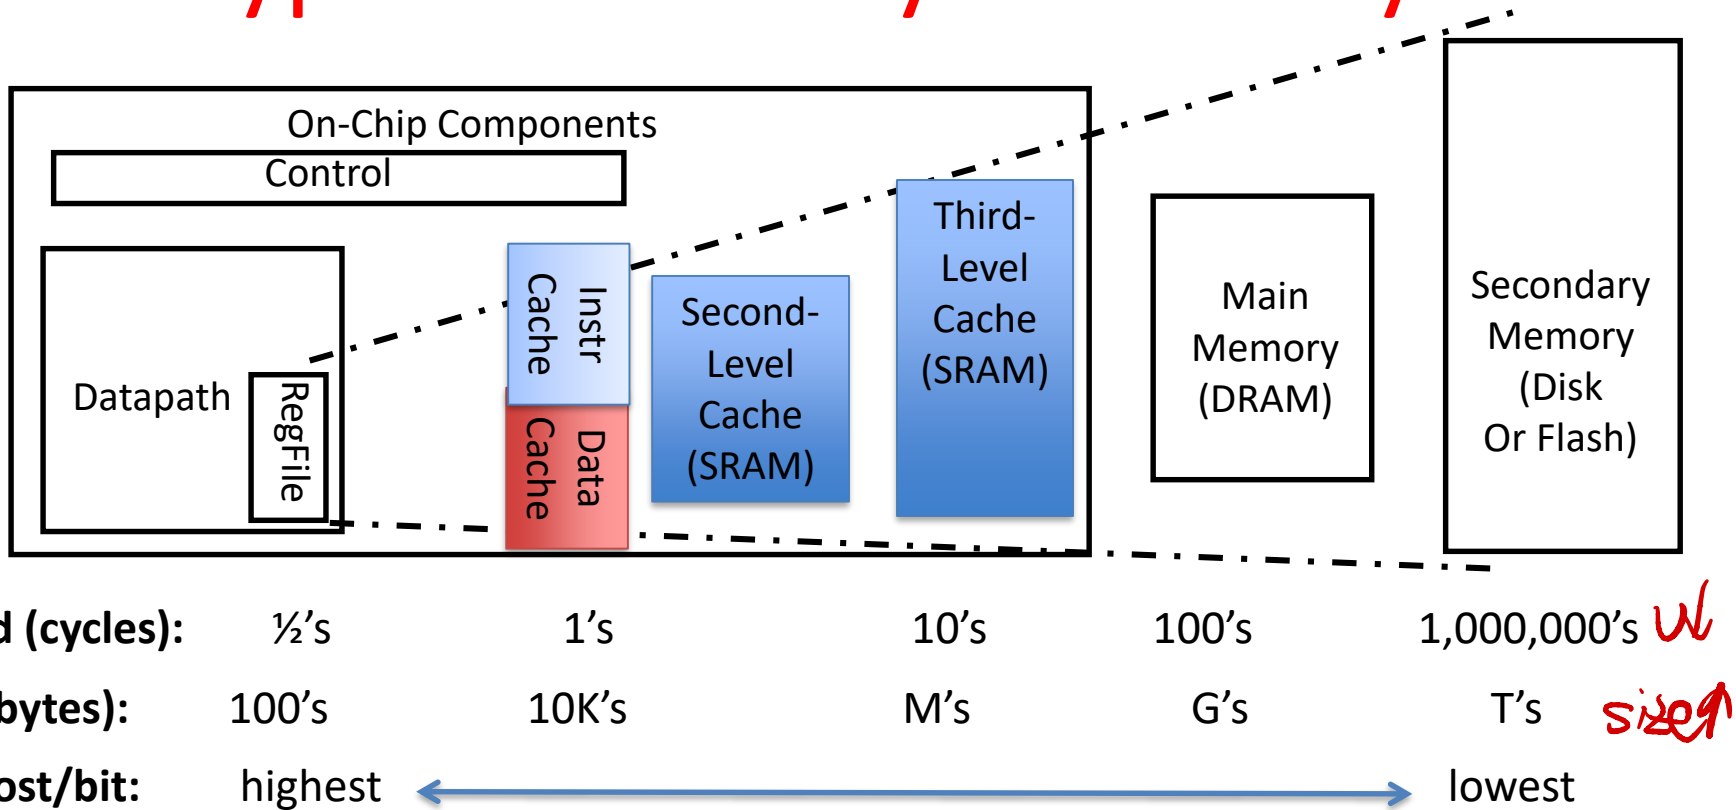
- Principle of locality + memory hierarchy presents programmer with ≈ as much memory as is available in the *cheapest* technology at the ≈ speed offered by the *fastest* technology
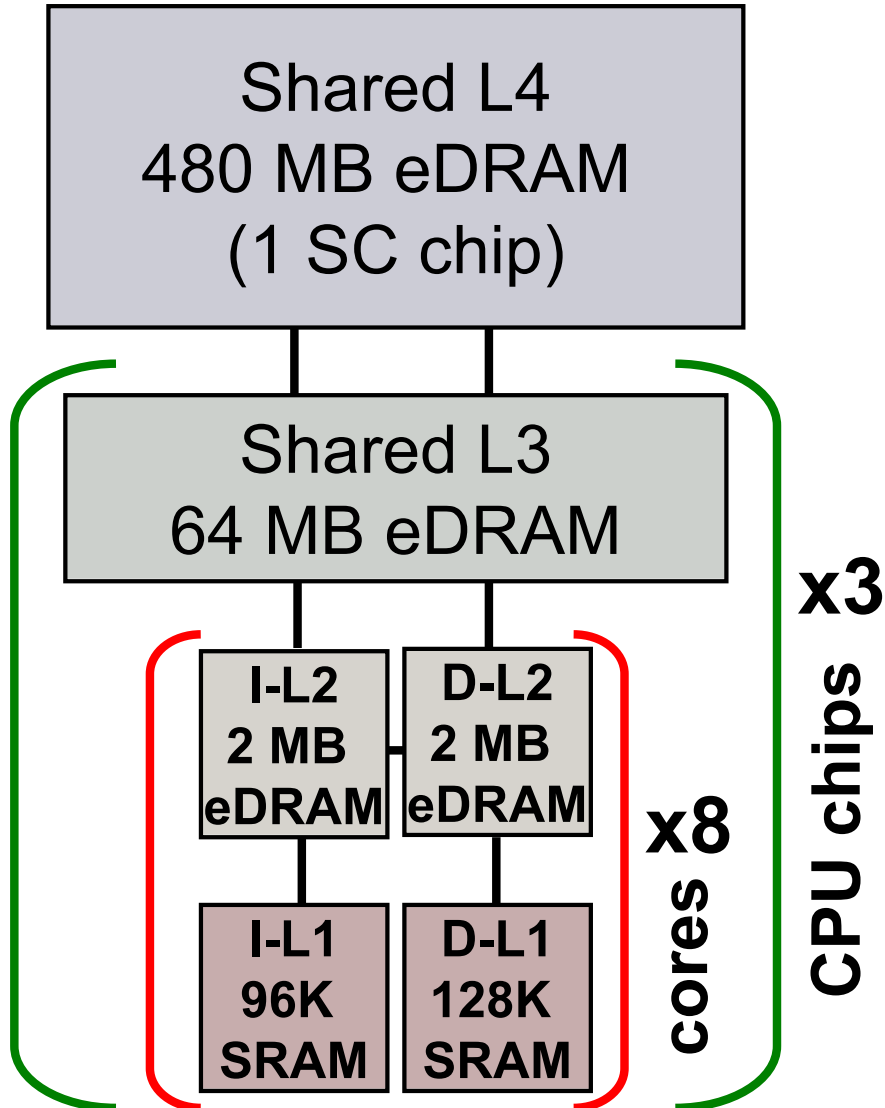
28

How is the Hierarchy Managed?

- registers ↔ memory
  - By compiler (or assembly level programmer)
- cache ↔ main memory
  - By the cache controller hardware

OS管不了cache

- main memory ↔ disks (secondary storage)
  - By the operating system (virtual memory)
  - Virtual to physical address mapping assisted by the hardware ('translation lookaside buffer' or TLB)
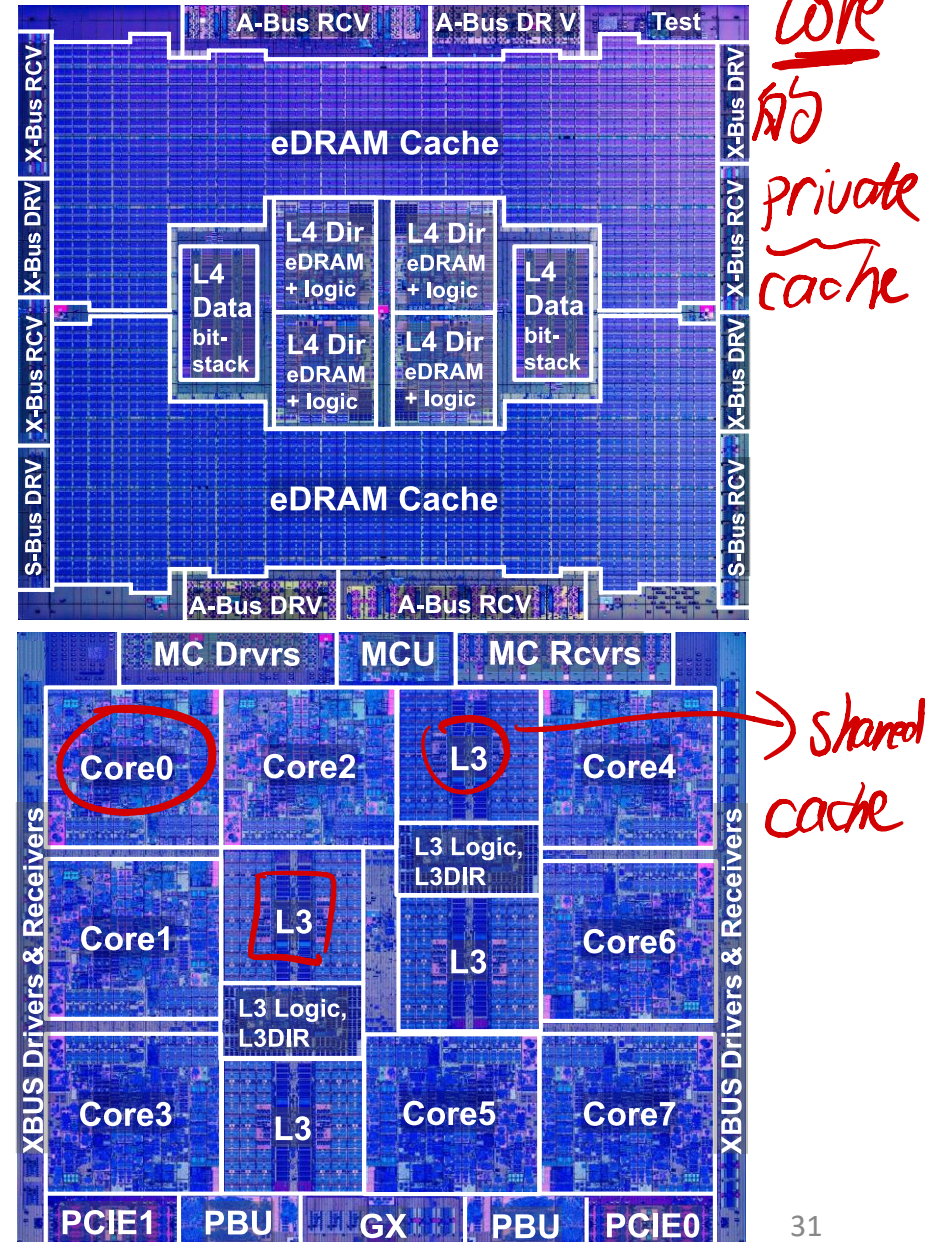 - By the programmer (files)

2015 IBM CPU

- z13 designed in 22nm SOI technology with seventeen metal layers, 4 billion transistors/chip

- 8 cores/chip, with 2MB L2 cache, 64MB L3 cache, and 480MB L4 off-chip cache.

- 5GHz clock rate, 6 instructions per cycle, 2 threads/core

- Up to 24 processor chips in shared memory node

IBM z13 Memory Hierarchy

L_1, L_2:
core 的 private cache

Shared L4
480 MB eDRAM
(1 SC chip)

Shared L3
64 MB eDRAM

x3 CPU chips

I-L2
2 MB
eDRAM

D-L2
2 MB
eDRAM

x8 cores

I-L1
96K
SRAM

D-L1
128K
SRAM

每个 core 都有 private L_1, L_2

→ shared cache

31

Local vs. Global Miss Rates

本层的

- *Local miss rate* – the fraction of references to one level of a cache that miss

- Local Miss rate L2$ = L2$ Misses / L1$ Misses

  L2的所有访问

  非L1的miss

  = L2$ Misses / total_L2_accesses

- *Global miss rate* – the fraction of references that miss in all levels of a multilevel cache

  - L2$ local miss rate >> than the global miss rate

L1, L2均 miss

L1 Cache: 32KB I$, 32KB D$
L2 Cache: 256 KB
L3 Cache: 4 MB

hit   miss
$L_1 > L_2 > L_3$

FIGURE 5.47 The L1, L2, and L3 data cache miss rates for the Intel Core i7 920 running the full integer SPECCPU2006 benchmarks.

Local vs. Global Miss Rates

- *Local miss rate* – the fraction of references to one level of a cache that miss
- Local Miss rate L2$ = $L2 Misses / L1$ Misses
- *Global miss rate* – the fraction of references that miss in all levels of a multilevel cache
  - L2$ local miss rate >> than the global miss rate
- Global Miss rate = L2$ Misses / Total Accesses
  = (L2$ Misses / L1$ Misses) × (L1$ Misses / Total Accesses)
  = Local Miss rate L2$ × Local Miss rate L1$

- AMAT = Time for a hit + Miss rate × Miss penalty
- AMAT = Time for a L1$ hit + (local) Miss rate L1$ ×
  (Time for a L2$ hit + (local) Miss rate L2$ × L2$ Miss penalty)

Characteristic	Intel Nehalem	AMD Opteron X4 (Barcelona)
L1 cache organization	Split instruction and data caches	Split instruction and data caches
L1 cache size	32 KB each for instructions/data per core	64 KB each for instructions/data per core
L1 block size	64 bytes	64 bytes
L1 write policy	Write-back, Write-allocate	Write-back, Write-allocate
L1 hit time (load-use)	Not Available	3 clock cycles
L2 cache organization	Unified (instruction and data) per core	Unified (instruction and data) per core
L2 cache size	256 KB (0.25 MB)	512 KB (0.5 MB)
L2 block size	64 bytes	64 bytes
L2 write policy	Write-back, Write-allocate	Write-back, Write-allocate
L2 hit time	Not Available	9 clock cycles
L3 cache organization	Unified (instruction and data)	Unified (instruction and data)
L3 cache size	8192 KB (8 MB), shared	2048 KB (2 MB), shared
L3 block size	64 bytes	64 bytes
L3 write policy	Write-back, Write-allocate	Write-back, Write-allocate
L3 hit time	Not Available	38 (?)clock cycles

CPI/Miss Rates/DRAM Access SpecInt2006

Name	CPI	L1 D cache misses/1000 instr Data Only	L2 D cache misses/1000 instr Data Only	DRAM accesses/1000 instr Instructions and Data
perl	0.75	3.5	1.1	1.3
bzip2	0.85	11.0	5.8	2.5
gcc	1.72	24.3	13.4	14.8
mcf	10.00	106.8	88.0	88.5
go	1.09	4.5	1.4	1.7
hmmer	0.80	4.4	2.5	0.6
sjeng	0.96	1.9	0.6	0.8
libquantum	1.61	33.0	33.1	47.7
h264avc	0.80	8.8	1.6	0.2
omnetpp	2.94	30.9	27.7	29.8
astar	1.79	16.3	9.2	8.2
xalancbmk	2.70	38.0	15.8	11.4
Median	1.35	13.6	7.5	5.4

In Conclusion, Cache Design Space

- Several interacting dimensions
  - Cache size
  - Block size
  - Associativity
  - Replacement policy
  - Write-through vs. write-back
  - Write-allocation
- Optimal choice is a compromise
  - Depends on access characteristics
    - Workload
    - Use (I-cache, D-cache)
  - Depends on technology / cost
- Simplicity often wins

Cache Size

Associativity

Block Size

Bad

Good

Factor A   Factor B

Less   **More**