

Course Info

- No Lab/discussion this week
- Project 2.2 will be available soon.
- HW5 available.
- MID-TERM II this Thursday!!!



信息科学与技术学院

School of Information Science and Technology

CS 110

Computer Architecture

Heterogeneous Computing & FPGA

异构计算

Instructors:

Siting Liu & Chundong Wang

Course website: <https://toast-lab.sist.shanghaitech.edu.cn/courses/>

[CS110@ShanghaiTech/Spring-2023/index.html](https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html)

School of Information Science and Technology (SIST)

ShanghaiTech University

2023/2/7

Hardware vs. Software

- How to perform an addition?

CPU

```
lw t1, address
lw t2, address
add t0, t1, t2
sw t0, address
```

IF → ID → EX → MEM → WB

300 ps 100 ps 200 ps 300 ps 100 ps

Very complex circuitry

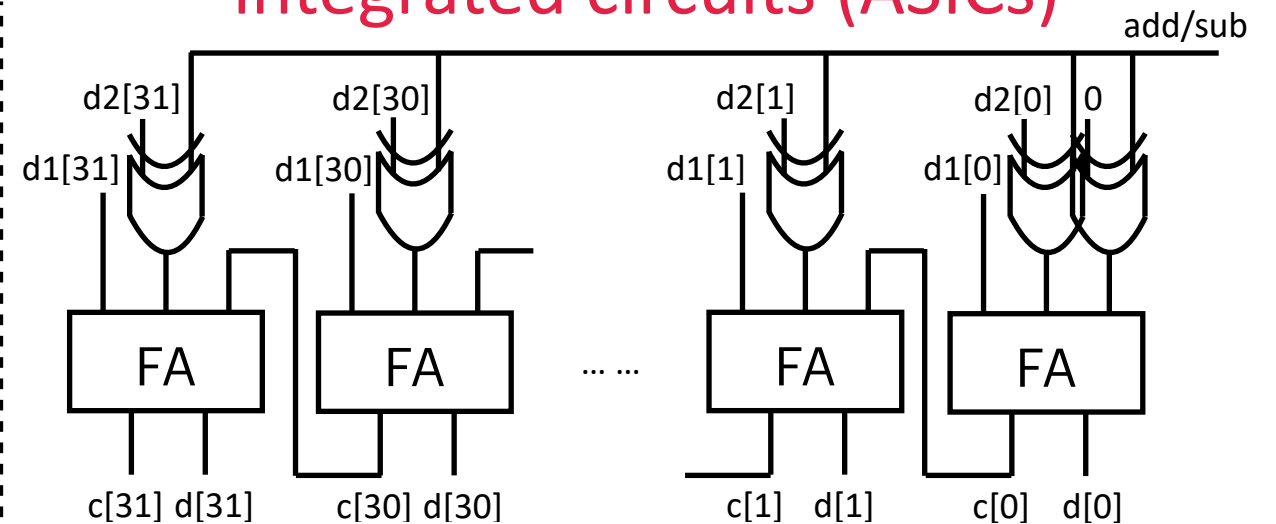
Relatively high power consumption

High flexibility

Software-programmable

可编程性

Application-specific
integrated circuits (ASICs)



Only EX

200 ps

Only use hundreds of gates

Low-power

Only perform addition/subtraction

Not programmable

Hardware vs. Software

- For a task, we can always write software code or build ASIC hardware.

Software solution	Hardware solution
+/add instruction	Build an adder
Shift-and-add	RV-M extension & add multiplier
Insert "nop" to avoid stall	Add hardware to avoid stall
...

CPU

DSA

ASIC

Software-programmable
General-purpose
Less efficient

*NRE
non-recurring engineering*

Not programmable
or limited
programmability
Application-specific
More efficient
High *NRE* cost

Domain-Specific Architecture

- Build hardware for an application “domain” instead of a certain task
- Moderate software-programmable, with relatively higher efficiency compared with CPU

GPU

For graphics

Rasterization/texture/rendering, etc.

Vector/Matrix operations

CUDA (2006), openCL, etc.

AI-accelerators/NPU/AI chips

For neural networks

Tensor operations, etc.

No unified programming language yet

Still a fast-developing field

GPGPU

通用GPU

CPU

DSA

针对某一特定场景

ASIC

Software-programmable

General-purpose

Less efficient

Low NRE cost

Software-programmable

Domain-specific

Efficient

Low NRE cost

Not programmable
or limited
programmability

Application-specific

More efficient


High NRE cost

Heterogeneous Computing

异构计算
多个不同架构共同组合

- Usually cannot work independently
 - **Parallel to CPU, as an I/O device**
 - Integrated in an SoC, as a co-processor

Heterogeneous computing refers to systems that use more than one kind of processor or core.



Heterogeneous Computing

- Usually cannot work independently
 - Parallel to CPU, as an I/O device
 - **Integrated in an SoC, as a co-processor**
 - Historically, math processor
 - Integrated GPU in CPU
 - etc.
- Image signal processors (ISP)

Huawei Kirin 990

Die shot source: [TechInsights](#) - Labelling & Custom contrast: AnandTech



Heterogeneous Computing

Apple M1


Die shot source: [TechInsights](#)

Image signal processors (ISP)

AES encryption/security

Video decoder/encoder

Co-processor examples



Heterogeneous Computing

Company	Name	Process	CPU	GPU	NPU
HiSilicon	Kirin 980 [1]	TSMC 7nm	2 big Cortex-A76 @ 1.92 GHz 2 big Cortex-A76 @ 2.6 GHz 4 little Cortex-A55 @ 1.8 GHz	ARM 10-core Mali-G76 (FP32 480 GFLOPs)	Dual NPU
Apple	M1Max [2][3]	TSMC 5nm	2 HE cores @ 2.064 GHz 8 HP cores @ 3.228 GHz	24/32-core GPU (FP32 10.4 TFLOPs)	16-core neural engine (11 TOPs)

Special cases:
SoC with different
CPU cores also
considered
heterogeneous

Another Dimension—Hardware Programmable

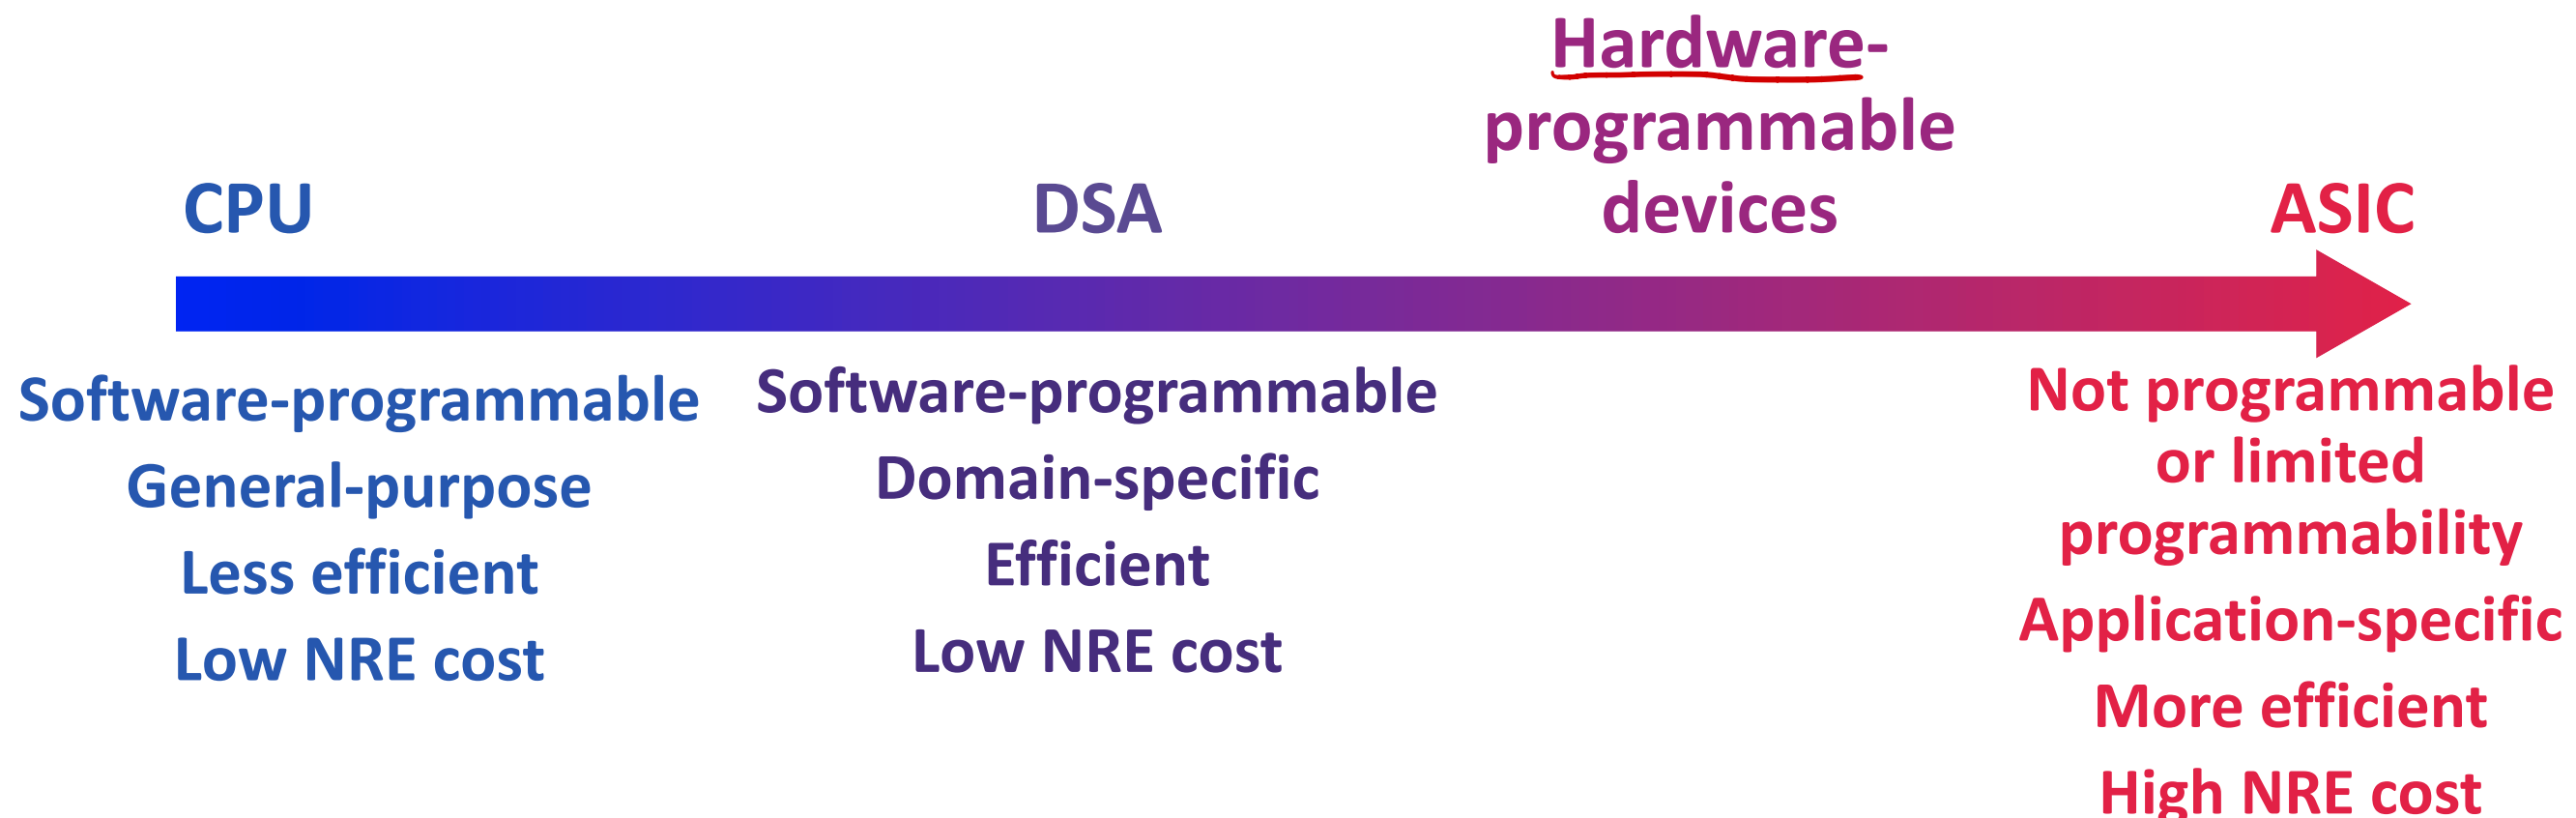
Image signal processors (ISP)

AES encryption/security

Video decoder/encoder

- Cannot continue the list forever (algorithms are evolving)
- NRE cost non-negligible

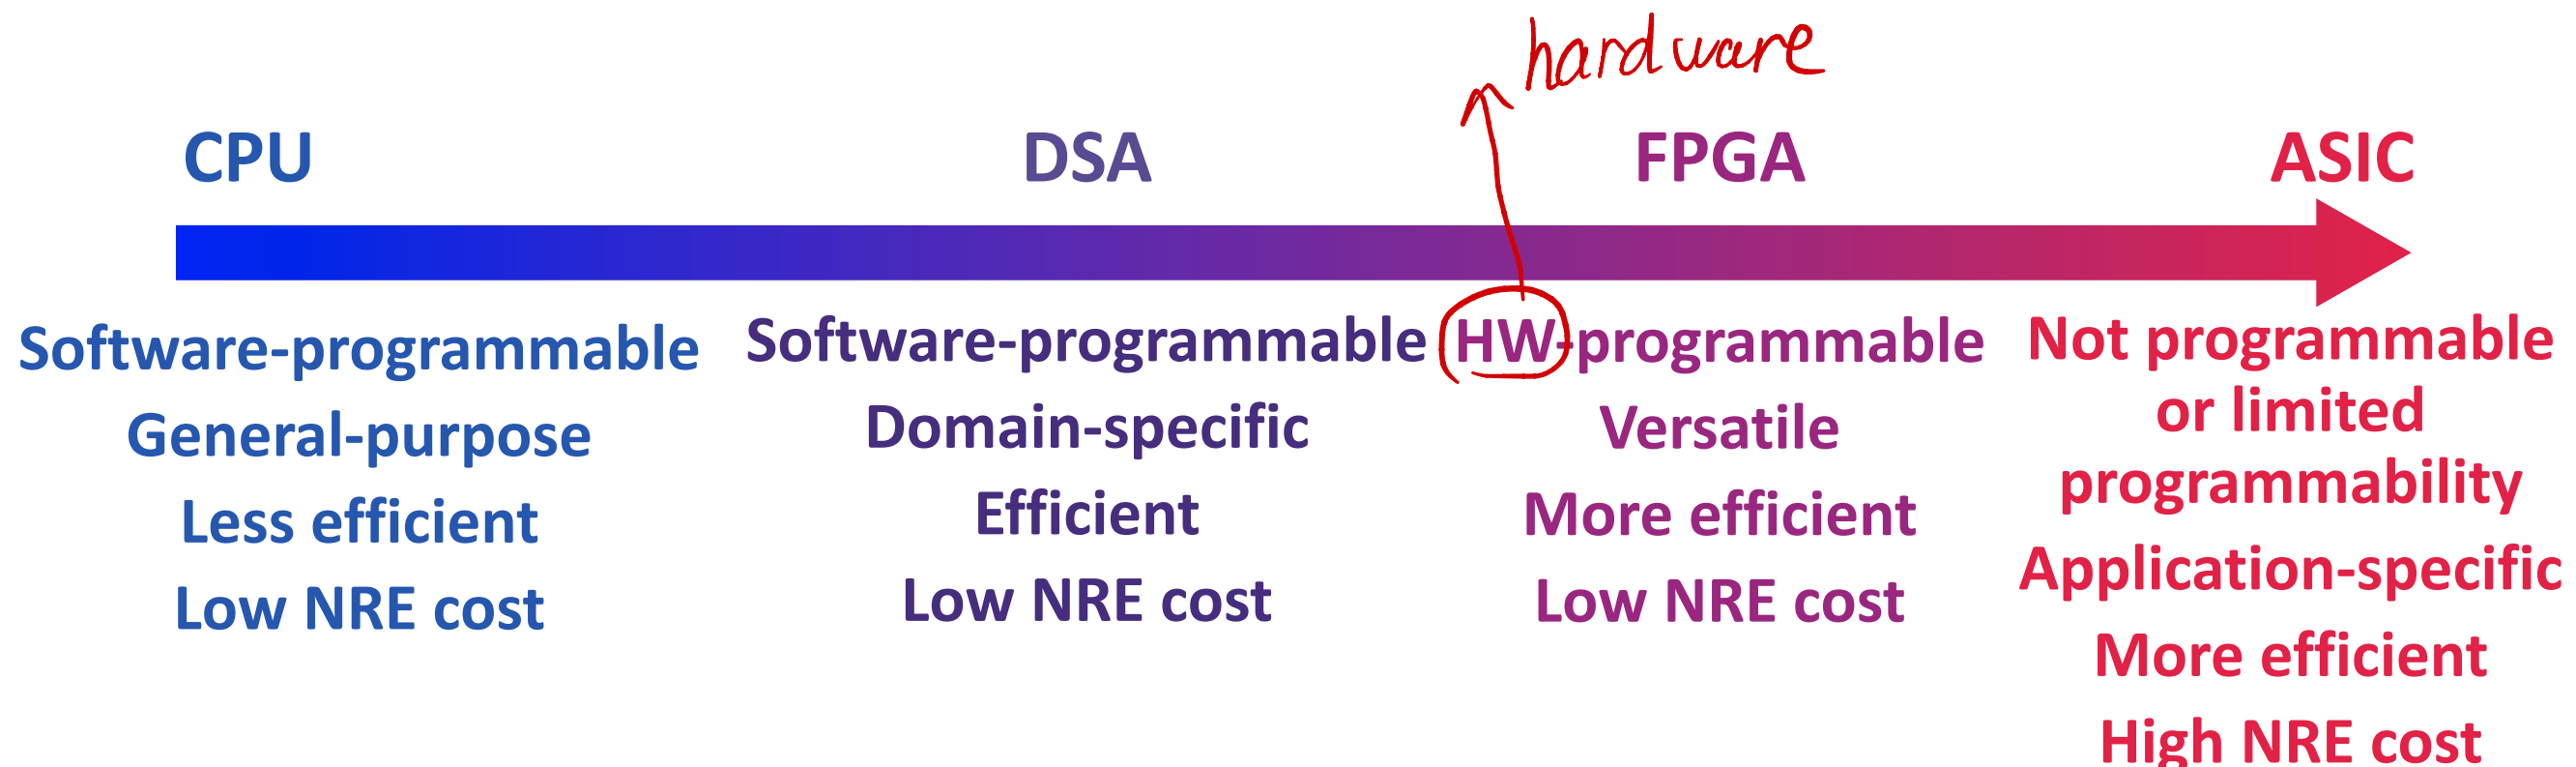
Co-processor examples



Field-Programmable Gate Array (FPGA)

现场可编程阵列

- The hardware functionality can be changed by programming (mostly HDLs such as Verilog & VHDL)
- FPGA can implement any digital circuits with a certain size
- Shorter time-to-market vs. ASIC 开发过程短
- Heterogeneous w.r.t. CPU/GPU



Another Dimension—Hardware Programmable

Hardware-
programmable


Providers	CPU	GPU	FPGA	ASIC (DSA)
Alibaba Cloud	X86/ARM/RISC-V	Nvidia/AMD	Intel (Altera)/ AMD (Xilinx)	AliNPU
AWS (Amazon)	Graviton (ARM) / X86	Nvidia/AMD	Xilinx	AWS Trainium
Azure (MS)	X86	Nvidia	Certain DNN models	
Baidu Cloud	X86	Nvidia	Xilinx	Kunlun
Google Cloud	X86	Nvidia	N/A	TPU
Huawei Cloud	Kunpeng (ARM)/X86	Nvidia & Ascend	Xilinx	Ascend
Tencent Cloud	X86	Nvidia & Xinghai	Xilinx	Enflame-tech (燧原)

FPGA Applications

- Communication (decoding/encoding algorithms, etc.)
 - Smart networking device
 - SmartNIC
- FPGA trading systems (High-frequency trading)
- AI tasks (MS Project Brainwave/Xilinx Vitis AI)
- Also in embedded systems 嵌入式系统
 - Digital signal processing
 - Image signal processing
 - Control logics
- Also in IC design, for hardware emulation

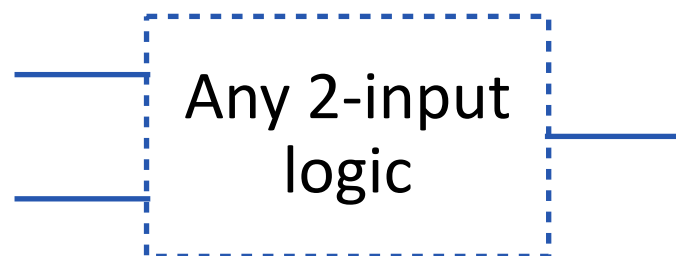


How?



Configurable Logic Blocks

Programmable/
Configurable
logic blocks



Combinational
circuits


In1	In2	Out
0	0	O0
0	1	O1
1	0	O2
1	1	O3

We store the truth table and it can implement **any 2-input logic**

SRAM-based FPGA


- Look-up table (LUT)

SRAM cells
(1-bit each)



In1	In2	Out
0	0	O0
0	1	O1
1	0	O2
1	1	O3

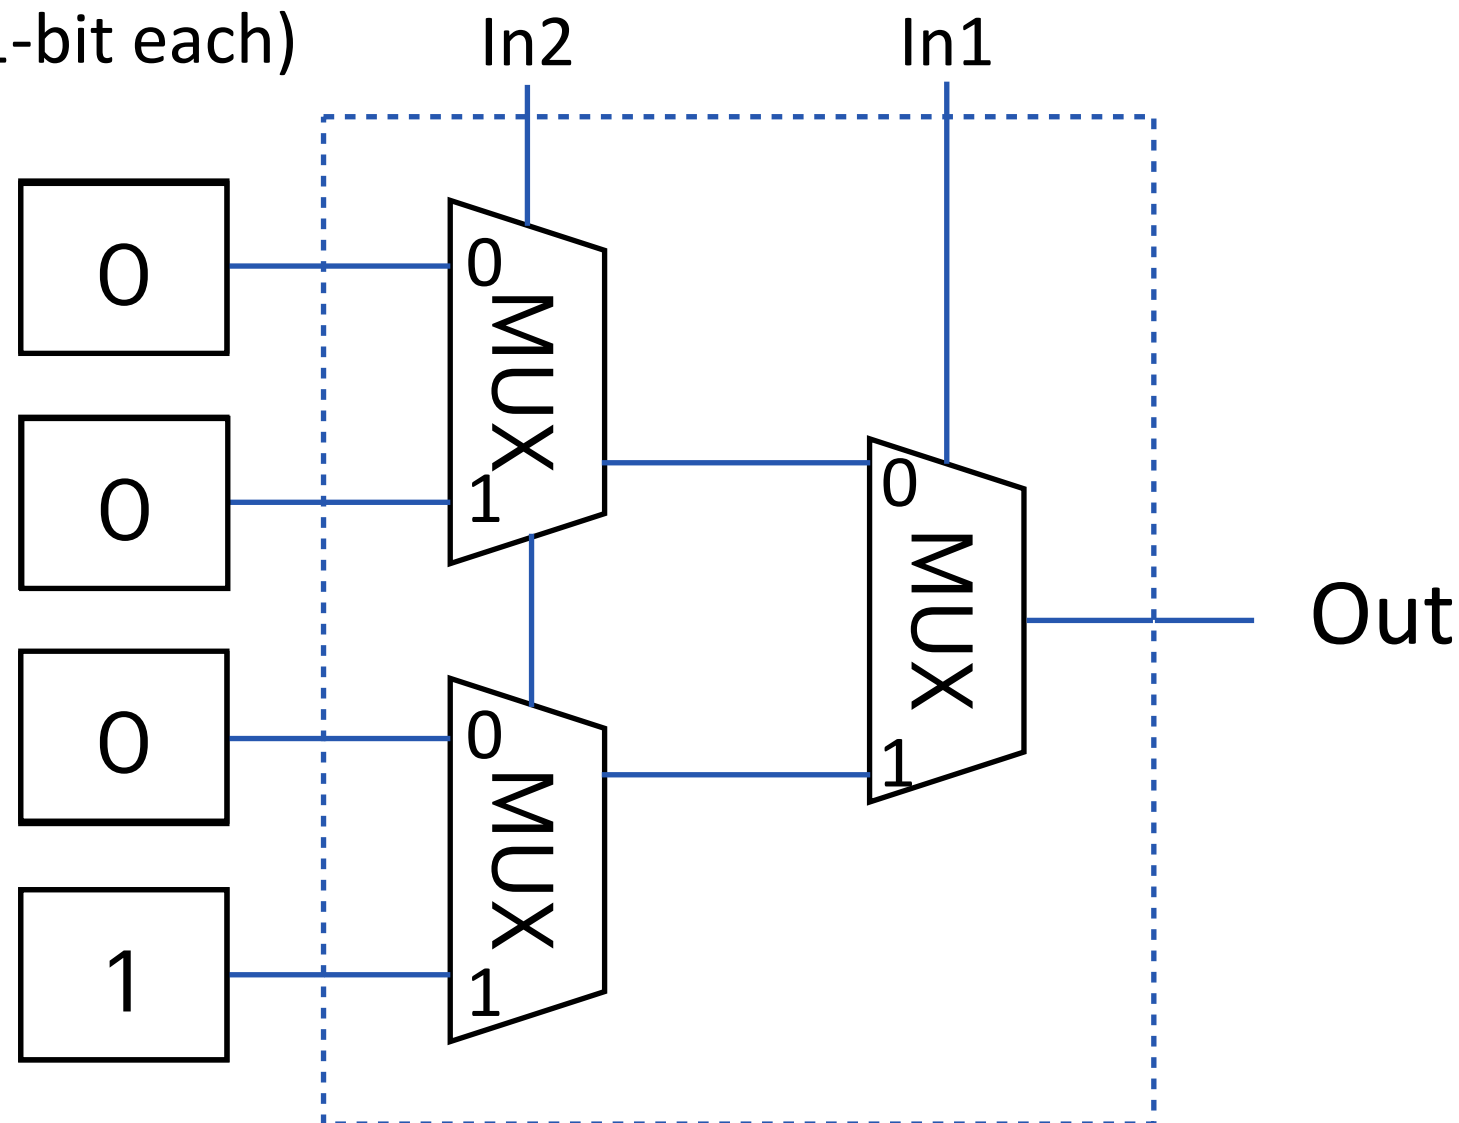
Equivalent to



SRAM-based FPGA


- Example: 2-input LUT implementing AND logic

SRAM cells
(1-bit each)




In1	In2	Out
0	0	0
0	1	0
1	0	0
1	1	1

Equivalent to



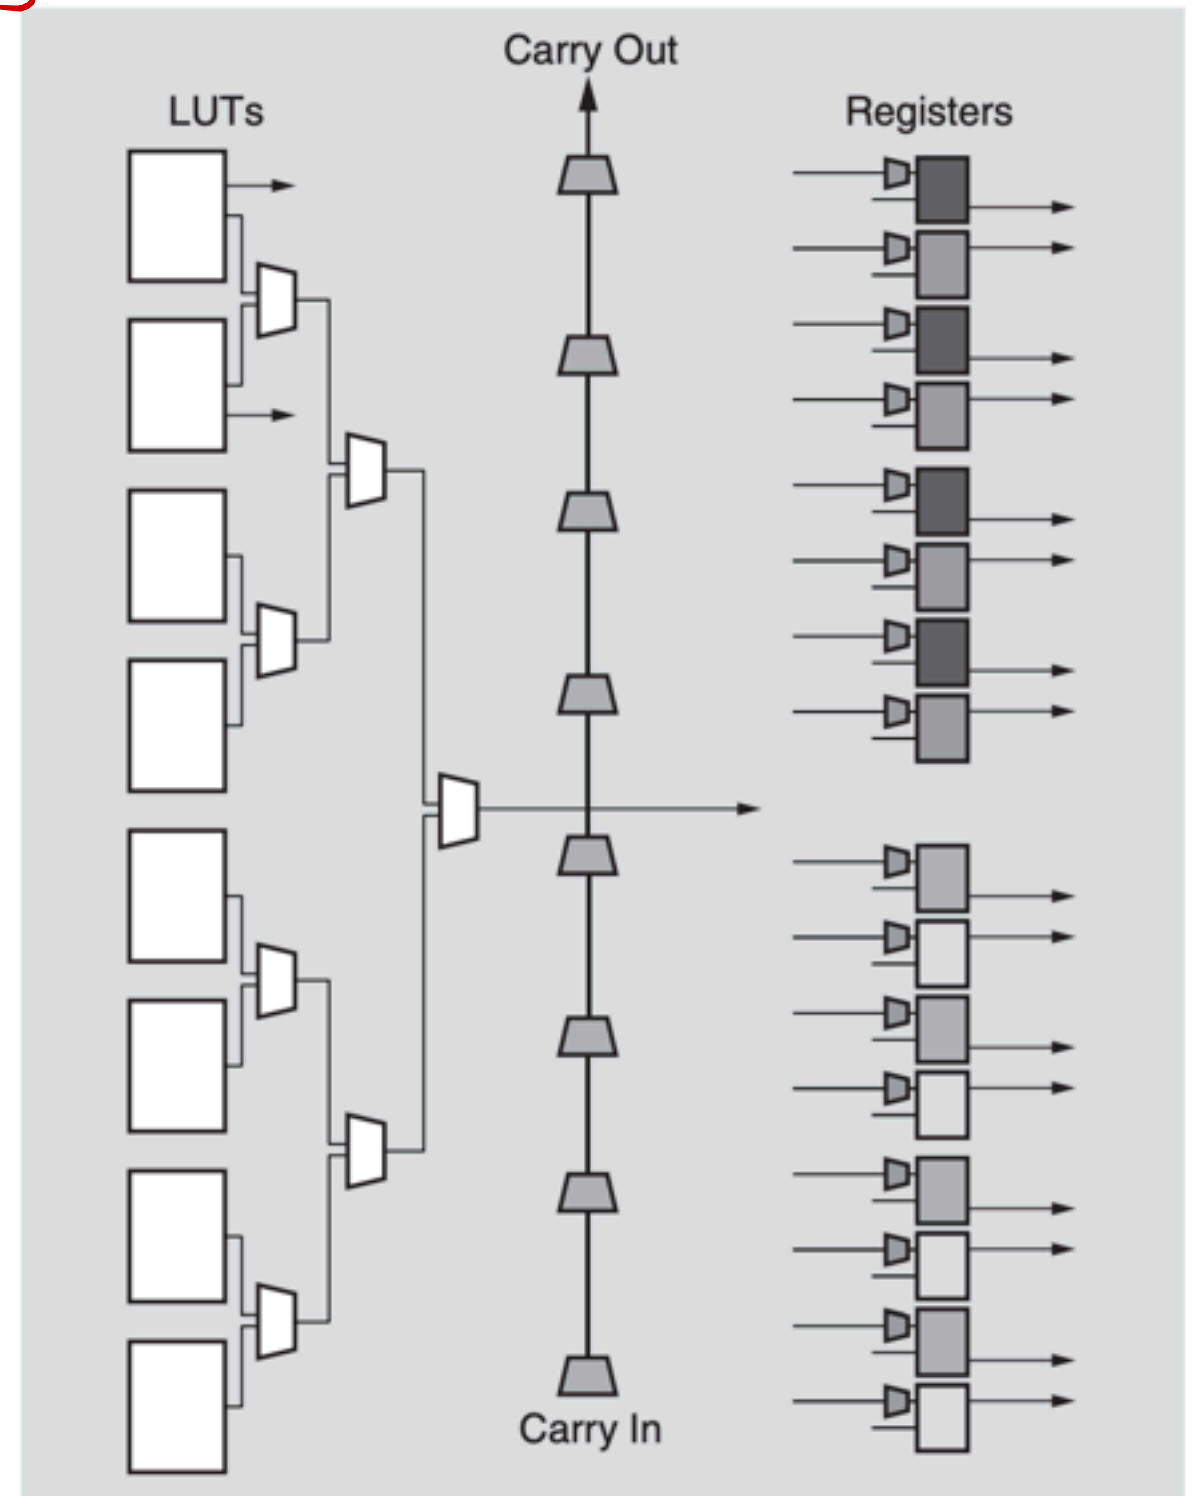
SRAM-based FPGA

- Reality: LUT with larger number of inputs are more capable



6-input LUT
with bypass path & registers

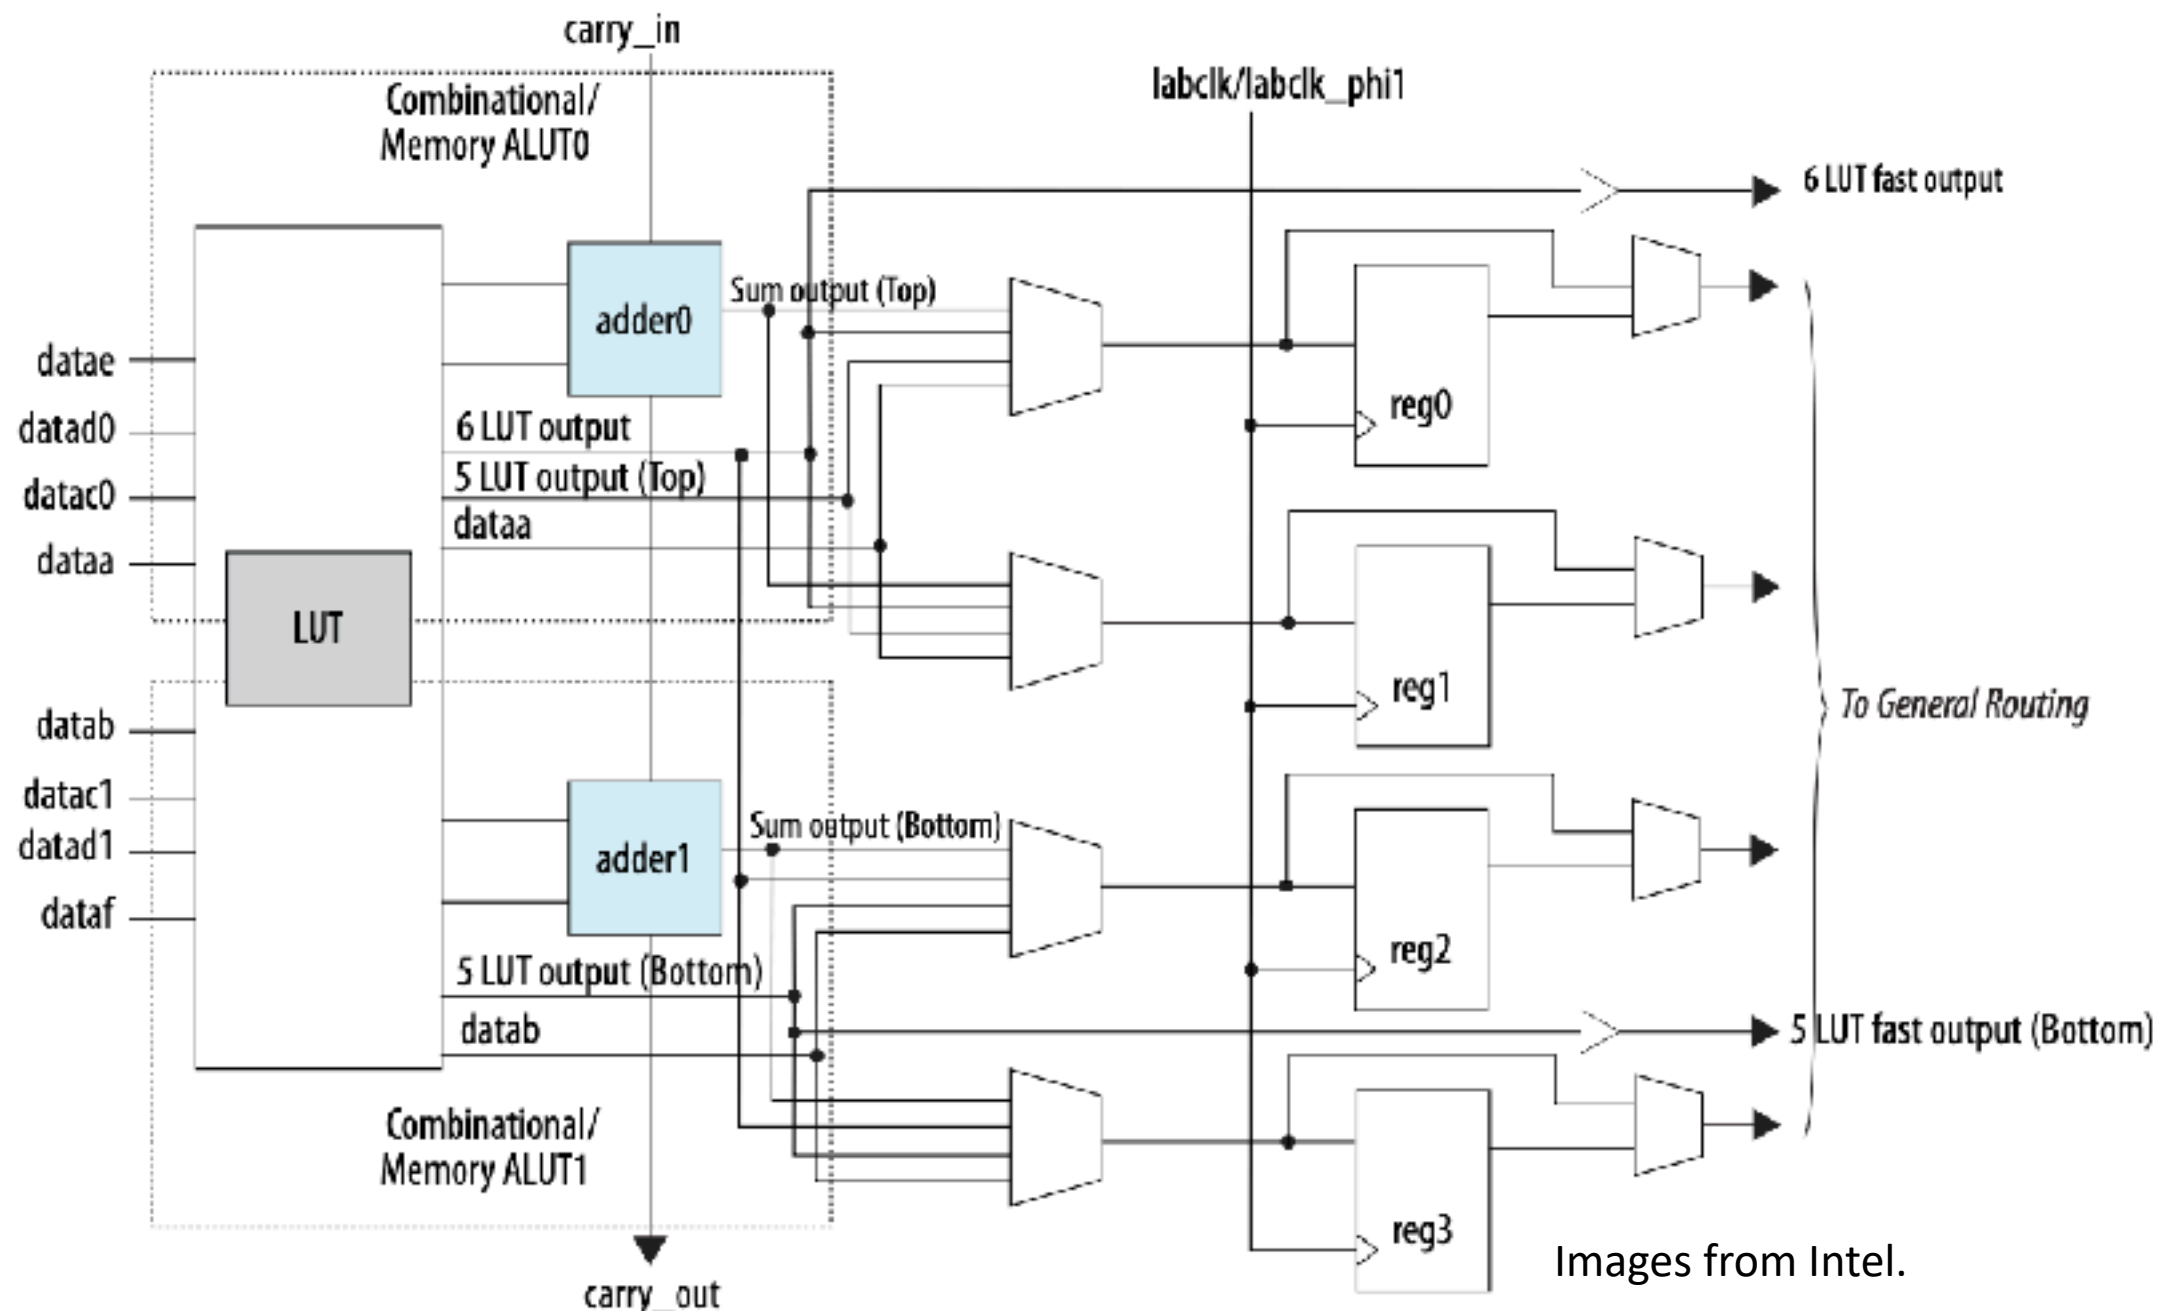
CLB
One configurable logic block
(CLB in Xilinx/AMD FPGA)
consists of many LUTs, registers
and carry chain (for arithmetic)



SRAM-based FPGA


- Reality: LUT with larger number of inputs are more capable

Intel Agilex 7 ALM High-Level Block Diagram (ALM in Altera/Intel FPGA)



Intel Agilex 7 LAB Structure and Interconnects Overview

This figure shows an overview of the Intel Agilex 7 LAB and MLAB structure with the LAB interconnects.



Images from Intel.

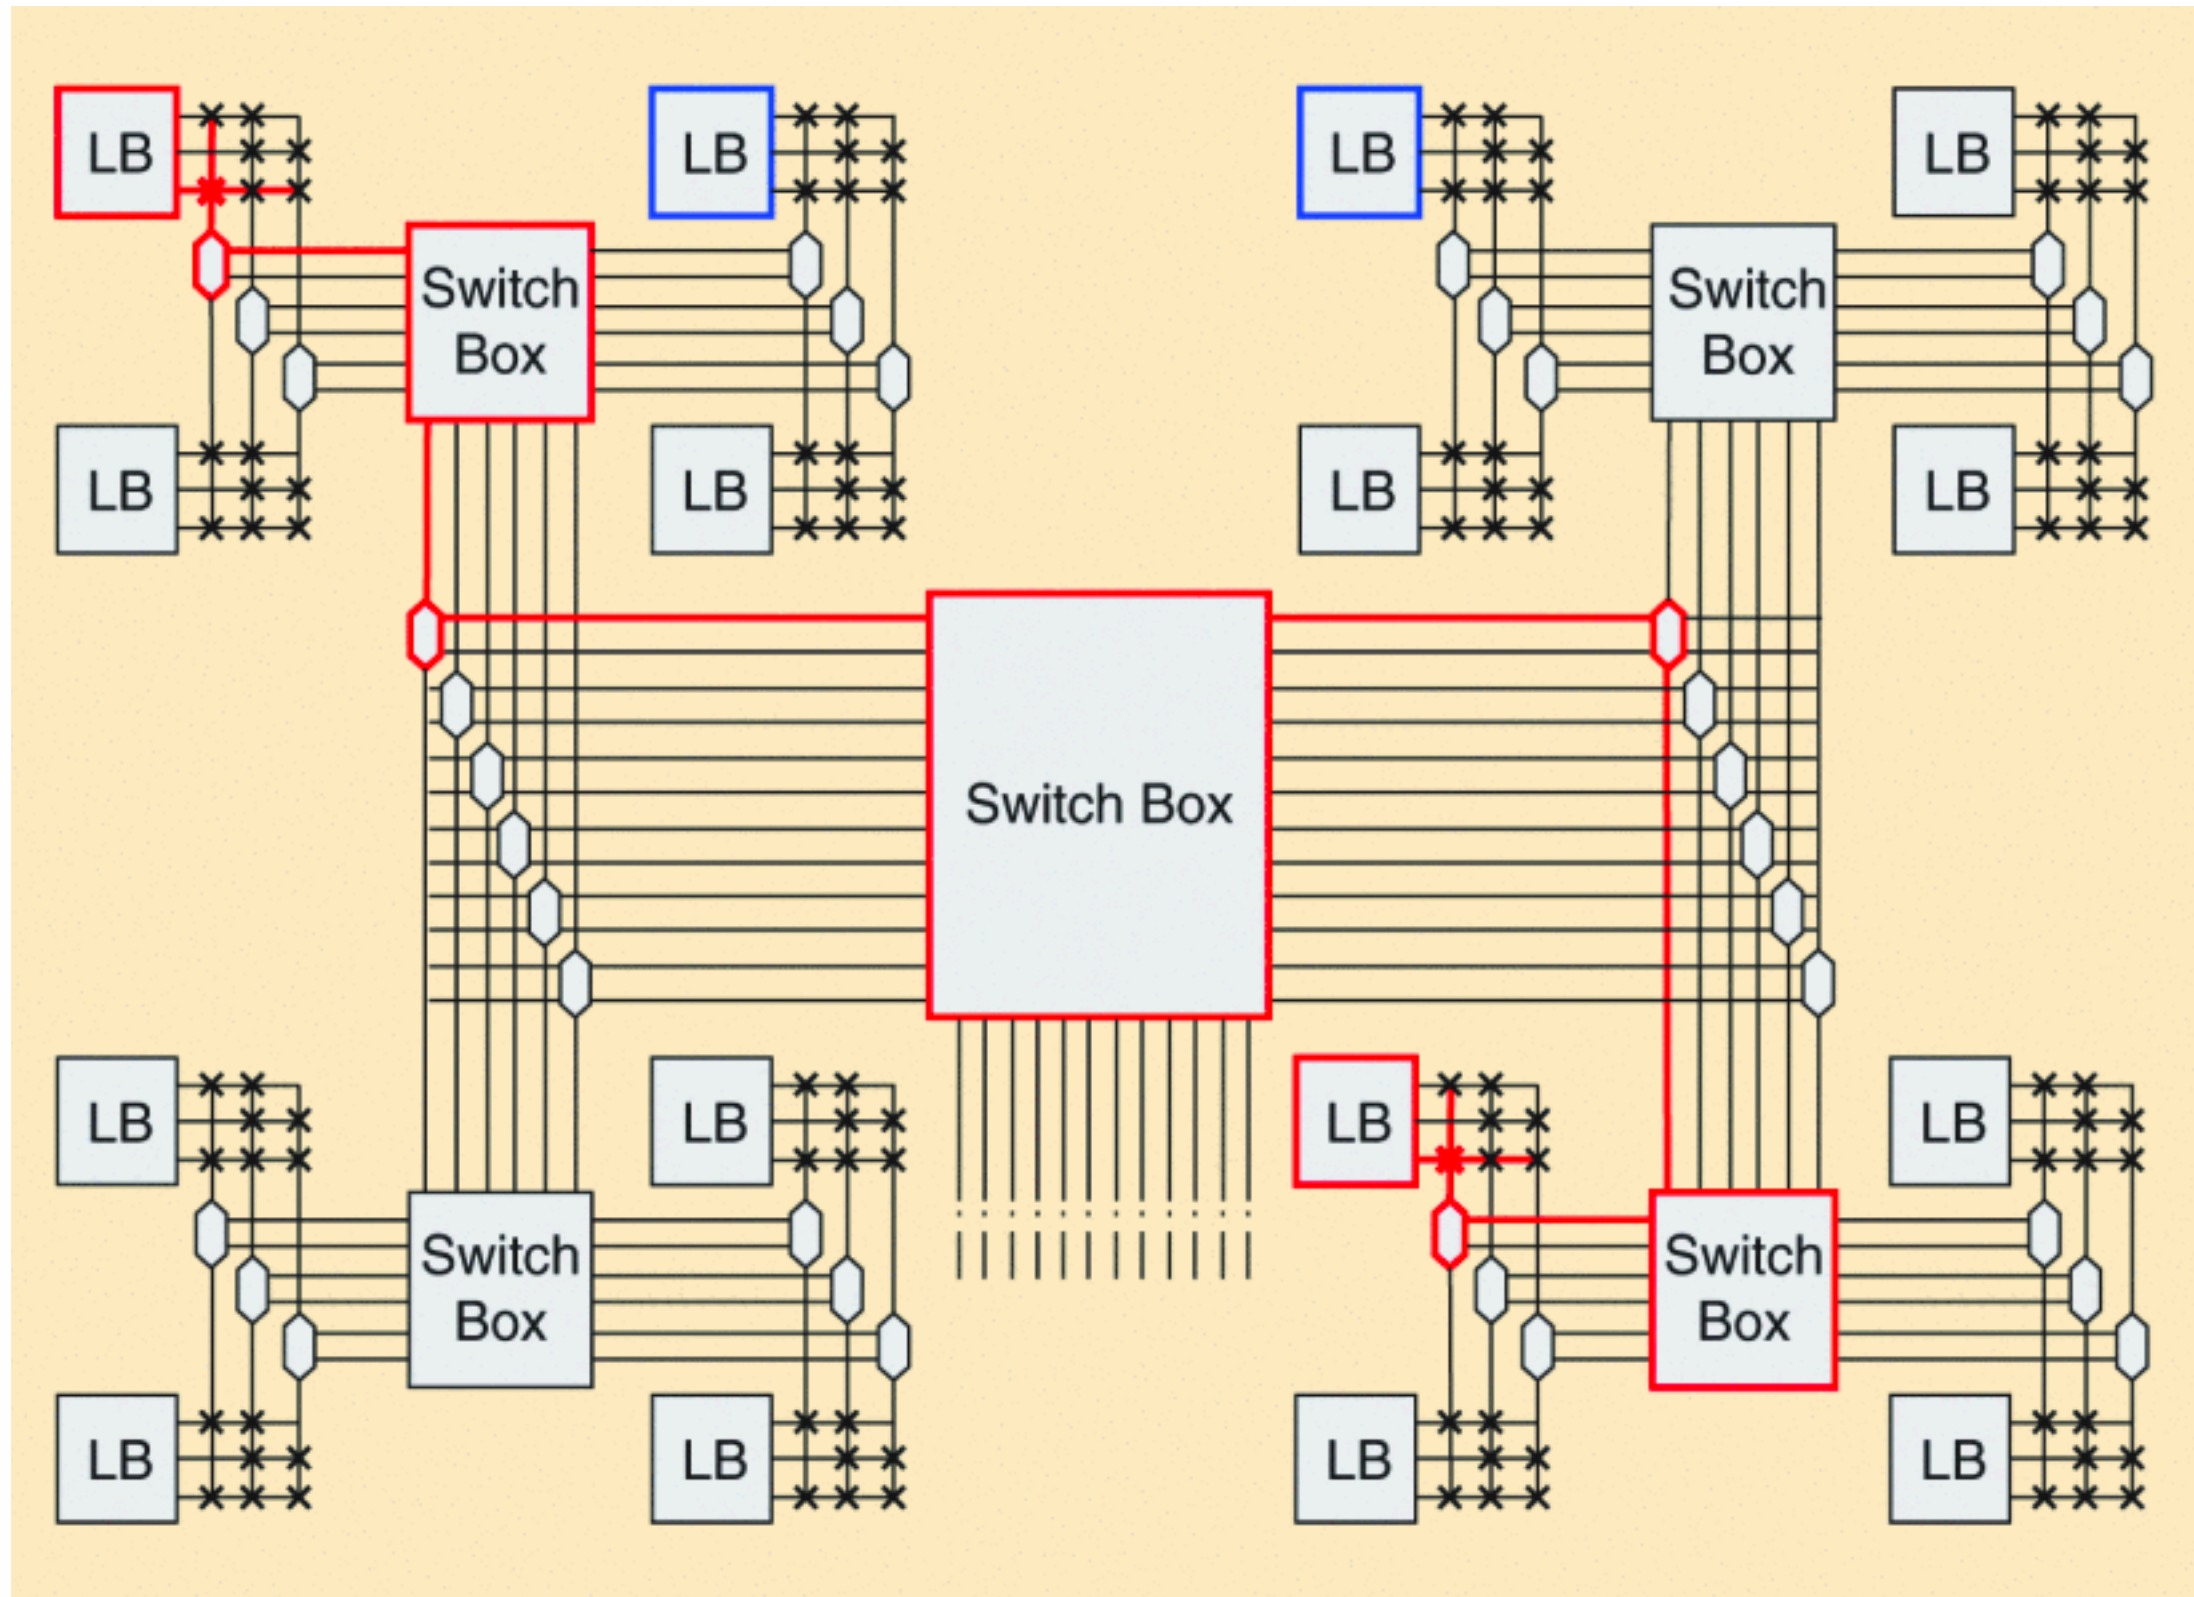
Other than CLB/LAB

- Routing, aka, interconnecting
 - Through programmable wires and switches
 - Between logic blocks (CLB/ALMs), and between I/O blocks and logic blocks
- Routing is a challenging problem
 - Routing technique used in an FPGA largely decides the amount of area used by wire segments and programmable switches, as compared to area consumed by functional blocks.
 - Inferior routing may lead to congestion or failure of signals.


Other than CLB/LAB

- Routing, aka, interconnecting
 - Through programmable wires and switches
 - Between logic blocks (CLB/ALMs), and between I/O blocks and logic blocks
- Routing is a challenging problem
 - Routing technique used in an FPGA largely decides the amount of area used by wire segments and programmable switches, as compared to area consumed by functional blocks.
 - Inferior routing may lead to congestion or failure of signals.
- Different FPGA routing architecture
 - Hierarchical FPGA
 - Island-style routing architecture


Hierarchical FPGA



Island-style FPGA




FPGA Routing Wires



- Some FPGAs contain routing architectures that include different lengths of wires.
- The length of a wire is the number of functional blocks it spans.
- Long wires introduce shorter delays for long interconnections since fewer switch blocks will be passed.

Programmable Switches

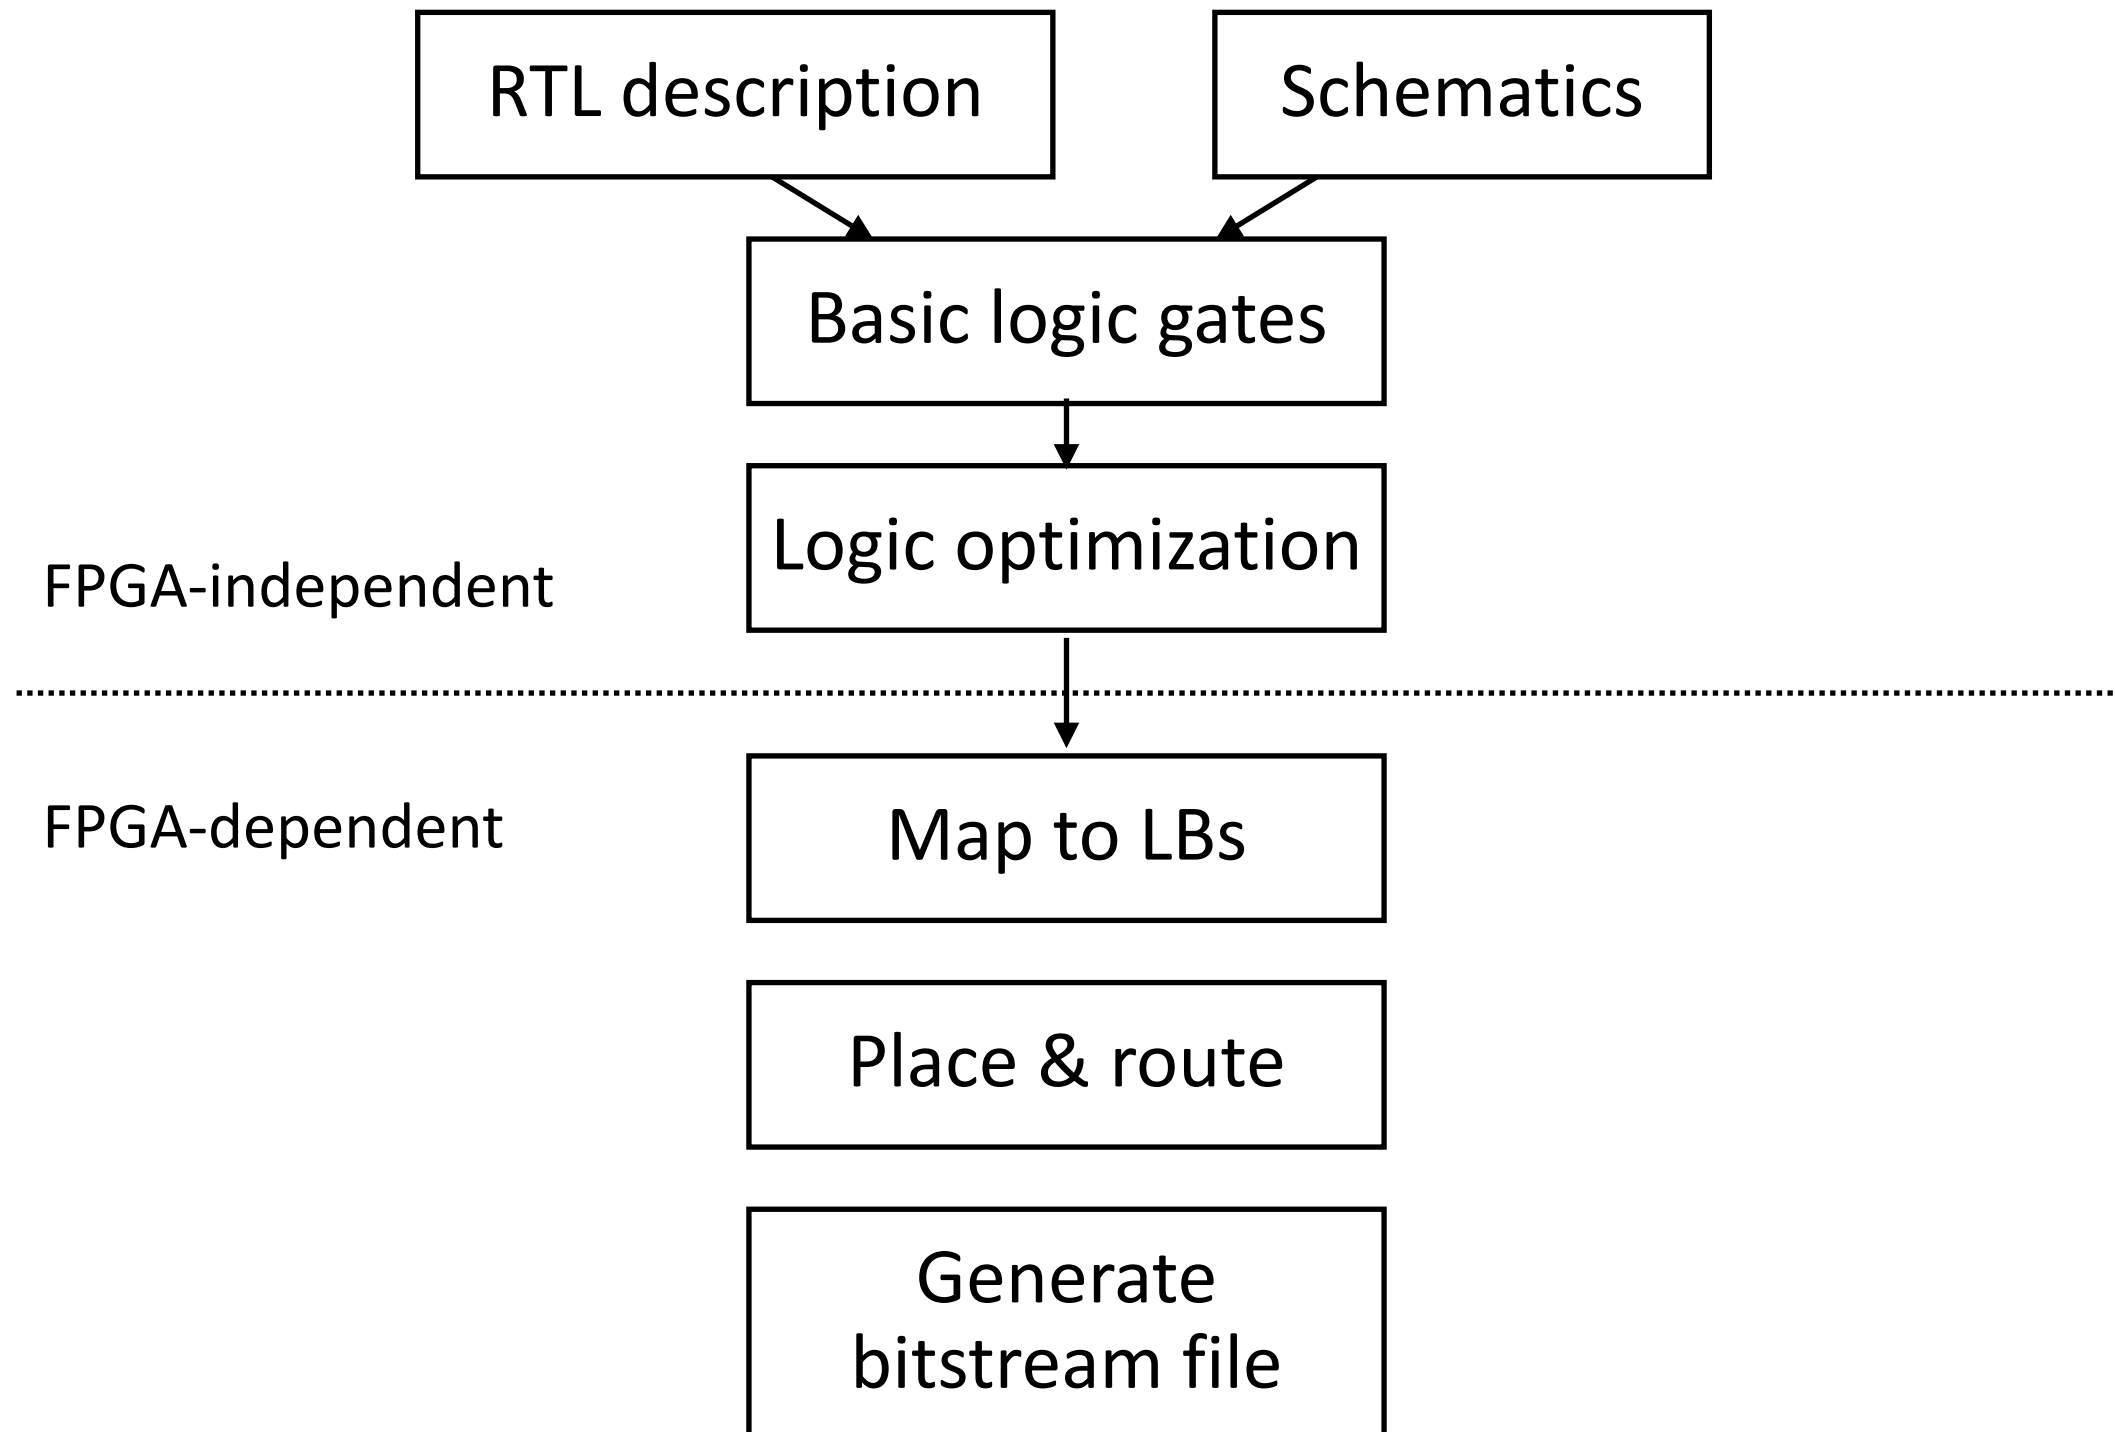
High resistance "Z"






Software Side: FPGA Design Flow

Register-transfer level (RTL)
model: describe the dataflow

A visual representation of the
design (such as Logisim)

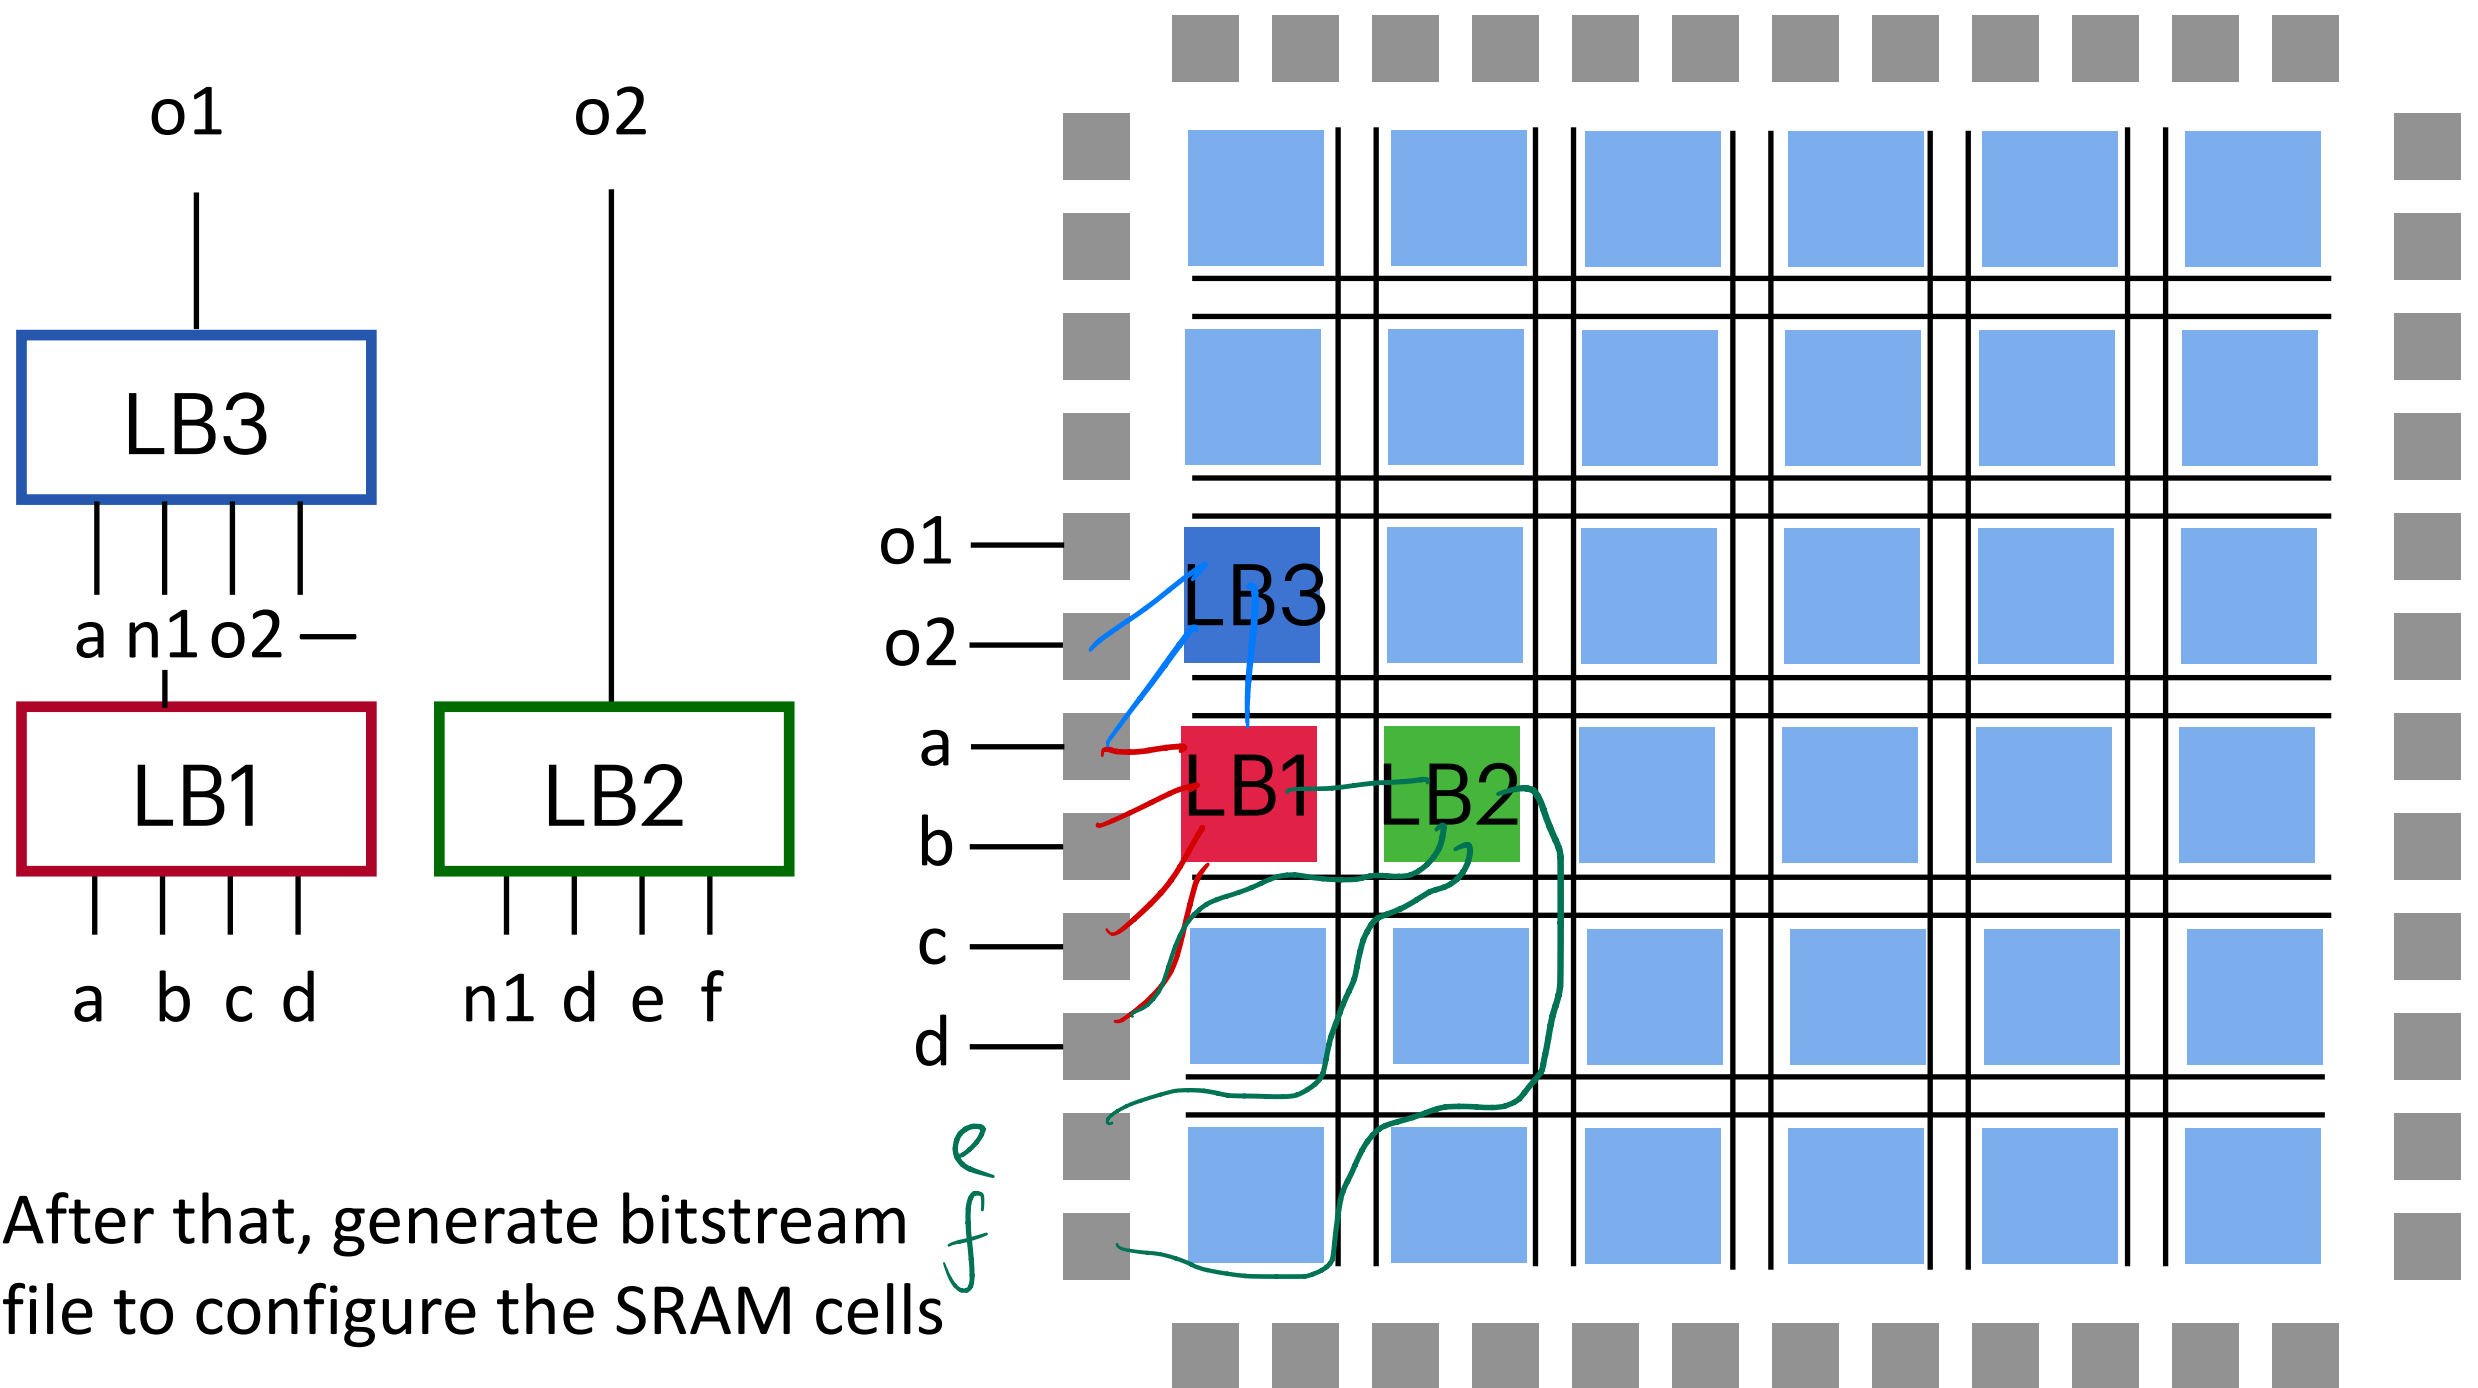


FPGA Mapping




4输入

FPGA Placement & Routing



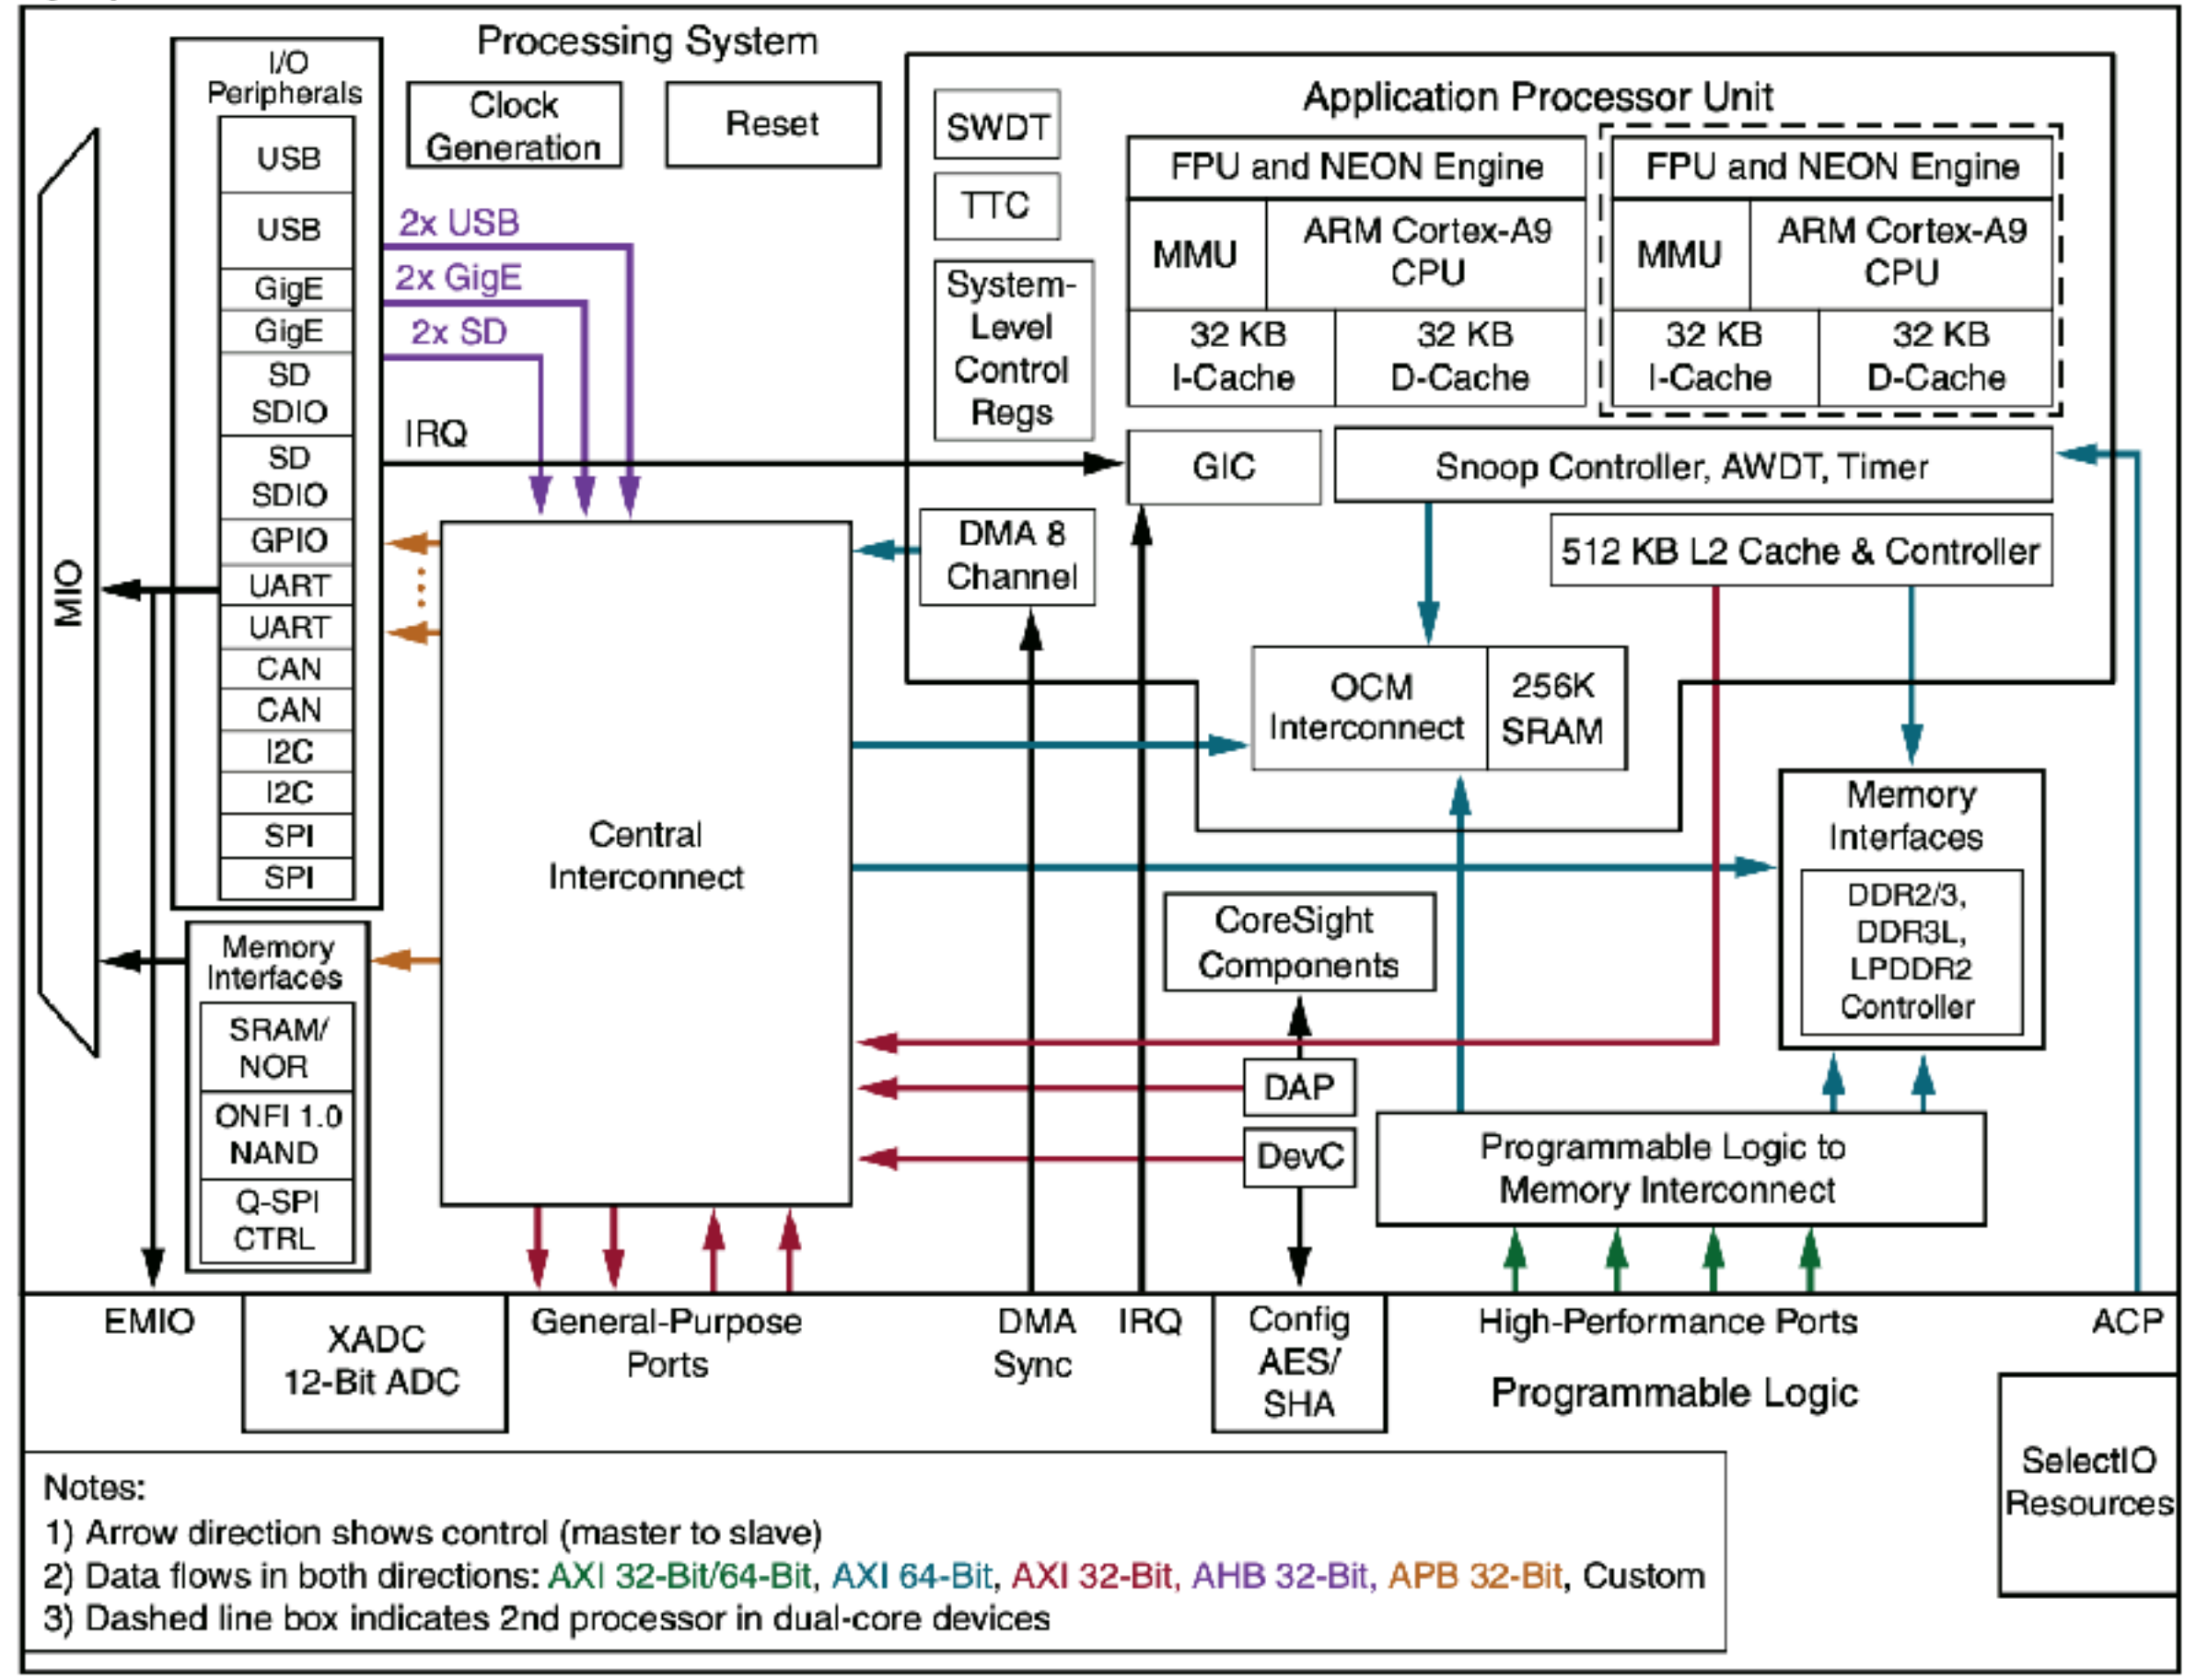
Modern FPGAs

- More like SoC (system-on-chip)
- Logic blocks (functional blocks), DSP slices, block/distributed RAM, I/O and even embedded CPUs (usually ARM core) & GPUs



AMD/Xilinx

Zynq-7000 SoC



Question: True or False

- Given enough resources (LUTs, logic blocks and RAMs), an FPGA can implement a RISC-V CPU (e.g., RV32I).

T

Hardware Description Language

- A way to document the hardware design, which has become IEEE standard (Verilog HDL & VHDL).

```
module alu(opA, opB, aluop, result, zero);  
parameter width=32;  
input [1:0] aluop; input [width-1:0] opA, opB;  
output reg [width-1:0] result; output reg zero;  
always @(*) zero = (result == 0);  
always @(opA, opB, aluop) begin  
case (aluop) 0: result = opA + opB;  
1: result = opA - opB;  
2: result = opA & opB;  
3: result = opA | opB;  
default: result = 0;  
endcase  
end  
endmodule
```

An ALU

Hardware Description Language

- A way to document the hardware design, which has become IEEE standard (Verilog HDL & VHDL).

```
module rf(reg1, reg2, wr, data, reg_wr, d1, d2, clk);  
parameter reg_width=32, num_reg=32;  
input [4:0] reg1, reg2, reg_wr;  
input wr, clk; input [reg_width-1:0] data;  
output reg [reg_width-1:0] d1, d2;  
reg [reg_width-1:0] reg_file [0:num_reg-1];  
  
always @(*)  
begin  
d1 = reg_file[reg1];  
d2 = reg_file[reg2];  
end  
always @(posedge clk)  
begin  
if (wr)  
reg_file[reg_wr] <= data;  
end  
initial $readmemh("rf.txt", reg_file);  
endmodule
```

A Regfile

Hardware Description Language

- A way to document the hardware design, which has become IEEE standard (Verilog HDL & VHDL).
- Can be used both for FPGA design & ASIC design.
- New HDLs like SpinalHDL & Chisel HDL

HDL vs. Software PL

Hardware	Software
Concurrent execution of tasks. This demands all tasks and events to operate in coherence with a timing reference signal called clock	Sequential execution of tasks and instructions. There is no concept of synchronization to clock reference
Very fast execution. Functional timing in nanosecond scale units is achievable in hardware. And therefore, time critical functions are designed to be in hardware	Slow execution. Minimum timing resolution is 100s of microsecond
Can be parallel	Sequential though it can appear to be parallel for the user
Physical and costs are exorbitant if it has to be redone	Can be recompiled
Need to be first time success.	Can be corrected and recompiled without much effort
Hardware can be one time developed as platform and reused for lifetime if the functionality is the same	Can be redone easily.
Development from paper specification to physical system on chip	Need processing hardware platform for sw development
Need to verify fully imagining all scenario ahead of fabrication and hence verification and validation are unavoidable	Verification is necessary to prove the intent of the design but in the case of minor defects, it can be corrected.

Source: <https://link.springer.com/book/10.1007%2F978-3-030-23049-4>

FPGA vs. ASIC

Field-programmable gate array	Application-specific integrated circuit
Fast time-to-market	Low cost when production volumes are high
Reconfigured without costly mask changes	Higher efficiency
Testing done by FPGA vendor	High risk due to high cost of correcting design errors
Cannot exploit 100% the hardware	Greater design/verification/production/test cost (Engineering costs, masks, packaging)
Less protection against design theft	
HDL	Domain-specific language

[1] Bruce F. Cockburn and Jie Han. Chapter 1. Review of Classical Sequential Logic Design. ECE 511 2013.