

CS 110

Computer Architecture

An Introduction to *Operating Systems*

Instructors:
Chundong Wang & Siting Liu

<https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html>

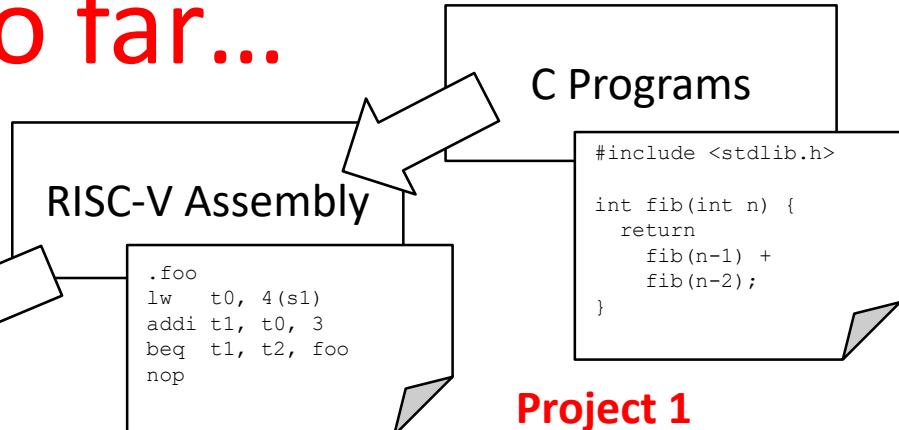
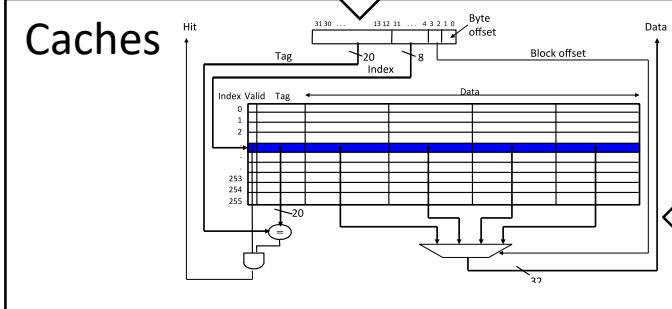
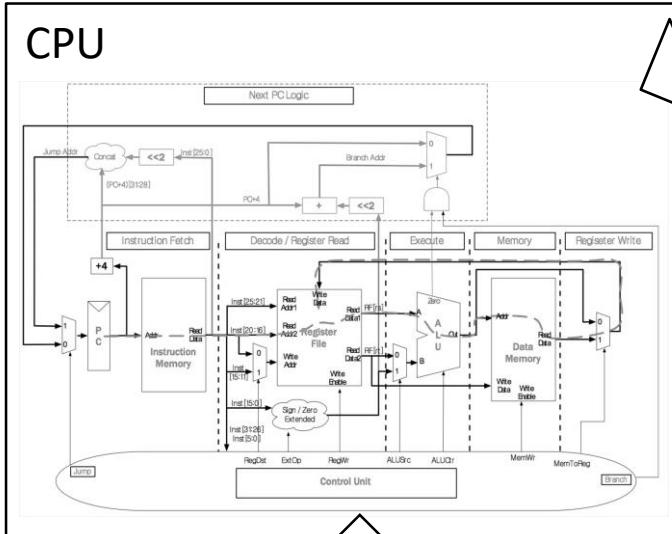
School of Information Science and Technology SIST

ShanghaiTech University

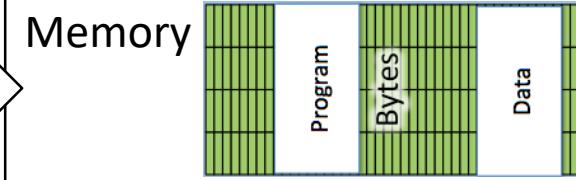
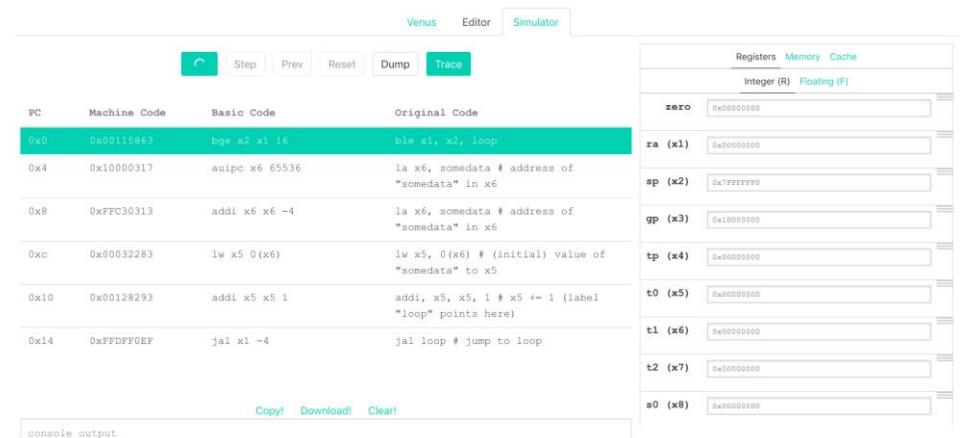
Slides based on UC Berkeley's CS61C

CA so far...

Project 2



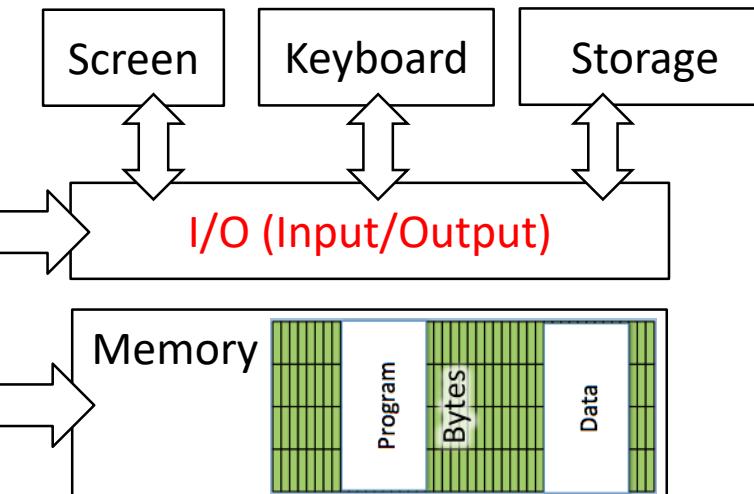
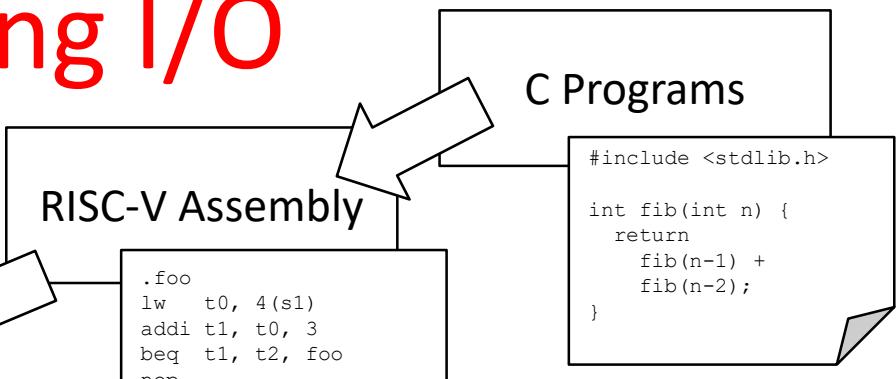
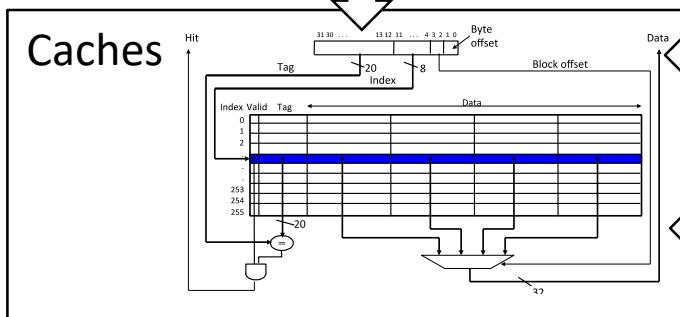
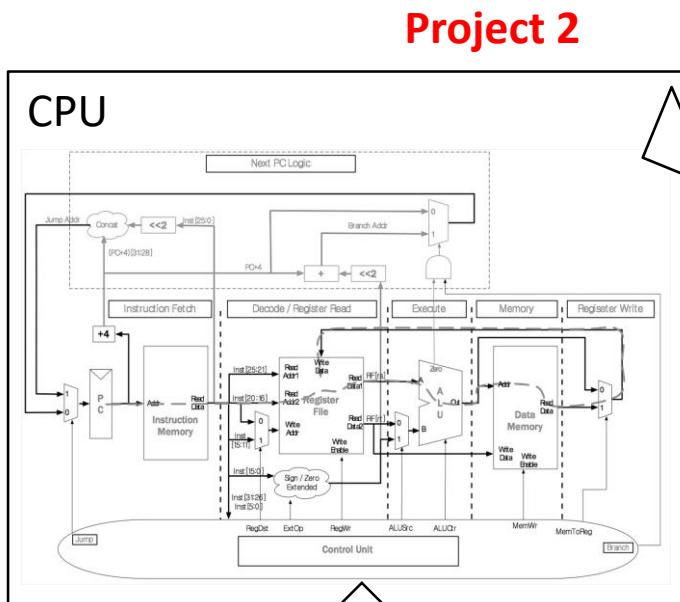
Project 1



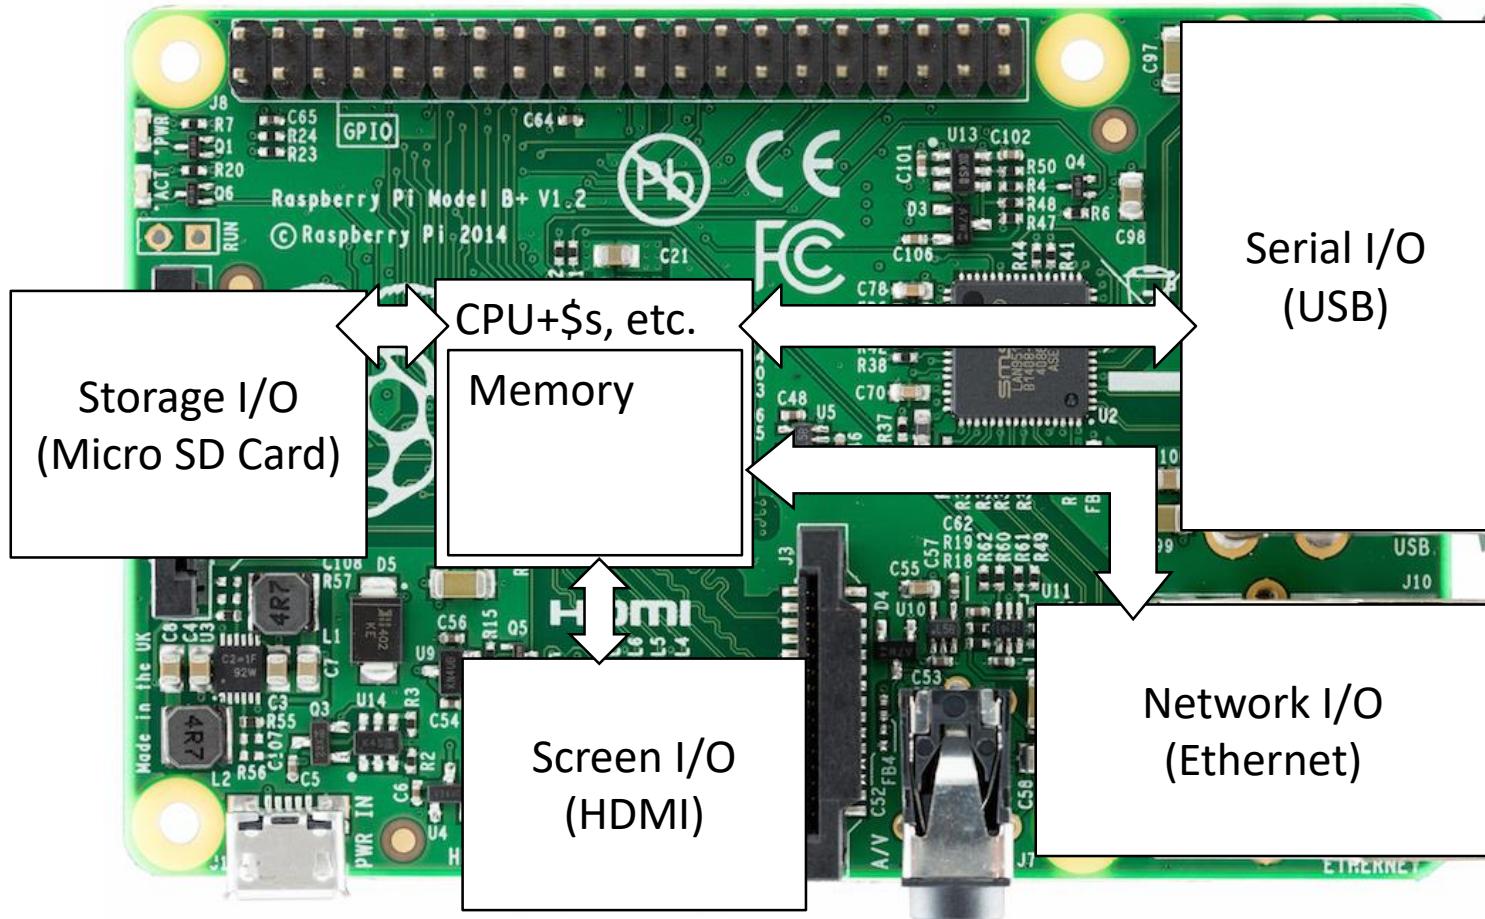
So how is this any different?



Adding I/O



Raspberry Pi ARM 架构

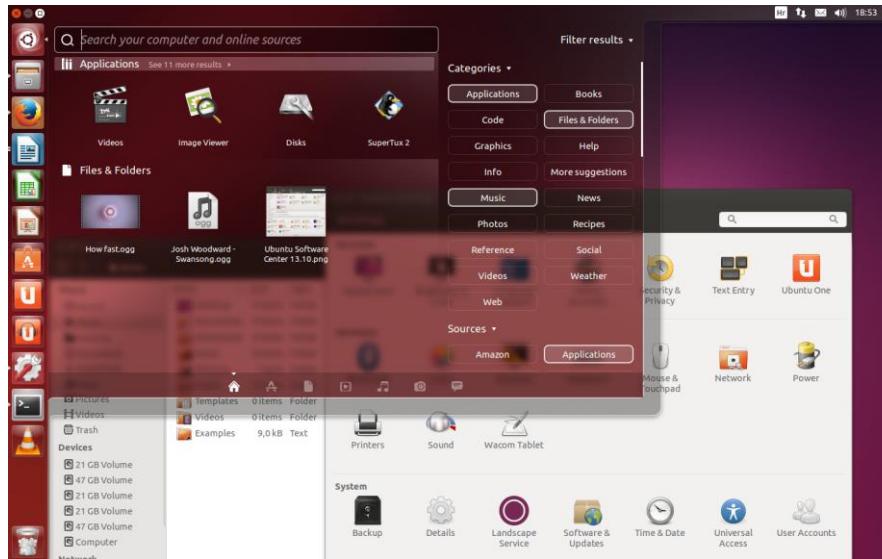


It's a real computer!



But wait...

- That's not the same! Our CS 110 experience isn't like the real world. When we run VENUS, it only executes one program and then stops.
- When I switch on my computer, I get this:



Yes, but that's just software! **The Operating System (OS)**

最基本的基础出软件



Well, “just software”

OS: the biggest piece of software on the machine

- The biggest piece of software on the machine?
- How many lines of code?



Year
1994
1996
2001
2003
2004
2015
Apr 2020
May 2021
May 2022
Apr 2023

Say No to Pirated Products
(拒绝盗版)

Size of zipped file
1MB
6MB
23MB
40MB
92MB
118MB
166MB
179MB
189MB
209MB

纯文本

1.4GB
unzip →

What does the OS do?

- One of the first things that runs when your computer starts (right after firmware/ bootloader)
- Loads, runs and manages programs:
 - Multiple programs at the same time (time-sharing)
 - Isolate programs from each other (isolation)
 - Multiplex resources between applications (e.g., devices)
- Services: File System, Network stack, printer, etc.
- Finds and controls all the devices in the machine in a general way (using “device drivers”)

What does the core of OS need to do?

交互

- Provide **interaction** with the outside world

- Interact with “devices”

- Disk, screen, keyboard, mouse, network, etc.

隔离

- Provide **isolation** between running programs (processes)

- Each program runs in its own little world

- Virtual memory

每个程序都有独立的内存空间
(虚拟内存)

Agenda

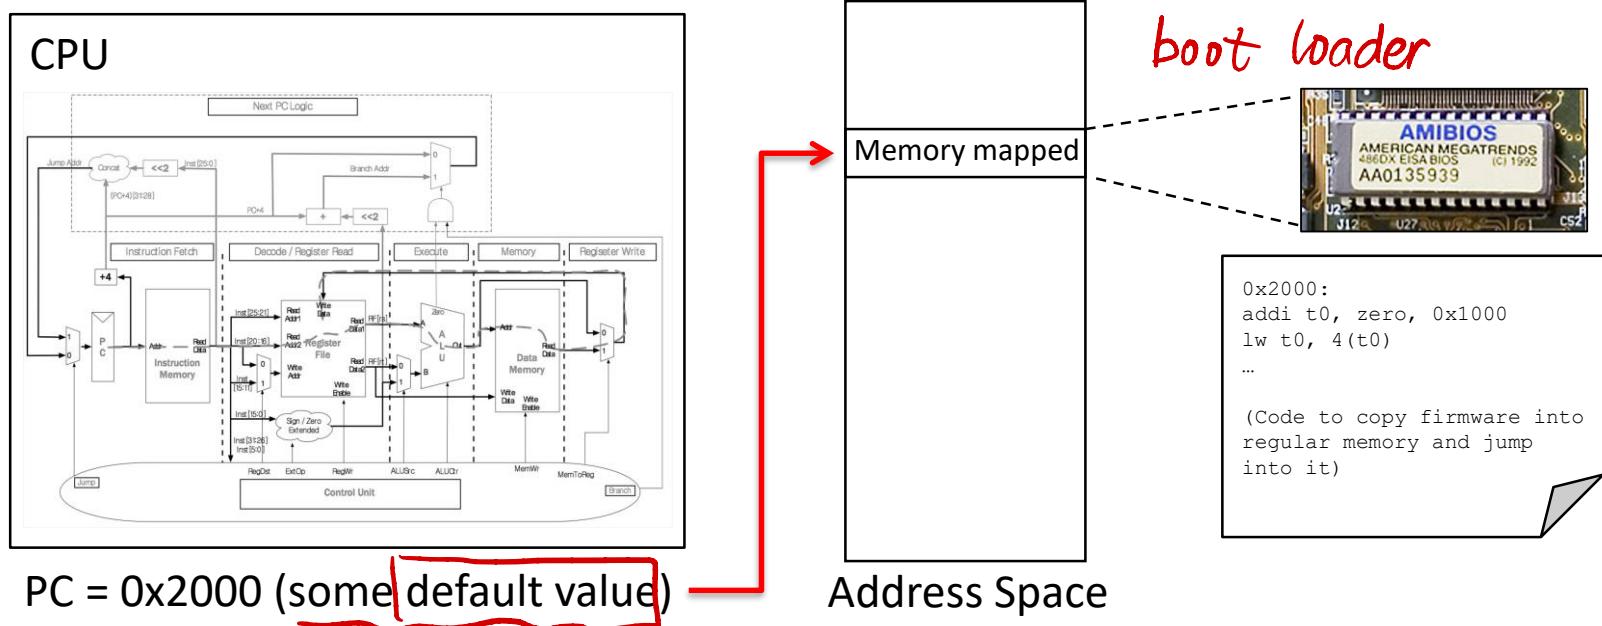
- OS Boot Sequence and Operation
- Devices and I/O, interrupt and traps
- Application, Multiprogramming/time-sharing

Agenda

- OS Boot Sequence and Operation
- Devices and I/O, interrupt and traps
- Application, Multiprogramming/time-sharing

What happens at boot?

- When the computer switches on, it does the same as Venus: the CPU executes instructions from some start address (stored in Flash ROM)



- Bootstrapping:

<https://en.wikipedia.org/wiki/Bootstrapping>

What happens at boot?

- When the computer switches on, it does the same as Venus: the CPU executes instructions from some start address (stored in Flash ROM)

BIOS → bootloader → kernel

1. BIOS: Find a storage device and load first sector (block of data)

Diskette Drive B : None Serial Port(s) : 3FO ZFO										
Pri. Master Disk : IBM-MAT 100, 250GB Parallel Port(s) : 3FO										
Pri. Slave Disk : IBM-MAT 100, 250GB DMA at Bank(s) : 0 1 2										
Sec. Master Disk : None										
Sec. Slave Disk : None										
PCI Devices Listing ...										
Bus	Dev	Fun	Vendor	Device	SUB	SSID	Class	Device Class	IRQ	
0	27	0	8006	266B	1450	0005	0003	Multimedia Device	5	
0	29	1	8006	265B	1450	265B	0003	USB 1.1 Host Contrlr		
0	29	1	8006	265B	1450	265B	0003	USB 1.1 Host Contrlr		
0	29	2	8006	265B	1450	265B	0003	USB 1.1 Host Contrlr		
0	29	3	8006	265B	1450	265B	0003	USB 1.1 Host Contrlr		
0	29	7	8006	265C	1450	265C	0003	USB 1.1 Host Contrlr		
0	31	2	0001	2651	1450	2651	1001	IDE Contrlr		
0	31	3	8006	266B	1450	266B	0003	Stbus Contrlr	11	
1	0	0	10DE	0421	10DE	0479	0300	Display Contrlr	5	
2	0	0	1203	0212	1000	0000	0180	Mass Storage Contrlr	10	
2	5	0	11AB	432B	1450	E000	0200	Network Contrlr	12	
								ACPI Controller	9	

2. Bootloader (stored on, e.g., disk): Load the OS kernel from disk into a location in memory and jump into it.

跳转进OS

```
QUESTION 3:  
conv: <speedup> x  
relu: <speedup> x  
pool: <speedup> x  
fc: <speedup> x  
softmax: <speedup> x  
  
Which layer should we optimize?  
which layers?  
  
c3:04:03 Wed Apr 15 2015 ~src/proj3/proj3_starter $ ls src/  
answers.txt cnn.cml  
cnn.py data LICENSE Makefile test web  
  
c3:04:03 Wed Apr 15 2015 cs61c-ti@hive22 Linux x86_64  
~/src/proj3/proj3_starter $ ls src/  
cnn.c main.c python.c util.c  
  
c3:04:16 Wed Apr 15 2015 cs61c-ti@hive22 Linux x86_64  
~/src/proj3/proj3_starter $ make cnn  
make: 'cnn' is up to date.  
  
c3:04:20 Wed Apr 15 2015 cs61c-ti@hive22 Linux x86_64  
~/src/proj3/proj3_starter $
```



4. Init: Launch an application that waits for input in loop (e.g., Terminal/Desktop/...)

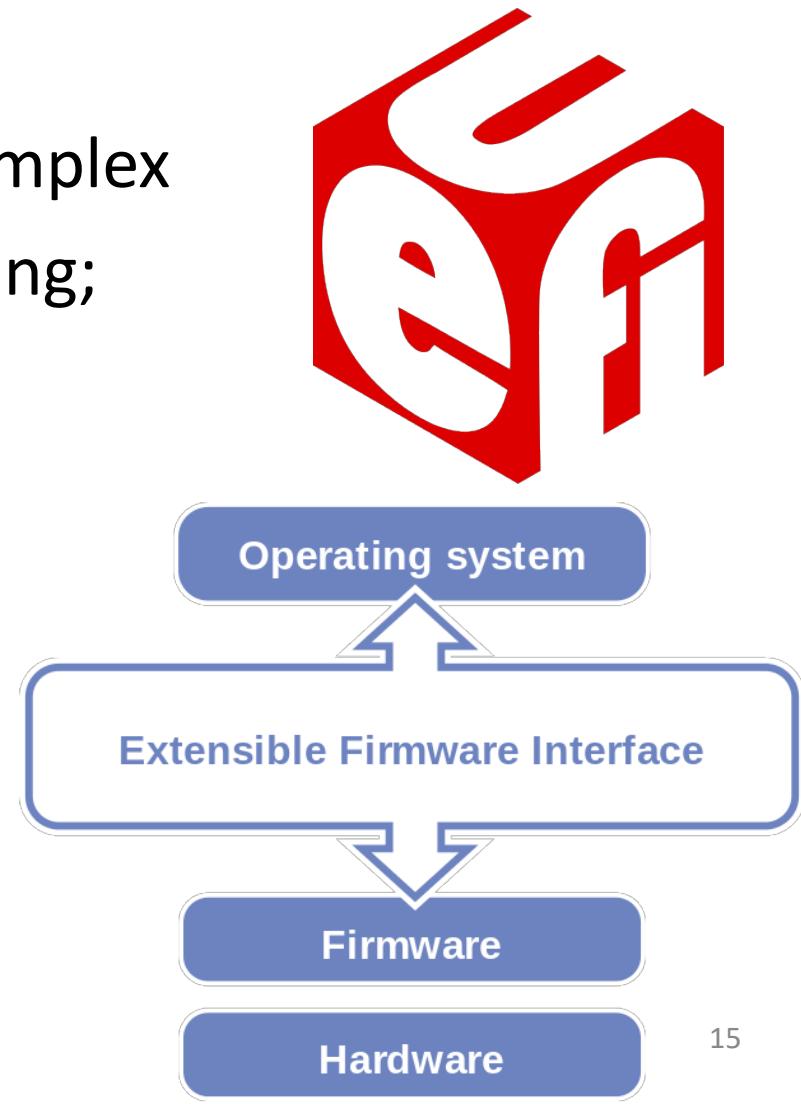
```
Welcome to the KNOPPIX live GNU/Linux on DVD!  
  
Running Linux Kernel 2.6.24.4  
Memory available: 124132KB, Memory free: 118180KB.  
For USB-Firewire devices... Done.  
DMA acceleration for: hdc [QEMU CD-ROM]  
Found primary KNOPPIX compressed image at /cdrom/KNOPPIX/KNOPPIX.  
Creating /randisk (atomic size=99304K) on shared memory...Done.  
Mounting initial filesystem and splitfs on /randisk...  
=> Read-only DVD system successfully merged with read-write /randisk.  
Done.  
Starting INIT (process 1).  
INIT: version 2.86 booting  
Configuring for Linux Kernel 2.6.24.4.  
Processor 0 is Pentium II (Klamath) 1662MHz, 128 KB Cache  
apmd1600BI apmd12.1.1 interfacing with apm driver 1.16ac and APM BIOS 1.2  
Processor 0 found power management functions enabled.  
USB Friend, managed by apm  
sure were found, managed by udev  
Starting udev hot-plug hardware detection... Started.  
toconfiguring devices...
```

3. OS Boot: Initialize services, drivers, etc.

UEFI

Unified Extensible Firmware Interface

- Successor of BIOS
- Much more powerful and complex
- E.g. graphics menu; networking; browsers
- All modern Intel & AMD based computer use UEFI

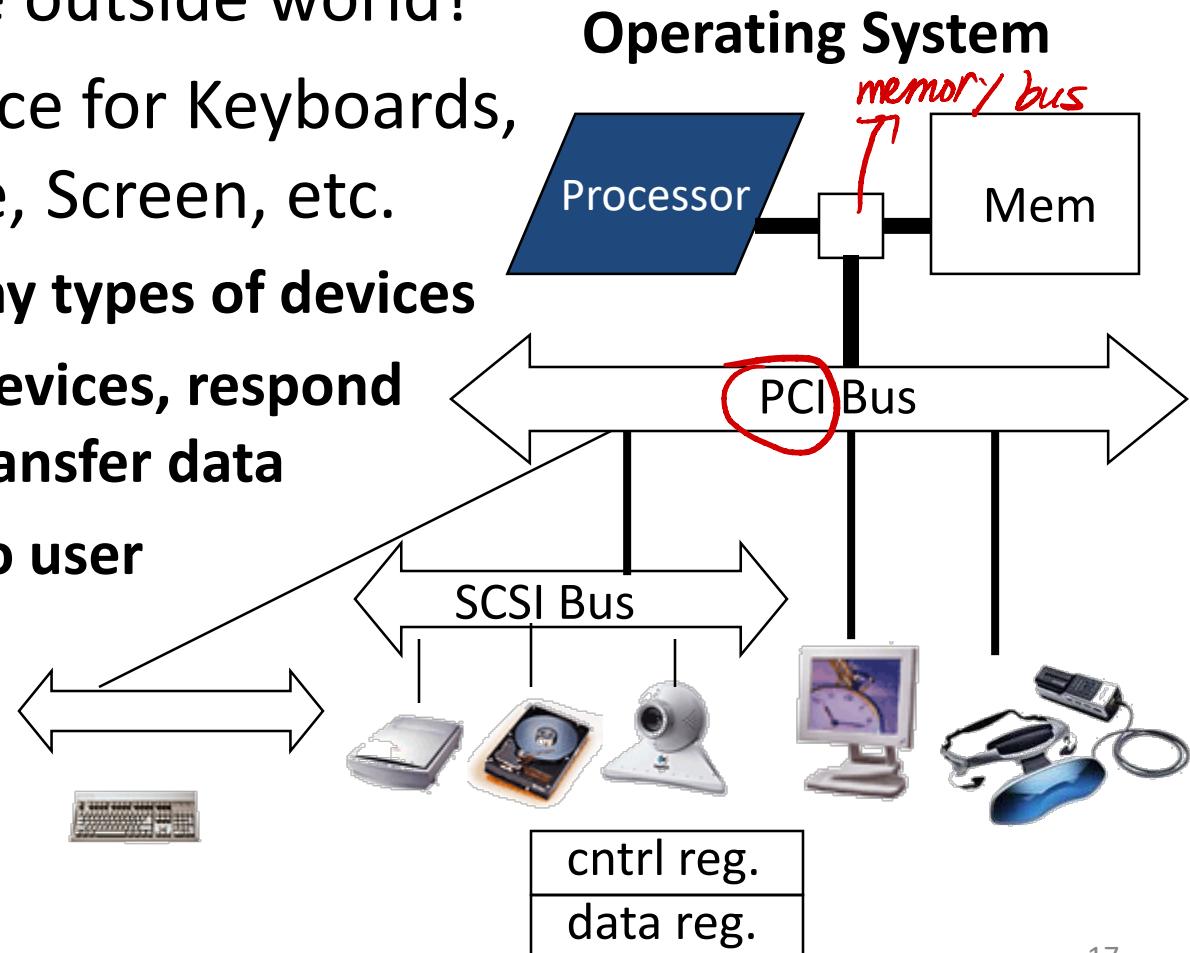


Agenda

- OS Boot Sequence and Operation
- Devices and I/O, interrupt and traps
- Application, Multiprogramming/time-sharing

How to interact with devices?

- Assume a program running on a CPU. How does it interact with the outside world?
- Need I/O interface for Keyboards, Network, Mouse, Screen, etc.
 - Connect to many types of devices
 - Control these devices, respond to them, and transfer data
 - Present them to user programs so they are useful

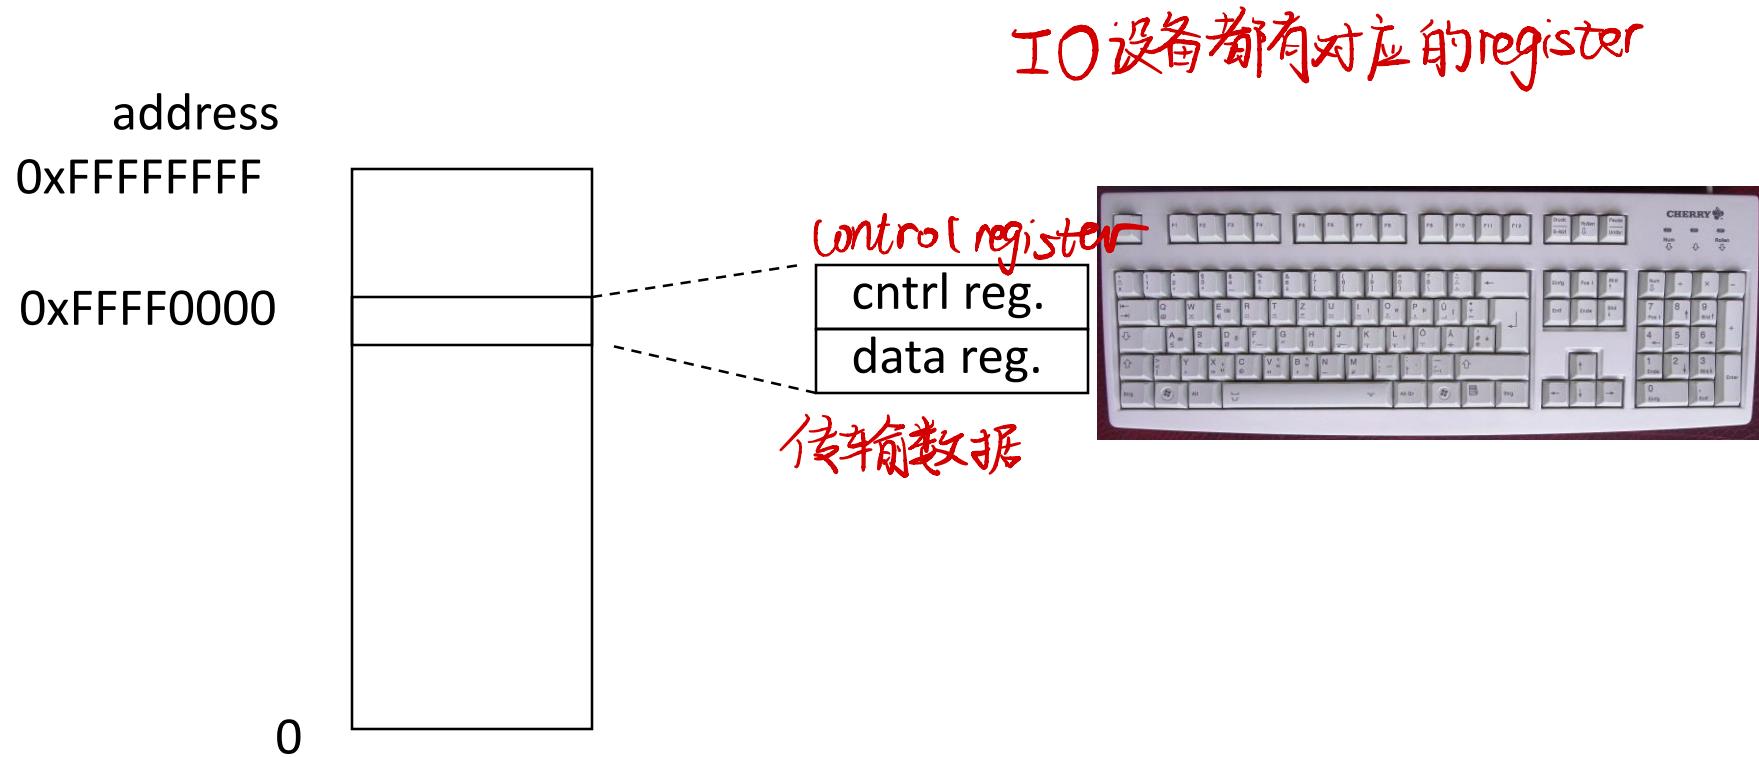


Instruction Set Architecture for I/O

- What must the processor do for I/O?
 - Input: reads a sequence of bytes
 - Output: writes a sequence of bytes
- Interface options
 - Some processors have special input/output instructions
 - **Memory Mapped Input/Output** (used by RISC-V):
 - Use normal load/store instructions, e.g., lw/sw, for input/output
 - In small pieces
 - A portion of the address space dedicated to IO
 - I/O device registers there (no memory there)

Memory Mapped I/O

- Certain addresses are not regular memory
- Instead, they correspond to registers in I/O devices



Processor-I/O Speed Mismatch

- 1GHz microprocessor can execute 1 billion load or store instructions per second, or $4,000,000 \text{ KB/s}$ data rate
 - I/O data rates range from 0.01 KB/s to $1,250,000 \text{ KB/s}$
- Input: device may not be ready to send data as fast as the processor loads it
 - Also, might be waiting for human to act
- Output: device not be ready to accept data as fast as processor stores it
- What to do?

每指令字节 | word $\Rightarrow 4 \text{ Byte}$

10^9

4×10^9

10^3

$I/O < \text{load/store}$

Processor Checks Status before Acting

- Path to a device generally has **2 registers**:
 - **Control Register**, says it's OK to read/write (I/O ready) [think of a flagman on a road] (*flag*)
 - **Data Register**, contains data
- Processor reads from Control Register in loop, waiting for device to set **Ready** bit in Control reg ($0 \Rightarrow 1$) to say it's OK
- Processor then loads from (input) or writes to (output) data register
 - Load from or Store into Data Register resets Ready bit ($1 \Rightarrow 0$) of Control Register
- This is called “**Polling**”

input/output via CPU的视角

I/O Example (polling)

- Input: Read from keyboard into a0

Waitloop:

```
    li      t0, 0xffff0000 #ffff0000
    lw      t1, 0(t0)      #control
    & andi  t1, t1,0x1
    beq   t1, zero, Waitloop
    lw    a0, 4(t0)      #data
```

→ control register

- Output: Write to display from a0

Waitloop:

```
    li      t0, 0xffff0000 #ffff0000
    lw      t1, 8(t0) ?    #control
    & andi  t1, t1,0x1
    beq   t1, zero, Waitloop
    sw    a0, 12(t0)     #data
```

“Ready” bit is from processor’s point of view!

Cost of Polling?

- Assume for a processor with a 1GHz clock it takes 400 clock cycles for a polling operation (call polling routine, accessing the device, and returning).
Determine % of processor time for polling
 - Mouse: polled 30 times/sec so as not to miss user movement

% Processor time to poll

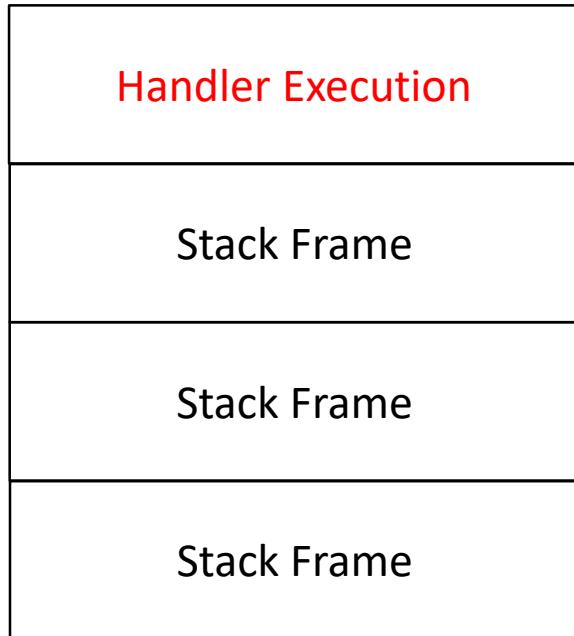
- Mouse Polling [clocks/sec]
= 30 [polls/s] * 400 [clocks/poll] = 12K [clocks/s]
- % Processor for polling:
 $12K \cdot (12 \cdot 10^3) [\text{clocks/s}] / (1 \cdot 10^9) [\text{clocks/s}] = 0.0012\%$
=> Polling mouse **little** impact on processor

What is the alternative to polling?

磁盘 \leftrightarrow 圆圈

- Wasteful to have processor spend most of its time “spin-waiting” for I/O to be ready
- Would like an unplanned procedure call that would be invoked only when I/O device is ready
- Solution: use **exception mechanism** to help I/O.
 - **Interrupt** program when I/O ready, return when done with data transfer
- Allow to register (post) **interrupt handlers**: functions that are called when an interrupt is triggered

Interrupt-driven I/O



1. Incoming interrupt suspends instruction stream
2. Looks up the vector (function address) of a handler in **an interrupt vector table** stored within the CPU
3. Perform a jal to the handler (**needs to store any state**)
4. Handler run on current stack and returns on finish
(thread doesn't notice that a handler was run)

```
handler: li t0, 0xfffff0000
          lw t1, 0(t0)
          andi t1, t1, 0x1
          lw a0, 4(t0)
          sw t1, 8(t0)
          ret
```

Label: sll t1,s3,2
addu t1,t1,s5
lw t1,0(t1)
add s1,s1,t1
addu s3,s3,s4
bne s3,s2,abel

处理完中断，回到原处

Interrupt(SPI0)

CPU Interrupt Table		
SPI0	handler	
...	...	

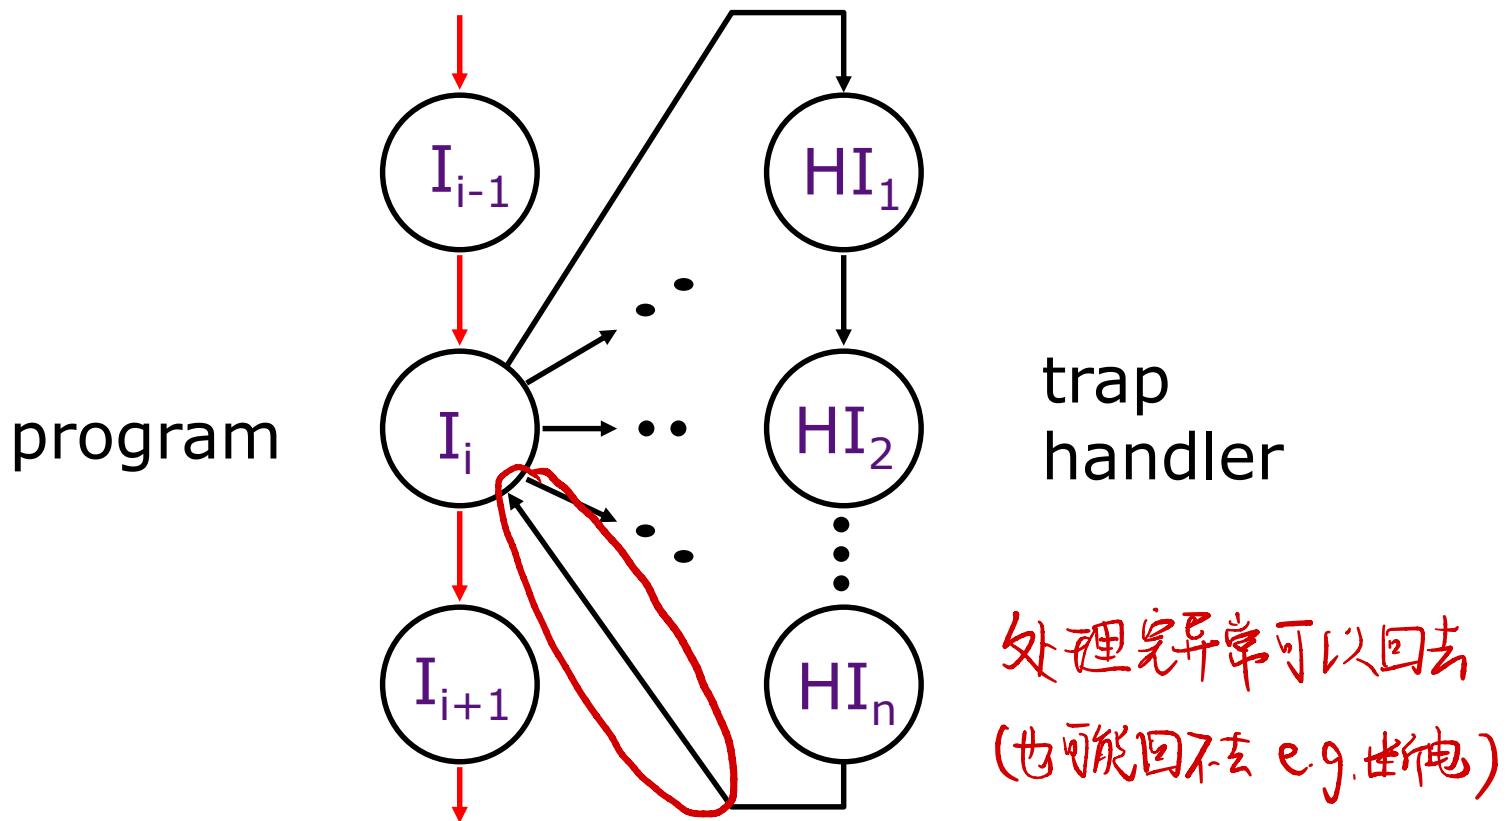
Terminology

In CA (you'll see other definitions in use elsewhere):

- Interrupt – caused by an event *external* to current running program (e.g. key press, mouse activity)
 - Asynchronous to current program, can handle interrupt on any convenient instruction
- Exception – caused by some event during execution of one instruction of current running program (e.g., page fault, bus error, illegal instruction)
 - Synchronous, must handle exception on instruction that causes exception
- Trap – action of servicing interrupt or exception by hardware jump to “trap handler” code

Traps/Interrupts/Exceptions:

altering the normal flow of control



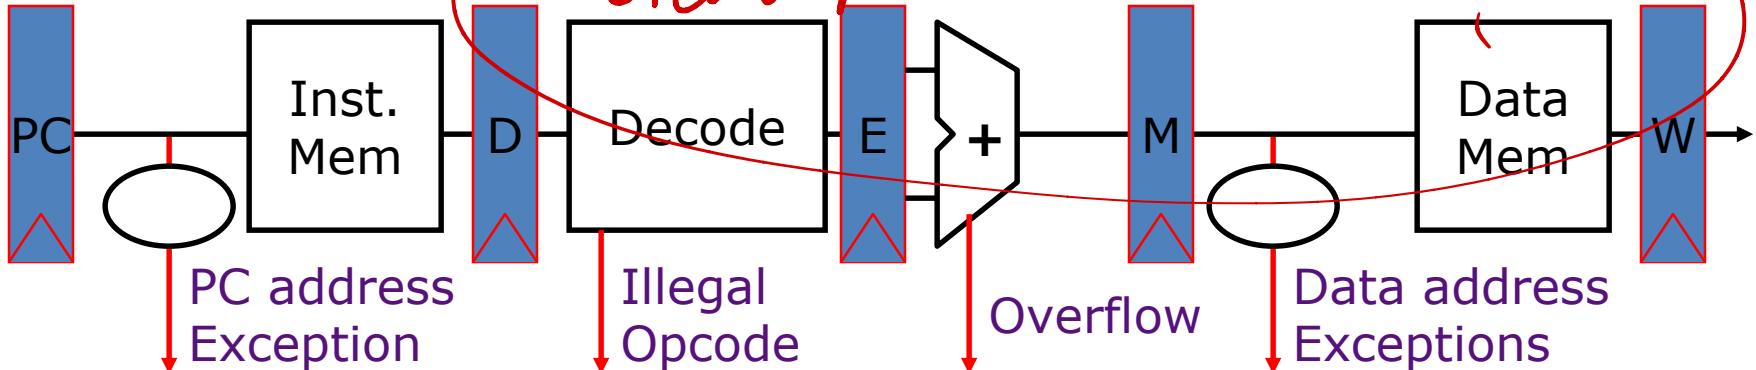
An *external or internal event* that needs to be processed - by another program – the OS. The event is often unexpected from original program's point of view.

Precise Traps

- *Trap handler's view of machine state is that every instruction prior to the trapped one has completed, and no instruction after the trap has executed.*
- Implies that handler can return from an interrupt by restoring user registers and jumping back to interrupted instruction (**SEPC** register will hold the instruction address)
 - Interrupt handler software doesn't need to understand the pipeline of the machine, or what program was doing!
 - More complex to handle trap caused by an exception than interrupt
- Providing precise traps is tricky in a pipelined superscalar out-of-order processor!
 - But handling imprecise interrupts in software is even worse.

Trap Handling in 5-Stage Pipeline

branch prediction 猜分支

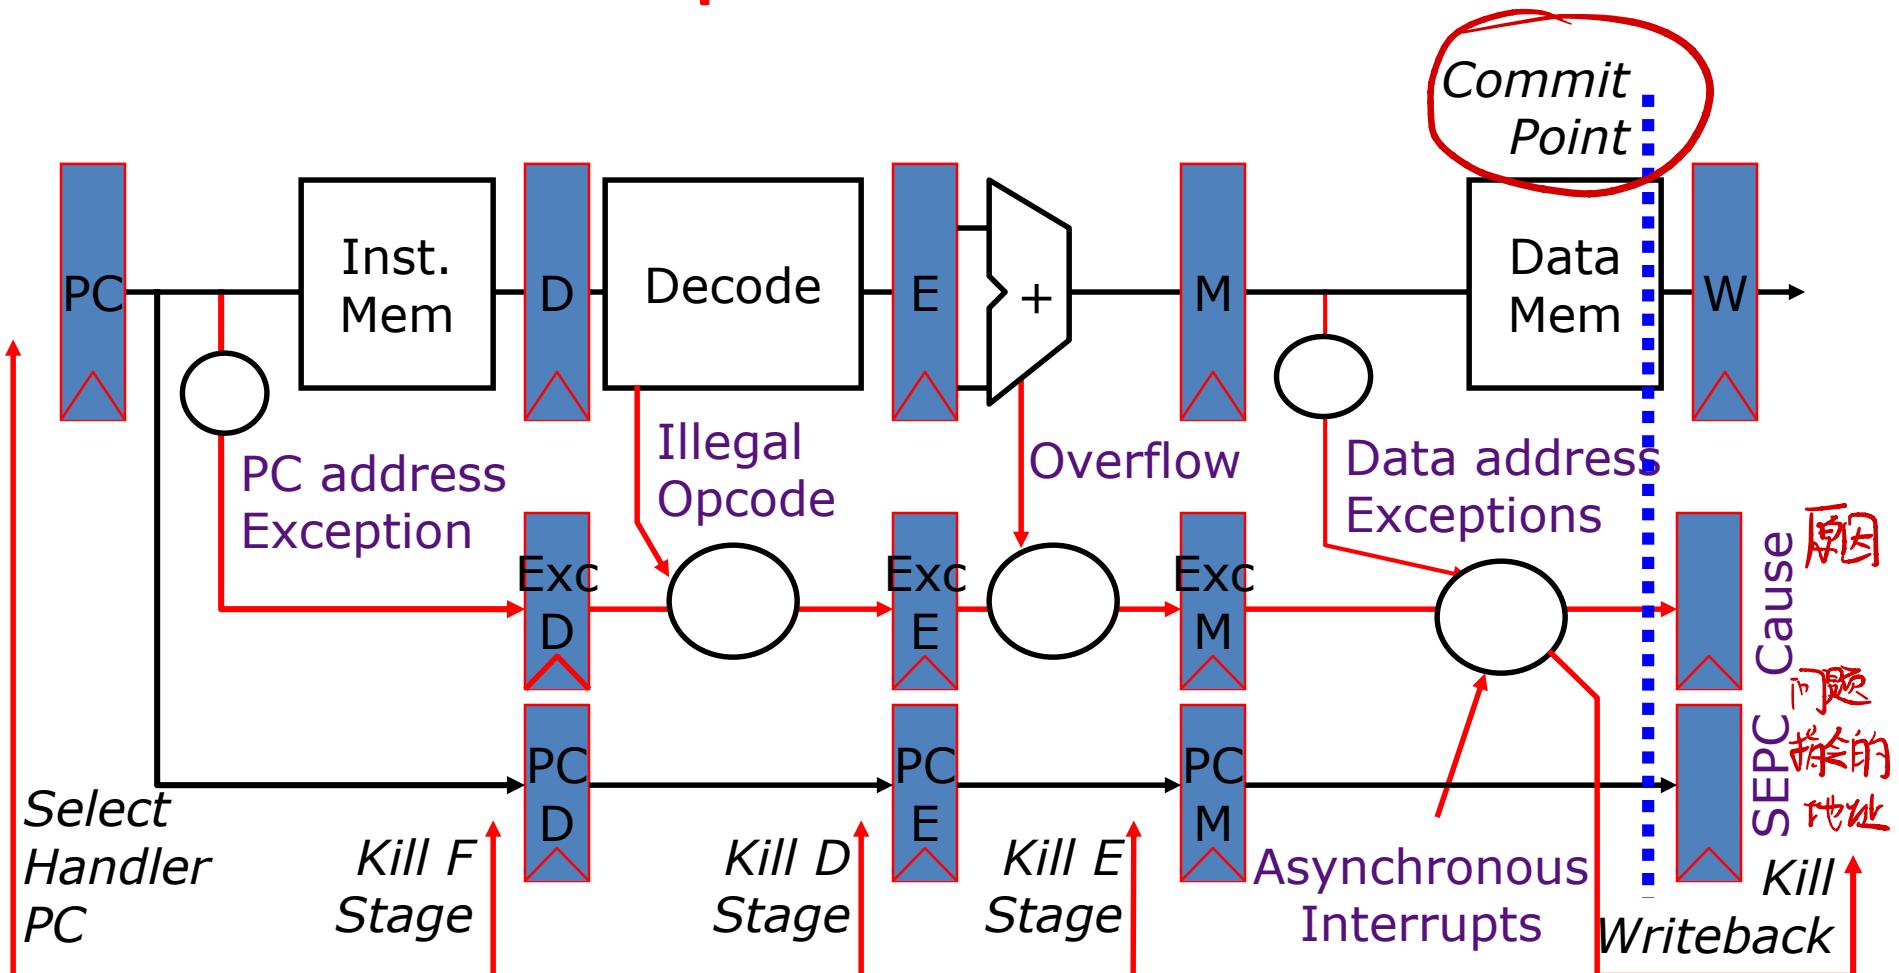


- How to handle multiple simultaneous exceptions in different pipeline stages? 多个阶段都出问题
- How and where to handle external asynchronous interrupts? 外部异步中断

先处理最先发生的异常

write back阶段(commit point), 统一处理exception

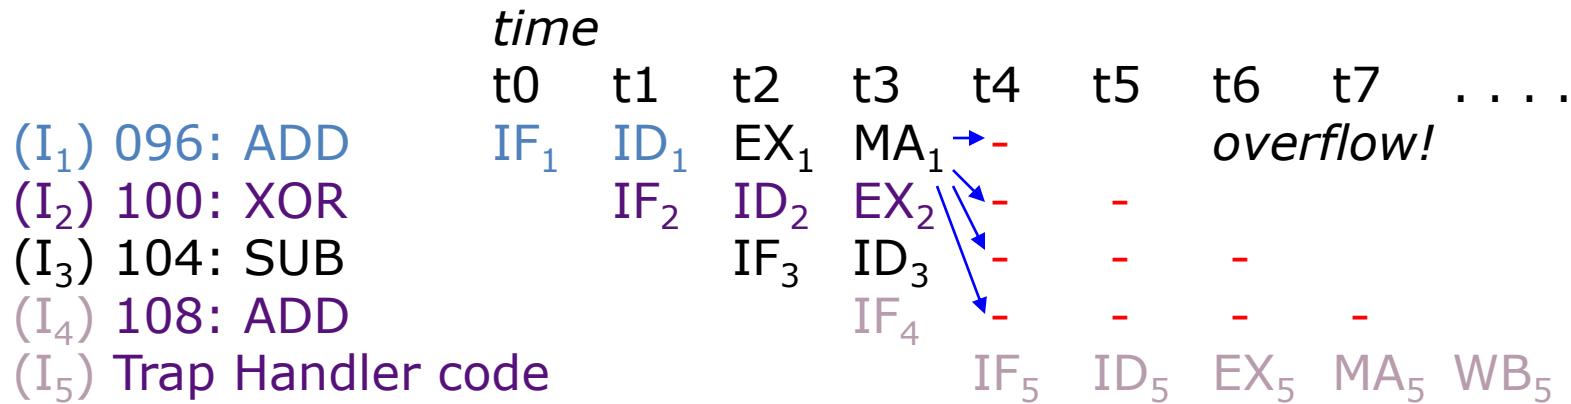
Save Exceptions Until Commit



Handling Traps in In-Order Pipeline

- Hold exception flags in pipeline until **commit point** (M stage)
- Exceptions in earlier instructions override exceptions in later instructions
- Exceptions in earlier pipe stages override later exceptions *for a given instruction*
- Inject external interrupts at commit point
- If exception/interrupt at commit: update Cause and SEPC registers, kill all stages, inject handler PC into fetch stage

Trap Pipeline Diagram



Agenda

- OS Boot Sequence and Operation
- Devices and I/O, interrupt and trap
- Application, Multiprogramming/time-sharing

Launching Applications

- Applications are called “processes” in most OSs.
 - Process: separate memory space;
 - Thread: shared memory space.
- Created by another process calling into an OS routine (using a “syscall”, more details later).
 - Depends on OS, but Linux uses `fork` to create a new process, and `execve` to load application.
- Loads executable file from disk (using the file system service) and puts instructions & data into memory (.text, .data sections), prepare stack and heap.
- Set argc and argv, jump into the main function.

Supervisor Mode

- If something goes wrong in an application, it could crash the entire machine.
 - And what about malware, etc.?
- The OS may need to enforce resource constraints to applications (e.g., access to devices).
- To help protect the OS from the application, CPUs have a **supervisor mode** bit.
 - When not in supervisor mode (user mode), a process can only access a subset of instructions and (physical) memory.
 - Process can enter the supervisor mode by using an **interrupt**, and change out of supervisor mode using a special instruction.

Syscalls

- What if we want to call into an OS routine? (e.g., to read a file, launch a new process, send data, etc.)
 - Need to perform a **syscall**: set up function arguments in registers, and then raise **software interrupt**
 - OS will perform the operation and return to user mode
- This way, the OS can mediate access to all resources, including devices and the CPU itself.

Multiprogramming

- The OS runs multiple applications at the same time.
- But not really (unless you have a core per process)
 - Time-sharing processor
- When jumping into process, set timer interrupt.
 - When it expires, store PC, registers, etc. (process state).
 - Pick a different process to run and load its state.
 - Set timer, change to user mode, jump to the new PC.
- Switches between processes very quickly. This is called a “context switch”.
- Deciding what process to run is called **scheduling**.

Protection, Translation, Paging

- Supervisor mode does not fully isolate applications from each other or from the OS.
 - Application could overwrite another application's memory.
 - Also, may want to address more memory than we actually have (e.g., for sparse data structures).
- Solution: [Virtual Memory](#). Gives each process the illusion of a full memory address space that it has completely for itself.

In Conclusion

- Once we have a basic machine, it's mostly up to the OS to use it and define application interfaces.
- Hardware helps by providing the right abstractions and features (e.g., Virtual Memory, I/O).