

CS 110

Computer Architecture

Advanced Caches

Instructors:

Chundong Wang & Siting Liu

<https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html>

School of Information Science and Technology SIST

ShanghaiTech University

Slides based on UC Berkeley's CS61C (2015)

Review

- Virtual Memory
- Page Table
 - Multi-level Page Table
- Page Fault
- TLB

访问cache的地址

VIVT

VIPT → 最常用

PIPT

V: virtual P: physical

I: index T: tag



Unlimited?

```
wangc@64G:~$ ulimit -s
8192
wangc@64G:~$ ulimit -a
core file size          (blocks, -c) 0
data seg size           (kbytes, -d) unlimited
scheduling priority     (-e) 0
file size               (blocks, -f) unlimited
pending signals         (-i) 256820
max locked memory       (kbytes, -l) 16384
max memory size         (kbytes, -m) unlimited
open files              (-n) 1024
pipe size                (512 bytes, -p) 8
POSIX message queues    (bytes, -q) 819200
real-time priority      (-r) 0
stack size              (kbytes, -s) 8192
cpu time                (seconds, -t) unlimited
max user processes      (-u) 256820
virtual memory          (kbytes, -v) unlimited
file locks              (-x) unlimited
wangc@64G:~$
```

Remember: Out of Memory

- Insufficient free memory: `malloc()` returns **NULL**

```
1 /*
2     This is a test for CS 110. All copyrights ...
3 */
4
5 #include <stdio.h>
6 #include <stdlib.h>
7
8 int main(int argc, char **argv) {
9     const int G = 1024 * 1024 * 1024;
10    for (int n = 0; ; n++) {
11        char *p = malloc(G * sizeof(char)); // 1GB every time
12        if (p == NULL) {
13            fprintf(stderr,
14                "failed to allocate > %g TeraBytes\n",
15                n / 1024.0);
16            return 1;
17        }
18        // no free, keep allocating until out of memory
19    }
20    return 0;
21 }
```

```
wangc@64G:~/TT$ gcc test.c -o t -Wall -O3
wangc@64G:~/TT$ ./t
failed to allocate > 127.99 TeraBytes
wangc@64G:~/TT$
```

Limited VM Space with x86-64

- 64-bit Linux allows up to **128TB** of virtual address space for individual processes, and can address approximately 64 TB of physical memory, subject to processor and system limitations.
- For Windows 64-bit versions, both 32- and 64-bit applications, if not linked with “*large address aware*”, are limited to **2GB** of virtual address space; otherwise, **128TB** for Windows 8.1 and Windows Server 2012 R2 or later.

AMD 48bit for address translation only

- Still provides plenty of space!
- Higher bits “sign extended”:
“canonical form”
- Convention: “Higher half” for
the Operating System
- Intel has plans (“whitepaper”) for
56 bit translation – no hardware yet

高位区 → OS
低 → 进程



- https://en.wikipedia.org/wiki/X86-64#Virtual_address_space_details

Using 128TB of Memory!?

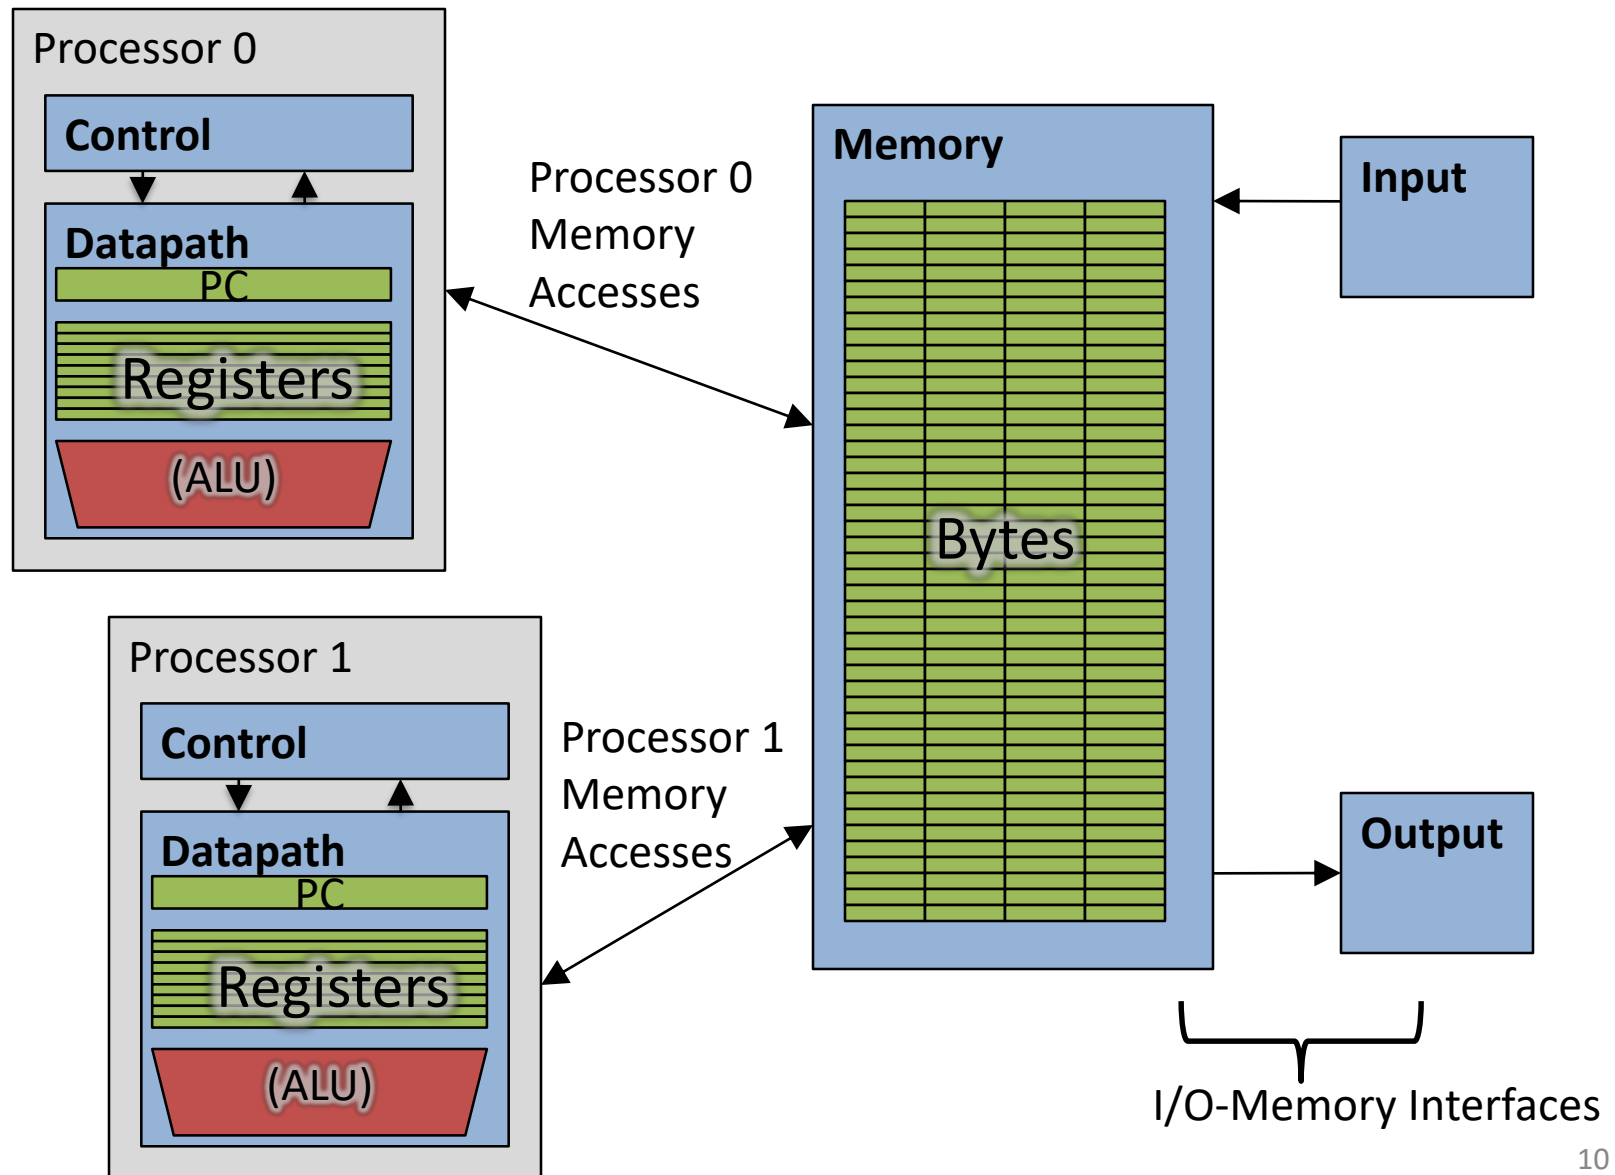
- A lazy allocation of virtual memory
 - Not used → not allocated
 - Try reading and writing from those pointers: works!
 - Even writing Gigabaytes of memory: works!
- Memory Compression! 压缩
 - Take not-recently used pages, compress them => free the physical page
- <https://www.lifewire.com/understanding-compressed-memory-os-x-2260327>

Process Name	Memory	Threads	Ports	PID	User	Compressed M...	Real Mem
a.out	60.51 GB	1	10	22329	schwertl	54.30 GB	6.22 GB

CACHE COHERENCE


高速缓存一致性

Simple Multi-core Processor



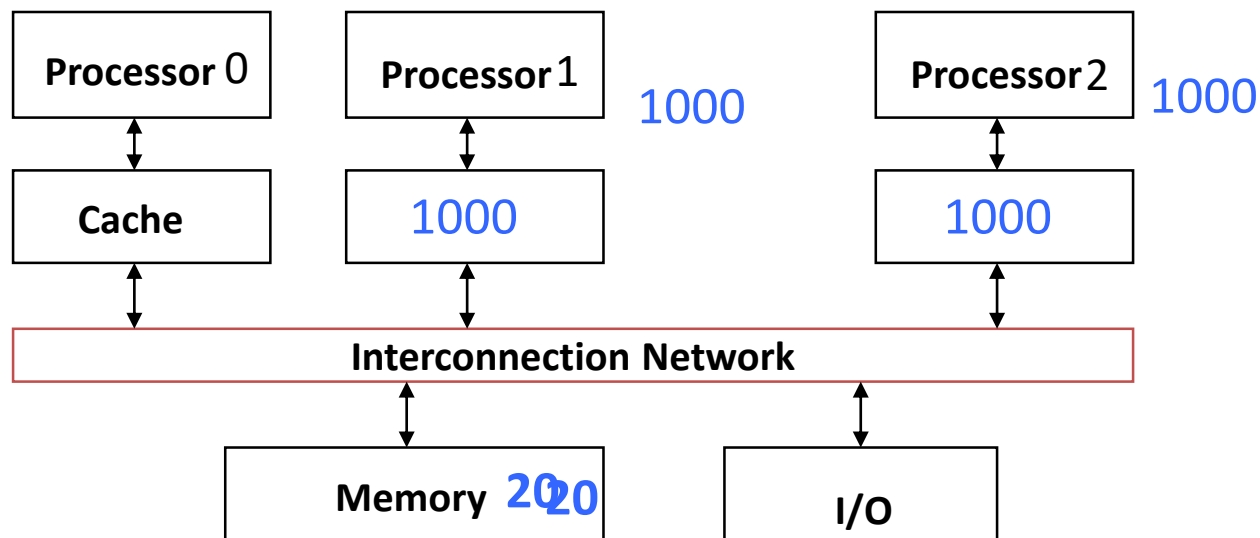
Multiprocessor Caches

- Memory is a performance bottleneck even with one processor
- Use caches to reduce bandwidth demands on main memory
- Each core has a local **private** cache holding data it has accessed recently 离CPU最近的东西: 最常访问的data
- Only cache misses have to access the **shared** common memory



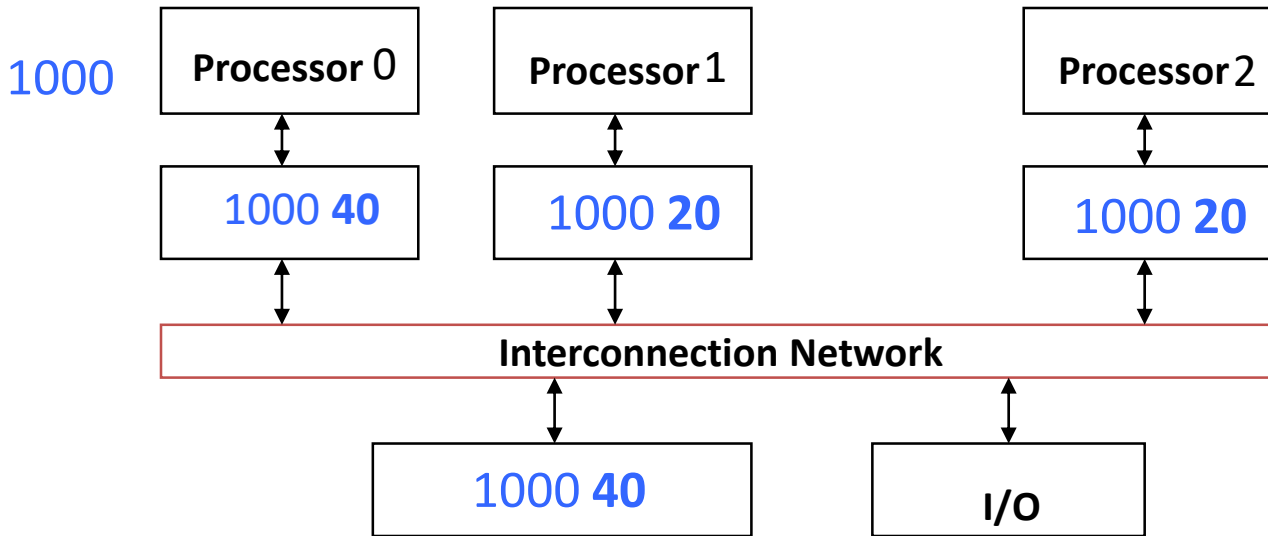
Shared Memory and Caches

- What if?
 - Processors 1 and 2 **read** Memory[1000] (value 20)



Shared Memory and Caches

- Now:
 - Processor 0 **writes** Memory[1000] with 40



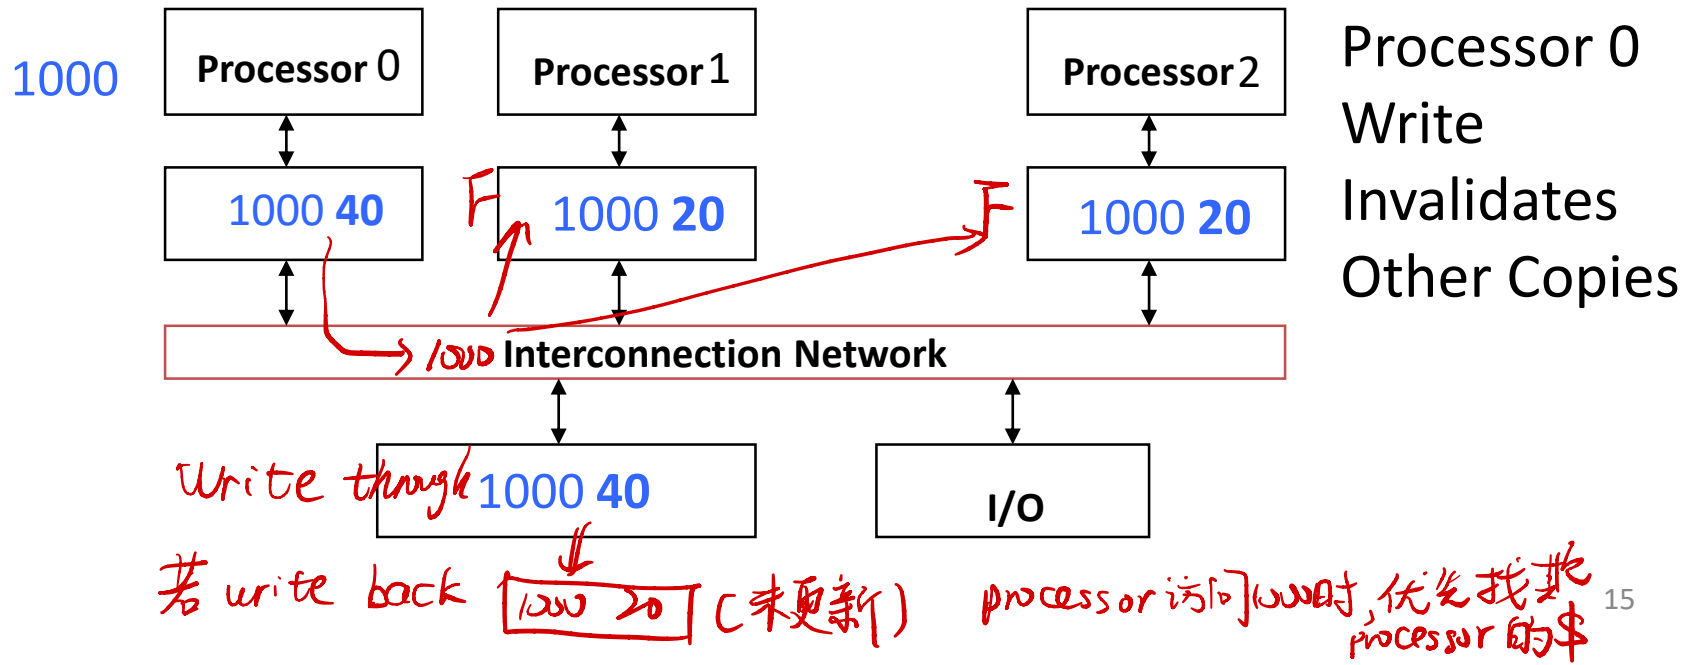
Problem?

Keeping Multiple Caches Coherent

- Architect's job: shared memory
=> keep cache values coherent
 - Idea: When any processor has cache miss or writes, notify other processors via interconnection network
 - If only reading, many processors can have copies
 - If a processor writes, invalidate any other copies
 - Write transactions from one processor, other caches “snoop” the common interconnect checking for tags they hold
 - Invalidate any copies of same address modified in other cache
- 修改 ⇒ 把 tag 相同的其它 \$ → invalid

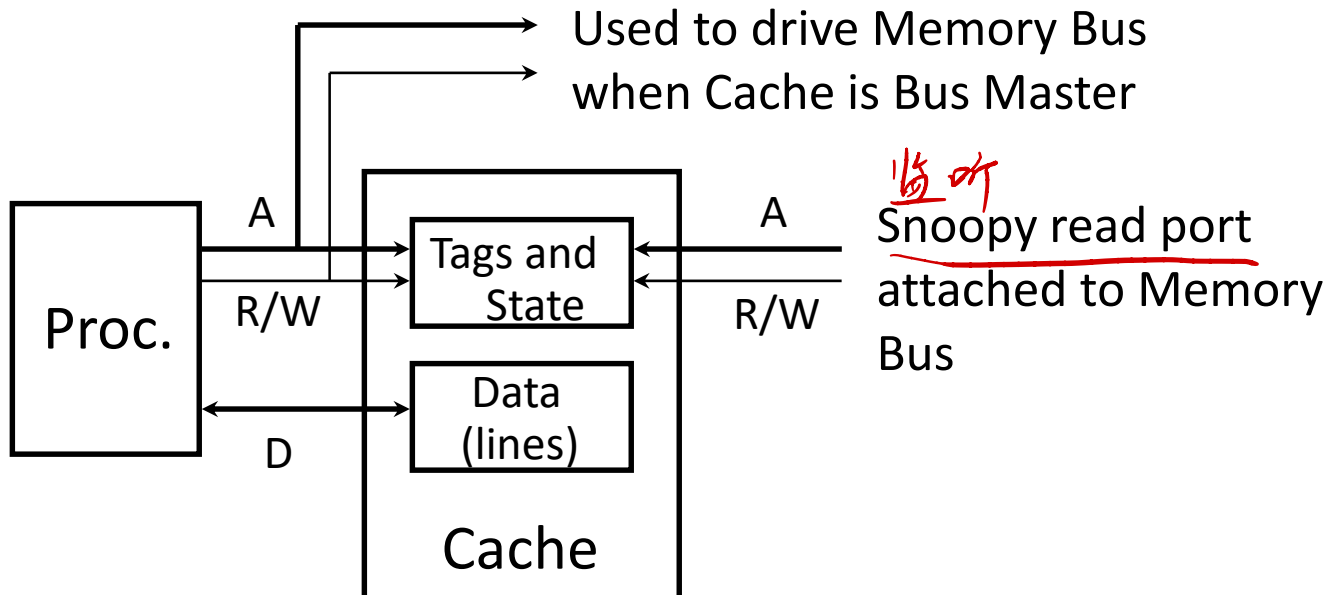
Shared Memory and Caches

- Example, now with cache coherence
 - Processors 1 and 2 read Memory[1000]
 - Processor 0 writes Memory[1000] with 40




Snoopy Cache, *Goodman 1983*

- Idea: Have cache watch (or snoop upon) other memory transactions, and then “do the right thing”
- Snoopy cache tags are dual-ported
双端



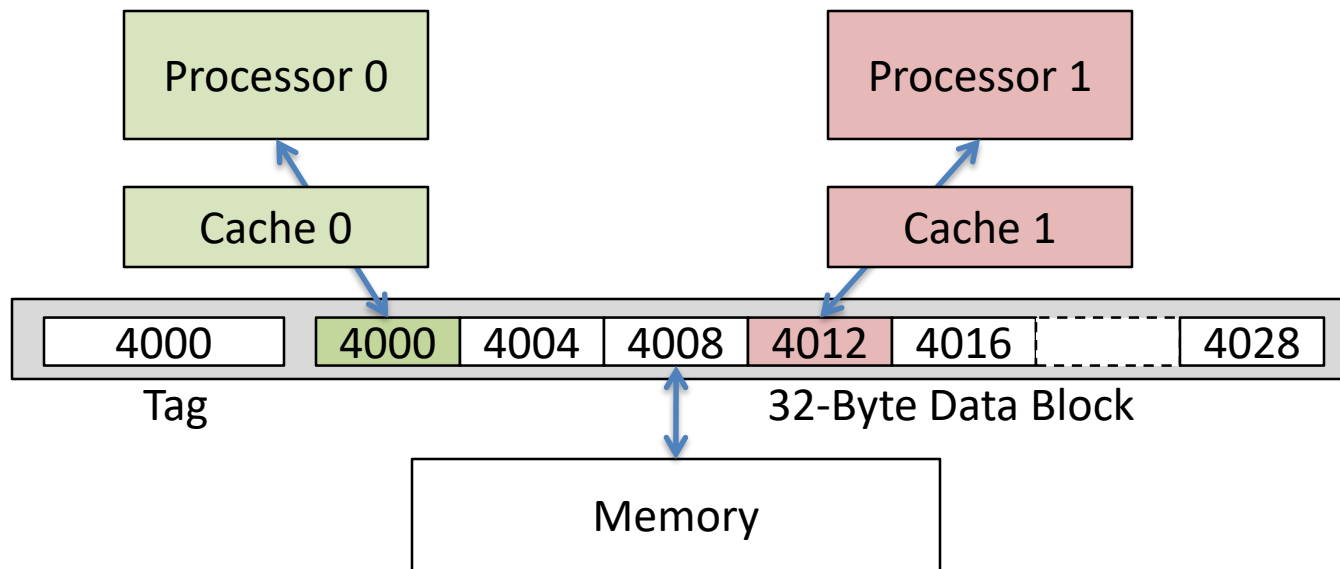
Optimized Snoop with Level-2 Caches



L₂ 负责监听

- Processors often have two-level caches
 - small L1, large L2 (usually both on chip now)
- Inclusion property: entries in L1 must be in L2
 - invalidation in L2 => invalidation in L1
- Snooping on L2 does not affect CPU-L1 bandwidth

Cache Coherency Tracked by Block



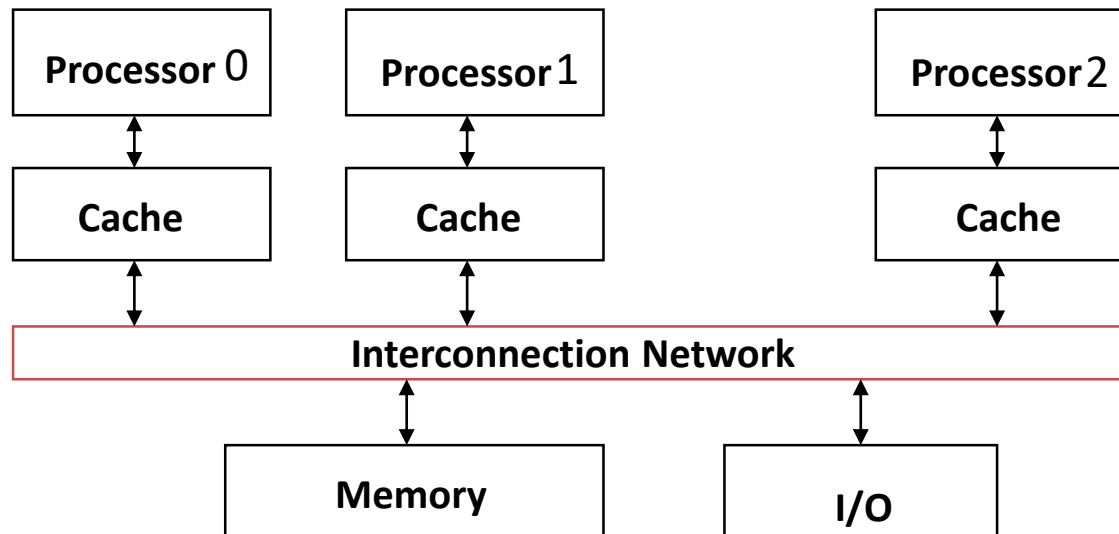
- Suppose block size is 32 bytes *8 words*
- Suppose Processor 0 reading and writing variable X, Processor 1 reading and writing variable Y
- Suppose in X location 4000, Y in 4012
- What will happen?

Coherency Tracked by Cache Block

- Block ping-pongs between two caches even though processors are accessing disjoint variables
- Effect called *false sharing*
- How can you prevent it?
 - Keep variables far apart (at least block size (64 byte))

Shared Memory and Caches

- Use valid bit to “unload” cache lines (in Processors 1 and 2)
- Dirty bit tells me: “I am the only one using this cache line”! => no need to announce on Network!



Review: Understanding Cache Misses: The 3Cs

- **Compulsory** (cold start or process migration, 1st reference):
 - First access to block, impossible to avoid; small effect for long-running programs
 - Solution: increase block size (increases miss penalty; very large blocks could increase miss rate)
- **Capacity** (not compulsory and...)
 - Cache cannot contain all blocks accessed by the program ***even with perfect replacement policy in fully associative cache***
 - Solution: increase cache size (may increase access time)
- **Conflict** (not compulsory or capacity and...):
 - Multiple memory locations map to the same cache location
 - Solution 1: increase cache size
 - Solution 2: increase associativity (may increase access time)
 - Solution 3: improve replacement policy, e.g.. LRU

Fourth “C” of Cache Misses: *Coherence Misses*

- Misses caused by coherence traffic with other processor
- Also known as *communication* misses because represents data moving between processors working together on a parallel program
- For some parallel programs, coherence misses can dominate total misses

Coherence Protocols 协议

- The cache coherence protocols ensure that there is a coherent view of data, with migration and replication.
 - A cache line has a state 移动 复制
- MSI
 - Modified, Shared, Invalid
- MESI
 - MSI + Exclusive
- MOESI
 - MESI + O ↪ Owner
 - e.g., AMD processor family
- MESIF
 - MESI + F
 - e.g., Intel Xeon processors

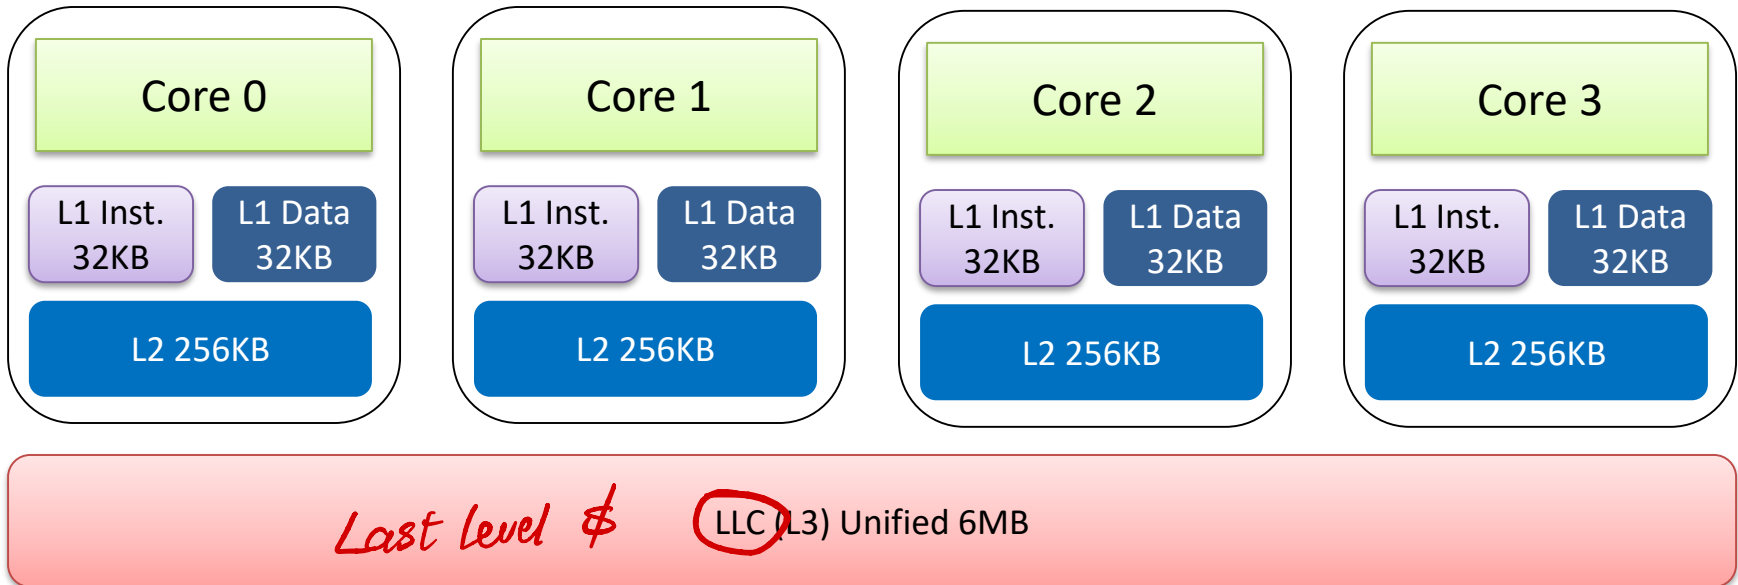
Advanced Caches:

MRU is LRU

most recently used

Cache Inclusion

- Multilevel caches



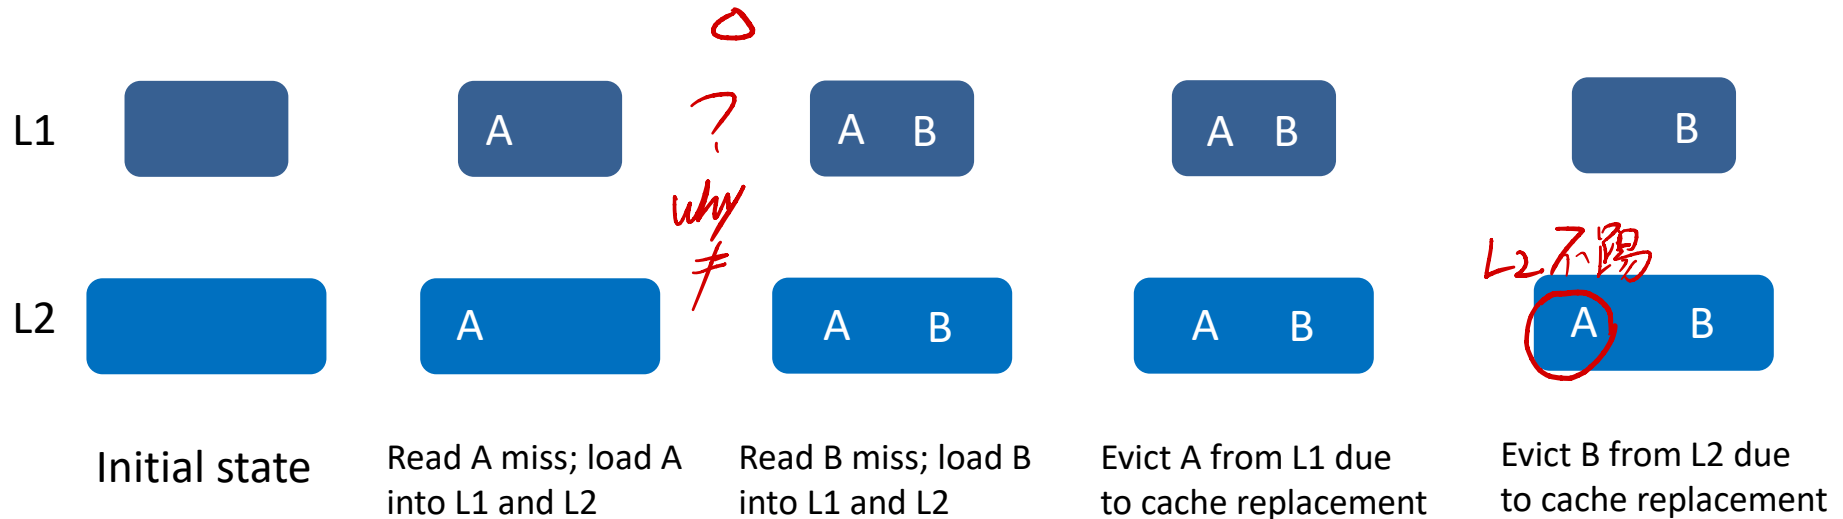
Intel Ivy Bridge Cache Architecture (Core i5-3470)

If all blocks in the higher level cache are also present in the lower level cache, then the lower level cache is said to be **inclusive** of the higher level cache.

↳ L₁中有, 则 L₂, L₃中也有

Inclusive

$$L_n \subseteq L_{n+1} \quad (n \geq 1)$$

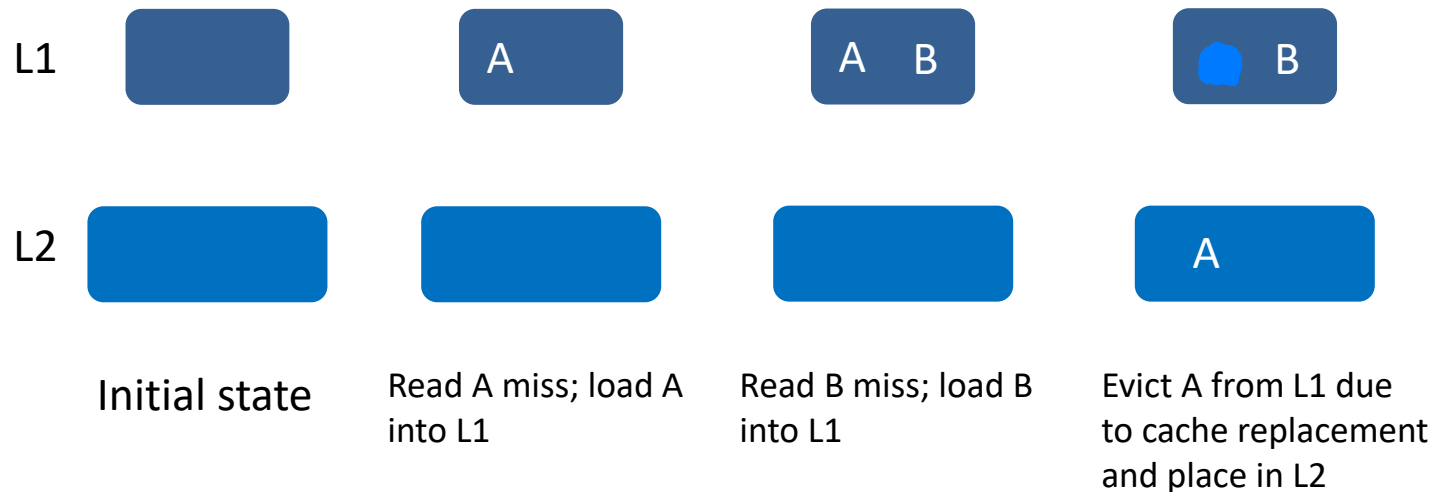


L2不陽

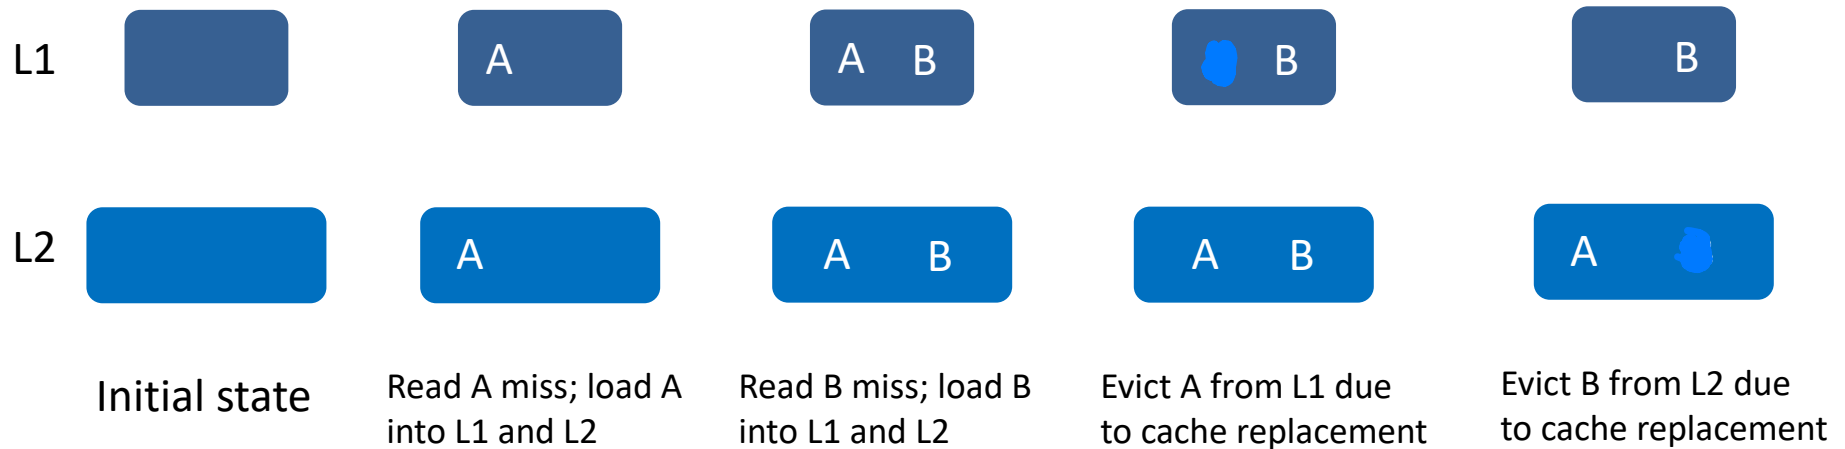
Back
invalidation

Exclusive

$$\underline{L_n \cap L_{n+1} = \emptyset \ (n \geq 1)}$$



Non-inclusive



Real-world CPUs

- Intel Processors
 - Sandy bridge, inclusive
 - Haswell, inclusive
 - Skylake-S, inclusive
 - Skylake-X, non-inclusive
- ARM Processors
 - ARMv7, non-inclusive
 - ARMv8, non-inclusive
- AMD
 - K6, exclusive
 - Zen, inclusive
 - Shanghai, LLC non-inclusive

Inclusive, or not?

- Inclusive cache eases coherence
 - A cache block in a higher-level surely existing in lower-level(s)
 - A non-inclusive LLC, say L2 cache, which needs to evict a block, **must** ask L1 cache if it has the block, because such information is not present in LLC.
- Non-inclusive cache yields higher performance though, why?
 - No back invalidation
 - More data can be cached ← larger capacity


'Sneaky' LRU for Inclusive Cache

CPU
Core

L1



A is frequently used



A is frequently hit in L1 cache. It is **MRU** in L1 cache.



Inclusive
LLC



A is evicted for replacement, in both L1 and L2



In LLC, A is **LRU**



In LLC, A is not frequently hit

A在L1里MRU, 在L2里LRU

As a result, MRU block that should be retained might be evicted, which causes performance penalty.

What if LLC is non-inclusive? *没问题*

Should you be interested, you can click <https://doi.org/10.1109/MICRO.2010.52> to read the related research paper for details.


Advanced Caches:

Reduce the size of LLC

Reduce LLC for high performance

- Problem

More than 83.8% LLC lines not productive



- A considerable portion of the shared LLC is dead

放入 LLC 后不再使用

dead block

- Why?

- LLC accesses, caused by L1 and L2 misses
 - Locality not accurate due to filtering by L1 and L2
 - LLC uniformly handles any access request for line allocation/deallocation

- How to resolve?

- Leverage the reuse locality to selectively allocate LLC lines

根据重用需求

Selective allocation upon reuse locality

- Reuse locality

- Selective allocation


- Tag and cache line decoupled

- Conventionally, one tag for one cache line

- Now, more tags than cache lines

- Some place holders

- Only keeping reused cache line

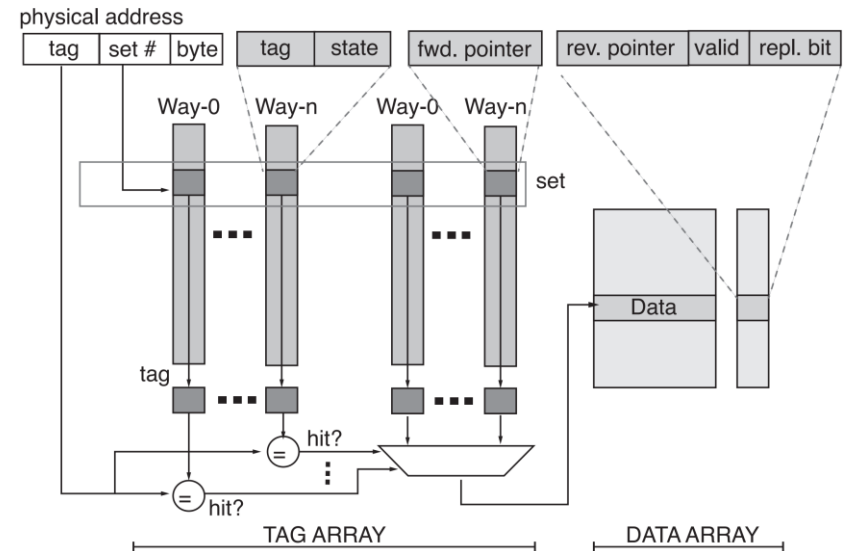
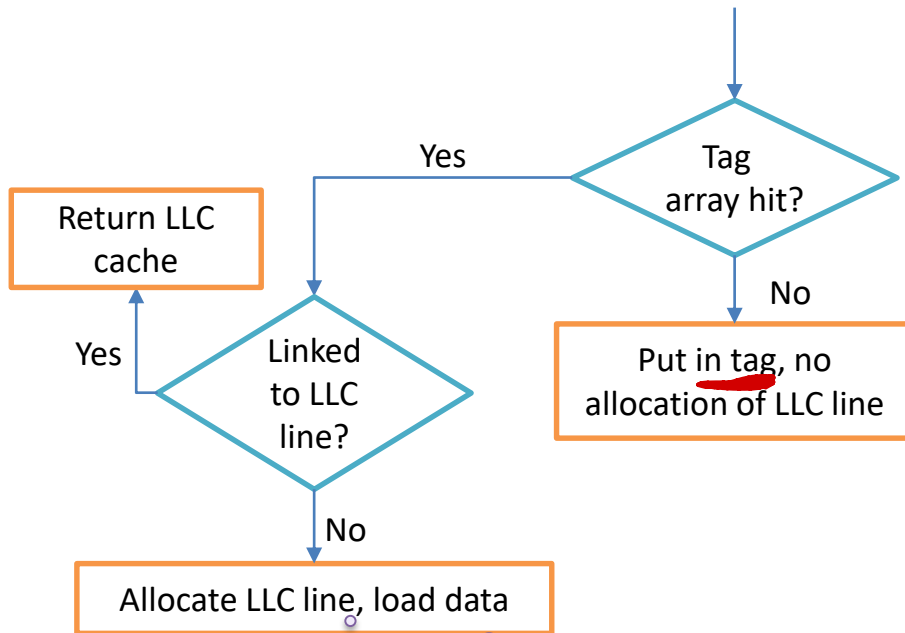


(b) Distribution of hits among all lines loaded (or reloaded) into the LRU SLLC during their stay. Each group represents 0.5% of the loaded lines

cache tag 分开

标记 & line 是否被重用

Allocation policy

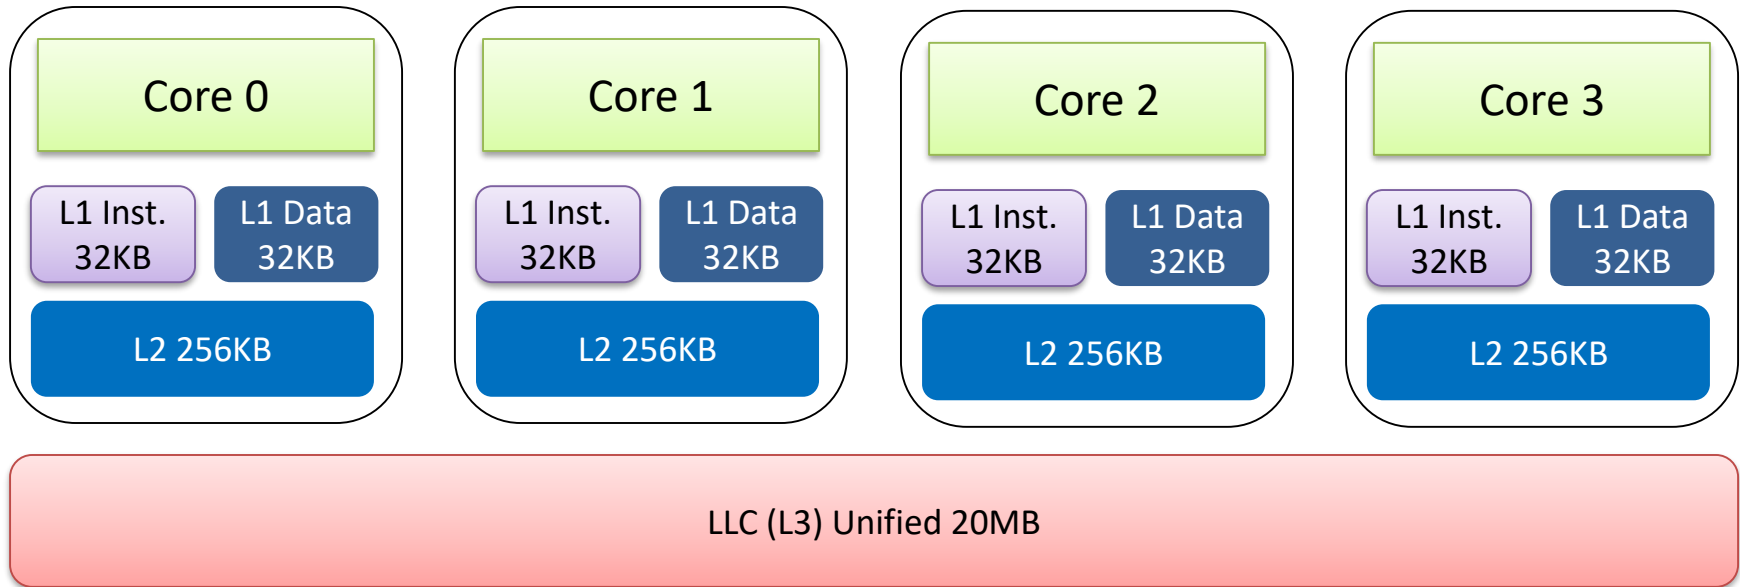


On replacement, the tag of evicted LLC line is kept. Once hit again, i.e., reused, data reloaded again.

Advanced Caches:
LLC is not monolithic

LLC is not monolithic

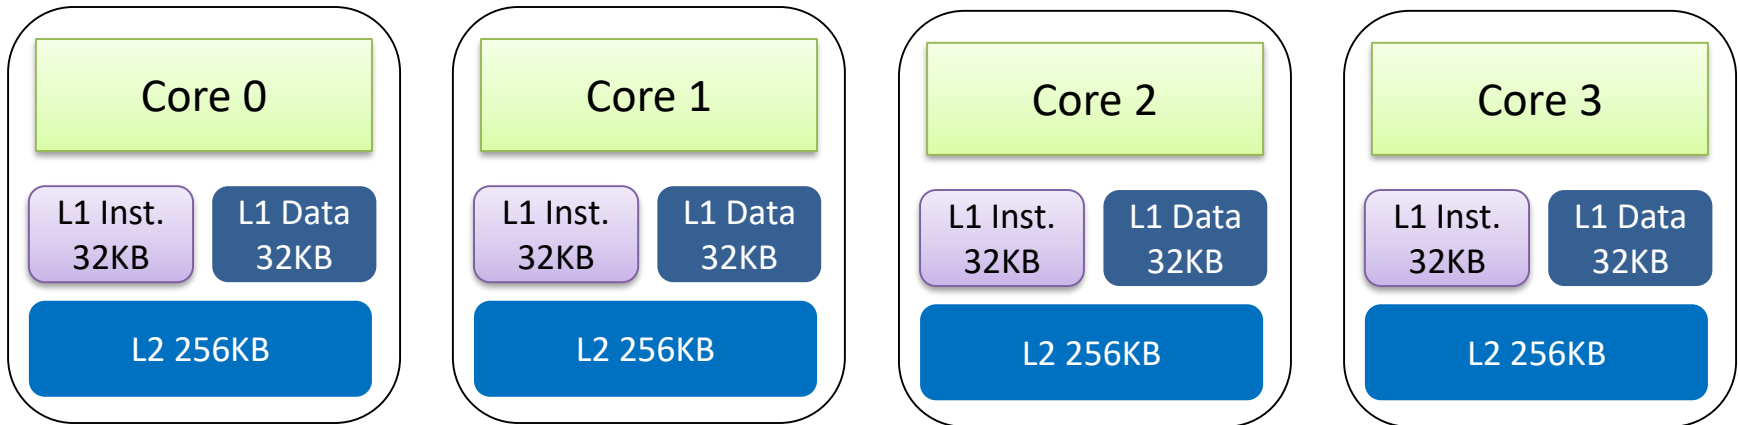
Intel® Xeon® Processor E5-2667 v3



Previously, it's considered that, to CPU cores, LLC is monolithic. No matter where a cache block in the LLC, a core would load it into private L2 and L1 cache with **the same** time cost.

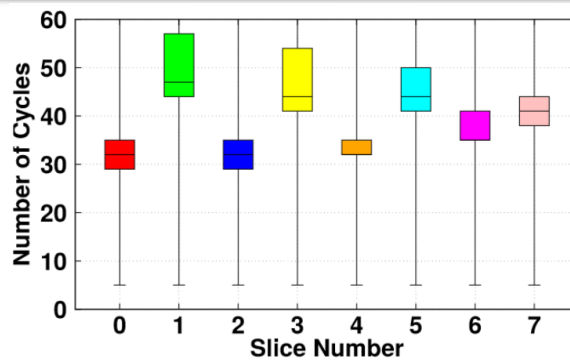
LLC is fine-grained

Intel® Xeon® Processor E5-2667 v3

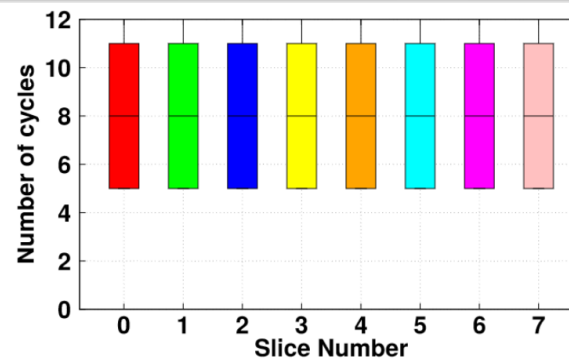


A

A



(a) Read.



(b) Write.

Slice-aware memory management

- The idea seems simple
 - Put your data closer to your program (core)
- But it not *EASY* to do so
 - Cache management is undocumented, not to mention fine-grained slices
 - Researchers did a lot of efforts
 - Click <https://doi.org/10.1145/3302424.3303977> for details
 - They managed to improve the average performance by 12.2% for GET operations of a key-value store.
 - 12.2% is a lot, if you consider the huge transactions every day for Google, Taobao, Tencent, JD, etc.

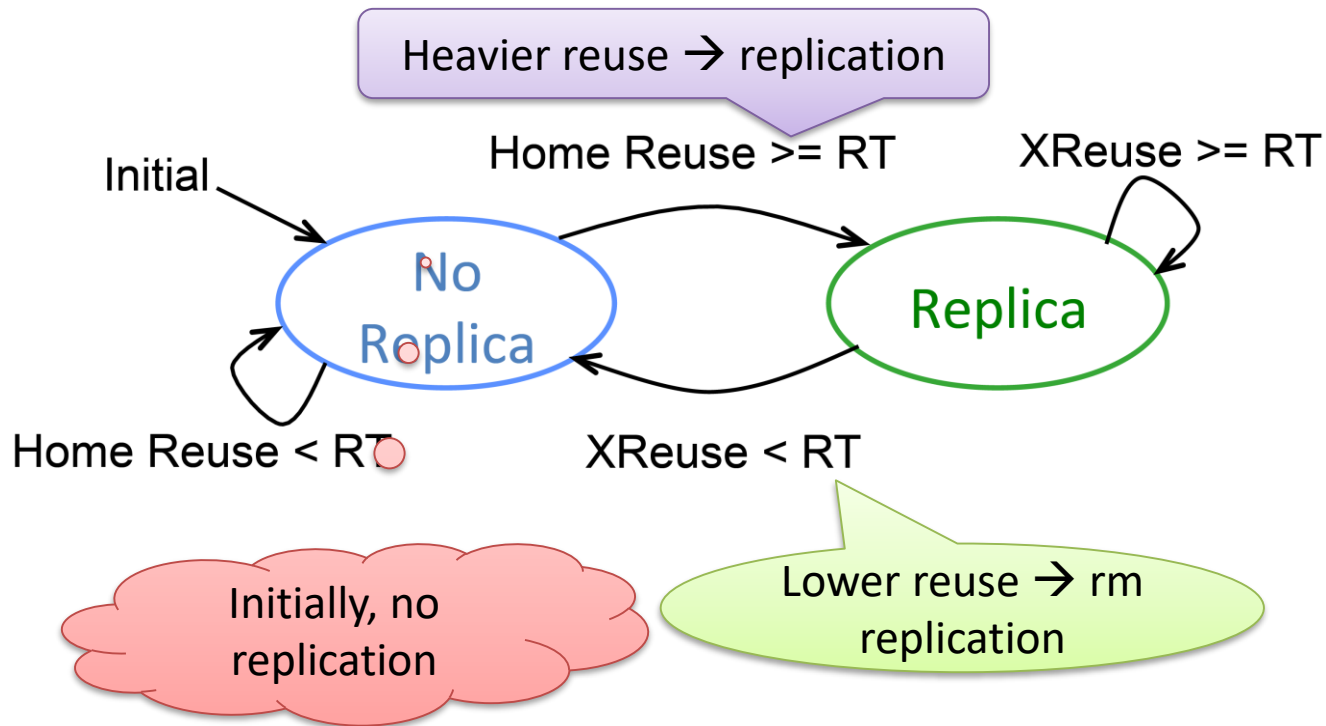
Advanced Caches: Replicating data in LLC

NUCA

- Non-uniform access cache
- The time cost to different locations in LLC is different for a particular core
 - In past slides, move data to a core using it
- Now replicate data at the **mircoarch** level
 - Towards a core using it

How to replicate cache lines?

- Reuse



How to replicate cache lines?

- Reuse

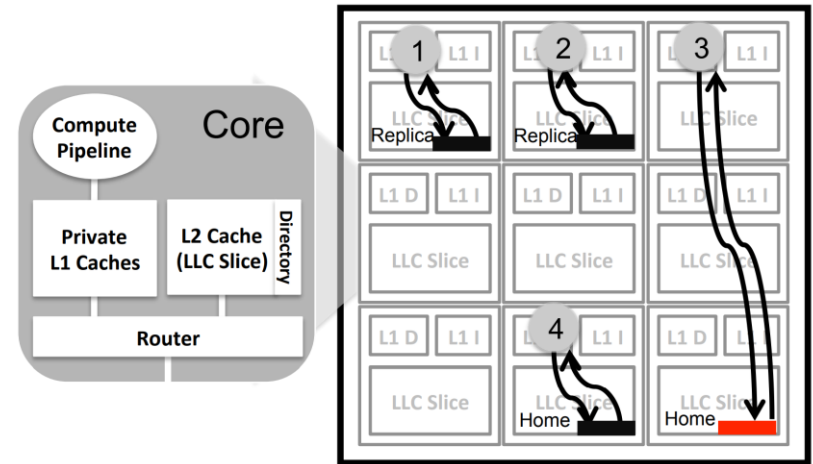
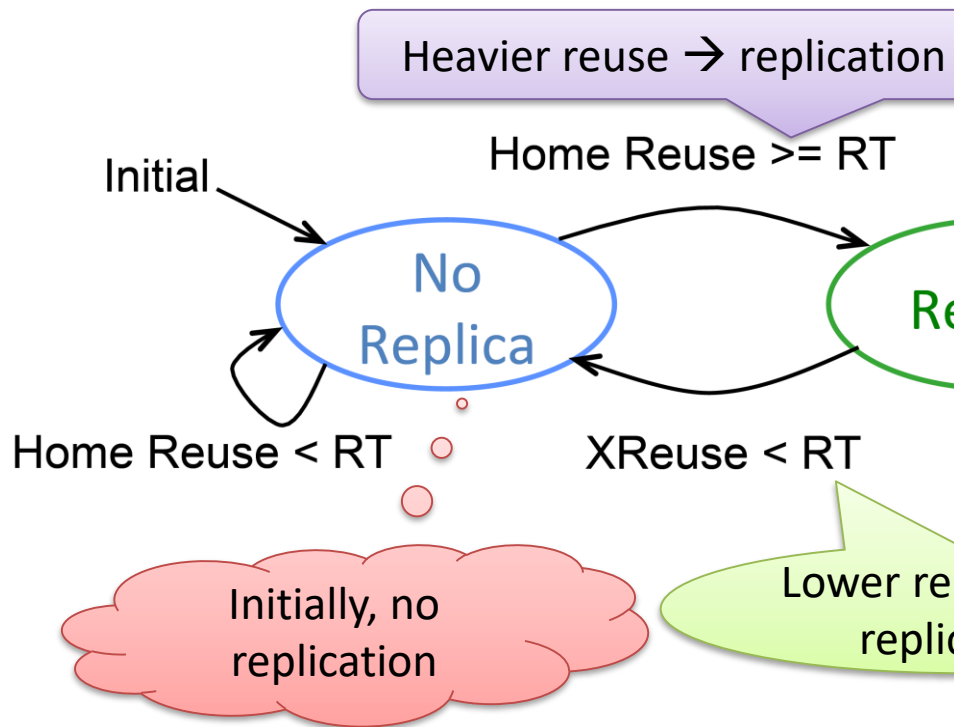


Figure 2. ① – ④ are mockup requests showing the *locality-aware LLC replication* protocol. The *black* data block has high reuse and a local LLC replica is allowed that services requests from ① and ②. The low-reuse *red* data block is not allowed to be replicated at the LLC, and the request from ③ that misses in the L1, must access the LLC slice at its home core. The home core for each data block can also service local private cache misses (e.g., ④).

Conclusion

- There are many interesting facts of CPU cache
- To make the best of cache can boost your program's performance!