

CS 110

Computer Architecture

Summary

Instructors:

Chundong Wang & Siting Liu

<https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2023/index.html>

School of Information Science and Technology SIST

ShanghaiTech University

Slides based on UC Berkley's CS61C

Let us review CA now.

New School Computer Architecture (1/3)



Personal
Mobile
Devices

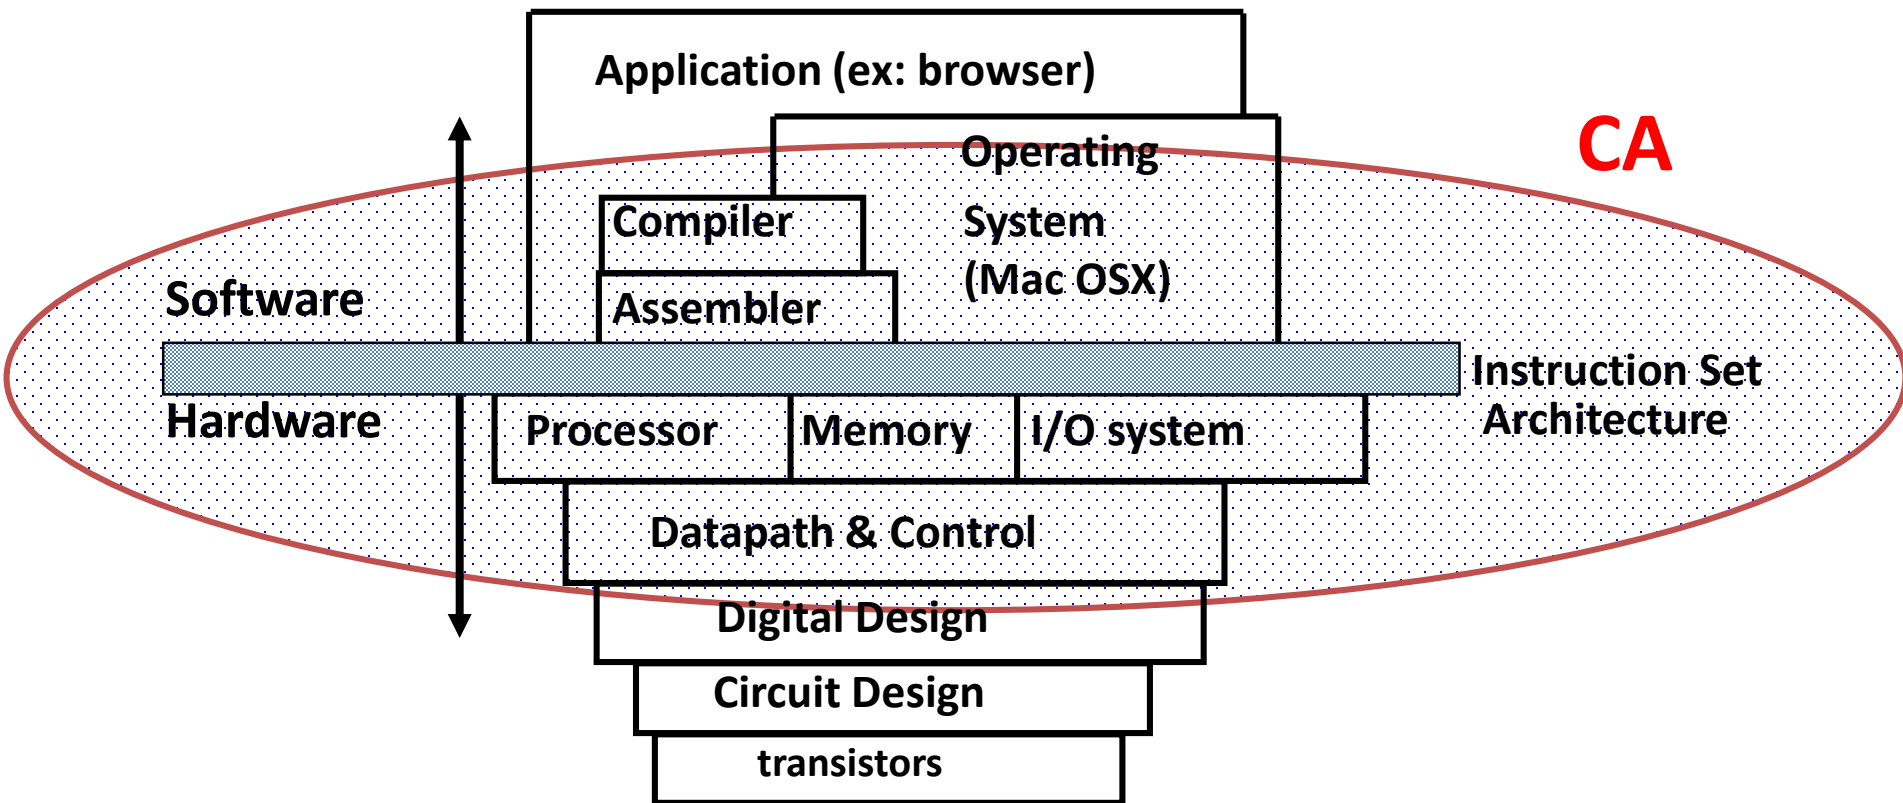
New School Computer Architecture (2/3)



New School Computer Architecture (3/3)



Old Machine Structures



New-School Machine Structures (It's a bit more complicated!)

Software

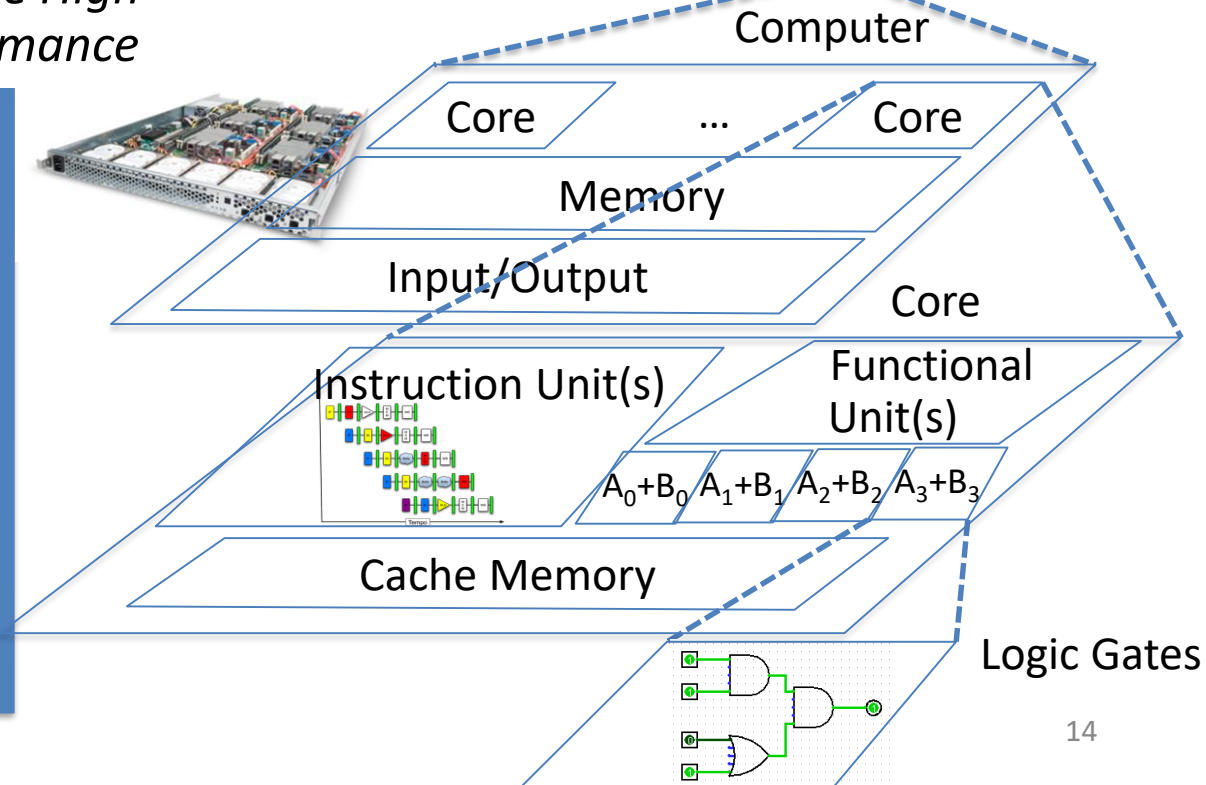
Hardware

Warehouse
Scale
Computer

Smart
Phone



*Leverage
Parallelism &
Achieve High
Performance*



- Parallel Requests
Assigned to computer
e.g., Search “Avatar 2”
- Parallel Threads
Assigned to core
e.g., Lookup, Ads
- Parallel Instructions
>1 instruction @ one time
e.g., 5 pipelined instructions
- Parallel Data
>1 data item @ one time
e.g., Add of 4 pairs of words
- Hardware descriptions
All gates functioning in
parallel at same time
- Programming Languages

Great Ideas in Computer Architecture

1. Design for Moore's Law
2. Abstraction to Simplify Design *eg. network*
3. Make the Common Case Fast
4. Dependability via Redundancy
5. Memory Hierarchy
6. Performance via
Parallelism/Pipelining/Prediction


Powers of Ten inspired CA Overview

Going Top-Down cover 3 Views

1. Architecture (when possible)
 2. Physical Implementation of that architecture
 3. Programming system for that architecture and implementation (when possible)
- See <http://www.powersof10.com/film>

Earth


10^7 meters




10^7 meters



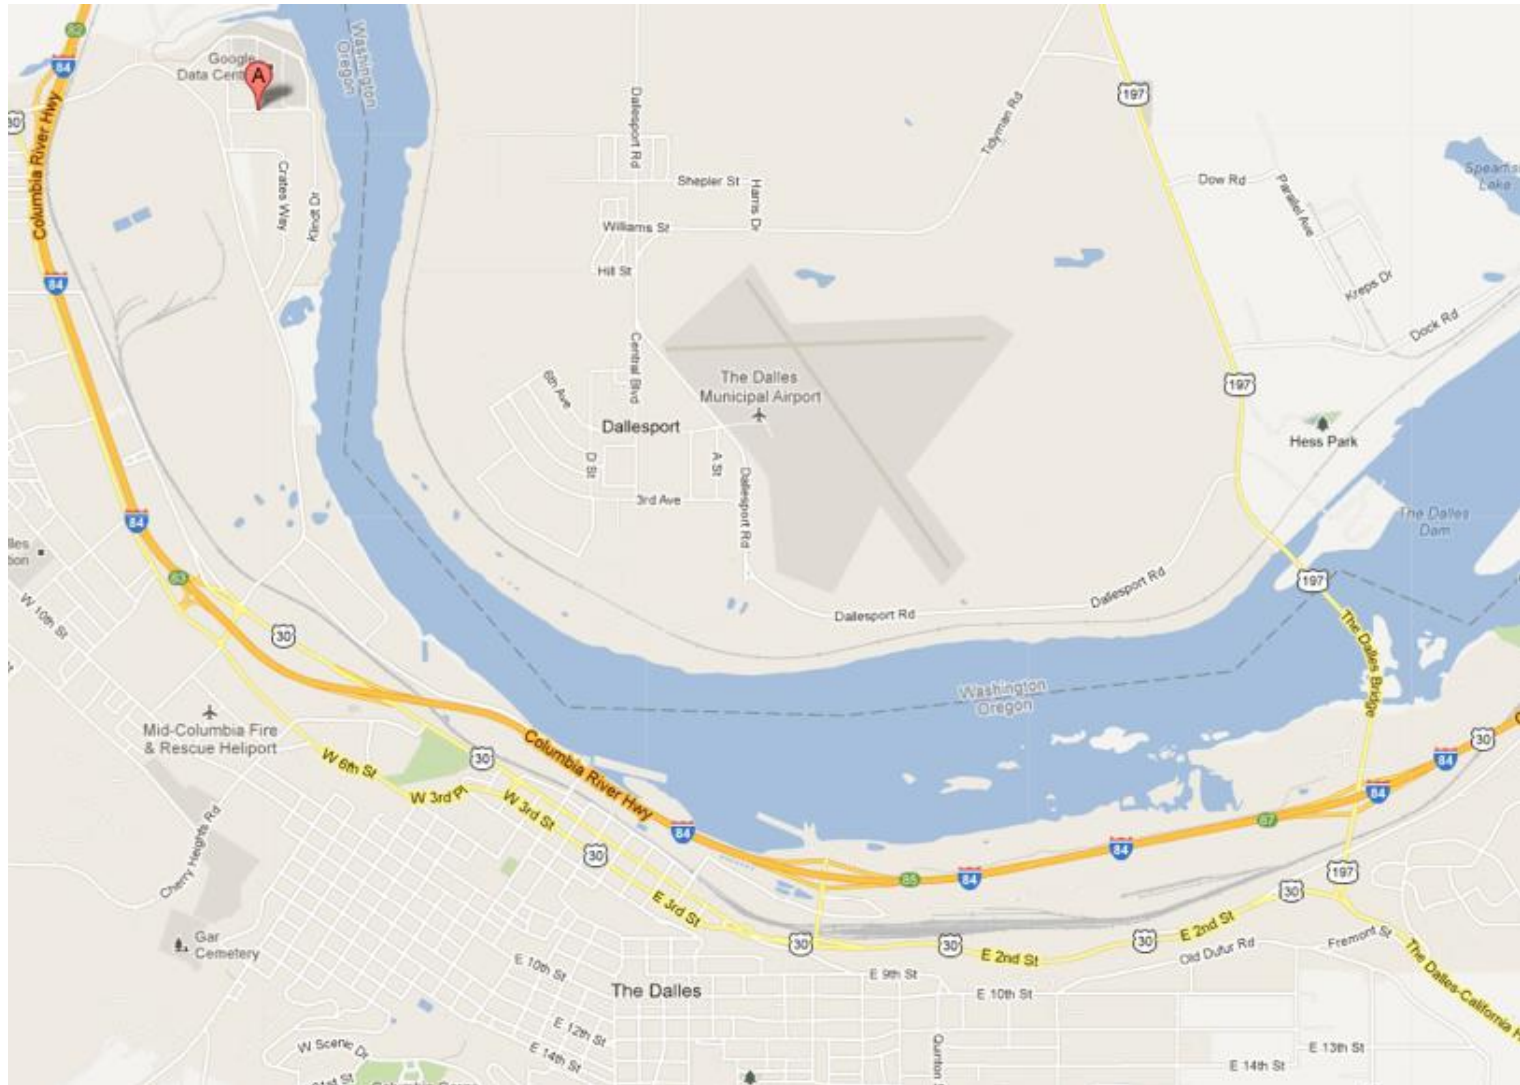
10^6 meters





10^5 meters

The Dalles, Oregon 10^4 meters



The Dalles, Oregon 10^4 meters



Google's Oregon WSC 10^3 meters



Google's Oregon WSC 10^4 meters

10 kilometers



10^2 meters



10^3 meters



Google Warehouse

- 90 meters by 75 meters, 10 Megawatts
- Contains 40,000 servers, 190,000 disks
- Power Utilization Effectiveness: 1.23
 - 85% of 0.23 overhead goes to cooling losses
 - 15% of 0.23 overhead goes to power losses
- Contains 45, 40-foot long containers
 - 8 feet × 9.5 feet × 40 feet
- 30 stacked as double layer, 15 as single layer

Containers in WSCs

10^2 meters



100 meters

Google Container

10¹ meters





FIG. 1

Google Container

10⁰ meters

10 meters



- 2 long rows, each with 29 racks
- Cooling below raised floor
- Hot air returned behind racks


Equipment Inside a Container



Server (in rack format):



7 foot Rack: servers + Ethernet local area network switch in middle (“rack switch”)




Array (aka cluster):
server racks + larger local area network switch (“array switch”) 10X faster => cost 100X: cost $f(N^2)$

Great Ideas in Computer Architecture

1. *Design for Moore's Law*
 - *WSC, Container, Rack*
2. Abstraction to Simplify Design
3. Make the Common Case Fast
4. *Dependability via Redundancy*
 - *Multiple WSCs, Multiple Racks, Multiple Switches*
5. Memory Hierarchy
6. *Performance via Parallelism/Pipelining/Prediction*
 - *Task level Parallelism, Data Level Parallelism*

Google Server Internals^{10⁻¹ meters}

10 centimeters



Facebook Datacenter

facebook



Software: Often uses MapReduce

- Simple data-parallel ***programming model*** and ***implementation*** for processing large datasets
- Users specify the computation in terms of
 - a ***map*** function, and
 - a ***reduce*** function
- Underlying runtime system
 - Automatically ***parallelize*** the computation across large scale clusters of machines
 - ***Handles*** machine ***failure***
 - ***Schedule*** inter-machine communication to make efficient use of the networks

} 多台机器

Programming Multicore Microprocessor: OpenMP

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000;
int value[num_steps];
int reduce()
{
    int i;
    int sum = 0;
    #pragma omp parallel for private(x) reduction(+:sum)
    for (i=1; i<= num_steps; i++) {
        sum = sum + value[i];
    }
}
```

Great Ideas in Computer Architecture

1. *Design for Moore's Law*
 - *More transistors = Multicore + SIMD*
2. Abstraction to Simplify Design
3. Make the Common Case Fast
4. Dependability via Redundancy
5. *Memory Hierarchy*
 - *More transistors = Cache Memories*
6. *Performance via Parallelism/Pipelining/Prediction*
 - *Thread-level Parallelism*

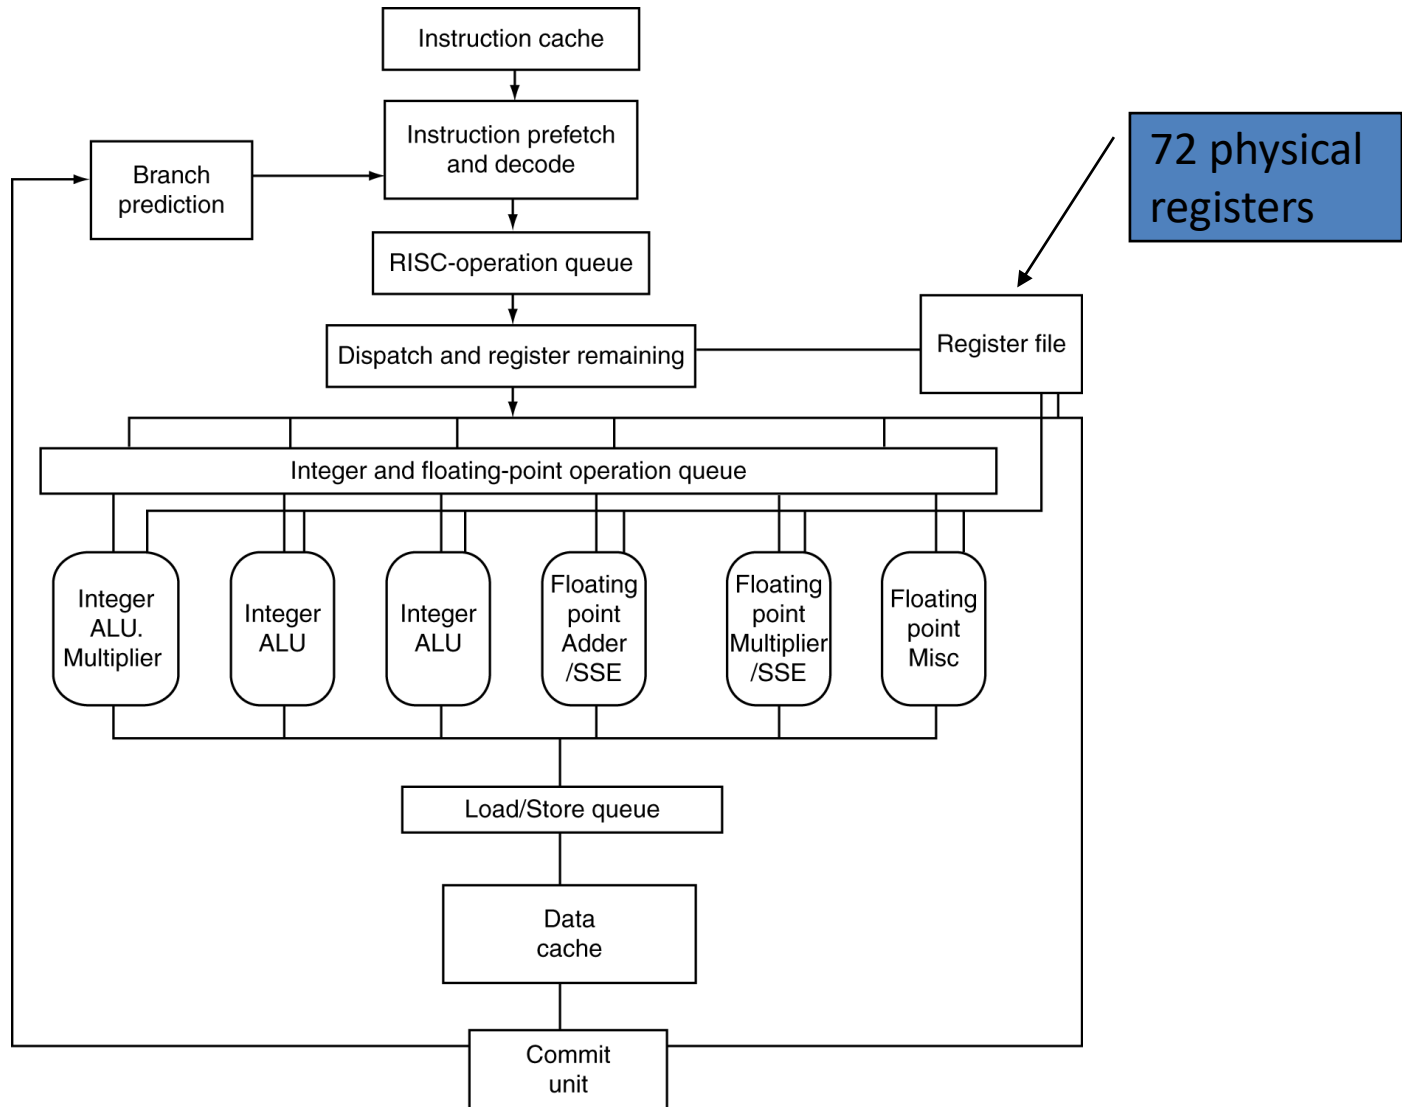
10^{-2} meters

AMD Opteron Microprocessor

centimeters

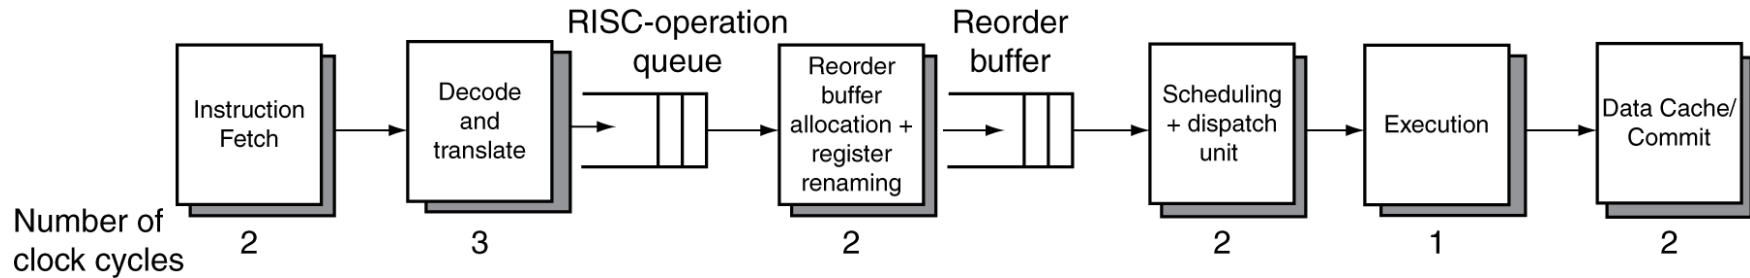


AMD Opteron Microarchitecture



AMD Opteron Pipeline Flow

- For integer operations

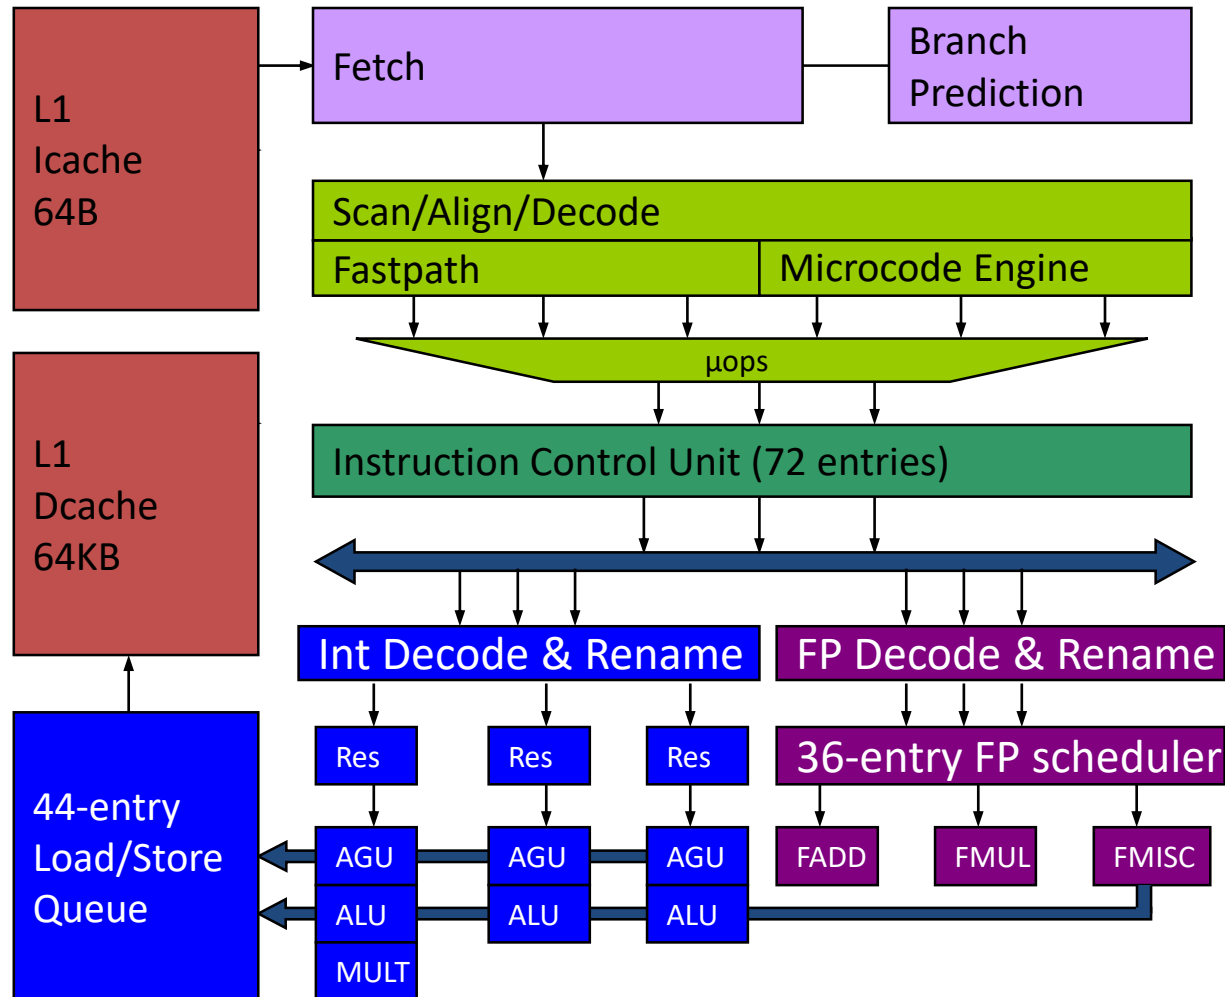


12级流水

12 stages (Floating Point is 17 stages)

Up to 106 RISC-ops in progress


AMD Opteron Block Diagram



10^{-2} meters

AMD Opteron Microprocessor


centimeters



10^{-3} meters

AMD Opteron Core

millimeters



Zoom into a Microchip

Zoom into a Microchip


From a Digital Camera
To a Scanning Electron
Microscope

Produced by
NISE Net

10^{-3} meters

AMD Opteron Core

millimeters



Programming One Core: C with Intrinsics

```
void mmult(int n, float *A, float *B, float *C)
{
    for ( int i = 0; i < n; i+=4 )
        for ( int j = 0; j < n; j++ )
        {
            __m128 c0 = _mm_load_ps(C+i+j*n);
            for( int k = 0; k < n; k++ )
                c0 = _mm_add_ps(c0,
                                _mm_mul_ps(_mm_load_ps(A+i+k*n),
                                              _mm_load1_ps(B+k+j*n)));
            _mm_store_ps(C+i+j*n, c0);
        }
}
```

p: pack

s: single 单精度

d: double 双精度

What are *p*
and *s* for?

$$C[i][j] = A[i][k] * B[k][j]$$

Inner loop from gcc -O -S

Assembly snippet from innermost loop:


```
movaps (%rax), %xmm9
mulps  %xmm0, %xmm9
addps  %xmm9, %xmm8
movaps 16(%rax), %xmm9
mulps  %xmm0, %xmm9
addps  %xmm9, %xmm7
movaps 32(%rax), %xmm9
mulps  %xmm0, %xmm9
addps  %xmm9, %xmm6
movaps 48(%rax), %xmm9
mulps  %xmm0, %xmm9
addps  %xmm9, %xmm5
```


Great Ideas in Computer Architecture

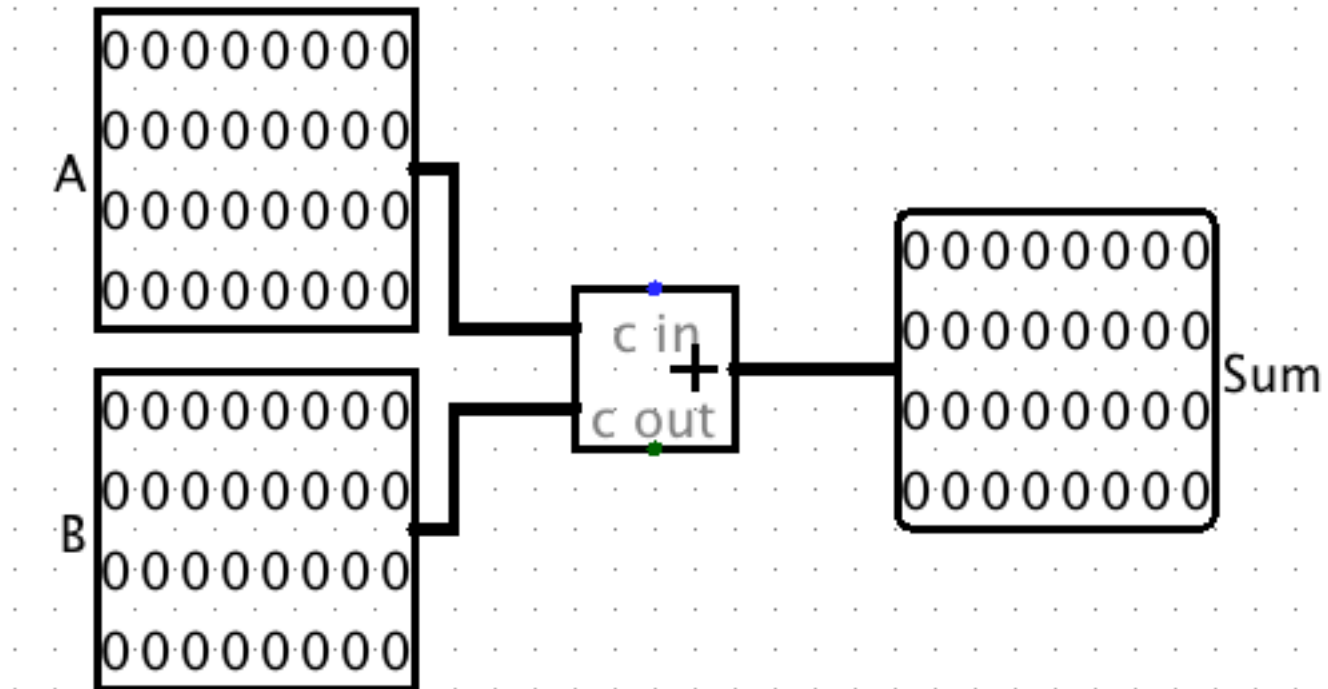
1. Design for Moore's Law
2. *Abstraction to Simplify Design*
 - *Instruction Set Architecture, Micro-operations*
3. Make the Common Case Fast
4. Dependability via Redundancy
5. Memory Hierarchy
6. *Performance via Parallelism/Pipelining/Prediction*
 - *Instruction-level Parallelism (superscalar, pipelining)*
 - *Data-level Parallelism*

SIMD Adder


- Four 32-bit adders that operate in parallel
 - Data Level Parallelism




One 32-bit Adder



1 bit of 32-bit Adder




Complementary MOS Transistors (NMOS and PMOS) of NAND Gate



x	y	z
0 volts	0 volts	3 volts
0 volts	3 volts	3 volts
3 volts	0 volts	3 volts
3 volts	3 volts	0 volts

Physical Layout of NAND Gate 10^{-7} meters


100 nanometers




10^{-7} meters

Scanning Electron Microscope

100 nanometers



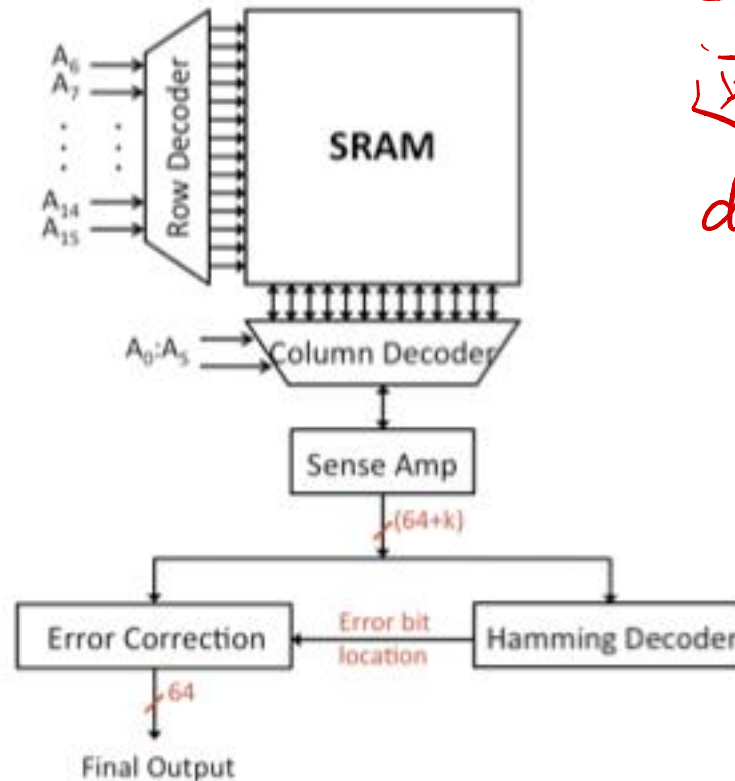
Top View



Cross Section


How to make a CMOS chip?

Block Diagram of Static RAM



SRAM
速度快
density大


1 Bit SRAM in 6 Transistors



10^{-7} meters

Physical Layout of SRAM Bit


100 nanometers



10^{-7} meters

SRAM Cross Section

100 nanometers




DIMM Module

- DDR = Double Data Rate
 - Transfers bits on Falling AND Rising Clock Edge
- Has Single Error Correcting, Double Error Detecting Redundancy (SEC/DED)
 - 72 bits to store 64 bits of data
 - Uses “Chip kill” organization so that if single DRAM chip fails can still detect failure
- Average server has 22,000 correctable errors and 1 uncorrectable error per year

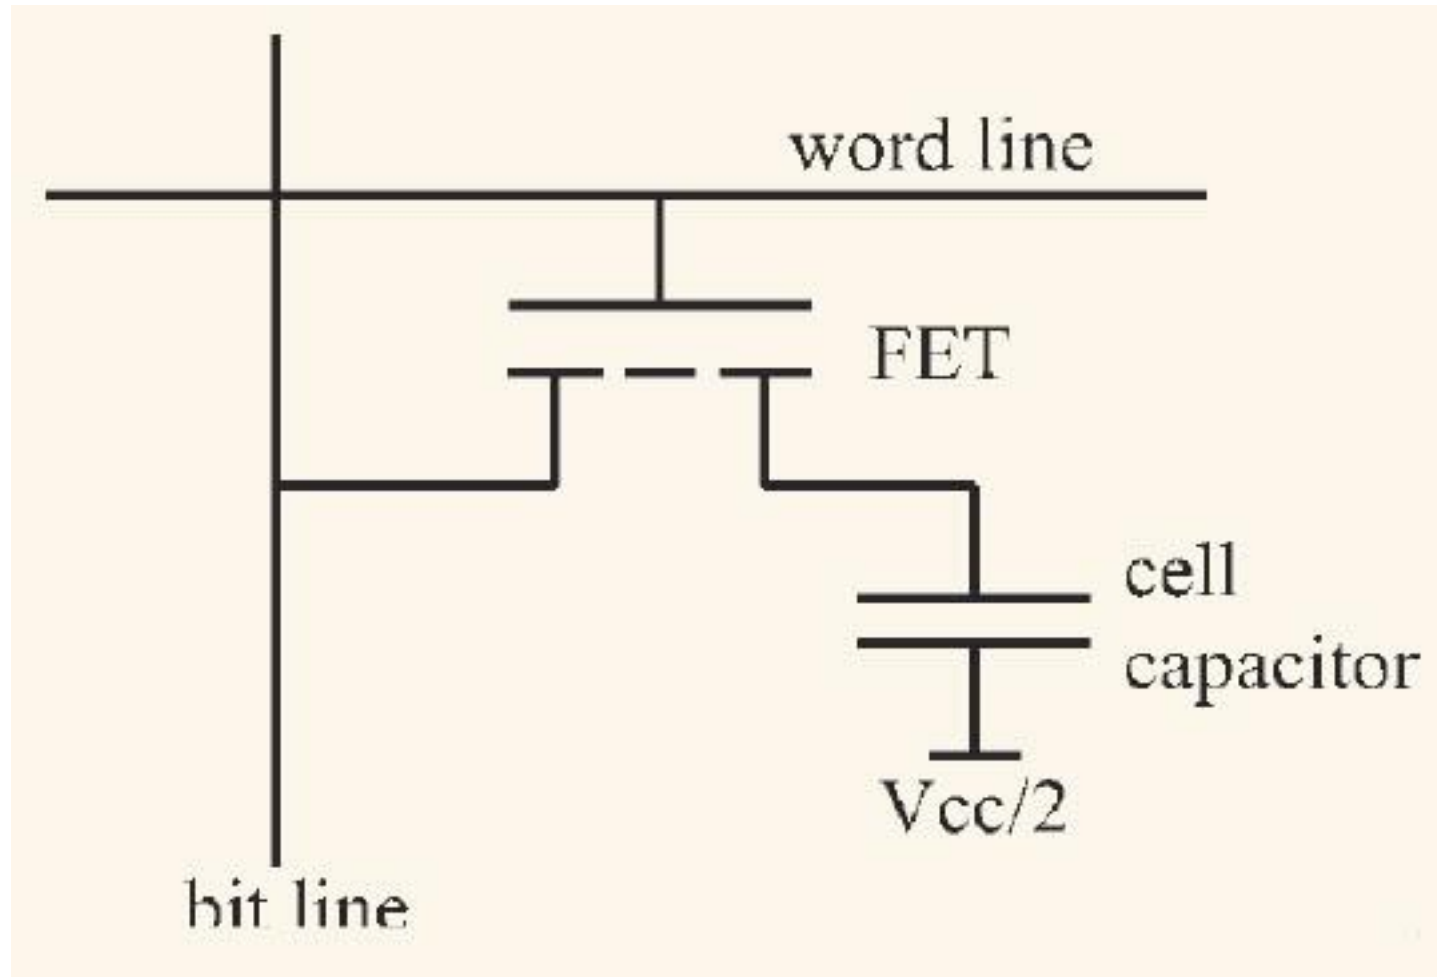
10^{-6} meters

DRAM Bits


1 micron



DRAM Cell in Transistors




Physical Layout of DRAM Bit



10^{-7} meters

Cross Section of DRAM Bits

100 nanometers



AMD Opteron Dependability

- L1 cache data is SEC/DED protected
- L2 cache and tags are SEC/DED protected
- DRAM is SEC/DED protected with chipkill
- On-chip and off-chip ECC protected arrays include autonomous, background hardware scrubbers
- Remaining arrays are parity protected
 - Instruction cache, tags and TLBs
 - Data tags and TLBs
 - Generally read only data that can be recovered from lower levels

Programming Memory Hierarchy: Cache Blocked Algorithm

- The blocked version of the i-j-k algorithm is written simply as (A,B,C are submatrices of a, b, c)

```
for (i=0;i<N/r;i++)  
  for (j=0;j<N/r;j++)  
    for (k=0;k<N/r;k++)  
      C[i][j] += A[i][k]*B[k][j]
```

- r = block (sub-matrix) size (Assume r divides N)
- $X[i][j]$ = a sub-matrix of X , defined by block row i and block column j

Great Ideas in Computer Architecture

1. *Design for Moore's Law*
 - *Higher capacities caches and DRAM*
2. *Abstraction to Simplify Design*
3. *Make the Common Case Fast*
4. *Dependability via Redundancy*
 - *Parity, SEC/DEC*
5. *Memory Hierarchy*
 - *Caches, TLBs*
6. *Performance via Parallelism/Pipelining/Prediction*
 - *Data-level Parallelism*

Course Summary

- As the field changes, Computer Architecture courses change, too!
- It is still about the software-hardware interface
 - Programming for performance!
 - Parallelism: Task-, Thread-, Instruction-, and Data-MapReduce, OpenMP, C, SSE Intrinsics
 - Understanding the memory hierarchy and its impact on application performance